# Oracle® Database

Advanced Replication Management API Reference

11*g* Release 2 (11.2)

**E10707-05**

June 2013

ORACLE®

Oracle Database Advanced Replication Management API Reference, 11*g* Release 2 (11.2)

E10707-05

Copyright © 1996, 2013, Oracle and/or its affiliates. All rights reserved.

Primary Author: Randy Urbano

Contributors: N. Arora, S. Balaraman, Y. Chan, A. Downing, C. Elsbernd, Y. Feng, J. Galagali, D. Goddard, L. Kaplan, V. Krishnamurthy, A. Lakshminath, P. Lane, J. Liu, E. Lu, P. McElroy, V. Moore, M. Pratt, A. Rajaram, N. Shodhan, W. Smith, J. Stamos, J. Stern, M. Subramaniam, L. Wong, D. Zhang

This software and related documentation are provided under a license agreement containing restrictions on use and disclosure and are protected by intellectual property laws. Except as expressly permitted in your license agreement or allowed by law, you may not use, copy, reproduce, translate, broadcast, modify, license, transmit, distribute, exhibit, perform, publish, or display any part, in any form, or by any means. Reverse engineering, disassembly, or decompilation of this software, unless required by law for interoperability, is prohibited.

The information contained herein is subject to change without notice and is not warranted to be error-free. If you find any errors, please report them to us in writing.

If this is software or related documentation that is delivered to the U.S. Government or anyone licensing it on behalf of the U.S. Government, the following notice is applicable:

U.S. GOVERNMENT END USERS: Oracle programs, including any operating system, integrated software, any programs installed on the hardware, and/or documentation, delivered to U.S. Government end users are "commercial computer software" pursuant to the applicable Federal Acquisition Regulation and agency-specific supplemental regulations. As such, use, duplication, disclosure, modification, and adaptation of the programs, including any operating system, integrated software, any programs installed on the hardware, and/or documentation, shall be subject to license terms and license restrictions applicable to the programs. No other rights are granted to the U.S. Government.

This software or hardware is developed for general use in a variety of information management applications. It is not developed or intended for use in any inherently dangerous applications, including applications that may create a risk of personal injury. If you use this software or hardware in dangerous applications, then you shall be responsible to take all appropriate fail-safe, backup, redundancy, and other measures to ensure its safe use. Oracle Corporation and its affiliates disclaim any liability for any damages caused by use of this software or hardware in dangerous applications.

Oracle and Java are registered trademarks of Oracle and/or its affiliates. Other names may be trademarks of their respective owners.

Intel and Intel Xeon are trademarks or registered trademarks of Intel Corporation. All SPARC trademarks are used under license and are trademarks or registered trademarks of SPARC International, Inc. AMD, Opteron, the AMD logo, and the AMD Opteron logo are trademarks or registered trademarks of Advanced Micro Devices. UNIX is a registered trademark of The Open Group.

This software or hardware and documentation may provide access to or information on content, products, and services from third parties. Oracle Corporation and its affiliates are not responsible for and expressly disclaim all warranties of any kind with respect to third-party content, products, and services. Oracle Corporation and its affiliates will not be responsible for any loss, costs, or damages incurred due to your access to or use of third-party content, products, or services.

# Contents

**Preface** ..... xv

Audience ..... xv
Documentation Accessibility ..... xvi
Related Documents ..... xvi
Conventions ..... xvi

## Part I Configuring Your Replication Environment

### 1 Overview of Advanced Replication

**Overview of Creating a Replication Environment** ..... 1-1
**Before You Start** ..... 1-2

### 2 Configuring the Replication Sites

**Overview of Setting Up Replication Sites** ..... 2-1
**Setting Up Master Sites** ..... 2-3
Setting Up orc1.example.com ..... 2-4
Setting Up orc2.example.com ..... 2-7
Setting Up orc3.example.com ..... 2-10
Creating Scheduled Links Between the Master Sites ..... 2-13
**Setting Up Materialized View Sites** ..... 2-16
Setting Up mv1.example.com ..... 2-16
Setting Up mv2.example.com ..... 2-21

### 3 Creating a Master Group

**Overview of Creating a Master Group** ..... 3-1
Before You Start ..... 3-2
**Creating a Master Group** ..... 3-4

### 4 Creating a Deployment Template

**Oracle Deployment Templates Concepts** ..... 4-1
**Before Creating the Deployment Template** ..... 4-2
**Creating a Deployment Template** ..... 4-2
**Packaging a Deployment Template for Instantiation** ..... 4-9
Packaging a Deployment Template ..... 4-10

Packaging a Deployment Template for Offline Instantiation ..... 4-11
Packaging a Deployment Template for Online Instantiation ..... 4-11
Saving an Instantiation Script to File ..... 4-12
**Distributing Instantiation Files** ..... 4-14
**Instantiating a Deployment Template** ..... 4-14
**Refreshing a Refresh Group After Instantiation** ..... 4-16

## 5 Creating a Materialized View Group

**Overview of Creating a Materialized View Group** ..... 5-1
**Creating a Materialized View Group** ..... 5-2
Creating the Materialized View Group at mv1.example.com ..... 5-3
Creating the Materialized View Group at mv2.example.com ..... 5-8

## 6 Configuring Conflict Resolution

**Preparing for Conflict Resolution** ..... 6-1
**Creating Conflict Resolution Methods for Update Conflicts** ..... 6-2
Overwrite and Discard Conflict Resolution Methods ..... 6-2
Minimum and Maximum Conflict Resolution Methods ..... 6-4
Timestamp Conflict Resolution Methods ..... 6-6
Additive and Average Conflict Resolution Methods ..... 6-9
Priority Groups Conflict Resolution Methods ..... 6-11
Site Priority Conflict Resolution Methods ..... 6-15
**Creating Conflict Resolution Methods for Uniqueness Conflicts** ..... 6-19
**Creating Conflict Avoidance Methods for Delete Conflicts** ..... 6-23
**Using Dynamic Ownership Conflict Avoidance** ..... 6-26
Workflow ..... 6-27
Token Passing ..... 6-27
Locating the Owner of a Row ..... 6-28
Obtaining Ownership ..... 6-29
Applying the Change ..... 6-30
**Auditing Successful Conflict Resolution** ..... 6-30
Collecting Conflict Resolution Statistics ..... 6-30
Viewing Conflict Resolution Statistics ..... 6-30
Canceling Conflict Resolution Statistics ..... 6-30
Clearing Statistics Information ..... 6-31

# Part II Managing and Monitoring Your Replication Environment

## 7 Managing a Master Replication Environment

**Changing the Master Definition Site** ..... 7-1
Option 1: All Master Sites Are Available ..... 7-1
Option 2: The Old Master Definition Site Is Not Available ..... 7-2
**Adding New Master Sites** ..... 7-2
Adding New Master Sites without Quiescing the Master Group ..... 7-3
Using Full Database Export/Import or Change-Based Recovery ..... 7-5
Using Object-Level Export/Import ..... 7-14

Adding New Master Sites to a Quiesced Master Group ..... 7-23
Adding New Master Sites Using the ADD_MASTER_DATABASE Procedure ..... 7-23
Adding New Master Sites with Offline Instantiation Using Export/Import ..... 7-25
**Removing a Master Site from a Master Group** ..... 7-31
Removing an Unavailable Master Site ..... 7-32
**Updating the Comments Fields in Data Dictionary Views** ..... 7-33
**Using Procedural Replication** ..... 7-34
Restrictions on Procedural Replication ..... 7-34
User-Defined Types and Procedural Replication ..... 7-36
Serializing Transactions ..... 7-37
Generating Support for Replicated Procedures ..... 7-37

## 8 Managing a Materialized View Replication Environment

**Refreshing Materialized Views** ..... 8-1
**Changing a Materialized View Group's Master Site** ..... 8-2
**Dropping Materialized View Groups and Objects** ..... 8-2
Dropping a Materialized View Group Created with a Deployment Template ..... 8-3
Using the Public Version of DROP_SITE_INSTANTIATION ..... 8-3
Using the Private Version of DROP_SITE_INSTANTIATION ..... 8-4
Dropping a Materialized View Group or Objects Created Manually ..... 8-6
Dropping a Materialized View Group Created Manually ..... 8-6
Dropping Objects at a Materialized View Site ..... 8-7
Cleaning Up a Master Site or Master Materialized View Site ..... 8-8
Cleaning Up After Dropping a Materialized View Group ..... 8-8
Cleaning Up Individual Materialized View Support ..... 8-10
**Managing Materialized View Logs** ..... 8-11
Altering Materialized View Logs ..... 8-12
Altering Materialized View Log Storage Parameters ..... 8-12
Altering a Materialized View Log to Add Columns ..... 8-12
Managing Materialized View Log Space ..... 8-12
Purging Rows from a Materialized View Log ..... 8-13
Truncating a Materialized View Log ..... 8-13
Reorganizing Master Tables that Have Materialized View Logs ..... 8-14
Reorganization Notification ..... 8-15
Truncating Masters ..... 8-15
Methods of Reorganizing a Database Table ..... 8-15
Dropping a Materialized View Log ..... 8-17
**Performing an Offline Instantiation of a Materialized View Site Using Export/Import** ..... 8-17
**Using a Group Owner for a Materialized View Group** ..... 8-27

## 9 Managing Replication Objects and Queues

**Altering a Replicated Object in a Quiesced Master Group** ..... 9-1
**Modifying Tables without Replicating the Modifications** ..... 9-4
Disabling Replication ..... 9-5
Reenabling Replication ..... 9-5
Ensuring that Replicated Triggers Fire Only Once ..... 9-6

**Converting a LONG Column to a LOB Column in a Replicated Table** ..... 9-6
**Determining Differences Between Replicated Tables** ..... 9-7
Using the DIFFERENCES Procedure ..... 9-8
Using the RECTIFY Procedure ..... 9-8
**Managing the Deferred Transactions Queue** ..... 9-11
Pushing the Deferred Transaction Queue ..... 9-11
Purging the Deferred Transaction Queue ..... 9-12
Using the ANYDATA Type to Determine the Value of an Argument in a Deferred Call ..... 9-13
**Managing the Error Queue** ..... 9-15
Reexecuting Error Transaction as the Receiver ..... 9-15
Reexecuting Error Transaction as Alternate User ..... 9-15

## 10 Monitoring a Replication Environment

**Monitoring Master Replication Environments** ..... 10-1
Monitoring Master Sites ..... 10-1
Listing General Information About a Master Site ..... 10-2
Monitoring Master Groups ..... 10-2
Listing the Master Sites Participating in a Master Group ..... 10-3
Listing General Information About Master Groups ..... 10-3
Monitoring Masters ..... 10-4
Listing Information About Materialized Views Based on a Master ..... 10-4
Listing Information About the Materialized View Logs at a Master ..... 10-5
Listing the Materialized Views that Use a Materialized View Log ..... 10-6
Listing Information About the Deployment Templates at a Master ..... 10-7
**Monitoring Materialized View Sites** ..... 10-8
Listing General Information About a Materialized View Site ..... 10-8
Listing General Information About Materialized View Groups ..... 10-8
Listing Information About Materialized Views ..... 10-9
Listing Master Information For Materialized Views ..... 10-9
Listing the Properties of Materialized Views ..... 10-10
Listing Information About the Refresh Groups at a Materialized View Site ..... 10-10
Determining the Job ID for Each Refresh Job at a Materialized View Site ..... 10-11
Determining Which Materialized Views Are Currently Refreshing ..... 10-12
**Monitoring Administrative Requests** ..... 10-12
Listing General Information About Administrative Requests ..... 10-12
Determining the Cause of Administrative Request Errors ..... 10-13
Listing General Information About the Job that Executes Administrative Requests ..... 10-14
Checking the Definition of Each do_deferred_repcat_admin Job ..... 10-14
**Monitoring the Deferred Transactions Queue** ..... 10-15
Monitoring Transaction Propagation ..... 10-15
Listing the Number of Deferred Transactions for Each Destination Master Site ..... 10-15
Listing General Information About the Push Jobs at a Replication Site ..... 10-15
Determining the Next Start Time and Interval for the Push Jobs ..... 10-16
Determining the Total Number of Transactions Queued for Propagation ..... 10-17
Monitoring Purges of Successfully Propagated Transactions ..... 10-17
Listing General Information About the Purge Job ..... 10-17
Checking the Definition of the Purge Job ..... 10-18

Determining the Amount of Time Since the Last Purge ..... 10-18
Determining the Total Number of Purged Transactions ..... 10-18
**Monitoring the Error Queue** ..... 10-18
Listing General Information About the Error Transactions at a Replication Site ..... 10-19
Determining the Percentage of Error Transactions ..... 10-19
Listing the Number of Error Transactions from Each Origin Master Site ..... 10-20
Listing the Error Messages for the Error Transactions at a Replication Site ..... 10-20
Determining the Error Operations at a Replication Site ..... 10-20
**Monitoring Performance in a Replication Environment** ..... 10-21
Tracking the Average Number of Row Changes in a Replication Transaction ..... 10-21
Tracking the Rate of Transactions Entering the Deferred Transactions Queue ..... 10-22
Determining the Average Network Traffic Created to Propagate a Transaction ..... 10-22
Determining the Average Amount of Time to Apply Transactions at Remote Sites ..... 10-23
Determining the Percentage of Time the Parallel Propagation Job Spends Sleeping ..... 10-23
Clearing the Statistics for a Remote Master Site in the DEFSCHEDULE View ..... 10-24
Monitoring Parallel Propagation of Deferred Transactions Using V$REPLPROP ..... 10-24
Determining the Databases to Which You Are Propagating Deferred Transactions .. 10-24
Determining the Transactions Currently Being Propagated to a Remote Master ..... 10-24

## Part III Replication Management API Packages Reference

## 11 Introduction to the Replication Management API Reference

**Examples of Using Oracle's Replication Management API** ..... 11-1
**Issues to Consider When Using the Replication Management API** ..... 11-2
**The Advanced Replication Interface and the Replication Management API** ..... 11-2
**Abbreviations for Datetime and Interval Data Types** ..... 11-2

## 12 DBMS_DEFER

**Summary of DBMS_DEFER Subprograms** ..... 12-2
CALL Procedure ..... 12-3
COMMIT_WORK Procedure ..... 12-4
*datatype*_ARG Procedure ..... 12-5
TRANSACTION Procedure ..... 12-7

## 13 DBMS_DEFER_QUERY

**Summary of DBMS_DEFER_QUERY Subprograms** ..... 13-2
GET_ARG_FORM Function ..... 13-3
GET_ARG_TYPE Function ..... 13-4
GET_CALL_ARGS Procedure ..... 13-6
GET_*datatype*_ARG Function ..... 13-8
GET_OBJECT_NULL_VECTOR_ARG Function ..... 13-10

## 14 DBMS_DEFER_SYS

**Summary of DBMS_DEFER_SYS Subprograms** ..... 14-2
ADD_DEFAULT_DEST Procedure ..... 14-4

CLEAR_PROP_STATISTICS Procedure ..... 14-5
DELETE_DEFAULT_DEST Procedure ..... 14-6
DELETE_DEF_DESTINATION Procedure ..... 14-7
DELETE_ERROR Procedure ..... 14-8
DELETE_TRAN Procedure ..... 14-9
DISABLED Function ..... 14-10
EXCLUDE_PUSH Function ..... 14-11
EXECUTE_ERROR Procedure ..... 14-12
EXECUTE_ERROR_AS_USER Procedure ..... 14-13
PURGE Function ..... 14-14
PUSH Function ..... 14-16
REGISTER_PROPAGATOR Procedure ..... 14-19
SCHEDULE_PURGE Procedure ..... 14-20
SCHEDULE_PUSH Procedure ..... 14-22
SET_DISABLED Procedure ..... 14-24
UNREGISTER_PROPAGATOR Procedure ..... 14-26
UNSCHEDULE_PURGE Procedure ..... 14-27
UNSCHEDULE_PUSH Procedure ..... 14-28

## 15 DBMS_OFFLINE_OG

**Summary of DBMS_OFFLINE_OG Subprograms** ..... 15-2
BEGIN_INSTANTIATION Procedure ..... 15-3
BEGIN_LOAD Procedure ..... 15-5
END_INSTANTIATION Procedure ..... 15-6
END_LOAD Procedure ..... 15-7
RESUME_SUBSET_OF_MASTERS Procedure ..... 15-9

## 16 DBMS_RECTIFIER_DIFF

**Summary of DBMS_RECTIFIER_DIFF Subprograms** ..... 16-2
DIFFERENCES Procedure ..... 16-3
RECTIFY Procedure ..... 16-6

## 17 DBMS_REFRESH

**Summary of DBMS_REFRESH Subprograms** ..... 17-2
ADD Procedure ..... 17-3
CHANGE Procedure ..... 17-4
DESTROY Procedure ..... 17-6
MAKE Procedure ..... 17-7
REFRESH Procedure ..... 17-9
SUBTRACT Procedure ..... 17-10

## 18 DBMS_REPCAT

**Summary of DBMS_REPCAT Subprograms** ..... 18-2
ADD_GROUPED_COLUMN Procedure ..... 18-6
ADD_MASTER_DATABASE Procedure ..... 18-7
ADD_NEW_MASTERS Procedure ..... 18-8

ADD_PRIORITY_*datatype* Procedure ..... 18-13
ADD_SITE_PRIORITY_SITE Procedure ..... 18-15
ADD_*conflicttype*_RESOLUTION Procedure ..... 18-16
ALTER_CATCHUP_PARAMETERS Procedure ..... 18-20
ALTER_MASTER_PROPAGATION Procedure ..... 18-22
ALTER_MASTER_REPOBJECT Procedure ..... 18-23
ALTER_MVIEW_PROPAGATION Procedure ..... 18-25
ALTER_PRIORITY Procedure ..... 18-26
ALTER_PRIORITY_*datatype* Procedure ..... 18-27
ALTER_SITE_PRIORITY Procedure ..... 18-28
ALTER_SITE_PRIORITY_SITE Procedure ..... 18-29
CANCEL_STATISTICS Procedure ..... 18-30
COMMENT_ON_COLUMN_GROUP Procedure ..... 18-31
COMMENT_ON_MVIEW_REPSITES Procedure ..... 18-32
COMMENT_ON_PRIORITY_GROUP Procedures ..... 18-33
COMMENT_ON_REPGROUP Procedure ..... 18-34
COMMENT_ON_REPOBJECT Procedure ..... 18-35
COMMENT_ON_REPSITES Procedure ..... 18-36
COMMENT_ON_SITE_PRIORITY Procedure ..... 18-37
COMMENT_ON_*conflicttype*_RESOLUTION Procedure ..... 18-38
COMPARE_OLD_VALUES Procedure ..... 18-40
CREATE_MASTER_REPGROUP Procedure ..... 18-42
CREATE_MASTER_REPOBJECT Procedure ..... 18-43
CREATE_MVIEW_REPGROUP Procedure ..... 18-46
CREATE_MVIEW_REPOBJECT Procedure ..... 18-48
DEFINE_COLUMN_GROUP Procedure ..... 18-51
DEFINE_PRIORITY_GROUP Procedure ..... 18-52
DEFINE_SITE_PRIORITY Procedure ..... 18-53
DO_DEFERRED_REPCAT_ADMIN Procedure ..... 18-54
DROP_COLUMN_GROUP Procedure ..... 18-55
DROP_GROUPED_COLUMN Procedure ..... 18-56
DROP_MASTER_REPGROUP Procedure ..... 18-57
DROP_MASTER_REPOBJECT Procedure ..... 18-58
DROP_MVIEW_REPGROUP Procedure ..... 18-59
DROP_MVIEW_REPOBJECT Procedure ..... 18-60
DROP_PRIORITY Procedure ..... 18-61
DROP_PRIORITY_GROUP Procedure ..... 18-62
DROP_PRIORITY_*datatype* Procedure ..... 18-63
DROP_SITE_PRIORITY Procedure ..... 18-64
DROP_SITE_PRIORITY_SITE Procedure ..... 18-65
DROP_*conflicttype*_RESOLUTION Procedure ..... 18-66
EXECUTE_DDL Procedure ..... 18-68
GENERATE_MVIEW_SUPPORT Procedure ..... 18-69
GENERATE_REPLICATION_SUPPORT Procedure ..... 18-71
MAKE_COLUMN_GROUP Procedure ..... 18-73
PREPARE_INSTANTIATED_MASTER Procedure ..... 18-74
PURGE_MASTER_LOG Procedure ..... 18-76

PURGE_STATISTICS Procedure ..... 18-77
REFRESH_MVIEW_REPGROUP Procedure ..... 18-78
REGISTER_MVIEW_REPGROUP Procedure ..... 18-80
REGISTER_STATISTICS Procedure ..... 18-81
RELOCATE_MASTERDEF Procedure ..... 18-82
REMOVE_MASTER_DATABASES Procedure ..... 18-84
RENAME_SHADOW_COLUMN_GROUP Procedure ..... 18-85
REPCAT_IMPORT_CHECK Procedure ..... 18-86
RESUME_MASTER_ACTIVITY Procedure ..... 18-87
RESUME_PROPAGATION_TO_MDEF Procedure ..... 18-88
SEND_OLD_VALUES Procedure ..... 18-89
SET_COLUMNS Procedure ..... 18-91
SPECIFY_NEW_MASTERS Procedure ..... 18-93
STREAMS_MIGRATION Procedure ..... 18-95
SUSPEND_MASTER_ACTIVITY Procedure ..... 18-96
SWITCH_MVIEW_MASTER Procedure ..... 18-97
UNDO_ADD_NEW_MASTERS_REQUEST Procedure ..... 18-98
UNREGISTER_MVIEW_REPGROUP Procedure ..... 18-100
VALIDATE Function ..... 18-101
WAIT_MASTER_LOG Procedure ..... 18-103

## 19 DBMS_REPCAT_INSTANTIATE

**Summary of DBMS_REPCAT_INSTANTIATE Subprograms** ..... 19-2
DROP_SITE_INSTANTIATION Procedure ..... 19-3
INSTANTIATE_OFFLINE Function ..... 19-4
INSTANTIATE_ONLINE Function ..... 19-6

## 20 DBMS_REPCAT_ADMIN

**Summary of DBMS_REPCAT_ADMIN Subprograms** ..... 20-2
GRANT_ADMIN_ANY_SCHEMA Procedure ..... 20-3
GRANT_ADMIN_SCHEMA Procedure ..... 20-4
REGISTER_USER_REPGROUP Procedure ..... 20-5
REVOKE_ADMIN_ANY_SCHEMA Procedure ..... 20-7
REVOKE_ADMIN_SCHEMA Procedure ..... 20-8
UNREGISTER_USER_REPGROUP Procedure ..... 20-9

## 21 DBMS_REPCAT_RGT

**Summary of DBMS_REPCAT_RGT Subprograms** ..... 21-2
ALTER_REFRESH_TEMPLATE Procedure ..... 21-4
ALTER_TEMPLATE_OBJECT Procedure ..... 21-6
ALTER_TEMPLATE_PARM Procedure ..... 21-8
ALTER_USER_AUTHORIZATION Procedure ..... 21-10
ALTER_USER_PARM_VALUE Procedure ..... 21-11
COMPARE_TEMPLATES Function ..... 21-13
COPY_TEMPLATE Function ..... 21-14
CREATE_OBJECT_FROM_EXISTING Function ..... 21-16

CREATE_REFRESH_TEMPLATE Function ... 21-18
CREATE_TEMPLATE_OBJECT Function ... 21-20
CREATE_TEMPLATE_PARM Function ... 21-22
CREATE_USER_AUTHORIZATION Function ... 21-24
CREATE_USER_PARM_VALUE Function ... 21-25
DELETE_RUNTIME_PARMS Procedure ... 21-27
DROP_ALL_OBJECTS Procedure ... 21-28
DROP_ALL_TEMPLATE_PARMS Procedure ... 21-29
DROP_ALL_TEMPLATE_SITES Procedure ... 21-30
DROP_ALL_TEMPLATES Procedure ... 21-31
DROP_ALL_USER_AUTHORIZATIONS Procedure ... 21-32
DROP_ALL_USER_PARM_VALUES Procedure ... 21-33
DROP_REFRESH_TEMPLATE Procedure ... 21-34
DROP_SITE_INSTANTIATION Procedure ... 21-35
DROP_TEMPLATE_OBJECT Procedure ... 21-36
DROP_TEMPLATE_PARM Procedure ... 21-37
DROP_USER_AUTHORIZATION Procedure ... 21-38
DROP_USER_PARM_VALUE Procedure ... 21-39
GET_RUNTIME_PARM_ID Function ... 21-40
INSERT_RUNTIME_PARMS Procedure ... 21-41
INSTANTIATE_OFFLINE Function ... 21-43
INSTANTIATE_ONLINE Function ... 21-45
LOCK_TEMPLATE_EXCLUSIVE Procedure ... 21-47
LOCK_TEMPLATE_SHARED Procedure ... 21-48

## 22 DBMS_REPUTIL

**Summary of DBMS_REPUTIL Subprograms** ... 22-2
REPLICATION_OFF Procedure ... 22-3
REPLICATION_ON Procedure ... 22-4
REPLICATION_IS_ON Function ... 22-5
FROM_REMOTE Function ... 22-6
GLOBAL_NAME Function ... 22-7
MAKE_INTERNAL_PKG Procedure ... 22-8
SYNC_UP_REP Procedure ... 22-9

## Part IV Replication Data Dictionary Reference

## 23 Replication Catalog Views

**Summary of Replication Catalog Views** ... 23-2
DBA_REGISTERED_MVIEW_GROUPS ... 23-5
ALL_REPCAT_REFRESH_TEMPLATES ... 23-6
ALL_REPCAT_TEMPLATE_OBJECTS ... 23-7
ALL_REPCAT_TEMPLATE_PARMS ... 23-9
ALL_REPCAT_TEMPLATE_SITES ... 23-11
ALL_REPCAT_USER_AUTHORIZATIONS ... 23-12
ALL_REPCAT_USER_PARM_VALUES ... 23-13

ALL_REPCATLOG ..... 23-15
ALL_REPCOLUMN ..... 23-16
ALL_REPCOLUMN_GROUP ..... 23-18
ALL_REPCONFLICT ..... 23-19
ALL_REPDDL ..... 23-20
ALL_REPGENOBJECTS ..... 23-21
ALL_REPGROUP ..... 23-22
ALL_REPGROUP_PRIVILEGES ..... 23-23
ALL_REPGROUPED_COLUMN ..... 23-24
ALL_REPKEY_COLUMNS ..... 23-25
ALL_REPOBJECT ..... 23-26
ALL_REPPARAMETER_COLUMN ..... 23-28
ALL_REPPRIORITY ..... 23-29
ALL_REPPRIORITY_GROUP ..... 23-30
ALL_REPPROP ..... 23-31
ALL_REPRESOL_STATS_CONTROL ..... 23-32
ALL_REPRESOLUTION ..... 23-33
ALL_REPRESOLUTION_METHOD ..... 23-34
ALL_REPRESOLUTION_STATISTICS ..... 23-35
ALL_REPSITES ..... 23-36
DBA_REPCAT_REFRESH_TEMPLATES ..... 23-37
DBA_REPCAT_TEMPLATE_OBJECTS ..... 23-38
DBA_REPCAT_TEMPLATE_PARMS ..... 23-39
DBA_REPCAT_TEMPLATE_SITES ..... 23-40
DBA_REPCAT_USER_AUTHORIZATIONS ..... 23-41
DBA_REPCAT_USER_PARM_VALUES ..... 23-42
DBA_REPCATLOG ..... 23-43
DBA_REPCOLUMN ..... 23-44
DBA_REPCOLUMN_GROUP ..... 23-45
DBA_REPCONFLICT ..... 23-46
DBA_REPDDL ..... 23-47
DBA_REPEXTENSIONS ..... 23-48
DBA_REPGENOBJECTS ..... 23-50
DBA_REPGROUP ..... 23-51
DBA_REPGROUP_PRIVILEGES ..... 23-52
DBA_REPGROUPED_COLUMN ..... 23-53
DBA_REPKEY_COLUMNS ..... 23-54
DBA_REPOBJECT ..... 23-55
DBA_REPPARAMETER_COLUMN ..... 23-56
DBA_REPPRIORITY ..... 23-57
DBA_REPPRIORITY_GROUP ..... 23-58
DBA_REPPROP ..... 23-59
DBA_REPRESOL_STATS_CONTROL ..... 23-60
DBA_REPRESOLUTION ..... 23-61
DBA_REPRESOLUTION_METHOD ..... 23-62
DBA_REPRESOLUTION_STATISTICS ..... 23-63
DBA_REPSITES ..... 23-64

DBA_REPSITES_NEW ..... 23-65
USER_REPCAT_REFRESH_TEMPLATES ..... 23-66
USER_REPCAT_TEMPLATE_OBJECTS ..... 23-67
USER_REPCAT_TEMPLATE_PARMS ..... 23-68
USER_REPCAT_TEMPLATE_SITES ..... 23-69
USER_REPCAT_USER_AUTHORIZATION ..... 23-70
USER_REPCAT_USER_PARM_VALUES ..... 23-71
USER_REPCATLOG ..... 23-72
USER_REPCOLUMN ..... 23-73
USER_REPCOLUMN_GROUP ..... 23-74
USER_REPCONFLICT ..... 23-75
USER_REPDDL ..... 23-76
USER_REPGENOBJECTS ..... 23-77
USER_REPGROUP ..... 23-78
USER_REPGROUP_PRIVILEGES ..... 23-79
USER_REPGROUPED_COLUMN ..... 23-80
USER_REPKEY_COLUMNS ..... 23-81
USER_REPOBJECT ..... 23-82
USER_REPPARAMETER_COLUMN ..... 23-83
USER_REPPRIORITY ..... 23-84
USER_REPPRIORITY_GROUP ..... 23-85
USER_REPPROP ..... 23-86
USER_REPRESOL_STATS_CONTROL ..... 23-87
USER_REPRESOLUTION ..... 23-88
USER_REPRESOLUTION_METHOD ..... 23-89
USER_REPRESOLUTION_STATISTICS ..... 23-90
USER_REPSITES ..... 23-91

## 24 Replication Dynamic Performance Views

**V$MVREFRESH** ..... 24-2
**V$REPLPROP** ..... 24-3
**V$REPLQUEUE** ..... 24-5

## 25 Materialized View and Refresh Group Views

## 26 Deferred Transaction Views

**DEFCALL** ..... 26-2
**DEFCALLDEST** ..... 26-3
**DEFDEFAULTDEST** ..... 26-4
**DEFERRCOUNT** ..... 26-5
**DEFERROR** ..... 26-6
**DEFLOB** ..... 26-7
**DEFPROPAGATOR** ..... 26-8
**DEFSCHEDULE** ..... 26-9
**DEFTRAN** ..... 26-12
**DEFTRANDEST** ..... 26-13

## Part V Appendixes

## A Security Options

**Security Setup for Multimaster Replication** ..... A-1
Trusted Compared with Untrusted Security ..... A-2
**Security Setup for Materialized View Replication** ..... A-5
Trusted Compared with Untrusted Security ..... A-6

## B User-Defined Conflict Resolution Methods

**User-Defined Conflict Resolution Methods** ..... B-1
Conflict Resolution Method Parameters ..... B-1
Resolving Update Conflicts ..... B-2
Resolving Uniqueness Conflicts ..... B-2
Resolving Delete Conflicts ..... B-2
Multitier Materialized Views and User-Defined Conflict Resolution Methods ..... B-3
Restrictions for User-Defined Conflict Resolution Methods ..... B-3
SQL Statement Restrictions for User-Defined Conflict Resolution Methods ..... B-3
Column Subsetting Restrictions for User-Defined Conflict Resolution Methods ..... B-3
Examples of User-Defined Conflict Resolution Method ..... B-4
Maximum User Function ..... B-4
Additive User Function ..... B-5
**User-Defined Conflict Notification Methods** ..... B-5
Creating a Conflict Notification Log ..... B-6
Sample Conflict Notification Log Table ..... B-6
Creating a Conflict Notification Package ..... B-6
Sample Conflict Notification Package ..... B-6
**Viewing Conflict Resolution Information** ..... B-8

## Index

# Preface

*Oracle Database Advanced Replication Management API Reference* contains information that describes the features and functionality of the replication management API. Specifically, the *Oracle Database Advanced Replication Management API Reference* contains reference information for the packages in the replication management API, as well as examples of their use.

In addition, Oracle Database Advanced Replication Management API Reference contains reference information about the replication catalog and other data dictionary views that are important for replication.

This Preface contains these topics:


- Audience
- Documentation Accessibility
- Related Documents
- Conventions


## Audience

*Oracle Database Advanced Replication Management API Reference* is intended for database administrators and application developers who develop and maintain replication environments. These administrators and application developers perform one or more of the following tasks:

- Configure replication sites
- Create master groups
- Create deployment templates
- Create materialized view groups
- Configure conflict resolution
- Manage replication environments
- Use the replication management API
- Monitor replication environments using data dictionary views
- Plan and configure security options

To use this document, you must be familiar with relational database concepts, distributed database administration, PL/SQL (if using procedural replication), and the operating system under which you run an Advanced Replication environment.

## Documentation Accessibility

For information about Oracle's commitment to accessibility, visit the Oracle Accessibility Program website at http://www.oracle.com/pls/topic/lookup?ctx=acc&id=docacc.

### Access to Oracle Support

Oracle customers have access to electronic support through My Oracle Support. For information, visit http://www.oracle.com/pls/topic/lookup?ctx=acc&id=info or visit http://www.oracle.com/pls/topic/lookup?ctx=acc&id=trs if you are hearing impaired.

## Related Documents

For more information, see these Oracle resources:

- *Oracle Database Advanced Replication*
- *Oracle Database Concepts*
- *Oracle Database Administrator's Guide*
- *Oracle Database SQL Language Reference*
- *Oracle Database PL/SQL Language Reference*
- *Oracle Streams Replication Administrator's Guide* for information about migrating your Advanced Replication environment to Oracle Streams

Many of the examples in this book use the sample schemas of the sample database, which is installed by default when you install Oracle Database. Refer to *Oracle Database Sample Schemas* for information about how these schemas were created and how you can use them yourself.

## Conventions

The following text conventions are used in this document:

| Convention | Meaning |
|---|---|
| **boldface** | Boldface type indicates graphical user interface elements associated with an action, or terms defined in text or the glossary. |
| *italic* | Italic type indicates book titles, emphasis, or placeholder variables for which you supply particular values. |
| `monospace` | Monospace type indicates commands within a paragraph, URLs, code in examples, text that appears on the screen, or text that you enter. |

# Part I

# Configuring Your Replication Environment

Part I contains instructions for using the replication management API to set up both multimaster replication and materialized view replication. This part also contains instructions for configuring conflict resolution methods using the replication management API.

Part I contains the following chapters:


- Chapter 1, "Overview of Advanced Replication"
- Chapter 2, "Configuring the Replication Sites"
- Chapter 3, "Creating a Master Group"
- Chapter 4, "Creating a Deployment Template"
- Chapter 5, "Creating a Materialized View Group"
- Chapter 6, "Configuring Conflict Resolution"

# 1

# Overview of Advanced Replication

This chapter reviews the process of building a replication environment with the replication management API.

This chapter contains these topics:


- Overview of Creating a Replication Environment
- Before You Start


## Overview of Creating a Replication Environment

Figure 1–1 illustrates the basic steps required to build a replication environment. Regardless of the type of replication site or sites that you are building, you begin by setting up the replicated site.

After you have set up your replication sites, you are ready to begin building your master groups and materialized view groups. After you have built your replication environment, ensure that you review Chapter 6 and the chapters in Part II, "Managing and Monitoring Your Replication Environment", to learn about conflict resolution and managing your replication environment.

***Figure 1–1 Create Replication Environment Process***

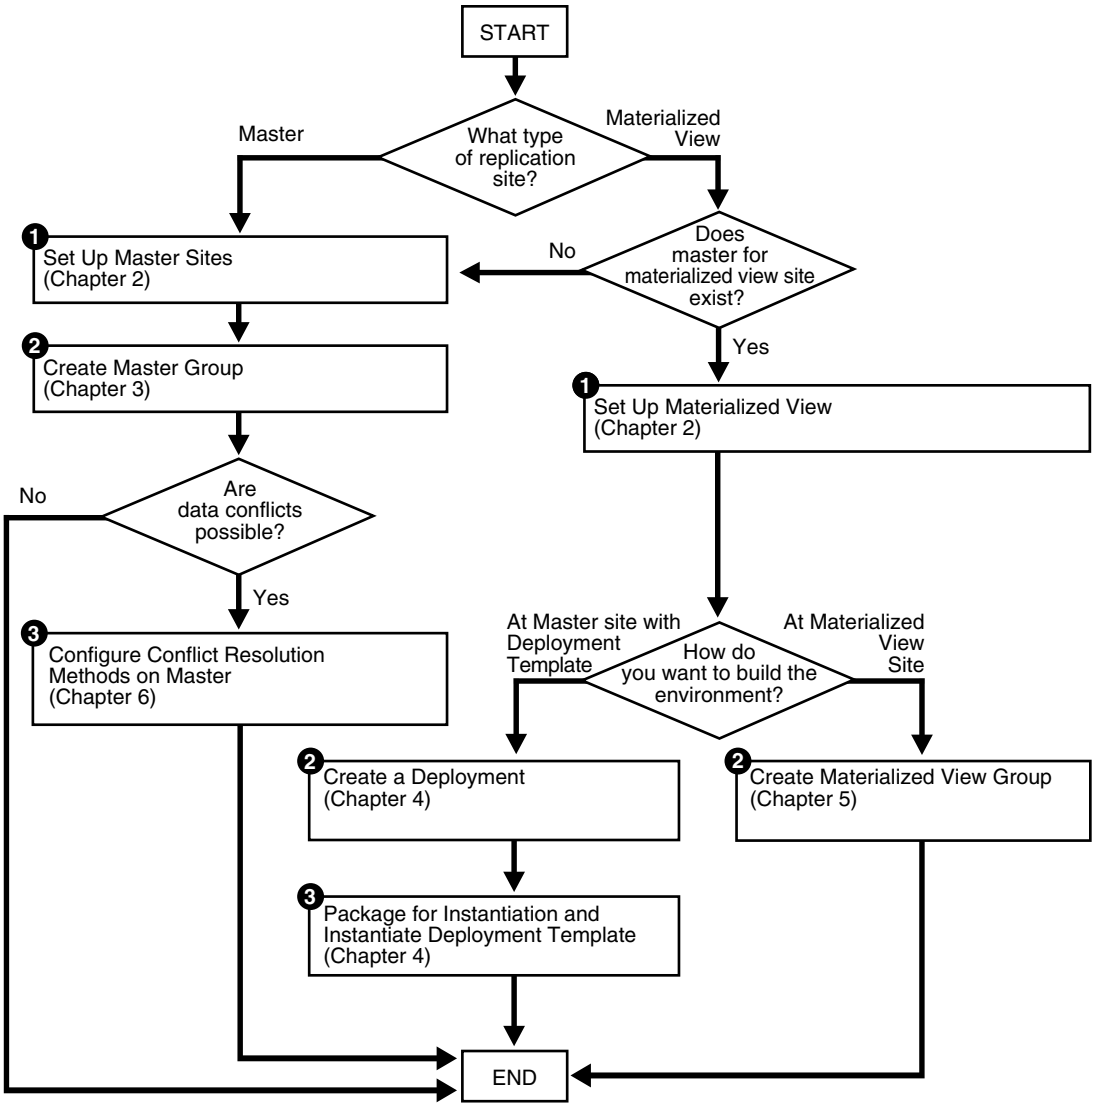


# Before You Start

Before you begin setting up your replication site, ensure that you plan your replication environment so that it meets your needs. Planning considerations include:

- Designing your replicated database objects
- Deciding on the settings of initialization parameters that are important for replication
- Deciding whether you want to create a multimaster replication environment or a materialized view replication environment, or if you want to combine both types of replication environments into a hybrid environment
- Deciding how you want to configure your scheduled links
- Deciding how you want to configure your scheduled purges
- Deciding whether you want to use serial or parallel propagation
- If you use parallel propagation, then deciding on the degree of parallelism

- If you plan to create a materialized view environment, then deciding whether you want to use deployment templates to create the environment
- Analyzing your environment for possible conflicts and, if conflicts are possible, then deciding which conflict resolution methods to use
- Configuring security for your replication environment
- Designing your replication environment for survivability

> **See Also:** *Oracle Database Advanced Replication* for more information planning your replication environment

# 2

# Configuring the Replication Sites

This chapter illustrates how to set up both a master site and a materialized view replication site using the replication management API.

This chapter contains these topics:


- Overview of Setting Up Replication Sites
- Setting Up Master Sites
- Setting Up Materialized View Sites


## Overview of Setting Up Replication Sites

Before you build your replication environment, you must set up the sites that will participate in the replication environment. As illustrated in Figure 2–2 and Figure 2–3, there are separate processes for setting up a master site versus setting up a materialized view site.

The examples in this book use the following databases:

- `orc1.example.com`
- `orc2.example.com`
- `orc3.example.com`
- `orc4.example.com`
- `orc5.example.com`
- `mv1.example.com`
- `mv2.example.com`

Chapters 2 - 6 work with the replication environment illustrated in Figure 2–1. You start to create this environment using the instructions in this chapter. Notice that `mv2.example.com` is a materialized view based on the `mv1.example.com` materialized view, creating a multitier materialized view environment. The arrows in Figure 2–1 represent database links.

*Figure 2–1 Three Master Sites and Two Materialized View Sites*

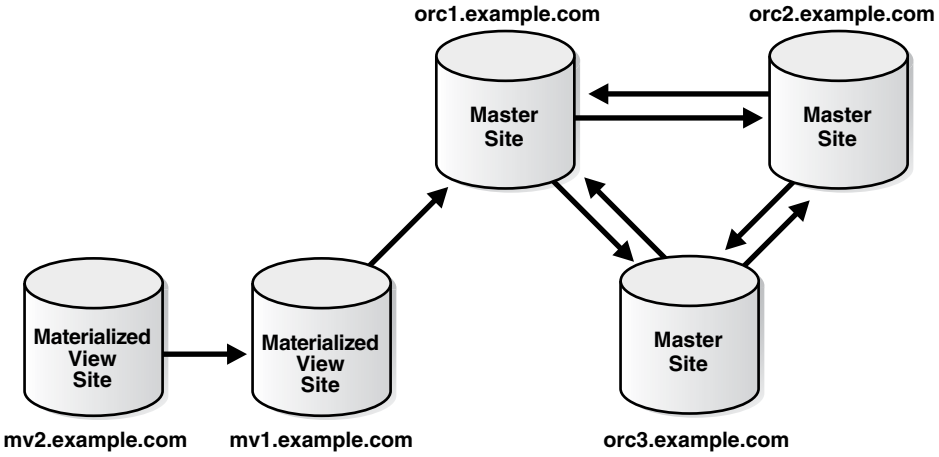


Follow the procedures identified in Figure 2–2 when you build a new master site or in Figure 2–3 when you build a new materialized view site.

***Figure 2–2 Setting Up Master Sites***

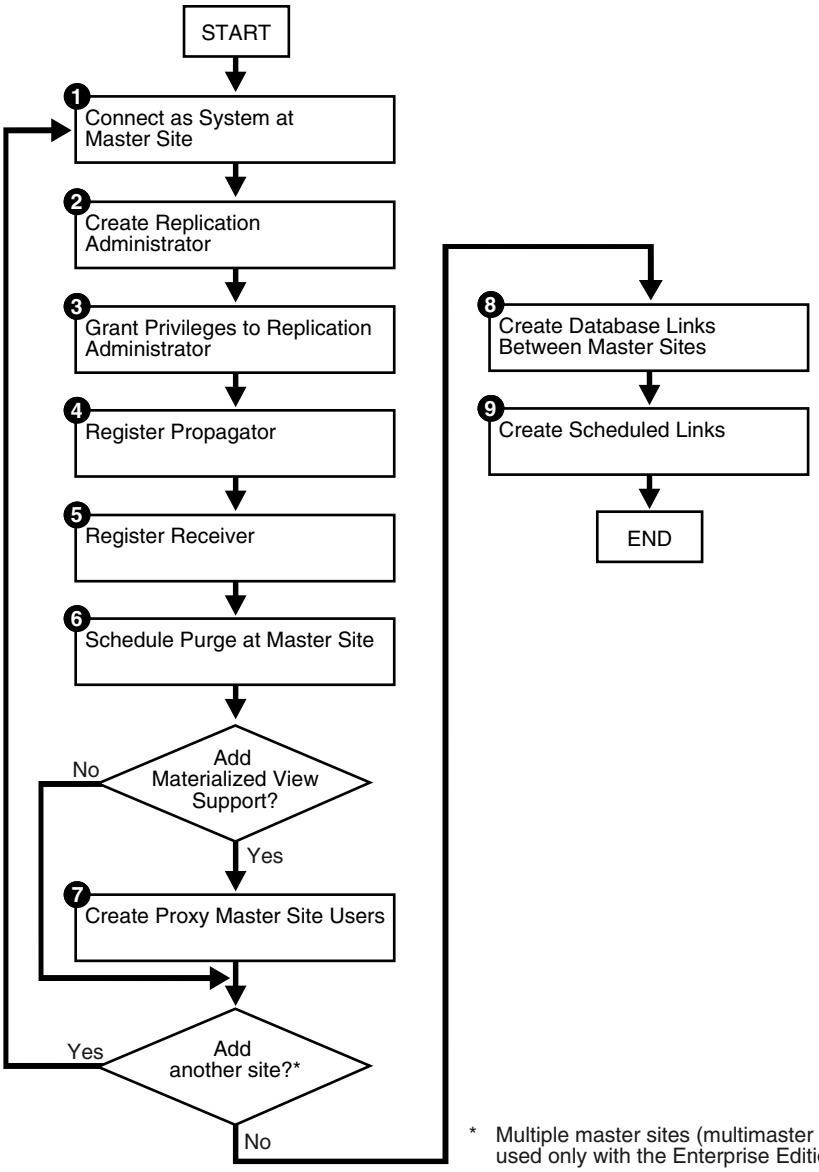


# Setting Up Master Sites

The following sections contain step-by-step instructions for setting up the three master sites in our sample replication environment: `orc1.example.com`, `orc2.example.com`, and `orc3.example.com`. Before you set up the master sites, configure your network and Oracle Net so that all three databases can communicate with each other.

---

**Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

---

```
/************************ BEGINNING OF SCRIPT ******************************
```

## Setting Up orc1.example.com

Complete the following steps to set up the `orc1.example.com` master site.

### Step 1 Connect as SYSTEM at a master site at orc1.example.com.

Connect as `SYSTEM` to the database that you want to set up for replication. After you set up `orc1.example.com`, begin again with Step 1 for site `orc2.example.com` on page 2-7 and Step 1 for site `orc3.example.com` on page 2-10.

```
*/

SET ECHO ON

SPOOL setup_masters.out

CONNECT system@orc1.example.com

/*
```

### Step 2 Create the replication administrator at orc1.example.com.

The replication administrator must be granted the necessary privileges to create and manage a replication environment. The replication administrator must be created at each database that participates in the replication environment.

```
*/

ACCEPT password PROMPT 'Enter password for user: ' HIDE

CREATE USER repadmin IDENTIFIED BY &password;

/*
```

---

**Note:** Enter an appropriate password for the administrative user.

---

**See Also:** *Oracle Database Security Guide* for guidelines for choosing passwords

### Step 3 Grant privileges to the replication administrator at orc1.example.com.

Execute the `GRANT_ADMIN_ANY_SCHEMA` procedure to grant the replication administrator powerful privileges to create and manage a replicated environment.

```
*/

BEGIN
   DBMS_REPCAT_ADMIN.GRANT_ADMIN_ANY_SCHEMA (
      username => 'repadmin');
END;
/

/*
```

If you want your `repadmin` to be able to create materialized view logs for any replicated table, then grant `COMMENT ANY TABLE` and `LOCK ANY TABLE` to `repadmin`:

```
*/

GRANT COMMENT ANY TABLE TO repadmin;
GRANT LOCK ANY TABLE TO repadmin;
```

```
/*
```

If you want your repadmin to be able to connect to the Advanced Replication interface in Oracle Enterprise Manager, then grant SELECT ANY DICTIONARY to repadmin:

```
*/

GRANT SELECT ANY DICTIONARY TO repadmin;

/*
```

### Step 4 Register the propagator at orc1.example.com.

The propagator is responsible for propagating the deferred transaction queue to other master sites.

```
*/

BEGIN
   DBMS_DEFER_SYS.REGISTER_PROPAGATOR (
      username => 'repadmin');
END;
/

/*
```

### Step 5 Register the receiver at orc1.example.com.

The receiver receives the propagated deferred transactions sent by the propagator from other master sites.

```
*/

BEGIN
   DBMS_REPCAT_ADMIN.REGISTER_USER_REPGROUP (
      username => 'repadmin',
      privilege_type => 'receiver',
      list_of_gnames => NULL);
END;
/

/*
```

### Step 6 Schedule purge at master site orc1.example.com.

To keep the size of the deferred transaction queue in check, purge successfully completed deferred transactions. The SCHEDULE_PURGE procedure automates the purge process for you. You must execute this procedure as the replication administrator.

---

**Note:** Date expressions are used for the NEXT_DATE and INTERVAL parameters. For example:

- Now is specified as: SYSDATE
- An interval of one hour is specified as: SYSDATE + 1/24
- An interval of seven days could be specified as: SYSDATE + 7

---

```
*/

CONNECT repadmin@orc1.example.com
```

```
BEGIN
   DBMS_DEFER_SYS.SCHEDULE_PURGE (
      next_date => SYSDATE,
      interval => 'SYSDATE + 1/24',
      delay_seconds => 0);
END;
/

/*
```

> **See Also:** *Oracle Database Administrator's Guide* and *Oracle Database SQL Language Reference* for more information about date expressions

### Step 7 Create proxy master site users at orc1.example.com.

If you plan to create materialized view sites based on this master site, then create proxy master site users at `orc1.example.com` that correspond to users at the materialized view site.

Create the proxy materialized view administrator.

The proxy materialized view administrator performs tasks at the target master site on behalf of the materialized view administrator at the materialized view site.

```
*/

CONNECT system@orc1.example.com

CREATE USER proxy_mviewadmin IDENTIFIED BY &password;

BEGIN
   DBMS_REPCAT_ADMIN.REGISTER_USER_REPGROUP (
      username => 'proxy_mviewadmin',
      privilege_type => 'proxy_snapadmin',
      list_of_gnames => NULL);
END;
/

-- Place GRANT SELECT_CATALOG_ROLE statement here if necessary.

/*
```

If you want your materialized view administrator at materialized view sites to be able to perform administrative operations using the Advanced Replication interface in Oracle Enterprise Manager, then grant `SELECT_CATALOG_ROLE` to `proxy_mviewadmin`:

```
GRANT SELECT_CATALOG_ROLE TO proxy_mviewadmin;
```

Granting this privilege to the `proxy_mviewadmin` is not required if you do not plan to use the Advanced Replication interface in Oracle Enterprise Manager. However, if you plan to use the Advanced Replication interface, then move the `GRANT` statement to the line directly after the previous `REGISTER_USER_REPGROUP` statement.

Create the proxy refresher.

The proxy refresher performs tasks at the master site on behalf of the refresher at the materialized view site.

```
*/
```

```
CREATE USER proxy_refresher IDENTIFIED BY &password;

GRANT CREATE SESSION TO proxy_refresher;
GRANT SELECT ANY TABLE TO proxy_refresher;

/*
```

**See Also:** "Security Setup for Materialized View Replication" on page A-5

## Setting Up orc2.example.com

Complete the following steps to set up the `orc2.example.com` master site.

### Step 1 Connect as SYSTEM at orc2.example.com.

You must connect as `SYSTEM` to the database that you want to set up for replication. After you set up `orc2.example.com`, begin with Step 1 for site `orc3.example.com` on page 2-10.

```
*/

CONNECT system@orc2.example.com

/*
```

### Step 2 Create the replication administrator at orc2.example.com.

The replication administrator must be granted the necessary privileges to create and manage a replication environment. The replication administrator must be created at each database that participates in the replication environment.

```
*/

CREATE USER repadmin IDENTIFIED BY &password;

/*
```

**Note:** Enter an appropriate password for the administrative user.

**See Also:** *Oracle Database Security Guide* for guidelines for choosing passwords

### Step 3 Grant privileges to replication administrator at orc2.example.com.

Execute the `GRANT_ADMIN_ANY_SCHEMA` procedure to grant the replication administrator powerful privileges to create and manage a replicated environment.

```
*/

BEGIN
   DBMS_REPCAT_ADMIN.GRANT_ADMIN_ANY_SCHEMA (
      username => 'repadmin');
END;
/

/*
```

If you want your repadmin to be able to create materialized view logs for any replicated table, then grant COMMENT ANY TABLE and LOCK ANY TABLE privileges to repadmin:

```
*/

GRANT COMMENT ANY TABLE TO repadmin;
GRANT LOCK ANY TABLE TO repadmin;

/*
```

If you want your repadmin to be able to connect to the Advanced Replication interface in Oracle Enterprise Manager, then grant SELECT ANY DICTIONARY to repadmin:

```
*/

GRANT SELECT ANY DICTIONARY TO repadmin;

/*
```

**Step 4 Register the propagator at orc2.example.com.**

The propagator is responsible for propagating the deferred transaction queue to other master sites.

```
*/

BEGIN
   DBMS_DEFER_SYS.REGISTER_PROPAGATOR (
      username => 'repadmin');
END;
/

/*
```

**Step 5 Register the receiver at orc2.example.com.**

The receiver receives the propagated deferred transactions sent by the propagator from the other master sites.

```
*/

BEGIN
   DBMS_REPCAT_ADMIN.REGISTER_USER_REPGROUP (
      username => 'repadmin',
      privilege_type => 'receiver',
      list_of_gnames => NULL);
END;
/

/*
```

**Step 6 Schedule purge at master site at orc2.example.com.**

To keep the size of the deferred transaction queue in check, purge successfully completed deferred transactions. The SCHEDULE_PURGE procedure automates the purge process for you. You must execute this procedure as the replication administrator.

```
*/

CONNECT repadmin@orc2.example.com

BEGIN
```

```
   DBMS_DEFER_SYS.SCHEDULE_PURGE (
      next_date => SYSDATE,
      interval => 'SYSDATE + 1/24',
      delay_seconds => 0);
END;
/

/*
```

### Step 7 Create proxy master site users at orc2.example.com.

If you plan to create materialized view sites based on this master site, then create proxy master site users at `orc2.example.com` that correspond to users at the materialized view site.

Create the proxy materialized view administrator.

The proxy materialized view administrator performs tasks at the target master site on behalf of the materialized view administrator at the materialized view site.

```
*/

CONNECT system@orc2.example.com

CREATE USER proxy_mviewadmin IDENTIFIED BY &password;

BEGIN
   DBMS_REPCAT_ADMIN.REGISTER_USER_REPGROUP (
      username => 'proxy_mviewadmin',
      privilege_type => 'proxy_snapadmin',
      list_of_gnames => NULL);
END;
/

-- Place GRANT SELECT_CATALOG_ROLE statement here if necessary.

/*
```

If you want your materialized view administrator at materialized view sites to be able to perform administrative operations using the Advanced Replication interface in Oracle Enterprise Manager, then grant `SELECT_CATALOG_ROLE` to `proxy_mviewadmin`:

```
*/

GRANT SELECT_CATALOG_ROLE TO proxy_mviewadmin;

/*
```

Granting this privilege to the `proxy_mviewadmin` is not required if you do not plan to use the Advanced Replication interface in Oracle Enterprise Manager. However, if you plan to use the Advanced Replication interface, then move the `GRANT` statement to the line directly after the previous `REGISTER_USER_REPGROUP` statement.

Create the proxy refresher.

The proxy refresher performs tasks at the master site on behalf of the refresher at the materialized view site.

```
*/

CREATE USER proxy_refresher IDENTIFIED BY &password;
```

```
GRANT CREATE SESSION TO proxy_refresher;
GRANT SELECT ANY TABLE TO proxy_refresher;

/*
```

**See Also:** "Security Setup for Materialized View Replication" on page A-5

## Setting Up orc3.example.com

Complete the following steps to set up the `orc3.example.com` master site.

### Step 1 Connect as SYSTEM at orc3.example.com.

You must connect as `SYSTEM` to the database that you want to set up for replication.

```
*/

CONNECT system@orc3.example.com

/*
```

### Step 2 Create the replication administrator at orc3.example.com.

The replication administrator must be granted the necessary privileges to create and manage a replication environment. The replication administrator must be created at each database that participates in the replication environment.

```
*/

CREATE USER repadmin IDENTIFIED BY &password;

/*
```

**Note:** Enter an appropriate password for the administrative user.

**See Also:** *Oracle Database Security Guide* for guidelines for choosing passwords

### Step 3 Grant privileges to replication administrator at orc3.example.com.

Execute the `GRANT_ADMIN_ANY_SCHEMA` procedure to grant the replication administrator powerful privileges to create and manage a replicated environment.

```
*/

BEGIN
   DBMS_REPCAT_ADMIN.GRANT_ADMIN_ANY_SCHEMA (
      username => 'repadmin');
END;
/

/*
```

If you want your `repadmin` to be able to create materialized view logs for any replicated table, then grant `COMMENT ANY TABLE` and `LOCK ANY TABLE` to `repadmin`:

```
*/

GRANT COMMENT ANY TABLE TO repadmin;
GRANT LOCK ANY TABLE TO repadmin;
```

```
/*
```

If you want your repadmin to be able to connect to the Advanced Replication interface in Oracle Enterprise Manager, then grant SELECT ANY DICTIONARY to repadmin:

```
*/

GRANT SELECT ANY DICTIONARY TO repadmin;

/*
```

### Step 4 Register the propagator at orc3.example.com.

The propagator is responsible for propagating the deferred transaction queue to other master sites.

```
*/

BEGIN
   DBMS_DEFER_SYS.REGISTER_PROPAGATOR (
      username => 'repadmin');
END;
/

/*
```

### Step 5 Register the receiver at orc3.example.com.

The receiver receives the propagated deferred transactions sent by the propagator from the other master sites.

```
*/

BEGIN
   DBMS_REPCAT_ADMIN.REGISTER_USER_REPGROUP (
      username => 'repadmin',
      privilege_type => 'receiver',
      list_of_gnames => NULL);
END;
/

/*
```

### Step 6 Schedule purge at master site at orc3.example.com.

To keep the size of the deferred transaction queue in check, purge successfully completed deferred transactions. The SCHEDULE_PURGE API automates the purge process for you. You must execute this procedure as the replication administrator.

```
*/

CONNECT repadmin@orc3.example.com

BEGIN
   DBMS_DEFER_SYS.SCHEDULE_PURGE (
      next_date => SYSDATE,
      interval => 'SYSDATE + 1/24',
      delay_seconds => 0);
END;
/

/*
```

**Step 7 Create proxy master site users at orc1.example.com.**

If you plan to create materialized view sites based on this master site, then create proxy master site users at `orc1.example.com` that correspond to users at the materialized view site.

Create the proxy materialized view administrator.

The proxy materialized view administrator performs tasks at the target master site on behalf of the materialized view administrator at the materialized view site.

```
*/

CONNECT system@orc3.example.com

CREATE USER proxy_mviewadmin IDENTIFIED BY &password;

BEGIN
   DBMS_REPCAT_ADMIN.REGISTER_USER_REPGROUP (
      username => 'proxy_mviewadmin',
      privilege_type => 'proxy_snapadmin',
      list_of_gnames => NULL);
END;
/

-- Place GRANT SELECT_CATALOG_ROLE statement here if necessary.

/*
```

If you want your materialized view administrator at materialized view sites to be able to perform administrative operations using the Advanced Replication interface in Oracle Enterprise Manager, then grant `SELECT_CATALOG_ROLE` to `proxy_mviewadmin`:

```
*/

GRANT SELECT_CATALOG_ROLE TO proxy_mviewadmin;

/*
```

Granting this privilege to the `proxy_mviewadmin` is not required if you do not plan to use the Advanced Replication interface in Oracle Enterprise Manager. However, if you plan to use the Advanced Replication interface, then move the `GRANT` statement to the line directly after the previous `REGISTER_USER_REPGROUP` statement.

Create proxy refresher.

The proxy refresher performs tasks at the master site on behalf of the refresher at the materialized view site.

```
*/

CREATE USER proxy_refresher IDENTIFIED BY &password;

GRANT CREATE SESSION TO proxy_refresher;
GRANT SELECT ANY TABLE TO proxy_refresher;

/*
```

**See Also:** "Security Setup for Materialized View Replication" on page A-5

## Creating Scheduled Links Between the Master Sites

Complete the following steps to create scheduled links between the master sites.

**Step 1 Create database links between master sites.**

The database links provide the necessary distributed mechanisms to allow the different replication sites to replicate data among themselves. Before you create any private database links, you must create the public database links that each private database link will use. You then must create a database link between all replication administrators at each of the master sites that you have set up.

> **See Also:** *Oracle Database Administrator's Guide* for more information about database links

```
*/

CONNECT system@orc1.example.com
CREATE PUBLIC DATABASE LINK orc2.example.com USING 'orc2.example.com';
CREATE PUBLIC DATABASE LINK orc3.example.com USING 'orc3.example.com';

CONNECT repadmin@orc1.example.com
CREATE DATABASE LINK orc2.example.com CONNECT TO repadmin
 IDENTIFIED BY &password;
CREATE DATABASE LINK orc3.example.com CONNECT TO repadmin
 IDENTIFIED BY &password;

CONNECT system@orc2.example.com
CREATE PUBLIC DATABASE LINK orc1.example.com USING 'orc1.example.com';
CREATE PUBLIC DATABASE LINK orc3.example.com USING 'orc3.example.com';

CONNECT repadmin@orc2.example.com
CREATE DATABASE LINK orc1.example.com CONNECT TO repadmin
 IDENTIFIED BY &password;
CREATE DATABASE LINK orc3.example.com CONNECT TO repadmin
 IDENTIFIED BY &password;

CONNECT system@orc3.example.com
CREATE PUBLIC DATABASE LINK orc1.example.com USING 'orc1.example.com';
CREATE PUBLIC DATABASE LINK orc2.example.com USING 'orc2.example.com';

CONNECT repadmin@orc3.example.com
CREATE DATABASE LINK orc1.example.com CONNECT TO repadmin
 IDENTIFIED BY &password;
CREATE DATABASE LINK orc2.example.com CONNECT TO repadmin
 IDENTIFIED BY &password;

/*
```

**Step 2 Define a schedule for each database link to create scheduled links.**

Create a scheduled link by defining a database link when you execute the `SCHEDULE_PUSH` procedure. The scheduled link determines how often your deferred transaction queue is propagated to each of the other master sites. you must execute the `SCHEDULE_PUSH` procedure for each database link that you created in Step 1. The database link is specified in the destination parameter of the `SCHEDULE_PUSH` procedure.

Even when using Oracle's asynchronous replication mechanisms, you can configure a scheduled link to simulate continuous, real-time replication. The scheduled links in this example simulate continuous replication.

**See Also:** *Oracle Database Advanced Replication* for more information about simulating continuous replication

```
*/

CONNECT repadmin@orc1.example.com

BEGIN
   DBMS_DEFER_SYS.SCHEDULE_PUSH (
      destination => 'orc2.example.com',
      interval => 'SYSDATE + (1/144)',
      next_date => SYSDATE,
      parallelism => 1,
      execution_seconds => 1500,
      delay_seconds => 1200);
END;
/

BEGIN
   DBMS_DEFER_SYS.SCHEDULE_PUSH (
      destination => 'orc3.example.com',
      interval => 'SYSDATE + (1/144)',
      next_date => SYSDATE,
      parallelism => 1,
      execution_seconds => 1500,
      delay_seconds => 1200);
END;
/

CONNECT repadmin@orc2.example.com

BEGIN
   DBMS_DEFER_SYS.SCHEDULE_PUSH (
      destination => 'orc1.example.com',
      interval => 'SYSDATE + (1/144)',
      next_date => SYSDATE,
      parallelism => 1,
      execution_seconds => 1500,
      delay_seconds => 1200);
END;
/

BEGIN
   DBMS_DEFER_SYS.SCHEDULE_PUSH (
      destination => 'orc3.example.com',
      interval => 'SYSDATE + (1/144)',
      next_date => SYSDATE,
      parallelism => 1,
      execution_seconds => 1500,
      delay_seconds => 1200);
END;
/

CONNECT repadmin@orc3.example.com

BEGIN
   DBMS_DEFER_SYS.SCHEDULE_PUSH (
      destination => 'orc1.example.com',
      interval => 'SYSDATE + (1/144)',
      next_date => SYSDATE,
```

```
      parallelism => 1,
      execution_seconds => 1500,
      delay_seconds => 1200);
END;
/

BEGIN
   DBMS_DEFER_SYS.SCHEDULE_PUSH (
      destination => 'orc2.example.com',
      interval => 'SYSDATE + (1/144)',
      next_date => SYSDATE,
      parallelism => 1,
      execution_seconds => 1500,
      delay_seconds => 1200);
END;
/

SET ECHO OFF

SPOOL OFF

/*************************END OF SCRIPT**********************************/
```

# Setting Up Materialized View Sites

***Figure 2–3 Setting Up Materialized View Sites***

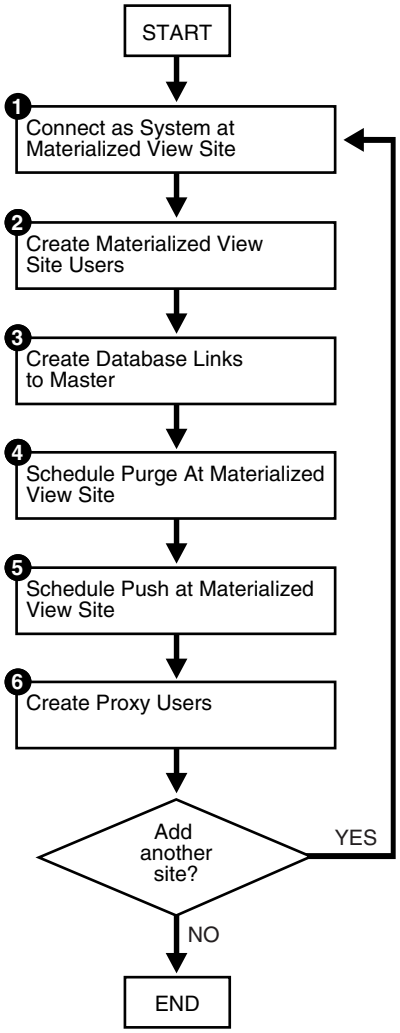


## Setting Up mv1.example.com

Complete the following steps to set up the `mv1.example.com` master materialized view site. `mv1.example.com` is a master materialized view site because `mv2.example.com` will be based on it. Before you set up the materialized sites, configure your network and Oracle Net so that all `mv1.example.com` can communicate with `orc1.example.com` and `mv2.example.com` can communicate with `mv1.example.com`.

---

**Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

---

```
/************************ BEGINNING OF SCRIPT ******************************
```

### Step 1 Connect as SYSTEM at materialized view site at mv1.example.com.

You must connect as SYSTEM to the database that you want to set up as a materialized view site.

```
*/

SET ECHO ON

SPOOL setup_mvs.out

CONNECT system@mv1.example.com

/*
```

### Step 2 Create materialized view site users at mv1.example.com.

Several users must be created at the materialized view site. These users are:

- Materialized view administrator
- Propagator
- Refresher
- Receiver (if the site will serve as a master materialized view site for other materialized views, as mv1.example.com is)

Complete the following tasks to create these users.

Create the materialized view administrator.

The materialized view administrator is responsible for creating and managing the materialized view site. Execute the GRANT_ADMIN_ANY_SCHEMA procedure to grant the materialized view administrator the appropriate privileges.

```
*/

ACCEPT password PROMPT 'Enter password for user: ' HIDE

CREATE USER mviewadmin IDENTIFIED BY &password;

BEGIN
   DBMS_REPCAT_ADMIN.GRANT_ADMIN_ANY_SCHEMA (
      username => 'mviewadmin');
END;
/

GRANT COMMENT ANY TABLE TO mviewadmin;

GRANT LOCK ANY TABLE TO mviewadmin;

/*
```

If you want your mviewadmin to be able to connect to the Advanced Replication interface in Oracle Enterprise Manager, then grant SELECT ANY DICTIONARY to mviewadmin:

```
*/

GRANT SELECT ANY DICTIONARY TO mviewadmin;

/*
```

Create the propagator.

The propagator is responsible for propagating the deferred transaction queue to the target master site.

```
*/

CREATE USER propagator IDENTIFIED BY &password;

BEGIN
   DBMS_DEFER_SYS.REGISTER_PROPAGATOR (
      username => 'propagator');
END;
/

/*
```

Create the refresher.

The refresher is responsible for "pulling" changes made to the replicated tables at the target master site to the materialized view site. This user refreshes one or more materialized views. If you want the mviewadmin user to be the refresher, then this step is not required.

```
*/

CREATE USER refresher IDENTIFIED BY &password;

GRANT CREATE SESSION TO refresher;

GRANT ALTER ANY MATERIALIZED VIEW TO refresher;

/*
```

Register the receiver.

The receiver receives the propagated deferred transactions sent by the propagator from materialized view sites. The receiver is necessary only if the site will function as a master materialized view site for other materialized view sites.

```
*/

BEGIN
   DBMS_REPCAT_ADMIN.REGISTER_USER_REPGROUP (
      username => 'mviewadmin',
      privilege_type => 'receiver',
      list_of_gnames => NULL);
END;
/

/*
```

---

**Note:** Enter appropriate passwords for the administrative users.

---

**See Also:** *Oracle Database Security Guide* for guidelines for choosing passwords

**Step 3 Create database links to the master site.**

Create a public database link.

```
*/

CONNECT system@mv1.example.com

CREATE PUBLIC DATABASE LINK orc1.example.com USING 'orc1.example.com';

/*
```

Create the materialized view administrator database link.

You must create a database link from the materialized view administrator at the materialized view site to the proxy materialized view administrator at the master site.

```
*/

CONNECT mviewadmin@mv1.example.com;

CREATE DATABASE LINK orc1.example.com
  CONNECT TO proxy_mviewadmin IDENTIFIED BY &password;

/*
```

Create the propagator/receiver database link.

You must create a database link from the propagator at the materialized view site to the receiver at the master site. The receiver was defined when you created the master site.

```
*/

CONNECT propagator@mv1.example.com

CREATE DATABASE LINK orc1.example.com
  CONNECT TO repadmin IDENTIFIED BY &password;

/*
```

**See Also:** Step 5 on page 2-5

**Step 4 Schedule purge at the mv1.example.com materialized view site.**

To keep the size of the deferred transaction queue in check, purge successfully completed deferred transactions. The `SCHEDULE_PURGE` procedure automates the purge process for you. If your materialized view site only contains "read-only" materialized views, then you do not need to execute this procedure.

```
*/

CONNECT mviewadmin@mv1.example.com

BEGIN
   DBMS_DEFER_SYS.SCHEDULE_PURGE (
   next_date => SYSDATE,
   interval => 'SYSDATE + 1/24',
   delay_seconds => 0,
   rollback_segment => '');
END;
/

/*
```

**Step 5 Schedule push at the mv1.example.com materialized view site (optional).**

If the materialized view site has a constant connection to its master site, then you optionally can schedule push at the `mv1.example.com` materialized view site. If the materialized view site is disconnected from its master site for extended periods of time, then it is typically better not to schedule push and refresh on demand, which pushes changes to the master site.

The `SCHEDULE_PUSH` procedure schedules when the deferred transaction queue should be propagated to the target master site.

```
*/

BEGIN
   DBMS_DEFER_SYS.SCHEDULE_PUSH (
      destination => 'orc1.example.com',
      interval => 'SYSDATE + 1/24',
      next_date => SYSDATE,
      stop_on_error => FALSE,
      delay_seconds => 0,
      parallelism => 0);
END;
/

/*
```

**Step 6 Create proxy users at the mv1.example.com materialized view site.**

Create the proxy materialized view administrator.

The proxy materialized view administrator performs tasks at the target master materialized view site on behalf of the materialized view administrator at the materialized view sites based on this materialized view site. This user is not required if the site will not function as a master materialized view site for other materialized view sites.

```
*/

CONNECT system@mv1.example.com

CREATE USER proxy_mviewadmin IDENTIFIED BY &password;

BEGIN
   DBMS_REPCAT_ADMIN.REGISTER_USER_REPGROUP (
      username => 'proxy_mviewadmin',
      privilege_type => 'proxy_snapadmin',
      list_of_gnames => NULL);
END;
/

-- Place GRANT SELECT_CATALOG_ROLE statement here if necessary.

/*
```

If you want your materialized view administrator at materialized view sites based on this materialized view site to be able to perform administrative operations using the Advanced Replication interface in Oracle Enterprise Manager, then grant `SELECT_CATALOG_ROLE` to `proxy_mviewadmin`:

```
GRANT SELECT_CATALOG_ROLE TO proxy_mviewadmin;
```

Granting this privilege to the proxy_mviewadmin is not required if you do not plan to use the Advanced Replication interface in Oracle Enterprise Manager. However, if you plan to use the Advanced Replication interface, then move the GRANT statement to the line directly after the previous REGISTER_USER_REPGROUP statement.

Create the proxy refresher.

The proxy refresher performs tasks at the master materialized view site on behalf of the refresher at the materialized view sites based on this materialized view site. This user is not required if the site will not function as a master materialized view site for other materialized view sites.

```
*/

CREATE USER proxy_refresher IDENTIFIED BY &password;

GRANT CREATE SESSION TO proxy_refresher;
GRANT SELECT ANY TABLE TO proxy_refresher;

/*
```

**See Also:** "Security Setup for Materialized View Replication" on page A-5

## Setting Up mv2.example.com

Complete the following steps to set up the mv2.example.com materialized view site. mv2.example.com is part of a multitier materialized view configuration because it is based on mv1.example.com, another materialized view.

### Step 1 Connect as SYSTEM at level 2 materialized view site mv2.example.com.

You must connect as SYSTEM to the database that you want to set up as a level 2 materialized view site. This site, mv2.example.com, will be a materialized view site that is based on mv1.example.com.

```
*/

CONNECT system@mv2.example.com

/*
```

### Step 2 Create level 2 materialized view site users at mv2.example.com.

Several users must be created at the level 2 materialized view site. These users are:

- Materialized view administrator
- Propagator
- Refresher

Complete the following tasks to create these users.

Create the materialized view administrator.

The materialized view administrator is responsible for creating and managing the level 2 materialized view site. Execute the GRANT_ADMIN_ANY_SCHEMA procedure to grant the materialized view administrator the appropriate privileges.

```
*/

CREATE USER mviewadmin IDENTIFIED BY &password;
```

```
BEGIN
   DBMS_REPCAT_ADMIN.GRANT_ADMIN_ANY_SCHEMA (
      username => 'mviewadmin');
END;
/

/*
```

If you want your `mviewadmin` to be able to connect to the Advanced Replication interface in Oracle Enterprise Manager, then grant `SELECT ANY DICTIONARY` to `mviewadmin`:

```
*/

GRANT SELECT ANY DICTIONARY TO mviewadmin;

/*
```

Create the propagator.

The propagator is responsible for propagating the deferred transaction queue to the target master materialized view site.

```
*/

CREATE USER propagator IDENTIFIED BY &password;

BEGIN
   DBMS_DEFER_SYS.REGISTER_PROPAGATOR (
      username => 'propagator');
END;
/

/*
```

Create the refresher.

The refresher is responsible for "pulling" changes made to the replicated materialized views at the target master materialized view site to the level 2 materialized view site.

```
*/

CREATE USER refresher IDENTIFIED BY &password;

GRANT CREATE SESSION TO refresher;
GRANT ALTER ANY MATERIALIZED VIEW TO refresher;

/*
```

---

**Note:** Enter appropriate passwords for the administrative users.

---

**See Also:** *Oracle Database Security Guide* for guidelines for choosing passwords

**Step 3 Create database links to master materialized view site.**

Create a public database link.

```
*/

CONNECT system@mv2.example.com
```

```
CREATE PUBLIC DATABASE LINK mv1.example.com USING 'mv1.example.com';

/*
```

Create materialized view administrator database link.

You must create a database link from the materialized view administrator at the level 2 materialized view site to the proxy materialized view administrator at the master materialized view site.

```
*/

CONNECT mviewadmin@mv2.example.com;

CREATE DATABASE LINK mv1.example.com
  CONNECT TO proxy_mviewadmin IDENTIFIED BY &password;

/*
```

Create a propagator/receiver database link.

You must create a database link from the propagator at the level 2 materialized view site to the receiver at the master materialized view site. The receiver was defined when you created the master materialized view site.

```
*/

CONNECT propagator@mv2.example.com

CREATE DATABASE LINK mv1.example.com
  CONNECT TO mviewadmin IDENTIFIED BY &password;

/*
```

**Step 4 Schedule purge at level 2 materialized view site at mv2.example.com.**

To keep the size of the deferred transaction queue in check, purge successfully completed deferred transactions. The `SCHEDULE_PURGE` procedure automates the purge process for you. If your level 2 materialized view site only contains "read-only" materialized views, then you do not need to execute this procedure.

```
*/

CONNECT mviewadmin@mv2.example.com

BEGIN
   DBMS_DEFER_SYS.SCHEDULE_PURGE (
     next_date => SYSDATE,
     interval => 'SYSDATE + 1/24',
     delay_seconds => 0,
     rollback_segment => '');
END;
/

/*
```

**Step 5 Schedule push at the mv2.example.com materialized view site (optional).**

If the materialized view site has a constant connection to its master materialized view site, then you optionally can schedule push at the `mv2.example.com` materialized view site. If the materialized view site is disconnected from its master materialized view site

for extended periods of time, then it is typically better not to schedule push and refresh on demand, which pushes changes to the master materialized view site.

The SCHEDULE_PUSH procedure schedules when the deferred transaction queue should be propagated to the target master materialized view site.

```
*/

CONNECT mviewadmin@mv2.example.com

BEGIN
   DBMS_DEFER_SYS.SCHEDULE_PUSH (
      destination => 'mv1.example.com',
      interval => 'SYSDATE + 1/24',
      next_date => SYSDATE,
      stop_on_error => FALSE,
      delay_seconds => 0,
      parallelism => 0);
END;
/

SET ECHO OFF

SPOOL OFF

/************************ END OF SCRIPT *********************************/
```

# 3

# Creating a Master Group

This chapter illustrates how to create a master group at a master replication site.

This chapter contains these topics:

- Overview of Creating a Master Group
- Creating a Master GroupCreating a Master Group

## Overview of Creating a Master Group

After you have set up your master sites, you are ready to build a master group. As illustrated in Figure 3–2, you must follow a specific sequence to successfully build a replication environment.

> **See Also:** "Configuring the Replication Sites" on page 2-1 for information about setting up master sites

In this chapter, you create the `hr_repg` master group and replicate the objects illustrated in Figure 3–1.

***Figure 3–1 Replicate the Tables in the hr Schema Between All Sites***

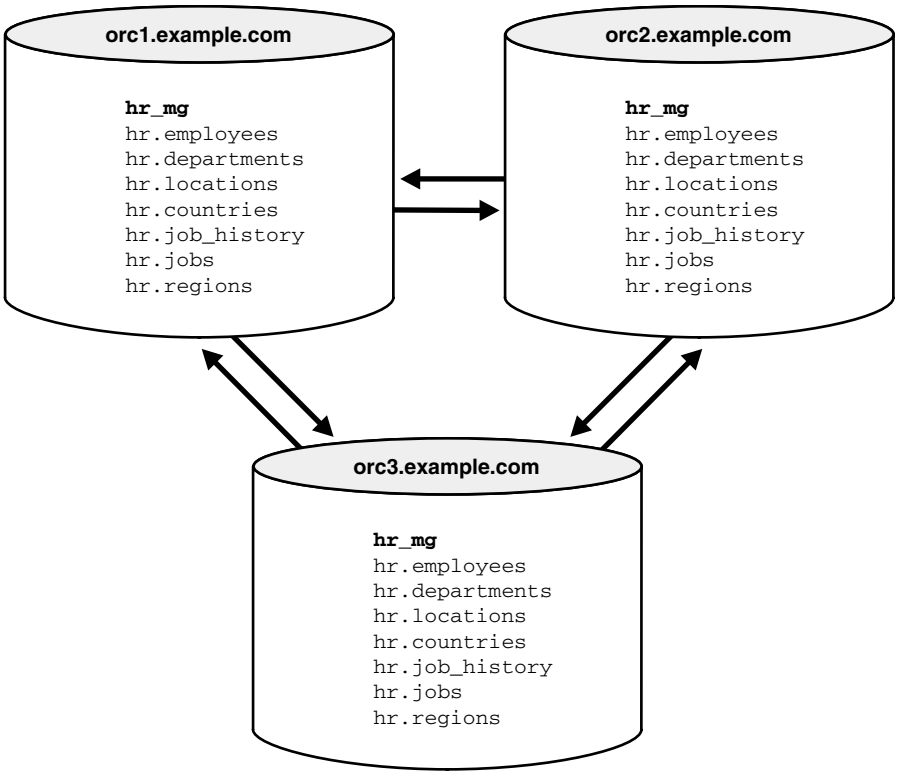


## Before You Start

In order for the script in this chapter to work as designed, it is assumed that the `hr` schema exists at `orc1.example.com`, `orc2.example.com`, and `orc3.example.com`. The `hr` schema includes the following database objects:

- `countries` table
- `departments` table
- `employees` table
- `jobs` table
- `job_history` table
- `locations` table
- `regions` table
- `dept_location_ix` index
- `emp_department_ix` index
- `emp_job_ix` index
- `emp_manager_ix` index
- `jhist_department_ix` index
- `jhist_employee_ix` index
- `jhist_job_ix` index
- `loc_country_ix` index

The indexes listed are the indexes based on foreign key columns in the hr schema. When replicating tables with foreign key referential constraints, Oracle recommends that you always index foreign key columns and replicate these indexes, unless no updates and deletes are allowed in the parent table. Indexes are not replicated automatically.

By default, the hr schema is installed automatically when you install Oracle Database. The example script in this chapter assumes that the hr schema exists at all master sites and that the schema contains all of these database objects at each site. The example script also assumes that the tables contain the data that is inserted automatically during Oracle installation. If the hr schema is not installed at your replication sites, then you can install it manually.

***Figure 3–2 Creating a Master Group***

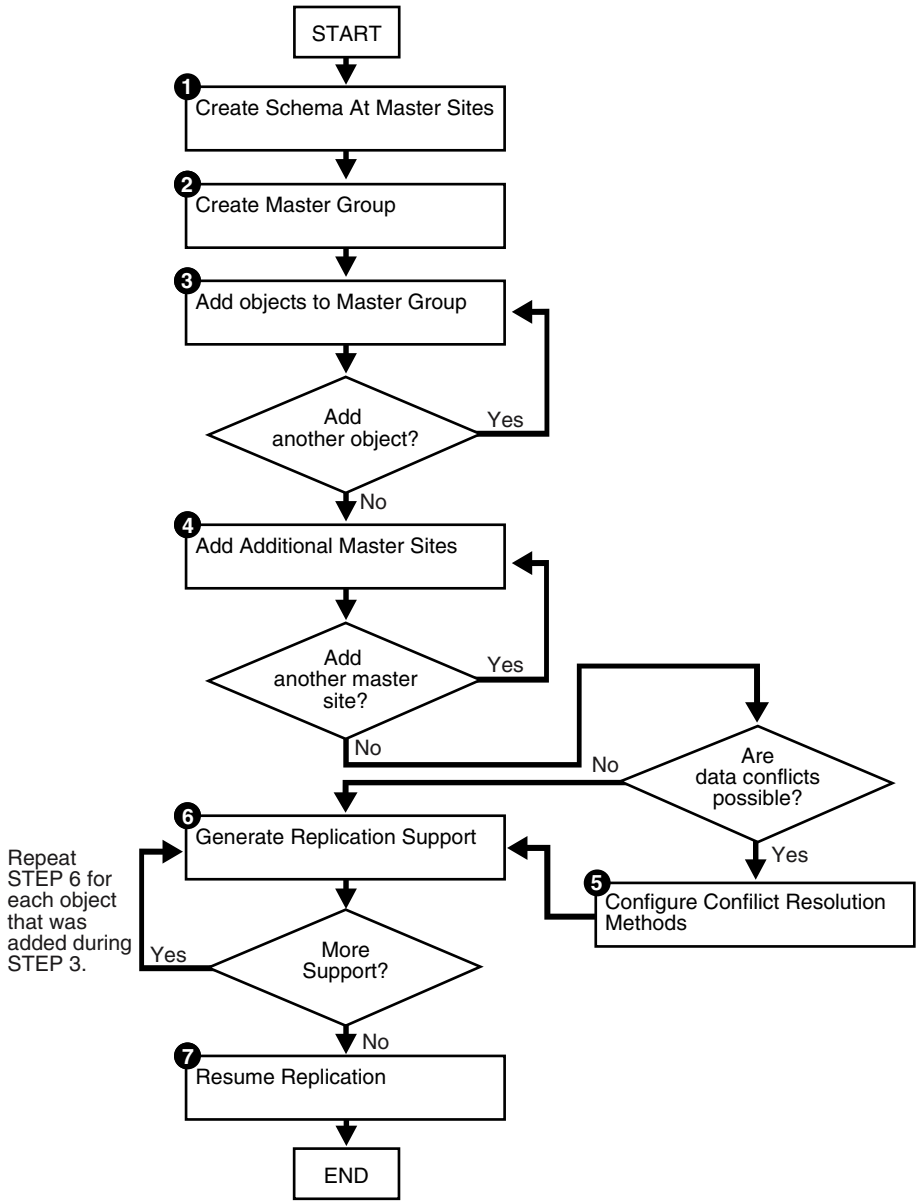

**See Also:** *Oracle Database Sample Schemas* for information about the hr schema and the other sample schemas, and for information about installing the sample schemas manually

## Creating a Master Group

Complete the following steps to create the hr_repg master group.

**Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

```
/************************* BEGINNING OF SCRIPT ****************************/

SET ECHO ON

SPOOL create_mg.out

CONNECT repadmin@orc1.example.com

/*
```

**Step 1 Create the schema at master sites.**

If the schema does not exist at all of the master sites participating in the master group, then create the schema now and grant it all of the necessary privileges. This example uses the hr schema, which is one of the sample schemas that are installed by default when you install Oracle. So, the hr schema should exist at all master sites.

```
*/

PAUSE Press <RETURN> to continue when the schema exists at all master sites.

/*
```

**Step 2 Create the master group.**

Use the CREATE_MASTER_REPGROUP procedure to define a new master group. When you add an object to your master group or perform other replication administrative tasks, you reference the master group name defined during this step. This step must be completed by the replication administrator.

```
*/

BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPGROUP (
      gname => 'hr_repg');
END;
/

/*
```

**Step 3 Add objects to master group.**

Use the CREATE_MASTER_REPOBJECT procedure to add an object to your master group. In most cases, you probably will be adding tables and indexes to your master group, but you can also add procedures, views, synonyms, and so on.

```
*/

BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
      gname => 'hr_repg',
      type => 'TABLE',
      oname => 'countries',
      sname => 'hr',
      use_existing_object => TRUE,
      copy_rows => FALSE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
      gname => 'hr_repg',
      type => 'TABLE',
      oname => 'departments',
      sname => 'hr',
      use_existing_object => TRUE,
      copy_rows => FALSE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
      gname => 'hr_repg',
      type => 'TABLE',
      oname => 'employees',
      sname => 'hr',
      use_existing_object => TRUE,
      copy_rows => FALSE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
      gname => 'hr_repg',
      type => 'TABLE',
      oname => 'jobs',
      sname => 'hr',
      use_existing_object => TRUE,
      copy_rows => FALSE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
      gname => 'hr_repg',
      type => 'TABLE',
      oname => 'job_history',
      sname => 'hr',
      use_existing_object => TRUE,
      copy_rows => FALSE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
      gname => 'hr_repg',
```

```
      type => 'TABLE',
      oname => 'locations',
      sname => 'hr',
      use_existing_object => TRUE,
      copy_rows => FALSE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
      gname => 'hr_repg',
      type => 'TABLE',
      oname => 'regions',
      sname => 'hr',
      use_existing_object => TRUE,
      copy_rows => FALSE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
      gname => 'hr_repg',
      type => 'INDEX',
      oname => 'dept_location_ix',
      sname => 'hr',
      use_existing_object => TRUE,
      copy_rows => FALSE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
      gname => 'hr_repg',
      type => 'INDEX',
      oname => 'emp_department_ix',
      sname => 'hr',
      use_existing_object => TRUE,
      copy_rows => FALSE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
      gname => 'hr_repg',
      type => 'INDEX',
      oname => 'emp_job_ix',
      sname => 'hr',
      use_existing_object => TRUE,
      copy_rows => FALSE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
      gname => 'hr_repg',
      type => 'INDEX',
      oname => 'emp_manager_ix',
      sname => 'hr',
      use_existing_object => TRUE,
      copy_rows => FALSE);
```

```
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
      gname => 'hr_repg',
      type => 'INDEX',
      oname => 'jhist_department_ix',
      sname => 'hr',
      use_existing_object => TRUE,
      copy_rows => FALSE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
      gname => 'hr_repg',
      type => 'INDEX',
      oname => 'jhist_employee_ix',
      sname => 'hr',
      use_existing_object => TRUE,
      copy_rows => FALSE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
      gname => 'hr_repg',
      type => 'INDEX',
      oname => 'jhist_job_ix',
      sname => 'hr',
      use_existing_object => TRUE,
      copy_rows => FALSE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
      gname => 'hr_repg',
      type => 'INDEX',
      oname => 'loc_country_ix',
      sname => 'hr',
      use_existing_object => TRUE,
      copy_rows => FALSE);
END;
/

/*
```

### Step 4 Add additional master sites.

After you have defined your master group at the master definition site (the site where the master group was created becomes the master definition site by default), you can define the other sites that will participate in the replication environment. You might have guessed that you will be adding the `orc2.example.com` and `orc3.example.com` sites to the replication environment. This example creates the master group at all master sites, but you have the option of creating the master group at one master site now and adding additional master sites later without quiescing the database. In this case, you can skip this step.

> **See Also:** "Adding New Master Sites without Quiescing the Master Group" on page 7-3 for more information

In this example, the use_existing_objects parameter in the ADD_MASTER_DATABASE procedure is set to TRUE because it is assumed that the hr schema exists at all master sites. In other words, it is assumed that the objects in the hr schema are precreated at all master sites. Also, the copy_rows parameter is set to FALSE because it is assumed that the identical data is stored in the tables at each master site.

> **Note:** When adding a master site to a master group that contains tables with circular dependencies or a table that contains a self-referential constraint, you must precreate the table definitions and manually load the data at the new master site. The following is an example of a circular dependency: Table A has a foreign key constraint on table B, and table B has a foreign key constraint on table A.

```
*/

BEGIN
   DBMS_REPCAT.ADD_MASTER_DATABASE (
      gname => 'hr_repg',
      master => 'orc2.example.com',
      use_existing_objects => TRUE,
      copy_rows => FALSE,
      propagation_mode => 'ASYNCHRONOUS');
END;
/

/*
```

> **Note:** You should wait until orc2.example.com appears in the DBA_REPSITES view before continuing. Execute the following SELECT statement in another SQL*Plus session to ensure that orc2.example.com has appeared:
>
> ```
> SELECT DBLINK FROM DBA_REPSITES WHERE GNAME = 'HR_REPG';
> ```

```
*/

PAUSE Press <RETURN> to continue.

BEGIN
   DBMS_REPCAT.ADD_MASTER_DATABASE (
      gname => 'hr_repg',
      master => 'orc3.example.com',
      use_existing_objects => TRUE,
      copy_rows => FALSE,
      propagation_mode => 'ASYNCHRONOUS');
END;
/

/*
```

> **Note:** You should wait until `orc3.example.com` appears in the `DBA_REPSITES` view before continuing. Execute the following `SELECT` statement in another SQL*Plus session to ensure that `orc3.example.com` has appeared:
>
> ```
> SELECT DBLINK FROM DBA_REPSITES WHERE GNAME = 'HR_REPG';
> ```

```
*/

PAUSE Press <RETURN> to continue.

/*
```

**Step 5 If conflicts are possible, then configure conflict resolution methods.**

> **See Also:** Chapter 6, "Configuring Conflict Resolution" for information about configuring conflict resolution methods

```
*/

PAUSE Press <RETURN> to continue after configuring conflict resolution methods
or if no conflict resolution methods are required.

/*
```

**Step 6 Generate replication support.**

```
*/

BEGIN
    DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT (
      sname => 'hr',
      oname => 'countries',
      type => 'TABLE',
      min_communication => TRUE);
END;
/

BEGIN
    DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT (
      sname => 'hr',
      oname => 'departments',
      type => 'TABLE',
      min_communication => TRUE);
END;
/

BEGIN
    DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT (
      sname => 'hr',
      oname => 'employees',
      type => 'TABLE',
      min_communication => TRUE);
END;
/

BEGIN
    DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT (
```

```
      sname => 'hr',
      oname => 'jobs',
      type => 'TABLE',
      min_communication => TRUE);
END;
/

BEGIN
    DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT (
      sname => 'hr',
      oname => 'job_history',
      type => 'TABLE',
      min_communication => TRUE);
END;
/

BEGIN
    DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT (
      sname => 'hr',
      oname => 'locations',
      type => 'TABLE',
      min_communication => TRUE);
END;
/

BEGIN
    DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT (
      sname => 'hr',
      oname => 'regions',
      type => 'TABLE',
      min_communication => TRUE);
END;
/

/*
```

**Note:** You should wait until the DBA_REPCATLOG view is empty before resuming master activity. Execute the following SELECT statement to monitor your DBA_REPCATLOG view:

```
SELECT COUNT(*) FROM DBA_REPCATLOG WHERE GNAME = 'HR_REPG';
```

```
*/

PAUSE Press <RETURN> to continue.

/*
```

### Step 7 Start replication.

After creating your master group, adding replication objects, generating replication support, and adding additional master databases, you must start replication activity. Use the RESUME_MASTER_ACTIVITY procedure to "turn on" replication for the specified master group.

```
*/

BEGIN
   DBMS_REPCAT.RESUME_MASTER_ACTIVITY (
```

```
      gname => 'hr_repg');
END;
/

SET ECHO OFF

SPOOL OFF

/************************ END OF SCRIPT *********************************/
```

# 4

# Creating a Deployment Template

This chapter illustrates how to build a deployment template using the replication management API.

This chapter contains these topics:


- Oracle Deployment Templates Concepts
- Before Creating the Deployment Template
- Creating a Deployment Template
- Packaging a Deployment Template for Instantiation
- Distributing Instantiation Files
- Instantiating a Deployment Template
- Refreshing a Refresh Group After Instantiation


Before you build materialized view environments, you must set up your master site, create a master group, and set up your intended materialized view sites. Also, if conflicts are possible at the master site due to activity at the materialized view sites you are creating, then configure conflict resolution for the master tables of the materialized views before you create the materialized view group.

**See Also:**

- "Setting Up Master Sites" on page 2-3
- "Overview of Creating a Master Group" on page 3-1
- "Setting Up Materialized View Sites" on page 2-16
- Chapter 6, "Configuring Conflict Resolution"

## Oracle Deployment Templates Concepts

Oracle offers deployment templates to allow the database administrator to package a materialized view environment for easy, custom, and secure distribution and installation. A deployment template can be simple (for example, it can contain a single materialized view with a fixed data set), or complex (for example, it can contain hundreds of materialized views with a dynamic data set based on one or more variables). The goal is to define the environment once and deploy the deployment template as often as necessary. Oracle deployment templates feature:

- Central control
- Repeated deployment of a materialized view environment
- Data subsetting at remote sites using template parameters

- Authorized user list to control template instantiation and data access

To prepare a materialized view environment for deployment, the DBA creates a deployment template at the master site. This template stores all of the information needed to deploy a materialized view environment, including the DDL to create the objects at the remote site and the target refresh group. This template also maintains links to user security information and template parameters for custom materialized view creation.

You cannot use deployment templates to instantiate the following types of objects:

- User-defined types
- User-defined type bodies
- User-defined operators
- Indextypes

Nor can you use deployment templates to instantiate any objects based on these types of objects.

> **See Also:** *Oracle Database Advanced Replication* for more conceptual information about deployment templates

## Before Creating the Deployment Template

If you want one of your master sites to support a materialized views that can be fast refreshed, then you must create materialized view logs for each master table that is replicated to a materialized view.

The example in this chapter uses the `hr` sample schema. Enter the following to create materialized view logs for the tables in the `hr` schema:

```
CONNECT hr@orc3.example.com
Enter password: password

CREATE MATERIALIZED VIEW LOG ON hr.countries;
CREATE MATERIALIZED VIEW LOG ON hr.departments;
CREATE MATERIALIZED VIEW LOG ON hr.employees;
CREATE MATERIALIZED VIEW LOG ON hr.jobs;
CREATE MATERIALIZED VIEW LOG ON hr.job_history;
CREATE MATERIALIZED VIEW LOG ON hr.locations;
CREATE MATERIALIZED VIEW LOG ON hr.regions;
```

> **See Also:** The `CREATE MATERIALIZED VIEW LOG` statement in the *Oracle Database SQL Language Reference* for detailed information about this SQL statement

## Creating a Deployment Template

This section contains a complete script example of how to construct a deployment template using the replication management API.

> **See Also:** *Oracle Database Advanced Replication* for conceptual and architectural information about deployment templates

***Figure 4–1 Creating a Deployment Template***

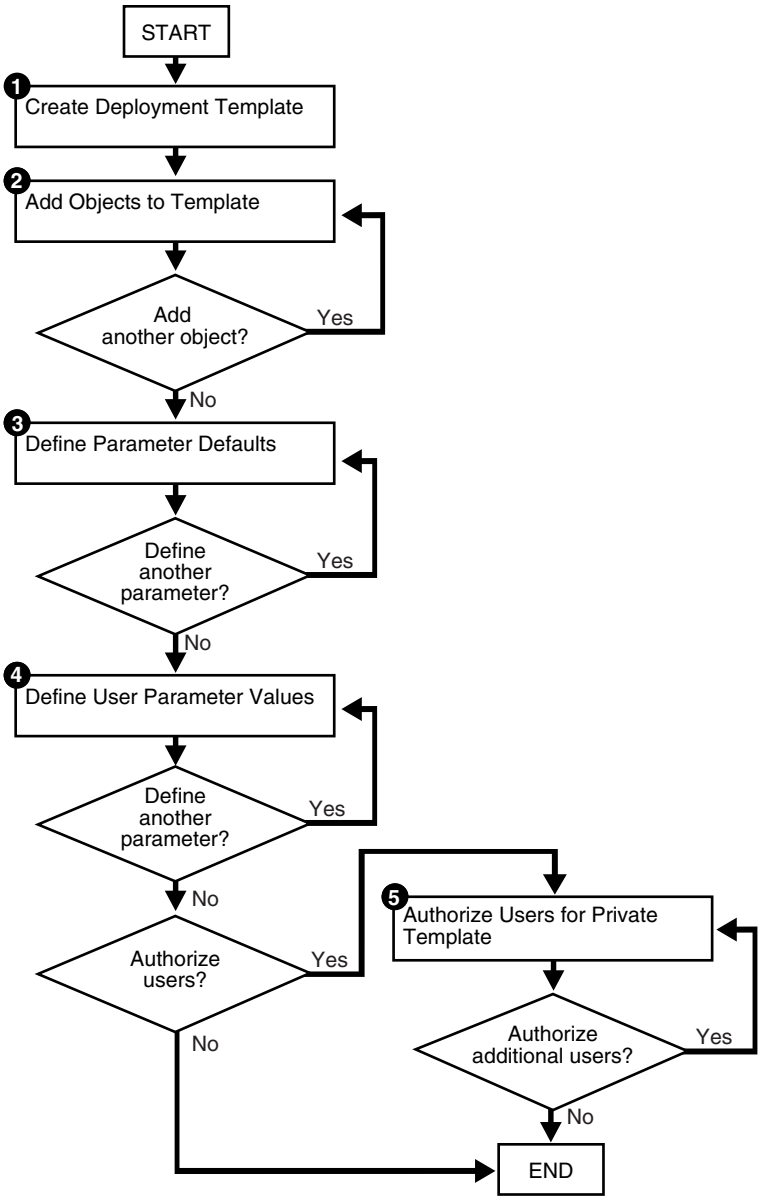


Be sure to read the comments contained within the scripts, as they contain important and useful information about building templates with the replication management API.

**Note:**

- You must use the Advanced Replication interface in Oracle Enterprise Manager to create materialized views with a subset of the columns in their master tables. See *Oracle Database Advanced Replication* and the Advanced Replication interface online Help for more information about column subsetting.
- If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

```
/************************* BEGINNING OF SCRIPT ******************************
```

This script creates a private deployment template that contains four template objects, two template parameters, a set of user parameter values, and an authorized user. Complete the following steps to build a template:

### Step 1 Create the deployment template.

Before assembling the components of your deployment template, use the `CREATE_REREFRESH_TEMPLATE` procedure to define the name of your deployment template, along with several other template characteristics (Public/Private status, target refresh group, and owner).

```
*/

SET ECHO ON

SPOOL create_dt.out

CONNECT repadmin@orc3.example.com

DECLARE
   a NUMBER;
BEGIN
   a := DBMS_REPCAT_RGT.CREATE_REFRESH_TEMPLATE (
            owner => 'hr',
            refresh_group_name => 'hr_refg',
            refresh_template_name => 'hr_refg_dt',
            template_comment => 'Human Resources Deployment Template',
            public_template => 'N');
END;
/

/*
```

### Step 2 Add objects to template.

Create `countries_mv` materialized view.

```
*/

DECLARE
   tempstring VARCHAR2(3000);
   a NUMBER;
BEGIN
   tempstring := 'CREATE MATERIALIZED VIEW hr.countries_mv
      REFRESH FAST WITH PRIMARY KEY FOR UPDATE AS SELECT
```

```
      country_id, country_name, region_id
      FROM hr.countries@:dblink';
   a := DBMS_REPCAT_RGT.CREATE_TEMPLATE_OBJECT (
         refresh_template_name => 'hr_refg_dt',
         object_name => 'countries_mv',
         object_type => 'MATERIALIZED VIEW',
         ddl_text => tempstring,
         master_rollback_seg => 'rbs');
END;
/

/*
```

Whenever you create a materialized view, always specify the schema name of the table owner in the query for the materialized view. In the example previously, `hr` is specified as the owner of the `countries` table.

Create `departments_mv` materialized view.

```
*/

DECLARE
   tempstring VARCHAR2(3000);
   a NUMBER;
BEGIN
   tempstring := 'CREATE MATERIALIZED VIEW hr.departments_mv
      REFRESH FAST WITH PRIMARY KEY FOR UPDATE AS SELECT
      department_id, department_name, manager_id, location_id
      FROM hr.departments@:dblink';
   a := DBMS_REPCAT_RGT.CREATE_TEMPLATE_OBJECT (
         refresh_template_name => 'hr_refg_dt',
         object_name => 'departments_mv',
         object_type => 'MATERIALIZED VIEW',
         ddl_text => tempstring,
         master_rollback_seg => 'rbs');
END;
/

/*
```

Create `employees_mv` materialized view.

```
*/

DECLARE
   tempstring VARCHAR2(3000);
   a NUMBER;
BEGIN
   tempstring := 'CREATE MATERIALIZED VIEW hr.employees_mv
      REFRESH FAST WITH PRIMARY KEY FOR UPDATE AS SELECT
      employee_id, first_name, last_name, email, phone_number,
      hire_date, job_id, salary, commission_pct, manager_id,
      department_id
      FROM hr.employees@:dblink WHERE department_id = :dept';
   a := DBMS_REPCAT_RGT.CREATE_TEMPLATE_OBJECT (
         refresh_template_name => 'hr_refg_dt',
         object_name => 'employees_mv',
         object_type => 'MATERIALIZED VIEW',
         ddl_text => tempstring,
         master_rollback_seg => 'rbs');
END;
```

```
/

/*
```

Create jobs_mv materialized view.

```
*/

DECLARE
   tempstring VARCHAR2(3000);
   a NUMBER;
BEGIN
   tempstring := 'CREATE MATERIALIZED VIEW hr.jobs_mv
      REFRESH FAST WITH PRIMARY KEY FOR UPDATE AS SELECT
      job_id, job_title, min_salary, max_salary
      FROM hr.jobs@:dblink';
   a := DBMS_REPCAT_RGT.CREATE_TEMPLATE_OBJECT (
            refresh_template_name => 'hr_refg_dt',
            object_name => 'jobs_mv',
            object_type => 'MATERIALIZED VIEW',
            ddl_text => tempstring,
            master_rollback_seg => 'rbs');
END;
/

/*
```

Create job_history_mv materialized view.

```
*/

DECLARE
   tempstring VARCHAR2(3000);
   a NUMBER;
BEGIN
   tempstring := 'CREATE MATERIALIZED VIEW hr.job_history_mv
      REFRESH FAST WITH PRIMARY KEY FOR UPDATE AS SELECT
      employee_id, start_date, end_date, job_id, department_id
      FROM hr.job_history@:dblink';
   a := DBMS_REPCAT_RGT.CREATE_TEMPLATE_OBJECT (
            refresh_template_name => 'hr_refg_dt',
            object_name => 'job_history_mv',
            object_type => 'MATERIALIZED VIEW',
            ddl_text => tempstring,
            master_rollback_seg => 'rbs');
 END;
/

/*
```

Create locations_mv materialized view.

```
*/

DECLARE
   tempstring VARCHAR2(3000);
   a NUMBER;
BEGIN
   tempstring := 'CREATE MATERIALIZED VIEW hr.locations_mv
      REFRESH FAST WITH PRIMARY KEY FOR UPDATE AS SELECT
      location_id, street_address, postal_code, city,
```

```
      state_province, country_id
      FROM hr.locations@:dblink';
   a := DBMS_REPCAT_RGT.CREATE_TEMPLATE_OBJECT (
         refresh_template_name => 'hr_refg_dt',
         object_name => 'locations_mv',
         object_type => 'MATERIALIZED VIEW',
         ddl_text => tempstring,
         master_rollback_seg => 'rbs');
END;
/

/*
```

Create `regions_mv` materialized view.

```
*/

DECLARE
   tempstring VARCHAR2(3000);
   a NUMBER;
BEGIN
   tempstring := 'CREATE MATERIALIZED VIEW hr.regions_mv
      REFRESH FAST WITH PRIMARY KEY FOR UPDATE AS SELECT
      region_id, region_name
      FROM hr.regions@:dblink';
   a := DBMS_REPCAT_RGT.CREATE_TEMPLATE_OBJECT (
         refresh_template_name => 'hr_refg_dt',
         object_name => 'regions_mv',
         object_type => 'MATERIALIZED VIEW',
         ddl_text => tempstring,
         master_rollback_seg => 'rbs');
END;
/

/*
```

**Step 3 Define parameter defaults.**

Rather than using the `CREATE_*` functions and procedures as in the other steps, use the `ALTER_TEMPLATE_PARM` procedure to define a template parameter value and prompt string. You use the `ALTER_*` procedure because the actual parameter was created in Step 1. Recall that you defined the `:dblink` and `:dept` template parameters in the `ddl_text` parameter. Oracle detects these parameters in the DDL and automatically creates the template parameter. Use the `ALTER_TEMPLATE_PARM` procedure to define the remainder of the template parameter information (that is, default parameter value and prompt string).

Complete the following tasks to define parameter defaults.

Define the default value for the `dept` parameter.

```
*/

BEGIN
   DBMS_REPCAT_RGT.ALTER_TEMPLATE_PARM (
      refresh_template_name => 'hr_refg_dt',
      parameter_name => 'dept',
      new_default_parm_value => '30',
      new_prompt_string => 'Enter your department number:',
      new_user_override => 'Y');
END;
/
```

```
/*
```

Define the default value for the `dblink` parameter.

```
*/

BEGIN
   DBMS_REPCAT_RGT.ALTER_TEMPLATE_PARM (
      refresh_template_name => 'hr_refg_dt',
      parameter_name => 'dblink',
      new_default_parm_value => 'orc3.example.com',
      new_prompt_string => 'Enter your master site:',
      new_user_override => 'Y');
END;
/

/*
```

**Step 4 Define user parameter values.**

To automate the instantiation of custom data sets at individual remote materialized view sites, you can define user parameter values that will be used automatically when the specified user instantiates the target template. The `CREATE_USER_PARM_VALUE` procedure enables you to assign a value to a parameter for a user.

Complete the following tasks to define user parameter values.

Define `dept` user parameter value for user `hr`.

```
*/

DECLARE
   a NUMBER;
BEGIN
   a := DBMS_REPCAT_RGT.CREATE_USER_PARM_VALUE (
            refresh_template_name => 'hr_refg_dt',
            parameter_name => 'dept',
            user_name => 'hr',
            parm_value => '20');
END;
/

/*
```

Define `dblink` user parameter value for user `hr`.

```
*/

DECLARE
   a NUMBER;
BEGIN
   a := DBMS_REPCAT_RGT.CREATE_USER_PARM_VALUE (
            refresh_template_name => 'hr_refg_dt',
            parameter_name => 'dblink',
            user_name => 'hr',
            parm_value => 'orc3.example.com');
END;
/

/*
```

**Step 5 Authorize users for private template.**

Because this is a private template (`public_template => 'N'` in the `DBMS_REPCAT_RGT.CREATE_REFRESH_TEMPLATE` function defined in Step 1 on page 4-4), you must authorize users to instantiate the `dt_personnel` deployment template. Use the `CREATE_USER_AUTHORIZATION` function in the `DBMS_REPCAT_RGT` package to create authorized users.

```
*/

DECLARE
   a NUMBER;
BEGIN
   a := DBMS_REPCAT_RGT.CREATE_USER_AUTHORIZATION (
           user_name => 'hr',
           refresh_template_name => 'hr_refg_dt');
END;
/

COMMIT;

SET ECHO OFF

SPOOL OFF

/************************ END OF SCRIPT *********************************/
```

# Packaging a Deployment Template for Instantiation

After building your deployment template, you must package the template for instantiation. This example illustrates how to use both the online and offline instantiation procedures. Notice that the instantiation procedures are very similar: you simply use either the `INSTANTIATE_ONLINE` function or `INSTANTIATE_OFFLINE` function according to your needs. This section describes two tasks: create the instantiation script and save the instantiation script to a file.

*Figure 4–2 Packaging and Instantiating a Deployment Template*

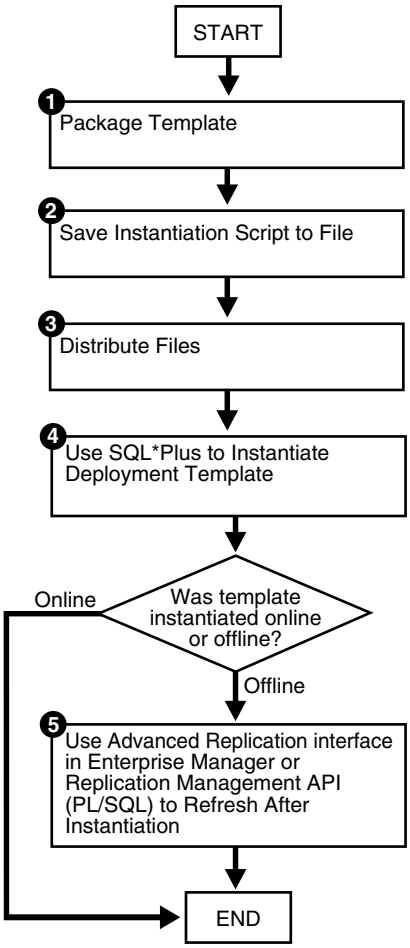


## Packaging a Deployment Template

When you execute either the `INSTANTIATE_OFFLINE` or the `INSTANTIATE_ONLINE` function, Oracle populates the `USER_REPCAT_TEMP_OUTPUT` data dictionary view with the script to create the remote materialized view environment. Both online and offline scripts contain the SQL statements to create the objects specified in the deployment template. The difference is that an offline instantiation script also contains the data to populate the objects. The online instantiation script does not contain the data. Rather, during online instantiation, the materialized view site connects to the master site to download the data.

Complete the steps in either the "Packaging a Deployment Template for Offline Instantiation" or "Packaging a Deployment Template for Online Instantiation" according to your needs.

---

**Note:** If you must execute either the `INSTANTIATE_OFFLINE` or the `INSTANTIATE_ONLINE` function more than once for a particular materialized view site, then run the `DROP_SITE_INSTANTIATION` procedure in the `DBMS_REPCAT_RGT` package before you attempt to repackage a template for the site. Otherwise, Oracle returns an error stating that there is a duplicate template site.

---

### Packaging a Deployment Template for Offline Instantiation

The INSTANTIATE_OFFLINE function creates a script that creates the materialized view environment according to the contents of a specified deployment template. In addition to containing the DDL (CREATE statements) to create the materialized view environment, this script also contains the DML (INSERT statements) to populate the materialized view environment with the appropriate data set.

> **Note:** If you are packaging your template at the same master site that contains the target master objects for your deployment template, then you must create a loopback database link.

```
--Use the INSTANTIATE_OFFLINE function to package the
--template for offline instantiation by a remote materialized view
--site. Executing this procedure both creates a script that
--creates that materialized view environment and populates the
--environment with the proper data set. This script is stored
--in the temporary USER_REPCAT_TEMP_OUTPUT view.

CONNECT repadmin@orc3.example.com
Enter password: password

SET SERVEROUTPUT ON
DECLARE
   dt_num NUMBER;
BEGIN
   dt_num := DBMS_REPCAT_RGT.INSTANTIATE_OFFLINE(
               refresh_template_name => 'hr_refg_dt',
               user_name => 'hr',
               site_name => 'mv4.example.com',
               next_date => SYSDATE,
               interval => 'SYSDATE + (1/144)');
   DBMS_OUTPUT.PUT_LINE('Template ID = ' || dt_num);
END;
/
COMMIT;
/
```

Make a note of the number that is returned for the dt_num variable. You must use this number when you select from the USER_REPCAT_TEMP_OUTPUT data dictionary view to retrieve the generated script. Ensure that you complete the steps in "Saving an Instantiation Script to File" on page 4-12 after you complete this section. This script is unique to an individual materialized view site and cannot be used for other materialized view sites.

### Packaging a Deployment Template for Online Instantiation

The INSTANTIATE_ONLINE function creates a script that creates the materialized view environment according to the contents of a specified deployment template. When this script is executed at the remote materialized view site, Oracle creates the materialized view site according to the DDL (CREATE statements) in the script and populates the environment with the appropriate data set from the master site. This requires that the remote materialized view site has a "live" connection to the master site.

> **See Also:** *Oracle Database Advanced Replication* for additional materialized view site requirements

```
--Use the INSTANTIATE_ONLINE function to "package" the
--template for online instantiation by a remote materialized view
--site. Executing this procedure creates a script which can
--then be used to create a materialized view environment. This script
--is stored in the temporary USER_REPCAT_TEMP_OUTPUT view.

CONNECT repadmin@orc3.example.com
Enter password: password

SET SERVEROUTPUT ON
DECLARE
   dt_num NUMBER;
BEGIN
   dt_num := DBMS_REPCAT_RGT.INSTANTIATE_ONLINE(
               refresh_template_name => 'hr_refg_dt',
               user_name => 'hr',
               site_name => 'mv4.example.com',
               next_date => SYSDATE,
               interval => 'SYSDATE + (1/144)');
   DBMS_OUTPUT.PUT_LINE('Template ID = ' || dt_num);
END;
/
COMMIT;
/
```

Make a note of the number that is returned for the `dt_num` variable. You must use this number when you select from the `USER_REPCAT_TEMP_OUTPUT` data dictionary view to retrieve the generated script. Ensure that you complete the steps in "Saving an Instantiation Script to File" after you complete this task.

## Saving an Instantiation Script to File

The best way to save the contents of the `USER_REPCAT_TEMP_OUTPUT` data dictionary view is to use the `UTL_FILE` package to save the contents of the `TEXT` column in the `USER_REPCAT_TEMP_OUTPUT` view to a file.

These contents are saved to a directory that corresponds to a directory object. To create a directory object, the `CREATE ANY DIRECTORY` privilege is required. If the replication administrator does not have this privilege, then connect as an administrative user who can grant privileges. For example:

```
GRANT CREATE ANY DIRECTORY TO repadmin;
```

---

**Note:** The following action must be performed immediately after you have called either the `INSTANTIATE_OFFLINE` or `INSTANTIATE_ONLINE` functions, because the contents of the `USER_REPCAT_TEMP_OUTPUT` data dictionary view are temporary. If you have not completed the steps in "Packaging a Deployment Template" on page 4-10, then do so now and then complete the following action.

---

**See Also:** *Oracle Database PL/SQL Packages and Types Reference* for more information about the `UTL_FILE` package

Enter the following to save the deployment template script to a file.

```
DECLARE
  fh UTL_FILE.FILE_TYPE;
```

```
  CURSOR ddlcursor(myid NUMBER) IS
  SELECT TEXT FROM USER_REPCAT_TEMP_OUTPUT WHERE OUTPUT_ID = myid ORDER BY LINE;
BEGIN
  fh := UTL_FILE.FOPEN ('file_location', 'file_name', 'w');
  UTL_FILE.PUT_LINE (fh, 'SET ECHO OFF;');
  FOR myrec IN ddlcursor(template_id) LOOP
    UTL_FILE.PUT_LINE(fh, myrec.text);
  END LOOP;
  UTL_FILE.PUT_LINE (fh, 'SET ECHO ON;');
  UTL_FILE.FFLUSH(fh);
  UTL_FILE.FCLOSE(fh);
END;
/
```

Notice that *file_location*, *file_name*, and *template_id* are placeholders. Substitute the correct values for your environment:

- Replace the *file_location* placeholder with the name of a directory object that represents the directory where you want to save the template script.
- Replace the *file_name* placeholder with name you want to use for the template script.
- Replace the *template_id* placeholder with the number returned by the INSTANTIATE_OFFLINE or INSTANTIATE_ONLINE function when you packaged the template previously.

For example, suppose you have the following values:

| Placeholder | Value |
|---|---|
| *file_location* | /home/gen_files/ |
| *file_name* | sf.sql |
| *template_id* | 18 |

Given these values, connect to the master site as the replication administrator and run the following procedure to save the template script to a file:

```
CONNECT repadmin@orc3.example.com
Enter password: password

CREATE DIRECTORY GFILES AS '/home/gen_files';

DECLARE
  fh UTL_FILE.FILE_TYPE;
  CURSOR ddlcursor(myid NUMBER) IS
  SELECT TEXT FROM USER_REPCAT_TEMP_OUTPUT WHERE OUTPUT_ID = myid
  ORDER BY LINE;
BEGIN
  fh := UTL_FILE.FOPEN ('GFILES', 'sf.sql', 'w');
  UTL_FILE.PUT_LINE (fh, 'SET ECHO OFF;');
  FOR myrec IN ddlcursor(18) LOOP
    UTL_FILE.PUT_LINE(fh, myrec.text);
  END LOOP;
  UTL_FILE.PUT_LINE (fh, 'SET ECHO ON;');
  UTL_FILE.FFLUSH(fh);
  UTL_FILE.FCLOSE(fh);
END;
/
```

## Distributing Instantiation Files

After creating the instantiation script and saving it to a file, you must distribute this file to the remote materialized view sites that must instantiate the template. You can distribute this file by posting the file on an FTP site or saving the file to a CD-ROM, floppy disk, or other distribution medium. You can also transfer the file using the DBMS_FILE_TRANSFER package.

## Instantiating a Deployment Template

After the instantiation script has been distributed to the remote materialized view sites, you are ready to instantiate the deployment template at the remote materialized view site. Ensure that you have set up the materialized view site before you instantiate the deployment template. The following script demonstrates how to complete the instantiation process at a remote materialized view site.

**See Also:**

- *Oracle Database Advanced Replication* for materialized view site requirements that must be met before instantiating your deployment template
- "Setting Up Materialized View Sites" on page 2-16

---

**Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

---

```
/************************* BEGINNING OF SCRIPT ******************************
```

**Step 1 If it does not exist, then create the schema at materialized view site.**

Before executing the instantiation script at the remote materialized view site, you must create the schema that contains the replicated objects.

The following illustrates creating the hr schema. This schema might exist in your database. In this case, the schema might need additional privileges, such as CREATE MATERIALIZED VIEW, ALTER ANY MATERIALIZED VIEW, and CREATE DATABASE LINK.

```
*/

SET ECHO ON

SPOOL instant_mv.out

CONNECT system@mv4.example.com

CREATE TABLESPACE demo_mv
 DATAFILE 'demo_mv.dbf' SIZE 10M AUTOEXTEND ON
 EXTENT MANAGEMENT LOCAL AUTOALLOCATE;

CREATE TEMPORARY TABLESPACE temp_mv
 TEMPFILE 'temp_mv.dbf' SIZE 5M AUTOEXTEND ON;
ACCEPT password PROMPT 'Enter password for user: ' HIDE

CREATE USER hr IDENTIFIED BY &password;
```

```
ALTER USER hr DEFAULT TABLESPACE demo_mv
              QUOTA UNLIMITED ON demo_mv;

ALTER USER hr TEMPORARY TABLESPACE temp_mv;

GRANT
  CREATE SESSION,
  CREATE TABLE,
  CREATE PROCEDURE,
  CREATE SEQUENCE,
  CREATE TRIGGER,
  CREATE VIEW,
  CREATE SYNONYM,
  ALTER SESSION,
  CREATE MATERIALIZED VIEW,
  ALTER ANY MATERIALIZED VIEW,
  CREATE DATABASE LINK
 TO hr;

/*
```

**Step 2 If they do not exist, then create database links for the schema.**

Before instantiating the deployment template, you must ensure that the necessary database links exist for the replicated schema. The owner of the materialized views needs a database link pointing to the `proxy_refresher` that was created when the master site was set up.

```
*/

CREATE PUBLIC DATABASE LINK orc3.example.com USING 'orc3.example.com';

CONNECT hr@mv4.example.com

CREATE DATABASE LINK orc3.example.com
   CONNECT TO proxy_refresher IDENTIFIED BY &password;

/*
```

> **See Also:** Step 7 on page 2-6 for more information about creating proxy master site users

**Step 3 Execute the instantiation script.**

```
*/

CONNECT mviewadmin@mv4.example.com

@d:\sf.sql

SET ECHO OFF

SPOOL OFF

/*
```

Depending on the size of the materialized view environment created and the amount of data loaded, the instantiation procedure might take a substantial amount of time.

```
************************* END OF SCRIPT *********************************/
```

## Refreshing a Refresh Group After Instantiation

If you have just instantiated a deployment template using the offline instantiation method, then you should perform a refresh of the refresh group as soon as possible by issuing the following execute statement:

```
CONNECT hr@mv4.example.com
Enter password: password

EXECUTE DBMS_REFRESH.REFRESH ('hr_refg');
```

# 5

# Creating a Materialized View Group

This chapter illustrates how to create a materialized view group at a remote materialized view replication site.

This chapter contains these topics:

- Overview of Creating a Materialized View Group
- Creating a Materialized View Group

Before you build materialized view environments, you must set up your master site, create a master group, and set up your intended materialized view sites. Also, if conflicts are possible at the master site due to activity at the materialized view sites you are creating, then configure conflict resolution for the master tables of the materialized views before you create the materialized view group.

**See Also:**

- "Setting Up Master Sites" on page 2-3
- "Overview of Creating a Master Group" on page 3-1
- "Setting Up Materialized View Sites" on page 2-16
- Chapter 6, "Configuring Conflict Resolution"

## Overview of Creating a Materialized View Group

After setting up your materialized view site and creating at least one master group, you are ready to create a materialized view group at a remote materialized view site. Figure 5–1 illustrates the process of creating a materialized view group.

**See Also:** Chapter 2, "Configuring the Replication Sites" for information about setting up a materialized view site, and see Chapter 3, "Creating a Master Group" for information about creating a master group.

***Figure 5–1 Creating a Materialized View Group***

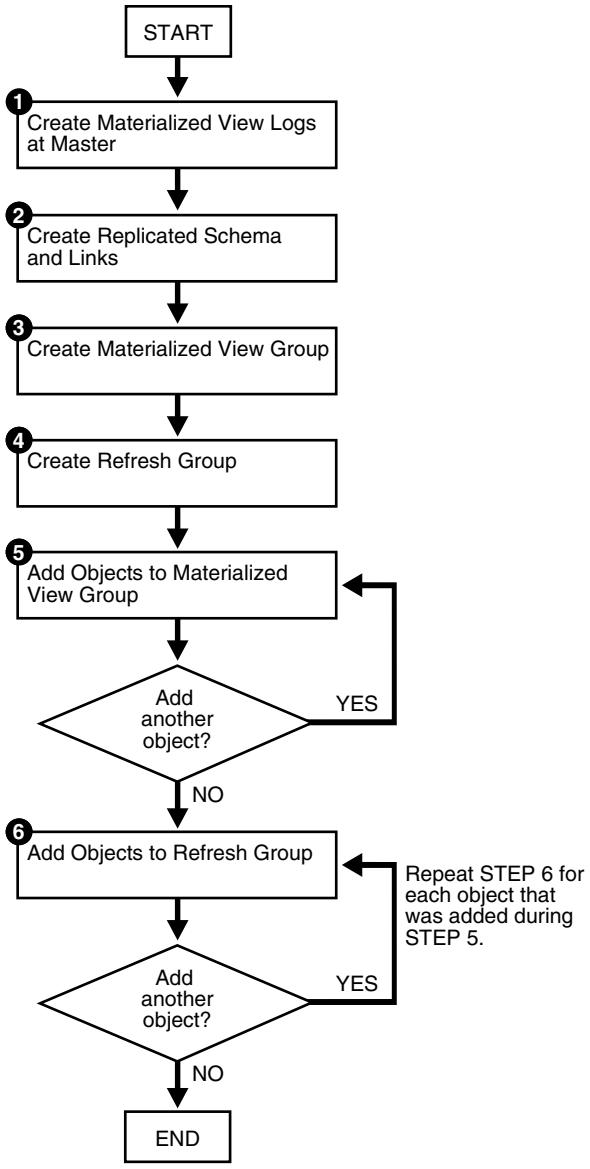


# Creating a Materialized View Group

This chapter guides you through the process of creating two materialized view groups at two different materialized view sites: `mv1.example.com` and `mv2.example.com`:

- The materialized view group at `mv1.example.com` is based on the objects in the `hr_repg` master group at the `orc1.example.com` master site.
- The materialized view group at `mv2.example.com` is based on the objects in the `hr_repg` materialized view group at the `mv1.example.com` materialized view site.

Therefore, the examples in this chapter illustrate how to create a multitier materialized view environment, where one or more materialized views are based on other materialized views.

Complete the following steps to create these two materialized view groups.

**Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

```
/************************ BEGINNING OF SCRIPT ****************************
```

## Creating the Materialized View Group at mv1.example.com

Complete the following steps to create the hr_repg materialized view group at the mv1.example.com materialized view site. This materialized view group is based on the hr_repg master group at the orc1.example.com master site.

### Step 1 Create materialized view logs at the master site.

If you want one of your master sites to support a materialized view site, then you must create materialized view logs for each master table that is replicated to a materialized view. Recall from Figure 2–1 on page 2-2 that orc1.example.com serves as the target master site for the mv1.example.com materialized view site. The required materialized view logs must be created at orc1.example.com.

```
*/

SET ECHO ON

SPOOL create_mv_group.out

CONNECT hr@orc1.example.com

CREATE MATERIALIZED VIEW LOG ON hr.countries;
CREATE MATERIALIZED VIEW LOG ON hr.departments;
CREATE MATERIALIZED VIEW LOG ON hr.employees;
CREATE MATERIALIZED VIEW LOG ON hr.jobs;
CREATE MATERIALIZED VIEW LOG ON hr.job_history;
CREATE MATERIALIZED VIEW LOG ON hr.locations;
CREATE MATERIALIZED VIEW LOG ON hr.regions;

/*
```

**See Also:** The CREATE MATERIALIZED VIEW LOG statement in the *Oracle Database SQL Language Reference* for detailed information about this SQL statement

### Step 2 If they do not exist, then create the replicated schema and its database link.

Before building your materialized view group, you must ensure that the replicated schema exists at the remote materialized view site and that the necessary database links have been created.

In this example, if the hr schema does not exist, then create the schema. If the hr schema exists at the materialized view site, then grant any necessary privileges and go to the next task in this step.

```
*/

CONNECT system@mv1.example.com

CREATE TABLESPACE demo_mv1
```

```
 DATAFILE 'demo_mv1.dbf' SIZE 10M AUTOEXTEND ON
 EXTENT MANAGEMENT LOCAL AUTOALLOCATE;

CREATE TEMPORARY TABLESPACE temp_mv1
 TEMPFILE 'temp_mv1.dbf' SIZE 5M AUTOEXTEND ON;

ACCEPT password PROMPT 'Enter password for user: ' HIDE

CREATE USER hr IDENTIFIED BY &password;

ALTER USER hr DEFAULT TABLESPACE demo_mv1
              QUOTA UNLIMITED ON demo_mv1;

ALTER USER hr TEMPORARY TABLESPACE temp_mv1;

GRANT
  CREATE SESSION,
  CREATE TABLE,
  CREATE PROCEDURE,
  CREATE SEQUENCE,
  CREATE TRIGGER,
  CREATE VIEW,
  CREATE SYNONYM,
  ALTER SESSION,
  CREATE MATERIALIZED VIEW,
  ALTER ANY MATERIALIZED VIEW,
  CREATE DATABASE LINK
 TO hr;

/*
```

If it does not exist, then create the database link for the replicated schema.

Before building your materialized view group, you must ensure that the necessary database links exist for the replicated schema. The owner of the materialized views needs a database link pointing to the `proxy_refresher` that was created when the master site was set up.

```
*/

CONNECT hr@mv1.example.com

CREATE DATABASE LINK orc1.example.com
   CONNECT TO proxy_refresher IDENTIFIED BY &password;

/*
```

**Step 3 Create the materialized view group.**

The following procedures must be executed by the materialized view administrator at the remote materialized view site.

```
*/

CONNECT mviewadmin@mv1.example.com

/*
```

The master group that you specify in the `gname` parameter must match the name of the master group that you are replicating at the target master site.

```
*/
```

```
BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPGROUP (
      gname => 'hr_repg',
      master => 'orc1.example.com',
      propagation_mode => 'ASYNCHRONOUS');
END;
/

/*
```

### Step 4 Create the refresh group.

All materialized views that are added to a particular refresh group are refreshed at the same time. This ensures transactional consistency between the related materialized views in the refresh group.

```
*/

BEGIN
   DBMS_REFRESH.MAKE (
      name => 'mviewadmin.hr_refg',
      list => '',
      next_date => SYSDATE,
      interval => 'SYSDATE + 1/24',
      implicit_destroy => FALSE,
      rollback_seg => '',
      push_deferred_rpc => TRUE,
      refresh_after_errors => FALSE);
END;
/

/*
```

### Step 5 Add objects to the materialized view group.

Create the materialized views based on the master tables.

Whenever you create a materialized view, always specify the schema name of the table owner in the query for the materialized view. In the following examples, `hr` is specified as the owner of the table in each query.

```
*/

CREATE MATERIALIZED VIEW hr.countries_mv1
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
  AS SELECT * FROM hr.countries@orc1.example.com;

CREATE MATERIALIZED VIEW hr.departments_mv1
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
  AS SELECT * FROM hr.departments@orc1.example.com;

CREATE MATERIALIZED VIEW hr.employees_mv1
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
  AS SELECT * FROM hr.employees@orc1.example.com;

CREATE MATERIALIZED VIEW hr.jobs_mv1
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
  AS SELECT * FROM hr.jobs@orc1.example.com;

CREATE MATERIALIZED VIEW hr.job_history_mv1
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
```

```
  AS SELECT * FROM hr.job_history@orc1.example.com;

CREATE MATERIALIZED VIEW hr.locations_mv1
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
  AS SELECT * FROM hr.locations@orc1.example.com;

CREATE MATERIALIZED VIEW hr.regions_mv1
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
  AS SELECT * FROM hr.regions@orc1.example.com;

/*
```

Add the objects to the materialized view group.

```
*/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'countries_mv1',
      type => 'SNAPSHOT',
      min_communication => TRUE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'departments_mv1',
      type => 'SNAPSHOT',
      min_communication => TRUE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'employees_mv1',
      type => 'SNAPSHOT',
      min_communication => TRUE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'jobs_mv1',
      type => 'SNAPSHOT',
      min_communication => TRUE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'job_history_mv1',
```

```
      type => 'SNAPSHOT',
      min_communication => TRUE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'locations_mv1',
      type => 'SNAPSHOT',
      min_communication => TRUE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'regions_mv1',
      type => 'SNAPSHOT',
      min_communication => TRUE);
END;
/

/*
```

**Step 6 Add objects to the refresh group.**

All of the materialized view group objects that you add to the refresh group are refreshed at the same time to preserve referential integrity between related materialized views.

```
*/

BEGIN
   DBMS_REFRESH.ADD (
      name => 'mviewadmin.hr_refg',
      list => 'hr.countries_mv1',
      lax => TRUE);
END;
/

BEGIN
   DBMS_REFRESH.ADD (
      name => 'mviewadmin.hr_refg',
      list => 'hr.departments_mv1',
      lax => TRUE);
END;
/

BEGIN
   DBMS_REFRESH.ADD (
      name => 'mviewadmin.hr_refg',
      list => 'hr.employees_mv1',
      lax => TRUE);
END;
/

BEGIN
   DBMS_REFRESH.ADD (
```

```
      name => 'mviewadmin.hr_refg',
      list => 'hr.jobs_mv1',
      lax => TRUE);
END;
/

BEGIN
   DBMS_REFRESH.ADD (
      name => 'mviewadmin.hr_refg',
      list => 'hr.job_history_mv1',
      lax => TRUE);
END;
/

BEGIN
   DBMS_REFRESH.ADD (
      name => 'mviewadmin.hr_refg',
      list => 'hr.locations_mv1',
      lax => TRUE);
END;
/

BEGIN
   DBMS_REFRESH.ADD (
      name => 'mviewadmin.hr_refg',
      list => 'hr.regions_mv1',
      lax => TRUE);
END;
/

/*
```

## Creating the Materialized View Group at mv2.example.com

Complete the following steps to create the `hr_repg` materialized view group at the `mv2.example.com` materialized view site. This materialized view group is based on the `hr_repg` materialized view group at the `mv1.example.com` materialized view site.

### Step 1 Create materialized view logs at the master materialized view site.

If you want one of your master materialized view sites to support another materialized view site, then you must create materialized view logs for each materialized view that is replicated to another materialized view site. Recall from Figure 2–1 on page 2-2 that `mv1.example.com` serves as the target master internalized view site for the `mv2.example.com` materialized view site. The required materialized view logs must be created at `mv1.example.com`.

```
*/

CONNECT hr@mv1.example.com

CREATE MATERIALIZED VIEW LOG ON hr.countries_mv1;
CREATE MATERIALIZED VIEW LOG ON hr.departments_mv1;
CREATE MATERIALIZED VIEW LOG ON hr.employees_mv1;
CREATE MATERIALIZED VIEW LOG ON hr.jobs_mv1;
CREATE MATERIALIZED VIEW LOG ON hr.job_history_mv1;
CREATE MATERIALIZED VIEW LOG ON hr.locations_mv1;
CREATE MATERIALIZED VIEW LOG ON hr.regions_mv1;

/*
```

**See Also:** The CREATE MATERIALIZED VIEW LOG statement in the *Oracle Database SQL Language Reference* for detailed information about this SQL statement

**Step 2 If they do not exist, then create the replicated schema and its database link.**

Before building your materialized view group, you must ensure that the replicated schema exists at the remote materialized view site and that the necessary database links have been created.

For this example, if the hr schema does not exist, then create the schema. If the hr schema exists at the materialized view site, then go to the next task in this step.

```
*/

CONNECT system@mv2.example.com
CREATE TABLESPACE demo_mv2
 DATAFILE 'demo_mv2.dbf' SIZE 10M AUTOEXTEND ON
 EXTENT MANAGEMENT LOCAL AUTOALLOCATE;

CREATE TEMPORARY TABLESPACE temp_mv2
 TEMPFILE 'temp_mv2.dbf' SIZE 5M AUTOEXTEND ON;

ACCEPT password PROMPT 'Enter password for user: ' HIDE

CREATE USER hr IDENTIFIED BY &password;

ALTER USER hr DEFAULT TABLESPACE demo_mv2
              QUOTA UNLIMITED ON demo_mv2;

ALTER USER hr TEMPORARY TABLESPACE temp_mv2;

GRANT
  CREATE SESSION,
  CREATE TABLE,
  CREATE PROCEDURE,
  CREATE SEQUENCE,
  CREATE TRIGGER,
  CREATE VIEW,
  CREATE SYNONYM,
  ALTER SESSION,
  CREATE MATERIALIZED VIEW,
  ALTER ANY MATERIALIZED VIEW,
  CREATE DATABASE LINK
 TO hr;

/*
```

If it does not exist, then create the database link for the replicated schema.

Before building your materialized view group, you must ensure that the necessary database links exist for the replicated schema. The owner of the materialized views needs a database link pointing to the proxy_refresher that was created when the master materialized view site was set up.

```
*/

CONNECT hr@mv2.example.com

CREATE DATABASE LINK mv1.example.com
```

```
   CONNECT TO proxy_refresher IDENTIFIED BY &password;

/*
```

> **See Also:** Step 6 on page 2-12 for more information about creating proxy master materialized view site users

### Step 3 Create the materialized view group.

The following procedures must be executed by the materialized view administrator at the remote materialized view site.

```
*/

CONNECT mviewadmin@mv2.example.com

/*
```

The replication group that you specify in the `gname` parameter must match the name of the replication group that you are replicating at the target master materialized view site.

```
*/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPGROUP (
      gname => 'hr_repg',
      master => 'mv1.example.com',
      propagation_mode => 'ASYNCHRONOUS');
END;
/

/*
```

### Step 4 Create the refresh group.

All materialized views that are added to a particular refresh group are refreshed at the same time. This ensures transactional consistency between the related materialized views in the refresh group.

```
*/

BEGIN
   DBMS_REFRESH.MAKE (
      name => 'mviewadmin.hr_refg',
      list => '',
      next_date => SYSDATE,
      interval => 'SYSDATE + 1/24',
      implicit_destroy => FALSE,
      rollback_seg => '',
      push_deferred_rpc => TRUE,
      refresh_after_errors => FALSE);
END;
/

/*
```

### Step 5 Add objects to the materialized view group.

Create the materialized views based on the master materialized views.

Whenever you create a materialized view that is based on another materialized view, always specify the schema name of the materialized view owner in the query for the materialized view. In the following examples, `hr` is specified as the owner of the materialized view in each query.

```
*/

CREATE MATERIALIZED VIEW hr.countries_mv2
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
  AS SELECT * FROM hr.countries_mv1@mv1.example.com;

CREATE MATERIALIZED VIEW hr.departments_mv2
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
  AS SELECT * FROM hr.departments_mv1@mv1.example.com;

CREATE MATERIALIZED VIEW hr.employees_mv2
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
  AS SELECT * FROM hr.employees_mv1@mv1.example.com;

CREATE MATERIALIZED VIEW hr.jobs_mv2
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
  AS SELECT * FROM hr.jobs_mv1@mv1.example.com;

CREATE MATERIALIZED VIEW hr.job_history_mv2
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
  AS SELECT * FROM hr.job_history_mv1@mv1.example.com;

CREATE MATERIALIZED VIEW hr.locations_mv2
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
  AS SELECT * FROM hr.locations_mv1@mv1.example.com;

CREATE MATERIALIZED VIEW hr.regions_mv2
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
  AS SELECT * FROM hr.regions_mv1@mv1.example.com;

/*
```

Add the materialized views to the materialized view group.

```
*/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'countries_mv2',
      type => 'SNAPSHOT',
      min_communication => TRUE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'departments_mv2',
      type => 'SNAPSHOT',
      min_communication => TRUE);
END;
/
```

```
BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'employees_mv2',
      type => 'SNAPSHOT',
      min_communication => TRUE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'jobs_mv2',
      type => 'SNAPSHOT',
      min_communication => TRUE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'job_history_mv2',
      type => 'SNAPSHOT',
      min_communication => TRUE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'locations_mv2',
      type => 'SNAPSHOT',
      min_communication => TRUE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'regions_mv2',
      type => 'SNAPSHOT',
      min_communication => TRUE);
END;
/

/*
```

**Step 6 Add objects to the refresh group.**

All of the materialized view group objects that you add to the refresh group are refreshed at the same time to preserve referential integrity between related materialized views.

```
*/

BEGIN
```

```
   DBMS_REFRESH.ADD (
      name => 'mviewadmin.hr_refg',
      list => 'hr.countries_mv2',
      lax => TRUE);
END;
/

BEGIN
   DBMS_REFRESH.ADD (
      name => 'mviewadmin.hr_refg',
      list => 'hr.departments_mv2',
      lax => TRUE);
END;
/

BEGIN
   DBMS_REFRESH.ADD (
      name => 'mviewadmin.hr_refg',
      list => 'hr.employees_mv2',
      lax => TRUE);
END;
/

BEGIN
   DBMS_REFRESH.ADD (
      name => 'mviewadmin.hr_refg',
      list => 'hr.jobs_mv2',
      lax => TRUE);
END;
/

BEGIN
   DBMS_REFRESH.ADD (
      name => 'mviewadmin.hr_refg',
      list => 'hr.job_history_mv2',
      lax => TRUE);
END;
/

BEGIN
   DBMS_REFRESH.ADD (
      name => 'mviewadmin.hr_refg',
      list => 'hr.locations_mv2',
      lax => TRUE);
END;
/

BEGIN
   DBMS_REFRESH.ADD (
      name => 'mviewadmin.hr_refg',
      list => 'hr.regions_mv2',
      lax => TRUE);
END;
/

SET ECHO OFF

SPOOL OFF

/************************ END OF SCRIPT *********************************/
```

# 6

# Configuring Conflict Resolution

This chapter illustrates how to define conflict resolution methods for your replication environment.

This chapter contains these topics:


- Preparing for Conflict Resolution
- Creating Conflict Resolution Methods for Update Conflicts
- Creating Conflict Resolution Methods for Uniqueness Conflicts
- Creating Conflict Avoidance Methods for Delete Conflicts
- Using Dynamic Ownership Conflict Avoidance
- Auditing Successful Conflict Resolution


## Preparing for Conflict Resolution

Though you might design your database and front-end application to avoid conflicts between multiple sites in a replication environment, you might not be able to eliminate the possibility of conflicts. One of the most important aspects of replication is to ensure data convergence at all sites participating in the replication environment.

When data conflicts occur, you need a mechanism to ensure that the conflict is resolved in accordance with your business rules and that the data converges correctly at all sites.

Advanced Replication lets you define a conflict resolution system for your database that resolves conflicts in accordance with your business rules. If you have a unique situation that Oracle's prebuilt conflict resolution methods cannot resolve, then you have the option of building and using your own conflict resolution methods.

Before you begin implementing conflict resolution methods for your replicated tables, analyze the data in your system to determine where the most conflicts can occur. For example, static data such as an employee number might change very infrequently and is not subject to a high occurrence of conflicts. An employee's customer assignments, however, might change often and would therefore be prone to data conflicts.

After you have determined where the conflicts are most likely to occur, you must determine how to resolve the conflict. For example, do you want the latest change to have precedence, or should one site have precedence over another?

As you read each of the sections describing the different conflict resolution methods, you will learn what each method is best suited for. So, read each section and then think about how your business would want to resolve any potential conflicts.

After you have identified the potential problem areas and have determined what business rules would resolve the problem, use Oracle's conflict resolution methods (or one of your own) to implement a conflict resolution system.

> **See Also:** *Oracle Database Advanced Replication* for conceptual information about conflict resolution methods and detailed information about data convergence for each method

# Creating Conflict Resolution Methods for Update Conflicts

The most common data conflict occurs when the same row at two or more different sites are updated at nearly the same time, or before the deferred transaction from one site was successfully propagated to the other sites.

One method to avoid update conflicts is to implement a synchronous replication environment, though this solution requires large network resource.

The other solution is to use the Oracle conflict resolution methods to deal with update conflicts that can occur when the same row receives two or more updates.

## Overwrite and Discard Conflict Resolution Methods

The overwrite and discard methods ignore the values from either the originating or destination site and therefore can never guarantee convergence with more than one master site. These methods are designed to be used by a single master site and multiple materialized view sites, or with some form of a user-defined notification facility.

The overwrite method replaces the current value at the destination site with the new value from the originating site. Conversely, the discard method ignores the new value from the originating site.

> **See Also:** "ADD_conflicttype_RESOLUTION Procedure" on page 18-16 and *Oracle Database Advanced Replication* for more information about overwrite and discard

Complete the following steps to create an overwrite or discard conflict resolution method. This example illustrates the use of the discard conflict resolution method at the master site. Therefore, in the event of a conflict, the data from a materialized view site is discarded and the master site data remains.

> **Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

```
/************************ BEGINNING OF SCRIPT ******************************
```

### Step 1 Connect as the replication administrator.

The procedures in the following steps must be executed by the replication administrator.

```
*/

SET ECHO ON
```

```
SPOOL discard_conflictres.out

CONNECT repadmin@orc1.example.com

/*
```

**Step 2 Quiesce the master group that contains the table to which you want to apply the conflict resolution method.**

Before you define overwrite or discard conflict resolution methods, quiesce the master group that contains the table to which you want to apply the conflict resolution method. In a single master replication environment, quiescing the master group might not be required.

```
*/

BEGIN
   DBMS_REPCAT.SUSPEND_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

/*
```

**Step 3 Create a column group for your target table.**

All Oracle conflict resolution methods are based on logical column groupings called column groups.

```
*/

BEGIN
   DBMS_REPCAT.MAKE_COLUMN_GROUP (
      sname => 'hr',
      oname => 'departments',
      column_group => 'dep_cg',
      list_of_column_names => 'manager_id,location_id');
END;
/

/*
```

**Step 4 Define the conflict resolution method for a specified table.**

This example creates an `OVERWRITE` conflict resolution method.

```
*/

BEGIN
   DBMS_REPCAT.ADD_UPDATE_RESOLUTION (
      sname => 'hr',
      oname => 'departments',
      column_group => 'dep_cg',
      sequence_no => 1,
      method => 'DISCARD',
      parameter_column_name => 'manager_id,location_id');
END;
/

/*
```

**Step 5 Regenerate replication support for the table that received the conflict resolution method.**

```
*/

BEGIN
    DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT (
      sname => 'hr',
      oname => 'departments',
      type => 'TABLE',
      min_communication => TRUE);
END;
/

/*
```

**Step 6 Resume master activity after replication support has been regenerated.**

```
*/

BEGIN
   DBMS_REPCAT.RESUME_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

SET ECHO OFF

SPOOL OFF

/************************ END OF SCRIPT *********************************/
```

## Minimum and Maximum Conflict Resolution Methods

When Advanced Replication detects a conflict with a column group and calls either the minimum or maximum value conflict resolution methods, it compares the new value from the originating site with the current value from the destination site for a designated column in the column group. You must designate this column when you define your conflict resolution method.

If the new value of the designated column is *less than* or *greater than* (depending on the method used) the current value, then the column group values from the originating site are applied at the destination site, if all other errors were successfully resolved for the row. Otherwise the rows remain unchanged.

Complete the following steps to create an maximum or minimum conflict resolution method.

> **Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

```
/************************ BEGINNING OF SCRIPT ****************************
```

**Step 1 Connect as the replication administrator.**

The procedures in the following steps must be executed by the replication administrator.

```
*/

SET ECHO ON

SPOOL min_conflictres.out

CONNECT repadmin@orc1.example.com

/*
```

**Step 2 Quiesce the master group that contains the table to which you want to apply the conflict resolution method.**

Before you define maximum or minimum conflict resolution methods, quiesce the master group that contains the table to which you want to apply the conflict resolution method. In a single master replication environment, quiescing the master group might not be required.

```
*/

BEGIN
   DBMS_REPCAT.SUSPEND_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

/*
```

**Step 3 Create a column group for your target table.**

All Oracle conflict resolution methods are based on logical column groupings called column groups.

```
*/

BEGIN
   DBMS_REPCAT.MAKE_COLUMN_GROUP (
      sname => 'hr',
      oname => 'jobs',
      column_group => 'job_minsal_cg',
      list_of_column_names => 'min_salary');
END;
/

/*
```

**Step 4 Define the conflict resolution method for a specified table.**

This example creates a `MINIMUM` conflict resolution method.

```
*/

BEGIN
   DBMS_REPCAT.ADD_UPDATE_RESOLUTION (
      sname => 'hr',
      oname => 'jobs',
      column_group => 'job_minsal_cg',
      sequence_no => 1,
      method => 'MINIMUM',
      parameter_column_name => 'min_salary');
END;
/
```

```
/*
```

**Step 5 Regenerate replication support for the table that received the conflict resolution method.**

```
*/

BEGIN
    DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT (
      sname => 'hr',
      oname => 'jobs',
      type => 'TABLE',
      min_communication => TRUE);
END;
/

/*
```

**Step 6 Resume replication activity.**

```
*/

BEGIN
   DBMS_REPCAT.RESUME_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

SET ECHO OFF

SPOOL OFF

/************************ END OF SCRIPT *********************************/
```

## Timestamp Conflict Resolution Methods

The earliest time stamp and latest time stamp methods are variations on the minimum and maximum value methods. To use the time stamp method, you must designate a column in the replicated table of type `DATE`. When an application updates any column in a column group, the application must also update the value of the designated time stamp column with the local `SYSDATE`. For a change applied from another site, the time stamp value should be set to the time stamp value from the originating site.

Two elements are needed to make time stamp conflict resolution work well:

- Synchronized time settings between computers
- Timestamp field and trigger to automatically record time stamp

Complete the following steps to create a time stamp conflict resolution method.

> **Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

```
/************************ BEGINNING OF SCRIPT ****************************
```

**Step 1 Connect as the replication administrator.**

The procedures in the following steps must be executed by the replication administrator.

```
*/

SET ECHO ON

SPOOL timestamp_conflictres.out

CONNECT repadmin@orc1.example.com

/*
```

**Step 2 Quiesce the master group that contains the table to which you want to apply the conflict resolution method.**

Before defining time stamp conflict resolution methods, quiesce the master group that contains the table to which you want to apply the conflict resolution method. In a single master replication environment, quiescing the master group might not be required.

```
*/

BEGIN
   DBMS_REPCAT.SUSPEND_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

/*
```

**Step 3 Add a column to your table to record the timestamp value when a row is inserted or updated.**

If the target table does not already contain a time stamp field, then add a column to your table to record the time stamp value when a row is inserted or updated. You must use the `ALTER_MASTER_REPOBJECT` procedure to apply the DDL to the target table. Simply issuing the DDL might cause the replicated object to become invalid.

```
*/

BEGIN
   DBMS_REPCAT.ALTER_MASTER_REPOBJECT (
      sname => 'hr',
      oname => 'countries',
      type => 'TABLE',
      ddl_text => 'ALTER TABLE hr.countries ADD (timestamp DATE)');
END;
/

/*
```

**Step 4 Regenerate replication support for the altered table.**

```
*/

BEGIN
    DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT (
      sname => 'hr',
      oname => 'countries',
      type => 'TABLE',
```

```
      min_communication => TRUE);
END;
/

/*
```

### Step 5 Create a trigger that records the timestamp when a row is either inserted or updated.

This recorded value is used in the resolution of conflicts based on the Timestamp method. Instead of directly executing the DDL, use the `DBMS_REPCAT.CREATE_MASTER_REPOBJECT` procedure to create the trigger and add it to your master group.

> **Note:** You cannot use columns of `datetime` and `interval` data types for priority group conflict resolution.

```
*/

BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
      gname => 'hr_repg',
      type => 'TRIGGER',
      oname => 'insert_time',
      sname => 'hr',
      ddl_text => 'CREATE TRIGGER hr.insert_time
                      BEFORE
                         INSERT OR UPDATE ON hr.countries FOR EACH ROW
                      BEGIN
                         IF DBMS_REPUTIL.FROM_REMOTE = FALSE THEN
                            :NEW.TIMESTAMP := SYSDATE;
                         END IF;
                      END;');
END;
/

/*
```

### Step 6 Create a column group for your target table.

All Oracle conflict resolution methods are based on logical column groupings called column groups.

```
*/

BEGIN
   DBMS_REPCAT.MAKE_COLUMN_GROUP (
      sname => 'hr',
      oname => 'countries',
      column_group => 'countries_timestamp_cg',
      list_of_column_names => 'country_name,region_id,timestamp');
END;
/

/*
```

### Step 7 Define the conflict resolution method for a specified table.

This example specifies the `LATEST TIMESTAMP` conflict resolution method using the `timestamp` column that you created earlier.

```
*/

BEGIN
   DBMS_REPCAT.ADD_UPDATE_RESOLUTION (
      sname => 'hr',
      oname => 'countries',
      column_group => 'countries_timestamp_cg',
      sequence_no => 1,
      method => 'LATEST TIMESTAMP',
      parameter_column_name => 'timestamp');
END;
/

/*
```

**Step 8 Regenerate replication support for the table that received the conflict resolution method.**

```
*/

BEGIN
    DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT (
      sname => 'hr',
      oname => 'countries',
      type => 'TABLE',
      min_communication => TRUE);
END;
/

/*
```

**Step 9 Resume replication activity.**

```
*/

BEGIN
   DBMS_REPCAT.RESUME_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

SET ECHO OFF

SPOOL OFF

/************************ END OF SCRIPT *********************************/
```

## Additive and Average Conflict Resolution Methods

The additive and average methods work with column groups consisting of a single numeric column only. Instead of "accepting" one value over another, this conflict resolution method either adds the two compared values together or takes an average of the two compared values.

Complete the following steps to create an additive or average conflict resolution method. This example averages the commission percentage for an employee in the event of a conflict.

**Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

```
/************************* BEGINNING OF SCRIPT ******************************
```

### Step 1 Connect as the replication administrator.

The procedures in the following steps must be executed by the replication administrator.

```
*/

SET ECHO ON

SPOOL average_conflictres.out

CONNECT repadmin@orc1.example.com

/*
```

### Step 2 Quiesce the master group that contains the table to which you want to apply the conflict resolution method.

Before you define additive and average conflict resolution methods, quiesce the master group that contains the table to which you want to apply the conflict resolution method. In a single master replication environment, quiescing the master group might not be required.

```
*/

BEGIN
   DBMS_REPCAT.SUSPEND_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

/*
```

### Step 3 Create a column group for your target table.

All Oracle conflict resolution methods are based on logical column groupings called column groups.

```
*/

BEGIN
   DBMS_REPCAT.MAKE_COLUMN_GROUP (
      sname => 'hr',
      oname => 'employees',
      column_group => 'commission_average_cg',
      list_of_column_names => 'commission_pct');
END;
/

/*
```

**Step 4 Define the conflict resolution method for a specified table.**

This example specifies the `AVERAGE` conflict resolution method using the `sal` column.

```
*/

BEGIN
   DBMS_REPCAT.ADD_UPDATE_RESOLUTION (
      sname => 'hr',
      oname => 'employees',
      column_group => 'commission_average_cg',
      sequence_no => 1,
      method => 'AVERAGE',
      parameter_column_name => 'commission_pct');
END;
/

/*
```

**Step 5 Regenerate replication support for the table that received the conflict resolution method.**

```
*/

BEGIN
    DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT (
      sname => 'hr',
      oname => 'employees',
      type => 'TABLE',
      min_communication => TRUE);
END;
/

/*
```

**Step 6 Resume replication activity.**

```
*/

BEGIN
   DBMS_REPCAT.RESUME_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

SET ECHO OFF

SPOOL OFF

/************************ END OF SCRIPT *********************************/
```

## Priority Groups Conflict Resolution Methods

Priority groups allow you to assign a priority level to each possible value of a particular column. If Oracle detects a conflict, then Oracle updates the table whose "priority" column has a lower value using the data from the table with the higher priority value.

Complete the following steps to create a priority groups conflict resolution method.

**Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

```
/************************* BEGINNING OF SCRIPT ******************************
```

### Step 1 Connect as the replication administrator.

The procedures in the following steps must be executed by the replication administrator.

```
*/

SET ECHO ON

SPOOL priority_groups_conflictres.out

CONNECT repadmin@orc1.example.com

/*
```

### Step 2 Quiesce the master group that contains the table to which you want to apply the conflict resolution method.

Before you define a priority groups conflict resolution method, quiesce the master group that contains the table to which you want to apply the conflict resolution method. In a single master replication environment, quiescing the master group might not be required.

```
*/

BEGIN
   DBMS_REPCAT.SUSPEND_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

/*
```

### Step 3 Ensure that the job_id column is part of the column group for which your site priority conflict resolution mechanism is used.

Use the `ADD_GROUPED_COLUMN` procedure to add this column to an existing column group. If you do not already have a column group, then you can create a new column group using the `DBMS_REPCAT.MAKE_COLUMN_GROUP` procedure.

```
*/

BEGIN
   DBMS_REPCAT.MAKE_COLUMN_GROUP (
      sname => 'hr',
      oname => 'employees',
      column_group => 'employees_priority_cg',
      list_of_column_names => 'manager_id,hire_date,salary,job_id');
END;
/

/*
```

**Step 4 Before you begin assigning a priority value to the values in your table, create a priority group that holds the values you defined.**

```
*/

BEGIN
   DBMS_REPCAT.DEFINE_PRIORITY_GROUP (
      gname => 'hr_repg',
      pgroup => 'job_pg',
      datatype => 'VARCHAR2');
END;
/

/*
```

**Step 5 Define a priority value for all possible table values.**

The DBMS_REPCAT.ADD_PRIORITY_*datatype* procedure is available in several different versions. There is a version for each available data type (NUMBER, VARCHAR2, and so on). Execute this procedure as often as necessary until you have defined a priority value for all possible table values.

> **See Also:** "ADD_PRIORITY_datatype Procedure" on page 18-13 for more information

```
*/

BEGIN
   DBMS_REPCAT.ADD_PRIORITY_VARCHAR2(
      gname => 'hr_repg',
      pgroup => 'job_pg',
      value => 'ad_pres',
      priority => 100);
END;
/

BEGIN
   DBMS_REPCAT.ADD_PRIORITY_VARCHAR2(
      gname => 'hr_repg',
      pgroup => 'job_pg',
      value => 'sa_man',
      priority => 80);
END;
/

BEGIN
   DBMS_REPCAT.ADD_PRIORITY_VARCHAR2(
      gname => 'hr_repg',
      pgroup => 'job_pg',
      value => 'sa_rep',
      priority => 60);
END;
/

BEGIN
   DBMS_REPCAT.ADD_PRIORITY_VARCHAR2(
      gname => 'hr_repg',
      pgroup => 'job_pg',
      value => 'pu_clerk',
      priority => 40);
END;
```

```
/

BEGIN
   DBMS_REPCAT.ADD_PRIORITY_VARCHAR2(
      gname => 'hr_repg',
      pgroup => 'job_pg',
      value => 'st_clerk',
      priority => 20);
END;
/

/*
```

**Step 6 Add the PRIORITY GROUP resolution method to your replicated table.**

The following example shows that it is the second conflict resolution method for the specified column group (sequence_no parameter).

```
*/

BEGIN
   DBMS_REPCAT.ADD_UPDATE_RESOLUTION (
      sname => 'hr',
      oname => 'employees',
      column_group => 'employees_priority_cg',
      sequence_no => 2,
      method => 'PRIORITY GROUP',
      parameter_column_name => 'job_id',
      priority_group => 'job_pg');
END;
/

/*
```

**Step 7 Regenerate replication support for the table that received the conflict resolution method.**

```
*/

BEGIN
    DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT (
      sname => 'hr',
      oname => 'employees',
      type => 'TABLE',
      min_communication => TRUE);
END;
/

/*
```

**Step 8 Resume replication activity.**

```
*/

BEGIN
   DBMS_REPCAT.RESUME_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

SET ECHO OFF
```

```
SPOOL OFF

/************************ END OF SCRIPT *******************************/
```

## Site Priority Conflict Resolution Methods

Site priority is a specialized form of a priority group. Therefore, many of the procedures associated with site priority behave similarly to the procedures associated with priority groups. Instead of resolving a conflict based on the priority of a field's value, the conflict is resolved based on the priority of the sites involved.

For example, if you assign `orc2.example.com` a higher priority value than `orc1.example.com` and a conflict arises between these two sites, then the value from `orc2.example.com` is used.

Complete the following steps to create a site priority conflict resolution method.

> **Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

```
/************************ BEGINNING OF SCRIPT ****************************
```

### Step 1 Connect as the replication administrator.

The procedures in the following steps must be executed by the replication administrator.

```
*/

SET ECHO ON

SPOOL site_priority_conflictres.out

CONNECT repadmin@orc1.example.com

/*
```

### Step 2 Quiesce the master group that contains the table to which you want to apply the conflict resolution method.

Before you define a site priority conflict resolution method, quiesce the master group that contains the table to which you want to apply the conflict resolution method. In a single master replication environment, quiescing the master group might not be required.

```
*/

BEGIN
   DBMS_REPCAT.SUSPEND_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

/*
```

**Step 3 Add a site column to your table to store the site value.**

Use the DBMS_REPCAT.ALTER_MASTER_REPOBJECT procedure to apply the DDL to the target table. Simply issuing the DDL might cause the replicated object to become invalid.

```
*/

BEGIN
   DBMS_REPCAT.ALTER_MASTER_REPOBJECT (
      sname => 'hr',
      oname => 'regions',
      type => 'TABLE',
      ddl_text => 'ALTER TABLE hr.regions ADD (site VARCHAR2(20))');
END;
/

/*
```

**Step 4 Regenerate replication support for the affected object.**

```
*/

BEGIN
    DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT (
      sname => 'hr',
      oname => 'regions',
      type => 'TABLE',
      min_communication => TRUE);
END;
/

/*
```

**Step 5 Create a trigger that records the global name of the site when a row is either inserted or updated.**

This recorded value is used in the resolution of conflicts based on the site priority method. Instead of directly executing the DDL, use the DBMS_REPCAT.CREATE_MASTER_REPOBJECT procedure to create the trigger and add it to your master group.

```
*/

BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
      gname => 'hr_repg',
      type => 'TRIGGER',
      oname => 'insert_site',
      sname => 'hr',
      ddl_text => 'CREATE TRIGGER hr.insert_site
                     BEFORE
                       INSERT OR UPDATE ON hr.regions FOR EACH ROW
                     BEGIN
                       IF DBMS_REPUTIL.FROM_REMOTE = FALSE THEN
                         SELECT global_name INTO :NEW.SITE FROM GLOBAL_NAME;
                       END IF;
                     END;');
END;
/

/*
```

**Step 6 Ensure that the new column is part of the column group for which your site priority conflict resolution mechanism is used.**

Use the ADD_GROUPED_COLUMN procedure to add this column to an existing column group. If you do not already have a column group, then you can create a new column group using the DBMS_REPCAT.MAKE_COLUMN_GROUP procedure.

```
*/

BEGIN
   DBMS_REPCAT.MAKE_COLUMN_GROUP (
      sname => 'hr',
      oname => 'regions',
      column_group => 'regions_sitepriority_cg',
      list_of_column_names => 'region_id,region_name,site');
END;
/

/*
```

**Step 7 Before assigning a site priority value to the sites in your replicated environment, create a site priority group that holds the values you defined.**

```
*/

BEGIN
   DBMS_REPCAT.DEFINE_SITE_PRIORITY (
      gname => 'hr_repg',
      name => 'regions_sitepriority_pg');
END;
/

/*
```

**Step 8 Define the priority value for each of the sites in your replication environment.**

Execute this procedure as often as necessary until you have defined a site priority value for each of the sites in our replication environment.

```
*/

BEGIN
   DBMS_REPCAT.ADD_SITE_PRIORITY_SITE (
      gname => 'hr_repg',
      name => 'regions_sitepriority_pg',
      site => 'orc1.example.com',
      priority => 100);
END;
/

BEGIN
   DBMS_REPCAT.ADD_SITE_PRIORITY_SITE (
      gname => 'hr_repg',
      name => 'regions_sitepriority_pg',
      site => 'orc2.example.com',
      priority => 50);
END;
/

BEGIN
   DBMS_REPCAT.ADD_SITE_PRIORITY_SITE (
```

```
      gname => 'hr_repg',
      name => 'regions_sitepriority_pg',
      site => 'orc3.example.com',
      priority => 25);
END;
/

/*
```

**Step 9 Add the SITE PRIORITY resolution method to your replicated table.**

The following example shows that it is the third conflict resolution method for the specified column group (`sequence_no` parameter).

```
*/

BEGIN
   DBMS_REPCAT.ADD_UPDATE_RESOLUTION (
      sname => 'hr',
      oname => 'regions',
      column_group => 'regions_sitepriority_cg',
      sequence_no => 1,
      method => 'SITE PRIORITY',
      parameter_column_name => 'site',
      priority_group => 'regions_sitepriority_pg');
END;
/

/*
```

**Step 10 Regenerate replication support for the table that received the conflict resolution method.**

```
*/

BEGIN
    DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT (
      sname => 'hr',
      oname => 'regions',
      type => 'TABLE',
      min_communication => TRUE);
END;
/

/*
```

**Step 11 Resume replication activity.**

```
*/

BEGIN
   DBMS_REPCAT.RESUME_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

SET ECHO OFF

SPOOL OFF

/************************ END OF SCRIPT *********************************/
```

# Creating Conflict Resolution Methods for Uniqueness Conflicts

In a replication environment, you might have situations where you encounter a conflict on a unique constraint, often resulting from an insert. If your business rules allow you to delete the duplicate row, then you can define a resolution method with Oracle's prebuilt conflict resolution methods.

More often, however, you probably want to modify the conflicting value so that it no longer violates the unique constraint. Modifying the conflicting value ensures that you do not lose important data. Oracle's prebuilt uniqueness conflict resolution method can make the conflicting value unique by appending a site name or a sequence number to the value.

An additional component that accompanies the uniqueness conflict resolution method is a notification facility. The conflicting information is modified by Oracle so that it can be inserted into the table, but you should be notified so that you can analyze the conflict to determine whether the record should be deleted, or the data merged into another record, or a completely new value be defined for the conflicting data.

A uniqueness conflict resolution method detects and resolves conflicts encountered on columns with a `UNIQUE` constraint. The example in this section uses the `employees` table in the `hr` sample schema, which has the unique constraint `emp_email_uk` on the `email` column.

> **Note:** To add unique conflict resolution method for a column, the name of the unique index on the column must match the name of the unique or primary key constraint.

Complete the following steps to create a uniqueness conflict resolution method.

> **Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

```
/************************ BEGINNING OF SCRIPT ******************************
```

**Step 1 Connect as the replication administrator.**

```
*/

SET ECHO ON

SPOOL unique_conflictres.out

CONNECT repadmin@orc1.example.com

/*
```

**Step 2 Quiesce the master group that contains the table to which you want to apply the conflict resolution method.**

Before you define a uniqueness conflict resolution method, ensure that the master group that contains the table to which you want to apply the conflict resolution method is quiesced.

```
*/
```

```
BEGIN
   DBMS_REPCAT.SUSPEND_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

/*
```

**Step 3 Create a table that stores the messages received from your notification facility.**

In this example, the table name is `conf_report`.

```
*/

BEGIN
   DBMS_REPCAT.EXECUTE_DDL (
      gname => 'hr_repg',
      ddl_text => 'CREATE TABLE hr.conf_report (
                   line NUMBER(2),
                   txt VARCHAR2(80),
                   timestamp DATE,
                   table_name VARCHAR2(30),
                   table_owner VARCHAR2(30),
                   conflict_type VARCHAR2(7))');
END;
/

/*
```

**Step 4 Connect as the owner of the table you created in Step 3.**

```
*/

CONNECT hr@orc1.example.com

/*
```

**Step 5 Create a package that sends a notification to the conf_report table when a conflict is detected.**

In this example, the package name is `notify`.

> **See Also:** Appendix B, "User-Defined Conflict Resolution Methods" describes the conflict resolution notification package that is created in this script

```
*/

CREATE OR REPLACE PACKAGE notify AS
   FUNCTION emp_unique_violation (email IN OUT VARCHAR2,
      discard_new_values IN OUT BOOLEAN)
   RETURN BOOLEAN;
END notify;
/

CREATE OR REPLACE PACKAGE BODY notify AS
   TYPE message_table IS TABLE OF VARCHAR2(80) INDEX BY BINARY_INTEGER;
   PROCEDURE report_conflict(conflict_report IN MESSAGE_TABLE,
      report_length IN NUMBER,
```

```
      conflict_time IN DATE,
      conflict_table IN VARCHAR2,
      table_owner IN VARCHAR2,
      conflict_type IN VARCHAR2) IS
      BEGIN
         FOR idx IN 1..report_length LOOP
            BEGIN
              INSERT INTO hr.conf_report
              (line, txt, timestamp, table_name, table_owner, conflict_type)
              VALUES (idx, SUBSTR(conflict_report(idx),1,80), conflict_time,
              conflict_table, table_owner, conflict_type);
            EXCEPTION WHEN others THEN NULL;
            END;
         END LOOP;
      END report_conflict;
   FUNCTION emp_unique_violation(email IN OUT VARCHAR2,
      discard_new_values IN OUT BOOLEAN)
   RETURN BOOLEAN IS
      local_node VARCHAR2(128);
      conf_report MESSAGE_TABLE;
      conf_time DATE := SYSDATE;
   BEGIN
      BEGIN
         SELECT global_name INTO local_node FROM global_name;
      EXCEPTION WHEN others THEN local_node := '?';
      END;
      conf_report(1) := 'UNIQUENESS CONFLICT DETECTED IN EMPLOYEES ON ' ||
         TO_CHAR(conf_time, 'MM-DD-YYYY HH24:MI:SS');
      conf_report(2) := ' AT NODE ' || local_node;
      conf_report(3) := 'ATTEMPTING TO RESOLVE CONFLICT USING' ||
         ' APPEND SITE NAME METHOD';
      conf_report(4) := 'EMAIL: ' || email;
      conf_report(5) := NULL;
      report_conflict(conf_report,5,conf_time,'employees','hr','UNIQUE');
      discard_new_values := FALSE;
      RETURN FALSE;
   END emp_unique_violation;
END notify;
/

/*
```

**Step 6 Connect as the replication administrator.**

```
*/

CONNECT repadmin@orc1.example.com

/*
```

**Step 7 Replicate the package you created in Step 5 to all of the master sites in your replication environment.**

This step ensures that the notification facility is available at all master sites.

```
*/

BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
      gname => 'hr_repg',
      type =>  'PACKAGE',
```

```
      oname => 'notify',
      sname => 'hr');
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
      gname => 'hr_repg',
      type =>  'PACKAGE BODY',
      oname => 'notify',
      sname => 'hr');
END;
/

/*
```

**Step 8 Add the notification facility as one of your conflict resolution methods.**

Add it even though it only notifies of a conflict. The following example demonstrates adding the notification facility as a `USER FUNCTION`.

```
*/

BEGIN
   DBMS_REPCAT.ADD_UPDATE_RESOLUTION (
      sname => 'hr',
      oname => 'employees',
      constraint_name => 'emp_email_uk',
      sequence_no => 1,
      method => 'USER FUNCTION',
      comment => 'Notify DBA',
      parameter_column_name => 'email',
      function_name => 'hr.notify.emp_unique_violation');
END;
/

/*
```

**Step 9 Add the actual conflict resolution method to your table.**

The following example demonstrates adding the `APPEND SITE NAME` uniqueness conflict resolution method to your replicated table.

```
*/

BEGIN
   DBMS_REPCAT.ADD_UPDATE_RESOLUTION (
      sname => 'hr',
      oname => 'employees',
      constraint_name => 'emp_email_uk',
      sequence_no => 2,
      method => 'APPEND SITE NAME',
      parameter_column_name => 'email');
END;
/

/*
```

**Step 10 Regenerate replication support for the table that received the conflict resolution methods.**

```
*/
```

```
BEGIN
    DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT (
      sname => 'hr',
      oname => 'employees',
      type => 'TABLE',
      min_communication => TRUE);
END;
/

/*
```

**Step 11 Resume replication activity.**

```
*/

BEGIN
   DBMS_REPCAT.RESUME_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

SET ECHO OFF

SPOOL OFF

/************************ END OF SCRIPT *********************************/
```

# Creating Conflict Avoidance Methods for Delete Conflicts

Unlike update conflicts, where there are two values to compare, simply deleting a row makes the update conflict resolution methods described in the previous section ineffective because only one value would exist.

The best way to deal with deleting rows in a replication environment is to avoid the conflict by marking a row for deletion and periodically purging the table of all marked records. Because you are not physically removing this row, your data can converge at all master sites if a conflict arises because you still have two values to compare, if no other errors have occurred. After you are sure that your data has converged, you can purge marked rows using a replicated purge procedure.

When developing the front-end application for your database, you probably want to filter out the rows that have been marked for deletion, because doing so makes it appear to your users as though the row was physically deleted. Simply exclude the rows that have been marked for deletion in the `SELECT` statement for your data set.

For example, a select statement for a current employee listing might be similar to the following:

```
SELECT * FROM hr.locations WHERE remove_date IS NULL;
```

This section describes how to prepare your replicated table to avoid delete conflicts. You also learn how to use procedural replication to purge those records that have been marked for deletion.

Complete the following steps to create a conflict avoidance method for delete conflicts.

**Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

```
/************************* BEGINNING OF SCRIPT ******************************
```

**Step 1 Connect as the replication administrator at the master definition site.**

```
*/

SET ECHO ON

SPOOL delete_conflictres.out

CONNECT repadmin@orc1.example.com

/*
```

**Step 2 Quiesce the master group that contains the table to which you want to apply the conflict resolution method.**

```
*/

BEGIN
   DBMS_REPCAT.SUSPEND_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

/*
```

**Step 3 Add a column to the replicated table that stores the mark for deleted records.**

It is advisable to use a time stamp to mark your records for deletion (time stamp reflects when the record was marked for deletion). Because you are using a time stamp, the new column can be a `DATE` data type. Use the `DBMS_REPCAT.ALTER_MASTER_REPOBJECT` procedure to add the `remove_date` column to your existing replicated table.

```
*/

BEGIN
   DBMS_REPCAT.ALTER_MASTER_REPOBJECT (
      sname => 'hr',
      oname => 'locations',
      type => 'TABLE',
      ddl_text => 'ALTER TABLE hr.locations ADD (remove_date DATE)');
END;
/

/*
```

**Step 4 Regenerate replication support for the altered table.**

```
*/

BEGIN
    DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT (
      sname => 'hr',
```

```
      oname => 'locations',
      type => 'TABLE',
      min_communication => TRUE);
END;
/

/*
```

**Step 5 Create a package that is replicated to all of the master sites in your replication environment.**

This package purges all marked records from the specified table.

```
*/

BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
      gname => 'hr_repg',
      type => 'PACKAGE',
      oname => 'purge',
      sname => 'hr',
      ddl_text => 'CREATE OR REPLACE PACKAGE hr.purge AS
                      PROCEDURE remove_locations(purge_date DATE);
                   END;');
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
      gname => 'hr_repg',
      type => 'PACKAGE BODY',
      oname => 'purge',
      sname => 'hr',
      ddl_text => 'CREATE OR REPLACE PACKAGE BODY hr.purge AS
                      PROCEDURE remove_locations(purge_date IN DATE) IS
                      BEGIN
                         DBMS_REPUTIL.REPLICATION_OFF;
                         LOCK TABLE hr.locations IN EXCLUSIVE MODE;
                         DELETE hr.locations WHERE remove_date IS NOT NULL
                            AND remove_date < purge_date;
                         DBMS_REPUTIL.REPLICATION_ON;
                      EXCEPTION WHEN others THEN
                         DBMS_REPUTIL.REPLICATION_ON;
                      END;
                   END;');
END;
/

/*
```

**Step 6 Generate replication support for each package and package body.**

After generating replication support, a synonym is created for you and added to your master group as a replicated object. This synonym is labeled as `defer_purge.remove_locations`.

```
*/

BEGIN
    DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT (
      sname => 'hr',
      oname => 'purge',
```

```
      type => 'PACKAGE',
      min_communication => TRUE);
END;
/

BEGIN
    DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT (
      sname => 'hr',
      oname => 'purge',
      type => 'PACKAGE BODY',
      min_communication => TRUE);
END;
/

/*
```

**Step 7 In a separate terminal window, manually push any administrative requests at all other master sites.**

You might need to execute the `DO_DEFERRED_REPCAT_ADMIN` procedure in the `DBMS_REPCAT` package several times, because some administrative operations have multiple steps. The following is an example:

```
*/

BEGIN
   DBMS_REPCAT.DO_DEFERRED_REPCAT_ADMIN (
      gname => 'hr_repg',
      all_sites => FALSE);
END;
/

*/

PAUSE Press <RETURN> to continue when you have verified that there are no
pending administrative requests in the DBA_REPCATLOG data dictionary view.

/*
```

**Step 8 Resume replication activity.**

```
*/

BEGIN
   DBMS_REPCAT.RESUME_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

SET ECHO OFF

SPOOL OFF

/************************ END OF SCRIPT *********************************/
```

# Using Dynamic Ownership Conflict Avoidance

This section describes a more advanced method of designing your applications to avoid conflicts. This method, known as **token passing,** is similar to the workflow method described in the following sections. Although this section describes how to

use this method to control the ownership of an entire row, you can use a modified form of this method to control ownership of the individual column groups within a row.

Both workflow and token passing allow dynamic ownership of data. With dynamic ownership, only one site at a time is allowed to update a row, but ownership of the row can be passed from site to site. Both workflow and token passing use the value of one or more "identifier" columns to determine who is currently allowed to update the row.

## Workflow

With workflow partitioning, you can think of data ownership as being "pushed" from site to site. Only the current owner of the row is allowed to push the ownership of the row to another site, by changing the value of the "identifier" columns.

Take the simple example of separate sites for ordering, shipping, and billing. Here, the identifier columns are used to indicate the status of an order. The status determines which site can update the row. After a user at the ordering site has entered the order, the user updates the status of this row to `ship`. Users at the ordering site are no longer allowed to modify this row — ownership has been pushed to the shipping site.

After shipping the order, the user at the shipping site updates the status of this row to `bill`, thus pushing ownership to the billing site, and so on.

To successfully avoid conflicts, applications implementing dynamic data ownership must ensure that the following conditions are met:

- Only the owner of the row can update the row.
- The row is never owned by more than one site.
- Ordering conflicts can be successfully resolved at all sites.

With workflow partitioning, only the current owner of the row can push the ownership of the row to the next site by updating the "identifier" columns. No site is given ownership unless another site has given up ownership; thus ensuring there is never more than one owner.

Because the flow of work is ordered, ordering conflicts can be resolved by applying the change from the site that occurs latest in the flow of work. Any ordering conflicts can be resolved using a form of the priority conflict resolution method, where the priority value increases with each step in the work flow process. The priority conflict resolution method successfully converges for more than one master site as long as the priority value is always increasing.

## Token Passing

Token passing uses a more generalized approach to meeting these criteria. To implement token passing, instead of the "identifier" columns, your replicated tables must have owner and epoch columns. The owner column stores the global database name of the site currently believed to own the row.

Once you have designed a token passing mechanism, you can use it to implement a variety of forms of dynamic partitioning of data ownership, including workflow.

You should design your application to implement token passing for you automatically. You should not allow the owner or epoch columns to be updated outside this application.

Whenever you attempt to update a row, your application should:

1. Locate the current owner of the row.
2. Establish ownership of the row.
3. Lock the row to prevent updates while ownership is changing.
4. Perform the update.

Oracle releases the lock when you commit your transaction.

For example, Figure 6–1 illustrates how ownership of employee `100` passes from the `acct_sf` database to the `acct_ny` database.

***Figure 6–1 Grabbing the Token***

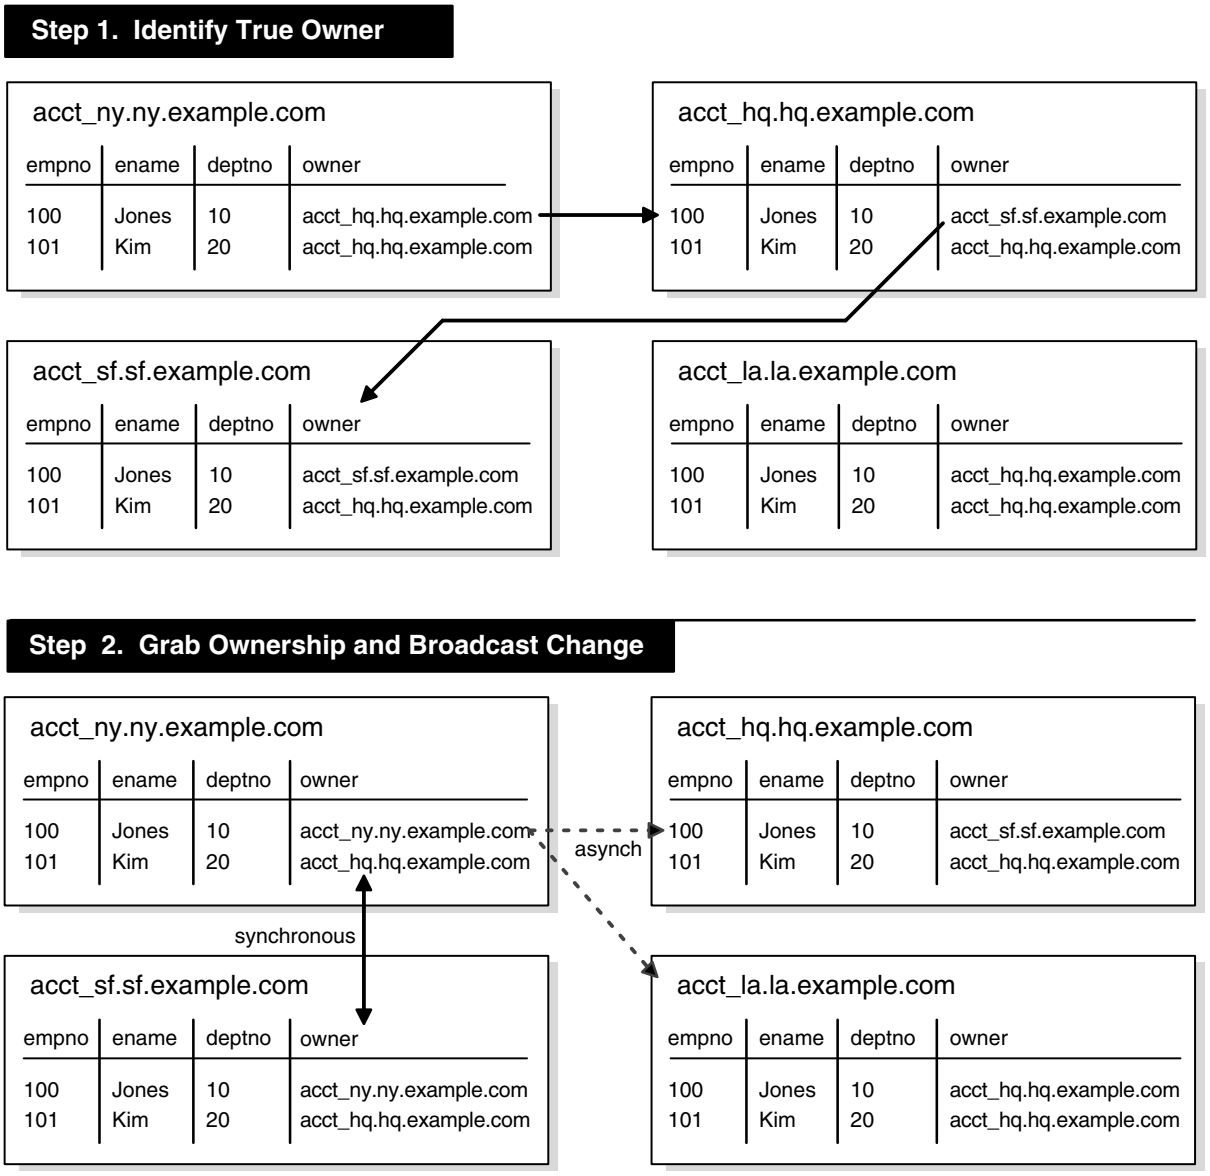


## Locating the Owner of a Row

To obtain ownership, the `acct_ny` database uses a simple recursive algorithm to locate the owner of the row. The sample code for this algorithm is shown as follows:

```
-- Sample code for locating the token owner.
```

```
-- This is for a table TABLE_NAME with primary key PK.
-- Initial call should initialize loc_epoch to 0 and loc_owner
-- to the local global name.
get_owner(PK IN primary_key_type, loc_epoch IN OUT NUMBER,
          loc_owner IN OUT VARCHAR2)
{
  -- use dynamic SQL (dbms_sql) to perform a select similar to
  -- the following:
  SELECT owner, epoch into rmt_owner, rmt_epoch
     FROM TABLE_NAME@loc_owner
     WHERE primary_key = PK FOR UPDATE;
  IF rmt_owner = loc_owner AND rmt_epoch >= loc_epoch THEN
   loc_owner := rmt_owner;
   loc_epoch := rmt_epoch;
   RETURN;
  ELSIF rmt_epoch >= loc_epoch THEN
   get_owner(PK, rmt_epoch, rmt_owner);
   loc_owner := rmt_owner;
   loc_epoch := rmt_epoch;
   RETURN;
  ELSE
   raise_application_error(-20000, 'No owner for row');
  END IF;}
```

## Obtaining Ownership

After locating the owner of the row, the `acct_ny` site gets ownership from the `acct_sf` site by completing the following steps:

1. Lock the row at the `sf` site to prevent any changes from occurring while ownership is being exchanged.

   This operation ensures that only one site considers itself to be the owner at all times. The update at the `sf` site should not be replicated using `DBMS_REPUTIL.REPLICATION_OFF`. The replicated change of ownership at the `ny` site in Step 4 will ultimately be propagated to all other sites in the replication environment, including the `sf` site, where it will have no effect.

2. Synchronously update the owner information at both the `sf` and `ny` sites.

3. Update the row information at the new owner site, `ny`, with the information from the current owner site, `sf`.

   This data is guaranteed to be the most recent. This time, the change at the `ny` site should not be replicated. Any queued changes to this data at the `sf` site are propagated to all other sites in the usual manner. When the `sf` change is propagated to `ny`, it is ignored because of the values of the epoch numbers, as described in the next bullet point.

4. Update the epoch number at the new owner site to be one greater than the value at the previous site.

   Perform this update at the new owner only, and then asynchronously propagate this update to the other master sites. Incrementing the epoch number at the new owner site prevents ordering conflicts.

   When the `sf` changes (that were in the deferred queue in Step 2 preceding) are ultimately propagated to the `ny` site, the `ny` site ignores them because they have a lower epoch number than the epoch number at the `ny` site for the same data.

As another example, suppose the hq site received the sf changes after receiving the ny changes, the hq site would ignore the sf changes because the changes applied from the ny site would have the greater epoch number.

### Applying the Change

You should design your application to implement this method of token passing for you automatically whenever you perform an update. You should not allow the owner or epoch columns to be updated outside this application. The lock that you grab when you change ownership is released when you apply your actual update. The changed information, along with the updated owner and epoch information, are asynchronously propagated to the other sites in the usual manner.

## Auditing Successful Conflict Resolution

Whenever Oracle detects and successfully resolves an update, delete, or uniqueness conflict, you can view information about what method was used to resolve the conflict by querying the ALL_REPRESOLUTION_STATISTICS data dictionary view. This view is updated only if you have enabled conflict resolution statistics gathering for the table involved in the conflict.

> **See Also:** The ALL_REPRESOLUTION_STATISTICS view on page 23-35 for more information

### Collecting Conflict Resolution Statistics

Use the REGISTER_STATISTICS procedure in the DBMS_REPCAT package to collect information about the successful resolution of update, delete, and uniqueness conflicts for a table. The following example gathers statistics for the employees table in the hr schema:

```
BEGIN
    DBMS_REPCAT.REGISTER_STATISTICS (
      sname => 'hr',
      oname => 'employees');
END;
/
```

### Viewing Conflict Resolution Statistics

After calling REGISTER_STATISTICS for a table, each conflict that is successfully resolved for that table is logged in the ALL_REPRESOLUTION_STATISTICS data dictionary view. Information about unresolved conflicts is always logged in the DEFERROR view, whether the object is registered or not.

> **See Also:** The ALL_REPRESOLUTION_STATISTICS view on page 23-35 and the DEFERROR view on page 26-6 for more information

### Canceling Conflict Resolution Statistics

Use the CANCEL_STATISTICS procedure in the DBMS_REPCAT package if you no longer want to collect information about the successful resolution of update, delete, and uniqueness conflicts for a table. The following example cancels statistics gathering on the employees table in the hr schema:

```
BEGIN
```

```
    DBMS_REPCAT.CANCEL_STATISTICS (
      sname => 'hr',
      oname => 'employees');
END;
/
```

## Clearing Statistics Information

If you registered a table to log information about the successful resolution of update, delete, and uniqueness conflicts, then you can remove this information from the DBA_REPRESOLUTION_STATISTICS data dictionary view by calling the PURGE_STATISTICS procedure in the DBMS_REPCAT package.

The following example purges the statistics gathered about conflicts resolved due to inserts, updates, and deletes on the employees table between January 1 and March 31:

```
BEGIN
    DBMS_REPCAT.PURGE_STATISTICS (
      sname  => 'hr',
      oname  => 'employees',
      start_date => '01-JAN-2001',
      end_date => '31-MAR-2001');
END;
/
```

# Part II

# Managing and Monitoring Your Replication Environment

Part II contains instructions on using the replication management API to manage your replication environment, as well as instructions on using the data dictionary to monitor your replication environment.

Part II contains the following chapters:


- Chapter 7, "Managing a Master Replication Environment"
- Chapter 8, "Managing a Materialized View Replication Environment"
- Chapter 9, "Managing Replication Objects and Queues"
- Chapter 10, "Monitoring a Replication Environment"

# 7

# Managing a Master Replication Environment

As your data delivery needs change due to growth, shrinkage, or emergencies, you are undoubtedly going to need to change the configuration of your replication environment. This chapter discusses managing the master sites of your replication environment. Specifically, this section describes altering and reconfiguring your master sites.

This chapter contains these topics:


- Changing the Master Definition Site
- Adding New Master Sites
- Removing a Master Site from a Master Group
- Updating the Comments Fields in Data Dictionary Views
- Using Procedural Replication


## Changing the Master Definition Site

Many replication administrative tasks can be performed only from the master definition site. Use the `RELOCATE_MASTERDEF` procedure in the `DBMS_REPCAT` package to move the master definition site to another master site. This API is especially useful when the master definition site becomes unavailable and you must specify a new master definition site (see "Option 2: The Old Master Definition Site Is Not Available" on page 7-2).

### Option 1: All Master Sites Are Available

Perform the actions in this section to change the master definition site if all master sites are available. Meet the following requirements to complete these actions:

**Executed As**: Replication Administrator

**Executed At**: Any Master Site

**Replication Status**: Running Normally (Not Quiesced)

Complete the following steps:

#### Step 1 In SQL*Plus, connect to a master site as the replication administrator.

See *Oracle Database Administrator's Guide* for information about connecting to a database in SQL*Plus.

#### Step 2 Relocate the master definition site.

```
BEGIN
```

```
   DBMS_REPCAT.RELOCATE_MASTERDEF (
      gname => 'hr_repg',
      old_masterdef => 'orc1.example.com',
      new_masterdef => 'orc2.example.com',
      notify_masters => TRUE,
      include_old_masterdef => TRUE);
END;
/
```

### Option 2: The Old Master Definition Site Is Not Available

Perform the actions in this section to change the master definition site if the old master definition site is *not* available. Meet the following requirements to complete these actions:

**Executed As**: Replication Administrator

**Executed At**: Any Master Site

**Replication Status**: Normal

Complete the following steps:

#### Step 1 In SQL*Plus, connect to a master site as the replication administrator.

See *Oracle Database Administrator's Guide* for information about connecting to a database in SQL*Plus.

#### Step 2 Relocate the master definition site.

```
BEGIN
   DBMS_REPCAT.RELOCATE_MASTERDEF (
      gname => 'hr_repg',
      old_masterdef => 'orc1.example.com',
      new_masterdef => 'orc2.example.com',
      notify_masters => TRUE,
      include_old_masterdef => FALSE);
END;
/
```

## Adding New Master Sites

As your replication environment expands, you might need to add new master sites to a master group. You can either add new master sites to a master group that is running normally or to a master group that is quiesced. If the master group is not quiesced, then users can perform data manipulation language (DML) operations on the data while the new master sites are being added. However, more administrative actions are required when adding new master sites if the master group is not quiesced.

> **Note:** When adding a master site to a master group that contains tables with circular dependencies or a table that contains a self-referential constraint, you must precreate the table definitions and manually load the data at the new master site. The following is an example of a circular dependency: Table A has a foreign key constraint on table B, and table B has a foreign key constraint on table A.

Follow the instructions in the appropriate section to add new master sites to a master group:

- Adding New Master Sites without Quiescing the Master Group
- Adding New Master Sites to a Quiesced Master Group

## Adding New Master Sites without Quiescing the Master Group

This section contains procedures for adding new master sites to an existing master group that is not quiesced. These new sites might or might not already be replication sites (master sites or materialized view sites) in other replication groups.

You can use one of the following methods when you are adding a new master site without quiescing the master group:

- Use full database export/import or change-based recovery to add a new master site that does not currently have any replication groups. See "Using Full Database Export/Import or Change-Based Recovery" on page 7-5 for instructions.
- Use object-level export/import to add a new master site that already has other replication groups or to add a new master site that does not currently have any replication groups. See "Using Object-Level Export/Import" on page 7-14 for instructions.

Use full database export/import and change-based recovery to add all of the replication groups at the master definition site to the new master sites. When you use this method, the following conditions apply:

- The new master sites cannot have any existing replication groups.
- The master definition site cannot have any materialized view groups.
- The master definition site must be the same for all of the master groups. If one or more of these master groups have a different master definition site, then do not use full database export/import or change-based recovery. Use object-level export/import instead.
- The new master site must include all of the replication groups in the master definition site when the extension process is complete. That is, you cannot add a subset of the master groups at the master definition site to the new master site. All of the groups must be added.

If your environment does not meet all of these conditions, then you must use object-level export/import to add the new master sites. Figure 7–1 summarizes these conditions.

---

**Note:** To use change-based recovery, the existing master site and the new master site must be running under the same operating system, although the release of the operating system can differ. This condition does not apply to full database export/import.

---

***Figure 7–1 Determining Which Method to Use When Adding Master Sites***

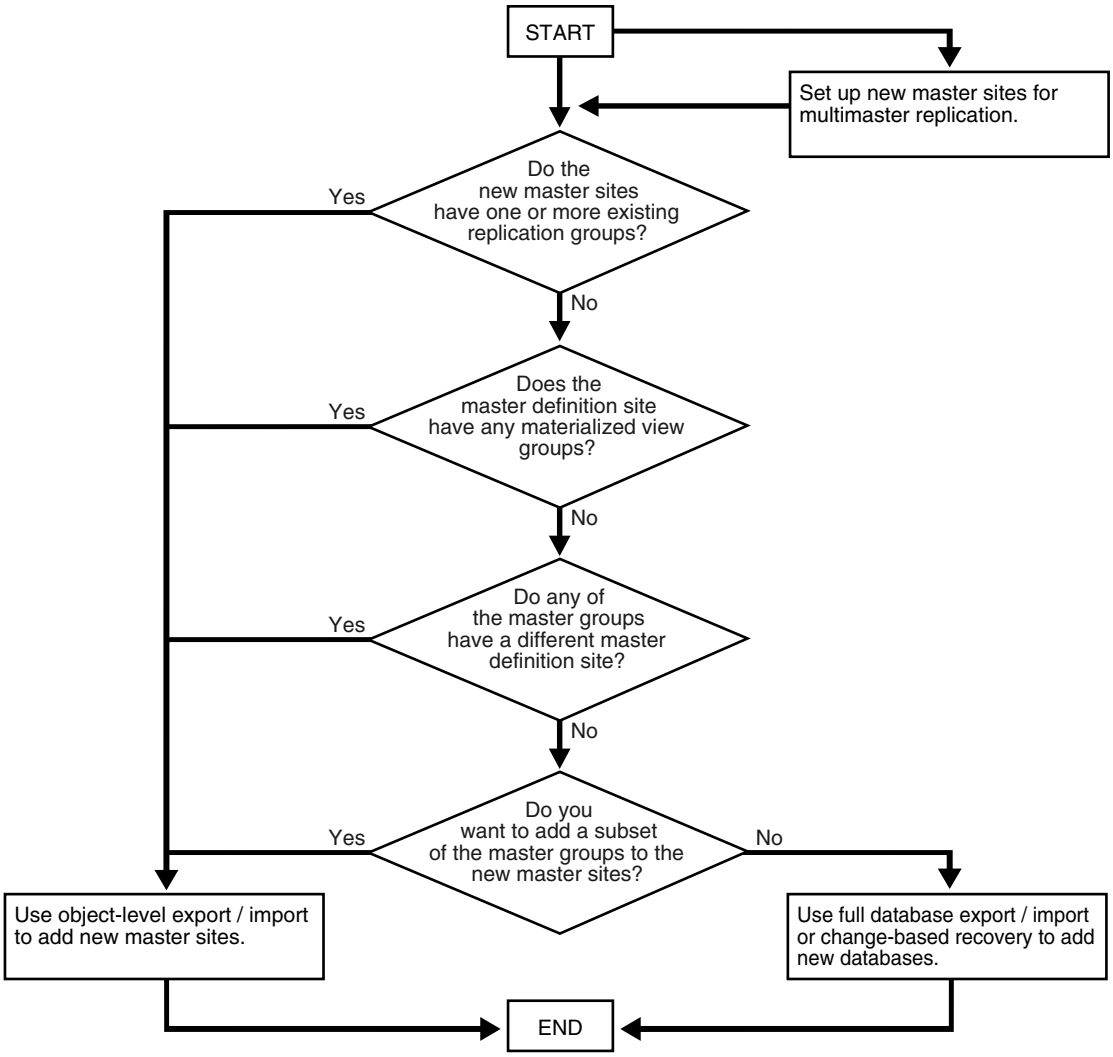


Use object-level export/import to add a master group to master sites that already have other replication groups or to add a master group to master sites that do not currently have any replication groups. This method can add one or more master groups to new master sites at a time, and you can choose a subset of the master groups at the master definition site to add to the new master sites during the operation.

If you use object-level export/import and there are integrity constraints that span multiple master groups, then you must temporarily disable these integrity constraints on the table being added to a new master site, if the other tables to which these constraints refer exist at the new master site. Initially, there are two rows in the `DEFSCHEDULE` data dictionary view that refer to the new master sites. When propagation is caught up, there is one row in this view, and when propagation from all the master sites to the new master site is caught up, you can reenable the integrity constraints you disabled.

Again, the two methods for adding new master sites without quiescing the master groups are the following:

- Full database export/import or change-based recovery
- Object-level export/import

When you use either method, propagation of deferred transactions to the new master site is partially or completely disabled while the new master sites are being added. Therefore, ensure that each existing master site has enough free space to store the largest unpropagated deferred transaction queue that you might encounter.

In addition, the following restrictions apply to both methods:

- All affected master groups must be using asynchronous replication. Synchronous replication is not allowed.
- All scheduled links must use parallel propagation with parallelism set to 1 or higher.
- Either the database links of all affected master groups must have no connection qualifier or they must all have the same connection qualifier.
- After you begin the process of adding new master sites to one or more master groups, you must wait until these new master sites are added before you begin to add another set of master sites to any of the affected master groups. If there is information about an affected master group in the `DBA_NEW_REPSITES` data dictionary view at the master definition site, then the process is started and is not yet complete for that master group.
- After you begin the process of adding new master sites to one or more master groups, you cannot relocate the master definition site for these master groups until the new master sites are added. If there is information about an affected master group in the `DBA_NEW_REPSITES` data dictionary view, then the process is started and is not yet complete for that master group.
- Only one add master site request at a time is allowed at a master site. For example, if `hq1.example.com` is the master definition site for `mgroup1` and `hq2.example.com` is the master definition site for `mgroup2`, then you cannot add `hq1.example.com` to `mgroup2` and `hq2.example.com` to `mgroup1` at the same time.
- If you are using object-level or full database export/import, then ensure that there is enough space in your rollback segments or undo tablespace for the export.

Also, before adding new master sites with either method, ensure that you properly set up your new master sites for multimaster replication.

---

**Note:** If progress appears to stop during one of the procedures described in the following sections, then check your trace files and the alert log for messages.

---

**See Also:**

- "Setting Up Master Sites" on page 2-3 for information about setting up your new master sites for multimaster replication
- *Oracle Database Administrator's Guide* for more information about trace files and the alert log
- *Oracle Database Administrator's Guide* for information about managing undo space

### Using Full Database Export/Import or Change-Based Recovery

Figure 7–2 shows the major steps for using full database export/import or change-based recovery to add new master sites to a master group without quiescing. The following example script adds the new master sites `orc4.example.com` and

orc5.example.com to the hr_repg master group. In this example, orc4.example.com is added using full database export/import and orc5.example.com is added using change-based recovery.

***Figure 7–2 Using Full Database Export/Import or Change-Based Recovery***

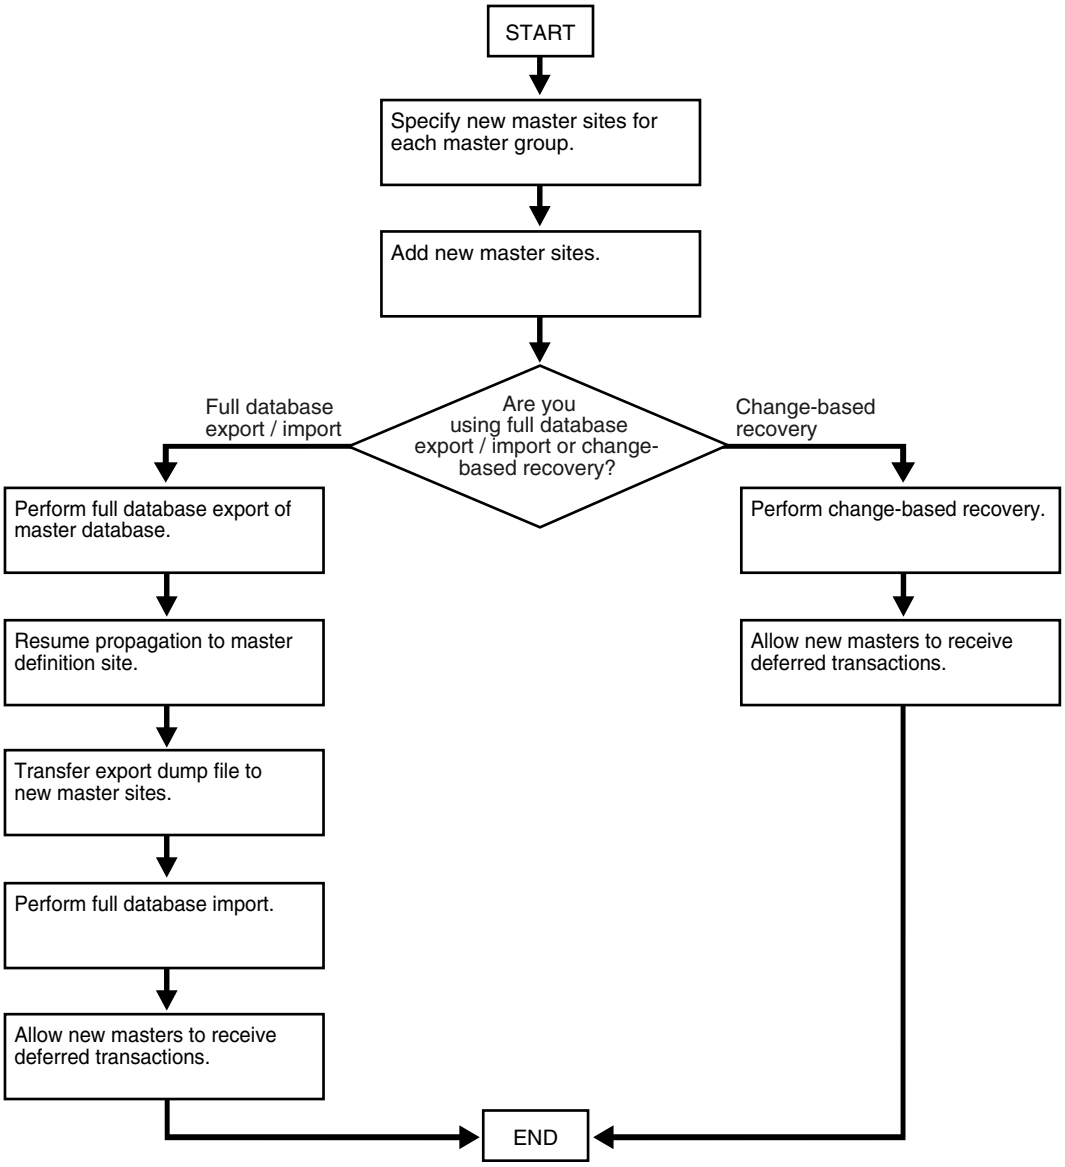


Meet the following requirements to complete these actions:

**Executed As**: Replication Administrator, unless specified otherwise

**Executed At**:

- Step 1 at Each New Master Site
- Steps 2 - 5 at Master Definition Site
- Step 6 at the Master Definition Site and at Each New Master Site
- Step 7 requires an export at the Master Definition site and a file transfer between sites.

- Steps 8 - 10 at Each New Master Site

**Replication Status**: Running Normally (Not Quiesced)

Complete the following steps to use full database export/import or change-based recovery to add sites to a master group.

> **Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

```
/************************ BEGINNING OF SCRIPT ****************************
```

**Step 1 If you are using full database export/import, then create the databases that you want to add to the master group.**

This step is not required if you are using change-based recovery.

> **See Also:** *Oracle Database Administrator's Guide* for information about creating a database

```
*/

SET ECHO ON

SPOOL add_masters_full.out

PAUSE Press <RETURN> when the databases for the new master sites are created.

/*
```

**Step 2 Set up each new master site as a replication site.**

Remember that you must configure the following:

```
*/

PAUSE Press <RETURN> to continue the new master sites have been setup and the
required scheduled links have been created.

/*
```

> **See Also:**
>
> - *Oracle Database Advanced Replication* for information about scheduled links
> - "Setting Up Master Sites" on page 2-3
> - "Creating Scheduled Links Between the Master Sites" on page 2-13

- The replication administrator at each new master site
- A scheduled link from each existing master site to each new master site
- A scheduled link from each new master site to each existing master site
- A schedule purge job at each new master site

**Step 3 Connect as the replication administrator to the master definition site.**

```
*/

CONNECT repadmin@orc1.example.com

/*
```

**Step 4 Specify new master sites for each master group.**

Before you begin, create the required scheduled links between existing master sites and each new master site if they do not exist.

**See Also:**

- *Oracle Database Advanced Replication* for information about scheduled links
- "Creating Scheduled Links Between the Master Sites" on page 2-13 for examples

```
*/

BEGIN
   DBMS_REPCAT.SPECIFY_NEW_MASTERS (
      gname => 'HR_REPG',
      master_list => 'orc4.example.com,orc5.example.com');
END;
/

/*
```

You can begin to track the extension process by querying the following data dictionary views in another SQL*Plus session:

- DBA_REPSITES_NEW
- DBA_REPEXTENSIONS

```
*/

PAUSE Press <RETURN> when you have completed the these steps.

/*
```

**Step 5 Add the new master sites.**

Before running the following procedure, ensure that there are an adequate number of background jobs running at each new master site. If you are using full database export/import, then ensure that there is enough space in your rollback segments or undo tablespace for the export before you run this procedure.

**See Also:**

- *Oracle Database Advanced Replication* for information about setting the JOB_QUEUE_PROCESSES initialization parameter properly for a replication environment
- *Oracle Database Administrator's Guide* for information about managing undo space

```
*/
```

```
VARIABLE masterdef_flashback_scn NUMBER;
VARIABLE extension_id VARCHAR2(32);
BEGIN
   DBMS_REPCAT.ADD_NEW_MASTERS (
      export_required => TRUE,
      available_master_list => NULL,
      masterdef_flashback_scn => :masterdef_flashback_scn,
      extension_id => :extension_id,
      break_trans_to_masterdef => FALSE,
      break_trans_to_new_masters => FALSE,
      percentage_for_catchup_mdef => 80,
      cycle_seconds_mdef => 60,
      percentage_for_catchup_new => 80,
      cycle_seconds_new => 60);
END;
/

/*
```

The values for `masterdef_flashback_scn` and `extension_id` are saved into variables to be used later in the process. To see these values, you can query the `DBA_REPSITES_NEW` and `DBA_REPEXTENSIONS` data dictionary views.

```
*/

PAUSE Press <RETURN> when you have completed the these steps.

/*
```

If you must undo the changes made to a particular master site by the `SPECIFY_NEW_MASTERS` and `ADD_NEW_MASTERS` procedures, then use the `DBMS_REPCAT.UNDO_ADD_NEW_MASTERS_REQUEST` procedure.

For the `export_required` parameter, `TRUE` is specified because `orc4.example.com` is being added using full database export/import. Although `orc5.example.com` is using change-based recovery, the `TRUE` setting is correct because at least one new master site is added using export/import.

After successfully executing this procedure, monitor its progress by querying the `DBA_REPCATLOG` data dictionary view in another SQL*Plus session. Do not proceed to Step 7 until there is no remaining information in this view about adding the new master sites. Assuming no extraneous information exists in `DBA_REPCATLOG` from other operations, you can enter the following statement:

```
SELECT COUNT(*) FROM DBA_REPCATLOG;
```

All of the processing is complete when this statement returns zero (0).

```
*/

PAUSE Press <RETURN> to continue when DBA_REPCATLOG is empty.

/*
```

**Step 6 If you are using full database export/import, then create a directory object at each database.**

For master sites being added using change-based recovery, this step is not required and you can proceed to Step 8 on page 7-12.

Each database involved in this operation must have a directory object to hold the Data Pump dump file, and the user who will perform the export or import must have `READ`

and WRITE privileges on this directory object. In this example, a Data Pump export is performed at the master definition site, and a Data Pump import is performed at each new master site.

If you are using full database export/import, then, while connected in SQL*Plus to the a database as an administrative user who can create directory objects using the SQL statement CREATE DIRECTORY, create a directory object to hold the Data Pump dump file and log files. For example:

```
*/

CONNECT system@orc1.example.com

CREATE DIRECTORY DPUMP_DIR AS '/usr/dpump_dir';

CONNECT system@orc4.example.com

CREATE DIRECTORY DPUMP_DIR AS '/usr/dpump_dir';

CONNECT system@orc5.example.com

CREATE DIRECTORY DPUMP_DIR AS '/usr/dpump_dir';

/*
```

In this example, SYSTEM user performs all exports and imports. If a user other than the user who created the directory object will perform the export or import, then grant this user READ and WRITE privileges on the directory object.

Ensure that you complete these actions at each database involved in the operation.

#### Step 7 Perform the following substeps for the master sites being added using full database export/import.

For master sites being added using change-based recovery, these substeps are not required and you can proceed to Step 8 on page 7-12.

Perform full database export of the master definition database. Use the system change number (SCN) returned by the masterdef_flashback_scn parameter in Step 5 for the FLASHBACK_SCN export parameter.

You can query the DBA_REPEXTENSIONS data dictionary view for the FLASHBACK_SCN value:

```
SELECT FLASHBACK_SCN FROM DBA_REPEXTENSIONS;
```

In this example, assume that the value returned by this query is 124723.

In this example, orc4.example.com is using full database export/import. Therefore, perform the full database export of the master definition database so that it can be imported into orc4.example.com during a later step. However, the orc5.example.com database is using change-based recovery. Therefore, the export would not be required if you were adding only orc5.example.com.

On a command line, perform the export. This example connects as the SYSTEM user. The following is an example Data Pump export command:

```
expdp system FULL=y DIRECTORY=DPUMP_DIR DUMPFILE=fulldb_orc1.dmp FLASHBACK_
SCN=124723
```

Consider the following when you run the Export utility:

- Only users with the DBA role or the EXP_FULL_DATABASE role can export in full database mode.
- Ensure that the UNDO_RETENTION initialization parameter is set correctly before performing the export.

**See Also:**

- *Oracle Database Utilities* for information about performing a Data Pump export
- *Oracle Database Administrator's Guide* for information about managing undo space and setting this parameter

```
*/

PAUSE Press <RETURN> to continue when the export is complete.

/*
```

Resume propagation to the master definition site.

Running the following procedure indicates that export is effectively finished and propagation can be enabled for both extended and unaffected master groups at the master sites.

```
*/

BEGIN
   DBMS_REPCAT.RESUME_PROPAGATION_TO_MDEF (
      extension_id => :extension_id);
END;
/

/*
```

You can find the extension_id by querying the DBA_REPSITES_NEW data dictionary view.

Transfer the export dump file to the new master sites.

Using the DBMS_FILE_TRANSFER package, FTP, or some other method, transfer the export dump file to the other new master sites that are being added with full database export/import. You will need this export dump file at each new site to perform the import described in the next step.

```
*/

PAUSE Press <RETURN> to continue after transferring the dump file.

/*
```

Set the JOB_QUEUE_PROCESSES initialization parameter to zero for each new master site.

```
*/

PAUSE Press <RETURN> to continue after JOB_QUEUE_PROCESSES is set to zero at each
new master site.

/*
```

**Step 8 Perform import or change-based recovery at each new master site.**

If you are using full database export/import, then complete the full database import of the database you exported in Step 7 at each new master site that is being added with full database export/import.

Perform the import. This example connects as the SYSTEM user to perform the import at orc4.example.com. The following is an example import command:

```
impdp system FULL=y DIRECTORY=DPUMP_DIR DUMPFILE=fulldb_orc1.dmp
```

Only users with the DBA role or the IMP_FULL_DATABASE role can import in full database mode.

> **See Also:** *Oracle Database Utilities* for information about performing a Data Pump import

```
*/

PAUSE Press <RETURN> to continue when the import is complete.

/*
```

If you are using change-based recovery, then perform change-based recovery using the system change number (SCN) returned by the masterdef_flashback_scn parameter in Step 5. You can query the DBA_REPEXTENSIONS data dictionary view for the masterdef_flashback_scn value.

You can perform a change-based recovery in one of the following ways:

- Using the SQL*Plus RECOVER command. See the *Oracle Database Backup and Recovery User's Guide* for instructions.
- Using the Recovery Manager (RMAN) DUPLICATE command. See the *Oracle Database Backup and Recovery User's Guide* for instructions.

Connect to the site where you will perform the change-based recovery:

```
*/

CONNECT repadmin@orc5.example.com

PAUSE Press <RETURN> to continue when the change-based recovery is complete. You
can use a separate terminal window to perform the change-based recovery.

/*
```

**Step 9 Configure the new sites for multimaster replication by completing the following steps:**

1. Ensure that the database structures, such as the data files, exist for the replicated schemas at each new master site. In this example, the replicated schema is hr.
2. Set the global name for each new master site. The global name for each new master site must match the global names specified in the SPECIFY_NEW_MASTERS procedure that you ran in Step 4. You can query the DBLINK column in the DBA_REPSITES_NEW data dictionary view to see the global name for each new master site.

   You can set the global name using the ALTER DATABASE statement, as in the following example:

   ```
   ALTER DATABASE RENAME GLOBAL_NAME TO orc4.example.com;
   ```

3. Create the appropriate scheduled links between the new master sites and the existing master sites, including the master definition site.

   **See Also:** "Creating Scheduled Links Between the Master Sites" on page 2-13 for information

```
*/

PAUSE Press <RETURN> when you have completed the these steps.

/*
```

### Step 10 Allow new masters to receive deferred transactions.

The following procedure enables the propagation of deferred transactions from other prepared new master sites and existing master sites to the invocation master site. This procedure also enables the propagation of deferred transactions from the invocation master site to the other new master sites and existing master sites.

> **Caution:** Do not invoke this procedure until instantiation (export/import or change-based recovery) of the new master site is complete.
>
> Do not allow any data manipulation language (DML) statements directly on the objects in the extended master group in the new master site until execution of this procedure returns successfully, because these DML statements might not be replicated.

```
*/

CONNECT repadmin@orc4.example.com

BEGIN
   DBMS_REPCAT.PREPARE_INSTANTIATED_MASTER (
      extension_id  => :extension_id);
END;
/

CONNECT repadmin@orc5.example.com

BEGIN
   DBMS_REPCAT.PREPARE_INSTANTIATED_MASTER (
      extension_id  => :extension_id);
END;
/

SET ECHO OFF

SPOOL OFF

/*
```

**Note:** You can find the `extension_id` by querying the `DBA_REPSITES_NEW` data dictionary view.

```
************************ END OF SCRIPT *******************************/
```

### Using Object-Level Export/Import

Figure 7–3 shows the major steps for using object-level export/import to add new master sites to a master group without quiescing. The following example procedure adds the new master sites `orc4.example.com` and `orc5.example.com` to the `hr_repg` master group. An object-level export/import involves exporting and importing the tables in a master group. When you export and import the tables, other dependent database objects, such as indexes, are exported and imported as well.

If you have an integrity constraint that spans two master groups, then you have a child table in one master group (the child master group) and a parent table in a different master group (the parent master group). In this case, Oracle recommends that you add new master sites to both master groups at the same time. However, if you cannot do this, then you must quiesce the child master group before adding new master sites to it. Here, the child table includes a foreign key, which makes it dependent on the values in the parent table. If you do not quiesce the child master group, then conflicts might result when you add master sites to it. You can still add master sites to the parent master group without quiescing it.

***Figure 7–3 Using Object-Level Export/Import***

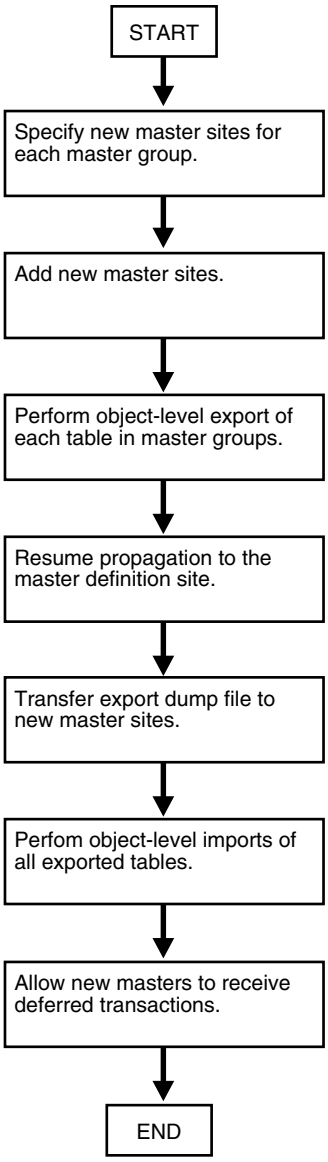


Meet the following requirements to complete these actions:

**Executed As**: Replication Administrator, unless specified otherwise

**Executed At**:

- Steps 1 - 6 at Master Definition Site
- Step 7 at the Master Definition Site and at Each New Master Site
- Steps 8 - 9 at Master Definition Site
- Step 10 requires a file transfer between sites.
- Steps 11 - 12 at Each New Master Site

**Replication Status**: Running Normally (Not Quiesced)

Complete the following steps to use object-level export/import to add sites to a master group.

**Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

```
/************************* BEGINNING OF SCRIPT ******************************
```

**Step 1 If the users for the replicated schemas do not exist at the new master sites, then create them now.**

In this example, the replicated schema is `hr`. This schema probably exists at the new master sites because it is a sample schema that is installed when you install Oracle.

**See Also:** *Oracle Database Sample Schemas* for general information about the sample schemas and for information about installing them

```
*/

SET ECHO ON

SPOOL add_masters_object.out

PAUSE Press <RETURN> to continue when the users are created at the new master
sites.

/*
```

**Step 2 If any of the tables in the master group have circular dependencies, then precreate these tables at the new master sites.**

Failure to precreate these tables will result in errors later in the procedure. If there are no circular dependencies, then this step is not required, and you can proceed to Step 3.

Some tables in the `hr` schema contain circular dependencies. Therefore, in this example, the tables in the `hr` schema must be precreated at each new master site. Again, the `hr` schema tables are typically created during Oracle installation and so might exist at the new master sites.

If you must precreate tables, then disable referential integrity constraints for these tables at the new master sites before the import. Referential integrity constraints can cause errors when you import data into existing tables. This example disables the referential integrity constraints for the precreated tables in the `hr` schema at the new master sites.

Further, the precreated tables at the new master sites should not contain any data. This example truncates the tables in the `hr` schema at the new master sites to ensure that they do not contain any data.

**See Also:**

- The note under "Adding New Master Sites" on page 7-2 for more information about circular dependencies
- *Oracle Database Utilities* for information about importing data into existing tables

```
*/
```

```
PAUSE Press <RETURN> to continue when the tables are precreated at the new
master sites, if table precreation is required. After the tables are
precreated, the following statements disable the referential integrity
constraints related to the hr schema and truncate the tables in the hr schema
at the new site.

CONNECT oe@orc4.example.com

ALTER TABLE oe.warehouses
  DISABLE CONSTRAINT warehouses_location_fk;

ALTER TABLE oe.customers
  DISABLE CONSTRAINT customers_account_manager_fk;

ALTER TABLE oe.orders
  DISABLE CONSTRAINT orders_sales_rep_fk;

CONNECT hr@orc4.example.com

ALTER TABLE hr.countries
  DISABLE CONSTRAINT countr_reg_fk;

ALTER TABLE hr.departments
  DISABLE CONSTRAINT dept_mgr_fk
  DISABLE CONSTRAINT dept_loc_fk;

ALTER TABLE hr.employees
  DISABLE CONSTRAINT emp_dept_fk
  DISABLE CONSTRAINT emp_job_fk
  DISABLE CONSTRAINT emp_manager_fk;

ALTER TABLE hr.job_history
  DISABLE CONSTRAINT jhist_job_fk
  DISABLE CONSTRAINT jhist_emp_fk
  DISABLE CONSTRAINT jhist_dept_fk;

ALTER TABLE hr.locations
  DISABLE CONSTRAINT loc_c_id_fk;

TRUNCATE TABLE hr.countries;
TRUNCATE TABLE hr.departments;
TRUNCATE TABLE hr.employees;
TRUNCATE TABLE hr.jobs;
TRUNCATE TABLE hr.job_history;
TRUNCATE TABLE hr.locations;
TRUNCATE TABLE hr.regions;

CONNECT oe@orc5.example.com

ALTER TABLE oe.warehouses
  DISABLE CONSTRAINT warehouses_location_fk;

ALTER TABLE oe.customers
  DISABLE CONSTRAINT customers_account_manager_fk;

ALTER TABLE oe.orders
  DISABLE CONSTRAINT orders_sales_rep_fk;

CONNECT hr@orc5.example.com
```

```
ALTER TABLE hr.countries
  DISABLE CONSTRAINT countr_reg_fk;

ALTER TABLE hr.departments
  DISABLE CONSTRAINT dept_mgr_fk
  DISABLE CONSTRAINT dept_loc_fk;

ALTER TABLE hr.employees
  DISABLE CONSTRAINT emp_dept_fk
  DISABLE CONSTRAINT emp_job_fk
  DISABLE CONSTRAINT emp_manager_fk;

ALTER TABLE hr.job_history
  DISABLE CONSTRAINT jhist_job_fk
  DISABLE CONSTRAINT jhist_emp_fk
  DISABLE CONSTRAINT jhist_dept_fk;

ALTER TABLE hr.locations
  DISABLE CONSTRAINT loc_c_id_fk;

TRUNCATE TABLE hr.countries;
TRUNCATE TABLE hr.departments;
TRUNCATE TABLE hr.employees;
TRUNCATE TABLE hr.jobs;
TRUNCATE TABLE hr.job_history;
TRUNCATE TABLE hr.locations;
TRUNCATE TABLE hr.regions;

/*
```

**Step 3 Set up each new master site as a replication site.**

Remember that you must configure the following:

- The replication administrator at each new master site
- A scheduled link from each existing master site to each new master site
- A scheduled link from each new master site to each existing master site
- A schedule purge job at each new master site

```
*/

PAUSE Press <RETURN> to continue the new master sites have been setup and the
required scheduled links have been created.

/*
```

**See Also:**

- *Oracle Database Advanced Replication* for information about scheduled links
- "Setting Up Master Sites" on page 2-3
- "Creating Scheduled Links Between the Master Sites" on page 2-13

**Step 4 Connect to the master definition site as the replication administrator.**

```
*/

CONNECT repadmin@orc1.example.com
```

```
/*
```

**Step 5 Specify new master sites for each master group.**

```
*/

BEGIN
   DBMS_REPCAT.SPECIFY_NEW_MASTERS (
      gname => 'hr_repg',
      master_list => 'orc4.example.com,orc5.example.com');
END;
/

/*
```

You can begin to track the extension process by querying the following data dictionary views in another SQL*Plus session:

- DBA_REPSITES_NEW
- DBA_REPEXTENSIONS

**Step 6 Add the new master sites.**

Before running the following procedure, ensure that there are an adequate number of background jobs running at each new master site. Also, ensure that there is enough space in your rollback segments or undo tablespace for the export before you run this procedure.

**See Also:**

- *Oracle Database Advanced Replication* for information about setting the JOB_QUEUE_PROCESSES initialization parameter properly for a replication environment
- *Oracle Database Administrator's Guide* for information about managing undo space

```
*/

VARIABLE masterdef_flashback_scn NUMBER;
VARIABLE extension_id VARCHAR2(32);
BEGIN
   DBMS_REPCAT.ADD_NEW_MASTERS (
      export_required => TRUE,
      available_master_list => 'orc4.example.com,orc5.example.com',
      masterdef_flashback_scn => :masterdef_flashback_scn,
      extension_id => :extension_id,
      break_trans_to_masterdef => FALSE,
      break_trans_to_new_masters => FALSE,
      percentage_for_catchup_mdef => 80,
      cycle_seconds_mdef => 60,
      percentage_for_catchup_new => 80,
      cycle_seconds_new => 60);
END;
/

/*
```

The sites specified for the available_master_list parameter must be same as the sites specified in the SPECIFY_NEW_MASTERS procedure in Step 5.

The values for masterdef_flashback_scn and extension_id are saved into variables to be used later in the process. To see these values, you can also query the DBA_REPSITES_NEW and DBA_REPEXTENSIONS data dictionary views.

If you must undo the changes made to a particular master site by the SPECIFY_NEW_MASTERS and ADD_NEW_MASTERS procedures, then use the UNDO_ADD_NEW_MASTERS_REQUEST procedure.

After successfully executing this procedure, monitor its progress by querying the DBA_REPCATLOG data dictionary view in another SQL*Plus session. Do not proceed to Step 8 until there is no remaining information in this view about adding the new master sites. Assuming there is no extraneous information in DBA_REPCATLOG from other operations, you can enter the following statement:

```
SELECT COUNT(*) FROM DBA_REPCATLOG;
```

All of the processing is complete when this statement returns zero (0).

```
*/

PAUSE Press <RETURN> to continue when DBA_REPCATLOG is empty.

/*
```

**Step 7 Create a directory object at each database.**

Each database involved in this operation must have a directory object to hold the Data Pump dump file, and the user who will perform the export or import must have READ and WRITE privileges on this directory object. In this example, a Data Pump export is performed at the master definition site, and a Data Pump import is performed at each new master site.

While connected in SQL*Plus to the a database as an administrative user who can create directory objects using the SQL statement CREATE DIRECTORY, create a directory object to hold the Data Pump dump file and log files. For example:

```
*/

CONNECT system@orc1.example.com

CREATE DIRECTORY DPUMP_DIR AS '/usr/dpump_dir';

CONNECT system@orc4.example.com

CREATE DIRECTORY DPUMP_DIR AS '/usr/dpump_dir';

CONNECT system@orc5.example.com

CREATE DIRECTORY DPUMP_DIR AS '/usr/dpump_dir';

/*
```

In this example, SYSTEM user performs all exports and imports. If a user other than the user who created the directory object will perform the export or import, then grant this user READ and WRITE privileges on the directory object.

Ensure that you complete these actions at each database involved in the operation.

### Step 8 Perform object-level export of tables at master definition database.

At the master definition database, perform an object-level export for each master table in the master groups that will be created at the new master sites. An object-level export includes exports performed in table mode, user mode, or tablespace mode.

Use the system change number (SCN) returned by the `masterdef_flashback_scn` parameter in Step 6 for the `FLASHBACK_SCN` export parameter. You can query the `DBA_REPEXTENSIONS` data dictionary view for the `FLASHBACK_SCN` value:

```
SELECT FLASHBACK_SCN FROM DBA_REPEXTENSIONS;
```

In this example, assume that the SCN value is `3456871`.

On a command line, perform the export. This example connects as the `SYSTEM` user. The following is an example Data Pump export command:

```
expdp system TABLES=HR.COUNTRIES,HR.DEPARTMENTS,HR.EMPLOYEES,
HR.JOB_HISTORY,HR.JOBS,HR.LOCATIONS,HR.REGIONS DIRECTORY=DPUMP_DIR
DUMPFILE=hr_tables.dmp CONTENT=data_only FLASHBACK_SCN=3456871
```

The `CONTENT` parameter is used in this example because the tables exist at the import sites. You might not need to specify this parameter.

Ensure that the `UNDO_RETENTION` initialization parameter is set correctly before performing the export.

**See Also:**

- *Oracle Database Utilities* for information about performing a Data Pump export
- *Oracle Database Administrator's Guide* for information about managing undo space and setting the `UNDO_RETENTION` initialization parameter

```
*/

PAUSE Press <RETURN> to continue when the export is complete.

/*
```

### Step 9 Resume propagation to the master definition site.

Running the following procedure indicates that export is effectively finished and propagation can be enabled for both extended and unaffected master groups at the master sites.

```
*/

CONNECT repadmin@orc1.example.com

BEGIN
   DBMS_REPCAT.RESUME_PROPAGATION_TO_MDEF (
      extension_id => :extension_id);
END;
/

/*
```

You can find the `extension_id` by querying the `DBA_REPSITES_NEW` data dictionary view.

**Step 10 Transfer the export dump files to the new master sites.**

Using the DBMS_FILE_TRANSFER package, FTP, or some other method, transfer the export dump files to the other new master sites that are being added with object-level export/import. You will need these export dump files at each new site to perform the import described in the next step.

```
*/

PAUSE Press <RETURN> to continue when the export dump files have been
transfered to the new master sites that are being added with object-level
export/import.

/*
```

**Step 11 Perform object-level imports at each new master site of each object you exported in Step 8.**

On a command line, perform the import. This example connects as the SYSTEM user. The following is an example import command:

```
impdp system TABLES=HR.COUNTRIES,HR.DEPARTMENTS,HR.EMPLOYEES,
HR.JOB_HISTORY,HR.JOBS,HR.LOCATIONS,HR.REGIONS DIRECTORY=DPUMP_DIR
DUMPFILE=hr_tables.dmp CONTENT=data_only TABLE_EXISTS_ACTION=append
```

Other objects, such as the indexes based on the tables, are imported automatically. The CONTENT and TABLE_EXISTS_ACTION parameters are used in this example because the tables exist at the import sites. You might not need to specify these parameters.

> **See Also:** *Oracle Database Utilities* for information about performing a Data Pump import

Perform the object-level imports at each site:

```
*/

PAUSE Press <RETURN> to continue when the imports are complete at each site. You
can use a separate terminal window to perform the object-level imports.

/*
```

**Step 12 Allow new masters to receive deferred transactions.**

The following procedure enables the propagation of deferred transactions from other prepared new master sites and existing master sites to the invocation master site. This procedure also enables the propagation of deferred transactions from the invocation master site to the other new master sites and existing master sites.

> **Caution:** Do not invoke this procedure until object-level export/import for the new master site is complete.
>
> Do not allow any data manipulation language (DML) statements directly on the objects in the extended master group in the new master site until execution of this procedure returns successfully, because these DML statements might not be replicated.

```
*/

CONNECT repadmin@orc4.example.com
```

```
BEGIN
   DBMS_REPCAT.PREPARE_INSTANTIATED_MASTER (
      extension_id  => :extension_id);
END;
/

CONNECT repadmin@orc5.example.com

BEGIN
   DBMS_REPCAT.PREPARE_INSTANTIATED_MASTER (
      extension_id  => :extension_id);
END;
/

SET ECHO OFF

SPOOL OFF

/*
```

> **Note:** You can find the `extension_id` by querying the `DBA_REPSITES_NEW` data dictionary view.

```
************************* END OF SCRIPT *******************************/
```

## Adding New Master Sites to a Quiesced Master Group

You can add new master sites to a quiesced master group in one of the following ways:

- Adding New Master Sites Using the ADD_MASTER_DATABASE Procedure
- Adding New Master Sites with Offline Instantiation Using Export/Import

Typically, you should only use the `ADD_MASTER_DATABASE` procedure if you have a relatively small master group or if you plan to precreate the replication tables and load the data into them at the new master sites. If this is not the case, the `ADD_MASTER_DATABASE` procedure might not be a good option because the entire master group is copied over the network. For larger master groups, either precreate the objects in the master group at the new master sites or use offline instantiation.

### Adding New Master Sites Using the ADD_MASTER_DATABASE Procedure

You can use the `ADD_MASTER_DATABASE` procedure to add additional master sites to an existing master group that is quiesced. Executing this procedure replicates existing master objects to the new site.

Meet the following requirements to complete these actions:

**Executed As**: Replication Administrator

**Executed At**: Master Definition Site

**Replication Status**: Quiesced

Complete the following steps to use the `ADD_MASTER_DATABASE` procedure to add sites to a master group.

**Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

```
/************************* BEGINNING OF SCRIPT ******************************
```

### Step 1 Set up the new master site.

Ensure that the appropriate schema and database links have been created before adding your new master site. Be sure to create the database links from the new master site to each of the existing masters sites. Also, create a database link from each of the existing master sites to the new master site. After the database links have been created, ensure that you also define the scheduled links for each of the new database links.

**See Also:**

- "Setting Up Master Sites" on page 2-3
- "Creating Scheduled Links Between the Master Sites" on page 2-13

```
*/

SET ECHO ON

SPOOL add_masters_quiesced.out

PAUSE Press <RETURN> to the new master site has been set up.

/*
```

### Step 2 Connect to the master definition site as the replication administrator.

```
*/

CONNECT repadmin@orc1.example.com

/*
```

### Step 3 If the replication status is normal, then change the status to quiesced.

```
*/

BEGIN
   DBMS_REPCAT.SUSPEND_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

/*
```

### Step 4 Add the new master sites.

This example assumes that the replicated objects do not exist at the new master site. Therefore, the `copy_rows` parameter is set to `TRUE` to copy the rows in the replicated objects at the master definition site to the new master site, and the `use_existing_objects` parameter is set to `FALSE` so that Advanced Replication creates the replicated

objects at the new site. If the replicated objects exist at the new site but do not contain any data, then set use_existing_objects to TRUE.

```
*/

BEGIN
   DBMS_REPCAT.ADD_MASTER_DATABASE (
      gname => 'hr_repg',
      master => 'orc4.example.com',
      use_existing_objects => FALSE,
      copy_rows => TRUE,
      propagation_mode => 'ASYNCHRONOUS');
END;
/

/*
```

You should wait until the DBA_REPCATLOG view is empty. This view has temporary information that is cleared after successful execution. Execute the following SELECT statement in another SQL*Plus session to monitor the DBA_REPCATLOG view:

```
SELECT COUNT(*) FROM DBA_REPCATLOG WHERE GNAME = 'HR_REPG';
```

All of the processing is complete when this statement returns zero (0).

```
*/

PAUSE Press <RETURN> to continue when DBA_REPCATLOG is empty.

/*
```

#### Step 5 Resume replication activity.

```
*/

BEGIN
   DBMS_REPCAT.RESUME_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

SET ECHO OFF

SPOOL OFF

/************************ END OF SCRIPT *********************************/
```

### Adding New Master Sites with Offline Instantiation Using Export/Import

Expanding established replication environments can cause network traffic when you add a new master site to your replication environment using the ADD_MASTER_DATABASE procedure. This is caused by propagating the entire contents of the table or materialized view through the network to the new replicated site.

To minimize such network traffic, you can expand your replication environment by using the offline instantiation procedure. Offline instantiation takes advantage of Oracle's Export and Import utilities, which allow you to create an export file and transfer the data to the new site through another storage medium, such as CD-ROM, tape, and so on.

The following script is an example of how to perform an offline instantiation of a master site. This script can potentially eliminate large amounts of network traffic

caused by the other method of adding a new master site to an existing quiesced master group. The script assumes that the hr schema does not exist at the new master site and instantiates this schema at the new master site. The hr schema is created automatically when Oracle is installed. You can choose to drop the hr schema at the new master site before you start this example.

Meet the following requirements to complete these actions:

**Executed As**: Replication Administrator, unless specified otherwise

**Executed At**: Master Definition Site and New Master Site

**Replication Status**: Quiesced and Partial

Complete the following steps to use offline instantiation to add sites to a master group.

---

**Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

---

```
/************************* BEGINNING OF SCRIPT ******************************
```

### Step 1 Set up the new master site.

Ensure that the appropriate schema and database links have been created before performing the offline instantiation of your new master site. Be sure to create the database links from the new master site to each of the existing masters sites. Also, create a database link from each of the existing master sites to the new master site. After the database links have been created, ensure that you also define the scheduled links for each of the new database links.

**See Also:**

- "Setting Up Master Sites" on page 2-3
- "Creating Scheduled Links Between the Master Sites" on page 2-13

```
*/

SET ECHO ON

SPOOL add_masters_instant.out

PAUSE Press <RETURN> to the new master site has been set up.

/*
```

### Step 2 Connect to the master definition site as the replication administrator.

```
*/

CONNECT repadmin@orc1.example.com

/*
```

### Step 3 Suspend master activity.

You must suspend master activity for the existing master sites before exporting your master data and beginning the offline instantiation process.

```
*/

BEGIN
   DBMS_REPCAT.SUSPEND_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

/*
```

**Step 4 Verify that there are no pending transactions in a separate SQL*Plus session.**

This includes pushing any outstanding deferred transactions, resolving any error transactions, and pushing any administrative transactions. This step must be performed at each of the existing master sites.

Check the error transaction queue.

```
SELECT * FROM DEFERROR;
```

If any deferred transactions have been entered into the error queue, then you must resolve the error situation and then manually reexecute the deferred transaction. The following is an example:

```
BEGIN
   DBMS_DEFER_SYS.EXECUTE_ERROR (
      deferred_tran_id => '128323',
      destination => 'orc1.example.com');
END;
/
```

Check for outstanding administrative requests.

```
SELECT * FROM DBA_REPCATLOG;
```

If any administrative requests remain, then you can manually execute these requests or wait for them to be executed automatically. You might need to execute the `DBMS_REPCAT.DO_DEFERRED_REPCAT_ADMIN` procedure several times, because some administrative operations have multiple steps. The following is an example:

```
BEGIN
   DBMS_REPCAT.DO_DEFERRED_REPCAT_ADMIN (
      gname => 'hr_repg',
      all_sites => TRUE);
END;
/

*/

PAUSE Press <RETURN> to continue when you have verified that there are no pending
requests.

/*
```

**Step 5 Begin offline instantiation procedure.**

```
*/

BEGIN
   DBMS_OFFLINE_OG.BEGIN_INSTANTIATION (
      gname => 'hr_repg',
```

```
      new_site => 'orc4.example.com');
END;
/

/*
```

You should wait until the DBA_REPCATLOG view is empty. This view has temporary information that is cleared after successful execution. Execute the following SELECT statement in another SQL*Plus session to monitor the DBA_REPCATLOG view:

```
SELECT * FROM DBA_REPCATLOG WHERE GNAME = 'HR_REPG';

*/

PAUSE Press <RETURN> to continue when DBA_REPCATLOG is empty.

/*
```

**Step 6 Create a directory object at each database.**

Each database involved in this operation must have a directory object to hold the Data Pump dump file, and the user who will perform the export or import must have READ and WRITE privileges on this directory object. In this example, a Data Pump export is performed at the master definition site, and a Data Pump import is performed at the new master site.

While connected in SQL*Plus to a database as an administrative user who can create directory objects using the SQL statement CREATE DIRECTORY, create a directory object to hold the Data Pump dump file and log files. For example:

```
*/

CONNECT system@orc1.example.com

CREATE DIRECTORY DPUMP_DIR AS '/usr/dpump_dir';

CONNECT system@orc4.example.com

CREATE DIRECTORY DPUMP_DIR AS '/usr/dpump_dir';

/*
```

Ensure that you complete these actions at both databases involved in the operation. In this example, SYSTEM user creates the directory objects and performs all exports and imports. If a user who does not own the directory object will perform the export or import, then grant the user READ and WRITE privileges on the directory object.

**Step 7 In a separate terminal window, perform the export.**

On a command line, perform the export. This example connects as the SYSTEM user. The following is an example Data Pump export command:

```
expdp system SCHEMAS=hr DIRECTORY=DPUMP_DIR DUMPFILE=hr_schema.dmp
```

When you export tables, their indexes are exported automatically.

> **See Also:** *Oracle Database Utilities* for information about performing a Data Pump export

```
*/
```

```
PAUSE Press <RETURN> to continue when the export is complete.

/*
```

### Step 8 Resume partial replication activity.

Because it might take some time to complete the offline instantiation process, you can resume replication activity for the remaining master sites (excluding the new master site) by executing the RESUME_SUBSET_OF_MASTERS procedure in the DBMS_OFFLINE_OG package after the export is complete. In the following example, replication activity is resumed at all master sites except the new master site -- orc4.example.com.

```
*/

CONNECT repadmin@orc1.example.com

BEGIN
   DBMS_OFFLINE_OG.RESUME_SUBSET_OF_MASTERS (
      gname => 'hr_repg',
      new_site => 'orc4.example.com');
END;
/

/*
```

### Step 9 Transfer the export dump files to the new master site.

Using the DBMS_FILE_TRANSFER package, FTP, or some other method, transfer the export dump file to the new master site. You will need this export dump file at the new site to perform the import described in the next step.

```
*/

PAUSE Press <RETURN> to continue when the export dump file has been
transfered to the new master site.

/*
```

### Step 10 Connect to the new master site as the replication administrator.

```
*/

CONNECT repadmin@orc4.example.com

/*
```

### Step 11 Prepare the new master site.

You must prepare the new site to import the data in your export file. Ensure that you execute the following procedure at the new master site.

```
*/

BEGIN
   DBMS_OFFLINE_OG.BEGIN_LOAD (
      gname => 'hr_repg',
      new_site => 'orc4.example.com');
END;
/

/*
```

**Step 12 In a separate terminal window, import data from export dump file.**

On a command line, perform the import. This example connects as the SYSTEM user. The following is an example import command:

```
impdp system SCHEMAS=hr DIRECTORY=DPUMP_DIR DUMPFILE=hr_schema.dmp
```

Other objects, such as the indexes based on the tables, are imported automatically.

> **See Also:** *Oracle Database Utilities* for information about performing a Data Pump import

```
*/

PAUSE Press <RETURN> to continue when the import is complete.

/*
```

**Step 13 Complete the load process at new master site.**

After importing the export file, you are ready to complete the offline instantiation process at the new master site. Executing the DBMS_OFFLINE_OG.END_LOAD procedure prepares the new site for normal replication activity.

```
*/

BEGIN
   DBMS_OFFLINE_OG.END_LOAD (
      gname => 'hr_repg',
      new_site => 'orc4.example.com');
END;
/

/*
```

**Step 14 Connect to the master definition site as the replication administrator.**

```
*/

CONNECT repadmin@orc1.example.com

/*
```

**Step 15 Complete instantiation process.**

After completing the steps at the new master site, you are ready to complete the offline instantiation process. Executing the END_INSTANTIATION procedure in the DBMS_OFFLINE_OG package completes the process and resumes normal replication activity at all master sites. Ensure that you execute the following procedure at the master definition site.

```
*/

BEGIN
   DBMS_OFFLINE_OG.END_INSTANTIATION (
      gname => 'hr_repg',
      new_site => 'orc4.example.com');
END;
/

SET ECHO OFF
```

```
SPOOL OFF

/************************ END OF SCRIPT *******************************/
```

## Removing a Master Site from a Master Group

When it becomes necessary to remove a master site from a master group, use the REMOVE_MASTER_DATABASES procedure to drop one or more master sites.

Meet the following requirements to complete these actions:

**Executed As**: Replication Administrator

**Executed At**: Master Definition Site

**Replication Status**: Quiesced

Complete the following steps to remove a master site.

> **Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

```
/************************ BEGINNING OF SCRIPT ****************************
```

**Step 1 Connect to the master definition site as the replication administrator.**

```
*/

SET ECHO ON

SPOOL remove_masters.out

CONNECT repadmin@orc1.example.com

/*
```

**Step 2 If the replication status is normal for the master group, then change the status to quiesced.**

```
*/

BEGIN
   DBMS_REPCAT.SUSPEND_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

/*
```

**Step 3 Remove the master site.**

```
*/

BEGIN
   DBMS_REPCAT.REMOVE_MASTER_DATABASES (
      gname => 'hr_repg',
      master_list => 'orc4.example.com');
END;
/
```

```
/*
```

You should wait until the DBA_REPCATLOG view is empty. Execute the following SELECT statement in another SQL*Plus session to monitor the DBA_REPCATLOG view:

```
SELECT * FROM DBA_REPCATLOG WHERE GNAME = 'HR_REPG';

*/

PAUSE Press <RETURN> to continue when DBA_REPCATLOG is empty for the master group.

/*
```

**Step 4 Resume master activity for the master group.**

```
*/

BEGIN
   DBMS_REPCAT.RESUME_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

SET ECHO OFF

SPOOL OFF

/************************ END OF SCRIPT *********************************/
```

## Removing an Unavailable Master Site

The sites being removed from a master group do not have to be accessible. When a master site will not be available for an extended period of time due to a system or network failure, you might decide to drop the master site from the master group.

However, because the site is unavailable, you most likely cannot suspend replication activity for the master group. You can use the REMOVE_MASTER_DATABASES procedure in the DBMS_REPCAT package to remove master sites from a master group, even if the master group is not quiesced.

If this is the case, you are responsible for:

- Cleaning the deferred transaction queue
- Removing any data inconsistencies

Specifically, the next time that you suspend replication activity for a master group, you must complete the following steps as soon as possible after the unavailable master sites are removed:

**Step 1 Suspend replication activity for the master group.**

See "SUSPEND_MASTER_ACTIVITY Procedure" on page 18-96 for information.

**Step 2 Delete all deferred transactions from each master site where the destination for the transaction is a removed master site.**

See "DELETE_TRAN Procedure" on page 14-9 for information.

**Step 3 Remove all deferred transactions from removed master sites.**

See "DELETE_TRAN Procedure" on page 14-9 for information.

**Step 4 Reexecute or delete all error transactions at each remaining master site.**

See "Managing the Error Queue" on page 9-15 for information about reexecuting error transactions, and see "DELETE_TRAN Procedure" on page 14-9 for information about removing error transactions.

**Step 5 Ensure that no deferred or error transactions exist at each remaining master.**

If you cannot remove one or more deferred transactions from a remaining master, execute the `DBMS_DEFER_SYS.DELETE_TRAN` procedure at the master site.

**Step 6 Ensure that all replicated data is consistent.**

See Chapter 16, "DBMS_RECTIFIER_DIFF" for information about determining and correcting differences.

**Step 7 Resume replication activity for the master group.**

See "RESUME_MASTER_ACTIVITY Procedure" on page 18-87 for information.

---

**Note:** After dropping an unavailable master site from a master group, you should also remove the master group from the dropped site to finish the cleanup.

---

## Updating the Comments Fields in Data Dictionary Views

Several procedures in the `DBMS_REPCAT` package enable you to update the comment information in the various data dictionary views associated with replication. Table 7–1 lists the appropriate procedure to call for each view.

***Table 7–1 Updating Comments in Advanced Replication Views***

| View | DBMS_REPCAT Procedure | See for Parameter Information |
|---|---|---|
| `DBA_REPGROUP` | `COMMENT_ON_REPGROUP(`<br>` gname IN VARCHAR2,`<br>` comment IN VARCHAR2)` | "COMMENT_ON_REPGROUP Procedure" on page 18-34. |
| `DBA_REPOBJECT` | `COMMENT_ON_REPOBJECT(`<br>` sname IN VARCHAR2,`<br>` oname IN VARCHAR2,`<br>` type IN VARCHAR2,`<br>` comment IN VARCHAR2)` | "COMMENT_ON_REPOBJECT Procedure" on page 18-35. |
| `DBA_REPSITES` | `COMMENT_ON_REPSITES(`<br>` gname IN VARCHAR2,`<br>` master IN VARCHAR,`<br>` comment IN VARCHAR2)` | "COMMENT_ON_REPSITES Procedure" on page 18-36. |
| `DBA_REPCOLUMN_GROUP` | `COMMENT_ON_COLUMN_GROUP(`<br>` sname IN VARCHAR2,`<br>` oname IN VARCHAR2,`<br>` column_group IN VARCHAR2,`<br>` comment IN VARCHAR2)` | "COMMENT_ON_COLUMN_GROUP Procedure" on page 18-31. |
| `DBA_REPPRIORITY_GROUP` | `COMMENT_ON_PRIORITY_GROUP(`<br>` gname IN VARCHAR2,`<br>` pgroup IN VARCHAR2)`<br>` comment IN VARCHAR2)` | "COMMENT_ON_PRIORITY_GROUP Procedures" on page 18-33. |

*Table 7–1 (Cont.) Updating Comments in Advanced Replication Views*

| View | DBMS_REPCAT Procedure | See for Parameter Information |
|---|---|---|
| DBA_REPPRIORITY_GROUP (site priority group) | COMMENT_ON_SITE_PRIORITY( gname IN VARCHAR2, name IN VARCHAR2, comment IN VARCHAR2) | "COMMENT_ON_PRIORITY_GROUP Procedures" on page 18-33. |
| DBA_REPRESOLUTION (uniqueness conflicts) | COMMENT_ON_UNIQUE_RESOLUTION( sname IN VARCHAR2, oname IN VARCHAR2, constraint_name IN VARCHAR2, sequence_no IN NUMBER, comment IN VARCHAR2) | The parameters for the COMMENT_ON_UNIQUE_RESOLUTION procedures are described in "COMMENT_ON_conflicttype_RESOLUTION Procedure" on page 18-38. |
| DBA_REPRESOLUTION (update conflicts) | COMMENT_ON_UPDATE_RESOLUTION( sname IN VARCHAR2, oname IN VARCHAR2, column_group IN VARCHAR2, sequence_no IN NUMBER, comment IN VARCHAR2) | The parameters for the COMMENT_ON_UNIQUE_RESOLUTION procedures are described in "COMMENT_ON_conflicttype_RESOLUTION Procedure" on page 18-38. |
| DBA_REPRESOLUTION (delete conflicts) | COMMENT_ON_DELETE_RESOLUTION( sname IN VARCHAR2, oname IN VARCHAR2, sequence_no IN NUMBER, comment IN VARCHAR2) | The parameters for the COMMENT_ON_UNIQUE_RESOLUTION procedures are described in "COMMENT_ON_conflicttype_RESOLUTION Procedure" on page 18-38. |

# Using Procedural Replication

Procedural replication can offer performance advantages for large batch-oriented operations operating on large numbers of rows that can be run serially within a replication environment.

A good example of an appropriate application is a purge operation, also referred to as an archive operation, that you run infrequently (for example, once in each quarter) during off hours to remove old data, or data that was "logically" deleted from the online database. An example using procedural replication to purge deleted rows is described in the "Avoiding Delete Conflicts" section in Chapter 5, "Conflict Resolution Concepts and Architecture", of *Oracle Database Advanced Replication*.

## Restrictions on Procedural Replication

All parameters for a replicated procedure must be IN parameters; OUT and IN/OUT modes are not supported. The following data types are supported for these parameters:

- VARCHAR2
- NVARCHAR2
- NUMBER
- DATE
- TIMESTAMP
- TIMESTAMP WITH TIME ZONE
- TIMESTAMP LOCAL TIME ZONE
- INTERVAL YEAR TO MONTH

- INTERVAL DAY TO SECOND
- RAW
- ROWID
- CHAR
- NCHAR
- CLOB with BASICFILE storage
- NCLOB with BASICFILE storage
- BLOB with BASICFILE storage
- XMLType stored as CLOB
- User-defined types that do not use type inheritance or type evolution
- Oracle-supplied types that do not use type inheritance or type evolution

The following data types are not supported for these parameters:

- FLOAT
- BINARY_FLOAT
- BINARY_DOUBLE
- LONG
- LONG RAW
- CLOB with SECUREFILE storage
- NCLOB with SECUREFILE storage
- BLOB with SECUREFILE storage
- BFILE
- XMLType stored object relationally or as binary XML
- Expression type
- User-defined types that use type inheritance or type evolution
- Oracle-supplied types that use type inheritance or type evolution

Oracle cannot detect update conflicts produced by replicated procedures. Replicated procedures must detect and resolve conflicts themselves. Because of the difficulties involved in writing your own conflict resolution routines, it is best to simply avoid the possibility of conflicts altogether.

Adhering to the following guidelines helps you ensure that your tables remain consistent at all sites when you plan to use procedural replication:

- You must disable row-level replication within the body of the deferred procedure. See "Updating the Comments Fields in Data Dictionary Views" on page 7-33.
- Run only one replicated procedure at a time, as described in "Serializing Transactions" on page 7-37.
- Propagate deferred transactions serially. For more information about guidelines for scheduled links, see *Oracle Database Advanced Replication*.
- The replicated procedure must be packaged and the package cannot contain any functions. Standalone deferred procedures and standalone or packaged deferred functions are not currently supported.

- The deferred procedures must reference only locally owned data.
- The procedures should not use locally generated fields, values, or environmentally dependent SQL functions. For example, the procedure should not call SYSDATE.
- Your data ownership should be statically partitioned. That is, ownership of a row should not change between sites.
- If you have multiple master groups at a master site, and one or more master groups are quiesced, then you cannot perform procedural replication on any master group at the master site. This restriction is enforced because a procedure in one master group can update objects in another master group. You can only perform procedural replication when all of the master groups on a master site are replicating data normally (that is, when none of the master groups is quiesced).

  For example, if you have a procedure named sal_raise in master group A on master site db1, then you cannot run the sal_raise procedure if master group B on master site db1 is quiesced, even if master group A is replicating normally.
- When using procedural replication, a procedure call is only propagated to master replication sites. The procedure call is not propagated to materialized view sites. However, procedural replication can be initiated at a materialized view site. In this case, the procedure call is propagated to all of the master sites in the replication environment, but the procedure call is not propagated to any other materialized view sites. Other materialized view sites must pull changes made at the master site by performing a materialized view refresh.

  For example, suppose a replication environment includes two master sites named msite1 and msite2 and two materialized view sites named mview1 and mview2. If procedural replication is initiated at mview1, then the procedure is run at mview1 and the procedure call is propagated to the two master sites, msite1 and msite2, where the procedure is also run. However, the procedure call is not propagated to mview2. Therefore, during the next refresh, mview2 pulls down all of the changes made by the procedure at its master site.

## User-Defined Types and Procedural Replication

When using procedural replication, the user-defined types and the objects referenced in the procedure must meet the following conditions:

- For an object type, all replication sites must agree about the order of attributes in the object type. You establish the attribute order when you create the object type. Consider the following object type:

```
CREATE TYPE cust_address_typ AS OBJECT
     (street_address     VARCHAR2(40),
      postal_code        VARCHAR2(10),
      city               VARCHAR2(30),
      state_province     VARCHAR2(10),
      country_id         CHAR(2));
/
```

  At all replication sites, street_address must be the first attribute, postal_code must be the second attribute, city must be the third attribute, and so on.
- For an Oracle object, all replication sites must have the same object identifier (OID), schema owner, and type name for each replicated object type.

  You can meet these conditions by always using distributed schema management to create or modify any replicated object, including object types, tables with

column objects, and object tables. If you do not use distributed schema management to create and modify object types, then replication errors can result.

> **See Also:** *Oracle Database Advanced Replication* for more information about type agreement at replication sites

## Serializing Transactions

Serial execution ensures that your data remains consistent. The replication facility propagates and executes replicated transactions one at a time. For example, assume that you have two procedures, A and B, that perform updates on local data. Now assume that you perform the following actions, in order:

1. Execute A and B locally.
2. Queue requests to execute other replicas of A and B on other nodes.
3. Commit.

The replicas of A and B on the other nodes are executed completely serially, in the same order that they were committed at the originating site. If A and B execute concurrently at the originating site, however, then they can produce different results locally than they do remotely. Executing A and B serially at the originating site ensures that all sites have identical results. Propagating the transaction serially ensures that A and B are executing in serial order at the target site in all cases.

Alternatively, you could write the procedures carefully, to ensure serialization. For example, you could use `SELECT... FOR UPDATE` for queries to ensure serialization at the originating site and at the target site if you are using parallel propagation.

## Generating Support for Replicated Procedures

You must disable row-level replication support at the start of your procedure, and then reenable support at the end. This operation ensures that any updates that occur because of executing the procedure are not propagated to other sites. Row-level replication is enabled and disabled by calling the following procedures, respectively:

- `DBMS_REPUTIL.REPLICATION_ON`
- `DBMS_REPUTIL.REPLICATION_OFF`

> **See Also:**
>
> - "Disabling Replication" on page 9-5
> - "REPLICATION_ON Procedure" on page 22-4
> - "REPLICATION_OFF Procedure" on page 22-3

When you generate replication support for your replicated package, Oracle creates a wrapper package in the schema of the replication propagator.

> **Note:** Unregistering the current propagator drops all existing generated wrappers in the propagator's schema. Replication support for wrapped stored procedures must be regenerated after you register a new propagator.

The wrapper package has the same name as the original package, but its name is prefixed with the string you supply when you generate replication support for the

procedure. If you do not supply a prefix, then Oracle uses the default prefix, defer_. The wrapper procedure has the same parameters as the original, along with two additional parameters: call_local and call_remote. These two CHAR parameters determine where the procedure is executed. When call_local is 'Y', the procedure is executed locally. When call_remote is 'Y', the procedure will ultimately be executed at all other master sites in the replication environment.

The remote procedures are called directly if you are propagating changes synchronously, or calls to these procedures are added to the deferred transaction queue if you are propagating changes asynchronously. By default, call_local is 'N', and call_remote is 'Y'.

Oracle generates replication support for a package in two phases. The first phase creates the package specification at all sites. Phase two generates the package body at all sites. These two phases are necessary to support synchronous replication.

For example, suppose you create the package emp_mgmt containing the procedure new_dept, which takes one argument, email. To replicate this package to all master sites in your system, you can use the Advanced Replication interface in Oracle Enterprise Manager to add the package to a master group and then generate replication support for the object. After completing these steps, an application can call procedure in the replicated package as follows:

```
BEGIN
defer_emp_mgmt.new_dept( email        => 'jones',
                         call_local   => 'Y',
                         call_remote  => 'Y');
END;
/
```

> **See Also:** The Advanced Replication interface's online Help for more information about managing master groups and replicated objects using the Advanced Replication interface in Oracle Enterprise Manager

As shown in Figure 7–4, the logic of the wrapper procedure ensures that the procedure is called at the local site and subsequently at all remote sites. The logic of the wrapper procedure also ensures that when the replicated procedure is called at the remote sites, call_remote is FALSE, ensuring that the procedure is not further propagated.

If you are operating in a mixed replication environment with static partitioning of data ownership (that is, if you are not preventing row-level replication), then Advanced Replication preserves the order of operations at the remote node, because both row-level and procedural replication use the same asynchronous queue.

***Figure 7–4 Asynchronous Procedural Replication***

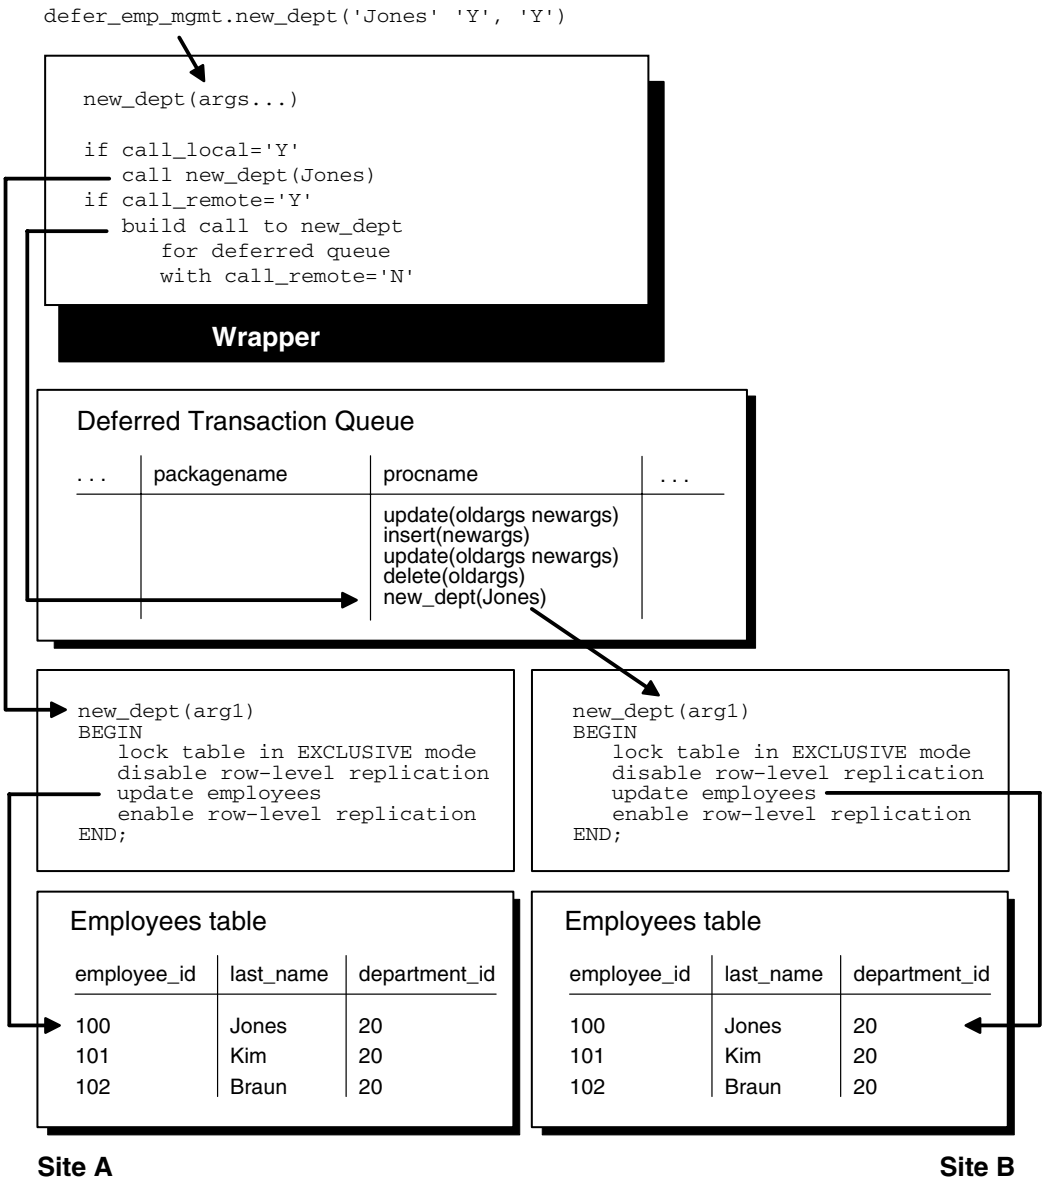

# 8

# Managing a Materialized View Replication Environment

Materialized view replication provides the flexibility to build data sets to meet the needs of your users and front-end applications, while still meeting the requirements of your security configuration. This chapter describes how to manage materialized view sites with the replication management API.

This chapter contains these topics:


- Refreshing Materialized Views
- Changing a Materialized View Group's Master Site
- Dropping Materialized View Groups and Objects
- Managing Materialized View Logs
- Performing an Offline Instantiation of a Materialized View Site Using Export/Import
- Using a Group Owner for a Materialized View Group


## Refreshing Materialized Views

Refreshing a materialized view synchronizes the data in the materialized view's master(s) and the data in the materialized view. You can either refresh all of the materialized views in a refresh group at once, or you can refresh materialized views individually. If you have applications that depend on multiple materialized views at a materialized view site, then Oracle recommends using refresh groups so that the data is transactionally consistent in all of the materialized views used by the application.

The following example refreshes the `hr_refg` refresh group:

```
EXECUTE DBMS_REFRESH.REFRESH ('hr_refg');
```

The following example refreshes the `hr.departments_mv` materialized view:

```
BEGIN
   DBMS_MVIEW.REFRESH (
     list   =>  'hr.departments_mv',
     method =>  '?');
END;
/
```

**Note:** Do not use the DBMS_MVIEW.REFRESH_ALL_MVIEWS or the DBMS_MVIEW.REFRESH_DEPENDENT procedure to refresh materialized views used in a replication environment. Instead, use the DBMS_REFRESH.REFRESH or the DBMS_MVIEW.REFRESH procedure to refresh materialized views in a replication environment.

**See Also:** *Oracle Database PL/SQL Packages and Types Reference* for more information about the DBMS_MVIEW package

## Changing a Materialized View Group's Master Site

To change the master site of a materialized view group at a level 1 materialized view site to another master site, call the SWITCH_MVIEW_MASTER procedure in the DBMS_REPCAT package, as shown in the following example:

```
BEGIN
   DBMS_REPCAT.SWITCH_MVIEW_MASTER (
      gname  => 'hr_repg',
      master => 'orc3.example.com');
END;
/
```

In this example, the master site for the hr_repg replication group is changed to the orc3.example.com master site. You must call this procedure at the materialized view site whose master site you want to change. The new database must be a master site in the replication environment. When you call this procedure, Oracle uses the new master to perform a full refresh of each materialized view in the local materialized view group. Ensure that you have set up the materialized view site to use the new master site before you run the SWITCH_MVIEW_MASTER procedure.

The entries in the SYS.SLOG$ table at the old master site for the switched materialized view are not removed. As a result, the materialized view log (MLOG$ table) of the switched updatable materialized view at the old master site has the potential to grow indefinitely, unless you purge it by calling DBMS_MVIEW.PURGE_LOG.

**Note:** You cannot switch the master of materialized views that are based on other materialized views (level 2 and greater materialized views). Such a materialized view must be dropped and re-created to base it on a different master.

**See Also:** "Setting Up Materialized View Sites" on page 2-16

## Dropping Materialized View Groups and Objects

You might need to drop replication activity at a materialized view site for several reasons. Perhaps the data requirements have changed or an employee has left the company.

This section contains the following sections:

- Dropping a Materialized View Group Created with a Deployment Template
- Dropping a Materialized View Group or Objects Created Manually
- Cleaning Up a Master Site or Master Materialized View Site

## Dropping a Materialized View Group Created with a Deployment Template

If a materialized view group was created with a deployment template, then, before you drop the materialized view group at the remote materialized view site, you must execute the DROP_SITE_INSTANTIATION procedure at the target master site of the materialized view group. In addition to removing the metadata relating to the materialized view group, this procedure also removes the related deployment template data regarding this site.

The DROP_SITE_INSTANTIATION procedure has a public and a private version. The public version allows the owner of the materialized view group to drop the materialized view site, while the private version allows the replication administrator to drop a materialized view site on behalf of the materialized view group owner.

### Using the Public Version of DROP_SITE_INSTANTIATION

Meet the following requirements to complete these actions:

**Executed As**:

- Materialized View Group Owner at Master Site
- Materialized View Administrator at Materialized View Site

**Executed At**:

- Master Site for Target Materialized View Site
- Materialized View Site

**Replication Status**: Normal

Complete the following steps to drop a materialized view group created with a deployment template.

> **Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

```
/************************ BEGINNING OF SCRIPT ******************************
```

**Step 1 Connect to the master site as the materialized view group owner.**

```
*/

SET ECHO ON

SPOOL drop_mv_group_public.out

CONNECT hr@orc3.example.com

/*
```

**Step 2 Drop the instantiated materialized view site from the master site.**

```
*/

BEGIN
   DBMS_REPCAT_INSTANTIATE.DROP_SITE_INSTANTIATION(
      refresh_template_name =>  'hr_refg_dt',
      site_name => 'mv4.example.com');
```

```
END;
/

/*
```

**Step 3 Connect to the remote materialized view site as the materialized view administrator.**

```
*/

CONNECT mviewadmin@mv4.example.com

/*
```

If you are not able to connect to the remote materialized view site, then the target materialized view group cannot refresh, but the existing data still remains at the materialized view site.

**Step 4 Drop the materialized view group.**

```
*/

BEGIN
   DBMS_REPCAT.DROP_MVIEW_REPGROUP (
      gname => 'hr_repg',
      drop_contents => TRUE);
END;
/

/*
```

To physically remove the contents of the materialized view group from the materialized view database, specify `TRUE` for the `drop_contents` parameter.

**Step 5 Remove the refresh group.**

Connect as the refresh group owner and remove the refresh group.

```
*/

CONNECT hr@mv4.example.com

BEGIN
   DBMS_REFRESH.DESTROY (
       name => 'hr_refg');
END;
/

SET ECHO OFF

SPOOL OFF

/************************ END OF SCRIPT *********************************/
```

## Using the Private Version of DROP_SITE_INSTANTIATION

The following steps are to be performed by the replication administrator on behalf of the materialized view group owner. Meet the following requirements to complete these actions:

**Executed As**:

- Replication Administrator at Master Site
- Materialized View Administrator at Materialized View Site

**Executed At**:

- Master Site for Target Materialized View Site
- Materialized View Site

**Replication Status**: Normal

Complete the following steps to drop a materialized view group created with a deployment template.

> **Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

```
/************************ BEGINNING OF SCRIPT ******************************
```

**Step 1 Connect to the master site as the replication administrator.**

```
*/

SET ECHO ON

SPOOL drop_mv_group_private.out

CONNECT repadmin@orc1.example.com

/*
```

**Step 2 Drop the instantiated materialized view site from the master site.**

```
*/

BEGIN
   DBMS_REPCAT_RGT.DROP_SITE_INSTANTIATION (
      refresh_template_name => 'hr_refg_dt',
      user_name             => 'hr',
      site_name             => 'mv4.example.com');
END;
/

/*
```

**Step 3 Connect to the remote materialized view site as the materialized view administrator.**

```
*/

CONNECT mviewadmin@mv4.example.com

/*
```

If you are unable to connect to the remote materialized view site, then the target materialized view group cannot refresh, but the existing data still remains at the materialized view site.

**Step 4 Drop the materialized view group.**

```
*/

BEGIN
   DBMS_REPCAT.DROP_MVIEW_REPGROUP (
      gname => 'hr_repg',
      drop_contents => TRUE,
      gowner => 'hr');
END;
/

/*
```

To physically remove the contents of the materialized view group from the materialized view database, specify `TRUE` for the `drop_contents` parameter.

**Step 5 Remove the refresh group.**

Connect as the refresh group owner and remove the refresh group.

```
*/

CONNECT hr@mv4.example.com

BEGIN
   DBMS_REFRESH.DESTROY (
       name => 'hr_refg');
END;
/

SET ECHO OFF

SPOOL OFF

/************************ END OF SCRIPT *********************************/
```

## Dropping a Materialized View Group or Objects Created Manually

The most secure method of removing replication support for a materialized view site is to physically drop the replicated objects or groups at the materialized view site. The following two sections describe how to drop these objects and groups while connected to the materialized view group.

Ideally, these procedures should be executed while the materialized view is connected to its target master site or master materialized view site. A connection ensures that any related metadata at the master site or master materialized view site is removed. If a connection to the master site or master materialized view site is not possible, then be sure to complete the procedure described in "Cleaning Up a Master Site or Master Materialized View Site" on page 8-8 to manually remove the related metadata.

### Dropping a Materialized View Group Created Manually

When it becomes necessary to remove a materialized view group from a materialized view site, use the `DROP_MVIEW_REPGROUP` procedure to drop a materialized view group. When you execute this procedure and are connected to the target master site or master materialized view site, the metadata for the target materialized view group at the master site or master materialized view site is removed. If you cannot connect, then see "Cleaning Up a Master Site or Master Materialized View Site" on page 8-8 for more information.

Meet the following requirements to complete these actions:

**Executed As**: Materialized View Administrator

**Executed At**: Remote Materialized View Site

**Replication Status**: N/A

Complete the following steps to drop a materialized view group at a materialized view site:

#### Step 1 Connect to the materialized view site as the materialized view administrator.

```
CONNECT mviewadmin@mv1.example.com
Enter password: password
```

#### Step 2 Drop the materialized view group.

```
BEGIN
   DBMS_REPCAT.DROP_MVIEW_REPGROUP (
      gname => 'hr_repg',
      drop_contents => TRUE);
END;
/
```

To physically remove the contents of the materialized view group from the materialized view database, specify `TRUE` for the `drop_contents` parameter.

### Dropping Objects at a Materialized View Site

When it becomes necessary to remove an individual materialized view from a materialized view site, use the `DROP_MVIEW_REPOBJECT` procedure API to drop a materialized view. When you execute this procedure and are connected to the target master site or master materialized view site, the metadata for the target materialized view at the master site or master materialized view site is removed. If you cannot connect, then see "Cleaning Up a Master Site or Master Materialized View Site" on page 8-8 for more information.

Meet the following requirements to complete these actions:

**Executed As**: Materialized View Administrator

**Executed At**: Remote Materialized View Site

**Replication Status**: N/A

Complete the following steps to drop an individual materialized view at a materialized view site.

#### Step 1 Connect to the materialized view site as the materialized view administrator.

```
CONNECT mviewadmin@mv1.example.com
Enter password: password
```

#### Step 2 Drop the materialized view.

```
BEGIN
   DBMS_REPCAT.DROP_MVIEW_REPOBJECT (
      sname => 'hr',
      oname => 'employees_mv1',
      type => 'SNAPSHOT',
      drop_objects => TRUE);
```

```
END;
/
```

To physically remove the contents of the materialized view from the materialized view database, specify `TRUE` for the `drop_contents` parameter.

## Cleaning Up a Master Site or Master Materialized View Site

If you are unable to drop a materialized view group or materialized view object while connected to the target master site or master materialized view site, then you must remove the related metadata at the master site or master materialized view site manually. Cleaning up the metadata also ensures that you are not needlessly maintaining master table or master materialized view changes to a materialized view log. The following sections describe how to clean up your master site or master materialized view site after dropping a materialized view group or object.

### Cleaning Up After Dropping a Materialized View Group

If you have executed the steps described in "Dropping a Materialized View Group Created Manually" on page 8-6 and were not connected to the master site or master materialized view site, then you are encouraged to complete the following steps to clean up the target master site or master materialized view site.

Meet the following requirements to complete these actions:

**Executed As**: Replication Administrator

**Executed At**: Master Site or Master Materialized View Site for Target Materialized View Site

**Replication Status**: Normal

Complete the following steps to clean up a master site or master materialized view site after dropping a materialized view group:

> **Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

```
/************************* BEGINNING OF SCRIPT ******************************
```

**Step 1 Connect to the master site or master materialized view site as the replication administrator.**

```
*/

SET ECHO ON

SPOOL cleanup_master1.out

CONNECT repadmin@orc1.example.com

/*
```

**Step 2 Unregister the materialized view groups.**

```
*/

BEGIN
```

```
   DBMS_REPCAT.UNREGISTER_MVIEW_REPGROUP (
      gname => 'hr_repg',
      mviewsite => 'mv1.example.com');
END;
/

/*
```

**Step 3 Purge the materialized view logs of the entries that were marked for the target materialized views.**

Execute the `PURGE_MVIEW_FROM_LOG` procedure for each materialized view that was in the materialized view groups you unregistered in Step 2.

> **Note:** If for some reason unregistering the materialized view group fails, then you should still complete this step.

> **See Also:** *Oracle Database PL/SQL Packages and Types Reference* for more information about the `DBMS_MVIEW` package

```
*/

BEGIN
   DBMS_MVIEW.PURGE_MVIEW_FROM_LOG (
      mviewowner => 'hr',
      mviewname => 'countries_mv1',
      mviewsite => 'mv1.example.com');
END;
/

BEGIN
   DBMS_MVIEW.PURGE_MVIEW_FROM_LOG (
      mviewowner => 'hr',
      mviewname => 'departments_mv1',
      mviewsite => 'mv1.example.com');
END;
/

BEGIN
   DBMS_MVIEW.PURGE_MVIEW_FROM_LOG (
      mviewowner => 'hr',
      mviewname => 'employees_mv1',
      mviewsite => 'mv1.example.com');
END;
/

BEGIN
   DBMS_MVIEW.PURGE_MVIEW_FROM_LOG (
      mviewowner => 'hr',
      mviewname => 'jobs_mv1',
      mviewsite => 'mv1.example.com');
END;
/

BEGIN
   DBMS_MVIEW.PURGE_MVIEW_FROM_LOG (
      mviewowner => 'hr',
      mviewname => 'job_history_mv1',
```

```
      mviewsite => 'mv1.example.com');
END;
/

BEGIN
   DBMS_MVIEW.PURGE_MVIEW_FROM_LOG (
      mviewowner => 'hr',
      mviewname => 'locations_mv1',
      mviewsite => 'mv1.example.com');
END;
/

BEGIN
   DBMS_MVIEW.PURGE_MVIEW_FROM_LOG (
      mviewowner => 'hr',
      mviewname => 'regions_mv1',
      mviewsite => 'mv1.example.com');
END;
/

SET ECHO OFF

SPOOL OFF

/************************ END OF SCRIPT *******************************/
```

### Cleaning Up Individual Materialized View Support

If you have executed the steps described in "Dropping Objects at a Materialized View Site" on page 8-7 and were not connected to the master site or master materialized view site, then you are encouraged to complete the following steps to clean up the target master site or master materialized view site.

Meet the following requirements to complete these actions:

**Executed As**: Replication Administrator

**Executed At**: Master Site or Master Materialized View Site for Target Materialized View Site

**Replication Status**: Normal

Complete the following steps to clean up a master site or master materialized view site after dropping an individual materialized view.

> **Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

```
/************************ BEGINNING OF SCRIPT ****************************
```

**Step 1 Connect to the master site or master materialized view site as the replication administrator.**

```
*/

SET ECHO ON

SPOOL cleanup_master2.out
```

```
CONNECT repadmin@orc1.example.com

/*
```

**Step 2 Unregister the materialized view.**

```
*/

BEGIN
   DBMS_MVIEW.UNREGISTER_MVIEW (
      mviewowner => 'hr',
      mviewname => 'employees_mv1',
      mviewsite => 'mv1.example.com');
END;
/

/*
```

> **See Also:** *Oracle Database PL/SQL Packages and Types Reference* for more information about the `DBMS_MVIEW` package

**Step 3 Purge the associated materialized view log of the entries that were marked for the target materialized views.**

> **Note:** If for some reason unregistering the materialized view fails, then you should still complete this step.

> **See Also:** *Oracle Database PL/SQL Packages and Types Reference* for more information about the `DBMS_MVIEW` package

```
*/

BEGIN
   DBMS_MVIEW.PURGE_MVIEW_FROM_LOG (
      mviewowner => 'hr',
      mviewname => 'employees_mv1',
      mviewsite => 'mv1.example.com');
END;
/

SET ECHO OFF

SPOOL OFF

/************************ END OF SCRIPT *********************************/
```

# Managing Materialized View Logs

The following sections explain how to manage materialized view logs:

- Altering Materialized View Logs
- Managing Materialized View Log Space
- Reorganizing Master Tables that Have Materialized View Logs
- Dropping a Materialized View Log

## Altering Materialized View Logs

After creating a materialized view log, you can alter its storage parameters and support for corresponding materialized views. The following sections explain more about altering materialized view logs. Only the following users can alter a materialized view log:

- The owner of the master table or master materialized view.
- A user with the `SELECT` privilege for the master table or master materialized view and `ALTER` privilege on the `MLOG$_`*master_name*, where *master_name* is the name of the master for the materialized view log. For example, if the master table is `employees`, then the materialized view log table name is `MLOG$_employees`.

### Altering Materialized View Log Storage Parameters

To alter a materialized view log's storage parameters, use the `ALTER MATERIALIZED VIEW LOG` statement. For example, the following statement alters a materialized view log on the `employees` table in the `hr` schema:

```
ALTER MATERIALIZED VIEW LOG ON hr.employees
  PCTFREE 25
  PCTUSED 40;
```

### Altering a Materialized View Log to Add Columns

To add new columns to a materialized view log, use the SQL statement `ALTER MATERIALIZED VIEW LOG`. For example, the following statement alters a materialized view log on the `customers` table in the `sales` schema:

```
ALTER MATERIALIZED VIEW LOG ON hr.employees
  ADD (department_id);
```

> **See Also:** *Oracle Database Advanced Replication* for more information about adding columns to a materialized view log

## Managing Materialized View Log Space

Oracle automatically tracks which rows in a materialized view log have been used during the refreshes of materialized views, and purges these rows from the log so that the log does not grow endlessly. Because multiple simple materialized views can use the same materialized view log, rows already used to refresh one materialized view might still be needed to refresh another materialized view. Oracle does not delete rows from the log until *all* materialized views have used them.

For example, suppose two materialized views were created against the `customers` table in a master site. Oracle refreshes the `customers` materialized view at the `spdb1` database. However, the server that manages the master table and associated materialized view log does not purge the materialized view log rows used during the refresh of this materialized view until the `customers` materialized view at the `spdb2` database also refreshes using these rows.

Because Oracle must wait for all dependent materialized views to refresh before purging rows from a materialized view log, unwanted situations can occur that cause a materialized view log to grow indefinitely when multiple materialized views are based on the same master table or master materialized view.

For example, such situations can occur when multiple materialized views are based on a master table or master materialized view and one of the following conditions is true:

- One materialized view is not configured for automatic refreshes and has not been manually refreshed for a long time.
- One materialized view has an infrequent refresh interval, such as every year (365 days).
- A network failure has prevented an automatic refresh of one or more of the materialized views based on the master table or master materialized view.
- A network or site failure has prevented a master table or master materialized view from becoming aware that a materialized view has been dropped.

> **Note:** If you purge or TRUNCATE a materialized view log before a materialized view has refreshed the changes that were deleted, then the materialized view must perform a complete refresh.

### Purging Rows from a Materialized View Log

Always try to keep a materialized view log as small as possible to minimize the database space that it uses. To remove rows from a materialized view log and make space for newer log records, you can perform one of the following actions:

- Refresh the materialized views associated with the log so that Oracle can purge rows from the materialized view log.
- Manually purge records in the log by deleting rows required only by the *n*th least recently refreshed materialized views.

To manually purge rows from a materialized view log, execute the PURGE_LOG procedure of the DBMS_MVIEW package at the database that contains the log. For example, to purge entries from the materialized view log of the customers table that are necessary only for the least recently refreshed materialized view, execute the following procedure:

```
BEGIN
   DBMS_MVIEW.PURGE_LOG (
      master => 'hr.employees',
      num    => 1,
      flag   => 'DELETE');
END;
/
```

Only the owner of a materialized view log or a user with the EXECUTE privilege for the DBMS_MVIEW package can purge rows from the materialized view log by executing the PURGE_LOG procedure.

> **See Also:** *Oracle Database PL/SQL Packages and Types Reference* for more information about the DBMS_MVIEW package

### Truncating a Materialized View Log

If a materialized view log grows and allocates many extents, then purging the log of rows does not reduce the amount of space allocated for the log. In such cases, you should truncate the materialized view log. Only the owner of a materialized view log or a user with the DELETE ANY TABLE system privilege can truncate a materialized view log.

To reduce the space allocated for a materialized view log by truncating it, complete the following steps:

**Step 1 Acquire an exclusive lock on the master table or master materialized view to prevent updates during the following process.**

For example, issue a statement similar to the following:

```
LOCK TABLE hr.employees IN EXCLUSIVE MODE;
```

**Step 2 Using a second database session, copy the rows in the materialized view log (in other words, the MLOG$ table) to a temporary table.**

For example, issue a statement similar to the following:

```
CREATE TABLE hr.templog AS SELECT * FROM hr.MLOG$_employees;
```

**Step 3 Using the second session, truncate the log using the SQL statement TRUNCATE TABLE.**

For example, issue a statement similar to the following:

```
TRUNCATE TABLE hr.MLOG$_employees;
```

**Step 4 Using the second session, reinsert the old rows.**

Perform this step so that you do not have to perform a complete refresh of the dependent materialized views.

For example, issue statements similar to the following:

```
INSERT INTO hr.MLOG$_employees SELECT * FROM hr.templog;

DROP TABLE hr.templog;
```

**Step 5 Using the first session, release the exclusive lock on the master table or master materialized view.**

You can accomplish this by performing a rollback:

```
ROLLBACK;
```

---

**Note:** Any changes made to the master table or master materialized view between the time you copy the rows to a new location and when you truncate the log do not appear until after you perform a complete refresh.

---

## Reorganizing Master Tables that Have Materialized View Logs

To improve performance and optimize disk use, you can periodically reorganize master tables. This section describes how to reorganize a master and preserve the fast refresh capability of associated materialized views.

---

**Note:** These sections do not discuss online redefinition of tables. Online redefinition is not allowed on master tables with materialized view logs, master materialized views, or materialized views. Online redefinition is allowed only on master tables that do not have materialized view logs. See the *Oracle Database Administrator's Guide* for more information about online redefinition of tables.

---

### Reorganization Notification

When you reorganize a table, any ROWID information of the materialized view log must be invalidated. Oracle detects a table reorganization automatically only if the table is truncated as part of the reorganization.

If the table is not truncated, then Oracle must be notified of the table reorganization. To support table reorganizations, two procedures in the DBMS_MVIEW package, BEGIN_TABLE_REORGANIZATION and END_TABLE_REORGANIZATION, notify Oracle that the specified table has been reorganized. The procedures perform clean-up operations, verify the integrity of the logs and triggers that the fast refresh mechanism needs, and invalidate the ROWID information in the table's materialized view log. The inputs are the owner and name of the master to be reorganized. There is no output.

> **See Also:** "Method 2 for Reorganizing Table employees" on page 8-16

### Truncating Masters

When a table is truncated, its materialized view log is also truncated. However, for primary key materialized views, you can preserve the materialized view log, allowing fast refreshes to continue. Although the information stored in a materialized view log is preserved, the materialized view log becomes invalid for rowids when the master is truncated. The rowid information in the materialized view log will seem to be newly created and cannot be used by rowid materialized views for fast refresh.

The PRESERVE MATERIALIZED VIEW LOG option is the default. Therefore, if you specify the PRESERVE MATERIALIZED VIEW LOG option or no option, then the information in the master's materialized view log is preserved, but current rowid materialized views can use the log for a fast refresh only *after* a complete refresh has been performed.

> **Note:** To ensure that any previously fast refreshable materialized view is still refreshable, follow the guidelines in "Methods of Reorganizing a Database Table" on page 8-15.

If the PURGE MATERIALIZED VIEW LOG option is specified, then the materialized view log is purged along with the master.

**Examples** Either of the following two statements preserves materialized view log information when the master table named employees is truncated:

```
TRUNCATE TABLE hr.employees PRESERVE MATERIALIZED VIEW LOG;
TRUNCATE TABLE hr.employees;
```

The following statement truncates the materialized view log along with the master table:

```
TRUNCATE TABLE hr.employees PURGE MATERIALIZED VIEW LOG;
```

### Methods of Reorganizing a Database Table

Oracle provides four table reorganization methods that preserve the capability for fast refresh. These appear in the following sections. Other reorganization methods require an initial complete refresh to enable subsequent fast refreshes.

**Note:** Do *not* use Direct Loader during a reorganization of a master. Direct Loader can cause reordering of the columns, which could invalidate the log information used in subquery and LOB materialized views.

**Method 1 for Reorganizing Table employees** Complete the following steps:

1. Call `DBMS_MVIEW.BEGIN_TABLE_REORGANIZATION` for table employees.
2. Rename table employees to employees_old.
3. Create table employees as SELECT * FROM employees_old.
4. Call `DBMS_MVIEW.END_TABLE_REORGANIZATION` for new table employees.

**Caution:** When a table is renamed, its associated PL/SQL triggers are also adjusted to the new name of the table.

Ensure that no transaction is issued against the reorganized table between calling `BEGIN_TABLE_REORGANIZATION` and `END_TABLE_REORGANIZATION`.

**Method 2 for Reorganizing Table employees** Complete the following steps:

1. Call `DBMS_MVIEW.BEGIN_TABLE_REORGANIZATION` for table employees.
2. Export table employees.
3. Truncate table employees with `PRESERVE MATERIALIZED VIEW LOG` option.
4. Import table employees using conventional path.
5. Call `DBMS_MVIEW.END_TABLE_REORGANIZATION` for new table employees.

**Caution:** When you truncate masters as part of a reorganization, you must use the `PRESERVE MATERIALIZED VIEW LOG` clause of the truncate table DDL.

Ensure that no transaction is issued against the reorganized table between calling `BEGIN_TABLE_REORGANIZATION` and `END_TABLE_REORGANIZATION`.

**Method 3 for Reorganizing Table employees** Complete the following steps:

1. Call `DBMS_MVIEW.BEGIN_TABLE_REORGANIZATION` for table employees.
2. Export table employees.
3. Rename table employees to employees_old.
4. Import table employees using conventional path.
5. Call `DBMS_MVIEW.END_TABLE_REORGANIZATION` for new table employees.

**Caution:** When a table is renamed, its associated PL/SQL triggers are also adjusted to the new name of the table.

Ensure that no transaction is issued against the reorganized table between calling `BEGIN_TABLE_REORGANIZATION` and `END_TABLE_REORGANIZATION`.

**Method 4 for Reorganizing Table employees** Complete the following steps:

1. Call `DBMS_MVIEW.BEGIN_TABLE_REORGANIZATION` for table employees.
2. Select contents of table employees to a flat file.
3. Rename table employees to employees_old.
4. Create table employees with the same shape as employees_old.
5. Run SQL*Loader using conventional path.
6. Call `DBMS_MVIEW.END_TABLE_REORGANIZATION` for new table employees.

> **Caution:** When a table is renamed, its associated PL/SQL triggers are also adjusted to the new name of the table.

Ensure that no transaction is issued against the reorganized table between calling `BEGIN_TABLE_REORGANIZATION` and `END_TABLE_REORGANIZATION`.

> **See Also:** *Oracle Database PL/SQL Packages and Types Reference* for more information about the `DBMS_MVIEW` package

## Dropping a Materialized View Log

You can delete a materialized view log regardless of its master or any existing materialized views. For example, you might decide to drop a materialized view log if one of the following conditions is true:

- All materialized views of a master have been dropped.
- All materialized views of a master are to be refreshed using complete refresh, not fast refresh.
- A master no longer supports materialized views that require fast refreshes.

Here, a master can be a master table or a master materialized view. To delete a materialized view log, execute the `DROP MATERIALIZED VIEW LOG` statement in SQL*Plus. For example, the following statement deletes the materialized view log for a table named `customers` in the `sales` schema:

```
DROP MATERIALIZED VIEW LOG ON hr.employees;
```

Only the owner of the master or a user with the `DROP ANY TABLE` system privilege can drop a materialized view log.

# Performing an Offline Instantiation of a Materialized View Site Using Export/Import

Adding a new materialized view site to your replication environment can cause network traffic. The network traffic is caused by propagating the entire contents of tables or materialized views through the network to the new materialized view site.

To minimize such network traffic, you can add a new materialized view site using offline instantiation procedure. With offline instantiation, you can create a new materialized view group at a materialized view site. Offline instantiation uses of Oracle's Export and Import utilities, which allow you to create an export file and transfer the data to the new site through a storage medium, such as CD-ROM, tape, and so on. Offline instantiation is especially useful for materialized views, because the target computer could be a laptop using a modem connection.

The following script performs an offline instantiation for a new materialized view group at a new materialized view site. The materialized view group is based on an existing master group at a master site. Meet the following requirements to complete these actions:

**Executed As**:

- Replication Administrator at Master Site
- Materialized View Administrator at New Materialized View Site

**Executed At**:

- Master Site for Target Materialized View Site
- New Materialized View Site

**Replication Status**: Normal

**Materialized View Site**:

- Set up materialized view site. In this example, the materialized view site is `mview.example.com` and the master site is `orc1.example.com`.
- Ensure that the appropriate schema has been created before the offline instantiation of your materialized view site.
- Create proxy users at the master site if they do not exist.

**See Also:**

- For more information about setting up a master site and creating proxy users at a master site, see "Setting Up Master Sites" on page 2-3
- For more information about setting up materialized view sites, see "Setting Up Materialized View Sites" on page 2-16

Complete the following steps to set up a materialized view site named `mview.example.com`.

**Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

```
/************************ BEGINNING OF SCRIPT ********************************
```

**Step 1 Connect to the master site as the replication administrator.**

```
*/

SET ECHO ON

SPOOL offline.out

CONNECT repadmin@orc1.example.com

/*
```

**Step 2 Create the necessary materialized view logs, if they do not exist.**

If materialized view logs do not exist for the relevant master tables, then create them at the master site.

```
*/

CREATE MATERIALIZED VIEW LOG ON hr.countries;
CREATE MATERIALIZED VIEW LOG ON hr.departments;
CREATE MATERIALIZED VIEW LOG ON hr.employees;
CREATE MATERIALIZED VIEW LOG ON hr.jobs;
CREATE MATERIALIZED VIEW LOG ON hr.job_history;
CREATE MATERIALIZED VIEW LOG ON hr.locations;
CREATE MATERIALIZED VIEW LOG ON hr.regions;

/*
```

**Step 3 Create a temporary schema at the master site for the materialized views.**

To prepare materialized views for export, you must create the schema that contains the replicated objects.

In this example, create a temporary schema `temp_schema`.

```
*/

CONNECT system@orc1.example.com

CREATE TABLESPACE offline_mview
 DATAFILE 'offline_mview.dbf' SIZE 10M AUTOEXTEND ON
 EXTENT MANAGEMENT LOCAL AUTOALLOCATE;

CREATE TEMPORARY TABLESPACE offline_temp_mview
 TEMPFILE 'offline_temp_mview.dbf' SIZE 5M AUTOEXTEND ON;

ACCEPT password PROMPT 'Enter password for user: ' HIDE

CREATE USER temp_schema IDENTIFIED BY &password;

ALTER USER temp_schema DEFAULT TABLESPACE offline_mview
              QUOTA UNLIMITED ON offline_mview;

ALTER USER temp_schema TEMPORARY TABLESPACE offline_temp_mview;

GRANT ALTER SESSION, CREATE CLUSTER, CREATE DATABASE LINK, CREATE SEQUENCE,
  CREATE SESSION, CREATE SYNONYM, CREATE TABLE, CREATE VIEW, CREATE INDEXTYPE,
  CREATE OPERATOR, CREATE PROCEDURE, CREATE TRIGGER, CREATE TYPE,
  CREATE MATERIALIZED VIEW, SELECT ANY TABLE
TO temp_schema;

CONNECT temp_schema@orc1.example.com;

/*
```

**Step 4 Create temporary materialized views at the master site in the separate schema you created in Step 3.**

These materialized views contain the data that you transfer to your new materialized view site using the Export utility.

**Note:** Ensure that the SELECT statements include the database link. In this example, the database link is orc1.example.com.

```
*/

CREATE MATERIALIZED VIEW temp_schema.countries
      REFRESH FAST WITH PRIMARY KEY FOR UPDATE AS SELECT *
      FROM hr.countries@orc1.example.com;

CREATE MATERIALIZED VIEW temp_schema.departments
      REFRESH FAST WITH PRIMARY KEY FOR UPDATE AS SELECT *
      FROM hr.departments@orc1.example.com;

CREATE MATERIALIZED VIEW temp_schema.employees
      REFRESH FAST WITH PRIMARY KEY FOR UPDATE AS SELECT *
      FROM hr.employees@orc1.example.com;

CREATE MATERIALIZED VIEW temp_schema.jobs
      REFRESH FAST WITH PRIMARY KEY FOR UPDATE AS SELECT *
      FROM hr.jobs@orc1.example.com;

CREATE MATERIALIZED VIEW temp_schema.job_history
      REFRESH FAST WITH PRIMARY KEY FOR UPDATE AS SELECT *
      FROM hr.job_history@orc1.example.com;

CREATE MATERIALIZED VIEW temp_schema.locations
      REFRESH FAST WITH PRIMARY KEY FOR UPDATE AS SELECT *
      FROM hr.locations@orc1.example.com;

CREATE MATERIALIZED VIEW temp_schema.regions
      REFRESH FAST WITH PRIMARY KEY FOR UPDATE AS SELECT *
      FROM hr.regions@orc1.example.com;

/*
```

### Step 5 Create a directory object at each database.

Each database involved in this operation must have a directory object to hold the Data Pump dump file, and the user who will perform the export or import must have READ and WRITE privileges on this directory object. In this example, a Data Pump export is performed at the master site, and a Data Pump import is performed at the materialized view site.

While connected in SQL*Plus to a database as an administrative user who can create directory objects using the SQL statement CREATE DIRECTORY, create a directory object to hold the Data Pump dump file and log files. For example:

```
*/

CONNECT system@orc1.example.com

CREATE DIRECTORY DPUMP_DIR AS '/usr/dpump_dir';

CONNECT system@mview.example.com

CREATE DIRECTORY DPUMP_DIR AS '/usr/dpump_dir';

/*
```

Ensure that you complete these actions at both databases involved in the operation. In this example, SYSTEM user creates the directory objects and performs all exports and imports. If a user who does not own the directory object will perform the export or import, then grant the user READ and WRITE privileges on the directory object.

**Step 6 Perform a schema-level export of the schema containing the materialized views.**

On a command line, perform the export that will contain all data and metadata for the materialized views. This example connects as the SYSTEM user. The following is an example Data Pump export command:

```
expdp system SCHEMAS=temp_schema DIRECTORY=DPUMP_DIR
DUMPFILE=temp_schema.dmp
```

> **See Also:** *Oracle Database Utilities* for information about performing a Data Pump export

```
*/

PAUSE Press <RETURN> to continue when the export is complete.

/*
```

**Step 7 Connect to the new materialized view site as SYSTEM user.**

```
*/

CONNECT system@mview.example.com

/*
```

**Step 8 Drop the hr User**

This example creates the materialized views in the hr schema at the materialized view site. This schema is created when Oracle is installed. This step drops the schema, but the schema will be re-created and populated with materialized views later in this example.

```
*/

DROP USER hr CASCADE;

/*
```

**Step 9 Create necessary schema and database link at the materialized view site, if they do not exist.**

Before you perform the offline instantiation of your materialized views, create the schema that will contain the materialized views at the new materialized view site and the database link from the materialized view site to the master site. The materialized views must be in the same schema that contains the master objects at the master site. If the schema exists, then grant the necessary privileges and create the database link.

```
*/

CREATE TABLESPACE demo_mview
 DATAFILE 'demo_mview.dbf' SIZE 10M AUTOEXTEND ON
 EXTENT MANAGEMENT LOCAL AUTOALLOCATE;
```

```
CREATE TEMPORARY TABLESPACE temp_mview
 TEMPFILE 'temp_mview.dbf' SIZE 5M AUTOEXTEND ON;

CREATE USER hr IDENTIFIED BY &password;

ALTER USER hr DEFAULT TABLESPACE demo_mview
              QUOTA UNLIMITED ON demo_mview;

ALTER USER hr TEMPORARY TABLESPACE temp_mview;

GRANT
  CREATE SESSION,
  CREATE TABLE,
  CREATE PROCEDURE,
  CREATE SEQUENCE,
  CREATE TRIGGER,
  CREATE VIEW,
  CREATE SYNONYM,
  ALTER SESSION,
  CREATE MATERIALIZED VIEW,
  ALTER ANY MATERIALIZED VIEW,
  CREATE DATABASE LINK
TO hr;

CONNECT hr@mview.example.com

CREATE DATABASE LINK orc1.example.com CONNECT TO hr IDENTIFIED by &password;

/*
```

**Step 10 Connect to the new materialized view site as the materialized view administrator.**

```
*/

CONNECT mviewadmin@mview.example.com

/*
```

**Step 11 Create an empty materialized view group.**

Run the `DBMS_REPCAT.CREATE_MVIEW_REPGROUP` procedure at the new materialized view site to create an empty materialized view group into which you will add your materialized views.

```
*/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPGROUP (
      gname            => 'hr_repg',
      master           => 'orc1.example.com',
      propagation_mode => 'ASYNCHRONOUS');
END;
/

/*
```

**Step 12 Create an empty refresh group.**

All materialized views that are added to a particular refresh group are refreshed at the same time. This ensures transactional consistency between the related materialized views in the refresh group.

```
*/

BEGIN
   DBMS_REFRESH.MAKE (
      name => 'mviewadmin.hr_refg',
      list => '',
      next_date => SYSDATE,
      interval => 'SYSDATE + 1/24',
      implicit_destroy => FALSE,
      rollback_seg => '',
      push_deferred_rpc => TRUE,
      refresh_after_errors => FALSE);
END;
/

/*
```

**Step 13 In a separate terminal window, transfer the export dump file to the new materialized view site.**

Using the `DBMS_FILE_TRANSFER` package, FTP or some other method, transfer the export dump file to the new materialized view site.

```
*/

PAUSE Press <RETURN> to continue after transferring the dump file.

/*
```

**Step 14 In a separate terminal window, import the materialized views to the owner at the new materialized view site.**

On a command line, perform the import of the file that you exported in Step 5. This example connects as the `SYSTEM` user.

If you use Data Pump, then ensure that you import your data using the `REMAP_SCHEMA` parameter to import the data from the temporary user you created at the master site to the owner of the materialized views at the materialized view site. In this example, the temporary user at the master site is `temp_schema` and the materialized view owner at the materialized view site is `hr`.

Also, if you use Data Pump, then you can use the `REMAP_TABLESPACE` parameter if the tablespace is different at the master site and the materialized view site. In this example, the tablespace at the master site is `offline_mview` (created in Step 3) and the tablespace at the materialized view site is `demo_mview` (created in Step 9).

The following is an example import command:

```
impdp system DIRECTORY=DPUMP_DIR DUMPFILE=temp_schema.dmp
REMAP_SCHEMA=temp_schema:hr REMAP_TABLESPACE=offline_mview:demo_mview
```

Only users with the `DBA` role or the `IMP_FULL_DATABASE` role can import using the `REMAP_SCHEMA` parameter.

> **See Also:** *Oracle Database Utilities* for information about performing a Data Pump import

```
*/

PAUSE Press <RETURN> to continue when the import is complete.

/*
```

### Step 15 Add materialized views to the materialized view group.

Execute the DBMS_REPCAT.CREATE_MVIEW_REPOBJECT procedure to add the materialized views to the materialized view group you created in Step 9.

```
*/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname             => 'hr_repg',
      sname             => 'hr',
      oname             => 'countries',
      type              => 'SNAPSHOT',
      min_communication => TRUE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname             => 'hr_repg',
      sname             => 'hr',
      oname             => 'departments',
      type              => 'SNAPSHOT',
      min_communication => TRUE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname             => 'hr_repg',
      sname             => 'hr',
      oname             => 'employees',
      type              => 'SNAPSHOT',
      min_communication => TRUE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname             => 'hr_repg',
      sname             => 'hr',
      oname             => 'jobs',
      type              => 'SNAPSHOT',
      min_communication => TRUE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname             => 'hr_repg',
      sname             => 'hr',
      oname             => 'job_history',
      type              => 'SNAPSHOT',
      min_communication => TRUE);
END;
```

```
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname             => 'hr_repg',
      sname             => 'hr',
      oname             => 'locations',
      type              => 'SNAPSHOT',
      min_communication => TRUE);
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname             => 'hr_repg',
      sname             => 'hr',
      oname             => 'regions',
      type              => 'SNAPSHOT',
      min_communication => TRUE);
END;
/

/*
```

**Step 16 Add the materialized views to the refresh group.**

All of the materialized view group objects that you add to the refresh group are refreshed at the same time to preserve referential integrity between related materialized views. Execute the `DBMS_REFRESH.ADD` procedure to add the materialized views to the refresh group you created in Step 12.

```
*/

BEGIN
   DBMS_REFRESH.ADD (
      name => 'mviewadmin.hr_refg',
      list => 'hr.countries',
      lax => TRUE);
END;
/

BEGIN
   DBMS_REFRESH.ADD (
      name => 'mviewadmin.hr_refg',
      list => 'hr.departments',
      lax => TRUE);
END;
/

BEGIN
   DBMS_REFRESH.ADD (
      name => 'mviewadmin.hr_refg',
      list => 'hr.employees',
      lax => TRUE);
END;
/

BEGIN
   DBMS_REFRESH.ADD (
      name => 'mviewadmin.hr_refg',
      list => 'hr.jobs',
```

```
      lax => TRUE);
END;
/

BEGIN
   DBMS_REFRESH.ADD (
      name => 'mviewadmin.hr_refg',
      list => 'hr.job_history',
      lax => TRUE);
END;
/

BEGIN
   DBMS_REFRESH.ADD (
      name => 'mviewadmin.hr_refg',
      list => 'hr.locations',
      lax => TRUE);
END;
/

BEGIN
   DBMS_REFRESH.ADD (
      name => 'mviewadmin.hr_refg',
      list => 'hr.regions',
      lax => TRUE);
END;
/

/*
```

**Step 17 Refresh materialized views to register them at master site.**

In addition to retrieving the latest changes from the master tables, refreshing the materialized views at the new materialized view site registers the offline instantiated materialized views at the target master site.

```
*/

EXECUTE DBMS_REFRESH.REFRESH ('hr_refg');

/*
```

**Step 18 Connect to the master site as SYSTEM user.**

```
*/

CONNECT system@orc1.example.com

/*
```

**Step 19 Drop the temporary schema to delete the temporary materialized views you created in Step 4 at the master site.**

```
*/

DROP USER temp_schema CASCADE;

SET ECHO OFF

SPOOL OFF

/************************ END OF SCRIPT *********************************/
```

# Using a Group Owner for a Materialized View Group

Specifying a group owner when you define a new materialized view group and its related objects enables you to create multiple materialized view groups based on the same replication group at a single materialized view site. Also, specifying group owners enables you to create multiple materialized view groups that are based on the same replication group at a master site or master materialized view site. You accomplish this by creating the materialized view groups under different schemas at the materialized view site. This example uses the schemas `bob` and `jane` as group owners and assumes that these schemas exist at the materialized view site.

**Executed As**:

- Materialized View Administrator at New Materialized View Site

**Executed At**:

- Materialized View Site

**Replication Status**: Normal

**Materialized View Site**:

- Set up materialized view site. In this example, the materialized view site is `mv1.example.com` and the master site is `orc1.example.com`.
- Create proxy users at the master site if they do not exist.
- Create materialized view logs for the tables in the `hr` schema at the master site if they do not exist.

Complete the following steps to use a group owner.

---

**Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

---

**See Also:**

- *Oracle Database Advanced Replication* for a complete description of using group owners and the advantages of using multiple data sets
- For more information about setting up a master site and creating proxy users at a master site, see "Setting Up Master Sites" on page 2-3
- For more information about setting up materialized view sites, see "Setting Up Materialized View Sites" on page 2-16

```
/************************ BEGINNING OF SCRIPT ******************************
```

### Step 1 Create a database link from the hr schema to the master site

Before building your materialized view group, you must ensure that the replicated schema exists at the remote materialized view site and that the necessary database links have been created.

In this example, if the `hr` schema does not exist, then create the schema. If the `hr` schema exists at the materialized view site, then grant any necessary privileges.

```
*/

CONNECT system@mv1.example.com

CREATE TABLESPACE demo_mv1
 DATAFILE 'demo_mv1.dbf' SIZE 10M AUTOEXTEND ON
 EXTENT MANAGEMENT LOCAL AUTOALLOCATE;

CREATE TEMPORARY TABLESPACE temp_mv1
 TEMPFILE 'temp_mv1.dbf' SIZE 5M AUTOEXTEND ON;

ACCEPT password PROMPT 'Enter password for user: ' HIDE

CREATE USER hr IDENTIFIED BY &password;

ALTER USER hr DEFAULT TABLESPACE demo_mv1
              QUOTA UNLIMITED ON demo_mv1;

ALTER USER hr TEMPORARY TABLESPACE temp_mv1;

GRANT
  CREATE SESSION,
  CREATE TABLE,
  CREATE PROCEDURE,
  CREATE SEQUENCE,
  CREATE TRIGGER,
  CREATE VIEW,
  CREATE SYNONYM,
  ALTER SESSION,
  CREATE MATERIALIZED VIEW,
  ALTER ANY MATERIALIZED VIEW,
  CREATE DATABASE LINK
 TO hr;

/*
```

If it does not exist, then create the database link for the replicated schema.

Before building your materialized view group, you must ensure that the necessary database links exist for the replicated schema. The owner of the materialized views needs a database link pointing to the `proxy_refresher` that was created when the master site was set up.

```
*/

SET ECHO ON

SPOOL mv_group_owner.out

CONNECT hr@mv1.example.com

CREATE DATABASE LINK orc1.example.com
   CONNECT TO proxy_refresher IDENTIFIED BY &password;

/*
```

**Step 2 Connect to the materialized view site as the materialized view administrator.**

```
*/
```

```
CONNECT mviewadmin@mv1.example.com

/*
```

### Step 3 Create materialized view group with group owner (gowner) bob using the CREATE_MVIEW_REPGROUP procedure.

The replication group that you specify in the `gname` parameter must match the name of the replication group that you are replicating at the target master site or master materialized view site. The `gowner` parameter enables you to specify an additional identifier that lets you create multiple materialized view groups based on the same replication group at the same materialized view site.

In this example, materialized view groups are created for the group owners `bob` and `jane`, and these two materialized view groups are based on the same replication group.

```
*/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPGROUP (
      gname => 'hr_repg',
      master => 'orc1.example.com',
      propagation_mode => 'ASYNCHRONOUS',
      gowner => 'bob');
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPGROUP (
      gname => 'hr_repg',
      master => 'orc1.example.com',
      propagation_mode => 'ASYNCHRONOUS',
      gowner => 'jane');
END;
/

/*
```

### Step 4 Create the materialized views owned by bob.

The `gowner` value used when creating your materialized view objects must match the `gowner` value specified when you created the materialized view group in the previous procedures. After creating the materialized view groups, you can create materialized views based on the same master in the `hr_repg` materialized view group owned by `bob` and `jane`. This example assumes that these users exist.

> **Caution:** Each object must have a unique name. When using a `gowner` to create multiple materialized view groups, duplicate object names could become a problem. To avoid any object-naming conflicts, you might want to append the `gowner` value to the end of the object name that you create, as illustrated in the following procedures (that is, `create materialized view hr.countries_bob`). Such a naming method ensures that you do not create any objects with conflicting names.

Whenever you create a materialized view, always specify the schema name of the table owner in the query for the materialized view. In the following examples, `hr` is specified as the owner of the table in each query.

```
*/

CREATE MATERIALIZED VIEW hr.countries_bob
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
  AS SELECT * FROM hr.countries@orc1.example.com;

CREATE MATERIALIZED VIEW hr.departments_bob
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
  AS SELECT * FROM hr.departments@orc1.example.com;

CREATE MATERIALIZED VIEW hr.employees_bob
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
  AS SELECT * FROM hr.employees@orc1.example.com;

CREATE MATERIALIZED VIEW hr.jobs_bob
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
  AS SELECT * FROM hr.jobs@orc1.example.com;

CREATE MATERIALIZED VIEW hr.job_history_bob
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
  AS SELECT * FROM hr.job_history@orc1.example.com;

CREATE MATERIALIZED VIEW hr.locations_bob
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
  AS SELECT * FROM hr.locations@orc1.example.com;

CREATE MATERIALIZED VIEW hr.regions_bob
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
  AS SELECT * FROM hr.regions@orc1.example.com;

/*
```

**Step 5 Create the materialized views owned by jane.**

```
*/

CREATE MATERIALIZED VIEW hr.departments_jane
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
  AS SELECT * FROM hr.departments@orc1.example.com;

CREATE MATERIALIZED VIEW hr.employees_jane
  REFRESH FAST WITH PRIMARY KEY FOR UPDATE
  AS SELECT * FROM hr.employees@orc1.example.com;

/*
```

**Step 6 Add the materialized views owned by bob to the materialized view group.**

```
*/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'countries_bob',
      type => 'SNAPSHOT',
      min_communication => TRUE,
```

```
      gowner => 'bob');
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'departments_bob',
      type => 'SNAPSHOT',
      min_communication => TRUE,
      gowner => 'bob');
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'employees_bob',
      type => 'SNAPSHOT',
      min_communication => TRUE,
      gowner => 'bob');
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'jobs_bob',
      type => 'SNAPSHOT',
      min_communication => TRUE,
      gowner => 'bob');
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'job_history_bob',
      type => 'SNAPSHOT',
      min_communication => TRUE,
      gowner => 'bob');
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'locations_bob',
      type => 'SNAPSHOT',
      min_communication => TRUE,
      gowner => 'bob');
END;
/

BEGIN
```

```
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'regions_bob',
      type => 'SNAPSHOT',
      min_communication => TRUE,
      gowner => 'bob');
END;
/

/*
```

**Step 7 Add the materialized views owned by jane to the materialized view group.**

```
*/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'departments_jane',
      type => 'SNAPSHOT',
      min_communication => TRUE,
      gowner => 'jane');
END;
/

BEGIN
   DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
      gname => 'hr_repg',
      sname => 'hr',
      oname => 'employees_jane',
      type => 'SNAPSHOT',
      min_communication => TRUE,
      gowner => 'jane');
END;
/

SET ECHO OFF

SPOOL OFF

/*
```

**Step 8 Add your materialized views to a refresh group.**

> **See Also:** Chapter 5, "Creating a Materialized View Group" (Step 6) for more information about adding materialized views to a refresh group

```
/************************ END OF SCRIPT *********************************/
```

# 9

# Managing Replication Objects and Queues

This chapter illustrates how to manage the replication objects and queues in your replication environment using the replication management API.

This chapter contains these topics:


- Altering a Replicated Object in a Quiesced Master Group
- Modifying Tables without Replicating the Modifications
- Converting a LONG Column to a LOB Column in a Replicated Table
- Determining Differences Between Replicated Tables
- Managing the Deferred Transactions Queue
- Managing the Error Queue


## Altering a Replicated Object in a Quiesced Master Group

As your database needs change, you might need to modify the characteristics of your replicated objects. It is important that you do not directly execute DDL to alter your replicated objects. Doing so might cause your replication environment to fail.

Use the `ALTER_MASTER_REPOBJECT` procedure in the `DBMS_REPCAT` package to alter the characteristics of your replicated objects in a quiesced master group. From the example following, notice that you simply include the necessary DDL within the procedure call (see the `ddl_text` parameter).

Meet the following requirements to complete these actions:

**Executed As**: Replication Administrator

**Executed At**: Master Definition Site

**Replication Status**: Quiesced

Complete the following steps to alter a replicated object in a quiesced master group.

> **Note:** If you are viewing this document online, then you can copy the text from the "BEGINNING OF SCRIPT" line after this note to the "END OF SCRIPT" line into a text editor and then edit the text to create a script for your environment.

```
/************************ BEGINNING OF SCRIPT ****************************
```

**Step 1 If you are altering a master table, and there are updatable materialized views based on the master table, then refresh the updatable materialized views.**

Refresh the updatable materialized views to push any changes from them to the master table. See "Refreshing Materialized Views" on page 8-1 for instructions.

```
*/

PAUSE Press <RETURN> to continue when all of the updatable materialized views that
are based on the master table have been refreshed.

/*
```

**Step 2 Connect to the master definition site as the replication administrator.**

```
*/

SET ECHO ON

SPOOL alter_rep_object.out

CONNECT repadmin@orc1.example.com

/*
```

**Step 3 If necessary, then quiesce the master group.**

See the "ALTER_MASTER_REPOBJECT Procedure" on page 18-23 for information about when quiesce is not required.

```
*/

BEGIN
   DBMS_REPCAT.SUSPEND_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

/*
```

**Step 4 In a separate SQL*Plus session, check the status of the master group you are quiescing.**

Do not proceed until the group's status is `QUIESCED`.

To check the status, run the following query:

```
SELECT GNAME, STATUS FROM DBA_REPGROUP;

*/

PAUSE Press <RETURN> to continue when the master group's status is QUIESCED.

/*
```

**Step 5 Alter the replicated object.**

```
*/

BEGIN
   DBMS_REPCAT.ALTER_MASTER_REPOBJECT (
      sname => 'hr',
      oname => 'employees',
```

```
      type => 'TABLE',
      ddl_text => 'ALTER TABLE hr.employees ADD (timestamp DATE)');
END;
/

/*
```

**Step 6 Regenerate replication support for the altered object.**

```
*/

BEGIN
    DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT (
      sname => 'hr',
      oname => 'employees',
      type => 'TABLE',
      min_communication => TRUE);
END;
/

/*
```

**Step 7 In a separate SQL*Plus session, check if DBA_REPCATLOG is empty.**

Do not proceed until this view is empty.

Execute the following `SELECT` statement in another SQL*Plus session to monitor the `DBA_REPCATLOG` view:

```
SELECT * FROM DBA_REPCATLOG WHERE GNAME = 'HR_REPG';

*/

PAUSE Press <RETURN> to continue when DBA_REPCATLOG is empty.

/*
```

**Step 8 Re-create the Materialized View Log for a Master Table, If Necessary**

If you altered a master table with fast-refreshable materialized views based on it, then the materialized view log must be dropped and re-created if any of the following conditions are met:

- The materialized view log must contain one or more new columns that were added to the master table. In this case, you can choose to alter the materialized view log to add the columns. See "Altering a Materialized View Log to Add Columns" on page 8-12.
- The materialized view log contains one or more columns in the master table that were altered.
- The materialized view log contains one more columns that were deleted from the master table.

```
*/

PAUSE Press <RETURN> to continue after the materialized view logs are re-created.

/*
```

#### Step 9 Re-create Materialized Views, If Necessary

If you altered a master table with updatable materialized views based on it, then all of these updatable materialized views must be re-created.

If you altered a master table with read-only materialized views based on it, then these read-only materialized views must be re-created if any of the following conditions are met:

- The read-only materialized view must reference one or more columns that were added to the master table.
- The read-only materialized view references one or more columns in the master table that were altered.
- The read-only materialized view references one or more columns that were deleted from the master table.

```
*/

PAUSE Press <RETURN> to continue after the materialized views are re-created.

/*
```

#### Step 10 Resume replication activity.

```
*/

BEGIN
   DBMS_REPCAT.RESUME_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/

SET ECHO OFF

SPOOL OFF

/************************* END OF SCRIPT *********************************/
```

## Modifying Tables without Replicating the Modifications

You might have a situation in which you must modify a replicated object, but you do not want this modification replicated to the other sites in the replication environment. For example, you might want to disable replication in the following situations:

- When you are using procedural replication to propagate a change, always disable row-level replication at the start of your procedure.
- You might need to disable replication in triggers defined on replicated tables to avoid replicating trigger actions multiple times. See "Ensuring that Replicated Triggers Fire Only Once" on page 9-6.
- Sometimes when you manually resolve a conflict, you might not want to replicate this modification to the other copies of the table.

You might need to do this, for example, if you must correct the state of a record at one site so that a conflicting replicated update will succeed when you reexecute the error transaction. Or, you might use an unreplicated modification to undo the effects of a transaction at its origin site because the transaction could not be applied at the destination site. In this example, you can use the Advanced Replication interface in

Oracle Enterprise Manager to delete the conflicting transaction from the destination site.

To modify tables without replicating the modifications, use the REPLICATION_ON and REPLICATION_OFF procedures in the DBMS_REPUTIL package. These procedures take no arguments and are used as flags by the generated replication triggers.

> **Note:** To enable and disable replication, you must have the EXECUTE privilege on the DBMS_REPUTIL package.

## Disabling Replication

The DBMS_REPUTIL.REPLICATION_OFF procedure sets the state of an internal replication variable for the current session to FALSE. Because all replicated triggers check the state of this variable before queuing any transactions, modifications made to the replicated tables that use row-level replication do not result in any queued deferred transactions.

> **Caution:** Turning replication on or off affects only the current session. That is, other users currently connected to the same server are not restricted from placing committed changes in the deferred transaction queue.

If you are using procedural replication, then call REPLICATION_OFF at the start of your procedure, as shown in the following example. This ensures that the replication facility does not attempt to use row-level replication to propagate the changes that you make.

```
CREATE OR REPLACE PACKAGE update_objects AS
  PROCEDURE update_emp(adjustment IN NUMBER);
END;
/

CREATE OR REPLACE PACKAGE BODY update_objects AS
  PROCEDURE update_emp(adjustment IN NUMBER) IS
  BEGIN
   --turn off row-level replication for set update
   DBMS_REPUTIL.REPLICATION_OFF;
   UPDATE emp . . .;
   --reenable replication
   DBMS_REPUTIL.REPLICATION_ON;
  EXCEPTION WHEN OTHERS THEN
   . . .
   DBMS_REPUTIL.REPLICATION_ON;
  END;
END;
/
```

## Reenabling Replication

After resolving any conflicts, or at the end of your replicated procedure, be certain to call DBMS_REPUTIL.REPLICATION_ON to resume normal replication of changes to your replicated tables or materialized views. This procedure takes no arguments. Calling REPLICATION_ON sets the internal replication variable to TRUE.

## Ensuring that Replicated Triggers Fire Only Once

If you have defined a replicated trigger on a replicated table, then you might need to ensure that the trigger fires only once for each change that you make. Typically, you only want the trigger to fire when the change is first made, and you do not want the remote trigger to fire when the change is replicated to the remote site.

You should check the value of the `DBMS_REPUTIL.FROM_REMOTE` package variable at the start of your trigger. The trigger should update the table only if the value of this variable is `FALSE`.

Alternatively, you can disable replication at the start of the trigger and reenable it at the end of the trigger when modifying rows other than the one that caused the trigger to fire. Using this method, only the original change is replicated to the remote sites. Then the replicated trigger fires at each remote site. Any updates performed by the replicated trigger are not pushed to any other sites.

Using this approach, conflict resolution is not invoked. Therefore, you must ensure that the changes resulting from the trigger do not affect the consistency of the data.

# Converting a LONG Column to a LOB Column in a Replicated Table

LOB columns using `BASICFILE` storage can be replicated, but `LONG` columns cannot be replicated. You can convert the data type of a `LONG` column to a `CLOB` column and the data type of a `LONG_RAW` column to a `BLOB` column.

Converting a `LONG` column to a LOB column can result in increased network bandwidth requirements because the data in such a column is replicated after conversion. Ensure that you have adequate network bandwidth before completing the procedure in this section.

---

**Note:** LOB columns using `SECUREFILE` storage cannot be replicated.

---

**See Also:** *Oracle Database SecureFiles and Large Objects Developer's Guide* for more information about applications and `LONG` to LOB conversion

Complete the following steps to convert a `LONG` column to a LOB column in a replicated table:

### Step 1 Ensure that the data in the LONG column is consistent at all replication sites.

If a table containing a `LONG` column is configured as a master table, then Oracle does not replicate changes to the data in the `LONG` column. Therefore, the data in the `LONG` column might not match at all of your replication sites. You must ensure that the data in the `LONG` column matches at all master sites before proceeding.

### Step 2 Connect to the master definition site as the replication administrator.

```
CONNECT repadmin@orc1.example.com
```

### Step 3 If the replication status is normal, then change the status to quiesced.

```
BEGIN
   DBMS_REPCAT.SUSPEND_MASTER_ACTIVITY (
      gname => 'sales_mg');
END;
```

```
/
```

**Step 4 Convert the LONG column to a LOB column.**

```
BEGIN
   DBMS_REPCAT.ALTER_MASTER_REPOBJECT (
      sname => 'staff',
      oname => 'positions',
      type => 'TABLE',
      ddl_text => 'ALTER TABLE staff.positions MODIFY (job_desc CLOB)');
END;
/
```

A `LONG_RAW` column can be converted to a `BLOB` column using a similar `ALTER TABLE` statement.

**Step 5 Regenerate replication support for the altered master table.**

```
BEGIN
    DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT (
      sname => 'staff',
      oname => 'positions',
      type => 'TABLE',
      min_communication => TRUE);
END;
/
```

**Step 6 Resume replication.**

```
BEGIN
   DBMS_REPCAT.RESUME_MASTER_ACTIVITY (
      gname => 'sales_mg');
END;
/
```

**Step 7 If materialized views are based on the altered table at any of the master sites, then rebuild these materialized views.**

Rebuild materialized views if necessary.

# Determining Differences Between Replicated Tables

It is possible for the differences to arise in replicated tables. When administering a replication environment, you might want to check, periodically, whether the contents of two replicated tables are identical. The following procedures in the `DBMS_RECTIFIER_DIFF` package let you identify, and optionally rectify, the differences between two tables.

---

**Note:** You can also determine differences between database objects and converge them using the `DBMS_COMPARISON` package.

---

**See Also:**

- *Oracle Database PL/SQL Packages and Types Reference* for information about the `DBMS_COMPARISON` package
- *Oracle Database 2 Day + Data Replication and Integration Guide* and *Oracle Streams Replication Administrator's Guide* for information about using the `DBMS_COMPARISON` package

## Using the DIFFERENCES Procedure

The DIFFERENCES procedure compares two replicas of a table, and determines all rows in the first replica that are not in the second and all rows in the second that are not in the first. The output of this procedure is stored in two user-created tables. The first table stores the values of the missing rows, and the second table is used to indicate which site contains each row.

## Using the RECTIFY Procedure

The RECTIFY procedure uses the information generated by the DIFFERENCES procedure to rectify the two tables. Any rows found in the first table and not in the second are inserted into the second table. Any rows found in the second table and not in the first are deleted from the second table.

To restore equivalency between all copies of a replicated table, complete the following steps:

**Step 1 Select one copy of the table to be the "reference" table.**

This copy will be used to update all other replicas of the table as needed.

**Step 2 Determine if it is necessary to check all rows and columns in the table for differences, or only a subset.**

For example, it might not be necessary to check rows that have not been updated since the last time that you checked for differences. Although it is not necessary to check all columns, your column list must include all columns that comprise the primary key (or that you designated as a substitute identity key) for the table.

**Step 3 After determining which columns you will be checking in the table, create two tables to hold the results of the comparison.**

You must create one table that can hold the data for the columns being compared. For example, if you decide to compare the employee_id, salary, and department_id columns of the employees table, then your CREATE statement must be similar to the following:

```
CREATE TABLE hr.missing_rows_data (
  employee_id     NUMBER(6),
  salary          NUMBER(8,2),
  department_id   NUMBER(4));
```

You must also create a table that indicates where the row is found. This table must contain three columns with the data types shown in the following example:

```
CREATE TABLE hr.missing_rows_location (
  present     VARCHAR2(128),
  absent      VARCHAR2(128),
  r_id        ROWID);
```

**Step 4 Suspend replication activity for the replication group containing the tables that you want to compare.**

Although suspending replication activity for the group is not a requirement, rectifying tables that were not quiesced first can result in inconsistencies in your data.

```
CONNECT repadmin

BEGIN
   DBMS_REPCAT.SUSPEND_MASTER_ACTIVITY (
      gname => 'hr_repg');
```

```
END;
/
```

**Step 5 At the site containing the "reference" table, call the DIFFERENCES procedure.**

For example, if you wanted to compare the `employees` tables at the New York and San Francisco sites, then your procedure call would look similar to the following:

```
BEGIN
   DBMS_RECTIFIER_DIFF.DIFFERENCES (
      sname1              =>   'hr',
      oname1              =>   'employees',
      reference_site      =>   'ny.example.com',
      sname2              =>   'hr',
      oname2              =>   'employees',
      comparison_site     =>   'mv4.example.com',
      where_clause        =>   '',
      column_list         =>   'employee_id,salary,department_id',
      missing_rows_sname  =>   'hr',
      missing_rows_oname1 =>   'missing_rows_data',
      missing_rows_oname2 =>   'missing_rows_location',
      missing_rows_site   =>   'ny.example.com',
      max_missing         =>    500,
      commit_rows         =>    50);
END;
/
```

Figure 9–1 shows an example of two replicas of the `employees` table and what the resulting missing rows tables would look like if you executed the `DIFFERENCES` procedure on these replicas.

***Figure 9–1 Determining Differences Between Replicas***

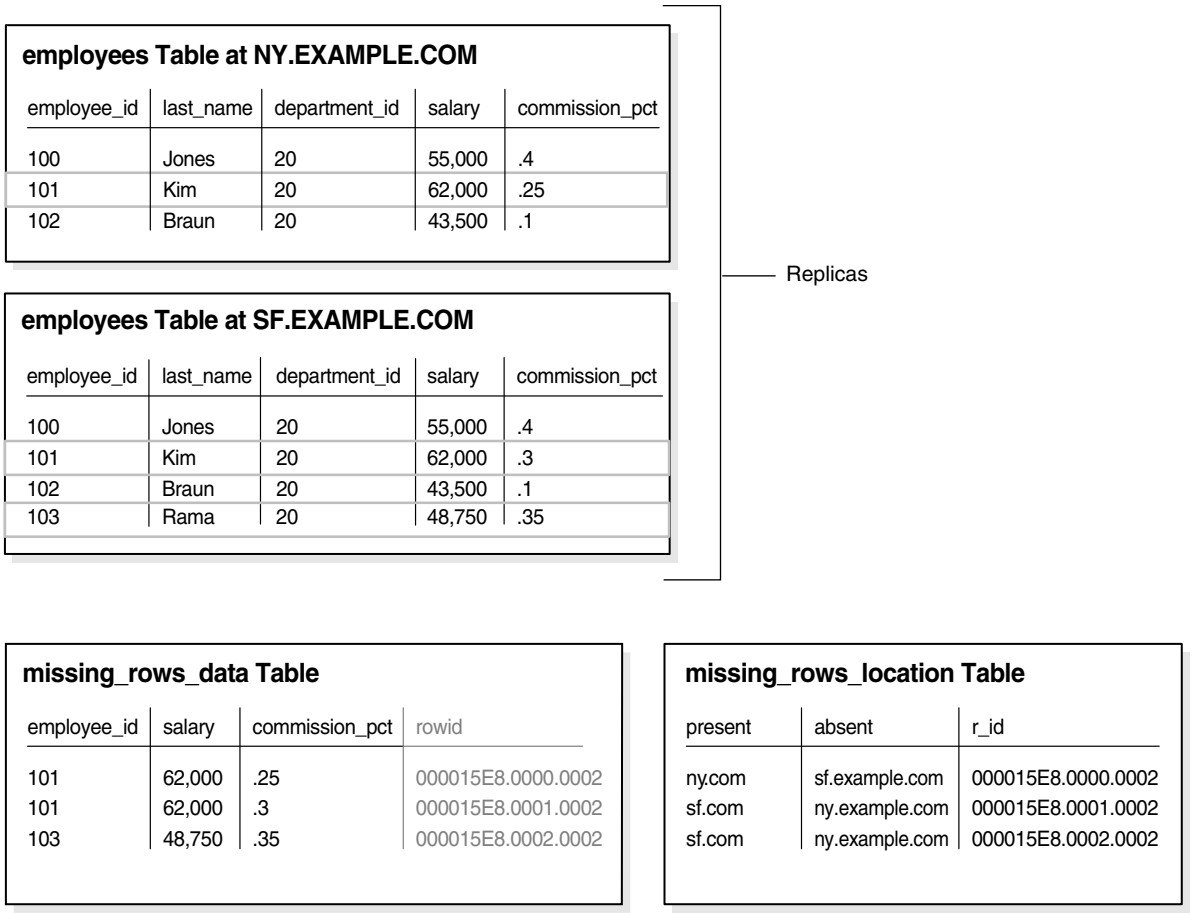


Notice that the two missing rows tables are related by the `ROWID` and `r_id` columns.

**Step 6 Rectify the table at the "comparison" site to be equivalent to the table at the "reference" site.**

```
BEGIN
   DBMS_RECTIFIER_DIFF.RECTIFY (
      sname1              =>   'hr',
      oname1              =>   'employees',
      reference_site      =>   'ny.example.com',
      sname2              =>   'hr',
      oname2              =>   'employees',
      comparison_site     =>   'mv4.example.com',
      column_list         =>   'employee_id,salary,department_id',
      missing_rows_sname  =>   'hr',
      missing_rows_oname1 =>   'missing_rows_data',
      missing_rows_oname2 =>   'missing_rows_location',
      missing_rows_site   =>   'ny.example.com',
      commit_rows         =>    50);
END;
/
```

The `RECTIFY` procedure temporarily disables replication at the "comparison" site while it performs the necessary insertions and deletions, as you would not want to propagate these changes. `RECTIFY` first performs all of the necessary `DELETE` operations

and then performs all of the INSERT operations. This ensures that there are no violations of a PRIMARY KEY constraint.

After you have successfully executed the RECTIFY procedure, your missing rows tables should be empty.

---

**Caution:** If you have any additional constraints on the "comparison" table, then you must ensure that they are not violated when you call RECTIFY. You might need to update the table directly using the information in the missing rows table. If so, then be sure to DELETE the appropriate rows from the missing rows tables.

---

#### Step 7 Repeat Steps 5 and 6 for the remaining copies of the replicated table.

Remember to use the same "reference" table each time to ensure that all copies are identical when you complete this procedure.

#### Step 8 Resume replication activity for the master group.

```
BEGIN
   DBMS_REPCAT.RESUME_MASTER_ACTIVITY (
      gname => 'hr_repg');
END;
/
```

## Managing the Deferred Transactions Queue

Typically, Advanced Replication is configured to push and purge the deferred transaction queue automatically. At times, however, you might need to push or purge the deferred transaction queue manually. The process for pushing the deferred transaction queue is the same at master sites and materialized view sites.

### Pushing the Deferred Transaction Queue

Master sites are configured to push the deferred transaction queue automatically at set intervals. At materialized view sites, if you do not automatically propagate the transactions in your deferred transaction queue during the refresh of your materialized view, then you must complete the following steps to propagate changes made to the updatable materialized view to its master table or master materialized view.

This example illustrates pushing the deferred transaction queue at a materialized view site, but the process is the same at master sites and materialized view sites.

**Executed As**: Materialized View Administrator

**Executed At**: Materialized View Site

Complete the following steps:

#### Step 1 Connect to the materialized view site as the materialized view administrator.

```
CONNECT mviewadmin@mv1.example.com
```

#### Step 2 Execute the following SELECT statement to view the deferred transactions and their destinations.

Propagation of the deferred transaction queue is based on the destination of the transaction. Each distinct destination and the number of transactions pending for the destination will be displayed.

```
SELECT DISTINCT(dblink), COUNT(deferred_tran_id)
   FROM deftrandest GROUP BY dblink;
```

#### Step 3 Execute the DBMS_DEFER_SYS.PUSH function for each site that is listed as a destination for a deferred transaction.

```
DECLARE
   temp INTEGER;
BEGIN
   temp := DBMS_DEFER_SYS.PUSH (
      destination => 'orc1.example.com',
      stop_on_error => FALSE,
      delay_seconds => 0,
      parallelism => 0);
END;
/
```

Run the `PUSH` procedure for each destination that was returned in the `SELECT` statement you ran in Step 2.

## Purging the Deferred Transaction Queue

If your system is not set to automatically purge the successfully propagated transactions in your deferred transaction queue periodically, then you must complete the following steps to purge them manually.

This example illustrates purging the deferred transaction queue at a materialized view site, but the process is the same at master sites and materialized view sites.

**Executed As**: Materialized View Administrator

**Executed At**: Materialized View Site

Complete the following steps:

#### Step 1 Connect to the materialized view site as the materialized view administrator.

```
CONNECT mviewadmin@mv1.example.com
```

#### Step 2 Purge the deferred transaction queue.

```
DECLARE
   temp INTEGER;
BEGIN
   temp := DBMS_DEFER_SYS.PURGE (
      purge_method => DBMS_DEFER_SYS.PURGE_METHOD_QUICK);
END;
/
```

**Note:** If you use the purge_method_quick parameter, deferred transactions and deferred procedure calls that have been successfully pushed can remain in the DEFTRAN and DEFCALL data dictionary views for longer than expected before they are purged. See the "Usage Notes" for DBMS_DEFER_SYS.PURGE on page 14-15 for details.

## Using the ANYDATA Type to Determine the Value of an Argument in a Deferred Call

If you are using column objects, collections, or REFs in a replicated table, then you can use the GET_ANYDATA_ARG function in the DBMS_DEFER_QUERY package to determine the value of an argument in a deferred call that involves one of these user-defined types.

The following example illustrates how to use the GET_ANYDATA_ARG function. This example uses the following user-defined types in the oe sample schema.

```
CREATE TYPE phone_list_typ AS VARRAY(5) OF VARCHAR2(25);
/

CREATE TYPE warehouse_typ AS OBJECT
    (warehouse_id       NUMBER(3),
     warehouse_name     VARCHAR2(35),
     location_id        NUMBER(4)
    );
/

CREATE TYPE inventory_typ AS OBJECT
    (product_id          NUMBER(6),
     warehouse           warehouse_typ,
     quantity_on_hand    NUMBER(8)
    );
/

CREATE TYPE inventory_list_typ AS TABLE OF inventory_typ;
/
```

The following procedure retrieves the argument value for collection, object, and REF instances of calls stored in the deferred transactions queue. This procedure assumes that the call number and transaction id are available.

The user who creates the procedure must have EXECUTE privilege on the DBMS_DEFER_QUERY package and must have CREATE PROCEDURE privilege. This example uses the oe sample schema. Therefore, to run the example, you must grant the oe user these privileges. Connect as an administrative user and enter the following:

```
GRANT EXECUTE ON DBMS_DEFER_QUERY TO oe;

GRANT CREATE PROCEDURE TO oe;

CONNECT oe@orc1.example.com

CREATE OR REPLACE PROCEDURE get_userdef_arg AS
  call_no      NUMBER := 0;
  txn_id       VARCHAR2(128) := 'xx.xx.xx';
  anydata_val  ANYDATA;
  t            ANYTYPE;
  data_pl      phone_list_typ;     -- varray
  data_ntt     inventory_list_typ; -- nested table type
  data_p       warehouse_typ;      -- object type
```

```
  ref1         REF inventory_typ;  -- REF type
  rval         PLS_INTEGER;        -- return value
  tc           PLS_INTEGER;        -- return value
  prec         PLS_INTEGER;        -- precision
  scale        PLS_INTEGER;        -- scale
  len          PLS_INTEGER;        -- length
  csid         PLS_INTEGER;        -- character set id
  csfrm        PLS_INTEGER;        -- character set form
  cnt          PLS_INTEGER;        -- count of varray elements or number of
                                   -- object attributes
  sname        VARCHAR2(35);       -- schema name
  type_name    VARCHAR2(35);       -- type name
  version      VARCHAR2(35);
BEGIN
   FOR i IN 1 .. 5 LOOP
     anydata_val := DBMS_DEFER_QUERY.GET_ANYDATA_ARG(call_no, i, txn_id);
     -- Get the type information, including type name.
     tc := anydata_val.GetType(t);
     tc := t.GetInfo(prec, scale, len, csid, csfrm, sname, type_name,
                     version, cnt);
     -- Based on the type name, convert the anydata value to the appropriate
     -- user-defined types.
     IF type_name = 'PHONE_LIST_TYP'
 THEN
       -- The anydata_val contains phone_list_typ varray instance.
       rval := anydata_val.GetCollection(data_pl);
       -- Do something with data_pl.
     ELSIF type_name = 'INVENTORY_LIST_TYP'
 THEN
       -- anydata_val contains inventory_list_typ nested table instance.
       rval := anydata_val.GetCollection(data_ntt);
       -- Do something with data_ntt.
     ELSIF type_name = 'WAREHOUSE_TYP'
 THEN
       -- The anydata_val contains warehouse_typ object instance.
       rval := anydata_val.GetObject(data_p);
       -- Do something with data_p.
     ELSIF type_name = 'INVENTORY_TYP'
 THEN
       -- The anydata_val contains a reference to inventory_typ object instance.
       rval := anydata_val.GetRef(ref1);
       -- Do something with ref1.
     END IF;
   END LOOP;
END;
/
```

**See Also:**

- "GET_datatype_ARG Function" on page 13-8
- *Oracle Database SQL Language Reference*, *Oracle Database Object-Relational Developer's Guide*, and *Oracle Database PL/SQL Packages and Types Reference* for more information about the `ANYDATA` data type

# Managing the Error Queue

As an administrator of a replication environment, you should regularly monitor the error queue to determine if any deferred transactions were not successfully applied at the target master site.

To check the error queue, issue the following SELECT statement (as the replication administrator) when connected to the target master site:

```
SELECT * FROM deferror;
```

If the error queue contains errors, then you should resolve the error condition and reexecute the deferred transaction. You have two options when reexecuting a deferred transaction: you can reexecute in the security context of the user who received the deferred transaction, or you can reexecute the deferred transaction with an alternate security context.

---

**Caution:** If you have multiple error transactions and you want to ensure that they are reexecuted in the correct order, then you can specify NULL for the deferred_tran_id parameter in the procedures in the following sections. If you do not specify NULL, then reexecuting individual transactions in the wrong order can cause conflicts.

---

## Reexecuting Error Transaction as the Receiver

The following procedure reexecutes a specified deferred transaction in the security context of the user who received the deferred transaction. This procedure should not be executed until the error situation has been resolved.

Meet the following requirements to complete these actions:

**Executed As**: Replication Administrator

**Executed At**: Site Containing Errors

**Replication Status**: Normal

Complete the following steps:

### Step 1 In SQL*Plus, connect to the master site as the replication administrator.

See *Oracle Database Administrator's Guide* for information about connecting to a database in SQL*Plus.

### Step 2 Reexecute the error transaction.

```
BEGIN
   DBMS_DEFER_SYS.EXECUTE_ERROR (
      deferred_tran_id => '1.12.2904',
      destination => 'orc2.example.com');
END;
/
```

## Reexecuting Error Transaction as Alternate User

The following procedure reexecutes a specified deferred transaction in the security context of the currently connected user. This procedure should not be executed until the error situation has been resolved.

Meet the following requirements to complete these actions:

**Executed As**: Connected User

**Executed At**: Site Containing Errors

**Replication Status**: Normal

Complete the following steps:

**Step 1 In SQL*Plus, connect to the master site as the alternate user.**

See *Oracle Database Administrator's Guide* for information about connecting to a database in SQL*Plus.

**Step 2 Reexecute the error transaction.**

```
BEGIN
   DBMS_DEFER_SYS.EXECUTE_ERROR_AS_USER (
      deferred_tran_id => '1.12.2904',
      destination => 'orc2.example.com');
END;
/
```

# 10

# Monitoring a Replication Environment

This chapter illustrates how to monitor a replication a replication environment using the data dictionary.

This chapter contains these topics:

- Monitoring Master Replication Environments
- Monitoring Materialized View Sites
- Monitoring Administrative Requests
- Monitoring the Deferred Transactions Queue
- Monitoring the Error Queue
- Monitoring Performance in a Replication Environment

---

**Note:** The Advanced Replication interface in Oracle Enterprise Manager is also an excellent way to monitor a replication environment. Most of the information obtained by the queries in this chapter can be found by using the Advanced Replication interface. See the Advanced Replication interface online Help for more information.

---

**See Also:** Part IV, "Replication Data Dictionary Reference"

## Monitoring Master Replication Environments

This section contains queries that you can run to display information about a master replication environment. The replication environment can be a multimaster environment, a master materialized view environment, or a hybrid environment that includes multiple master sites and materialized views.

This section contains the following topics:

- Monitoring Master Sites
- Monitoring Master Groups
- Monitoring Masters

### Monitoring Master Sites

This section contains queries that you can run to display information about master sites.

### Listing General Information About a Master Site

You can find the following general information about a master site by running the query in this section:

- The number of administrative requests.
- The number of administrative request errors.
- The number of unpropagated deferred transaction-destination pairs. Each deferred transaction can have multiple destinations to which it will be propagated, and each destination is a single deferred transaction-destination pair.

For example, if there are ten deferred transactions and each one must be propagated to three sites, then there are 30 deferred transaction-pairs returned by this query. After some time, if the first deferred transaction is propagated to two of the three destination sites, then there are still ten deferred transactions, but there are two fewer deferred-transaction pairs, and this query returns 28 unpropagated deferred transaction-pairs. In this case, the first deferred transaction only has one transaction-pair remaining.

- The number of deferred transaction errors (error transactions).
- The number of successfully propagated transactions that are still in the queue. These transactions should be purged from the queue.

Run the following query to list this information for the current master site:

```
COLUMN GLOBAL_NAME HEADING 'Database' FORMAT A25
COLUMN ADMIN_REQUESTS HEADING 'Admin|Reqests' FORMAT 9999
COLUMN STATUS HEADING 'Admin|Errors' FORMAT 9999
COLUMN TRAN HEADING 'Def|Trans|Pairs' FORMAT 9999
COLUMN ERRORS HEADING 'Def|Trans|Errors' FORMAT 9999
COLUMN COMPLETE HEADING 'Propagated|Trans' FORMAT 9999

SELECT G.GLOBAL_NAME, D.ADMIN_REQUESTS, E.STATUS, DT.TRAN, DE.ERRORS, C.COMPLETE
    FROM (SELECT GLOBAL_NAME FROM GLOBAL_NAME) G,
      (SELECT COUNT(ID) ADMIN_REQUESTS FROM DBA_REPCATLOG) D,
     (SELECT COUNT(STATUS) STATUS FROM DBA_REPCATLOG WHERE STATUS = 'ERROR') E,
      (SELECT COUNT(*) TRAN FROM DEFTRANDEST) DT,
      (SELECT COUNT(*) ERRORS FROM DEFERROR) DE,
      (SELECT COUNT(A.DEFERRED_TRAN_ID) COMPLETE FROM DEFTRAN A
    WHERE A.DEFERRED_TRAN_ID NOT IN (
      SELECT B.DEFERRED_TRAN_ID FROM DEFTRANDEST B)) C;
```

Your output looks similar to the following:

```
                                      Def    Def
                           Admin  Admin Trans  Trans Propagated
Database                 Reqests Errors Pairs Errors      Trans
------------------------ ------- ------ ----- ------ ----------
ORC1.EXAMPLE.COM               5      0    37      0         53
```

> **Note:** This query can be expensive if you have a large number of transactions in the deferred transactions queue.

## Monitoring Master Groups

This section contains queries that you can run to display information about the master groups at a replication site.

### Listing the Master Sites Participating in a Master Group

Run the following query to list the master sites for each master group at a replication site and indicate which master site is the master definition site for each master group:

```
COLUMN GNAME HEADING 'Master Group' FORMAT A20
COLUMN DBLINK HEADING 'Sites' FORMAT A25
COLUMN MASTERDEF HEADING 'Master|Definition|Site?' FORMAT A10

SELECT GNAME, DBLINK, MASTERDEF
    FROM DBA_REPSITES
    WHERE MASTER = 'Y'
    AND GNAME NOT IN (SELECT GNAME FROM DBA_REPSITES WHERE SNAPMASTER = 'Y')
    ORDER BY GNAME;
```

The subquery in the `SELECT` statement ensures that materialized view groups do not appear in the output. Your output looks similar to the following:

```
                                          Master
                                          Definition
Master Group         Sites                Site?
-------------------- -------------------- ----------
HR_REPG              ORC1.EXAMPLE.COM     Y
HR_REPG              ORC2.EXAMPLE.COM     N
HR_REPG              ORC3.EXAMPLE.COM     N
```

This list indicates that `orc1.example.com` is the master definition site for the `hr_repg` master group, which also includes the master sites `orc2.example.com` and `orc3.example.com`.

### Listing General Information About Master Groups

You can use the query in this section to list the following general information about the master groups at a master site:

- The name of each master group.
- The number of unpropagated deferred transaction-destination pairs. Each deferred transaction can have multiple destinations to which it will be propagated, and each destination is a single deferred transaction-destination pair.

For example, if there are ten deferred transactions and each one must be propagated to three sites, then there are 30 deferred transaction-pairs returned by this query. After some time, if the first deferred transaction is propagated to two of the three destination sites, then there are still ten deferred transactions, but there are two fewer deferred-transaction pairs, and this query returns 28 unpropagated deferred transaction-pairs. In this case, the first deferred transaction only has one transaction-pair remaining.

- The number of deferred transaction errors (error transactions) for each master group
- The number of administrative requests for each master group
- The number of administrative request errors for each master group

Run the following query to list this information:

```
COLUMN GNAME HEADING 'Master Group' FORMAT A15
COLUMN deftran HEADING 'Number of|Deferred|Transaction|Pairs' FORMAT 9999
COLUMN deftranerror HEADING 'Number of|Deferred|Transaction|Errors' FORMAT 9999
COLUMN adminreq HEADING 'Number of|Administrative|Requests' FORMAT 9999
COLUMN adminreqerror HEADING 'Number of|Administrative|Request|Errors'
```

```
COLUMN adminreqerror FORMAT 9999

SELECT G.GNAME,
       NVL(T.CNT1, 0) deftran,
       NVL(IE.CNT2, 0) deftranerror,
       NVL(A.CNT3, 0) adminreq,
       NVL(B.CNT4, 0) adminreqerror
    FROM
    (SELECT DISTINCT GNAME FROM DBA_REPGROUP WHERE MASTER='Y') G,
    (SELECT DISTINCT RO.GNAME, COUNT(DISTINCT D.DEFERRED_TRAN_ID) CNT1
        FROM  DBA_REPOBJECT RO, DEFCALL D, DEFTRANDEST TD
        WHERE RO.SNAME = D.SCHEMANAME
        AND RO.ONAME = D.PACKAGENAME
        AND RO.TYPE IN ('TABLE', 'PACKAGE', 'MATERIALIZED VIEW')
        AND TD.DEFERRED_TRAN_ID = D.DEFERRED_TRAN_ID
        GROUP BY RO.GNAME ) T,
    (SELECT DISTINCT RO.GNAME, COUNT(DISTINCT E.DEFERRED_TRAN_ID) CNT2
        FROM DBA_REPOBJECT RO, DEFCALL D, DEFERROR E
        WHERE RO.SNAME = D.SCHEMANAME
        AND RO.ONAME = D.PACKAGENAME
        AND RO.TYPE IN ('TABLE', 'PACKAGE', 'MATERIALIZED VIEW')
        AND E.DEFERRED_TRAN_ID = D.DEFERRED_TRAN_ID
        AND E.CALLNO = D.CALLNO
        GROUP BY RO.GNAME ) IE,
    (SELECT GNAME, COUNT(*) CNT3 FROM DBA_REPCATLOG GROUP BY GNAME) A,
    (SELECT GNAME, COUNT(*) CNT4 FROM DBA_REPCATLOG
        WHERE STATUS = 'ERROR'
        GROUP BY GNAME) B WHERE G.GNAME = IE.GNAME (+)
        AND G.GNAME = T.GNAME (+)
        AND G.GNAME = A.GNAME (+)
        AND G.GNAME = B.GNAME (+) ORDER BY G.GNAME;
```

Your output looks similar to the following:

```
                 Number of   Number of                    Number of
                  Deferred    Deferred     Number of Administrative
                Transaction Transaction Administrative        Request
Master Group          Pairs      Errors       Requests         Errors
--------------- ----------- ----------- -------------- --------------
HR_REPG                  54           0              0              0
OE_RG                    33           1              5              0
```

> **Note:** This query can be expensive if you have a large number of transactions waiting to be propagated.

## Monitoring Masters

A master can be either a master site or a master materialized view site. This section contains queries that you can run to display information about masters.

### Listing Information About Materialized Views Based on a Master

If you have materialized view sites based on a master, then you can use the query in this section to list the following information about the master:

- The number of replication groups at a master. The replication groups can be either master groups or materialized view groups.
- The number of registered materialized view groups based on the replication groups at the master.

- The number of registered materialized views based on objects at the master. The objects can be either master tables or master materialized views.
- The number of materialized view logs at the master.
- The number of deployment templates at the master.

Run the following query to list this information:

```
COLUMN repgroup HEADING 'Number of|Replication|Groups' FORMAT 9999
COLUMN mvgroup HEADING 'Number of|Registered|MV Groups' FORMAT 9999
COLUMN mv HEADING 'Number of|Registered MVs' FORMAT 9999
COLUMN mvlog HEADING 'Number of|MV Logs' FORMAT 9999
COLUMN template HEADING 'Number of|Templates' FORMAT 9999

SELECT A.REPGROUP repgroup,
       B.MVGROUP mvgroup,
       C.MV mv,
       D.MVLOG mvlog,
       E.TEMPLATE template
    FROM (SELECT COUNT(G.GNAME) REPGROUP
             FROM DBA_REPGROUP G, DBA_REPSITES S
             WHERE G.MASTER = 'Y'
             AND S.MASTER = 'Y'
             AND G.GNAME = S.GNAME
             AND S.MY_DBLINK = 'Y') A,
          (SELECT COUNT(*) MVGROUP
             FROM DBA_REGISTERED_MVIEW_GROUPS) B,
          (SELECT COUNT(*) MV
             FROM DBA_REGISTERED_MVIEWS) C,
          (SELECT COUNT(*) MVLOG
             FROM (SELECT 1 FROM DBA_MVIEW_LOGS
             GROUP BY LOG_OWNER, LOG_TABLE)) D,
          (SELECT COUNT(*) TEMPLATE FROM DBA_REPCAT_REFRESH_TEMPLATES) E;
```

Your output looks similar to the following:

```
  Number of  Number of
Replication Registered      Number of Number of Number of
     Groups  MV Groups Registered MVs   MV Logs Templates
----------- ---------- -------------- --------- ---------
          1          5             27         6         3
```

### Listing Information About the Materialized View Logs at a Master

A materialized view log enables you to perform a fast refresh on materialized views based on a master. A master can be a master table or a master materialized view. If you have materialized view logs based at a master, then you can use the query in this section to list the following information about them:

- The name of each log table that stores the materialized view log data
- The owner of each materialized view log
- The master on which each materialized view log is based
- Whether a materialized view log is a row id materialized view log
- Whether a materialized view log is a primary key materialized view log
- Whether the materialized view log is an object id materialized view log
- Whether a materialized view log has filter columns

Run the following query to list this information:

```
COLUMN LOG_TABLE HEADING 'Log Table' FORMAT A20
COLUMN LOG_OWNER HEADING 'Log|Owner' FORMAT A5
COLUMN MASTER HEADING 'Master' FORMAT A15
COLUMN ROWIDS HEADING 'Row|ID?' FORMAT A3
COLUMN PRIMARY_KEY HEADING 'Primary|Key?' FORMAT A7
COLUMN OBJECT_ID HEADING 'Object|ID?' FORMAT A6
COLUMN FILTER_COLUMNS HEADING 'Filter|Columns?' FORMAT A8

SELECT DISTINCT LOG_TABLE,
       LOG_OWNER,
       MASTER,
       ROWIDS,
       PRIMARY_KEY,
       OBJECT_ID,
       FILTER_COLUMNS
    FROM DBA_MVIEW_LOGS
    ORDER BY 1;
```

Your output looks similar to the following:

```
                    Log                   Row Primary Object Filter
Log Table           Owner Master          ID? Key?    ID?    Columns?
------------------- ----- --------------- --- ------- ------ --------
MLOG$_COUNTRIES     HR    COUNTRIES       NO  YES     NO     NO
MLOG$_DEPARTMENTS   HR    DEPARTMENTS     NO  YES     NO     NO
MLOG$_EMPLOYEES     HR    EMPLOYEES       NO  YES     NO     NO
MLOG$_JOBS          HR    JOBS            NO  YES     NO     NO
MLOG$_JOB_HISTORY   HR    JOB_HISTORY     NO  YES     NO     NO
MLOG$_LOCATIONS     HR    LOCATIONS       NO  YES     NO     NO
MLOG$_REGIONS       HR    REGIONS         NO  YES     NO     NO
```

> **See Also:** *Oracle Database Advanced Replication* for information about materialized view logs

### Listing the Materialized Views that Use a Materialized View Log

More than one materialized view can use a materialized view log. If you have materialized view logs based at a master, then you can use the query in this section to list the following the materialized views that use each log:

- The name of each log table that stores the materialized view log data
- The owner of each materialized view log
- The master on which each materialized view log is based
- The materialized view identification number of each materialized view that uses the materialized view log
- The name of each materialized view that uses the materialized view log

Run the following query to list this information:

```
COLUMN LOG_TABLE HEADING 'Mview|Log Table' FORMAT A20
COLUMN LOG_OWNER HEADING 'Mview|Log Owner' FORMAT A10
COLUMN MASTER HEADING 'Master' FORMAT A20
COLUMN MVIEW_ID HEADING 'Mview|ID' FORMAT 9999
COLUMN NAME HEADING 'Mview Name' FORMAT A20

SELECT L.LOG_TABLE, L.LOG_OWNER, B.MASTER, B.MVIEW_ID, R.NAME
FROM ALL_MVIEW_LOGS L, ALL_BASE_TABLE_MVIEWS B, ALL_REGISTERED_MVIEWS R
WHERE B.MVIEW_ID = R.MVIEW_ID
```

```
AND B.OWNER = L.LOG_OWNER
AND B.MASTER = L.MASTER;
```

Your output looks similar to the following:

```
Mview                Mview                          Mview
Log Table            Log Owner  Master                 ID Mview Name
-------------------- ---------- -------------------- ----- --------------------
MLOG$_COUNTRIES      HR         COUNTRIES              21 COUNTRIES_MV1
MLOG$_DEPARTMENTS    HR         DEPARTMENTS            22 DEPARTMENTS_MV1
MLOG$_EMPLOYEES      HR         EMPLOYEES              23 EMPLOYEES_MV1
MLOG$_JOBS           HR         JOBS                   24 JOBS_MV1
MLOG$_JOB_HISTORY    HR         JOB_HISTORY            25 JOB_HISTORY_MV1
MLOG$_LOCATIONS      HR         LOCATIONS              26 LOCATIONS_MV1
MLOG$_REGIONS        HR         REGIONS                27 REGIONS_MV1
```

### Listing Information About the Deployment Templates at a Master

Deployment templates enable you to create multiple materialized view environments quickly. They also enable you to use variables to customize each materialized view environment for its individual needs. You can use the query in this section to list the following information about the deployment templates at a master:

- The name of each deployment template
- The owner of each deployment template
- Whether a deployment template is public
- The number of instantiated materialized view sites based on each deployment template
- The comment associated with each deployment template

Run the following query to list this information:

```
COLUMN REFRESH_TEMPLATE_NAME HEADING 'Template|Name' FORMAT A10
COLUMN OWNER HEADING 'Owner' FORMAT A10
COLUMN PUBLIC_TEMPLATE HEADING 'Public?' FORMAT A7
COLUMN INSTANTIATED HEADING 'Number of|Instantiated|Sites' FORMAT 9999
COLUMN TEMPLATE_COMMENT HEADING 'Comment' FORMAT A35

SELECT DISTINCT RT.REFRESH_TEMPLATE_NAME,
       OWNER,
       PUBLIC_TEMPLATE,
       RS.INSTANTIATED,
       RT.TEMPLATE_COMMENT
    FROM DBA_REPCAT_REFRESH_TEMPLATES RT,
    (SELECT Y.REFRESH_TEMPLATE_NAME, COUNT(X.STATUS) INSTANTIATED
        FROM DBA_REPCAT_TEMPLATE_SITES X, DBA_REPCAT_REFRESH_TEMPLATES Y
        WHERE X.REFRESH_TEMPLATE_NAME(+) = Y.REFRESH_TEMPLATE_NAME
        GROUP BY Y.REFRESH_TEMPLATE_NAME) RS
        WHERE RT.REFRESH_TEMPLATE_NAME(+) = RS.REFRESH_TEMPLATE_NAME
        ORDER BY 1;
```

Your output looks similar to the following:

```
                                  Number of
Template                        Instantiated
Name       Owner      Public?        Sites Comment
---------- ---------- ------- ------------ ----------------------------------
HR_REFG_DT HR         N                  2 Human Resources Deployment Template
```

The N in the Public? column means that the deployment template is private. Therefore, it can only be instantiated by authorized users. A Y in this column means that the deployment template is public. Any user can instantiate a public deployment template.

# Monitoring Materialized View Sites

This section contains queries that you can run to display information about the materialized view sites. This section contains the following topics:

- Listing General Information About a Materialized View Site
- Listing General Information About Materialized View Groups
- Listing Information About Materialized Views
- Listing Information About the Refresh Groups at a Materialized View Site
- Determining the Job ID for Each Refresh Job at a Materialized View Site
- Determining Which Materialized Views Are Currently Refreshing

## Listing General Information About a Materialized View Site

You can use the query in this section to list the following general information about the current materialized view site:

- The number of materialized view groups at the site
- The number of materialized views at the site
- The number of refresh groups at the site

Run the following query to list this information:

```
COLUMN MVGROUP HEADING 'Number of|Materialized|View Groups' FORMAT 9999
COLUMN MV HEADING 'Number of|Materialized|Views' FORMAT 9999
COLUMN RGROUP HEADING 'Number of|Refresh Groups' FORMAT 9999

SELECT A.MVGROUP, B.MV, C.RGROUP
    FROM
    (SELECT COUNT(S.GNAME) MVGROUP
        FROM DBA_REPSITES S
        WHERE S.SNAPMASTER = 'Y') A,
     (SELECT COUNT(*) MV
        FROM DBA_MVIEWS) B,
     (SELECT COUNT(*) RGROUP
        FROM DBA_REFRESH) C;
```

Your output looks similar to the following:

```
   Number of    Number of
Materialized Materialized      Number of
 View Groups        Views Refresh Groups
------------ ------------ --------------
           5           25              5
```

## Listing General Information About Materialized View Groups

You can use the query in this section to list the following general information about the materialized view groups at the current materialized view site:

- The name of each materialized view group

- The master of each materialized view group
- The method of propagation to a materialized view group's master, either asynchronous or synchronous
- The comment associated with each materialized view group

Run the following query to list this information:

```
COLUMN GNAME HEADING 'Group Name' FORMAT A10
COLUMN DBLINK HEADING 'Master' FORMAT A25
COLUMN Propagation HEADING 'Propagation|Method' FORMAT A12
COLUMN SCHEMA_COMMENT HEADING 'Comment' FORMAT A30

SELECT S.GNAME,
       S.DBLINK,
       DECODE(S.PROP_UPDATES,
              0, 'ASYNCHRONOUS',
              1, 'SYNCHRONOUS') Propagation,
       G.SCHEMA_COMMENT
    FROM DBA_REPSITES S, DBA_REPGROUP G
    WHERE S.GNAME = G.GNAME
    AND S.SNAPMASTER = 'Y';
```

Your output looks similar to the following:

```
                                    Propagation
Group Name Master                   Method       Comment
---------- ------------------------ ------------ ------------------------------
HR_REPG    ORC1.EXAMPLE.COM         ASYNCHRONOUS
```

## Listing Information About Materialized Views

This section contains queries that you can run to display information about the materialized views at a replication site.

### Listing Master Information For Materialized Views

The following query shows the master for each materialized view at a replication site and whether the materialized view can be fast refreshed:

```
COLUMN MVIEW_NAME HEADING 'Materialized|View Name' FORMAT A15
COLUMN OWNER HEADING 'Owner' FORMAT A10
COLUMN MASTER_LINK HEADING 'Master Link' FORMAT A30
COLUMN Fast_Refresh HEADING 'Fast|Refreshable?' FORMAT A16

SELECT MVIEW_NAME,
       OWNER,
       MASTER_LINK,
       DECODE(FAST_REFRESHABLE,
              'NO', 'NO',
              'DML', 'YES',
              'DIRLOAD', 'DIRECT LOAD ONLY',
              'DIRLOAD_DML', 'YES',
              'DIRLOAD_LIMITEDDML', 'LIMITED') Fast_Refresh
    FROM DBA_MVIEWS;
```

Your output looks similar to the following:

```
Materialized                                              Fast
View Name       Owner      Master Link                    Refreshable?
--------------- ---------- ------------------------------ ----------------
```

```
COUNTRIES_MV1   HR         @ORC1.EXAMPLE.COM             YES
DEPARTMENTS_MV1 HR         @ORC1.EXAMPLE.COM             YES
EMPLOYEES_MV1   HR         @ORC1.EXAMPLE.COM             YES
JOBS_MV1        HR         @ORC1.EXAMPLE.COM             YES
JOB_HISTORY_MV1 HR         @ORC1.EXAMPLE.COM             YES
LOCATIONS_MV1   HR         @ORC1.EXAMPLE.COM             YES
REGIONS_MV1     HR         @ORC1.EXAMPLE.COM             YES
```

### Listing the Properties of Materialized Views

You can use the query in this section to list the following information about the materialized views at the current replication site:

- The name of each materialized view
- The owner of each materialized view
- The refresh method used by each materialized view: `COMPLETE`, `FORCE`, `FAST`, or `NEVER`
- Whether a materialized view is updatable
- The last date on which each materialized view was refreshed

Run the following query to list this information:

```
COLUMN MVIEW_NAME HEADING 'Materialized|View Name' FORMAT A15
COLUMN OWNER HEADING 'Owner' FORMAT A10
COLUMN REFRESH_METHOD HEADING 'Refresh|Method' FORMAT A10
COLUMN UPDATABLE HEADING 'Updatable?' FORMAT A10
COLUMN LAST_REFRESH_DATE HEADING 'Last|Refresh|Date'
COLUMN LAST_REFRESH_TYPE HEADING 'Last|Refresh|Type' FORMAT A15

SELECT MVIEW_NAME,
       OWNER,
       REFRESH_METHOD,
       UPDATABLE,
       LAST_REFRESH_DATE,
       LAST_REFRESH_TYPE
    FROM DBA_MVIEWS;
```

Your output looks similar to the following:

```
                                                 Last      Last
Materialized                Refresh              Refresh   Refresh
View Name       Owner       Method     Updatable? Date      Type
--------------- ----------- ---------- ---------- --------- ---------------
COUNTRIES_MV1   HR          FAST       Y          21-OCT-03 FAST
DEPARTMENTS_MV1 HR          FAST       Y          21-OCT-03 FAST
EMPLOYEES_MV1   HR          FAST       Y          21-OCT-03 FAST
JOBS_MV1        HR          FAST       Y          21-OCT-03 FAST
JOB_HISTORY_MV1 HR          FAST       Y          21-OCT-03 FAST
LOCATIONS_MV1   HR          FAST       Y          21-OCT-03 FAST
REGIONS_MV1     HR          FAST       Y          21-OCT-03 FAST
```

## Listing Information About the Refresh Groups at a Materialized View Site

Each refresh group at a materialized view site is associated with a refresh job that refreshes the materialized views in the refresh group at a set interval. You can query the `DBA_REFRESH` data dictionary view to list the following information about the refresh jobs at a materialized view site:

- The name of the refresh group.

- The owner of the refresh group.
- Whether the refresh job is broken.
- The next date and time when the refresh job will run.
- The current interval setting for the refresh job. The interval setting specifies the amount of time between the start of a job and the next start of the same job.

The following query displays this information:

```
COLUMN RNAME HEADING 'Refresh|Group|Name' FORMAT A10
COLUMN ROWNER HEADING 'Refresh|Group|Owner' FORMAT A10
COLUMN BROKEN HEADING 'Broken?' FORMAT A7
COLUMN next_refresh HEADING 'Next Refresh'
COLUMN INTERVAL HEADING 'Interval' FORMAT A20

SELECT RNAME,
       ROWNER,
       BROKEN,
       TO_CHAR(NEXT_DATE, 'DD-MON-YYYY HH:MI:SS AM') next_refresh,
       INTERVAL
    FROM DBA_REFRESH
    ORDER BY 1;
```

Your output looks similar to the following:

```
Refresh    Refresh
Group      Group
Name       Owner      Broken? Next Refresh            Interval
---------- ---------- ------- ----------------------- --------------------
HR_REFG    MVIEWADMIN N       24-OCT-2003 07:18:44 AM SYSDATE + 1/24
```

The `N` in the `Broken?` column means that the job is not broken. Therefore, the refresh job will run at the next start time. A `Y` in this column means that the job is broken.

## Determining the Job ID for Each Refresh Job at a Materialized View Site

You can use the query in this section to list the following information about the refresh jobs at a materialized view site:

- The job identification number of each refresh job. Each job created by the `DBMS_JOB` package is assigned a unique identification number.
- The privilege schema, which is the schema whose default privileges apply to the job.
- The schema that owns each refresh job. Typically, the materialized view administrator owns a refresh job. A common user name for the materialized view administrator is `mviewadmin`.
- The name of the refresh group that the job refreshes.
- The status of the refresh job, either normal or broken.

The following query displays this information:

```
COLUMN JOB HEADING 'Job ID' FORMAT 999999
COLUMN PRIV_USER HEADING 'Privilege|Schema' FORMAT A10
COLUMN RNAME HEADING 'Refresh|Group|Name' FORMAT A10
COLUMN ROWNER HEADING 'Refresh|Group|Owner' FORMAT A10
COLUMN BROKEN HEADING 'Broken?' FORMAT A7

SELECT J.JOB,
```

```
    J.PRIV_USER,
    R.ROWNER,
    R.RNAME,
    J.BROKEN
FROM DBA_REFRESH R, DBA_JOBS J
WHERE R.JOB = J.JOB
ORDER BY 1;
```

Your output looks similar to the following:

```
                  Refresh    Refresh
       Privilege  Group      Group
 Job ID Schema     Owner      Name       Broken?
------- ---------- ---------- ---------- -------
     21 MVIEWADMIN MVIEWADMIN HR_REFG    N
```

The `N` in the `Broken?` column means that the job is not broken. Therefore, the job will run at the next start time. A `Y` in this column means that the job is broken.

### Determining Which Materialized Views Are Currently Refreshing

The following query shows the materialized views that are currently refreshing:

```
COLUMN SID HEADING 'Session|Identifier' FORMAT 9999
COLUMN SERIAL# HEADING 'Serial|Number' FORMAT 999999
COLUMN CURRMVOWNER HEADING 'Owner' FORMAT A15
COLUMN CURRMVNAME HEADING 'Materialized|View' FORMAT A25

SELECT * FROM V$MVREFRESH;
```

Your output looks similar to the following:

```
   Session  Serial                 Materialized
Identifier  Number Owner           View
---------- ------- --------------- -------------------------
        19     233 HR              COUNTRIES_MV
         5     647 HR              EMPLOYEES_MV
```

> **Note:** The `V$MVREFRESH` dynamic performance view does not contain information about updatable materialized views when the materialized views' deferred transactions are being pushed to its master.

## Monitoring Administrative Requests

This section contains queries that you can run to display information about the administrative requests at a master site. This section contains the following topics:

- Listing General Information About Administrative Requests
- Listing General Information About the Job that Executes Administrative Requests
- Listing General Information About the Job that Executes Administrative Requests

### Listing General Information About Administrative Requests

You can use the query in this section to list the following general information about the administrative requests at a master site:

- The identification number of each administrative request
- The action requested by each administrative request
- The status of each request
- The master site where the request is being executed

The following query displays this information:

```
COLUMN ID HEADING 'Admin|Request|ID' FORMAT 999999
COLUMN REQUEST HEADING 'Request' FORMAT A25
COLUMN STATUS HEADING 'Status' FORMAT A15
COLUMN MASTER HEADING 'Master|Site' FORMAT A25

SELECT ID, REQUEST, STATUS, MASTER FROM DBA_REPCATLOG;
```

Your output looks similar to the following:

```
  Admin
Request                                          Master
     ID Request                  Status          Site
------- ------------------------ --------------- -------------------------
     44 RESUME_MASTER_ACTIVITY    AWAIT_CALLBACK  ORC3.EXAMPLE.COM
```

You can use the `DO_DEFERRED_REPCAT_ADMIN` procedure in the `DBMS_REPCAT` package to execute administrative requests.

## Determining the Cause of Administrative Request Errors

You can determine the cause of an administrative request error by displaying its error message. The following query displays the error message for each administrative request that resulted in an error:

```
COLUMN ID HEADING 'Admin|Request|ID' FORMAT 999999
COLUMN REQUEST HEADING 'Request' FORMAT A30
COLUMN ERRNUM HEADING 'Error|Number' FORMAT 999999
COLUMN MESSAGE HEADING 'Error|Message' FORMAT A32

SELECT ID, REQUEST, ERRNUM, MESSAGE
    FROM DBA_REPCATLOG WHERE STATUS = 'ERROR';
```

Your output looks similar to the following:

```
  Admin
Request                               Error Error
     ID Request                      Number Message
------- ------------------------------ ------- --------------------------------
     70 CREATE_MASTER_REPOBJECT          -2292 ORA-02292: integrity constrain
                                               t (HR.DEPT_LOC_FK) violated -
                                               child record found
                                               ORA-02266: unique/primary keys
                                                in table referenced by enable
                                               d foreign keys

     71 GENERATE_INTERNAL_PKG_SUPPORT   -23308 ORA-23308: object HR.LOCATIONS
                                                does not exist or is invalid
```

## Listing General Information About the Job that Executes Administrative Requests

Each master group is associated with a do_deferred_repcat_admin job that executes administrative requests. You can query the DBA_JOBS data dictionary view to list the following information about this job at a replication site:

- The job identification number of each do_deferred_repcat_admin job. Each job created by the DBMS_JOB package is assigned a unique identification number.
- The privilege schema, which is the schema whose default privileges apply to the job.
- The status of each do_deferred_repcat_admin job, either normal or broken.
- The next date and time when each do_deferred_repcat_admin job will run.
- The current interval setting for each do_deferred_repcat_admin job. The interval setting specifies the amount of time between the start of a job and the next start of the same job.

The following query displays this information:

```
COLUMN JOB HEADING 'Job ID' FORMAT 999999
COLUMN PRIV_USER HEADING 'Privilege|Schema' FORMAT A10
COLUMN BROKEN HEADING 'Broken?' FORMAT A7
COLUMN next_start HEADING 'Next Start'
COLUMN INTERVAL HEADING 'Interval' FORMAT A20

SELECT JOB,
       PRIV_USER,
       BROKEN,
       TO_CHAR(NEXT_DATE,'DD-MON-YYYY HH:MI:SS AM') next_start,
       INTERVAL
    FROM DBA_JOBS
    WHERE WHAT LIKE '%dbms_repcat.do_deferred_repcat_admin%'
    ORDER BY 1;
```

Your output looks similar to the following:

```
        Privilege
 Job ID Schema     Broken? Next Start              Interval
------- ---------- ------- ----------------------- --------------------
     24 REPADMIN   N       24-OCT-2003 07:23:48 AM SYSDATE + (1/144)
```

The N in the Broken? column means that the job is not broken. Therefore, the job will run at the next start time. A Y in this column means that the job is broken.

### Checking the Definition of Each do_deferred_repcat_admin Job

You can query the DBA_JOBS data dictionary view to show the definition of each do_deferred_repcat_admin job at a replication site. The following query shows the definitions:

```
COLUMN JOB HEADING 'Job ID' FORMAT 999999
COLUMN WHAT HEADING 'Definitions of Admin Req Jobs' FORMAT A70

SELECT JOB, WHAT
  FROM DBA_JOBS
  WHERE WHAT LIKE '%dbms_repcat.do_deferred_repcat_admin%'
  ORDER BY 1;
```

Your output looks similar to the following:

```
 Job ID Definitions of Admin Req Jobs
```

```
------- -------------------------------------------------------------------
    321 dbms_repcat.do_deferred_repcat_admin('"HR_REPG"', FALSE);
    342 dbms_repcat.do_deferred_repcat_admin('"OE_RG"', FALSE);
```

# Monitoring the Deferred Transactions Queue

This section contains queries that you can run to display information about the deferred transactions queue at a replication site. This section contains the following topics:

- Monitoring Transaction Propagation
- Monitoring Purges of Successfully Propagated Transactions

## Monitoring Transaction Propagation

This section contains queries that you can run to display information about propagation of transactions in the deferred transactions queue.

### Listing the Number of Deferred Transactions for Each Destination Master Site

You can find the number of unpropagated deferred transactions for each destination master site by running the query in this section. This query shows each master site to which the current master site is propagating deferred transactions and the number of deferred transactions to be propagated to each destination site.

Run the following query to see the number of deferred and error transactions:

```
COLUMN DEST HEADING 'Destination' FORMAT A45
COLUMN TRANS HEADING 'Def Trans' FORMAT 9999

SELECT DBLINK DEST, COUNT(*) TRANS
    FROM DEFTRANDEST D
    GROUP BY DBLINK;
```

Your output looks similar to the following:

```
Destination                                   Def Trans
--------------------------------------------- ---------
ORC2.EXAMPLE.COM                                      1
ORC3.EXAMPLE.COM                                      1
```

> **Note:** This query can be expensive if you have a large number of transactions waiting to be propagated.

### Listing General Information About the Push Jobs at a Replication Site

Each scheduled link at a replication site is associated with a push job that propagates deferred transactions in the deferred transaction queue to a destination site. You can use the query in this section to list the following information about the push jobs at a replication site:

- The job identification number of each push job. Each job created by the `DBMS_JOB` package is assigned a unique identification number.
- The privilege schema, which is the schema whose default privileges apply to the job.
- The destination site where the deferred transactions are pushed.

- The status of the push job, either normal or broken.

The following query displays this information:

```
COLUMN JOB HEADING 'Job ID' FORMAT 999999
COLUMN PRIV_USER HEADING 'Privilege|Schema' FORMAT A10
COLUMN DBLINK HEADING 'Destination' FORMAT A40
COLUMN BROKEN HEADING 'Broken?' FORMAT A7

SELECT J.JOB,
       J.PRIV_USER,
       S.DBLINK,
       J.BROKEN
    FROM DEFSCHEDULE S, DBA_JOBS J
    WHERE S.DBLINK != (SELECT GLOBAL_NAME FROM GLOBAL_NAME)
    AND S.JOB = J.JOB
    ORDER BY 1;
```

Your output looks similar to the following:

```
        Privilege
 Job ID Schema     Destination                              Broken?
------- ---------- ---------------------------------------- -------
     22 REPADMIN   ORC2.EXAMPLE.COM                         N
     23 REPADMIN   ORC3.EXAMPLE.COM                         N
```

The `N` in the `Broken?` column means that the job is not broken. Therefore, the job will run at the next start time. A `Y` in this column means that the job is broken.

### Determining the Next Start Time and Interval for the Push Jobs

Each scheduled link at a replication site is associated with a push job that propagates deferred transactions in the deferred transaction queue to a destination site. You can query the `DEFSCHEDULE` and `DBA_JOBS` data dictionary views to list the following information about the push jobs at a replication site:

- The job identification number of each push job. Each job created by the `DBMS_JOB` package is assigned a unique identification number.
- The destination site where the deferred transactions are pushed.
- The next date and time when the push job will run.
- The current interval setting for the push job. The interval setting specifies the amount of time between the start of a job and the next start of the same job.

The following query displays this information:

```
COLUMN JOB HEADING 'Job ID' FORMAT 999999
COLUMN DBLINK HEADING 'Destination' FORMAT A22
COLUMN next_start HEADING 'Next Start'
COLUMN INTERVAL HEADING 'Interval' FORMAT A25

SELECT JOB,
       DBLINK,
       TO_CHAR(NEXT_DATE, 'DD-MON-YYYY HH:MI:SS AM') next_start,
       INTERVAL
    FROM DEFSCHEDULE
    WHERE DBLINK != (SELECT GLOBAL_NAME FROM GLOBAL_NAME)
    AND JOB IS NOT NULL
    ORDER BY 1;
```

Your output looks similar to the following:

```
 Job ID Destination            Next Start              Interval
------- ---------------------- ---------------------- -------------------------
     22 ORC2.EXAMPLE.COM       24-OCT-2003 07:23:48 AM SYSDATE + (1/144)
     23 ORC3.EXAMPLE.COM       24-OCT-2003 07:23:48 AM SYSDATE + (1/144)
```

### Determining the Total Number of Transactions Queued for Propagation

Run the following query to display the total number of transactions in the deferred transaction queue that are waiting to be propagated:

```
SELECT COUNT(DISTINCT DEFERRED_TRAN_ID) "Transactions Queued"
    FROM DEFTRANDEST;
```

Your output looks similar to the following:

```
Transactions Queued
-------------------
                 37
```

> **Note:** This query can be expensive if you have a large number of transactions waiting to be propagated.

## Monitoring Purges of Successfully Propagated Transactions

This section contains queries that you can run to display information about purges of successfully propagated transactions from the deferred transactions queue.

### Listing General Information About the Purge Job

During standard setup of a replication site, you configure a purge job to remove successfully propagated transactions from the deferred transactions queue. You can query the `DBA_JOBS` data dictionary view to list the following information about the purge job at a replication site:

- The job identification number of the purge job. Each job created by the `DBMS_JOB` package is assigned a unique identification number.
- The privilege schema, which is the schema whose default privileges apply to the job.
- The status of the job, either normal or broken.
- The next date and time when the purge job will run.
- The current interval setting for the purge job. The interval setting specifies the amount of time between the start of a job and the next start of the same job.

The following query displays this information:

```
COLUMN JOB HEADING 'Job ID' FORMAT 999999
COLUMN PRIV_USER HEADING 'Privilege|Schema' FORMAT A10
COLUMN BROKEN HEADING 'Broken?' FORMAT A7
COLUMN next_start HEADING 'Next Start'
COLUMN INTERVAL HEADING 'Interval' FORMAT A25

SELECT JOB,
       PRIV_USER,
       BROKEN,
       TO_CHAR(NEXT_DATE, 'DD-MON-YYYY HH:MI:SS AM') next_start,
       INTERVAL
    FROM DBA_JOBS
```

```
    WHERE WHAT LIKE '%dbms_defer_sys.purge%'
    ORDER BY 1;
```

Your output looks similar to the following:

```
        Privilege
 Job ID Schema     Broken? Next Start              Interval
------- ---------- ------- ---------------------- ------------------------
     21 REPADMIN   N       24-OCT-2003 07:42:18 AM SYSDATE + 1/24
```

The `N` in the `Broken?` column means that the job is not broken. Therefore, the job will run at the next start time. A `Y` in this column means that the job is broken.

### Checking the Definition of the Purge Job

You can query the `DBA_JOBS` data dictionary view to show the definition of the purge job at a replication site. The following query shows the definition:

```
SELECT WHAT "Definition of the Purge Job"
    FROM DBA_JOBS
    WHERE WHAT LIKE '%dbms_defer_sys.purge%' ORDER BY 1;
```

Your output looks similar to the following:

```
Definition of the Purge Job
--------------------------------------------------------------------------------
declare rc binary_integer; begin rc := sys.dbms_defer_sys.purge( delay_seconds=>
0); end;
```

### Determining the Amount of Time Since the Last Purge

The following query shows the total amount of time, in minutes, since the successfully propagated transactions were purged from the deferred transactions queue:

```
SELECT ((SYSDATE - LAST_PURGE_TIME) / 60) "Minutes Since Last Purge"
    FROM V$REPLQUEUE;
```

Your output looks similar to the following:

```
Minutes Since Last Purge
------------------------
              13.43333
```

### Determining the Total Number of Purged Transactions

The following query shows the total number of successfully propagated transactions that have been purged from the deferred transaction queue since the instance was last started:

```
SELECT TXNS_PURGED "Transactions Purged"
    FROM V$REPLQUEUE;
```

Your output looks similar to the following:

```
Transactions Purged
-------------------
               6541
```

## Monitoring the Error Queue

This section contains queries that you can run to display information about the error queue at a replication site. The error queue contains deferred transactions that resulted

in an error at the destination site. These error transactions are placed in the error queue at the destination site.

This section contains the following topics:


- Listing General Information About the Error Transactions at a Replication Site
- Determining the Percentage of Error Transactions
- Listing the Number of Error Transactions from Each Origin Master Site
- Listing the Error Messages for the Error Transactions at a Replication Site
- Determining the Error Operations at a Replication Site


## Listing General Information About the Error Transactions at a Replication Site

The following query lists the general information about the error transactions at a replication site:

```
COLUMN DEFERRED_TRAN_ID HEADING 'Deferred|Transaction|ID' FORMAT A11
COLUMN ORIGIN_TRAN_DB HEADING 'Origin|Database' FORMAT A15
COLUMN DESTINATION HEADING 'Destination|Database' FORMAT A15
COLUMN TIME_OF_ERROR HEADING 'Time of|Error' FORMAT A22
COLUMN ERROR_NUMBER HEADING 'Oracle|Error|Number' FORMAT 999999

SELECT DEFERRED_TRAN_ID,
       ORIGIN_TRAN_DB,
       DESTINATION,
       TO_CHAR(START_TIME, 'DD-Mon-YYYY hh24:mi:ss') TIME_OF_ERROR,
       ERROR_NUMBER
    FROM DEFERROR ORDER BY START_TIME;
```

Your output looks similar to the following:

```
Deferred                                                        Oracle
Transaction Origin          Destination     Time of               Error
ID          Database        Database        Error                Number
----------- --------------- --------------- -------------------- -------
1.8.2470    ORC2.EXAMPLE.COM ORC1.EXAMPLE.COM 22-Oct-2003 07:19:14      1403
```

You can use the deferred transaction ID and the destination database to either attempt to rerun the transaction that caused the error or to delete the error.

For example, to attempt to rerun the transaction in the previous example, enter the following:

```
EXECUTE DBMS_DEFER_SYS.EXECUTE_ERROR('1.8.2470', 'ORC1.EXAMPLE.COM');
```

To delete the error in the previous example, enter the following:

```
EXECUTE DBMS_DEFER_SYS.DELETE_ERROR('1.8.2470', 'ORC1.EXAMPLE.COM');
```

Typically, you should delete an error only if you have resolved it manually.

## Determining the Percentage of Error Transactions

When propagating transactions to a remote master site, some transactions are propagated and applied successfully while other transactions can result in errors at the remote master site. Transactions that result in errors are called error transactions.

Run the following query to display the percentage of error transactions that resulted from propagation to the remote master site `orc2.example.com`:

```
SELECT DECODE(TOTAL_TXN_COUNT, 0, 'No Transactions',
       (TOTAL_ERROR_COUNT/TOTAL_TXN_COUNT)*100) "ERROR PERCENTAGE"
  FROM DEFSCHEDULE
  WHERE DBLINK = 'ORC2.EXAMPLE.COM';
```

Your output looks similar to the following:

```
Error Percentage
----------------
           3.265
```

> **Note:** If this query returns `'No transactions'`, then no transactions have been propagated to the specified remote site since the statistics were last cleared.

## Listing the Number of Error Transactions from Each Origin Master Site

You can find the number of transaction errors resulting from pushes by each origin master site by running the query in this section.

Run the following query to see the number of deferred and error transactions:

```
COLUMN SOURCE HEADING 'Origin' FORMAT A45
COLUMN ERRORS HEADING 'Def Trans Errors' FORMAT 9999

SELECT E.ORIGIN_TRAN_DB SOURCE, COUNT(*) ERRORS
    FROM DEFERROR E
    GROUP BY E.ORIGIN_TRAN_DB;
```

Your output looks similar to the following:

```
Origin                                      Def Trans Errors
------------------------------------------- ----------------
ORC2.EXAMPLE.COM                                           1
ORC3.EXAMPLE.COM                                           3
```

## Listing the Error Messages for the Error Transactions at a Replication Site

The following query lists the error messages for the error transactions at a replication site:

```
COLUMN DEFERRED_TRAN_ID HEADING 'Deferred|Transaction|ID' FORMAT A11
COLUMN ERROR_MSG HEADING 'Error Messages' FORMAT A68

SELECT DEFERRED_TRAN_ID, ERROR_MSG
    FROM DEFERROR;
```

Your output looks similar to the following:

```
Deferred
Transaction
ID          Error Messages
----------- --------------------------------------------------------------------
1.8.2470    ORA-01403: no data found
```

## Determining the Error Operations at a Replication Site

The following query lists the type of operation that was attempted for each call that caused an error at a replication site:

```
COLUMN CALLNO HEADING 'Call|Number' FORMAT 9999
COLUMN DEFERRED_TRAN_ID HEADING 'Deferred|Transaction|ID' FORMAT A11
COLUMN PACKAGENAME HEADING 'Package|Name' FORMAT A20
COLUMN PROCNAME HEADING 'Operation' FORMAT A15
COLUMN ORIGIN_TRAN_DB HEADING 'Origin|Database' FORMAT A16

SELECT /*+ ORDERED */
       C.CALLNO,
       C.DEFERRED_TRAN_ID,
       C.PACKAGENAME,
       C.PROCNAME, E.ORIGIN_TRAN_DB
    FROM DEFERROR E, DEFCALL C
    WHERE C.DEFERRED_TRAN_ID = E.DEFERRED_TRAN_ID
    AND C.CALLNO = E.CALLNO
    ORDER BY E.START_TIME;
```

Your output looks similar to the following:

```
       Deferred
  Call Transaction Package                          Origin
Number ID          Name                 Operation    Database
------ ----------- -------------------- ------------ ----------------
     0 1.8.2470    EMPLOYEES$RP         REP_UPDATE   ORC2.EXAMPLE.COM
```

# Monitoring Performance in a Replication Environment

This section contains queries that you can run to monitor the performance of your replication environment. This section contains the following topics:

- Tracking the Average Number of Row Changes in a Replication Transaction
- Determining the Average Amount of Time to Apply Transactions at Remote Sites
- Determining the Percentage of Time the Parallel Propagation Job Spends Sleeping
- Clearing the Statistics for a Remote Master Site in the DEFSCHEDULE View
- Monitoring Parallel Propagation of Deferred Transactions Using V$REPLPROP

## Tracking the Average Number of Row Changes in a Replication Transaction

The following query shows the average number of row changes in a replication transaction since instance startup:

```
SELECT DECODE(TXNS_ENQUEUED, 0, 'No Transactions Enqueued',
       (CALLS_ENQUEUED / TXNS_ENQUEUED)) "Average Number of Row Changes"
  FROM V$REPLQUEUE;
```

Your output looks similar to the following:

```
Average Number of Row Changes
-----------------------------
                        56.16
```

**Note:** If this query returns `'No Transactions Enqueued'`, then no transactions have been enqueued since the start of the instance.

## Tracking the Rate of Transactions Entering the Deferred Transactions Queue

The following query shows the average number of transactions for each second entering at the deferred transactions queue at the current site since instance startup:

```
SELECT (R.TXNS_ENQUEUED / ((SYSDATE - I.STARTUP_TIME)*24*60*60)) "Average TPS"
    FROM V$REPLQUEUE R, V$INSTANCE I;
```

Your output looks similar to the following:

```
Average TPS
-----------
        150
```

## Determining the Average Network Traffic Created to Propagate a Transaction

Propagation of deferred transactions creates a certain amount of traffic on your network. Here, the network traffic created by a transaction is the number of bytes being sent and received and the number of network round trips needed to propagate the transaction.

A round trip is one or more consecutively sent messages followed by one or more consecutively received messages. For example, both of the following scenarios constitute only one round trip:

- Site A sends one message to site B and then site B sends one message to site A.
- Site A sends 20 messages to site B and then site B sends one message to site A.

These scenarios illustrate that the number of messages is irrelevant when evaluating the number of round trips, because the number of round trips is the number of back and forth communications between sites.

The following query shows the average network traffic created when propagating a transaction to the `orc2.example.com` remote master site:

```
COLUMN AV_BYTES HEADING 'Average Bytes' FORMAT 999999999
COLUMN AV_TRIPS HEADING 'Average Round Trips' FORMAT 9999999

SELECT
 DECODE(TOTAL_TXN_COUNT, 0, 'No Transactions',
 ((TOTAL_BYTES_SENT + TOTAL_BYTES_RECEIVED) / TOTAL_TXN_COUNT)) AV_BYTES,
 DECODE(TOTAL_TXN_COUNT, 0, 'No Transactions',
 (TOTAL_ROUND_TRIPS / TOTAL_TXN_COUNT)) AV_TRIPS
 FROM DEFSCHEDULE WHERE DBLINK = 'ORC2.EXAMPLE.COM';
```

Your output looks similar to the following:

```
Average Bytes          Average Round Trips
---------------------- -------------------
69621.5                5
```

**Note:**

- If this query returns `'No transactions'` in both columns, then no transactions have been propagated to the specified remote site since the statistics were last cleared.
- This query returns results only if parallel propagation is used with the specified database link. To use parallel propagation, set the `parallelism` parameter to `1` or greater when you run the `SCHEDULE_PUSH` procedure in the `DBMS_DEFER_SYS` package.

**See Also:**

- *Oracle Database Advanced Replication* for information about parallel propagation
- "SCHEDULE_PUSH Procedure" on page 14-22

## Determining the Average Amount of Time to Apply Transactions at Remote Sites

Average latency is the average number of seconds between the first call of a transaction on the current site and the confirmation that the transaction was applied at the remote site. The first call begins when the user makes the first data manipulation language (DML) change, not when the transaction is committed.

The following query shows the average latency for applying transactions at the remote master site `orc2.example.com`:

```
SELECT AVG_LATENCY "Average Latency"
    FROM DEFSCHEDULE
    WHERE DBLINK='ORC2.EXAMPLE.COM';
```

Your output looks similar to the following:

```
Average Latency
---------------
           25.5
```

## Determining the Percentage of Time the Parallel Propagation Job Spends Sleeping

When the parallel propagation coordinator is inactive, it is sleeping. You control the amount of time that the propagation coordinator sleeps using the `delay_seconds` parameter in the `DBMS_DEFER_SYS.PUSH` procedure.

The following query shows the percentage of time that the parallel propagation coordinator spends sleeping when propagating transactions to the `orc2.example.com` remote master site:

```
SELECT DECODE(AVG_THROUGHPUT, 0, NULL,
       ((TOTAL_SLEEP_TIME / (TOTAL_TXN_COUNT / AVG_THROUGHPUT)) * 100))
       "Percent Sleep Time"
  FROM DEFSCHEDULE WHERE DBLINK = 'ORC2.EXAMPLE.COM';
```

Your output looks similar to the following:

```
Percent Sleep Time
------------------
                 2
```

In this case, the parallel propagation coordinator is active 98% of the time.

**Note:** If this query returns a NULL, then no transactions have been propagated to the specified remote site since the statistics were last cleared or since the last database startup.

## Clearing the Statistics for a Remote Master Site in the DEFSCHEDULE View

To clear the propagation statistics in the DEFSCHEDULE view for a particular remote master site, use the CLEAR_PROP_STATISTICS procedure in the DBMS_DEFER_SYS package. For example, to clear the propagation statistics for the orc2.example.com remote master site, run the following procedure:

```
BEGIN
  DBMS_DEFER_SYS.CLEAR_PROP_STATISTICS (
   dblink => 'ORC2.EXAMPLE.COM');
END;
/
```

## Monitoring Parallel Propagation of Deferred Transactions Using V$REPLPROP

The V$REPLPROP dynamic performance view provides information about current parallel propagation sessions.

**Note:** The V$REPLPROP dynamic performance view is only relevant if you are using parallel propagation of deferred transactions. If you are using serial propagation, then this view is empty.

### Determining the Databases to Which You Are Propagating Deferred Transactions

Run the following query to list the database link of each database to which you are currently propagating deferred transactions using parallel propagation:

```
SELECT DBLINK "Database Link"
    FROM V$REPLPROP
    WHERE NAME LIKE '%Coordinator%';
```

Your output looks similar to the following:

```
Database Link
-----------------
ORC2.EXAMPLE.COM
ORC3.EXAMPLE.COM
```

### Determining the Transactions Currently Being Propagated to a Remote Master

You can list the following information about the transactions that are currently being propagated to a specified remote master site using parallel propagation:

- The transaction identification number of each transaction.
- The number of calls in each transaction.
- The percentage of processed calls in each transaction. The number in this column becomes larger as the calls in the transaction are processed. When the number reaches 100, all of the calls are processed.

The following query displays this information:

```
SELECT /*+ ORDERED */ P.XID "Tran Being Propagated",
       (MAX(C.CALLNO) + 1) "Number of Calls in Tran",
```

```
    (P.SEQUENCE/MAX(C.CALLNO) + 1) * 100 "% Processed Calls"
FROM V$REPLPROP P, DEFCALL C
WHERE P.NAME LIKE '%SLAVE%'
AND P.DBLINK = 'mv4.example.com'
AND C.DEFERRED_TRAN_ID = P.XID
GROUP BY P.XID, P.SEQUENCE;
```

Your output looks similar to the following:

```
Tran Being Propagated  Number of Calls in Tran % Processed Calls
---------------------- ----------------------- -----------------
1.11.4264                                43357                78
1.15.4256                                23554                49
```

The transaction identification numbers should change as existing transactions are pushed and new transactions are processed. This query can be particularly useful if the any of the following conditions apply to your replication environment:

- You push a large number of transactions on a regular basis.
- You have some transactions that are very large.
- You are simulating continuous push using asynchronous propagation.

If the first two bullets apply to your replication environment, then you can run this query to check if the processes are pushing the transactions. In this type of environment, the processes do not exist when they are not pushing transactions.

In replication environments that are simulating continuous push, the processes exist whenever there are transactions to push in the deferred transactions queue. When there are no transactions to push, the processes might not exist. So, when there are transactions to push, you can use this query to ensure that the processes exist and are processing the transactions.

> **See Also:** *Oracle Database Advanced Replication* for more information about scheduling continuous push in your replication environment

# Part III

# Replication Management API Packages Reference

Part III includes reference information about the replication management API, including:

- The procedures and functions in each package
- The parameters for each packaged procedure or function
- Exceptions that each procedure or function can raise

> **Note:** Some PL/SQL procedures and functions described in the chapters in this part are overloaded. That is, two or more procedures or functions have the same name in a single package, but their formal parameters differ in number, order, or data type family. When a procedure or function is overloaded, it is noted in the description. See the *Oracle Database PL/SQL Language Reference* for more information about overloading and for more information about PL/SQL in general.

This part contains the following chapters:


- Chapter 11, "Introduction to the Replication Management API Reference"
- Chapter 12, "DBMS_DEFER"
- Chapter 13, "DBMS_DEFER_QUERY"
- Chapter 14, "DBMS_DEFER_SYS"
- Chapter 15, "DBMS_OFFLINE_OG"
- Chapter 16, "DBMS_RECTIFIER_DIFF"
- Chapter 17, "DBMS_REFRESH"
- Chapter 18, "DBMS_REPCAT"
- Chapter 19, "DBMS_REPCAT_INSTANTIATE"
- Chapter 20, "DBMS_REPCAT_ADMIN"
- Chapter 21, "DBMS_REPCAT_RGT"
- Chapter 22, "DBMS_REPUTIL"

# 11

# Introduction to the Replication Management API Reference

All installations of Advanced Replication include the replication management application programming interface (API). This **replication management API** is a collection of PL/SQL packages that administrators use to configure and manage replication features at each site. The Advanced Replication interface in Oracle Enterprise Manager also uses the procedures and functions of each site's replication management API to perform work.

This chapter contains these topics:

- Examples of Using Oracle's Replication Management API
- Issues to Consider When Using the Replication Management API
- The Advanced Replication Interface and the Replication Management API
- Abbreviations for Datetime and Interval Data Types

> **Note:** Some PL/SQL procedures and functions described in the chapters in this part are overloaded. That is, two or more procedures or functions have the same name in a single package, but their formal parameters differ in number, order, or data type family. When a procedure or function is overloaded, it is noted in the description. See the *Oracle Database PL/SQL Language Reference* for more information about overloading and for more information about PL/SQL in general.

## Examples of Using Oracle's Replication Management API

To use Oracle's replication management API, you issue procedure or function calls using a query tool such as SQL*Plus. For example, the following call to the `DBMS_REPCAT.CREATE_MASTER_REPOBJECT` procedure creates a new replicated table `hr.employees` in the `hr_repg` replication group:

```
BEGIN
   DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
      gname => 'hr_repg',
      type => 'TABLE',
      oname => 'employees',
      sname => 'hr',
      use_existing_object => TRUE,
      copy_rows => FALSE);
END;
/
```

To call a replication management API function, you must provide an environment to receive the return value of the function. For example, the following anonymous PL/SQL block calls the DBMS_DEFER_SYS.DISABLED function in an IF statement.

```
BEGIN
 IF DBMS_DEFER_SYS.DISABLED ('inst2') THEN
   DBMS_OUTPUT.PUT_LINE('Propagation to INST2 is disabled.');
 ELSE
   DBMS_OUTPUT.PUT_LINE('Propagation to INST2 is enabled.');
 END IF;
END;
/
```

## Issues to Consider When Using the Replication Management API

For many procedures and functions in the replication management API, there are important issues to consider. For example:

- Some procedures or functions are appropriate to call only from the master definition site in a multimaster configuration.
- To perform some administrative operations for master groups, you must first suspend replication activity for the group before calling replication management API procedures and functions.
- The order in which you call different procedures and functions in Oracle's replication management API is extremely important. See the next section for more information about learning how to correctly issue replication management calls.

## The Advanced Replication Interface and the Replication Management API

The Advanced Replication interface in Oracle Enterprise Manager uses the replication management API to perform most of its functions. Using the Advanced Replication interface is much more convenient than issuing replication management API calls individually because the utility:

- Provides a GUI interface to type in and adjust API call parameters
- Automatically orders numerous, related API calls in the proper sequence
- Displays output returned from API calls in message boxes and error files

## Abbreviations for Datetime and Interval Data Types

Many of the datetime and interval data types have names that are too long to be used with the procedures and functions in the replication management API. Therefore, you must use abbreviations for these data types instead of the full names. The following table lists each data type and its abbreviation. No abbreviation is necessary for the DATE and TIMESTAMP data types.

| Data Type | Abbreviation |
|---|---|
| TIMESTAMP WITH TIME ZONE | TSTZ |
| TIMESTAMP LOCAL TIME ZONE | TSLTZ |
| INTERVAL YEAR TO MONTH | IYM |
| INTERVAL DAY TO SECOND | IDS |

For example, to use the DBMS_DEFER_QUERY.GET_*datatype*_ARG function to determine the value of a TIMESTAMP LOCAL TIME ZONE argument in a deferred call, substitute TSLTZ for *datatype*. Therefore, you run the DBMS_DEFER_QUERY.GET_TSLTZ_ARG function.

# 12

# DBMS_DEFER

DBMS_DEFER is the user interface to a replicated transactional deferred remote procedure call facility. Replicated applications use the calls in this interface to queue procedure calls for later transactional execution at remote nodes.

These procedures are typically called from either after row triggers or application specified update procedures.

This chapter contains this topic:


- Summary of DBMS_DEFER Subprograms

# Summary of DBMS_DEFER Subprograms

*Table 12–1 DBMS_DEFER Package Subprograms*

| Subprogram | Description |
|---|---|
| "CALL Procedure" on page 12-3 | Builds a deferred call to a remote procedure. |
| "COMMIT_WORK Procedure" on page 12-4 | Performs a transaction commit after checking for well-formed deferred remote procedure calls. |
| "datatype_ARG Procedure" on page 12-5 | Provides the data that is to be passed to a deferred remote procedure call. |
| "TRANSACTION Procedure" on page 12-7 | Indicates the start of a new deferred transaction. |

## CALL Procedure

This procedure builds a deferred call to a remote procedure.

### Syntax

```
DBMS_DEFER.CALL (
   schema_name       IN   VARCHAR2,
   package_name      IN   VARCHAR2,
   proc_name         IN   VARCHAR2,
   arg_count         IN   NATURAL,
   { nodes           IN   node_list_t
   | group_name      IN   VARCHAR2 :=''});
```

**Note:** This procedure is overloaded. The `nodes` and `group_name` parameters are mutually exclusive.

### Parameters

***Table 12–2 CALL Procedure Parameters***

| Parameter | Description |
|---|---|
| `schema_name` | Name of the schema in which the stored procedure is located. |
| `package_name` | Name of the package containing the stored procedure. The stored procedure must be part of a package. Deferred calls to standalone procedures are not supported. |
| `proc_name` | Name of the remote procedure to which you want to defer a call. |
| `arg_count` | Number of parameters for the procedure. You must have one call to `DBMS_DEFER.`*datatype*`_ARG` for each of these parameters.<br>**Note:** You must include all of the parameters for the procedure, even if some parameters have defaults. |
| `nodes` | A PL/SQL associative array of fully qualified database names to which you want to propagate the deferred call. The table is indexed starting at position 1 and continuing until a `NULL` entry is found, or the `no_data_found` exception is raised. The data in the table is case insensitive. This parameter is optional. |
| `group_name` | Reserved for internal use. |

### Exceptions

***Table 12–3 CALL Procedure Exceptions***

| Exception | Description |
|---|---|
| `ORA-23304 (malformedcall)` | Previous call was not correctly formed. |
| `ORA-23319` | Parameter value is not appropriate. |
| `ORA-23352` | Destination list (specified by `nodes` or by a previous `DBMS_DEFER.TRANSACTION` call) contains duplicates. |

# COMMIT_WORK Procedure

This procedure performs a transaction commit after checking for well-formed deferred remote procedure calls.

## Syntax

```
DBMS_DEFER.COMMIT_WORK (
   commit_work_comment IN VARCHAR2);
```

## Parameters

***Table 12–4 COMMIT_WORK Procedure Parameters***

| Parameter | Description |
|---|---|
| commit_work_comment | Equivalent to the COMMIT COMMENT statement in SQL. |

## Exceptions

***Table 12–5 COMMIT_WORK Procedure Exceptions***

| Exception | Description |
|---|---|
| ORA-23304 (malformedcall) | Transaction was not correctly formed or terminated. |

## *datatype*_ARG Procedure

This procedure provides the data that is to be passed to a deferred remote procedure call. Depending upon the type of the data that you must pass to a procedure, you must call one of the following procedures for each argument to the procedure.

You must specify each parameter in your procedure using the *datatype*_ARG procedure after you execute DBMS_DEFER.CALL. That is, you cannot use the default parameters for the deferred remote procedure call. For example, suppose you have the following procedure:

```
CREATE OR REPLACE PACKAGE my_pack AS
   PROCEDURE my_proc(a VARCHAR2, b VARCHAR2 DEFAULT 'SALES');
END;
/
```

When you run the DBMS_DEFER.CALL procedure, you must include a separate procedure call for each parameter in the my_proc procedure:

```
CREATE OR REPLACE PROCEDURE load_def_tx IS
   node DBMS_DEFER.NODE_LIST_T;
BEGIN
   node(1) := 'MYCOMPUTER.EXAMPLE.COM';
   node(2) := NULL;
   DBMS_DEFER.TRANSACTION(node);
   DBMS_DEFER.CALL('PR', 'MY_PACK', 'MY_PROC', 2);
   DBMS_DEFER.VARCHAR2_ARG('TEST');
   DBMS_DEFER.VARCHAR2_ARG('SALES'); -- required, cannot omit to use default
END;
/
```

**Note:**

- The ANYDATA_ARG procedure supports the following user-defined types: object types, collections, and REFs. See *Oracle Database SQL Language Reference* and *Oracle Database Object-Relational Developer's Guide* for more information about the ANYDATA data type.
- This procedure uses abbreviations for some datetime and interval data types. For example, TSTZ is used for the TIMESTAMP WITH TIME ZONE data type. For information about these abbreviations, see "Abbreviations for Datetime and Interval Data Types" on page 11-2.

## Syntax

```
DBMS_DEFER.ANYDATA_ARG      (arg  IN ANYDATA);
DBMS_DEFER.NUMBER_ARG       (arg  IN NUMBER);
DBMS_DEFER.DATE_ARG         (arg  IN DATE);
DBMS_DEFER.VARCHAR2_ARG     (arg  IN VARCHAR2);
DBMS_DEFER.CHAR_ARG         (arg  IN CHAR);
DBMS_DEFER.ROWID_ARG        (arg  IN ROWID);
DBMS_DEFER.RAW_ARG          (arg  IN RAW);
DBMS_DEFER.BLOB_ARG         (arg  IN BLOB);
DBMS_DEFER.CLOB_ARG         (arg  IN CLOB);
DBMS_DEFER.NCLOB_ARG        (arg  IN NCLOB);
DBMS_DEFER.NCHAR_ARG        (arg  IN NCHAR);
DBMS_DEFER.NVARCHAR2_ARG    (arg  IN NVARCHAR2);
```

```
DBMS_DEFER.ANY_CLOB_ARG     (arg  IN CLOB);
DBMS_DEFER.ANY_VARCHAR2_ARG (arg  IN VARCHAR2);
DBMS_DEFER.ANY_CHAR_ARG     (arg  IN CHAR);
DBMS_DEFER.IDS_ARG          (arg  IN DSINTERVAL_UNCONSTRAINED);
DBMS_DEFER.IYM_ARG          (arg  IN YMINTERVAL_UNCONSTRAINED);
DBMS_DEFER.TIMESTAMP_ARG    (arg  IN TIMESTAMP_UNCONSTRAINED);
DBMS_DEFER.TSLTZ_ARG        (arg  IN TIMESTAMP_LTZ_UNCONSTRAINED);
DBMS_DEFER.TSTZ_ARG         (arg  IN TIMESTAMP_TZ_UNCONSTRAINED);
```

## Parameters

***Table 12–6 datatype_ARG Procedure Parameters***

| Parameter | Description |
|---|---|
| arg | Value of the parameter that you want to pass to the remote procedure to which you previously deferred a call. |

## Exceptions

***Table 12–7 datatype_ARG Procedure Exceptions***

| Exception | Description |
|---|---|
| ORA-23323 | Argument value is too long. |

## TRANSACTION Procedure

This procedure indicates the start of a new deferred transaction. If you omit this call, then Oracle considers your first call to DBMS_DEFER.CALL to be the start of a new transaction.

### Syntax

```
DBMS_DEFER.TRANSACTION (
   nodes  IN   node_list_t);
```

> **Note:** This procedure is overloaded. The behavior of the version without an input parameter is similar to that of the version with an input parameter, except that the former uses the nodes in the DEFDEFAULTDEST view instead of using the nodes in the nodes parameter.

### Parameters

***Table 12–8 TRANSACTION Procedure Parameters***

| Parameter | Description |
|---|---|
| nodes | A PL/SQL associative array of fully qualified database names to which you want to propagate the deferred calls of the transaction. The table is indexed starting at position 1 and continuing until a NULL entry is found, or the no_data_found exception is raised. The data in the table is case insensitive. |

### Exceptions

***Table 12–9 TRANSACTION Procedure Exceptions***

| Exception | Description |
|---|---|
| ORA-23304 (malformedcall) | Previous transaction was not correctly formed or terminated. |
| ORA-23319 | Parameter value is not appropriate. |
| ORA-23352 | Raised by DBMS_DEFER.CALL if the node list contains duplicates. |

# 13

# DBMS_DEFER_QUERY

DBMS_DEFER_QUERY enables querying the deferred transactions queue data that is not exposed through views.

This chapter contains this topic:

- Summary of DBMS_DEFER_QUERY Subprograms

# Summary of DBMS_DEFER_QUERY Subprograms

*Table 13–1 DBMS_DEFER_QUERY Package Subprograms*

| Subprogram | Description |
|---|---|
| "GET_ARG_TYPE Function" on page 13-4 | Determines the form of an argument in a deferred call. |
| "GET_ARG_TYPE Function" on page 13-4 | Determines the type of an argument in a deferred call. |
| "GET_CALL_ARGS Procedure" on page 13-6 | Returns the text version of the various arguments for the specified call. |
| "GET_datatype_ARG Function" on page 13-8 | Determines the value of an argument in a deferred call. |
| "GET_OBJECT_NULL_ VECTOR_ARG Function" on page 13-10 | Returns the type information for a column object. |

## GET_ARG_FORM Function

This function returns the character set form of a deferred call parameter.

> **See Also:** The Advanced Replication interface's online Help for information about displaying deferred transactions and error transactions in the Advanced Replication interface in Oracle Enterprise Manager

### Syntax

```
DBMS_DEFER_QUERY.GET_ARG_FORM (
   callno                IN   NUMBER,
   arg_no                IN   NUMBER,
   deferred_tran_id      IN   VARCHAR2)
  RETURN NUMBER;
```

### Parameters

***Table 13–2 GET_ARG_FORM Function Parameters***

| Parameter | Description |
|---|---|
| callno | Call identifier from the DEFCALL view. |
| arg_no | Position of desired parameter in calls argument list. Parameter positions are 1...*number* of parameters in call. |
| deferred_tran_id | Deferred transaction identification. |

### Exceptions

***Table 13–3 GET_ARG_FORM Function Exceptions***

| Exception | Description |
|---|---|
| NO_DATA_FOUND | Input parameters do not correspond to a parameter of a deferred call. |

### Returns

***Table 13–4 GET_ARG_FORM Function Returns***

| Constant Return Value | Return Value | Possible Data Type |
|---|---|---|
| DBMS_DEFER_QUERY.ARG_FORM_NONE | 0 | DATE<br>NUMBER<br>ROWID<br>RAW<br>BLOB<br>User-defined types |
| DBMS_DEFER_QUERY.ARG_FORM_IMPLICIT | 1 | CHAR<br>VARCHAR2<br>CLOB |
| DBMS_DEFER_QUERY.ARG_FORM_NCHAR | 2 | NCHAR<br>NVARCHAR2<br>NCLOB |

# GET_ARG_TYPE Function

This function determines the type of an argument in a deferred call. The type of the deferred remote procedure call (RPC) parameter is returned.

> **See Also:** The Advanced Replication interface's online Help for information about displaying deferred transactions and error transactions in the Advanced Replication interface in Oracle Enterprise Manager

## Syntax

```
DBMS_DEFER_QUERY.GET_ARG_TYPE (
   callno            IN   NUMBER,
   arg_no            IN   NUMBER,
   deferred_tran_id  IN   VARCHAR2)
  RETURN NUMBER;
```

## Parameters

***Table 13–5 GET_ARG_TYPE Function Parameters***

| Parameter | Description |
|---|---|
| callno | Identification number from the DEFCALL view of the deferred remote procedure call. |
| arg_no | Numerical position of the argument to the call whose type you want to determine. The first argument to a procedure is in position 1. |
| deferred_tran_id | Identifier of the deferred transaction. |

## Exceptions

***Table 13–6 GET_ARG_TYPE Function Exceptions***

| Exception | Description |
|---|---|
| NO_DATA_FOUND | Input parameters do not correspond to a parameter of a deferred call. |

## Returns

***Table 13–7 GET_ARG_TYPE Function Returns***

| Constant Return Value | Return Value | Corresponding Data Type |
|---|---|---|
| DBMS_DEFER_QUERY.ARG_TYPE_VARCHAR2 | 1 | VARCHAR2 |
| DBMS_DEFER_QUERY.ARG_TYPE_NUM | 2 | NUMBER |
| DBMS_DEFER_QUERY.ARG_TYPE_ROWID | 11 | ROWID |
| DBMS_DEFER_QUERY.ARG_TYPE_DATE | 12 | DATE |
| DBMS_DEFER_QUERY.ARG_TYPE_RAW | 23 | RAW |
| DBMS_DEFER_QUERY.ARG_TYPE_CHAR | 96 | CHAR |
| DBMS_DEFER_QUERY.ARG_TYPE_ANYDATA | 109 | ANYDATA |

*Table 13–7 (Cont.) GET_ARG_TYPE Function Returns*

| Constant Return Value | Return Value | Corresponding Data Type |
|---|---|---|
| DBMS_DEFER_QUERY.ARG_TYPE_CLOB | 112 | CLOB |
| DBMS_DEFER_QUERY.ARG_TYPE_BLOB | 113 | BLOB |
| DBMS_DEFER_QUERY.ARG_TYPE_BFIL | 114 | BFILE |
| DBMS_DEFER_QUERY.ARG_TYPE_OBJECT_NULL_VECTOR | 121 | OBJECT_NULL_VECTOR |
| DBMS_DEFER_QUERY.ARG_TYPE_TIMESTAMP | 180 | TIMESTAMP |
| DBMS_DEFER_QUERY.ARG_TYPE_TSTZ | 181 | TSTZ |
| DBMS_DEFER_QUERY.ARG_TYPE_IYM | 182 | IYM |
| DBMS_DEFER_QUERY.ARG_TYPE_IDS | 183 | IDS |
| DBMS_DEFER_QUERY.ARG_TYPE_TSLTZ | 231 | TSLTZ |

**Note:**

- The ANYDATA data type supports the following user-defined types: object types, collections, and REFs. See *Oracle Database SQL Language Reference* and *Oracle Database Object-Relational Developer's Guide* for more information about the ANYDATA data type.
- This function uses abbreviations for some datetime and interval data types. For example, TSTZ is used for the TIMESTAMP WITH TIME ZONE data type. For information about these abbreviations, see "Abbreviations for Datetime and Interval Data Types" on page 11-2.

# GET_CALL_ARGS Procedure

This procedure returns the text version of the various arguments for the specified call. The text version is limited to the first 2000 bytes.

**See Also:**

- "GET_datatype_ARG Function" on page 13-8
- *Oracle Database SQL Language Reference* and *Oracle Database Object-Relational Developer's Guide* for more information about the `ANYDATA` data type

## Syntax

```
DBMS_DEFER_QUERY.GET_CALL_ARGS (
   callno    IN  NUMBER,
   startarg  IN  NUMBER := 1,
   argcnt    IN  NUMBER,
   argsize   IN  NUMBER,
   tran_db   IN  VARCHAR2,
   tran_id   IN  VARCHAR2,
   date_fmt  IN  VARCHAR2,
   types     OUT TYPE_ARY,
   forms     OUT TYPE_ARY,
   vals      OUT VAL_ARY);

DBMS_DEFER_QUERY.GET_CALL_ARGS (
   callno         IN  NUMBER,
   startarg       IN  NUMBER := 1,
   argcnt         IN  NUMBER,
   argsize        IN  NUMBER,
   tran_id        IN  VARCHAR2,
   date_fmt       IN  VARCHAR2,
   time_fmt       IN  VARCHAR2,
   ttz_fmt        IN  VARCHAR2,
   timestamp_fmt  IN  VARCHAR2,
   tstz_fmt       IN  VARCHAR2,
   types          OUT TYPE_ARY,
   forms          OUT TYPE_ARY,
   vals           OUT VAL_ARY);

DBMS_DEFER_QUERY.GET_CALL_ARGS (
   callno         IN  NUMBER,
   startarg       IN  NUMBER := 1,
   argcnt         IN  NUMBER,
   argsize        IN  NUMBER,
   tran_id        IN  VARCHAR2,
   date_fmt       IN  VARCHAR2,
   time_fmt       IN  VARCHAR2,
   ttz_fmt        IN  VARCHAR2,
   timestamp_fmt  IN  VARCHAR2,
   tstz_fmt       IN  VARCHAR2,
   tsltz_fmt      IN  VARCHAR2,
   types          OUT TYPE_ARY,
   forms          OUT TYPE_ARY,
   vals           OUT VAL_ARY);
```

**Note:** This procedure is overloaded.

## Parameters

*Table 13–8 GET_CALL_ARGS Procedure Parameters*

| Parameter | Description |
|---|---|
| callno | Identification number from the DEFCALL view of the deferred remote procedure call (RPC). |
| startarg | Numerical position of the first argument you want described. |
| argcnt | Number of arguments in the call. |
| argsize | Maximum size of returned argument. |
| tran_db | Global name of the database where the transaction originated. |
| tran_id | Identifier of the deferred transaction. |
| date_fmt | Format in which the date is returned. |
| time_fmt | Format in which the time is returned. |
| ttz_fmt | Format in which the time with time zone is returned. |
| timestamp_fmt | Format in which the timestamp is returned. |
| tstz_fmt | Format in which the timestamp with time zone is returned. |
| tsltz_fmt | Format in which the timestamp with local time zone is returned. |
| types | Array containing the types of arguments. |
| forms | Array containing the character set forms of arguments. |
| vals | Array containing the values of the arguments in a textual form. |

## Exceptions

*Table 13–9 GET_CALL_ARGS Procedure Exceptions*

| Exception | Description |
|---|---|
| NO_DATA_FOUND | Input parameters do not correspond to a parameter of a deferred call. |

# GET_*datatype*_ARG Function

This function determines the value of an argument in a deferred call.

The ANYDATA type supports the following user-defined types: object types, collections and REFs. Not all types supported by this function can be enqueued by the ANYDATA_ARG procedure in the DBMS_DEFER package.

The returned text for type arguments includes the following values: type owner, type name, type version, length, precision, scale, character set identifier, character set form, and number of elements for collections or number of attributes for object types. These values are separated by a colon (:).

**See Also:**

- "datatype_ARG Procedure" on page 12-5
- The Advanced Replication interface's online Help for information about displaying deferred transactions and error transactions in the Advanced Replication interface in Oracle Enterprise Manager
- *Oracle Database SQL Language Reference* and *Oracle Database Object-Relational Developer's Guide* for more information about the ANYDATA data type
- This function uses abbreviations for some datetime and interval data types. For example, TSTZ is used for the TIMESTAMP WITH TIME ZONE data type. For information about these abbreviations, see "Abbreviations for Datetime and Interval Data Types" on page 11-2.

## Syntax

Depending upon the type of the argument value that you want to retrieve, the syntax for the appropriate function is as follows. Each of these functions returns the value of the specified argument.

```
DBMS_DEFER_QUERY.GET_datatype_ARG (
   callno             IN   NUMBER,
   arg_no             IN   NUMBER,
   deferred_tran_id   IN   VARCHAR2 DEFAULT NULL)
  RETURN datatype;
```

where *datatype* is:

```
{ ANYDATA
| NUMBER
| VARCHAR2
| CHAR
| DATE
| RAW
| ROWID
| BLOB
| CLOB
| NCLOB
| NCHAR
| NVARCHAR2
| IDS
| IYM
| TIMESTAMP
```

```
| TSLTZ
| TSTZ }
```

## Parameters

***Table 13–10 GET_datatype_ARG Function Parameters***

| Parameter | Description |
|---|---|
| callno | Identification number from the DEFCALL view of the deferred remote procedure call. |
| arg_no | Numerical position of the argument to the call whose value you want to determine. The first argument to a procedure is in position 1. |
| deferred_tran_id | Identifier of the deferred transaction. Defaults to the last transaction identifier passed to the GET_ARG_TYPE function. The default is NULL. |

## Exceptions

***Table 13–11 GET_datatype_ARG Function Exceptions***

| Exception | Description |
|---|---|
| NO_DATA_FOUND | Input parameters do not correspond to a parameter of a deferred call. |
| ORA-26564 | Argument in this position is not of the specified type or is not one of the types supported by the ANYDATA type. |

# GET_OBJECT_NULL_VECTOR_ARG Function

This function returns the type information for a column object, including the type owner, name, and hashcode.

## Syntax

```
DBMS_DEFER_QUERY.GET_OBJECT-NULL_VECTOR_ARG (
   callno                IN   NUMBER,
   arg_no                IN   NUMBER,
   deferred_tran_id      IN   VARCHAR2)
  RETURN SYSTEM.REPCAT$_OBJECT_NULL_VECTOR;
```

## Parameters

***Table 13–12 GET_OBJECT_NULL_VECTOR_ARG Function Parameters***

| Parameter | Description |
|---|---|
| callno | Call identifier from the DEFCALL view. |
| arg_no | Position of desired parameter in calls argument list. Parameter positions are 1...*number* of parameters in call. |
| deferred_tran_id | Deferred transaction identification. |

## Exceptions

***Table 13–13 GET_OBJECT_NULL_VECTOR_ARG Function Exceptions***

| Exception | Description |
|---|---|
| NO_DATA_FOUND | Input parameters do not correspond to a parameter of a deferred call. |
| ORA-26564 | Parameter is not an object_null_vector type. |

## Returns

***Table 13–14 GET_OBJECT_NULL_VECTOR_ARG Function Returns***

| Return Value | Type Definition |
|---|---|
| SYSTEM.REPCAT$_OBJECT_NULL_VECTOR type | CREATE TYPE<br>SYSTEM.REPCAT$_OBJECT_NULL_VECTOR<br>AS OBJECT (<br>  type_owner  VARCHAR2(30),<br>  type_name  VARCHAR2(30),<br>  type_hashcode  RAW(17),<br>  null_vector  RAW(2000)); |

# 14

# DBMS_DEFER_SYS

DBMS_DEFER_SYS procedures manage default replication node lists. This package is the system administrator interface to a replicated transactional deferred remote procedure call facility. Administrators and replication daemons can execute transactions queued for remote nodes using this facility, and administrators can control the nodes to which remote calls are destined.

This chapter contains this topic:


- Summary of DBMS_DEFER_SYS Subprograms

# Summary of DBMS_DEFER_SYS Subprograms

*Table 14–1 DBMS_DEFER_SYS Package Subprograms*

| Subprogram | Description |
|---|---|
| "ADD_DEFAULT_DEST Procedure" on page 14-4 | Adds a destination database to the DEFDEFAULTDEST view. |
| "CLEAR_PROP_ STATISTICS Procedure" on page 14-5 | Clears the propagation statistics in the DEFSCHEDULE data dictionary view. |
| "DELETE_DEFAULT_ DEST Procedure" on page 14-6 | Removes a destination database from the DEFDEFAULTDEST view. |
| "DELETE_DEF_ DESTINATION Procedure" on page 14-7 | Removes a destination database from the DEFSCHEDULE view. |
| "DELETE_ERROR Procedure" on page 14-8 | Deletes a transaction from the DEFERROR view. |
| "DELETE_TRAN Procedure" on page 14-9 | Deletes a transaction from the DEFTRANDEST view. |
| "DISABLED Function" on page 14-10 | Determines whether propagation of the deferred transaction queue from the current site to a specified site is enabled. |
| EXCLUDE_PUSH Function on page 14-11 | Acquires an exclusive lock that prevents deferred transaction PUSH. |
| "EXECUTE_ERROR Procedure" on page 14-12 | Reexecutes a deferred transaction that did not initially complete successfully in the security context of the original receiver of the transaction. |
| "EXECUTE_ERROR_AS_ USER Procedure" on page 14-13 | Reexecutes a deferred transaction that did not initially complete successfully in the security context of the user who executes this procedure. |
| "PURGE Function" on page 14-14 | Purges pushed transactions from the deferred transaction queue at your current master site or materialized view site. |
| "PUSH Function" on page 14-16 | Forces a deferred remote procedure call queue at your current master site or materialized view site to be pushed to a remote site. |
| "REGISTER_ PROPAGATOR Procedure" on page 14-19 | Registers the specified user as the propagator for the local database. |
| "SCHEDULE_PURGE Procedure" on page 14-20 | Schedules a job to purge pushed transactions from the deferred transaction queue at your current master site or materialized view site. |
| "SCHEDULE_PUSH Procedure" on page 14-22 | Schedules a job to push the deferred transaction queue to a remote site. |
| "SET_DISABLED Procedure" on page 14-24 | Disables or enables propagation of the deferred transaction queue from the current site to a specified destination site. |
| "UNREGISTER_ PROPAGATOR Procedure" on page 14-26 | Unregisters a user as the propagator from the local database. |
| "UNSCHEDULE_PURGE Procedure" on page 14-27 | Stops automatic purges of pushed transactions from the deferred transaction queue at a master site or materialized view site. |

*Table 14–1 (Cont.) DBMS_DEFER_SYS Package Subprograms*

| Subprogram | Description |
|---|---|
| "UNSCHEDULE_PUSH Procedure" on page 14-28 | Stops automatic pushes of the deferred transaction queue from a master site or materialized view site to a remote site. |

# ADD_DEFAULT_DEST Procedure

This procedure adds a destination database to the DEFDEFAULTDEST data dictionary view.

## Syntax

```
DBMS_DEFER_SYS.ADD_DEFAULT_DEST (
   dblink   IN   VARCHAR2);
```

## Parameters

***Table 14–2 ADD_DEFAULT_DEST Procedure Parameters***

| Parameter | Description |
|---|---|
| dblink | The fully qualified database name of the node that you want to add to the DEFDEFAULTDEST view. |

## Exceptions

***Table 14–3 ADD_DEFAULT_DEST Procedure Exceptions***

| Exception | Description |
|---|---|
| ORA-23352 | The dblink that you specified is already in the default list. |

## CLEAR_PROP_STATISTICS Procedure

This procedure clears the propagation statistics in the DEFSCHEDULE data dictionary view. When this procedure is executed successfully, all statistics in this view are returned to zero and statistic gathering starts fresh.

Specifically, this procedure clears statistics from the following columns in the DEFSCHEDULE data dictionary view:

- TOTAL_TXN_COUNT
- AVG_THROUGHPUT
- AVG_LATENCY
- TOTAL_BYTES_SENT
- TOTAL_BYTES_RECEIVED
- TOTAL_ROUND_TRIPS
- TOTAL_ADMIN_COUNT
- TOTAL_ERROR_COUNT
- TOTAL_SLEEP_TIME

### Syntax

```
DBMS_DEFER_SYS.CLEAR_PROP_STATISTICS (
   dblink   IN   VARCHAR2);
```

### Parameters

*Table 14–4 CLEAR_PROP_STATISTICS Procedure Parameters*

| Parameter | Description |
|---|---|
| dblink | The fully qualified database name of the node whose statistics you want to clear. The statistics to be cleared are the statistics for propagation of deferred transactions from the current node to the node you specify for dblink. |

# DELETE_DEFAULT_DEST Procedure

This procedure removes a destination database from the DEFDEFAULTDEST view.

## Syntax

```
DBMS_DEFER_SYS.DELETE_DEFAULT_DEST (
   dblink   IN   VARCHAR2);
```

## Parameters

***Table 14–5 DELETE_DEFAULT_DEST Procedure Parameters***

| Parameter | Description |
|---|---|
| dblink | The fully qualified database name of the node that you want to delete from the DEFDEFAULTDEST view. If Oracle does not find this dblink in the view, then no action is taken. |

## DELETE_DEF_DESTINATION Procedure

This procedure removes a destination database from the DEFSCHEDULE view.

### Syntax

```
DBMS_DEFER_SYS.DELETE_DEF_DESTINATION (
   destination   IN   VARCHAR2,
   force         IN   BOOLEAN := FALSE);
```

### Parameters

***Table 14–6 DELETE_DEF_DESTINATION Procedure Parameters***

| Parameter | Description |
|---|---|
| destination | The fully qualified database name of the destination that you want to delete from the DEFSCHEDULE view. If Oracle does not find this destination in the view, then no action is taken. |
| force | When set to TRUE, Oracle ignores all safety checks and deletes the destination. |

# DELETE_ERROR Procedure

This procedure deletes a transaction from the DEFERROR view.

## Syntax

```
DBMS_DEFER_SYS.DELETE_ERROR(
   deferred_tran_id     IN   VARCHAR2,
   destination          IN   VARCHAR2);
```

## Parameters

***Table 14–7 DELETE_ERROR Procedure Parameters***

| Parameter | Description |
|---|---|
| deferred_tran_id | Identification number from the DEFERROR view of the deferred transaction that you want to remove from the DEFERROR view. If this parameter is NULL, then all transactions meeting the requirements of the other parameter are removed. |
| destination | The fully qualified database name from the DEFERROR view of the database to which the transaction was originally queued. If this parameter is NULL, then all transactions meeting the requirements of the other parameter are removed from the DEFERROR view. |

## DELETE_TRAN Procedure

This procedure deletes a transaction from the DEFTRANDEST view. If there are no other DEFTRANDEST or DEFERROR entries for the transaction, then the transaction is deleted from the DEFTRAN and DEFCALL views as well.

### Syntax

```
DBMS_DEFER_SYS.DELETE_TRAN (
   deferred_tran_id     IN   VARCHAR2,
   destination          IN   VARCHAR2);
```

### Parameters

***Table 14–8 DELETE_TRAN Procedure Parameters***

| Parameter | Description |
|---|---|
| deferred_tran_id | Identification number from the DEFTRAN view of the deferred transaction that you want to delete. If this is NULL, then all transactions meeting the requirements of the other parameter are deleted. |
| destination | The fully qualified database name from the DEFTRANDEST view of the database to which the transaction was originally queued. If this is NULL, then all transactions meeting the requirements of the other parameter are deleted. |

# DISABLED Function

This function determines whether propagation of the deferred transaction queue from the current site to a specified site is enabled. The DISABLED function returns TRUE if the deferred remote procedure call (RPC) queue is disabled for the specified destination.

## Syntax

```
DBMS_DEFER_SYS.DISABLED (
   destination  IN   VARCHAR2)
  RETURN BOOLEAN;
```

## Parameters

***Table 14–9 DISABLED Function Parameters***

| Parameter | Description |
|---|---|
| destination | The fully qualified database name of the node whose propagation status you want to check. |

## Exceptions

***Table 14–10 DISABLED Function Exceptions***

| Exception | Description |
|---|---|
| NO_DATA_FOUND | Specified destination does not appear in the DEFSCHEDULE view. |

## Returns

***Table 14–11 DISABLED Function Return Values***

| Value | Description |
|---|---|
| TRUE | Propagation to this site from the current site is disabled. |
| FALSE | Propagation to this site from the current site is enabled. |

## EXCLUDE_PUSH Function

This function acquires an exclusive lock that prevents deferred transaction PUSH (either serial or parallel). This function performs a commit when acquiring the lock. The lock is acquired with RELEASE_ON_COMMIT => TRUE, so that pushing of the deferred transaction queue can resume after the next commit.

### Syntax

```
DBMS_DEFER_SYS.EXCLUDE_PUSH (
   timeout   IN   INTEGER)
  RETURN INTEGER;
```

### Parameters

*Table 14–12 EXCLUDE_PUSH Function Parameters*

| Parameter | Description |
|---|---|
| timeout | Timeout in seconds. If the lock cannot be acquired within this time period (either because of an error or because a PUSH is currently under way), then the call returns a value of 1. A timeout value of DBMS_LOCK.MAXWAIT waits indefinitely. |

### Returns

*Table 14–13 EXCLUDE_PUSH Function Return Values*

| Value | Description |
|---|---|
| 0 | Success, lock acquired. |
| 1 | Timeout, no lock acquired. |
| 2 | Deadlock, no lock acquired. |
| 4 | Already own lock. |

# EXECUTE_ERROR Procedure

This procedure reexecutes a deferred transaction that did not initially complete successfully in the security context of the original receiver of the transaction.

## Syntax

```
DBMS_DEFER_SYS.EXECUTE_ERROR (
   deferred_tran_id IN   VARCHAR2,
   destination      IN   VARCHAR2);
```

## Parameters

***Table 14–14 EXECUTE_ERROR Procedure Parameters***

| Parameter | Description |
|---|---|
| deferred_tran_id | Identification number from the DEFERROR view of the deferred transaction that you want to reexecute. If this is NULL, then all transactions queued for destination are reexecuted. |
| destination | The fully qualified database name from the DEFERROR view of the database to which the transaction was originally queued. This must not be NULL. If the provided database name is not fully qualified or is invalid, no error will be raised. |

## Exceptions

***Table 14–15 EXECUTE_ERROR Procedure Exceptions***

| Exception | Description |
|---|---|
| ORA-24275 error | Illegal combinations of NULL and non-NULL parameters were used. |
| badparam | Parameter value missing or invalid (for example, if destination is NULL). |
| missinguser | Invalid user. |

## EXECUTE_ERROR_AS_USER Procedure

This procedure reexecutes a deferred transaction that did not initially complete successfully. Each transaction is executed in the security context of the connected user.

### Syntax

```
DBMS_DEFER_SYS.EXECUTE_ERROR_AS_USER (
   deferred_tran_id IN   VARCHAR2,
   destination      IN   VARCHAR2);
```

### Parameters

*Table 14–16 EXECUTE_ERROR_AS_USER Procedure Parameters*

| Parameter | Description |
|---|---|
| deferred_tran_id | Identification number from the DEFERROR view of the deferred transaction that you want to reexecute. If this is NULL, then all transactions queued for destination are reexecuted. |
| destination | The fully qualified database name from the DEFERROR view of the database to which the transaction was originally queued. This must not be NULL. |

### Exceptions

*Table 14–17 EXECUTE_ERROR_AS_USER Procedure Exceptions*

| Exception | Description |
|---|---|
| ORA-24275 error | Illegal combinations of NULL and non-NULL parameters were used. |
| badparam | Parameter value missing or invalid (for example, if destination is NULL). |
| missinguser | Invalid user. |

# PURGE Function

This function purges pushed transactions from the deferred transaction queue at your current master site or materialized view site.

## Syntax

```
DBMS_DEFER_SYS.PURGE (
   purge_method         IN  BINARY_INTEGER := purge_method_quick,
   rollback_segment     IN  VARCHAR2       := NULL,
   startup_seconds      IN  BINARY_INTEGER := 0,
   execution_seconds    IN  BINARY_INTEGER := seconds_infinity,
   delay_seconds        IN  BINARY_INTEGER := 0,
   transaction_count    IN  BINARY_INTEGER := transactions_infinity,
   write_trace          IN  BOOLEAN        := NULL);
  RETURN BINARY_INTEGER;
```

## Parameters

***Table 14–18 PURGE Function Parameters***

| Parameter | Description |
|---|---|
| purge_method | Controls how to purge the deferred transaction queue: purge_method_quick costs less, while purge_method_precise offers better precision.<br>Specify the following for this parameter to use purge_method_quick:<br>dbms_defer_sys.purge_method_quick<br>Specify the following for this parameter to user purge_method_precise:<br>dbms_defer_sys.purge_method_precise<br>If you use purge_method_quick, deferred transactions and deferred procedure calls that have been successfully pushed can remain in the DEFTRAN and DEFCALL data dictionary views for longer than expected before they are purged. See "Usage Notes" on page 14-15 for more information. |
| rollback_segment | Name of rollback segment to use for the purge, or NULL for default. |
| startup_seconds | Maximum number of seconds to wait for a previous purge of the same deferred transaction queue. |
| execution_seconds | If > 0, then stop purge cleanly after the specified number of seconds of real time. |
| delay_seconds | Stop purge cleanly after the deferred transaction queue has no transactions to purge for delay_seconds. |
| transaction_count | If > 0, then shut down cleanly after purging transaction_count number of transactions. |
| write_trace | When set to TRUE, Oracle records the result value returned by the PURGE function in the server's trace file. When set to FALSE, Oracle does not record the result value. |

## Returns

*Table 14–19 Purge Function Returns*

| Value | Description |
|---|---|
| result_ok | OK, terminated after delay_seconds expired. |
| result_startup_seconds | Terminated by lock timeout while starting. |
| result_execution_seconds | Terminated by exceeding execution_seconds. |
| result_transaction_count | Terminated by exceeding transaction_count. |
| result_errors | Terminated after errors. |
| result_split_del_order_limit | Terminated after failing to acquire the enqueue in exclusive mode. If you receive this return code, then retry the purge. If the problem persists, then contact Oracle Support Services. |
| result_purge_disabled | Queue purging is disabled internally for synchronization when adding new master sites without quiesce. |

## Exceptions

*Table 14–20 PURGE Function Exceptions*

| Exception | Description |
|---|---|
| argoutofrange | Parameter value is out of a valid range. |
| executiondisabled | Execution of purging is disabled. |
| defererror | Internal error. |

## Usage Notes

When you use the purge_method_quick for the purge_method parameter in the DBMS_DEFER_SYS.PURGE function, deferred transactions and deferred procedure calls can remain in the DEFCALL and DEFTRAN data dictionary views after they have been successfully pushed. This behavior occurs in replication environments that have multiple database links and the push is executed to only one database link.

To purge the deferred transactions and deferred procedure calls, perform one of the following actions:

- Use purge_method_precise for the purge_method parameter instead of the purge_method_quick. Using purge_method_precise is more expensive, but it ensures that the deferred transactions and procedure calls are purged after they have been successfully pushed.
- Using purge_method_quick for the purge_method parameter, push the deferred transactions to all database links. The deferred transactions and deferred procedure calls are purged efficiently when the push to the last database link is successful.

# PUSH Function

This function forces a deferred remote procedure call (RPC) queue at your current master site or materialized view site to be pushed (propagated) to a remote site using either serial or parallel propagation.

## Syntax

```
DBMS_DEFER_SYS.PUSH (
   destination          IN  VARCHAR2,
   parallelism          IN  BINARY_INTEGER := 0,
   heap_size            IN  BINARY_INTEGER := 0,
   stop_on_error        IN  BOOLEAN        := FALSE,
   write_trace          IN  BOOLEAN        := FALSE,
   startup_seconds      IN  BINARY_INTEGER := 0,
   execution_seconds    IN  BINARY_INTEGER := seconds_infinity,
   delay_seconds        IN  BINARY_INTEGER := 0,
   transaction_count    IN  BINARY_INTEGER := transactions_infinity,
   delivery_order_limit IN  NUMBER         := delivery_order_infinity)
  RETURN BINARY_INTEGER;
```

## Parameters

***Table 14–21 PUSH Function Parameters***

| Parameter | Description |
|---|---|
| `destination` | The fully qualified database name of the master site or master materialized view site to which you are forwarding changes. |
| `parallelism` | 0 specifies serial propagation.<br>*n* > 1 specifies parallel propagation with *n* parallel processes.<br>1 specifies parallel propagation using only one parallel process. |
| `heap_size` | Maximum number of transactions to be examined simultaneously for parallel propagation scheduling. Oracle automatically calculates the default setting for optimal performance.<br>**Note:** Do not set the parameter unless so directed by Oracle Support Services. |
| `stop_on_error` | The default, `FALSE`, indicates that the executor should continue even if errors, such as conflicts, are encountered. If `TRUE`, then stops propagation at the first indication that a transaction encountered an error at the destination site.<br>**Note:** If `stop_on_error` is set to `TRUE` and the `parallelism` parameter is greater than `0` (zero), then transactions might continue to be propagated and applied for a period of time after an error is encountered. |
| `write_trace` | When set to `TRUE`, Oracle records the result value returned by the function in the server's trace file. When set to `FALSE`, Oracle does not record the result value. |
| `startup_seconds` | Maximum number of seconds to wait for a previous push to the same destination. |

*Table 14–21 (Cont.) PUSH Function Parameters*

| Parameter | Description |
|---|---|
| execution_seconds | If > 0, then stop push cleanly after the specified number of seconds of real time. If transaction_count and execution_seconds are zero (the default), then transactions are executed until there are no more in the queue.<br><br>The execution_seconds parameter only controls the duration of time that operations can be started. It does not include the amount of time that the transactions require at remote sites. Therefore, the execution_seconds parameter is not intended to be used as a precise control to stop the propagation of transactions to a remote site. If a precise control is required, use the transaction_count or delivery_order parameters. |
| delay_seconds | Do not return before the specified number of seconds have elapsed, even if the queue is empty. Useful for reducing execution overhead if PUSH is called from a tight loop. |
| transaction_count | If > 0, then the maximum number of transactions to be pushed before stopping. If transaction_count and execution_seconds are zero (the default), then transactions are executed until there are no more in the queue that must be pushed. |
| delivery_order_limit | Stop execution cleanly before pushing a transaction where delivery_order >= delivery_order_limit. |

## Returns

*Table 14–22 PUSH Function Returns*

| Value | Description |
|---|---|
| result_ok | OK, terminated after delay_seconds expired. |
| result_startup_seconds | Terminated by lock timeout while starting. |
| result_execution_seconds | Terminated by exceeding execution_seconds. |
| result_transaction_count | Terminated by exceeding transaction_count. |
| result_delivery_order_limit | Terminated by exceeding delivery_order_limit. |
| result_errors | Terminated after errors. |
| result_push_disabled | Push was disabled internally. Typically, this return value means that propagation to the destination was set to disabled internally by Oracle for propagation synchronization when adding a new master site to a master group without quiescing the master group. Oracle will enable propagation automatically at a later time. |
| result_split_del_order_limit | Terminated after failing to acquire the enqueue in exclusive mode. If you receive this return code, then retry the push. If the problem persists, then contact Oracle Support Services. |

## Exceptions

*Table 14–23 PUSH Function Exceptions*

| Exception | Description |
|---|---|
| incompleteparallelpush | Serial propagation requires that parallel propagation shuts down cleanly. |
| executiondisabled | Execution of deferred remote procedure calls (RPCs) is disabled at the destination. |
| crt_err_err | Error while creating entry in DEFERROR. |
| deferred_rpc_quiesce | Replication activity for replication group is suspended. |
| commfailure | Communication failure during deferred remote procedure call (RPC). |
| missingpropagator | A propagator does not exist. |

## REGISTER_PROPAGATOR Procedure

This procedure registers the specified user as the propagator for the local database. It also grants the following privileges to the specified user (so that the user can create wrappers):

- CREATE SESSION
- CREATE PROCEDURE
- CREATE DATABASE LINK
- EXECUTE ANY PROCEDURE

### Syntax

```
DBMS_DEFER_SYS.REGISTER_PROPAGATOR (
   username  IN  VARCHAR2);
```

### Parameter

*Table 14–24 REGISTER_PROPAGATOR Procedure Parameter*

| Parameter | Description |
|---|---|
| username | Name of the user. |

### Exceptions

*Table 14–25 REGISTER_PROPAGATOR Procedure Exceptions*

| Exception | Description |
|---|---|
| missinguser | Specified user does not exist. |
| alreadypropagator | Specified user is already the propagator. |
| duplicatepropagator | There is already a different propagator. |

# SCHEDULE_PURGE Procedure

This procedure schedules a job to purge pushed transactions from the deferred transaction queue at your current master site or materialized view site. You should schedule one purge job.

> **See Also:** *Oracle Database Advanced Replication* for information about using this procedure to schedule continuous or periodic purge of your deferred transaction queue

## Syntax

```
DBMS_DEFER_SYS.SCHEDULE_PURGE (
   interval             IN  VARCHAR2,
   next_date            IN  DATE,
   reset                IN  BOOLEAN        := NULL,
   purge_method         IN  BINARY_INTEGER := NULL,
   rollback_segment     IN  VARCHAR2       := NULL,
   startup_seconds      IN  BINARY_INTEGER := NULL,
   execution_seconds    IN  BINARY_INTEGER := NULL,
   delay_seconds        IN  BINARY_INTEGER := NULL,
   transaction_count    IN  BINARY_INTEGER := NULL,
   write_trace          IN  BOOLEAN        := NULL);
```

## Parameters

***Table 14–26 SCHEDULE_PURGE Procedure Parameters***

| Parameter | Description |
|---|---|
| interval | Allows you to provide a function to calculate the next time to purge. This value is stored in the interval field of the DEFSCHEDULE view and calculates the next_date field of this view. If you use the default value for this parameter, NULL, then the value of this field remains unchanged. If the field had no previous value, it is created with a value of NULL. If you do not supply a value for this field, you must supply a value for next_date. |
| next_date | Allows you to specify a time to purge pushed transactions from the site's queue. This value is stored in the next_date field of the DEFSCHEDULE view. If you use the default value for this parameter, NULL, then the value of this field remains unchanged. If this field had no previous value, it is created with a value of NULL. If you do not supply a value for this field, then you must supply a value for interval. |
| reset | Set to TRUE to reset LAST_TXN_COUNT, LAST_ERROR, and LAST_MSG to NULL. |

*Table 14–26 (Cont.) SCHEDULE_PURGE Procedure Parameters*

| Parameter | Description |
|---|---|
| purge_method | Controls how to purge the deferred transaction queue: purge_method_quick costs less, while purge_method_precise offers better precision.<br>Specify the following for this parameter to use purge_method_quick:<br>dbms_defer_sys.purge_method_quick<br>Specify the following for this parameter to user purge_method_precise:<br>dbms_defer_sys.purge_method_precise<br>If you use purge_method_quick, deferred transactions and deferred procedure calls that have been successfully pushed can remain in the DEFTRAN and DEFCALL data dictionary views for longer than expected before they are purged. For more information, see "Usage Notes" on page 14-15. These usage notes are for the DBMS_DEFER_SYS.PURGE function, but they also apply to the DBMS_DEFER_SYS.SCHEDULE_PURGE procedure. |
| rollback_segment | Name of rollback segment to use for the purge, or NULL for default. |
| startup_seconds | Maximum number of seconds to wait for a previous purge of the same deferred transaction queue. |
| execution_seconds | If >0, then stop purge cleanly after the specified number of seconds of real time. |
| delay_seconds | Stop purge cleanly after the deferred transaction queue has no transactions to purge for delay_seconds. |
| transaction_count | If > 0, then shut down cleanly after purging transaction_count number of transactions. |
| write_trace | When set to TRUE, Oracle records the result value returned by the PURGE function in the server's trace file. |

# SCHEDULE_PUSH Procedure

This procedure schedules a job to push the deferred transaction queue to a remote site. This procedure performs a COMMIT.

> **See Also:** *Oracle Database Advanced Replication* for information about using this procedure to schedule continuous or periodic push of your deferred transaction queue

## Syntax

```
DBMS_DEFER_SYS.SCHEDULE_PUSH (
   destination          IN  VARCHAR2,
   interval             IN  VARCHAR2,
   next_date            IN  DATE,
   reset                IN  BOOLEAN        := FALSE,
   parallelism          IN  BINARY_INTEGER := NULL,
   heap_size            IN  BINARY_INTEGER := NULL,
   stop_on_error        IN  BOOLEAN        := NULL,
   write_trace          IN  BOOLEAN        := NULL,
   startup_seconds      IN  BINARY_INTEGER := NULL,
   execution_seconds    IN  BINARY_INTEGER := NULL,
   delay_seconds        IN  BINARY_INTEGER := NULL,
   transaction_count    IN  BINARY_INTEGER := NULL);
```

## Parameters

***Table 14–27 SCHEDULE_PUSH Procedure Parameters***

| Parameter | Description |
|---|---|
| destination | The fully qualified database name of the master site or master materialized view site to which you are forwarding changes. |
| interval | Allows you to provide a function to calculate the next time to push. This value is stored in the interval field of the DEFSCHEDULE view and calculates the next_date field of this view. If you use the default value for this parameter, NULL, then the value of this field remains unchanged. If the field had no previous value, it is created with a value of NULL. If you do not supply a value for this field, then you must supply a value for next_date. |
| next_date | Allows you to specify a time to push deferred transactions to the remote site. This value is stored in the next_date field of the DEFSCHEDULE view. If you use the default value for this parameter, NULL, then the value of this field remains unchanged. If this field had no previous value, then it is created with a value of NULL. If you do not supply a value for this field, then you must supply a value for interval. |
| reset | Set to TRUE to reset LAST_TXN_COUNT, LST_ERROR, and LAST_MSG to NULL. |
| parallelism | 0 specifies serial propagation.<br>*n* > 1 specifies parallel propagation with *n* parallel processes.<br>1 specifies parallel propagation using only one parallel process. |
| heap_size | Maximum number of transactions to be examined simultaneously for parallel propagation scheduling. Oracle automatically calculates the default setting for optimal performance.<br>**Note:** Do not set the parameter unless so directed by Oracle Support Services. |

*Table 14–27 (Cont.) SCHEDULE_PUSH Procedure Parameters*

| Parameter | Description |
|---|---|
| stop_on_error | The default, FALSE, indicates that the executor should continue even if errors, such as conflicts, are encountered. If TRUE, then stops propagation at the first indication that a transaction encountered an error at the destination site.<br><br>**Note:** If stop_on_error is set to TRUE and the parallelism parameter is greater than 0 (zero), then transactions might continue to be propagated and applied for a period of time after an error is encountered. |
| write_trace | When set to TRUE, Oracle records the result value returned by the function in the server's trace file. |
| startup_seconds | Maximum number of seconds to wait for a previous push to the same destination. |
| execution_seconds | If >0, then stop execution cleanly after the specified number of seconds of real time. If transaction_count and execution_seconds are zero (the default), then transactions are executed until there are no more in the queue. |
| delay_seconds | Do not return before the specified number of seconds have elapsed, even if the queue is empty. Useful for reducing execution overhead if PUSH is called from a tight loop. |
| transaction_count | If > 0, then the maximum number of transactions to be pushed before stopping. If transaction_count and execution_seconds are zero (the default), then transactions are executed until there are no more in the queue that must be pushed. |

# SET_DISABLED Procedure

To disable or enable propagation of the deferred transaction queue from the current site to a specified destination site. If the disabled parameter is TRUE, then the procedure disables propagation to the specified destination and future invocations of PUSH do not push the deferred remote procedure call (RPC) queue. SET_DISABLED eventually affects a session already pushing the queue to the specified destination, but does not affect sessions appending to the queue with DBMS_DEFER.

If the disabled parameter is FALSE, then the procedure enables propagation to the specified destination and, although this does not push the queue, it permits future invocations of PUSH to push the queue to the specified destination. Whether the disabled parameter is TRUE or FALSE, a COMMIT is required for the setting to take effect in other sessions.

## Syntax

```
DBMS_DEFER_SYS.SET_DISABLED (
   destination   IN   VARCHAR2,
   disabled      IN   BOOLEAN := TRUE,
   catchup       IN   RAW := '00',
   override      IN   BOOLEAN := FALSE);
```

## Parameters

***Table 14–28 SET_DISABLED Procedure Parameters***

| Parameter | Description |
|---|---|
| destination | The fully qualified database name of the node whose propagation status you want to change. |
| disabled | By default, this parameter disables propagation of the deferred transaction queue from your current site to the specified destination. Set this to FALSE to enable propagation. |
| catchup | The extension identifier for adding new master sites to a master group without quiescing the master group. The new master site is the destination. Query the DEFSCHEDULE data dictionary view for the existing extension identifiers. |
| override | A FALSE setting, the default, specifies that Oracle raises the cantsetdisabled exception if the disabled parameter is set to FALSE and propagation was disabled internally by Oracle.<br>A TRUE setting specifies that Oracle ignores whether the disabled state was set internally for synchronization and always tries to set the state as specified by the disabled parameter.<br>**Note:** Do not set this parameter unless directed to do so by Oracle Support Services. |

## Exceptions

***Table 14–29 SET_DISABLED Procedure Exceptions***

| Exception | Description |
|---|---|
| NO_DATA_FOUND | No entry was found in the DEFSCHEDULE view for the specified destination. |

*Table 14–29 (Cont.) SET_DISABLED Procedure Exceptions*

| Exception | Description |
|---|---|
| `cantsetdisabled` | The disabled status for this site is set internally by Oracle for synchronization during adding a new master site to a master group without quiescing the master group. Ensure that adding a new master site without quiescing finished before invoking this procedure. |

# UNREGISTER_PROPAGATOR Procedure

To unregister a user as the propagator from the local database. This procedure:

- Deletes the specified propagator from DEFPROPAGATOR.
- Revokes privileges granted by REGISTER_PROPAGATOR from the specified user (including identical privileges granted independently).
- Drops any generated wrappers in the schema of the specified propagator, and marks them as dropped in the replication catalog.

## Syntax

```
DBMS_DEFER_SYS.UNREGISTER_PROPAGATOR (
   username  IN  VARCHAR2
   timeout   IN  INTEGER DEFAULT DBMS_LOCK.MAXWAIT);
```

## Parameters

***Table 14–30 UNREGISTER_PROPAGATOR Procedure Parameters***

| Parameter | Description |
|---|---|
| username | Name of the propagator user. |
| timeout | Timeout in seconds. If the propagator is in use, then the procedure waits until timeout. The default is DBMS_LOCK.MAXWAIT. |

## Exceptions

***Table 14–31 UNREGISTER_PROPAGATOR Procedure Exceptions***

| Parameter | Description |
|---|---|
| missingpropagator | Specified user is not a propagator. |
| propagator_inuse | Propagator is in use, and thus cannot be unregistered. Try later. |

## UNSCHEDULE_PURGE Procedure

This procedure stops automatic purges of pushed transactions from the deferred transaction queue at a master site or materialized view site.

### Syntax

```
DBMS_DEFER_SYS.UNSCHEDULE_PURGE();
```

### Parameters

None

# UNSCHEDULE_PUSH Procedure

This procedure stops automatic pushes of the deferred transaction queue from a master site or materialized view site to a remote site.

## Syntax

```
DBMS_DEFER_SYS.UNSCHEDULE_PUSH (
   dblink   IN   VARCHAR2);
```

## Parameters

***Table 14–32 UNSCHEDULE_PUSH Procedure Parameters***

| Parameter | Description |
|---|---|
| dblink | Fully qualified path name for the database at which you want to unschedule periodic execution of deferred remote procedure calls. |

## Exceptions

***Table 14–33 UNSCHEDULE_PUSH Procedure Exceptions***

| Exception | Description |
|---|---|
| NO_DATA_FOUND | No entry was found in the DEFSCHEDULE view for the specified dblink. |

# 15

# DBMS_OFFLINE_OG

The DBMS_OFFLINE_OG package contains public APIs for offline instantiation of master groups.

This chapter contains this topic:

- Summary of DBMS_OFFLINE_OG Subprograms

---

**Note:** These procedures are used in performing an offline instantiation of a master table in a multimaster replication environment.

These procedure should not be confused with the procedures in the DBMS_REPCAT_INSTANTIATE package (used for instantiating a deployment template). See the documentation for this package for more information about its use.

---

# Summary of DBMS_OFFLINE_OG Subprograms

*Table 15–1 DBMS_OFFLINE_OG Package Subprograms*

| Subprogram | Description |
|---|---|
| "BEGIN_INSTANTIATION Procedure" on page 15-3 | Starts offline instantiation of a master group. |
| "BEGIN_LOAD Procedure" on page 15-5 | Disables triggers while data is imported to new master site as part of offline instantiation. |
| "END_INSTANTIATION Procedure" on page 15-6 | Completes offline instantiation of a master group. |
| "END_LOAD Procedure" on page 15-7 | Reenables triggers after importing data to new master site as part of offline instantiation. |
| "RESUME_SUBSET_OF_ MASTERS Procedure" on page 15-9 | Resumes replication activity at all existing sites except the new site during offline instantiation of a master group. |

## BEGIN_INSTANTIATION Procedure

This procedure starts offline instantiation of a master group. You must call this procedure from the master definition site.

---

**Note:** This procedure is used to perform an offline instantiation of a master table in a multimaster replication environment.

This procedure should not be confused with the procedures in the DBMS_REPCAT_INSTANTIATE package (used for instantiating a deployment template). See the documentation for this package for more information about its use.

---

**See Also:** "Adding New Master Sites with Offline Instantiation Using Export/Import" on page 7-25 for information about adding a new master site to a master group by performing an offline instantiation of a master site

### Syntax

```
DBMS_OFFLINE_OG.BEGIN_INSTANTIATION (
   gname     IN   VARCHAR2,
   new_site  IN   VARCHAR2
   fname     IN   VARCHAR2);
```

### Parameters

***Table 15–2 BEGIN_INSTANTIATION Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Name of the replication group that you want to replicate to the new site. |
| new_site | The fully qualified database name of the new site to which you want to replicate the replication group. |
| fname | This parameter is for internal use only.<br>**Note:** Do not set this parameter unless directed to do so by Oracle Support Services. |

### Exceptions

***Table 15–3 BEGIN_INSTANTIATION Procedure Exceptions***

| Exception | Description |
|---|---|
| badargument | NULL or empty string for replication group or new master site name. |
| dbms_repcat.nonmasterdef | This procedure must be called from the master definition site. |
| sitealreadyexists | Specified site is already a master site for this replication group. |
| wrongstate | Status of master definition site must be quiesced. |
| dbms_repcat.missingrepgroup | gname does not exist as a master group. |

*Table 15–3 (Cont.) BEGIN_INSTANTIATION Procedure Exceptions*

| Exception | Description |
|---|---|
| `dbms_repcat.missing_flavor` | If you receive this exception, contact Oracle Support Services. |

## BEGIN_LOAD Procedure

This procedure disables triggers while data is imported to the new master site as part of offline instantiation. You must call this procedure from the new master site.

---

**Note:** This procedure is used to perform an offline instantiation of a master table in a multimaster replication environment.

This procedure should not be confused with the procedures in the DBMS_REPCAT_INSTANTIATE package (used for instantiating a deployment template). See the documentation for this package for more information about its use.

---

**See Also:** "Adding New Master Sites with Offline Instantiation Using Export/Import" on page 7-25 for information about adding a new master site to a master group by performing an offline instantiation of a master site

### Syntax

```
DBMS_OFFLINE_OG.BEGIN_LOAD (
   gname     IN   VARCHAR2,
   new_site  IN   VARCHAR2);
```

### Parameters

***Table 15–4 BEGIN_LOAD Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Name of the replication group whose members you are importing. |
| new_site | The fully qualified database name of the new site at which you will be importing the replication group members. |

### Exceptions

***Table 15–5 BEGIN_LOAD Procedure Exceptions***

| Exception | Description |
|---|---|
| badargument | NULL or empty string for replication group or new master site name. |
| wrongsite | This procedure must be called from the new master site. |
| unknownsite | Specified site is not recognized by replication group. |
| wrongstate | Status of the new master site must be quiesced. |
| dbms_repcat.missingrepgroup | gname does not exist as a master group. |

# END_INSTANTIATION Procedure

This procedure completes offline instantiation of a master group. You must call this procedure from the master definition site.

---

**Note:** This procedure is used to perform an offline instantiation of a master table in a multimaster replication environment.

This procedure should not be confused with the procedures in the DBMS_REPCAT_INSTANTIATE package (used for instantiating a deployment template). See the documentation for this package for more information about its use.

---

**See Also:** "Adding New Master Sites with Offline Instantiation Using Export/Import" on page 7-25 for information about adding a new master site to a master group by performing an offline instantiation of a master site

## Syntax

```
DBMS_OFFLINE_OG.END_INSTANTIATION (
  gname     IN  VARCHAR2,
  new_site  IN  VARCHAR2);
```

## Parameters

***Table 15–6 END_INSTANTIATION Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Name of the replication group that you are replicating to the new site. |
| new_site | The fully qualified database name of the new site to which you are replicating the replication group. |

## Exceptions

***Table 15–7 END_INSTANTIATION Procedure Exceptions***

| Exception | Description |
|---|---|
| badargument | NULL or empty string for replication group or new master site name. |
| dbms_repcat.nonmasterdef | This procedure must be called from the master definition site. |
| unknownsite | Specified site is not recognized by replication group. |
| wrongstate | Status of master definition site must be quiesced. |
| dbms_repcat.missingrepgroup | gname does not exist as a master group. |

## END_LOAD Procedure

This procedure reenables triggers after importing data to new master site as part of offline instantiation. You must call this procedure from the new master site.

---

**Note:** This procedure is used to perform an offline instantiation of a master table in a multimaster replication environment.

This procedure should not be confused with the procedures in the DBMS_REPCAT_INSTANTIATE package (used for instantiating a deployment template). See the documentation for this package for more information about its use.

---

**See Also:** "Adding New Master Sites with Offline Instantiation Using Export/Import" on page 7-25 for information about adding a new master site to a master group by performing an offline instantiation of a master site

### Syntax

```
DBMS_OFFLINE_OG.END_LOAD (
   gname     IN   VARCHAR2,
   new_site  IN   VARCHAR2
   fname     IN   VARCHAR2);
```

### Parameters

***Table 15–8 END_LOAD Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Name of the replication group whose members you have finished importing. |
| new_site | The fully qualified database name of the new site at which you have imported the replication group members. |
| fname | This parameter is for internal use only.<br>**Note:** Do not set this parameter unless directed to do so by Oracle Support Services. |

### Exceptions

***Table 15–9 END_LOAD Procedure Exceptions***

| Exception | Description |
|---|---|
| badargument | NULL or empty string for replication group or new master site name. |
| wrongsite | This procedure must be called from the new master site. |
| unknownsite | Specified site is not recognized by replication group. |
| wrongstate | Status of the new master site must be quiesced. |
| dbms_repcat.missingrepgroup | gname does not exist as a master group. |
| dbms_repcat.flavor_noobject | If you receive this exception, contact Oracle Support Services. |

*Table 15–9 (Cont.) END_LOAD Procedure Exceptions*

| Exception | Description |
|---|---|
| `dbms_repcat.flavor_contains` | If you receive this exception, contact Oracle Support Services. |

## RESUME_SUBSET_OF_MASTERS Procedure

When you add a new master site to a master group by performing an offline instantiation of a master site, it might take some time to complete the offline instantiation process. This procedure resumes replication activity at all existing sites, except the new site, during offline instantiation of a master group. You typically execute this procedure after executing the DBMS_OFFLINE_OG.BEGIN_INSTANTIATION procedure. You must call this procedure from the master definition site.

---

**Note:** This procedure is used to perform an offline instantiation of a master table in a multimaster replication environment.

This procedure should not be confused with the procedures in the DBMS_REPCAT_INSTANTIATE package (used for instantiating a deployment template). See the documentation for this package for more information about its use.

---

**See Also:** "Adding New Master Sites with Offline Instantiation Using Export/Import" on page 7-25 for information about adding a new master site to a master group by performing an offline instantiation of a master site

### Syntax

```
DBMS_OFFLINE_OG.RESUME_SUBSET_OF_MASTERS (
   gname     IN  VARCHAR2,
   new_site  IN  VARCHAR2
   override  IN  BOOLEAN := FALSE);
```

### Parameters

***Table 15–10 RESUME_SUBSET_OF_MASTERS Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Name of the replication group that you are replicating to the new site. |
| new_site | The fully qualified database name of the new site to which you are replicating the replication group. |
| override | If this is TRUE, then any pending administrative requests are ignored and normal replication activity is restored at each master as quickly as possible. The override parameter should be set to TRUE only in emergency situations.<br>If this is FALSE, then normal replication activity is restored at each master only when there is no pending administrative request for gname at that master. |

## Exceptions

*Table 15–11 RESUME_SUBSET_OF_MASTERS Procedure Exceptions*

| Exception | Description |
|---|---|
| badargument | NULL or empty string for replication group or new master site name. |
| dbms_repcat.nonmasterdef | This procedure must be called from the master definition site. |
| unknownsite | Specified site is not recognized by replication group. |
| wrongstate | Status of master definition site must be quiesced. |
| dbms_repcat.missingrepgroup | gname does not exist as a master group. |

# 16

# DBMS_RECTIFIER_DIFF

The DBMS_RECTIFIER_DIFF package contains APIs used to detect and resolve data inconsistencies between two replicated sites.

This chapter contains this topic:

- Summary of DBMS_RECTIFIER_DIFF Subprograms

---

**Note:** You can also determine differences between database objects and converge them using the DBMS_COMPARISON package.

---

**See Also:**

- *Oracle Database PL/SQL Packages and Types Reference* for information about the DBMS_COMPARISON package
- *Oracle Database 2 Day + Data Replication and Integration Guide* and *Oracle Streams Replication Administrator's Guide* for information about using the DBMS_COMPARISON package

# Summary of DBMS_RECTIFIER_DIFF Subprograms

*Table 16–1 DBMS_RECTIFIER_DIFF Package Subprograms*

| Subprogram | Description |
|---|---|
| "DIFFERENCES Procedure" on page 16-3 | Determines the differences between two tables. |
| "RECTIFY Procedure" on page 16-6 | Resolves the differences between two tables. |

## DIFFERENCES Procedure

This procedure determines the differences between two tables. It accepts the storage table of a nested table.

> **Note:** This procedure cannot be used on LOB columns, nor on columns based on user-defined types.

### Syntax

```
DBMS_RECTIFIER_DIFF.DIFFERENCES (
   sname1               IN  VARCHAR2,
   oname1               IN  VARCHAR2,
   reference_site       IN  VARCHAR2 := '',
   sname2               IN  VARCHAR2,
   oname2               IN  VARCHAR2,
   comparison_site      IN  VARCHAR2 := '',
   where_clause         IN  VARCHAR2 := '',
   { column_list        IN  VARCHAR2 := '',
   | array_columns      IN  DBMS_UTILITY.NAME_ARRAY, }
   missing_rows_sname   IN  VARCHAR2,
   missing_rows_oname1  IN  VARCHAR2,
   missing_rows_oname2  IN  VARCHAR2,
   missing_rows_site    IN  VARCHAR2 := '',
   max_missing          IN  INTEGER,
   commit_rows          IN  INTEGER := 500);
```

> **Note:** This procedure is overloaded. The `column_list` and `array_columns` parameters are mutually exclusive.

### Parameters

***Table 16–2 DIFFERENCES Procedure Parameters***

| Parameter | Description |
|---|---|
| `sname1` | Name of the schema at `reference_site`. |
| `oname1` | Name of the table at `reference_site`. |
| `reference_site` | Name of the reference database site. The default, `NULL`, indicates the current site. |
| `sname2` | Name of the schema at `comparison_site`. |
| `oname2` | Name of the table at `comparison_site`. |
| `comparison_site` | Name of the comparison database site. The default, `NULL`, indicates the current site. |
| `where_clause` | Only rows satisfying this clause are selected for comparison. The default, `NULL`, indicates all rows are compared. |
| `column_list` | A comma-delimited list of one or more column names being compared for the two tables. You must not have any spaces before or after a comma. The default, `NULL`, indicates that all columns will be compared. |

*Table 16–2 (Cont.) DIFFERENCES Procedure Parameters*

| Parameter | Description |
|---|---|
| array_columns | A PL/SQL associative array of column names being compared for the two tables. Indexing begins at 1, and the final element of the array must be NULL. If position 1 is NULL, then all columns are used. |
| missing_rows_sname | Name of the schema containing the tables with the missing rows. |
| missing_rows_oname1 | Name of an existing table at missing_rows_site that stores information about the rows in the table at reference_site that are missing from the table at comparison_site, and information about the rows at comparison_site site that are missing from the table at reference_site. |
| missing_rows_oname2 | Name of an existing table at missing_rows_site that stores information about the missing rows. This table has three columns: the R_ID column shows the rowid of the row in the missing_rows_oname1 table, the PRESENT column shows the name of the site where the row is present, and the ABSENT column shows name of the site from which the row is absent. |
| missing_rows_site | Name of the site where the missing_rows_oname1 and missing_rows_oname2 tables are located. The default, NULL, indicates that the tables are located at the current site. |
| max_missing | Integer that specifies the maximum number of rows that should be inserted into the missing_rows_oname table. If more than max_missing rows are missing, then that many rows are inserted into missing_rows_oname, and the routine then returns normally without determining whether more rows are missing. This parameter is useful if the fragments are so different that the missing rows table has too many entries and there is no point in continuing. Raises exception badnumber if max_missing is less than 1 or NULL. |
| commit_rows | Maximum number of rows to insert to or delete from the reference or comparison table before a COMMIT occurs. By default, a COMMIT occurs after 500 inserts or 500 deletes. An empty string (' ') or NULL indicates that a COMMIT should be issued only after all rows for a single table have been inserted or deleted. |

## Exceptions

*Table 16–3 DIFFERENCES Procedure Exceptions*

| Exception | Description |
|---|---|
| nosuchsite | Database site could not be found. |
| badnumber | The commit_rows parameter is less than 1. |
| missingprimarykey | Column list must include primary key (or SET_COLUMNS equivalent). |
| badname | NULL or empty string for table or schema name. |
| cannotbenull | Parameter cannot be NULL. |
| notshapeequivalent | Tables being compared are not shape equivalent. Shape refers to the number of columns, their column names, and the column data types. |
| unknowncolumn | Column does not exist. |
| unsupportedtype | Type not supported. |
| dbms_repcat.commfailure | Remote site is inaccessible. |
| dbms_repcat.missingobject | Table does not exist. |

## Restrictions

The error ORA-00001 (unique constraint violated) is issued when there are any unique or primary key constraints on the missing rows table.

# RECTIFY Procedure

This procedure resolves the differences between two tables. It accepts the storage table of a nested table.

**Note:** This procedure cannot be used on LOB columns, nor on columns based on user-defined types.

## Syntax

```
DBMS_RECTIFIER_DIFF.RECTIFY (
   sname1               IN  VARCHAR2,
   oname1               IN  VARCHAR2,
   reference_site       IN  VARCHAR2 := '',
   sname2               IN  VARCHAR2,
   oname2               IN  VARCHAR2,
   comparison_site      IN  VARCHAR2 := '',
   { column_list        IN  VARCHAR2 := '',
   | array_columns      IN  dbms_utility.name_array, }
   missing_rows_sname   IN  VARCHAR2,
   missing_rows_oname1  IN  VARCHAR2,
   missing_rows_oname2  IN  VARCHAR2,
   missing_rows_site    IN  VARCHAR2 := '',
   commit_rows          IN  INTEGER := 500);
```

**Note:** This procedure is overloaded. The `column_list` and `array_columns` parameters are mutually exclusive.

## Parameters

***Table 16–4 RECTIFY Procedure Parameters***

| Parameter | Description |
|---|---|
| `sname1` | Name of the schema at `reference_site`. |
| `oname1` | Name of the table at `reference_site`. |
| `reference_site` | Name of the reference database site. The default, `NULL`, indicates the current site. |
| `sname2` | Name of the schema at `comparison_site`. |
| `oname2` | Name of the table at `comparison_site`. |
| `comparison_site` | Name of the comparison database site. The default, `NULL`, indicates the current site. |
| `column_list` | A comma-delimited list of one or more column names being compared for the two tables. You must not have any spaces before or after a comma. The default, `NULL`, indicates that all columns will be compared. |
| `array_columns` | A PL/SQL associative array of column names being compared for the two tables. Indexing begins at 1, and the final element of the array must be `NULL`. If position 1 is `NULL`, then all columns are used. |
| `missing_rows_sname` | Name of the schema containing the tables with the missing rows. |

*Table 16–4 (Cont.) RECTIFY Procedure Parameters*

| Parameter | Description |
|---|---|
| missing_rows_oname1 | Name of the table at missing_rows_site that stores information about the rows in the table at reference_site that are missing from the table at comparison_site, and information about the rows at comparison_site that are missing from the table at reference_site. |
| missing_rows_oname2 | Name of the table at missing_rows_site that stores information about the missing rows. This table has three columns: the rowid of the row in the missing_rows_oname1 table, the name of the site at which the row is present, and the name of the site from which the row is absent. |
| missing_rows_site | Name of the site where the missing_rows_oname1 and missing_rows_oname2 tables are located. The default, NULL, indicates that the tables are located at the current site. |
| commit_rows | Maximum number of rows to insert to or delete from the reference or comparison table before a COMMIT occurs. By default, a COMMIT occurs after 500 inserts or 500 deletes. An empty string ('') or NULL indicates that a COMMIT should be issued only after all rows for a single table have been inserted or deleted. |

## Exceptions

*Table 16–5 RECTIFY Procedure Exceptions*

| Exception | Description |
|---|---|
| nosuchsite | Database site could not be found. |
| badnumber | The commit_rows parameter is less than 1. |
| badname | NULL or empty string for table or schema name. |
| dbms_repcat.commfailure | Remote site is inaccessible. |
| dbms_repcat.missingobject | Table does not exist. |

# 17

# DBMS_REFRESH

DBMS_REFRESH enables you to create groups of materialized views that can be refreshed together to a transactionally consistent point in time.

This chapter contains this topic:


- Summary of DBMS_REFRESH Subprograms

# Summary of DBMS_REFRESH Subprograms

*Table 17–1 DBMS_REFRESH Package Subprograms*

| Subprogram | Description |
|---|---|
| "ADD Procedure" on page 17-3 | Adds materialized views to a refresh group. |
| "CHANGE Procedure" on page 17-4 | Changes the refresh interval for a refresh group. |
| "DESTROY Procedure" on page 17-6 | Removes all of the materialized views from a refresh group and deletes the refresh group. |
| "MAKE Procedure" on page 17-7 | Specifies the members of a refresh group and the time interval used to determine when to refresh the members of this group. |
| "REFRESH Procedure" on page 17-9 | Manually refreshes a refresh group. |
| "SUBTRACT Procedure" on page 17-10 | Removes materialized views from a refresh group. |

## ADD Procedure

This procedure adds materialized views to a refresh group.

> **See Also:** Step 6, "Add objects to the refresh group.", on page 5-7 and *Oracle Database Advanced Replication* for more information

### Syntax

```
DBMS_REFRESH.ADD (
   name     IN VARCHAR2,
   { list   IN VARCHAR2,
   | tab    IN DBMS_UTILITY.UNCL_ARRAY, }
   lax      IN BOOLEAN := FALSE);
```

> **Note:** This procedure is overloaded. The `list` and `tab` parameters are mutually exclusive.

### Parameters

***Table 17–2 ADD Procedures Parameters***

| Parameter | Description |
|---|---|
| `name` | Name of the refresh group to which you want to add members, specified as `[`*`schema_name`*`.]`*`refresh_group_name`*. If the schema is not specified, then the current user is the default. |
| `list` | Comma-delimited list of materialized views that you want to add to the refresh group. Synonyms are not supported.<br>Each materialized view is specified as `[`*`schema_name`*`.]`*`materialized_view_name`*. If the schema is not specified, then the refresh group owner is the default. |
| `tab` | Instead of a comma-delimited list, you can supply a PL/SQL associative array of type `DBMS_UTILITY.UNCL_ARRAY`, where each element is the name of a materialized view. The first materialized view should be in position 1. The last position must be `NULL`.<br>Each materialized view is specified as `[`*`schema_name`*`.]`*`materialized_view_name`*. If the schema is not specified, then the refresh group owner is the default. |
| `lax` | A materialized view can belong to only one refresh group at a time. If you are moving a materialized view from one group to another, then you must set the `lax` flag to `TRUE` to succeed. Oracle then automatically removes the materialized view from the other refresh group and updates its refresh interval to be that of its new group. Otherwise, the call to `ADD` generates an error message. |

# CHANGE Procedure

This procedure changes the refresh interval for a refresh group.

> **See Also:** *Oracle Database Advanced Replication* for more information about refresh groups

## Syntax

```
DBMS_REFRESH.CHANGE (
   name                 IN VARCHAR2,
   next_date            IN DATE           := NULL,
   interval             IN VARCHAR2       := NULL,
   implicit_destroy     IN BOOLEAN        := NULL,
   rollback_seg         IN VARCHAR2       := NULL,
   push_deferred_rpc    IN BOOLEAN        := NULL,
   refresh_after_errors IN BOOLEAN        := NULL,
   purge_option         IN BINARY_INTEGER := NULL,
   parallelism          IN BINARY_INTEGER := NULL,
   heap_size            IN BINARY_INTEGER := NULL);
```

## Parameters

***Table 17–3 CHANGE Procedures Parameters***

| Parameter | Description |
|---|---|
| `name` | Name of the refresh group for which you want to alter the refresh interval. |
| `next_date` | Next date that you want a refresh to occur. By default, this date remains unchanged. |
| `interval` | Function used to calculate the next time to refresh the materialized views in the refresh group. This interval is evaluated immediately before the refresh. Thus, select an interval that is greater than the time it takes to perform a refresh. By default, the interval remains unchanged. |
| `implicit_destroy` | Allows you to reset the value of the `implicit_destroy` flag. If this flag is set, then Oracle automatically deletes the group if it no longer contains any members. By default, this flag remains unchanged. |
| `rollback_seg` | Allows you to change the rollback segment used. By default, the rollback segment remains unchanged. To reset this parameter to use the default rollback segment, specify `NULL`, including the quotes. Specifying `NULL` without quotes indicates that you do not want to change the rollback segment currently being used. |
| `push_deferred_rpc` | Used by updatable materialized views only. Set this parameter to `TRUE` to push changes from the materialized view to its associated master table or master materialized view before refreshing the materialized view. Otherwise, these changes might appear to be temporarily lost. By default, this flag remains unchanged. |
| `refresh_after_errors` | Used by updatable materialized views only. Set this parameter to `TRUE` if you want the refresh to proceed even if there are outstanding conflicts logged in the `DEFERROR` view for the materialized view's master table or master materialized view. By default, this flag remains unchanged. |

***Table 17–3 (Cont.) CHANGE Procedures Parameters***

| Parameter | Description |
|---|---|
| `purge_option` | If you are using the parallel propagation mechanism (that is, `parallelism` is set to 1 or greater), then:<br>■ 0 = do not purge<br>■ 1 = lazy (default)<br>■ 2 = aggressive<br>In most cases, *lazy* purge is the optimal setting. Set purge to *aggressive* to trim back the queue if multiple master replication groups are pushed to different target sites, and updates to one or more replication groups are infrequent and infrequently pushed. If all replication groups are infrequently updated and pushed, then set purge to *do not purge* and occasionally execute `PUSH` with purge set to *aggressive* to reduce the queue. |
| `parallelism` | 0 specifies serial propagation.<br>*n* > 1 specifies parallel propagation with *n* parallel processes.<br>1 specifies parallel propagation using only one parallel process. |
| `heap_size` | Maximum number of transactions to be examined simultaneously for parallel propagation scheduling. Oracle automatically calculates the default setting for optimal performance.<br>**Note:** Do not set this parameter unless directed to do so by Oracle Support Services. |

## DESTROY Procedure

This procedure removes all of the materialized views from a refresh group and delete the refresh group.

> **See Also:** *Oracle Database Advanced Replication* for more information refresh groups

### Syntax

```
DBMS_REFRESH.DESTROY (
   name    IN    VARCHAR2);
```

### Parameters

***Table 17–4 DESTROY Procedure Parameters***

| Parameter | Description |
|---|---|
| name | Name of the refresh group that you want to destroy. |

## MAKE Procedure

This procedure specifies the members of a refresh group and the time interval used to determine when to refresh the members of this group.

> **See Also:** Step 4, "Create the refresh group.", on page 5-5 and *Oracle Database Advanced Replication* for more information

### Syntax

```
DBMS_REFRESH.MAKE (
   name                 IN    VARCHAR2
   { list               IN    VARCHAR2,
   | tab                IN    DBMS_UTILITY.UNCL_ARRAY,}
   next_date            IN    DATE,
   interval             IN    VARCHAR2,
   implicit_destroy     IN    BOOLEAN         := FALSE,
   lax                  IN    BOOLEAN         := FALSE,
   job                  IN    BINARY_INTEGER  := 0,
   rollback_seg         IN    VARCHAR2        := NULL,
   push_deferred_rpc    IN    BOOLEAN         := TRUE,
   refresh_after_errors IN    BOOLEAN         := FALSE
   purge_option         IN    BINARY_INTEGER  := NULL,
   parallelism          IN    BINARY_INTEGER  := NULL,
   heap_size            IN    BINARY_INTEGER  := NULL);
```

> **Note:** This procedure is overloaded. The `list` and `tab` parameters are mutually exclusive.

### Parameters

***Table 17–5 MAKE Procedure Parameters***

| Parameter | Description |
|---|---|
| `name` | Unique name used to identify the refresh group, specified as `[`*`schema_name`*`.]`*`refresh_group_name`*. If the schema is not specified, then the current user is the default. Refresh groups must follow the same naming conventions as tables. |
| `list` | Comma-delimited list of materialized views that you want to refresh. Synonyms are not supported. These materialized views can be located in different schemas and have different master tables or master materialized views. However, all of the listed materialized views must be in your current database.<br><br>Each materialized view is specified as `[`*`schema_name`*`.]`*`materialized_view_name`*. If the schema is not specified, then the refresh group owner is the default. |
| `tab` | Instead of a comma-delimited list, you can supply a PL/SQL associative array of names of materialized views that you want to refresh using the data type `DBMS_UTILITY.UNCL_ARRAY`. If the table contains the names of *n* materialized views, then the first materialized view should be in position 1 and the *n* + 1 position should be set to `NULL`.<br><br>Each materialized view is specified as `[`*`schema_name`*`.]`*`materialized_view_name`*. If the schema is not specified, then the refresh group owner is the default. |

***Table 17–5 (Cont.) MAKE Procedure Parameters***

| Parameter | Description |
|---|---|
| next_date | Next date that you want a refresh to occur. |
| interval | Function used to calculate the next time to refresh the materialized views in the group. This field is used with the next_date value.<br><br>For example, if you specify NEXT_DAY(SYSDATE+1, "MONDAY") as your interval, and if your next_date evaluates to Monday, then Oracle refreshes the materialized views every Monday. This interval is evaluated immediately before the refresh. Thus, select an interval that is greater than the time it takes to perform a refresh. |
| implicit_destroy | Set this to TRUE to delete the refresh group automatically when it no longer contains any members. Oracle checks this flag only when you call the SUBTRACT procedure. That is, setting this flag still enables you to create an empty refresh group. |
| lax | A materialized view can belong to only one refresh group at a time. If you are moving a materialized view from an existing group to a new refresh group, then you must set this to TRUE to succeed. Oracle then automatically removes the materialized view from the other refresh group and updates its refresh interval to be that of its new group. Otherwise, the call to MAKE generates an error message. |
| job | Needed by the Import utility. Use the default value, 0. |
| rollback_seg | Name of the rollback segment to use while refreshing materialized views. The default, NULL, uses the default rollback segment. |
| push_deferred_rpc | Used by updatable materialized views only. Use the default value, TRUE, to push changes from the materialized view to its associated master table or master materialized view before refreshing the materialized view. Otherwise, these changes might appear to be temporarily lost. |
| refresh_after_errors | Used by updatable materialized views only. Set this to 0 if you want the refresh to proceed even if there are outstanding conflicts logged in the DEFERROR view for the materialized view's master table or master materialized view. |
| purge_option | If you are using the parallel propagation mechanism (in other words, parallelism is set to 1 or greater), then 0 = do not purge; 1 = lazy (default); 2 = aggressive. In most cases, *lazy* purge is the optimal setting.<br><br>Set purge to *aggressive* to trim back the queue if multiple master replication groups are pushed to different target sites, and updates to one or more replication groups are infrequent and infrequently pushed. If all replication groups are infrequently updated and pushed, then set purge to *do not purge* and occasionally execute PUSH with purge set to *aggressive* to reduce the queue. |
| parallelism | 0 specifies serial propagation.<br><br>*n* > 1 specifies parallel propagation with *n* parallel processes.<br><br>1 specifies parallel propagation using only one parallel process. |
| heap_size | Maximum number of transactions to be examined simultaneously for parallel propagation scheduling. Oracle automatically calculates the default setting for optimal performance.<br><br>**Note:** Do not set this parameter unless directed to do so by Oracle Support Services. |

## REFRESH Procedure

This procedure manually refreshes a refresh group.

> **See Also:** *Oracle Database Advanced Replication* for more information about refresh groups

### Syntax

```
DBMS_REFRESH.REFRESH (
   name    IN    VARCHAR2);
```

### Parameter

***Table 17–6 REFRESH Procedure Parameter***

| Parameter | Description |
|---|---|
| name | Name of the refresh group that you want to refresh manually. |

# SUBTRACT Procedure

This procedure removes materialized views from a refresh group.

> **See Also:** *Oracle Database Advanced Replication* for more information about refresh groups

## Syntax

```
DBMS_REFRESH.SUBTRACT (
   name      IN    VARCHAR2,
   { list    IN    VARCHAR2,
   | tab     IN    DBMS_UTILITY.UNCL_ARRAY, }
   lax       IN    BOOLEAN := FALSE);
```

> **Note:** This procedure is overloaded. The `list` and `tab` parameters are mutually exclusive.

## Parameters

***Table 17–7 SUBTRACT Procedure Parameters***

| Parameter | Description |
|---|---|
| `name` | Name of the refresh group from which you want to remove members. |
| `list` | Comma-delimited list of materialized views that you want to remove from the refresh group. (Synonyms are not supported.) These materialized views can be located in different schemas and have different master tables or master materialized views. However, all of the listed materialized views must be in your current database. |
| `tab` | Instead of a comma-delimited list, you can supply a PL/SQL associative array of names of materialized views that you want to refresh using the data type `DBMS_UTILITY.UNCL_ARRAY`. If the table contains the names of *n* materialized views, then the first materialized view should be in position 1 and the *n* + 1 position should be set to `NULL`. |
| `lax` | Set this to `FALSE` if you want Oracle to generate an error message if the materialized view you are attempting to remove is not a member of the refresh group. |

# 18

# DBMS_REPCAT

DBMS_REPCAT provides routines to administer and update the replication catalog and environment.

This chapter contains this topic:

- Summary of DBMS_REPCAT Subprograms

# Summary of DBMS_REPCAT Subprograms

*Table 18–1 DBMS_REPCAT Package Subprograms*

| Subprogram | Description |
|---|---|
| "ADD_GROUPED_COLUMN Procedure" on page 18-6 | Adds members to an existing column group. |
| "ADD_MASTER_DATABASE Procedure" on page 18-7 | Adds another master site to your replication environment. |
| "ADD_NEW_MASTERS Procedure" on page 18-8 | Adds the master sites in the DBA_REPSITES_NEW data dictionary view to the replication catalog at all available master sites. |
| "ADD_PRIORITY_datatype Procedure" on page 18-13 | Adds a member to a priority group. |
| "ADD_SITE_PRIORITY_SITE Procedure" on page 18-15 | Adds a new site to a site priority group. |
| "ADD_conflicttype_RESOLUTION Procedure" on page 18-16 | Designates a method for resolving an update, delete, or uniqueness conflict. |
| "ALTER_CATCHUP_PARAMETERS Procedure" on page 18-20 | Alters the values for parameters stored in the DBA_REPEXTENSIONS data dictionary view. |
| "ALTER_MASTER_PROPAGATION Procedure" on page 18-22 | Alters the propagation method for a specified replication group at a specified master site. |
| "ALTER_MASTER_REPOBJECT Procedure" on page 18-23 | Alters an object in your replication environment. |
| "ALTER_MVIEW_PROPAGATION Procedure" on page 18-25 | Alters the propagation method for a specified replication group at the current materialized view site. |
| "ALTER_PRIORITY Procedure" on page 18-26 | Alters the priority level associated with a specified priority group member. |
| "ALTER_PRIORITY_datatype Procedure" on page 18-27 | Alters the value of a member in a priority group. |
| "ALTER_SITE_PRIORITY Procedure" on page 18-28 | Alters the priority level associated with a specified site. |
| "ALTER_SITE_PRIORITY_SITE Procedure" on page 18-29 | Alters the site associated with a specified priority level. |
| "CANCEL_STATISTICS Procedure" on page 18-30 | Stops collecting statistics about the successful resolution of update, uniqueness, and delete conflicts for a table. |
| "COMMENT_ON_COLUMN_GROUP Procedure" on page 18-31 | Updates the comment field in the ALL_REPCOLUMN_GROUP view for a column group. |
| "COMMENT_ON_MVIEW_REPSITES Procedure" on page 18-32 | Updates the SCHEMA_COMMENT field in the ALL_REPGROUP view for a materialized view site. |
| "COMMENT_ON_PRIORITY_GROUP Procedures" on page 18-33 | Updates the comment field in the ALL_REPPRIORITY_GROUP view for a priority group. |
| "COMMENT_ON_REPGROUP Procedure" on page 18-34 | Updates the comment field in the ALL_REPGROUP view for a master group. |
| "COMMENT_ON_REPOBJECT Procedure" on page 18-35 | Updates the comment field in the ALL_REPOBJECT view for a replicated object. |

*Table 18–1 (Cont.) DBMS_REPCAT Package Subprograms*

| Subprogram | Description |
|---|---|
| "COMMENT_ON_REPSITES Procedure" on page 18-36 | Updates the comment field in the ALL_REPSITE view for a replicated site. |
| "COMMENT_ON_SITE_PRIORITY Procedure" on page 18-37 | Updates the comment field in the ALL_REPPRIORITY_GROUP view for a site priority group. |
| "COMMENT_ON_conflicttype_RESOLUTION Procedure" on page 18-38 | Updates the comment field in the ALL_REPRESOLUTION view for a conflict resolution routine. |
| "COMPARE_OLD_VALUES Procedure" on page 18-40 | Specifies whether to compare old column values at each master site for each nonkey column of a replicated table for updates and deletes. |
| "CREATE_MASTER_REPGROUP Procedure" on page 18-42 | Creates a new, empty, quiesced master group. |
| "CREATE_MASTER_REPOBJECT Procedure" on page 18-43 | Specifies that an object is a replicated object. |
| "CREATE_MVIEW_REPGROUP Procedure" on page 18-46 | Creates a new, empty materialized view group in your local database. |
| "CREATE_MVIEW_REPOBJECT Procedure" on page 18-48 | Adds a replicated object to a materialized view group. |
| "DEFINE_COLUMN_GROUP Procedure" on page 18-51 | Creates an empty column group. |
| "DEFINE_PRIORITY_GROUP Procedure" on page 18-52 | Creates a new priority group for a master group. |
| "DEFINE_SITE_PRIORITY Procedure" on page 18-53 | Creates a new site priority group for a master group. |
| "DO_DEFERRED_REPCAT_ADMIN Procedure" on page 18-54 | Executes the local outstanding deferred administrative procedures for the specified master group at the current master site, or for all master sites. |
| "DROP_COLUMN_GROUP Procedure" on page 18-55 | Drops a column group. |
| "DROP_GROUPED_COLUMN Procedure" on page 18-56 | Removes members from a column group. |
| "DROP_MASTER_REPGROUP Procedure" on page 18-57 | Drops a master group from your current site. |
| "DROP_MASTER_REPOBJECT Procedure" on page 18-58 | Drops a replicated object from a master group. |
| "DROP_MVIEW_REPGROUP Procedure" on page 18-59 | Drops a replicated object from a master group. |
| "DROP_MVIEW_REPGROUP Procedure" on page 18-59 | Drops a materialized view site from your replication environment. |
| "DROP_MVIEW_REPOBJECT Procedure" on page 18-60 | Drops a replicated object from a materialized view site. |
| "DROP_PRIORITY Procedure" on page 18-61 | Drops a member of a priority group by priority level. |
| "DROP_PRIORITY_GROUP Procedure" on page 18-62 | Drops a priority group for a specified master group. |

*Table 18–1 (Cont.) DBMS_REPCAT Package Subprograms*

| Subprogram | Description |
|---|---|
| "DROP_PRIORITY_datatype Procedure" on page 18-63 | Drops a member of a priority group by value. |
| "DROP_SITE_PRIORITY Procedure" on page 18-64 | Drops a site priority group for a specified master group. |
| "DROP_SITE_PRIORITY_SITE Procedure" on page 18-65 | Drops a specified site, by name, from a site priority group. |
| "DROP_conflicttype_RESOLUTION Procedure" on page 18-66 | Drops an update, delete, or uniqueness conflict resolution method. |
| "EXECUTE_DDL Procedure" on page 18-68 | Supplies DDL that you want to have executed at each master site. |
| "GENERATE_MVIEW_SUPPORT Procedure" on page 18-69 | Activates triggers and generate packages needed to support the replication of updatable materialized views or procedural replication. |
| "GENERATE_REPLICATION_ SUPPORT Procedure" on page 18-71 | Generates the triggers, packages, and procedures needed to support replication for a specified object. |
| "MAKE_COLUMN_GROUP Procedure" on page 18-73 | Creates a new column group with one or more members. |
| "PREPARE_INSTANTIATED_MASTER Procedure" on page 18-74 | Changes the global name of the database you are adding to a master group. |
| "PURGE_MASTER_LOG Procedure" on page 18-76 | Removes local messages in the `DBA_REPCATLOG` associated with a specified identification number, source, or master group. |
| "PURGE_STATISTICS Procedure" on page 77 | Removes information from the `ALL_ REPRESOLUTION_STATISTICS` view. |
| "REFRESH_MVIEW_REPGROUP Procedure" on page 18-78 | Refreshes a materialized view group with the most recent data from its associated master site or master materialized view site. |
| REGISTER_MVIEW_REPGROUP Procedure on page 18-80 | Facilitates the administration of materialized views at their respective master sites or master materialized view sites by inserting, modifying, or deleting from `DBA_REGISTERED_MVIEW_GROUPS`. |
| "REGISTER_STATISTICS Procedure" on page 18-81 | Collects information about the successful resolution of update, delete, and uniqueness conflicts for a table. |
| "RELOCATE_MASTERDEF Procedure" on page 82 | Changes your master definition site to another master site in your replication environment. |
| "REMOVE_MASTER_DATABASES Procedure" on page 18-84 | Removes one or more master databases from a replication environment. |
| "RENAME_SHADOW_COLUMN_ GROUP Procedure" on page 18-85 | Renames the shadow column group of a replicated table to make it a named column group. |
| "REPCAT_IMPORT_CHECK Procedure" on page 18-86 | Ensures that the objects in the master group have the appropriate object identifiers and status values after you perform an export/import of a replicated object or an object used by the advanced replication facility. |
| "RESUME_MASTER_ACTIVITY Procedure" on page 18-87 | Resumes normal replication activity after quiescing a replication environment. |

*Table 18–1 (Cont.) DBMS_REPCAT Package Subprograms*

| Subprogram | Description |
|---|---|
| "RESUME_PROPAGATION_TO_MDEF Procedure" on page 18-88 | Indicates that export is effectively finished and propagation for both extended and unaffected replication groups existing at master sites can be enabled. |
| "SEND_OLD_VALUES Procedure" on page 18-89 | Specifies whether to send old column values for each nonkey column of a replicated table for updates and deletes. |
| "SET_COLUMNS Procedure" on page 18-91 | Specifies use of an alternate column or group of columns, instead of the primary key, to determine which columns of a table to compare when using row-level replication. |
| "SPECIFY_NEW_MASTERS Procedure" on page 18-93 | Specifies the master sites you intend to add to an existing replication group without quiescing the group. |
| "STREAMS_MIGRATION Procedure" on page 18-95 | Generates a migration script that migrates an Advanced Replication environment to a Streams environment. |
| "SUSPEND_MASTER_ACTIVITY Procedure" on page 18-96 | Suspends replication activity for a master group. |
| "SWITCH_MVIEW_MASTER Procedure" on page 18-97 | Changes the master site of a materialized view group to another master site. |
| "UNDO_ADD_NEW_MASTERS_ REQUEST Procedure" on page 18-98 | Undoes all of the changes made by the `SPECIFY_NEW_MASTERS` and `ADD_NEW_MASTERS` procedures for a specified `extension_id`. |
| "UNREGISTER_MVIEW_REPGROUP Procedure" on page 18-100 | Facilitates the administration of materialized views at their respective master sites and master materialized view sites by inserting, modifying, or deleting from `DBA_REGISTERED_MVIEW_GROUPS`. |
| "VALIDATE Function" on page 18-101 | Validates the correctness of key conditions of a multimaster replication environment. |
| "WAIT_MASTER_LOG Procedure" on page 18-103 | Determines whether changes that were asynchronously propagated to a master site have been applied. |

# ADD_GROUPED_COLUMN Procedure

This procedure adds members to an existing column group. You must call this procedure from the master definition site.

## Syntax

```
DBMS_REPCAT.ADD_GROUPED_COLUMN (
   sname                 IN   VARCHAR2,
   oname                 IN   VARCHAR2,
   column_group          IN   VARCHAR2,
   list_of_column_names  IN   VARCHAR2 | DBMS_REPCAT.VARCHAR2S);
```

## Parameters

***Table 18–2 ADD_GROUPED_COLUMN Procedure Parameters***

| Parameter | Description |
|---|---|
| sname | Schema in which the replicated table is located. |
| oname | Name of the replicated table with which the column group is associated. The table can be the storage table of a nested table. |
| column_group | Name of the column group to which you are adding members. |
| list_of_column_names | Names of the columns that you are adding to the designated column group. This can either be a comma-delimited list or a PL/SQL associative array of column names. The PL/SQL associative array must be of type DBMS_REPCAT.VARCHAR2S. Use the single value '*' to create a column group that contains all of the columns in your table.<br>You can specify column objects, but you cannot specify attributes of column objects.<br>If the table is an object, then you can specify SYS_NC_OID$ to add the object identifier column to the column group. This column tracks the object identifier of each row object.<br>If the table is a storage table of a nested table, then you can specify NESTED_TABLE_ID to add the column that tracks the identifier for each row of the nested table. |

## Exceptions

***Table 18–3 ADD_GROUPED_COLUMN Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missingobject | Specified table does not exist. |
| missinggroup | Specified column group does not exist. |
| missingcolumn | Specified column does not exist in the specified table. |
| duplicatecolumn | Specified column is already a member of another column group. |
| missingschema | Specified schema does not exist. |
| notquiesced | Replication group to which the specified table belongs is not quiesced. |

## ADD_MASTER_DATABASE Procedure

This procedure adds another master site to your replication environment. This procedure regenerates all the triggers and their associated packages at existing master sites. You must call this procedure from the master definition site.

### Syntax

```
DBMS_REPCAT.ADD_MASTER_DATABASE (
   gname                 IN   VARCHAR2,
   master                IN   VARCHAR2,
   use_existing_objects  IN   BOOLEAN  := TRUE,
   copy_rows             IN   BOOLEAN  := TRUE,
   comment               IN   VARCHAR2 := '',
   propagation_mode      IN   VARCHAR2 := 'ASYNCHRONOUS',
   fname                 IN   VARCHAR2 := NULL);
```

### Parameters

***Table 18–4 ADD_MASTER_DATABASE Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Name of the replication group being replicated. This replication group must exist at the master definition site. |
| master | Fully qualified database name of the new master database. |
| use_existing_objects | Indicate TRUE to reuse any objects of the same type and shape that exist in the schema at the new master site. |
| copy_rows | Indicate TRUE if you want the initial contents of a table at the new master site to match the contents of the table at the master definition site. |
| comment | This comment is added to the MASTER_COMMENT field of the DBA_REPSITES view. |
| propagation_mode | Method of forwarding changes to and receiving changes from new master database. Accepted values are synchronous and asynchronous. |
| fname | This parameter is for internal use only.<br>**Note:** Do not set this parameter unless directed to do so by Oracle Support Services. |

### Exceptions

***Table 18–5 ADD_MASTER_DATABASE Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| notquiesced | Replication has not been suspended for the master group. |
| missingrepgroup | Replication group does not exist at the specified database site. |
| commfailure | New master is not accessible. |
| typefailure | An incorrect propagation mode was specified. |
| duplrepgrp | Master site exists. |

# ADD_NEW_MASTERS Procedure

This procedure adds the master sites in the DBA_REPSITES_NEW data dictionary view to the master groups specified when the SPECIFY_NEW_MASTERS procedure was run. Information about these new master sites are added to the replication catalog at all available master sites.

All master sites instantiated with object-level export/import must be accessible at this time. Their new replication groups are added in the quiesced state. Master sites instantiated through full database export/import or through changed-based recovery do not need to be accessible.

Run this procedure after you run the SPECIFY_NEW_MASTERS procedure.

**Caution:** After running this procedure, do not disable or enable propagation of the deferred transactions queue until after the new master sites are added. The DBA_REPEXTENSIONS data dictionary view must be clear before you disable or enable propagation. You can use the Advanced Replication interface in Oracle Enterprise Manager or the SET_DISABLED procedure in the DBMS_DEFER_SYS package to disable or enable propagation.

**See Also:**

- "SPECIFY_NEW_MASTERS Procedure" on page 18-93
- "Adding New Master Sites" on page 7-2 for more information about adding master sites to a master group

## Syntax

```
DBMS_REPCAT.ADD_NEW_MASTERS (
   export_required               IN    BOOLEAN,
   { available_master_list       IN    VARCHAR2,
   | available_master_table      IN    DBMS_UTILITY.DBLINK_ARRAY,}
   masterdef_flashback_scn       OUT   NUMBER,
   extension_id                  OUT   RAW,
   break_trans_to_masterdef      IN    BOOLEAN   := FALSE,
   break_trans_to_new_masters    IN    BOOLEAN   := FALSE,
   percentage_for_catchup_mdef   IN    BINARY_INTEGER  := 100,
   cycle_seconds_mdef            IN    BINARY_INTEGER  := 60,
   percentage_for_catchup_new    IN    BINARY_INTEGER  := 100,
   cycle_seconds_new             IN    BINARY_INTEGER  := 60);
```

**Note:** This procedure is overloaded. The available_master_list and available_master_table parameters are mutually exclusive.

## Parameters

*Table 18–6 ADD_NEW_MASTERS Procedure Parameters*

| Parameter | Description |
|---|---|
| export_required | Set to TRUE if either object-level or full database export is required for at least one of the new master sites. Set to FALSE if you are using change-based recovery for all of the new master sites. |
| available_master_list | A comma-delimited list of the new master sites to be instantiated using object-level export/import. The sites listed must match the sites specified in the SPECIFY_NEW_MASTERS procedure. List only the new master sites, not the existing master sites. Do not put any spaces between site names.<br>Specify NULL if all masters will be instantiated using full database export/import or change-based recovery. |
| available_master_table | A table that lists the new master sites to be instantiated using object-level export/import. The sites in the table must match the sites specified in the SPECIFY_NEW_MASTERS procedure. Do not specify masters that will be instantiated using full database export/import or change-based recovery.<br>In the table that lists the master sites to be instantiated using object-level export/import, list only the new master sites for the master groups being extended. Do not list the existing master sites in the master groups being extended. The first master site should be at position 1, the second at position 2, and so on. |
| masterdef_flashback_scn | This OUT parameter returns a system change number (SCN) that must be used during export or change-based recovery. Use the value returned by this parameter for the FLASHBACK_SCN export parameter when you perform the export. You can find the flashback_scn value by querying the DBA_REPEXTENSIONS data dictionary view. |
| extension_id | This OUT parameter returns an identifier for the current pending request to add master databases without quiesce. You can find the extension_id by querying the DBA_REPSITES_NEW and DBA_REPEXTENSIONS data dictionary views. |

*Table 18–6 (Cont.) ADD_NEW_MASTERS Procedure Parameters*

| Parameter | Description |
|---|---|
| break_trans_to_masterdef | This parameter is meaningful only if export_required is set to TRUE.<br><br>If break_trans_to_masterdef is set to TRUE, then existing masters can continue to propagate their deferred transactions to the master definition site for replication groups that are not adding master sites. Deferred transactions for replication groups that are adding master sites cannot be propagated until the export completes.<br><br>Each deferred transaction is composed of one or more remote procedure calls (RPCs). If set to FALSE and a transaction occurs that references objects in both unaffected master groups and master groups that are being extended, then the transaction might be split into two parts and sent to a destination in two separate transactions at different times. Such transactions are called split-transactions. If split-transactions are possible, then you must disable integrity constraints that might be violated by this behavior until the new master sites are added.<br><br>If break_trans_to_masterdef is set to FALSE, then existing masters cannot propagate their deferred transactions to the master definition site. |
| break_trans_to_new_masters | If break_trans_to_new_masters is set to TRUE, then existing master sites can continue to propagate deferred transactions to the new master sites for replication groups that are not adding master sites.<br><br>Each deferred transaction is composed of one or more remote procedure calls (RPCs). If set to TRUE and a transaction occurs that references objects in both unaffected master groups and master groups that are being extended, then the transaction might be split into two parts and sent to a destination in two separate transactions at different times. Such transactions are called split-transactions. If split-transactions are possible, then you must disable integrity constraints that might be violated by this behavior until the new master sites are added.<br><br>If break_trans_to_new_masters is set to FALSE, then propagation of deferred transaction queues to the new masters is disabled. |
| percentage_for_catchup_mdef | This parameter is meaningful only if export_required and break_trans_to_masterdef are both set to TRUE.<br><br>The percentage of propagation resources to use for catching up propagation to the master definition site. Must be a multiple of 10 and must be between 0 and 100. |
| cycle_seconds_mdef | This parameter is meaningful when percentage_for_ catchup_mdef is both meaningful and set to a value between 10 and 90, inclusive. In this case, propagation to the master definition site alternates between replication groups that are not being extended and replication groups that are being extended, with one push to each during each cycle. This parameter indicates the length of the cycle in seconds. |

*Table 18–6 (Cont.) ADD_NEW_MASTERS Procedure Parameters*

| Parameter | Description |
|---|---|
| percentage_for_catchup_new | This parameter is meaningful only if break_trans_to_new_masters is set to TRUE.<br><br>The percentage of propagation resources to use for catching up propagation to new master sites. Must be a multiple of 10 and must be between 0 and 100. |
| cycle_seconds_new | This parameter is meaningful when percentage_for_catchup_new is both meaningful and set to a value between 10 and 90, inclusive. In this case, propagation to a new master alternates between replication groups that are not being extended and replication groups that are being extended, with one push to each during each cycle. This parameter indicates the length of the cycle in seconds. |

## Exceptions

*Table 18–7 ADD_NEW_MASTERS Procedure Exceptions*

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| typefailure | The parameter value specified for one of the parameters is not appropriate. |
| novalidextreq | No valid extension request. The extension_id is not valid. |
| nonewsites | No new master sites to be added for the specified extension request. |
| notanewsite | Not a new site for extension request. A site was specified that was not specified when you ran the SPECIFY_NEW_MASTERS procedure. |
| dbnotcompatible | Feature is incompatible with database version. All databases must be at 9.2.0 or higher compatibility level. |

## Usage Notes

For a new master site to be instantiated using change-based recovery or full database export/import, the following conditions apply:

- The new master sites cannot have any existing replication groups.
- The master definition site cannot have any materialized view groups.
- The master definition site must be the same for all of the master groups. If one or more of these master groups have a different master definition site, then do not use change-based recovery or full database export/import. Use object-level export/import instead.
- The new master site must include all of the replication groups in the master definition site when the extension process is complete. That is, you cannot add a subset of the master groups at the master definition site to the new master site; all of the groups must be added.

For object-level export/import, before importing ensure that all the requests in the DBA_REPCATLOG data dictionary view for the extended groups have been processed without any error.

**Note:** To use change-based recovery, the existing master site and the new master site must be running under the same operating system, although the release of the operating system can differ.

## ADD_PRIORITY_*datatype* Procedure

This procedure adds a member to a priority group. You must call this procedure from the master definition site. The procedure that you must call is determined by the data type of your `priority` column. You must call this procedure once for each of the possible values of the `priority` column.

> **See Also:** Chapter 6, "Configuring Conflict Resolution" and *Oracle Database Advanced Replication* for more information about conflict resolution methods

### Syntax

```
DBMS_REPCAT.ADD_PRIORITY_datatype (
   gname              IN   VARCHAR2,
   pgroup             IN   VARCHAR2,
   value              IN   datatype,
   priority           IN   NUMBER);
```

where *datatype*:

```
{ NUMBER
| VARCHAR2
| CHAR
| DATE
| RAW
| NCHAR
| NVARCHAR2 }
```

### Parameters

***Table 18–8 ADD_PRIORITY_datatype Procedure Parameters***

| Parameter | Description |
|---|---|
| `gname` | Master group for which you are creating a priority group. |
| `pgroup` | Name of the priority group. |
| `value` | Value of the priority group member. This is one of the possible values of the associated `priority` column of a table using this priority group. |
| `priority` | Priority of this value. The higher the number, the higher the priority. |

### Exceptions

***Table 18–9 ADD_PRIORITY_datatype Procedure Exceptions***

| Exception | Description |
|---|---|
| `nonmasterdef` | Invocation site is not the master definition site. |
| `duplicatevalue` | Specified value exists in the priority group. |
| `duplicatepriority` | Specified priority exists in the priority group. |
| `missingrepgroup` | Specified master group does not exist. |
| `missingprioritygroup` | Specified priority group does not exist. |
| `typefailure` | Specified value has the incorrect data type for the priority group. |

*Table 18–9 (Cont.) ADD_PRIORITY_datatype Procedure Exceptions*

| Exception | Description |
|---|---|
| notquiesced | Specified master group is not quiesced. |

## ADD_SITE_PRIORITY_SITE Procedure

This procedure adds a new site to a site priority group. You must call this procedure from the master definition site.

> **See Also:** Chapter 6, "Configuring Conflict Resolution" and *Oracle Database Advanced Replication* for more information about conflict resolution methods

### Syntax

```
DBMS_REPCAT.ADD_SITE_PRIORITY_SITE (
   gname          IN   VARCHAR2,
   name           IN   VARCHAR2,
   site           IN   VARCHAR2,
   priority       IN   NUMBER);
```

### Parameters

***Table 18–10 ADD_SITE_PRIORITY_SITE Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Master group for which you are adding a site to a group. |
| name | Name of the site priority group to which you are adding a member. |
| site | Global database name of the site that you are adding. |
| priority | Priority level of the site that you are adding. A higher number indicates a higher priority level. |

### Exceptions

***Table 18–11 ADD_SITE_PRIORITY_SITE Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missingrepgroup | Specified master group does not exist. |
| missingpriority | Specified site priority group does not exist. |
| duplicatepriority | Specified priority level exists for another site in the group. |
| duplicatevalue | Specified site exists in the site priority group. |
| notquiesced | Master group is not quiesced. |

# ADD_*conflicttype*_RESOLUTION Procedure

These procedures designate a method for resolving an update, delete, or uniqueness conflict. You must call these procedures from the master definition site. The procedure that you must call is determined by the type of conflict that the routine resolves.

***Table 18–12 ADD_conflicttype_RESOLUTION Procedures***

| Conflict Type | Procedure Name |
|---|---|
| update | ADD_UPDATE_RESOLUTION |
| uniqueness | ADD_UNIQUE_RESOLUTION |
| delete | ADD_DELETE_RESOLUTION |

**See Also:** Chapter 6, "Configuring Conflict Resolution" and *Oracle Database Advanced Replication* for more information about designating methods to resolve update conflicts, selecting uniqueness conflict resolution methods, and assigning delete conflict resolution methods

## Syntax

```
DBMS_REPCAT.ADD_UPDATE_RESOLUTION (
   sname                  IN   VARCHAR2,
   oname                  IN   VARCHAR2,
   column_group           IN   VARCHAR2,
   sequence_no            IN   NUMBER,
   method                 IN   VARCHAR2,
   parameter_column_name  IN   VARCHAR2
                               | DBMS_REPCAT.VARCHAR2S
                               | DBMS_UTILITY.LNAME_ARRAY,
   priority_group         IN   VARCHAR2     := NULL,
   function_name          IN   VARCHAR2     := NULL,
   comment                IN   VARCHAR2     := NULL);

DBMS_REPCAT.ADD_DELETE_RESOLUTION (
   sname                  IN   VARCHAR2,
   oname                  IN   VARCHAR2,
   sequence_no            IN   NUMBER,
   parameter_column_name  IN   VARCHAR2 | DBMS_REPCAT.VARCHAR2S,
   function_name          IN   VARCHAR2,
   comment                IN   VARCHAR2     := NULL
   method                 IN   VARCHAR2     := 'USER FUNCTION');

DBMS_REPCAT.ADD_UNIQUE_RESOLUTION(
   sname                 IN   VARCHAR2,
   oname                 IN   VARCHAR2,
   constraint_name       IN   VARCHAR2,
   sequence_no           IN   NUMBER,
   method                IN   VARCHAR2,
   parameter_column_name IN   VARCHAR2
                              | DBMS_REPCAT.VARCHAR2S
                              | DBMS_UTILITY.LNAME_ARRAY,
   function_name         IN   VARCHAR2     := NULL,
   comment               IN   VARCHAR2     := NULL);
```

## Parameters

*Table 18–13 ADD_conflicttype_RESOLUTION Procedure Parameters*

| Parameter | Description |
|---|---|
| sname | Name of the schema containing the table to be replicated. |
| oname | Name of the table to which you are adding a conflict resolution routine. The table can be the storage table of a nested table. |
| column_group | Name of the column group to which you are adding a conflict resolution routine. Column groups are required for update conflict resolution routines only. |
| constraint_name | Name of the unique constraint or unique index for which you are adding a conflict resolution routine. Use the name of the unique index if it differs from the name of the associated unique constraint. Constraint names are required for uniqueness conflict resolution routines only. |
| sequence_no | Order in which the designated conflict resolution methods should be applied. |
| method | Type of conflict resolution routine that you want to create. This can be the name of one of the standard routines provided with advanced replication, or, if you have written your own routine, choose user function, and provide the name of your method as the function_name parameter.<br>The standard methods supported in this release for update conflicts are:<br>■ minimum<br>■ maximum<br>■ latest timestamp<br>■ earliest timestamp<br>■ additive, average<br>■ priority group<br>■ site priority<br>■ overwrite<br>■ discard<br>The standard methods supported in this release for uniqueness conflicts are: append site name, append sequence, and discard. There are no built-in (Oracle supplied) methods for delete conflicts. |

*Table 18–13 (Cont.) ADD_conflicttype_RESOLUTION Procedure Parameters*

| Parameter | Description |
|---|---|
| parameter_column_name | Name of the columns used to resolve the conflict. The standard methods operate on a single column. For example, if you are using the latest timestamp method for a column group, then you should pass the name of the column containing the time stamp value as this parameter. If your are using a user function, then you can resolve the conflict using any number of columns.<br><br>For update or unique conflicts, this parameter accepts either a comma-delimited list of column names, or a PL/SQL associative array of type DBMS_REPCAT.VARCHAR2S or DBMS_UTILITY.LNAME_ARRAY. Use DBMS_UTILITY.LNAME_ARRAY if any column name is greater than or equal to 30 bytes, which might occur when you specify the attributes of column objects.<br><br>For delete conflicts, this parameter accepts either a comma-delimited list of column names or a PL/SQL associative array of type DBMS_REPCAT.VARCHAR2S.<br><br>The single value '*' indicates that you want to use all of the columns in the table (or column group, for update conflicts) to resolve the conflict. If you specify '*', then the columns are passed to your function in alphabetical order.<br><br>LOB columns cannot be specified for this parameter.<br><br>**See Also:** "Usage Notes" on page 18-19 if you are using column objects |
| priority_group | If you are using the priority group or site priority update conflict resolution method, then you must supply the name of the priority group that you have created.<br><br>See Chapter 6, "Configuring Conflict Resolution" and *Oracle Database Advanced Replication* for more information. If you are using a different method, you can use the default value for this parameter, NULL. This parameter is applicable to update conflicts only. |
| function_name | If you selected the user function method, or if you are adding a delete conflict resolution routine, then you must supply the name of the conflict resolution routine that you have written. If you are using one of the standard methods, then you can use the default value for this parameter, NULL. |
| comment | This user comment is added to the DBA_REPRESOLUTION view. |

## Exceptions

*Table 18–14 ADD_conflicttype_RESOLUTION Procedure Exceptions*

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missingobject | Specified object does not exist as a table in the specified schema using row-level replication. |
| missingschema | Specified schema does not exist. |
| missingcolumn | Column that you specified as part of the parameter_column_name parameter does not exist. |
| missinggroup | Specified column group does not exist. |
| missingprioritygroup | The priority group that you specified does not exist for the table. |
| invalidmethod | Resolution method that you specified is not recognized. |
| invalidparameter | Number of columns that you specified for the parameter_column_name parameter is invalid. (The standard routines take only one column name.) |
| missingfunction | User function that you specified does not exist. |
| missingconstraint | Constraint that you specified for a uniqueness conflict does not exist. |
| notquiesced | Replication group to which the specified table belongs is not quiesced. |
| duplicateresolution | Specified conflict resolution method is already registered. |
| duplicatesequence | The specified sequence number exists for the specified object. |
| invalidprioritygroup | The specified priority group does not exist. |
| paramtype | Type is different from the type assigned to the priority group. |

## Usage Notes

If you are using column objects, then whether you can specify the attributes of the column objects for the parameter_column_name parameter depends on whether the conflict resolution method is built-in (Oracle supplied) or user-created:

- If you are using a built-in conflict resolution method, then you can specify attributes of objects for this parameter. For example, if a column object named cust_address has street_address as an attribute, then you can specify cust_address.street_address for this parameter.
- If you are using a built-in conflict resolution method, the following types of columns cannot be specified for this parameter: LOB attribute of a column object, collection or collection attribute of a column object, REF, or an entire column object.
- If you are using a user-created conflict resolution method, then you must specify an entire column object. You cannot specify the attributes of a column object. For example, if a column object named cust_address has street_address as an attribute (among other attributes), then you can specify only cust_address for this parameter.

# ALTER_CATCHUP_PARAMETERS Procedure

This procedure alters the values for the following parameters stored in the DBA_REPEXTENSIONS data dictionary view:

- percentage_for_catchup_mdef
- cycle_seconds_mdef
- percentage_for_catchup_new
- cycle_seconds_new

These parameters were originally set by the ADD_NEW_MASTERS procedure. The new values you specify for these parameters are used during the remaining steps in the process of adding new master sites to a master group. These changes are only to the site at which it is executed. Therefore, it must be executed at each master site, including the master definition site, to alter parameters at all sites.

**See Also:**

- "ADD_NEW_MASTERS Procedure" on page 18-8
- "Adding New Master Sites" on page 7-2 for more information about adding master sites to a master group

## Syntax

```
DBMS_REPCAT.ALTER_CATCHUP_PARAMETERS (
   extension_id                  IN    RAW,
   percentage_for_catchup_mdef   IN    BINARY_INTEGER  := NULL,
   cycle_seconds_mdef            IN    BINARY_INTEGER  := NULL,
   percentage_for_catchup_new    IN    BINARY_INTEGER  := NULL,
   cycle_seconds_new             IN    BINARY_INTEGER  := NULL);
```

## Parameters

***Table 18–15 ALTER_CATCHUP_PARAMETERS Procedure Parameters***

| Parameter | Description |
|---|---|
| extension_id | The identifier for the current pending request to add master database without quiesce. You can find the extension_id by querying the DBA_REPSITES_NEW and DBA_REPEXTENSIONS data dictionary views. |
| percentage_for_catchup_mdef | The percentage of propagation resources to use for catching up propagation to the master definition site. Must be a multiple of 10 and must be between 0 and 100. |
| cycle_seconds_mdef | This parameter is meaningful when percentage_for_catchup_mdef is both meaningful and set to a value between 10 and 90, inclusive. In this case, propagation to the master definition site alternates between replication groups that are not being extended and replication groups that are being extended, with one push to each during each cycle. This parameter indicates the length of the cycle in seconds. |
| percentage_for_catchup_new | The percentage of propagation resources to use for catching up propagation to new master sites. Must be a multiple of 10 and must be between 0 and 100. |

*Table 18–15 (Cont.) ALTER_CATCHUP_PARAMETERS Procedure Parameters*

| Parameter | Description |
|---|---|
| cycle_seconds_new | This parameter is meaningful when percentage_for_catchup_new is both meaningful and set to a value between 10 and 90, inclusive. In this case, propagation to a new master alternates between replication groups that are not being extended and replication groups that are being extended, with one push to each during each cycle. This parameter indicates the length of the cycle in seconds. |

## Exceptions

*Table 18–16 ALTER_CATCHUP_PARAMETERS Procedure Exceptions*

| Exception | Description |
|---|---|
| typefailure | The parameter value specified for one of the parameters is not appropriate. |
| dbnotcompatible | Feature is incompatible with database version. All databases must be at 9.2.0 or higher compatibility level. |

# ALTER_MASTER_PROPAGATION Procedure

This procedure alters the propagation method for a specified replication group at a specified master site. This replication group must be quiesced. You must call this procedure from the master definition site. If the master appears in the dblink_list or dblink_table, then ALTER_MASTER_PROPAGATION ignores that database link. You cannot change the propagation mode from a master to itself.

## Syntax

```
DBMS_REPCAT.ALTER_MASTER_PROPAGATION (
   gname               IN   VARCHAR2,
   master              IN   VARCHAR2,
   { dblink_list       IN   VARCHAR2,
   | dblink_table      IN   DBMS_UTILITY.DBLINK_ARRAY,}
   propagation_mode    IN   VARCHAR2 : ='ASYNCHRONOUS',
   comment             IN   VARCHAR2 := '');
```

**Note:** This procedure is overloaded. The dblink_list and dblink_table parameters are mutually exclusive.

## Parameters

***Table 18–17 ALTER_MASTER_PROPAGATION Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Name of the replication group to which to alter the propagation mode. |
| master | Name of the master site at which to alter the propagation mode. |
| dblink_list | A comma-delimited list of database links for which to alter the propagation method. If NULL, then all masters except the master site being altered are used by default. |
| dblink_table | A PL/SQL associative array, indexed from position 1, of database links for which to alter propagation. |
| propagation_mode | Determines the manner in which changes from the specified master site are propagated to the sites identified by the list of database links. Appropriate values are synchronous and asynchronous. |
| comment | This comment is added to the DBA_REPPROP view. |

## Exceptions

***Table 18–18 ALTER_MASTER_PROPAGATION Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| notquiesced | Invocation site is not quiesced. |
| typefailure | Propagation mode specified was not recognized. |
| nonmaster | List of database links includes a site that is not a master site. |

## ALTER_MASTER_REPOBJECT Procedure

This procedure alters an object in your replication environment. You must call this procedure from the master definition site.

This procedure requires that you quiesce the master group of the object if either of the following conditions is true:

- You are altering a table in a multimaster replication environment.
- You are altering a table with the `safe_table_change` parameter set to `FALSE` in a single master replication environment.

You can use this procedure to alter non table objects without quiescing the master group.

### Syntax

```
DBMS_REPCAT.ALTER_MASTER_REPOBJECT (
   sname               IN   VARCHAR2,
   oname               IN   VARCHAR2,
   type                IN   VARCHAR2,
   ddl_text            IN   VARCHAR2,
   comment             IN   VARCHAR2     := '',
   retry               IN   BOOLEAN      := FALSE
   safe_table_change   IN   BOOLEAN      := FALSE);
```

### Parameters

***Table 18–19 ALTER_MASTER_REPOBJECT Procedure Parameters***

| Parameter | Description |
|---|---|
| `sname` | Schema containing the object that you want to alter. |
| `oname` | Name of the object that you want to alter. The object cannot be a storage table for a nested table. |
| `type` | Type of the object that you are altering. The following types are supported:<br>`FUNCTION` `SYNONYM`<br>`INDEX` `TABLE`<br>`INDEXTYPE` `TRIGGER`<br>`OPERATOR` `TYPE`<br>`PACKAGE` `TYPE BODY`<br>`PACKAGE BODY` `VIEW`<br>`PROCEDURE` |
| `ddl_text` | The DDL text that you want used to alter the object. Oracle does not parse this DDL before applying it. Therefore, you must ensure that your DDL text provides the appropriate schema and object name for the object being altered.<br>If the DDL is supplied without specifying a schema, then the default schema is the replication administrator's schema. Be sure to specify the schema if it is other than the replication administrator's schema. |
| `comment` | If not `NULL`, then this comment is added to the `COMMENT` field of the `DBA_REPOBJECT` view. |
| `retry` | If retry is `TRUE`, then `ALTER_MASTER_REPOBJECT` alters the object only at masters whose object status is not `VALID`. |

*Table 18–19 (Cont.) ALTER_MASTER_REPOBJECT Procedure Parameters*

| Parameter | Description |
|---|---|
| safe_table_change | Specify TRUE if the change to a table is safe. Specify FALSE if the change to a table is unsafe.<br>You can make safe changes to a master table in a single master replication environment without quiescing the master group that contains the table. To make unsafe changes, you must quiesce the master group.<br>Only specify this parameter for tables in single master replication environments. This parameter is ignored in multimaster replication environments and when the object specified is not a table. In multimaster replication environments, you must quiesce the master group to run the ALTER_MASTER_REPOBJECT procedure on a table.<br>The following are safe changes:<br>■ Changing storage and extent information<br>■ Making existing columns larger. For example, changing a VARCHAR2(20) column to a VARCHAR2(50) column.<br>■ Adding non primary key constraints<br>■ Altering non primary key constraints<br>■ Enabling and disabling non primary key constraints<br>The following are unsafe changes:<br>■ Changing the primary key by adding or deleting columns in the key<br>■ Adding or deleting columns<br>■ Making existing columns smaller. For example, changing a VARCHAR2(50) column to a VARCHAR2(20) column.<br>■ Disabling a primary key constraint<br>■ Changing the data type of an existing column<br>■ Dropping an existing column<br>If you are unsure whether a change is safe or unsafe, then quiesce the master group before you run the ALTER_MASTER_REPOBJECT procedure. |

## Exceptions

*Table 18–20 ALTER_MASTER_REPOBJECT Procedure Exceptions*

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| notquiesced | Associated replication group has not been suspended. |
| missingobject | Object identified by sname and oname does not exist. |
| typefailure | Specified type parameter is not supported. |
| ddlfailure | DDL at the master definition site did not succeed. |
| commfailure | At least one master site is not accessible. |

## ALTER_MVIEW_PROPAGATION Procedure

This procedure alters the propagation method for a specified replication group at the current materialized view site. This procedure pushes the deferred transaction queue at the materialized view site, locks the materialized views, and regenerates any triggers and their associated packages. You must call this procedure from the materialized view site.

### Syntax

```
DBMS_REPCAT.ALTER_MVIEW_PROPAGATION (
   gname                IN  VARCHAR2,
   propagation_mode     IN  VARCHAR2,
   comment              IN  VARCHAR2   := '',
   gowner               IN  VARCHAR2   := 'PUBLIC');
```

### Parameters

***Table 18–21 ALTER_MVIEW_PROPAGATION Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Name of the replication group for which to alter the propagation method. |
| propagation_mode | Manner in which changes from the current materialized view site are propagated to its associated master site or master materialized view site. Appropriate values are synchronous and asynchronous. |
| comment | This comment is added to the DBA_REPPROP view. |
| gowner | Owner of the materialized view group. |

### Exceptions

***Table 18–22 ALTER_MVIEW_PROPAGATION Procedure Exceptions***

| Exception | Description |
|---|---|
| missingrepgroup | Specified replication group does not exist. |
| typefailure | Propagation mode was specified incorrectly. |
| nonmview | Current site is not a materialized view site for the specified replication group. |
| commfailure | Cannot contact master site or master materialized view site. |
| failaltermviewrop | Materialized view group propagation can be altered only when there are no other materialized view groups with the same master site or master materialized view site sharing the materialized view site. |

# ALTER_PRIORITY Procedure

This procedure alters the priority level associated with a specified priority group member. You must call this procedure from the master definition site.

> **See Also:** Chapter 6, "Configuring Conflict Resolution" and *Oracle Database Advanced Replication* for more information about conflict resolution methods

## Syntax

```
DBMS_REPCAT.ALTER_PRIORITY (
   gname            IN   VARCHAR2,
   pgroup           IN   VARCHAR2,
   old_priority     IN   NUMBER,
   new_priority     IN   NUMBER);
```

## Parameters

***Table 18–23 ALTER_PRIORITY Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Master group with which the priority group is associated. |
| pgroup | Name of the priority group containing the priority that you want to alter. |
| old_priority | Current priority level of the priority group member. |
| new_priority | New priority level that you want assigned to the priority group member. |

## Exceptions

***Table 18–24 ALTER_PRIORITY Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| duplicatepriority | New priority level exists in the priority group. |
| missingrepgroup | Specified master group does not exist. |
| missingvalue | Value was not registered by a call to DBMS_REPCAT.ADD_PRIORITY_*datatype*. |
| missingprioritygroup | Specified priority group does not exist. |
| notquiesced | Specified master group is not quiesced. |

## ALTER_PRIORITY_*datatype* Procedure

This procedure alters the value of a member in a priority group. You must call this procedure from the master definition site. The procedure that you must call is determined by the data type of your `priority` column.

> **See Also:** Chapter 6, "Configuring Conflict Resolution" and *Oracle Database Advanced Replication* for more information about conflict resolution methods

### Syntax

```
DBMS_REPCAT.ALTER_PRIORITY_datatype (
   gname        IN   VARCHAR2,
   pgroup       IN   VARCHAR2,
   old_value    IN   datatype,
   new_value    IN   datatype);
```

where *datatype*:

```
{ NUMBER
| VARCHAR2
| CHAR
| DATE
| RAW
| NCHAR
| NVARCHAR2 }
```

### Parameters

***Table 18–25 ALTER_PRIORITY_datatype Procedure Parameters***

| Parameter | Description |
|---|---|
| `gname` | Master group with which the priority group is associated. |
| `pgroup` | Name of the priority group containing the value that you want to alter. |
| `old_value` | Current value of the priority group member. |
| `new_value` | New value that you want assigned to the priority group member. |

### Exceptions

***Table 18–26 ALTER_PRIORITY_datatype Procedure Exceptions***

| Exception | Description |
|---|---|
| `nonmasterdef` | Invocation site is not the master definition site. |
| `duplicatevalue` | New value exists in the priority group. |
| `missingrepgroup` | Specified master group does not exist. |
| `missingprioritygroup` | Specified priority group does not exist. |
| `missingvalue` | Old value does not exist. |
| `paramtype` | New value has the incorrect data type for the priority group. |
| `typefailure` | Specified value has the incorrect data type for the priority group. |
| `notquiesced` | Specified master group is not quiesced. |

# ALTER_SITE_PRIORITY Procedure

This procedure alters the priority level associated with a specified site. You must call this procedure from the master definition site.

> **See Also:** Chapter 6, "Configuring Conflict Resolution" and *Oracle Database Advanced Replication* for more information about conflict resolution methods

## Syntax

```
DBMS_REPCAT.ALTER_SITE_PRIORITY (
   gname         IN   VARCHAR2,
   name          IN   VARCHAR2,
   old_priority  IN   NUMBER,
   new_priority  IN   NUMBER);
```

## Parameters

***Table 18–27 ALTER_SITE_PRIORITY Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Master group with which the site priority group is associated. |
| name | Name of the site priority group whose member you are altering. |
| old_priority | Current priority level of the site whose priority level you want to change. |
| new_priority | New priority level for the site. A higher number indicates a higher priority level. |

## Exceptions

***Table 18–28 ALTER_SITE_PRIORITY Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missingrepgroup | Specified master group does not exist. |
| missingpriority | Old priority level is not associated with any group members. |
| duplicatepriority | New priority level exists for another site in the group. |
| missingvalue | Old value does not exist. |
| paramtype | New value has the incorrect data type for the priority group. |
| notquiesced | Master group is not quiesced. |

## ALTER_SITE_PRIORITY_SITE Procedure

This procedure alters the site associated with a specified priority level. You must call this procedure from the master definition site.

> **See Also:** Chapter 6, "Configuring Conflict Resolution" and *Oracle Database Advanced Replication* for more information about conflict resolution methods

### Syntax

```
DBMS_REPCAT.ALTER_SITE_PRIORITY_SITE (
   gname     IN   VARCHAR2,
   name      IN   VARCHAR2,
   old_site  IN   VARCHAR2,
   new_site  IN   VARCHAR2);
```

### Parameters

***Table 18–29 ALTER_SITE_PRIORITY_SITE Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Master group with which the site priority group is associated. |
| name | Name of the site priority group whose member you are altering. |
| old_site | Current global database name of the site to disassociate from the priority level. |
| new_site | New global database name that you want to associate with the current priority level. |

### Exceptions

***Table 18–30 ALTER_SITE_PRIORITY_SITE Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missingrepgroup | Specified master group does not exist. |
| missingpriority | Specified site priority group does not exist. |
| missingvalue | Old site is not a group member. |
| notquiesced | Master group is not quiesced. |

# CANCEL_STATISTICS Procedure

This procedure stops the collection of statistics about the successful resolution of update, uniqueness, and delete conflicts for a table.

## Syntax

```
DBMS_REPCAT.CANCEL_STATISTICS (
   sname    IN   VARCHAR2,
   oname    IN   VARCHAR2);
```

## Parameters

***Table 18–31 CANCEL_STATISTICS Procedure Parameters***

| Parameter | Description |
|---|---|
| sname | Name of the schema in which the table is located. |
| oname | Name of the table for which you do not want to gather conflict resolution statistics. |

## Exceptions

***Table 18–32 CANCEL_STATISTICS Procedure Exceptions***

| Exception | Description |
|---|---|
| missingschema | Specified schema does not exist. |
| missingobject | Specified table does not exist. |
| statnotreg | Specified table is not currently registered to collect statistics. |

## COMMENT_ON_COLUMN_GROUP Procedure

This procedure updates the comment field in the DBA_REPCOLUMN_GROUP view for a column group. This comment is not added at all master sites until the next call to DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT.

### Syntax

```
DBMS_REPCAT.COMMENT_ON_COLUMN_GROUP (
   sname           IN   VARCHAR2,
   oname           IN   VARCHAR2,
   column_group    IN   VARCHAR2,
   comment         IN   VARCHAR2);
```

### Parameters

*Table 18–33 COMMENT_ON_COLUMN_GROUP Procedure Parameters*

| Parameter | Description |
|---|---|
| sname | Name of the schema in which the object is located. |
| oname | Name of the replicated table with which the column group is associated. |
| column_group | Name of the column group. |
| comment | Text of the updated comment that you want included in the GROUP_COMMENT field of the DBA_REPCOLUMN_GROUP view. |

### Exceptions

*Table 18–34 COMMENT_ON_COLUMN_GROUP Procedure Exceptions*

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missinggroup | Specified column group does not exist. |
| missingobj | Object is missing. |

# COMMENT_ON_MVIEW_REPSITES Procedure

This procedure updates the SCHEMA_COMMENT field in the DBA_REPGROUP data dictionary view for the specified materialized view group. The group name must be registered locally as a replicated materialized view group. This procedure must be executed at the materialized view site.

## Syntax

```
DBMS_REPCAT.COMMENT_ON_MVIEW_REPSITES (
   gowner    IN   VARCHAR2,
   gname     IN   VARCHAR2,
   comment   IN   VARCHAR2);
```

## Parameters

***Table 18–35 COMMENT_ON_MVIEW_REPSITES Procedure Parameters***

| Parameter | Description |
|---|---|
| gowner | Owner of the materialized view group. |
| gname | Name of the materialized view group. |
| comment | Updated comment to include in the SCHEMA_COMMENT field of the DBA_REPGROUP view. |

## Exceptions

***Table 18–36 COMMENT_ON_MVIEW_REPSITES Procedure Exceptions***

| Parameter | Description |
|---|---|
| missingrepgroup | The materialized view group does not exist. |
| nonmview | The connected site is not a materialized view site. |

## COMMENT_ON_PRIORITY_GROUP Procedures

This procedure updates the comment field in the DBA_REPPRIORITY_GROUP view for a priority group. This comment is not added at all master sites until the next call to GENERATE_REPLICATION_SUPPORT.

### Syntax

```
DBMS_REPCAT.COMMENT_ON_PRIORITY_GROUP (
   gname       IN   VARCHAR2,
   pgroup      IN   VARCHAR2,
   comment     IN   VARCHAR2);
```

### Parameters

*Table 18–37 COMMENT_ON_PRIORITY_GROUP Procedure Parameters*

| Parameter | Description |
|---|---|
| gname | Name of the master group. |
| pgroup | Name of the priority group. |
| comment | Text of the updated comment that you want included in the PRIORITY_COMMENT field of the DBA_REPPRIORITY_GROUP view. |

### Exceptions

*Table 18–38 COMMENT_ON_PRIORITY_GROUP Procedure Exceptions*

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missingrepgroup | Specified master group does not exist. |
| missingprioritygroup | Specified priority group does not exist. |

# COMMENT_ON_REPGROUP Procedure

This procedure updates the comment field in the DBA_REPGROUP view for a master group. This procedure must be issued at the master definition site.

## Syntax

```
DBMS_REPCAT.COMMENT_ON_REPGROUP (
   gname     IN   VARCHAR2,
   comment   IN   VARCHAR2);
```

## Parameters

***Table 18–39 COMMENT_ON_REPGROUP Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Name of the replication group that you want to comment on. |
| comment | Updated comment to include in the SCHEMA_COMMENT field of the DBA_REPGROUP view. |

## Exceptions

***Table 18–40 COMMENT_ON_REPGROUP Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| commfailure | At least one master site is not accessible. |

## COMMENT_ON_REPOBJECT Procedure

This procedure updates the comment field in the DBA_REPOBJECT view for a replicated object in a master group. This procedure must be issued at the master definition site.

### Syntax

```
DBMS_REPCAT.COMMENT_ON_REPOBJECT (
   sname    IN   VARCHAR2,
   oname    IN   VARCHAR2,
   type     IN   VARCHAR2,
   comment  IN   VARCHAR2);
```

### Parameters

*Table 18–41 COMMENT_ON_REPOBJECT Procedure Parameters*

| Parameter | Description |
|---|---|
| sname | Name of the schema in which the object is located. |
| oname | Name of the object that you want to comment on. The object cannot be a storage table for a nested table. |
| type | Type of the object. The following types are supported:<br>FUNCTION<br>INDEX<br>INDEXTYPE<br>OPERATOR<br>PACKAGE<br>PACKAGE BODY<br>PROCEDURE<br>SYNONYM<br>TABLE<br>TRIGGER<br>TYPE<br>TYPE BODY<br>VIEW |
| comment | Text of the updated comment that you want to include in the OBJECT_COMMENT field of the DBA_REPOBJECT view. |

### Exceptions

*Table 18–42 COMMENT_ON_REPOBJECT Procedure Exceptions*

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missingobject | Specified object does not exist. |
| typefailure | Specified type parameter is not supported. |
| commfailure | At least one master site is not accessible. |

# COMMENT_ON_REPSITES Procedure

If the replication group is a master group, then this procedure updates the MASTER_COMMENT field in the DBA_REPSITES view for a master site. If the replication group is a materialized view group, this procedure updates the SCHEMA_COMMENT field in the DBA_REPGROUP view for a materialized view site.

This procedure can be executed at either a master site or a materialized view site. If you execute this procedure on a a materialized view site, then the materialized view group owner must be PUBLIC.

> **See Also:** "COMMENT_ON_conflicttype_RESOLUTION Procedure" on page 18-38 for instructions on placing a comment in the SCHEMA_COMMENT field of the DBA_REPGROUP view for a materialized view site if the materialized view group owner is not PUBLIC

## Syntax

```
DBMS_REPCAT.COMMENT_ON_REPSITES (
   gname       IN   VARCHAR2,
   [ master    IN   VARCHAR,]
   comment     IN   VARCHAR2);
```

## Parameters

***Table 18–43 COMMENT_ON_REPSITES Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Name of the replication group. This avoids confusion if a database is a master site in multiple replication environments. |
| master | The fully qualified database name of the master site on which you want to comment. If you are executing the procedure on a master site, then this parameter is required. To update comments at a materialized view site, omit this parameter. This parameter is optional. |
| comment | Text of the updated comment that you want to include in the comment field of the appropriate dictionary view. If the site is a master site, then this procedure updates the MASTER_COMMENT field of the DBA_REPSITES view. If the site is a materialized view site, then this procedure updates the SCHEMA_COMMENT field of the DBA_REPGROUP view. |

## Exceptions

***Table 18–44 COMMENT_ON_REPSITES Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| nonmaster | Invocation site is not a master site. |
| commfailure | At least one master site is not accessible. |
| missingrepgroup | Replication group does not exist. |
| commfailure | One or more master sites are not accessible. |
| corrupt | There is an inconsistency in the replication catalog views. |

## COMMENT_ON_SITE_PRIORITY Procedure

This procedure updates the comment field in the DBA_REPPRIORITY_GROUP view for a site priority group. This procedure is a wrapper for the COMMENT_ON_COLUMN_GROUP procedure and is provided as a convenience only. This procedure must be issued at the master definition site.

### Syntax

```
DBMS_REPCAT.COMMENT_ON_SITE_PRIORITY (
   gname       IN   VARCHAR2,
   name        IN   VARCHAR2,
   comment     IN   VARCHAR2);
```

### Parameters

*Table 18–45 COMMENT_ON_SITE_PRIORITY Procedure Parameters*

| Parameter | Description |
|---|---|
| gname | Name of the master group. |
| name | Name of the site priority group. |
| comment | Text of the updated comment that you want included in the PRIORITY_COMMENT field of the DBA_REPPRIORITY_GROUP view. |

### Exceptions

*Table 18–46 COMMENT_ON_SITE_PRIORITY Procedure Exceptions*

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missingrepgroup | Specified master group does not exist. |
| missingprioritygroup | Specified priority group does not exist. |

# COMMENT_ON_*conflicttype*_RESOLUTION Procedure

This procedure updates the RESOLUTION_COMMENT field in the DBA_REPRESOLUTION view for a conflict resolution routine. The procedure that you must call is determined by the type of conflict that the routine resolves. These procedures must be issued at the master definition site.

***Table 18–47 COMMENT_ON_conflicttype_RESOLUTION Procedures***

| Conflict Type | Procedure Name |
|---|---|
| update | COMMENT_ON_UPDATE_RESOLUTION |
| uniqueness | COMMENT_ON_UNIQUE_RESOLUTION |
| delete | COMMENT_ON_DELETE_RESOLUTION |

The comment is not added at all master sites until the next call to GENERATE_REPLICATION_SUPPORT.

## Syntax

```
DBMS_REPCAT.COMMENT_ON_UPDATE_RESOLUTION (
   sname                IN   VARCHAR2,
   oname                IN   VARCHAR2,
   column_group         IN   VARCHAR2,
   sequence_no          IN   NUMBER,
   comment              IN   VARCHAR2);

DBMS_REPCAT.COMMENT_ON_UNIQUE_RESOLUTION (
   sname                IN   VARCHAR2,
   oname                IN   VARCHAR2,
   constraint_name      IN   VARCHAR2,
   sequence_no          IN   NUMBER,
   comment              IN   VARCHAR2);

DBMS_REPCAT.COMMENT_ON_DELETE_RESOLUTION (
   sname                IN   VARCHAR2,
   oname                IN   VARCHAR2,
   sequence_no          IN   NUMBER,
   comment              IN   VARCHAR2);
```

## Parameters

*Table 18–48 COMMENT_ON_conflicttype_RESOLUTION Procedure Parameters*

| Parameter | Description |
|---|---|
| sname | Name of the schema. |
| oname | Name of the replicated table with which the conflict resolution routine is associated. |
| column_group | Name of the column group with which the update conflict resolution routine is associated. |
| constraint_name | Name of the unique constraint with which the uniqueness conflict resolution routine is associated. |
| sequence_no | Sequence number of the conflict resolution procedure. |
| comment | The text of the updated comment that you want included in the RESOLUTION_COMMENT field of the DBA_REPRESOLUTION view. |

## Exceptions

*Table 18–49 COMMENT_ON_conflicttype_RESOLUTION Procedure Exceptions*

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missingobject | Specified object does not exist. |
| missingresolution | Specified conflict resolution routine is not registered. |

# COMPARE_OLD_VALUES Procedure

This procedure specifies whether to compare old column values during propagation of deferred transactions at each master site for each nonkey column of a replicated table for updates and deletes. The default is to compare old values for all columns. You can change this behavior at all master sites and materialized view sites by invoking DBMS_REPCAT.COMPARE_OLD_VALUES at the master definition site.

When you use user-defined types, you can specify leaf attributes of a column object, or you can specify an entire column object. For example, if a column object named cust_address has street_address as an attribute, then you can specify cust_address.street_address for the column_list parameter or as part of the column_table parameter, or you can specify only cust_address.

When performing equality comparisons for conflict detection, Oracle treats objects as equal only if one of the following conditions is true:

- Both objects are atomically NULL (the entire object is NULL)
- All of the corresponding attributes are equal in the objects

Given these conditions, if one object is atomically NULL while the other is not, then Oracle does not consider the objects to be equal. Oracle does not consider MAP and ORDER methods when performing equality comparisons.

## Syntax

```
DBMS_REPCAT.COMPARE_OLD_VALUES(
   sname           IN  VARCHAR2,
   oname           IN  VARCHAR2,
   { column_list   IN  VARCHAR2,
   | column_table  IN  DBMS_UTILITY.VARCHAR2S | DBMS_UTILITY.LNAME_ARRAY,}
   operation       IN  VARCHAR2 := 'UPDATE',
   compare         IN  BOOLEAN  := TRUE );
```

**Note:** This procedure is overloaded. The column_list and column_table parameters are mutually exclusive.

## Parameters

***Table 18–50 COMPARE_OLD_VALUES Procedure Parameters***

| Parameter | Description |
|---|---|
| sname | Schema in which the table is located. |
| oname | Name of the replicated table. The table can be the storage table of a nested table. |
| column_list | A comma-delimited list of the columns in the table. There must be no spaces between entries. |
| column_table | Instead of a list, you can use a PL/SQL associative array of type DBMS_REPCAT.VARCHAR2S or DBMS_UTILITY.LNAME_ARRAY to contain the column names. The first column name should be at position 1, the second at position 2, and so on.<br>Use DBMS_UTILITY.LNAME_ARRAY if any column name is greater than or equal to 30 bytes, which might occur when you specify the attributes of column objects. |

*Table 18–50 (Cont.) COMPARE_OLD_VALUES Procedure Parameters*

| Parameter | Description |
|---|---|
| operation | Possible values are: update, delete, or the asterisk wildcard '*', which means update and delete. |
| compare | If compare is TRUE, the old values of the specified columns are compared when sent. If compare is FALSE, the old values of the specified columns are not compared when sent. Unspecified columns and unspecified operations are not affected. The specified change takes effect at the master definition site as soon as min_communication is TRUE for the table. The change takes effect at a master site or at a materialized view site the next time replication support is generated at that site with min_communication TRUE. |

**Note:** The operation parameter enables you to decide whether to compare old values for nonkey columns when rows are deleted or updated. If you do not compare the old value, then Oracle assumes the old value is equal to the current value of the column at the target side when the update or delete is applied.

See *Oracle Database Advanced Replication* for more information about reduced data propagation using the COMPARE_OLD_VALUES procedure before changing the default behavior of Oracle.

## Exceptions

*Table 18–51 COMPARE_OLD_VALUES Procedure Exceptions*

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missingobject | Specified object does not exist as a table in the specified schema waiting for row-level replication information. |
| missingcolumn | At least one column is not in the table. |
| notquiesced | Master group has not been quiesced. |
| typefailure | An illegal operation is specified. |
| keysendcomp | A specified column is a key column in a table. |
| dbnotcompatible | Feature is incompatible with database version. Typically, this exception arises when you are trying to compare the attributes of column objects. In this case, all databases must be at 9.2.0 or higher compatibility level. |

# CREATE_MASTER_REPGROUP Procedure

This procedure creates a new, empty, quiesced master group.

## Syntax

```
DBMS_REPCAT.CREATE_MASTER_REPGROUP (
   gname            IN   VARCHAR2,
   group_comment    IN   VARCHAR2     := '',
   master_comment   IN   VARCHAR2     := '',
   qualifier        IN   VARCHAR2     := '');
```

## Parameters

***Table 18–52 CREATE_MASTER_REPGROUP Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Name of the master group that you want to create. |
| group_comment | This comment is added to the DBA_REPGROUP view. |
| master_comment | This comment is added to the DBA_REPSITES view. |
| qualifier | Connection qualifier for master group. Be sure to use the @ sign. See *Oracle Database Advanced Replication* and *Oracle Database Administrator's Guide* for more information about connection qualifiers. |

## Exceptions

***Table 18–53 CREATE_MASTER_REPGROUP Procedure Exceptions***

| Exception | Description |
|---|---|
| duplicaterepgroup | Master group exists. |
| norepopt | Advanced replication option is not installed. |
| missingrepgroup | Master group name was not specified. |
| qualifiertoolong | Connection qualifier is too long. |

## CREATE_MASTER_REPOBJECT Procedure

This procedure makes an object a replicated object by adding the object to a master group. This procedure preserves the object identifier for user-defined types and object tables at all replication sites.

Replication of clustered tables is supported, but the `use_existing_object` parameter cannot be set to `FALSE` for clustered tables. In other words, you must create the clustered table at all master sites participating in the master group before you execute the `CREATE_MASTER_REPOBJECT` procedure. However, these tables do not need to contain the table data. So, the `copy_rows` parameter can be set to `TRUE` for clustered tables.

### Syntax

```
DBMS_REPCAT.CREATE_MASTER_REPOBJECT (
   sname                IN   VARCHAR2,
   oname                IN   VARCHAR2,
   type                 IN   VARCHAR2,
   use_existing_object  IN   BOOLEAN     := TRUE,
   ddl_text             IN   VARCHAR2    := NULL,
   comment              IN   VARCHAR2    := '',
   retry                IN   BOOLEAN     := FALSE,
   copy_rows            IN   BOOLEAN     := TRUE,
   gname                IN   VARCHAR2    := '');
```

### Parameters

*Table 18–54 CREATE_MASTER_REPOBJECT Procedure Parameters*

| Parameters | Description |
|---|---|
| `sname` | Name of the schema in which the object that you want to replicate is located. |
| `oname` | Name of the object you are replicating. If `ddl_text` is `NULL`, then this object must exist in the specified schema. To ensure uniqueness, table names should be a maximum of 27 bytes long, and package names should be no more than 24 bytes. The object cannot be a storage table for a nested table. |
| `type` | Type of the object that you are replicating. The following types are supported:<br>`FUNCTION` `SYNONYM`<br>`INDEX` `TABLE`<br>`INDEXTYPE` `TRIGGER`<br>`OPERATOR` `TYPE`<br>`PACKAGE` `TYPE BODY`<br>`PACKAGE BODY` `VIEW`<br>`PROCEDURE` |

*Table 18–54 (Cont.) CREATE_MASTER_REPOBJECT Procedure Parameters*

| Parameters | Description |
|---|---|
| use_existing_object | Indicate TRUE to reuse any objects of the same type and shape at the current master sites. See Table 18–56 for more information.<br><br>**Note:** This parameter must be set to TRUE for clustered tables. |
| ddl_text | If the object does not exist at the master definition site, then you must supply the DDL text necessary to create this object. PL/SQL packages, package bodies, procedures, and functions must have a trailing semicolon. SQL statements do not end with trailing semicolon. Oracle does not parse this DDL before applying it; therefore, you must ensure that your DDL text provides the appropriate schema and object name for the object being created.<br><br>If the DDL is supplied without specifying a schema (sname parameter), then the default schema is the replication administrator's schema. Be sure to specify the schema if it is other than the replication administrator's schema.<br><br>**Note:** Do not use the ddl_text parameter to add user-defined types or object tables. Instead, create the object first and then add the object. |
| comment | This comment is added to the OBJECT_COMMENT field of the DBA_REPOBJECT view. |
| retry | Indicate TRUE if you want Oracle to reattempt to create an object that it was previously unable to create. Use this if the error was transient or has since been rectified, or if you previously had insufficient resources. If this is TRUE, then Oracle creates the object only at master sites whose object status is not VALID. |
| copy_rows | Indicate TRUE if you want the initial contents of a newly replicated object to match the contents of the object at the master definition site. See Table 18–56 for more information. |
| gname | Name of the replication group in which you want to create the replicated object. The schema name is used as the default replication group name if none is specified, and a replication group with the same name as the schema must exist for the procedure to complete successfully in that case. |

## Exceptions

*Table 18–55 CREATE_MASTER_REPOBJECT Procedure Exceptions*

| Exceptions | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| notquiesced | Master group is not quiesced. |
| duplicateobject | Specified object exists in the master group and retry is FALSE, or if a name conflict occurs. |
| missingobject | Object identified by sname and oname does not exist and appropriate DDL has not been provided. |
| typefailure | Objects of the specified type cannot be replicated. |
| ddlfailure | DDL at the master definition site did not succeed. |
| commfailure | At least one master site is not accessible. |

## Object Creations

*Table 18–56 Object Creation at Master Sites*

| Object Already Exists? | COPY_ROWS | USE_EXISTING_OBJECTS | Result |
|---|---|---|---|
| yes | TRUE | TRUE | duplicatedobject message if objects do not match. For tables, use data from master definition site. |
| yes | FALSE | TRUE | duplicatedobject message if objects do not match. For tables, DBA must ensure contents are identical. |
| yes | TRUE/FALSE | FALSE | duplicatedobject message. |
| no | TRUE | TRUE/FALSE | Object is created. Tables populated using data from master definition site. |
| no | FALSE | TRUE/FALSE | Object is created. DBA must populate tables and ensure consistency of tables at all sites. |

# CREATE_MVIEW_REPGROUP Procedure

This procedure creates a new, empty materialized view group in your local database. CREATE_MVIEW_REPGROUP automatically calls REGISTER_MIEW_REPGROUP, but ignores any errors that might have happened during registration.

## Syntax

```
DBMS_REPCAT.CREATE_MVIEW_REPGROUP (
   gname              IN   VARCHAR2,
   master             IN   VARCHAR2,
   comment            IN   VARCHAR2     := '',
   propagation_mode   IN   VARCHAR2     := 'ASYNCHRONOUS',
   fname              IN   VARCHAR2     := NULL
   gowner             IN   VARCHAR2     := 'PUBLIC');
```

## Parameters

***Table 18–57 CREATE_MVIEW_REPGROUP Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Name of the replication group. This group must exist at the specified master site or master materialized view site. |
| master | Fully qualified database name of the database in the replication environment to use as the master site or master materialized view site. You can include a connection qualifier if necessary. See *Oracle Database Advanced Replication* and *Oracle Database Administrator's Guide* for information about using connection qualifiers. |
| comment | This comment is added to the DBA_REPGROUP view. |
| propagation_mode | Method of propagation for all updatable materialized views in the replication group. Acceptable values are synchronous and asynchronous. |
| fname | This parameter is for internal use only.<br>**Note:** Do not set this parameter unless directed to do so by Oracle Support Services. |
| gowner | Owner of the materialized view group. |

## Exceptions

***Table 18–58 CREATE_MVIEW_REPGROUP Procedure Exceptions***

| Exception | Description |
|---|---|
| duplicaterepgroup | Replication group exists at the invocation site. |
| nonmaster | Specified database is not a master site or master materialized view site. |
| commfailure | Specified database is not accessible. |
| norepopt | Advanced replication option is not installed. |
| typefailure | Propagation mode was specified incorrectly. |
| missingrepgroup | Replication group does not exist at master site. |
| invalidqualifier | Connection qualifier specified for the master site or master materialized view site is not valid for the replication group. |

*Table 18–58 (Cont.) CREATE_MVIEW_REPGROUP Procedure Exceptions*

| Exception | Description |
|---|---|
| `alreadymastered` | At the local site, there is another materialized view group with the same group name, but different master site or master materialized view site. |

# CREATE_MVIEW_REPOBJECT Procedure

This procedure adds a replicated object to a materialized view group.

## Syntax

```
DBMS_REPCAT.CREATE_MVIEW_REPOBJECT (
   sname                   IN   VARCHAR2,
   oname                   IN   VARCHAR2,
   type                    IN   VARCHAR2,
   ddl_text                IN   VARCHAR2  := '',
   comment                 IN   VARCHAR2  := '',
   gname                   IN   VARCHAR2  := '',
   gen_objs_owner          IN   VARCHAR2  := '',
   min_communication       IN   BOOLEAN   := TRUE,
   generate_80_compatible  IN   BOOLEAN   := TRUE,
   gowner                  IN   VARCHAR2  := 'PUBLIC');
```

## Parameters

***Table 18–59 CREATE_MVIEW_REPOBJECT Procedure Parameters***

| Parameter | Description |
|---|---|
| sname | Name of the schema in which the object is located. The schema must be same as the schema that owns the master table or master materialized view on which this materialized view is based. |
| oname | Name of the object that you want to add to the replicated materialized view group. |
| type | Type of the object that you are replicating. The following types are supported:<br>FUNCTION SNAPSHOT<br>INDEX SYNONYM<br>INDEXTYPE TRIGGER<br>OPERATOR TYPE<br>PACKAGE TYPE BODY<br>PACKAGE BODY VIEW<br>PROCEDURE<br>Use SNAPSHOT type of the object is a materialized view. |

*Table 18–59 (Cont.) CREATE_MVIEW_REPOBJECT Procedure Parameters*

| Parameter | Description |
|---|---|
| ddl_text | For objects of type MATERIALIZED VIEW, the DDL needed to create the object. For other types, use the default:<br>'' (an empty string)<br>If a materialized view with the same name exists, then Oracle ignores the DDL and registers the existing materialized view as a replicated object. If the master table or master materialized view for a materialized view does not exist in the replication group of the master designated for this schema, then Oracle raises a missingobject error.<br>If the DDL is supplied without specifying a schema, then the default schema is the replication administrator's schema. Be sure to specify the schema if it is other than the replication administrator's schema.<br>If the object is not of type MATERIALIZED VIEW, then the materialized view site connects to the master site or master materialized view site and pulls down the DDL text to create the object. If the object type is TYPE or TYPE BODY, then the object identifier (OID) for the object at the materialized view site is the same as the OID at the master site or master materialized view site. |
| comment | This comment is added to the OBJECT_COMMENT field of the DBA_REPOBJECT view. |
| gname | Name of the replicated materialized view group to which you are adding an object. The schema name is used as the default group name if none is specified, and a materialized view group with the same name as the schema must exist for the procedure to complete successfully. |
| gen_objs_owner | Name of the user you want to assign as owner of the transaction. |
| min_communication | This parameter is obsolete. Use the default value (TRUE). |
| generate_80_compatible | Set to FALSE because interoperability is not supported between Oracle8*i* databases and Oracle 11*g* or later databases. |
| gowner | Owner of the materialized view group. |

## Exceptions

*Table 18–60 CREATE_MVIEW_REPOBJECT Procedure Exceptions*

| Exception | Description |
|---|---|
| nonmview | Invocation site is not a materialized view site. |
| nonmaster | Master is no longer a master site or master materialized view site. |
| missingobject | Specified object does not exist in the master's replication group. |
| duplicateobject | Specified object exists with a different shape. |
| typefailure | Type is not an allowable type. |
| ddlfailure | DDL did not succeed. |
| commfailure | Master site or master materialized view site is not accessible. |
| missingschema | Schema does not exist as a database schema. |
| badmviewddl | DDL was executed but materialized view does not exist. |
| onlyonemview | Only one materialized view for master table or master materialized view can be created. |
| badmviewname | Materialized view differs from master table or master materialized view. |
| missingrepgroup | Replication group at the master does not exist. |

## DEFINE_COLUMN_GROUP Procedure

This procedure creates an empty column group. You must call this procedure from the master definition site.

> **See Also:** Chapter 6, "Configuring Conflict Resolution" and *Oracle Database Advanced Replication* for more information about conflict resolution methods

### Syntax

```
DBMS_REPCAT.DEFINE_COLUMN_GROUP (
   sname             IN   VARCHAR2,
   oname             IN   VARCHAR2,
   column_group      IN   VARCHAR2,
   comment           IN   VARCHAR2 := NULL);
```

### Parameters

***Table 18–61 DEFINE_COLUMN_GROUP Procedure Parameters***

| Parameter | Description |
|---|---|
| sname | Schema in which the replicated table is located. |
| oname | Name of the replicated table for which you are creating a column group. |
| column_group | Name of the column group that you want to create. |
| comment | This user text is displayed in the DBA_REPCOLUMN_GROUP view. |

### Exceptions

***Table 18–62 DEFINE_COLUMN_GROUP Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missingobject | Specified table does not exist. |
| duplicategroup | Specified column group exists for the table. |
| notquiesced | Replication group to which the specified table belongs is not quiesced. |

# DEFINE_PRIORITY_GROUP Procedure

This procedure creates a new priority group for a master group. You must call this procedure from the master definition site.

> **See Also:** Chapter 6, "Configuring Conflict Resolution" and *Oracle Database Advanced Replication* for more information about conflict resolution methods

## Syntax

```
DBMS_REPCAT.DEFINE_PRIORITY_GROUP (
   gname          IN   VARCHAR2,
   pgroup         IN   VARCHAR2,
   datatype       IN   VARCHAR2,
   fixed_length   IN   INTEGER := NULL,
   comment        IN   VARCHAR2 := NULL);
```

## Parameters

***Table 18–63 DEFINE_PRIORITY_GROUP Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Master group for which you are creating a priority group. |
| pgroup | Name of the priority group that you are creating. |
| datatype | Data type of the priority group members. The data types supported are: CHAR, VARCHAR2, NUMBER, DATE, RAW, NCHAR, and NVARCHAR2. |
| fixed_length | You must provide a column length for the CHAR data type. All other types can use the default, NULL. |
| comment | This user comment is added to the DBA_REPPRIORITY view. |

## Exceptions

***Table 18–64 DEFINE_PRIORITY_GROUP Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missingrepgroup | Specified master group does not exist. |
| duplicatepriority group | Specified priority group exists in the master group. |
| typefailure | Specified data type is not supported. |
| notquiesced | Master group is not quiesced. |

## DEFINE_SITE_PRIORITY Procedure

This procedure creates a new site priority group for a master group. You must call this procedure from the master definition site.

> **See Also:** Chapter 6, "Configuring Conflict Resolution" and *Oracle Database Advanced Replication* for more information about conflict resolution methods

### Syntax

```
DBMS_REPCAT.DEFINE_SITE_PRIORITY (
   gname        IN   VARCHAR2,
   name         IN   VARCHAR2,
   comment      IN   VARCHAR2 := NULL);
```

### Parameters

***Table 18–65 DEFINE_SITE_PRIORITY Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | The master group for which you are creating a site priority group. |
| name | Name of the site priority group that you are creating. |
| comment | This user comment is added to the DBA_REPPRIORITY view. |

### Exceptions

***Table 18–66 DEFINE_SITE_PRIORITY Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missingrepgroup | Specified master group does not exist. |
| duplicate prioritygroup | Specified site priority group exists in the master group. |
| notquiesced | Master group is not quiesced. |

# DO_DEFERRED_REPCAT_ADMIN Procedure

This procedure executes the local outstanding deferred administrative procedures for the specified master group at the current master site, or (with assistance from job queues) for all master sites.

DO_DEFERRED_REPCAT_ADMIN executes only those administrative requests submitted by the connected user who called DO_DEFERRED_REPCAT_ADMIN. Requests submitted by other users are ignored.

## Syntax

```
DBMS_REPCAT.DO_DEFERRED_REPCAT_ADMIN (
   gname          IN   VARCHAR2,
   all_sites      IN   BOOLEAN := FALSE);
```

## Parameters

*Table 18–67 DO_DEFERRED_REPCAT_ADMIN Procedure Parameters*

| Parameter | Description |
|---|---|
| gname | Name of the master group. |
| all_sites | If this is TRUE, then use a job to execute the local administrative procedures at each master site. |

## Exceptions

*Table 18–68 DO_DEFERRED_REPCAT_ADMIN Procedure Exceptions*

| Exception | Description |
|---|---|
| nonmaster | Invocation site is not a master site. |
| commfailure | At least one master site is not accessible and all_sites is TRUE. |

## DROP_COLUMN_GROUP Procedure

This procedure drops a column group. You must call this procedure from the master definition site.

> **See Also:** Chapter 6, "Configuring Conflict Resolution" and *Oracle Database Advanced Replication* for more information about conflict resolution methods

### Syntax

```
DBMS_REPCAT.DROP_COLUMN_GROUP (
   sname        IN   VARCHAR2,
   oname        IN   VARCHAR2,
   column_group IN   VARCHAR2);
```

### Parameters

***Table 18–69 DROP_COLUMN_GROUP Procedure Parameters***

| Parameter | Description |
|---|---|
| sname | Schema in which the replicated table is located. |
| oname | Name of the replicated table whose column group you are dropping. |
| column_group | Name of the column group that you want to drop. |

### Exceptions

***Table 18–70 DROP_COLUMN_GROUP Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| referenced | Specified column group is being used in conflict detection and resolution. |
| missingobject | Specified table does not exist. |
| missinggroup | Specified column group does not exist. |
| notquiesced | Master group to which the table belongs is not quiesced. |

# DROP_GROUPED_COLUMN Procedure

This procedure removes members from a column group. You must call this procedure from the master definition site.

> **See Also:** Chapter 6, "Configuring Conflict Resolution" and *Oracle Database Advanced Replication* for more information about conflict resolution methods

## Syntax

```
DBMS_REPCAT.DROP_GROUPED_COLUMN (
   sname                 IN   VARCHAR2,
   oname                 IN   VARCHAR2,
   column_group          IN   VARCHAR2,
   list_of_column_names  IN   VARCHAR2 | DBMS_REPCAT.VARCHAR2S);
```

## Parameters

***Table 18–71 DROP_GROUPED_COLUMN Procedure Parameters***

| Parameter | Description |
|---|---|
| sname | Schema in which the replicated table is located. |
| oname | Name of the replicated table in which the column group is located. The table can be the storage table of a nested table. |
| column_group | Name of the column group from which you are removing members. |
| list_of_column_names | Names of the columns that you are removing from the designated column group. This can either be a comma-delimited list or a PL/SQL associative array of column names. The PL/SQL associative array must be of type DBMS_REPCAT.VARCHAR2S.<br>You can specify column objects, but you cannot specify attributes of column objects.<br>If the table is an object, then you can specify SYS_NC_OID$ to add the object identifier column to the column group. This column tracks the object identifier of each row object.<br>If the table is a storage table of a nested table, then you can specify NESTED_TABLE_ID to add the column that tracks the identifier for each row of the nested table. |

## Exceptions

***Table 18–72 DROP_GROUPED_COLUMN Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missingobject | Specified table does not exist. |
| notquiesced | Master group that the table belongs to is not quiesced. |

## DROP_MASTER_REPGROUP Procedure

This procedure drops a master group from your current site. To drop the master group from all master sites, including the master definition site, you can call this procedure at the master definition site, and set `all_sites` to `TRUE`.

### Syntax

```
DBMS_REPCAT.DROP_MASTER_REPGROUP (
   gname             IN VARCHAR2,
   drop_contents     IN BOOLEAN    := FALSE,
   all_sites         IN BOOLEAN    := FALSE);
```

### Parameters

***Table 18–73 DROP_MASTER_REPGROUP Procedure Parameters***

| Parameter | Description |
|---|---|
| `gname` | Name of the master group that you want to drop from the current master site. |
| `drop_contents` | By default, when you drop the replication group at a master site, all of the objects remain in the database. They simply are no longer replicated. That is, the replicated objects in the replication group no longer send changes to, or receive changes from, other master sites. If you set this to `TRUE`, then any replicated objects in the master group are dropped from their associated schemas. |
| `all_sites` | If this is `TRUE` and if the invocation site is the master definition site, then the procedure synchronously multicasts the request to all masters. In this case, execution is immediate at the master definition site and might be deferred at all other master sites. |

### Exceptions

***Table 18–74 DROP_MASTER_REPGROUP Procedure Exceptions***

| Exception | Description |
|---|---|
| `nonmaster` | Invocation site is not a master site. |
| `nonmasterdef` | Invocation site is not the master definition site and `all_sites` is `TRUE`. |
| `commfailure` | At least one master site is not accessible and `all_sites` is `TRUE`. |
| `fullqueue` | Deferred remote procedure call (RPC) queue has entries for the master group. |
| `masternotremoved` | Master does not recognize the master definition site and `all_sites` is `TRUE`. |

# DROP_MASTER_REPOBJECT Procedure

This procedure drops a replicated object from a master group. You must call this procedure from the master definition site.

## Syntax

```
DBMS_REPCAT.DROP_MASTER_REPOBJECT (
   sname          IN   VARCHAR2,
   oname          IN   VARCHAR2,
   type           IN   VARCHAR2,
   drop_objects   IN   BOOLEAN     := FALSE);
```

## Parameters

***Table 18–75 DROP_MASTER_REPOBJECT Procedure Parameters***

| Parameter | Description |
|---|---|
| sname | Name of the schema in which the object is located. |
| oname | Name of the object that you want to remove from the master group. The object cannot be a storage table for a nested table. |
| type | Type of object that you want to drop. The following types are supported:<br>FUNCTION<br>INDEX<br>INDEXTYPE<br>OPERATOR<br>PACKAGE<br>PACKAGE BODY<br>PROCEDURE<br>SYNONYM<br>TABLE<br>TRIGGER<br>TYPE<br>TYPE BODY<br>VIEW |
| drop_objects | By default, the object remains in the schema, but is dropped from the master group. That is, any changes to the object are no longer replicated to other master and materialized view sites. To completely remove the object from the replication environment, set this parameter to TRUE. If the parameter is set to TRUE, the object is dropped from the database at each master site. |

## Exceptions

***Table 18–76 DROP_MASTER_REPOBJECT Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missingobject | Specified object does not exist. |
| typefailure | Specified type parameter is not supported. |
| commfailure | At least one master site is not accessible. |

## DROP_MVIEW_REPGROUP Procedure

This procedure drops a materialized view site from your replication environment. DROP_MVIEW_REPGROUP automatically calls UNREGISTER_MVIEW_REPGROUP at the master site or master materialized view site to unregister the materialized view, but ignores any errors that might have occurred during unregistration. If DROP_MVIEW_REPGROUP is unsuccessful, then connect to the master site or master materialized view site and run UNREGISTER_MVIEW_REPGROUP.

### Syntax

```
DBMS_REPCAT.DROP_MVIEW_REPGROUP (
   gname                 IN   VARCHAR2,
   drop_contents         IN   BOOLEAN   := FALSE,
   gowner                IN   VARCHAR2  := 'PUBLIC');
```

### Parameters

*Table 18–77 DROP_MVIEW_REPGROUP Procedure Parameters*

| Parameter | Description |
|---|---|
| gname | Name of the replication group that you want to drop from the current materialized view site. All objects generated to support replication, such as triggers and packages, are dropped. |
| drop_contents | By default, when you drop the replication group at a materialized view site, all of the objects remain in their associated schemas. They simply are no longer replicated. If you set this to TRUE, then any replicated objects in the replication group are dropped from their schemas. |
| gowner | Owner of the materialized view group. |

### Exceptions

*Table 18–78 DROP_MVIEW_REPGROUP Procedure Exceptions*

| Exception | Description |
|---|---|
| nonmview | Invocation site is not a materialized view site. |
| missingrepgroup | Specified replication group does not exist. |

# DROP_MVIEW_REPOBJECT Procedure

This procedure drops a replicated object from a materialized view site.

## Syntax

```
DBMS_REPCAT.DROP_MVIEW_REPOBJECT (
   sname          IN   VARCHAR2,
   oname          IN   VARCHAR2,
   type           IN   VARCHAR2,
   drop_objects   IN   BOOLEAN  := FALSE);
```

## Parameters

***Table 18–79 DROP_MVIEW_REPOBJECT Procedure Parameters***

| Parameter | Description |
|---|---|
| sname | Name of the schema in which the object is located. |
| oname | Name of the object that you want to drop from the replication group. |
| type | Type of the object that you want to drop. The following types are supported:<br>FUNCTION<br>INDEX<br>INDEXTYPE<br>OPERATOR<br>PACKAGE<br>PACKAGE BODY<br>PROCEDURE<br>SNAPSHOT<br>SYNONYM<br>TRIGGER<br>TYPE<br>TYPE BODY<br>VIEW<br>Use SNAPSHOT to drop a materialized view. |
| drop_objects | By default, the object remains in its associated schema, but is dropped from its associated replication group. To completely remove the object from its schema at the current materialized view site, set this parameter to TRUE. If the parameter is set to TRUE, the object is dropped from the database at the materialized view site. |

## Exceptions

***Table 18–80 DROP_MVIEW_REPOBJECT Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmview | Invocation site is not a materialized view site. |
| missingobject | Specified object does not exist. |
| typefailure | Specified type parameter is not supported. |

## DROP_PRIORITY Procedure

This procedure drops a member of a priority group by priority level. You must call this procedure from the master definition site.

> **See Also:** Chapter 6, "Configuring Conflict Resolution" and *Oracle Database Advanced Replication* for more information about conflict resolution methods

### Syntax

```
DBMS_REPCAT.DROP_PRIORITY(
      gname          IN   VARCHAR2,
      pgroup         IN   VARCHAR2,
      priority_num   IN   NUMBER);
```

### Parameters

***Table 18–81 DROP_PRIORITY Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Master group with which the priority group is associated. |
| pgroup | Name of the priority group containing the member that you want to drop. |
| priority_num | Priority level of the priority group member that you want to remove from the group. |

### Exceptions

***Table 18–82 DROP_PRIORITY Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missingrepgroup | Specified master group does not exist. |
| missingprioritygroup | Specified priority group does not exist. |
| notquiesced | Master group is not quiesced. |

# DROP_PRIORITY_GROUP Procedure

This procedure drops a priority group for a specified master group. You must call this procedure from the master definition site.

> **See Also:** Chapter 6, "Configuring Conflict Resolution" and *Oracle Database Advanced Replication* for more information about conflict resolution methods

## Syntax

```
DBMS_REPCAT.DROP_PRIORITY_GROUP (
   gname      IN   VARCHAR2,
   pgroup     IN   VARCHAR2);
```

## Parameters

***Table 18–83 DROP_PRIORITY_GROUP Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Master group with which the priority group is associated. |
| pgroup | Name of the priority group that you want to drop. |

## Exceptions

***Table 18–84 DROP_PRIORITY_GROUP Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missingrepgroup | Specified master group does not exist. |
| referenced | Specified priority group is being used in conflict resolution. |
| notquiesced | Specified master group is not quiesced. |

## DROP_PRIORITY_*datatype* Procedure

This procedure drops a member of a priority group by value. You must call this procedure from the master definition site. The procedure that you must call is determined by the data type of your `priority` column.

> **See Also:** Chapter 6, "Configuring Conflict Resolution" and *Oracle Database Advanced Replication* for more information about conflict resolution methods

### Syntax

```
DBMS_REPCAT.DROP_PRIORITY_datatype (
   gname    IN   VARCHAR2,
   pgroup   IN   VARCHAR2,
   value    IN   datatype);
```

where *datatype*:

```
{ NUMBER
| VARCHAR2
| CHAR
| DATE
| RAW
| NCHAR
| NVARCHAR2 }
```

### Parameters

***Table 18–85 DROP_PRIORITY_datatype Procedure Parameters***

| Parameter | Description |
|---|---|
| `gname` | Master group with which the priority group is associated. |
| `pgroup` | Name of the priority group containing the member that you want to drop. |
| `value` | Value of the priority group member that you want to remove from the group. |

### Exceptions

***Table 18–86 DROP_PRIORITY_datatype Procedure Exceptions***

| Exception | Description |
|---|---|
| `nonmasterdef` | Invocation site is not the master definition site. |
| `missingrepgroup` | Specified master group does not exist. |
| `missingprioritygroup` | Specified priority group does not exist. |
| `paramtype, typefailure` | Value has the incorrect data type for the priority group. |
| `notquiesced` | Specified master group is not quiesced. |

# DROP_SITE_PRIORITY Procedure

This procedure drops a site priority group for a specified master group. You must call this procedure from the master definition site.

> **See Also:** Chapter 6, "Configuring Conflict Resolution" and *Oracle Database Advanced Replication* for more information about conflict resolution methods

## Syntax

```
DBMS_REPCAT.DROP_SITE_PRIORITY (
   gname      IN   VARCHAR2,
   name       IN   VARCHAR2);
```

## Parameters

***Table 18–87 DROP_SITE_PRIORITY Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Master group with which the site priority group is associated. |
| name | Name of the site priority group that you want to drop. |

## Exceptions

***Table 18–88 DROP_SITE_PRIORITY Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missingrepgroup | Specified master group does not exist. |
| referenced | Specified site priority group is being used in conflict resolution. |
| notquiesced | Specified master group is not quiesced. |

## DROP_SITE_PRIORITY_SITE Procedure

This procedure drops a specified site, by name, from a site priority group. You must call this procedure from the master definition site.

> **See Also:** Chapter 6, "Configuring Conflict Resolution" and *Oracle Database Advanced Replication* for more information about conflict resolution methods

### Syntax

```
DBMS_REPCAT.DROP_SITE_PRIORITY_SITE (
   gname      IN   VARCHAR2,
   name       IN   VARCHAR2,
   site       IN   VARCHAR2);
```

### Parameters

***Table 18–89 DROP_SITE_PRIORITY_SITE Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Master group with which the site priority group is associated. |
| name | Name of the site priority group whose member you are dropping. |
| site | Global database name of the site you are removing from the group. |

### Exceptions

***Table 18–90 DROP_SITE_PRIORITY_SITE Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missingrepgroup | Specified master group does not exist. |
| missingpriority | Specified site priority group does not exist. |
| notquiesced | Specified master group is not quiesced. |

# DROP_*conflicttype*_RESOLUTION Procedure

This procedure drops an update, delete, or uniqueness conflict resolution routine. You must call these procedures from the master definition site. The procedure that you must call is determined by the type of conflict that the routine resolves.

## Conflict Resolution Routines

Table 18–91 shows the procedure name for each conflict resolution routine.

***Table 18–91 Conflict Resolution Routines***

| Routine | Procedure Name |
|---|---|
| update | DROP_UPDATE_RESOLUTION |
| uniqueness | DROP_UNIQUE_RESOLUTION |
| delete | DROP_DELETE_RESOLUTION |

## Syntax

```
DBMS_REPCAT.DROP_UPDATE_RESOLUTION (
   sname               IN   VARCHAR2,
   oname               IN   VARCHAR2,
   column_group        IN   VARCHAR2,
   sequence_no         IN   NUMBER);

DBMS_REPCAT.DROP_DELETE_RESOLUTION (
   sname               IN   VARCHAR2,
   oname               IN   VARCHAR2,
   sequence_no         IN   NUMBER);

DBMS_REPCAT.DROP_UNIQUE_RESOLUTION (
   sname               IN   VARCHAR2,
   oname               IN   VARCHAR2,
   constraint_name     IN   VARCHAR2,
   sequence_no         IN   NUMBER);
```

## Parameters

***Table 18–92 DROP_conflicttype_RESOLUTION Procedure Parameters***

| Parameter | Description |
|---|---|
| sname | Schema in which the table is located. |
| oname | Name of the table for which you want to drop a conflict resolution routine. |
| column_group | Name of the column group for which you want to drop an update conflict resolution routine. |
| constraint_name | Name of the unique constraint for which you want to drop a unique conflict resolution routine. |
| sequence_no | Sequence number assigned to the conflict resolution method that you want to drop. This number uniquely identifies the routine. |

## Exceptions

***Table 18–93 DROP_conflicttype_RESOLUTION Procedure Exceptions***

| Exception | Description |
|---|---|
| `nonmasterdef` | Invocation site is not the master definition site. |
| `missingobject` | Specified object does not exist as a table in the specified schema, or a conflict resolution routine with the specified sequence number is not registered. |
| `notquiesced` | Master group is not quiesced. |

# EXECUTE_DDL Procedure

This procedure supplies DDL that you want to have executed at some or all master sites. You can call this procedure only from the master definition site.

## Syntax

```
DBMS_REPCAT.EXECUTE_DDL (
   gname           IN   VARCHAR2,
   { master_list   IN   VARCHAR2     := NULL,
   | master_table  IN   DBMS_UTILITY.DBLINK_ARRAY,}
   DDL_TEXT        IN   VARCHAR2);
```

**Note:** This procedure is overloaded. The `master_list` and `master_table` parameters are mutually exclusive.

## Parameters

***Table 18–94 EXECUTE_DDL Procedure Parameters***

| Parameter | Description |
|---|---|
| `gname` | Name of the master group. |
| `master_list` | A comma-delimited list of master sites at which you want to execute the supplied DDL. Do not put any spaces between site names. The default value, `NULL`, indicates that the DDL should be executed at all sites, including the master definition site. |
| `master_table` | A table that lists the master sites where you want to execute the supplied DDL. The first master should be at position 1, the second at position 2, and so on. |
| `ddl_text` | The DDL that you want to execute at each of the specified master sites. If the DDL is supplied without specifying a schema, then the default schema is the replication administrator's schema. Be sure to specify the schema if it is other than the replication administrator's schema. |

## Exceptions

***Table 18–95 EXECUTE_DDL Procedure Exceptions***

| Exception | Description |
|---|---|
| `nonmasterdef` | Invocation site is not the master definition site. |
| `nonmaster` | At least one site is not a master site. |
| `ddlfailure` | DDL at the master definition site did not succeed. |
| `commfailure` | At least one master site is not accessible. |

## GENERATE_MVIEW_SUPPORT Procedure

This procedure activates triggers and generate packages needed to support the replication of updatable materialized views or procedural replication.You must call this procedure from the materialized view site.

> **Note:** CREATE_MVIEW_REPOBJECT automatically generates materialized view support for updatable materialized views.

### Syntax

```
DBMS_REPCAT.GENERATE_MVIEW_SUPPORT (
   sname                    IN VARCHAR2,
   oname                    IN VARCHAR2,
   type                     IN VARCHAR2,
   gen_objs_owner           IN VARCHAR2 := '',
   min_communication        IN BOOLEAN  := TRUE,
   generate_80_compatible   IN BOOLEAN  := TRUE);
```

### Parameters

***Table 18–96 GENERATE_MVIEW_SUPPORT Procedure Parameters***

| Parameter | Description |
|---|---|
| sname | Schema in which the object is located. |
| oname | The name of the object for which you are generating support. |
| type | Type of the object. The types supported are MATERIALIZED VIEW, PACKAGE, and PACKAGE BODY. |
| gen_objs_owner | For objects of type PACKAGE or PACKAGE BODY, the schema in which the generated object should be created. If NULL, the objects are created in SNAME. |
| min_communication | If TRUE, then the update trigger sends the new value of a column only if the update statement modifies the column. The update trigger sends the old value of the column only if it is a key column or a column in a modified column group. |
| generate_80_compatible | Set to FALSE because interoperability is not supported between Oracle8*i* databases and Oracle 11*g* or later databases. |

### Exceptions

***Table 18–97 GENERATE_MVIEW_SUPPORT Procedure Exceptions***

| Exceptions | Descriptions |
|---|---|
| nonmview | Invocation site is not a materialized view site. |
| missingobject | Specified object does not exist as a materialized view in the replicated schema waiting for row/column-level replication information or as a package (body) waiting for wrapper generation. |
| typefailure | Specified type parameter is not supported. |
| missingschema | Specified owner of generated objects does not exist. |

*Table 18–97 (Cont.) GENERATE_MVIEW_SUPPORT Procedure Exceptions*

| Exceptions | Descriptions |
|---|---|
| `missingremoteobject` | Object at master site or master materialized view site has not yet generated replication support. |
| `commfailure` | Master site or master materialized view site is not accessible. |

## GENERATE_REPLICATION_SUPPORT Procedure

This procedure generates the triggers and packages needed to support replication for a specified object. You must call this procedure from the master definition site.

### Syntax

```
DBMS_REPCAT.GENERATE_REPLICATION_SUPPORT (
  sname                     IN    VARCHAR2,
  oname                     IN    VARCHAR2,
  type                      IN    VARCHAR2,
  package_prefix            IN    VARCHAR2   := NULL,
  procedure_prefix          IN    VARCHAR2   := NULL,
  distributed               IN    BOOLEAN    := TRUE,
  gen_objs_owner            IN    VARCHAR2   := NULL,
  min_communication         IN    BOOLEAN    := TRUE,
  generate_80_compatible    IN    BOOLEAN    := TRUE);
```

### Parameters

***Table 18–98 GENERATE_REPLICATION_SUPPORT Procedure Parameters***

| Parameter | Description |
|---|---|
| sname | Schema in which the object is located. |
| oname | Name of the object for which you are generating replication support. |
| type | Type of the object. The types supported are: TABLE, PACKAGE, and PACKAGE BODY. |
| package_prefix | For objects of type PACKAGE or PACKAGE BODY this value is prepended to the generated wrapper package name. The default is DEFER_. |
| procedure_prefix | For objects of type PACKAGE or PACKAGE BODY, this value is prepended to the generated wrapper procedure names. By default, no prefix is assigned. |
| distributed | This must be set to TRUE. |
| gen_objs_owner | For objects of type PACKAGE or PACKAGE BODY, the schema in which the generated object should be created. If NULL, the objects are created in sname. |
| min_communication | This parameter is obsolete. Use the default value (TRUE). |
| generate_80_compatible | Set to FALSE because interoperability is not supported between Oracle8*i* databases and Oracle 11*g* or later databases. |

## Exceptions

*Table 18–99 GENERATE_REPLICATION_SUPPORT Procedure Exceptions*

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missingobject | Specified object does not exist as a table in the specified schema waiting for row-level replication information or as a package (body) waiting for wrapper generation. |
| typefailure | Specified type parameter is not supported. |
| notquiesced | Replication group has not been quiesced. |
| commfailure | At least one master site is not accessible. |
| missingschema | Schema does not exist. |
| duplicateobject | Object exists. |

## MAKE_COLUMN_GROUP Procedure

This procedure creates a new column group with one or more members. You must call this procedure from the master definition site.

> **See Also:** Chapter 6, "Configuring Conflict Resolution" and *Oracle Database Advanced Replication* for more information about conflict resolution methods

### Syntax

```
DBMS_REPCAT.MAKE_COLUMN_GROUP (
   sname                 IN   VARCHAR2,
   oname                 IN   VARCHAR2,
   column_group          IN   VARCHAR2,
   list_of_column_names  IN   VARCHAR2 | DBMS_REPCAT.VARCHAR2S);
```

### Parameters

***Table 18–100 MAKE_COLUMN_GROUP Procedure Parameters***

| Parameter | Description |
|---|---|
| sname | Schema in which the replicated table is located. |
| oname | Name of the replicated table for which you are creating a new column group. The table can be the storage table of a nested table. |
| column_group | Name that you want assigned to the column group that you are creating. |
| list_of_column_names | Names of the columns that you are grouping. This can either be a comma-delimited list or a PL/SQL associative array of column names. The PL/SQL associative array must be of type DBMS_REPCAT.VARCHAR2S. Use the single value '*' to create a column group that contains all of the columns in your table.<br>You can specify column objects, but you cannot specify attributes of column objects.<br>If the table is an object table, then you can specify SYS_NC_OID$ to add the object identifier column to the column group. This column tracks the object identifier of each row object.<br>If the table is the storage table of a nested table, then you can specify NESTED_TABLE_ID to add the column that tracks the identifier for each row of the nested table. |

### Exceptions

***Table 18–101 MAKE_COLUMN_GROUP Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| duplicategroup | Specified column group exists for the table. |
| missingobject | Specified table does not exist. |
| missingcolumn | Specified column does not exist in the designated table. |
| duplicatecolumn | Specified column is already a member of another column group. |
| notquiesced | Master group is not quiesced. |

# PREPARE_INSTANTIATED_MASTER Procedure

This procedure enables the propagation of deferred transactions from other prepared new master sites and existing master sites to the invocation master site. This procedure also enables the propagation of deferred transactions from the invocation master site to the other prepared new master sites and existing master sites.

If you performed a full database export/import or a change-based recovery, then the new master site includes all of the deferred transactions that were in the deferred transactions queue at the master definition site. Because these deferred transactions should not exist at the new master site, this procedure deletes all transactions in the deferred transactions queue and error queue if full database export/import or change-based recovery was used.

For object-level export/import, ensure that all the requests in the DBA_REPCATLOG data dictionary view for the extended groups have been processed without error before running this procedure.

**Caution:**

- Do not invoke this procedure until instantiation (export/import or change-based recovery) for the new master site is complete.
- Do not allow any data manipulation language (DML) statements directly on the objects in the extended master group in the new master site until execution of this procedure returns successfully. These DML statements might not be replicated.
- Do not use the DBMS_DEFER package to create deferred transactions until execution of this procedure returns successfully. These deferred transactions might not be replicated.

**Note:** To use change-based recovery, the existing master site and the new master site must be running under the same operating system, although the release of the operating system can differ.

## Syntax

```
DBMS_REPCAT.PREPARE_INSTANTIATED_MASTER (
   extension_id             IN       RAW);
```

## Parameters

***Table 18–102 PREPARE_INSTANTIATED_MASTER Procedure Parameters***

| Parameter | Description |
|---|---|
| extension_id | The identifier for the current pending request to add master databases to a master group without quiesce. You can find the extension_id by querying the DBA_REPSITES_NEW and DBA_REPEXTENSIONS data dictionary views. |

## Exceptions

*Table 18–103 PREPARE_INSTANTIATED_MASTER Procedure Exceptions*

| Exception | Description |
|---|---|
| `typefailure` | The parameter value specified for one of the parameters is not appropriate. |
| `dbnotcompatible` | Feature is incompatible with database version. All databases must be at 9.2.0 or higher compatibility level. |

# PURGE_MASTER_LOG Procedure

This procedure removes local messages in the DBA_REPCATLOG view associated with a specified identification number, source, or master group.

To purge all of the administrative requests from a particular source, specify NULL for the id parameter. To purge all administrative requests from all sources, specify NULL for both the id parameter and the source parameter.

## Syntax

```
DBMS_REPCAT.PURGE_MASTER_LOG (
   id      IN   BINARY_INTEGER,
   source  IN   VARCHAR2,
   gname   IN   VARCHAR2);
```

## Parameters

***Table 18–104 PURGE_MASTER_LOG Procedure Parameters***

| Parameter | Description |
|---|---|
| id | Identification number of the request, as it appears in the DBA_REPCATLOG view. |
| source | Master site from which the request originated. |
| gname | Name of the master group for which the request was made. |

## Exceptions

***Table 18–105 PURGE_MASTER_LOG Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmaster | gname is not NULL, and the invocation site is not a master site. |

## PURGE_STATISTICS Procedure

This procedure removes information from the DBA_REPRESOLUTION_STATISTICS view.

### Syntax

```
DBMS_REPCAT.PURGE_STATISTICS (
   sname      IN   VARCHAR2,
   oname      IN   VARCHAR2,
   start_date IN   DATE,
   end_date   IN   DATE);
```

### Parameters

*Table 18–106 PURGE_STATISTICS Procedure Parameters*

| Parameter | Description |
|---|---|
| sname | Name of the schema in which the replicated table is located. |
| oname | Name of the table whose conflict resolution statistics you want to purge. |
| start_date/end_date | Range of dates for which you want to purge statistics. If start_date is NULL, then purge all statistics up to the end_date. If end_date is NULL, then purge all statistics after the start_date. |

### Exceptions

*Table 18–107 PURGE_STATISTICS Procedure Exceptions*

| Exception | Description |
|---|---|
| missingschema | Specified schema does not exist. |
| missingobject | Specified table does not exist. |
| statnotreg | Table not registered to collect statistics. |

# REFRESH_MVIEW_REPGROUP Procedure

This procedure refreshes a materialized view group with the most recent data from its associated master site or master materialized view site.

## Syntax

```
DBMS_REPCAT.REFRESH_MVIEW_REPGROUP (
   gname                  IN   VARCHAR2,
   drop_missing_contents  IN   BOOLEAN    := FALSE,
   refresh_mviews         IN   BOOLEAN    := FALSE,
   refresh_other_objects  IN   BOOLEAN    := FALSE,
   gowner                 IN   VARCHAR2   := 'PUBLIC');
```

## Parameters

***Table 18–108 REFRESH_MVIEW_REPGROUP Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Name of the replication group. |
| drop_missing_contents | If an object was dropped from the replication group at the master site or master materialized view site, then it is not automatically dropped from the schema at the materialized view site. It is simply no longer replicated. That is, changes to this object are no longer sent to its associated master site or master materialized view site. Materialized views can continue to be refreshed from their associated master tables or master materialized views. However, any changes to an updatable materialized view are lost. When an object is dropped from the replication group, you can choose to have it dropped from the schema entirely by setting this parameter to TRUE. |
| refresh_mviews | Set to TRUE to refresh the contents of the materialized views in the replication group. |
| refresh_other_objects | Set this to TRUE to refresh the contents of the non materialized view objects in the replication group. Non materialized view objects can include the following:<br>■ Tables<br>■ Views<br>■ Indexes<br>■ PL/SQL packages and package bodies<br>■ PL/SQL procedures and functions<br>■ Triggers<br>■ Synonyms |
| gowner | Owner of the materialized view group. |

## Exceptions

*Table 18–109 REFRESH_MVIEW_REPGROUP Procedure Exceptions*

| Exception | Description |
|---|---|
| nonmview | Invocation site is not a materialized view site. |
| nonmaster | Master is no longer a master site or master materialized view site. |
| commfailure | Master site or master materialized view site is not accessible. |
| missingrepgroup | Replication group name not specified. |

# REGISTER_MVIEW_REPGROUP Procedure

This procedure facilitates the administration of materialized views at their respective master sites or master materialized view sites by inserting or modifying a materialized view group in DBA_REGISTERED_MVIEW_GROUPS.

## Syntax

```
DBMS_REPCAT.REGISTER_MVIEW_REPGROUP (
   gname            IN   VARCHAR2,
   mviewsite        IN   VARCHAR2,
   comment          IN   VARCHAR2  := NULL,
   rep_type         IN   NUMBER    := reg_unknown,
   fname            IN   VARCHAR2  := NULL,
   gowner           IN   VARCHAR2  := 'PUBLIC');
```

## Parameters

***Table 18–110 REGISTER_MVIEW_REPGROUP Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Name of the materialized view group to be registered. |
| mviewsite | Global name of the materialized view site. |
| comment | Comment for the materialized view site or update for an existing comment. |
| rep_type | Version of the materialized view group. Valid constants that can be assigned include the following:<br>■ dbms_repcat.reg_unknown (the default)<br>■ dbms_repcat.reg_v7_group<br>■ dbms_repcat.reg_v8_group |
| fname | This parameter is for internal use only.<br>**Note:** Do not set this parameter unless directed to do so by Oracle Support Services. |
| gowner | Owner of the materialized view group. |

## Exceptions

***Table 18–111 REGISTER_MVIEW_REPGROUP Procedure Exceptions***

| Exception | Description |
|---|---|
| failregmviewrepgroup | Registration of materialized view group failed. |
| missingrepgroup | Replication group name not specified. |
| nullsitename | A materialized view site was not specified. |
| nonmaster | Procedure must be executed at the materialized view's master site or master materialized view site. |
| duplicaterepgroup | Replication group exists. |

## REGISTER_STATISTICS Procedure

This procedure collects information about the successful resolution of update, delete, and uniqueness conflicts for a table.

### Syntax

```
DBMS_REPCAT.REGISTER_STATISTICS (
   sname IN   VARCHAR2,
   oname IN   VARCHAR2);
```

### Parameters

*Table 18–112 REGISTER_STATISTICS Procedure Parameters*

| Parameter | Description |
|---|---|
| sname | Name of the schema in which the table is located. |
| oname | Name of the table for which you want to gather conflict resolution statistics. |

### Exceptions

*Table 18–113 REGISTER_STATISTICS Procedure Exceptions*

| Exception | Description |
|---|---|
| missingschema | Specified schema does not exist. |
| missingobject | Specified table does not exist. |

# RELOCATE_MASTERDEF Procedure

This procedure changes your master definition site to another master site in your replication environment.

It is not necessary for either the old or new master definition site to be available when you call RELOCATE_MASTERDEF. In a planned reconfiguration, invoke RELOCATE_MASTERDEF with notify_masters set to TRUE and include_old_masterdef set to TRUE.

## Syntax

```
DBMS_REPCAT.RELOCATE_MASTERDEF (
   gname                  IN   VARCHAR2,
   old_masterdef          IN   VARCHAR2,
   new_masterdef          IN   VARCHAR2,
   notify_masters         IN   BOOLEAN    := TRUE,
   include_old_masterdef  IN   BOOLEAN    := TRUE,
   require_flavor_change  IN   BOOLEAN    := FALSE);
```

## Parameters

***Table 18–114 RELOCATE_MASTERDEF Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Name of the replication group whose master definition you want to relocate. |
| old_masterdef | Fully qualified database name of the current master definition site. |
| new_masterdef | Fully qualified database name of the existing master site that you want to make the new master definition site. |
| notify_masters | If this is TRUE, then the procedure synchronously multicasts the change to all masters (including old_masterdef only if include_old_masterdef is TRUE). If any master does not make the change, then roll back the changes at all masters.<br><br>If just the master definition site fails, then you should invoke RELOCATE_MASTERDEF with notify_masters set to TRUE and include_old_masterdef set to FALSE. If several master sites and the master definition site fail, then the administrator should invoke RELOCATE_MASTERDEF at each operational master with notify_masters set to FALSE. |
| include_old_masterdef | If notify_masters is TRUE and if include_old_masterdef is also TRUE, then the old master definition site is also notified of the change. |
| require_flavor_change | This parameter is for internal use only.<br><br>**Note:** Do not set this parameter unless directed to do so by Oracle Support Services. |

## Exceptions

*Table 18–115 RELOCATE_MASTERDEF Procedure Exceptions*

| Exception | Description |
|---|---|
| nonmaster | new_masterdef is not a master site or the invocation site is not a master site. |
| nonmasterdef | old_masterdef is not the master definition site. |
| commfailure | At least one master site is not accessible and notify_masters is TRUE. |

# REMOVE_MASTER_DATABASES Procedure

This procedure removes one or more master databases from a replication environment. This procedure regenerates the triggers and their associated packages at the remaining master sites. You must call this procedure from the master definition site.

## Syntax

```
DBMS_REPCAT.REMOVE_MASTER_DATABASES (
   gname           IN   VARCHAR2,
   master_list     IN   VARCHAR2 |
   master_table    IN   DBMS_UTILITY.DBLINK_ARRAY);
```

**Note:** This procedure is overloaded. The `master_list` and `master_table` parameters are mutually exclusive.

## Parameters

***Table 18–116 REMOVE_MASTER_DATABASES Procedure Parameters***

| Parameter | Description |
|---|---|
| `gname` | Name of the replication group associated with the replication environment. This prevents confusion if a master database is involved in multiple replication environments. |
| `master_list` | A comma-delimited list of fully qualified master database names that you want to remove from the replication environment. There must be no spaces between names in the list. |
| `master_table` | In place of a list, you can specify the database names in a PL/SQL associative array of type `DBMS_UTILITY.DBLINK_ARRAY`. |

## Exceptions

***Table 18–117 REMOVE_MASTER_DATABASES Procedure Exceptions***

| Exception | Description |
|---|---|
| `nonmasterdef` | Invocation site is not the master definition site. |
| `nonmaster` | At least one of the specified databases is not a master site. |
| `reconfigerror` | One of the specified databases is the master definition site. |
| `commfailure` | At least one remaining master site is not accessible. |

## RENAME_SHADOW_COLUMN_GROUP Procedure

This procedure renames the shadow column group of a replicated table to make it a named column group. The replicated table's master group does not need to be quiesced to run this procedure.

### Syntax

```
DBMS_REPCAT.RENAME_SHADOW_COLUMN_GROUP (
   sname                IN VARCHAR2,
   oname                IN VARCHAR2,
   new_col_group_name   IN VARCHAR2)
```

### Parameters

*Table 18–118 RENAME_SHADOW_COLUMN_GROUP Procedure Parameters*

| Parameter | Description |
|---|---|
| sname | Schema in which the replicated table is located. |
| oname | Name of the replicated table. |
| new_col_group_name | Name of the new column group. The columns currently in the shadow group are placed in a column group with the name you specify. |

### Exceptions

*Table 18–119 RENAME_SHADOW_COLUMN_GROUP Procedure Exceptions*

| Exception | Description |
|---|---|
| missmview | The specified schema does not exist. |
| nonmasterdef | Invocation site is not the master definition site. |
| missingobject | The specified object does not exist. |
| duplicategroup | The column group that was specified for creation exists. |

# REPCAT_IMPORT_CHECK Procedure

This procedure ensures that the objects in the master group have the appropriate object identifiers and status values after you perform an export/import of a replicated object or an object used by Advanced Replication.

## Syntax

```
DBMS_REPCAT.REPCAT_IMPORT_CHECK (
   gname      IN   VARCHAR2,
   master     IN   BOOLEAN,
   gowner     IN   VARCHAR2  := 'PUBLIC');
```

## Parameters

***Table 18–120 REPCAT_IMPORT_CHECK Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Name of the master group. If you omit both parameters, then the procedure checks all master groups at your current site. |
| master | Set this to TRUE if you are checking a master site and FALSE if you are checking a materialized view site. |
| gowner | Owner of the master group. |

## Exceptions

***Table 18–121 REPCAT_IMPORT_CHECK Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmaster | master is TRUE and either the database is not a master site for the replication group or the database is not the expected database. |
| nonmview | master is FALSE and the database is not a materialized view site for the replication group. |
| missingobject | A valid replicated object in the replication group does not exist. |
| missingrepgroup | The specified replicated replication group does not exist. |
| missingschema | The specified schema does not exist. |

## RESUME_MASTER_ACTIVITY Procedure

This procedure resumes normal replication activity after quiescing a replication environment.

### Syntax

```
DBMS_REPCAT.RESUME_MASTER_ACTIVITY (
   gname       IN  VARCHAR2,
   override    IN  BOOLEAN := FALSE);
```

### Parameters

***Table 18–122 RESUME_MASTER_ACTIVITY Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Name of the master group. |
| override | If this is TRUE, then it ignores any pending administrative requests and restores normal replication activity at each master as quickly as possible. This should be considered only in emergency situations.<br>If this is FALSE, then it restores normal replication activity at each master only when there is no pending administrative request for gname at that master. |

### Exceptions

***Table 18–123 RESUME_MASTER_ACTIVITY Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| notquiesced | Master group is not quiescing or quiesced. |
| commfailure | At least one master site is not accessible. |
| notallgenerated | Generate replication support before resuming replication activity. |

# RESUME_PROPAGATION_TO_MDEF Procedure

During the process of adding new master sites to a master group without quiesce, this procedure indicates that export is effectively finished and propagation to the master definition site for both extended and unaffected replication groups existing at master sites can be enabled. Run this procedure after the export required to add new master sites to a master group is complete.

> **See Also:** "Adding New Master Sites" on page 7-2 for more information about adding master sites to a master group

## Syntax

```
DBMS_REPCAT.RESUME_PROPAGATION_TO_MDEF (
   extension_id    IN    RAW);
```

## Parameters

***Table 18–124 RESUME_PROPAGATION_TO_MDEF Procedure Parameters***

| Parameter | Description |
|---|---|
| extension_id | The identifier for the current pending request to add master databases to a master group without quiesce. You can find the extension_id by querying the DBA_REPSITES_NEW and DBA_REPEXTENSIONS data dictionary views. |

## Exceptions

***Table 18–125 RESUME_PROPAGATION_TO_MDEF Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| extstinapp | Extension status is inappropriate. The extension status should be EXPORTING when you run this procedure. To check the extension status, query the DBA_REPEXTENSIONS data dictionary view. |
| dbnotcompatible | Feature is incompatible with database version. All databases must be at 9.2.0 or higher compatibility level. |

## SEND_OLD_VALUES Procedure

You have the option of sending old column values during propagation of deferred transactions for each nonkey column of a replicated table when rows are updated or deleted in the table. When `min_communication` is set to `TRUE`, the default is the following:

- For a deleted row, to send old values for all columns
- For an updated row, to send old values for key columns and the modified columns in a column group

You can change this behavior at all master sites and materialized view sites by invoking `DBMS_REPCAT.SEND_OLD_VALUES` at the master definition site. Then, generate replication support at all master sites and at each materialized view site.

When you use user-defined types, you can specify the leaf attributes of a column object, or an entire column object. For example, if a column object named `cust_address` has `street_address` as an attribute, then you can specify `cust_address.street_address` for the `column_list` parameter or as part of the `column_table` parameter, or you can specify only `cust_address`.

### Syntax

```
DBMS_REPCAT.SEND_OLD_VALUES(
   sname           IN  VARCHAR2,
   oname           IN  VARCHAR2,
   { column_list   IN  VARCHAR2,
   | column_table  IN  DBMS_UTILITY.VARCHAR2S | DBMS_UTILITY.LNAME_ARRAY,}
   operation       IN  VARCHAR2 := 'UPDATE',
   send            IN  BOOLEAN  := TRUE );
```

> **Note:** This procedure is overloaded. The `column_list` and `column_table` parameters are mutually exclusive.

### Parameters

***Table 18–126 SEND_OLD_VALUES Procedure Parameters***

| Parameter | Description |
|---|---|
| `sname` | Schema in which the table is located. |
| `oname` | Name of the replicated table. The table can be the storage table of a nested table. |
| `column_list` | A comma-delimited list of the columns in the table. There must be no spaces between entries. |
| `column_table` | Instead of a list, you can use a PL/SQL associative array of type `DBMS_REPCAT.VARCHAR2S` or `DBMS_UTILITY.LNAME_ARRAY` to contain the column names. The first column name should be at position 1, the second at position 2, and so on.<br>Use `DBMS_UTILITY.LNAME_ARRAY` if any column name is greater than or equal to 30 bytes, which might occur when you specify the attributes of column objects. |
| `operation` | Possible values are: `update`, `delete`, or the asterisk wildcard `'*'`, which means update and delete. |

*Table 18–126 (Cont.) SEND_OLD_VALUES Procedure Parameters*

| Parameter | Description |
|---|---|
| send | If TRUE, then the old values of the specified columns are sent. If FALSE, then the old values of the specified columns are not sent. Unspecified columns and unspecified operations are not affected.<br>The specified change takes effect at the master definition site as soon as min_communication is TRUE for the table. The change takes effect at a master site or at a materialized view site the next time replication support is generated at that site with min_communication TRUE. |

**Note:** The operation parameter enables you to specify whether to transmit old values for nonkey columns when rows are deleted or updated. If you do not send the old value, then Oracle sends a NULL for the old value and assumes the old value is equal to the current value of the column at the target side when the update or delete is applied.

See *Oracle Database Advanced Replication* for information about reduced data propagation using the SEND_OLD_VALUES procedure before changing the default behavior of Oracle.

## Exceptions

*Table 18–127 SEND_OLD_VALUES Procedure Exceptions*

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missingobject | Specified object does not exist as a table in the specified schema waiting for row-level replication information. |
| missingcolumn | At least one column is not in the table. |
| notquiesced | Master group has not been quiesced. |
| typefailure | An illegal operation is specified. |
| keysendcomp | A specified column is a key column in a table. |
| dbnotcompatible | Feature is incompatible with database version. Typically, this exception arises when you are trying to send the attributes of column objects. In this case, all databases must be at 9.2.0 or higher compatibility level. |

## SET_COLUMNS Procedure

This procedure enables you to use an alternate column or group of columns, instead of the primary key, to determine which columns of a table to compare when using row-level replication. You must call this procedure from the master definition site.

When you use column objects, if an attribute of a column object can be used as a primary key or part of a primary key, then the attribute can be part of an alternate key column. For example, if a column object named `cust_address` has `street_address` as a `VARCHAR2` attribute, then you can specify `cust_address.street_address` for the `column_list` parameter or as part of the `column_table` parameter. However, the entire column object, `cust_address`, cannot be specified.

For the storage table of a nested table column, this procedure accepts the `NESTED_TABLE_ID` as an alternate key column.

When you use object tables, you cannot specify alternate key columns. If the object identifier (OID) is system-generated for an object table, then Oracle uses the OID column in the object table as the key for the object table. If the OID is user-defined for an object table, then Oracle uses the primary key in the object table as the key.

The following types of columns cannot be alternate key columns:

- LOB or LOB attribute of a column object
- Collection or collection attribute of a column object
- `REF`
- An entire column object

> **See Also:** The *constraint_clause* in *Oracle Database SQL Language Reference* for more information about restrictions on primary key columns

### Syntax

```
DBMS_REPCAT.SET_COLUMNS (
   sname           IN    VARCHAR2,
   oname           IN    VARCHAR2,
   { column_list   IN    VARCHAR2
   | column_table  IN    DBMS_UTILITY.NAME_ARRAY | DBMS_UTILITY.LNAME_ARRAY } );
```

> **Note:** This procedure is overloaded. The `column_list` and `column_table` parameters are mutually exclusive.

## Parameters

*Table 18–128 SET_COLUMNS Procedure Parameters*

| Parameter | Description |
|---|---|
| sname | Schema in which the table is located. |
| oname | Name of the table. |
| column_list | A comma-delimited list of the columns in the table that you want to use as a primary key. There must be no spaces between entries. |
| column_table | Instead of a list, you can use a PL/SQL associative array of type DBMS_UTILITY.NAME_ARRAY or DBMS_UTILITY.LNAME_ARRAY to contain the column names. The first column name should be at position 1, the second at position 2, and so on.<br>Use DBMS_UTILITY.LNAME_ARRAY if any column name is greater than or equal to 30 bytes, which might occur when you specify the attributes of column objects. |

## Exceptions

*Table 18–129 SET_COLUMNS Procedure Exceptions*

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| missingobject | Specified object does not exist as a table in the specified schema waiting for row-level replication information. |
| missingcolumn | At least one column is not in the table. |
| notquiesced | Replication group is not quiescing or quiesced. |

## SPECIFY_NEW_MASTERS Procedure

This procedure specifies the master sites you intend to add to an existing replication group without quiescing the group. This procedure must be run at the master definition site of the specified master group.

If necessary, this procedure creates an `extension_id` that tracks the process of adding new master sites to a master group. You use this `extension_id` in the other procedures that you run at various stages in the process. You can view information about the `extension_id` in the `DBA_REPSITES_NEW` and `DBA_REPEXTENSIONS` data dictionary views.

This procedure adds the new master sites to the `DBA_REPSITES_NEW` data dictionary view for the specified replication group. This procedure can be run any number of times for a given replication group. If it is run more than once, then it replaces any masters in the local `DBA_REPSITES_NEW` data dictionary view for the specified replication group with the masters specified in the `master_list/master_table` parameters.

You must run this procedure before you run the `ADD_NEW_MASTERS` procedure. No new master sites are added to the master group until you run the `ADD_NEW_MASTERS` procedure.

**See Also:**

- "ADD_NEW_MASTERS Procedure" on page 18-8
- "Adding New Master Sites" on page 7-2 for more information about adding master sites to a master group

### Syntax

```
DBMS_REPCAT.SPECIFY_NEW_MASTERS (
   gname           IN   VARCHAR2,
   { master_list   IN   VARCHAR2
   | master_table  IN   DBMS_UTILITY.DBLINK_ARRAY});
```

**Note:** This procedure is overloaded. The `master_list` and `master_table` parameters are mutually exclusive.

## Parameters

*Table 18–130 SPECIFY_NEW_MASTERS Procedure Parameters*

| Parameter | Description |
|---|---|
| gname | Master group to which you are adding new master sites. |
| master_list | A comma-delimited list of new master sites that you want to add to the master group. List only the new master sites, not the existing master sites. Do not put any spaces between site names.<br>If master_list is NULL, all master sites for the given replication group are removed from the DBA_REPSITES_NEW data dictionary view. Specify NULL to indicate that the master group is not being extended. |
| master_table | A table that lists the new master sites that you want to add to the master group. In the table, list only the new master sites, not the existing master sites. The first master site should be at position 1, the second at position 2, and so on.<br>If the table is empty, then all master sites for the specified replication group are removed from the DBA_REPSITES_NEW data dictionary view. Use an empty table to indicate that the master group is not being extended. |

## Exceptions

*Table 18–131 SPECIFY_NEW_MASTERS Procedure Exceptions*

| Exception | Description |
|---|---|
| duplicaterepgroup | A master site that you are attempting to add is already part of the master group. |
| nonmasterdef | Invocation site is not the master definition site. |
| propmodenotallowed | Synchronous propagation mode not allowed for this operation. Only asynchronous propagation mode is allowed. |
| extstinapp | Extension request with status not allowed. There must either be no extension_id for the master group or the extension_id status must be READY. You can view the status for each extension_id at a master site in the DBA_REPEXTENSIONS data dictionary view. |
| dbnotcompatible | Feature is incompatible with database version. All databases must be at 9.2.0 or higher compatibility level. |
| notsamecq | Master groups do not have the same connection qualifier. |

## STREAMS_MIGRATION Procedure

Generates a migration script that migrates an Advanced Replication environment to a Streams environment. Specifically, this procedure generates a script that sets up a Streams environment for the specified replication groups. The generated script can be customized and run at each master site to perform the migration.

> **See Also:** *Oracle Streams Replication Administrator's Guide* for detailed information about migrating from Advanced Replication to Streams

### Syntax

```
DBMS_REPCAT.STREAMS_MIGRATION (
     gnames              IN   DBMS_UTILITY.NAME_ARRAY,
     file_location       IN   VARCHAR2,
     filename            IN   VARCHAR2);
```

### Parameters

***Table 18–132 STREAMS_MIGRATION Procedure Parameters***

| Parameter | Description |
|---|---|
| gnames | List of replication groups to migrate to Streams. The replication groups listed must all contain the same master sites. An error is raised if the replication groups have different masters. |
| file_location | Directory location of the migration script. The specified location must be a directory object that is accessible to PL/SQL. You can use the SQL statement CREATE DIRECTORY to create a directory object.<br>**See Also:** *Oracle Database SQL Language Reference* for more information about the CREATE DIRECTORY statement |
| filename | Name of the migration script. |

# SUSPEND_MASTER_ACTIVITY Procedure

This procedure suspends replication activity for a master group. You use this procedure to quiesce the master group. You must call this procedure from the master definition site.

## Syntax

```
DBMS_REPCAT.SUSPEND_MASTER_ACTIVITY (
   gname   IN   VARCHAR2);
```

## Parameters

***Table 18–133 SUSPEND_MASTER_ACTIVITY Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Name of the master group for which you want to suspend activity. |

## Exceptions

***Table 18–134 SUSPEND_MASTER_ACTIVITY Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmasterdef | Invocation site is not the master definition site. |
| notnormal | Master group is not in normal operation. |
| commfailure | At least one master site is not accessible. |

## SWITCH_MVIEW_MASTER Procedure

This procedure changes the master site of a materialized view group to another master site. This procedure does a full refresh of the affected materialized views and regenerates the triggers and their associated packages as needed. This procedure does not push the queue to the old master site before changing master sites.

---

**Note:** You cannot switch the master of materialized views that are based on other materialized views (level 2 and greater materialized views). Such a materialized view must be dropped and re-created to base it on a different master.

---

**See Also:** "GENERATE_MVIEW_SUPPORT Procedure" on page 18-69

### Syntax

```
DBMS_REPCAT.SWITCH_MVIEW_MASTER (
   gname          IN   VARCHAR2,
   master         IN   VARCHAR2,
   gowner         IN   VARCHAR2  := 'PUBLIC');
```

### Parameters

***Table 18–135 SWITCH_MVIEW_MASTER Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Name of the materialized view group for which you want to change the master site. |
| master | Fully qualified database name of the new master site to use for the materialized view group. |
| gowner | Owner of the materialized view group. |

### Exceptions

***Table 18–136 SWITCH_MVIEW_MASTER Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmview | Invocation site is not a materialized view site. |
| nonmaster | Specified database is not a master site. |
| commfailure | Specified database is not accessible. |
| missingrepgroup | Materialized view group does not exist. |
| qrytoolong | Materialized view definition query is greater 32 KB. |
| alreadymastered | At the local site, there is another materialized view group with the same group name mastered at the old master site. |

# UNDO_ADD_NEW_MASTERS_REQUEST Procedure

This procedure undoes all of the changes made by the SPECIFY_NEW_MASTERS and ADD_NEW_MASTERS procedures for a specified extension_id.

This procedure is executed at one master site, which can be the master definition site, and it only affects that master site. If you run this procedure at one master site affected by the request, you must run it at all new and existing master sites affected by the request. You can query the DBA_REPSITES_NEW data dictionary view to see the new master sites affected by the extension_id. This data dictionary view also lists the replication group name, and you must run this procedure at all existing master sites in the replication group.

---

**Caution:** This procedure is not normally called. Use this procedure only if the adding new masters without quiesce operation cannot proceed at one or more master sites. Run this procedure after you have already run the SPECIFY_NEW_MASTERS and ADD_NEW_MASTERS procedures, but *before* you have run the RESUME_PROPAGATION_TO_MDEF and PREPARE_INSTANTIATED_MASTER procedures.

*Do not* run this procedure after you have run either RESUME_PROPAGATION_TO_MDEF or PREPARE_INSTANTIATED_MASTER for a particular extension_id.

---

**See Also:**

- "SPECIFY_NEW_MASTERS Procedure" on page 18-93
- "ADD_NEW_MASTERS Procedure" on page 18-8
- "RESUME_PROPAGATION_TO_MDEF Procedure" on page 18-88
- "PREPARE_INSTANTIATED_MASTER Procedure" on page 18-74

## Syntax

```
DBMS_REPCAT.UNDO_ADD_NEW_MASTERS_REQUEST (
   extension_id    IN  RAW,
   drop_contents   IN  BOOLEAN := TRUE);
```

## Parameters

***Table 18–137 UNDO_ADD_NEW_MASTERS_REQUEST Procedure Parameters***

| Parameter | Description |
|---|---|
| extension_id | The identifier for the current pending request to add master databases to a master group without quiesce. You can find the extension_id by querying the DBA_REPSITES_NEW and DBA_REPEXTENSIONS data dictionary views. |
| drop_contents | Specify TRUE, the default, to drop the contents of objects in new replication groups being extended at the local site. Specify FALSE to retain the contents. |

## Exceptions

*Table 18–138 UNDO_ADD_NEW_MASTERS_REQUEST Procedure Exceptions*

| Exception | Description |
|---|---|
| dbnotcompatible | Feature is incompatible with database version. All databases must be at 9.2.0 or higher compatibility level. |
| typefail | A parameter value that you specified is not appropriate. |

# UNREGISTER_MVIEW_REPGROUP Procedure

This procedure facilitates the administration of materialized views at their respective master sites or master materialized view sites by deleting a materialized view group from DBA_REGISTERED_MVIEW_GROUPS. Run this procedure at the master site or master materialized view site.

## Syntax

```
DBMS_REPCAT.UNREGISTER_MVIEW_REPGROUP (
   gname      IN   VARCHAR2,
   mviewsite  IN   VARCHAR2,
   gowner     IN   VARCHAR2  := 'PUBLIC');
```

## Parameters

***Table 18–139 UNREGISTER_MVIEW_REPGROUP Procedure Parameters***

| Parameter | Description |
|---|---|
| gname | Name of the materialized view group to be unregistered. |
| mviewsite | Global name of the materialized view site. |
| gowner | Owner of the materialized view group. |

## VALIDATE Function

This function validates the correctness of key conditions of a multimaster replication environment.

### Syntax

```
DBMS_REPCAT.VALIDATE (
   gname                IN  VARCHAR2,
   check_genflags       IN  BOOLEAN := FALSE,
   check_valid_objs     IN  BOOLEAN := FALSE,
   check_links_sched    IN  BOOLEAN := FALSE,
   check_links          IN  BOOLEAN := FALSE,
   error_table          OUT DBMS_REPCAT.VALIDATE_ERR_TABLE)
  RETURN BINARY_INTEGER;

DBMS_REPCAT.VALIDATE (
   gname                IN  VARCHAR2,
   check_genflags       IN  BOOLEAN := FALSE,
   check_valid_objs     IN  BOOLEAN := FALSE,
   check_links_sched    IN  BOOLEAN := FALSE,
   check_links          IN  BOOLEAN := FALSE,
   error_msg_table      OUT DBMS_UTILITY.UNCL_ARRAY,
   error_num_table      OUT DBMS_UTILITY.NUMBER_ARRAY )
  RETURN BINARY_INTEGER;
```

**Note:** This function is overloaded. The return value of `VALIDATE` is the number of errors found. The function's `OUT` parameter returns any errors that are found. In the first interface function shown under "Syntax" on page 18-101, the `error_table` consists of an array of records. Each record has a `VARCHAR2` and a `NUMBER` in it. The string field contains the error message, and the number field contains the Oracle error number.

The second interface function shown under "Syntax" on page 18-101 is similar except that there are two `OUT` arrays: a `VARCHAR2` array with the error messages and a `NUMBER` array with the error numbers.

### Parameters

***Table 18–140 VALIDATE Function Parameters***

| Parameter | Description |
|---|---|
| `gname` | Name of the master group to validate. |
| `check_genflags` | Check whether all the objects in the group are generated. This must be done at the master definition site only. |
| `check_valid_objs` | Check that the underlying objects for objects in the group valid. This must be done at the master definition site only. The master definition site goes to all other sites and checks that the underlying objects are valid. The validity of the objects is checked within the schema of the connected user. |
| `check_links_sched` | Check whether the links are scheduled for execution. This should be invoked at each master site. |

*Table 18–140 (Cont.) VALIDATE Function Parameters*

| Parameter | Description |
|---|---|
| check_links | Check whether the connected user (repadmin), as well as the propagator, have correct links for replication to work properly. Checks that the links exist in the database and are accessible. This should be invoked at each master site. |
| error_table | Returns the messages and numbers of all errors that are found. |
| error_msg_table | Returns the messages of all errors that are found. |
| error_num_table | Returns the numbers of all errors that are found. |

## Exceptions

*Table 18–141 VALIDATE Function Exceptions*

| Exception | Description |
|---|---|
| missingdblink | Database link does not exist in the schema of the replication propagator or has not been scheduled. Ensure that the database link exists in the database, is accessible, and is scheduled for execution. |
| dblinkmismatch | Database link name at the local node does not match the global name of the database that the link accesses. Ensure that the GLOBAL_NAMES initialization parameter is set to TRUE and the link name matches the global name. |
| dblinkuidmismatch | User name of the replication administration user at the local node and the user name at the node corresponding to the database link are different. Advanced Replication expects the two users to be the same. Ensure that the user identification of the replication administration user at the local node and the user identification at the node corresponding to the database link are the same. |
| objectnotgenerated | Object has not been generated at other master sites or is still being generated. Ensure that the object is generated by calling GENERATE_REPLICATION_SUPPORT and DO_DEFERRED_REPCAT_ADMIN for the object at the master definition site. |

## Usage Notes

The return value of VALIDATE is the number of errors found. The function's OUT parameter returns any errors that are found. In the first interface function, the error_table consists of an array of records. Each record has a VARCHAR2 and a NUMBER in it. The string field contains the error message and the number field contains the Oracle error number.

The second interface is similar except that there are two OUT arrays. A VARCHAR2 array with the error messages and a NUMBER array with the error numbers.

## WAIT_MASTER_LOG Procedure

This procedure determines whether changes that were asynchronously propagated to a master site have been applied.

### Syntax

```
DBMS_REPCAT.WAIT_MASTER_LOG (
   gname          IN    VARCHAR2,
   record_count   IN    NATURAL,
   timeout        IN    NATURAL,
   true_count     OUT   NATURAL);
```

### Parameters

*Table 18–142 WAIT_MASTER_LOG Procedure Parameters*

| Parameter | Description |
|---|---|
| gname | Name of the master group. |
| record_count | Procedure returns whenever the number of incomplete activities is at or below this threshold. |
| timeout | Maximum number of seconds to wait before the procedure returns. |
| true_count (out parameter) | Returns the number of incomplete activities. |

### Exceptions

*Table 18–143 WAIT_MASTER_LOG Procedure Exceptions*

| Exception | Description |
|---|---|
| nonmaster | Invocation site is not a master site. |

# 19

# DBMS_REPCAT_INSTANTIATE

DBMS_REPCAT_INSTANTIATE package instantiates deployment templates.

This chapter contains this topic:

- Summary of DBMS_REPCAT_INSTANTIATE Subprograms

# Summary of DBMS_REPCAT_INSTANTIATE Subprograms

*Table 19–1 DBMS_REPCAT_INSTANTIATE Package Subprograms*

| Subprogram | Description |
|---|---|
| DROP_SITE_INSTANTIATION Procedure on page 19-3 | Public procedure that removes the target site from the DBA_REPCAT_TEMPLATE_SITES view. |
| INSTANTIATE_OFFLINE Function on page 19-4 | Public function that generates a script at the master site that is used to create the materialized view environment at the remote materialized view site while offline. |
| INSTANTIATE_ONLINE Function on page 19-6 | Public function that generates a script at the master site that is used to create the materialized view environment at the remote materialized view site while online. |

## DROP_SITE_INSTANTIATION Procedure

This procedure drops a template instantiation at a target site. This procedure removes all related metadata at the master site and disables the specified site from refreshing its materialized views. You must execute this procedure as the user who originally instantiated the template. To see who instantiated the template, query the ALL_REPCAT_TEMPLATE_SITES view.

### Syntax

```
DBMS_REPCAT_INSTANTIATE.DROP_SITE_INSTANTIATION(
     refresh_template_name  IN   VARCHAR2,
     site_name              IN   VARCHAR2);
```

### Parameters

*Table 19–2 DROP_SITE_INSTANTIATION Procedure Parameters*

| Parameter | Description |
|---|---|
| refresh_template_name | The name of the deployment template to be dropped. |
| site_name | Identifies the master site where you want to drop the specified template instantiation. |

# INSTANTIATE_OFFLINE Function

This function generates a file at the master site that is used to create the materialized view environment at the remote materialized view site while offline. This generated file is an offline instantiation file and should be used at remote materialized view sites that are not able to remain connected to the master site for an extended amount of time.

This is an ideal solution when the remote materialized view site is a laptop. Use the packaging interface in the Advanced Replication interface in Oracle Enterprise Manager to package the generated file and data into a single file that can be posted on an FTP site or loaded to a CD-ROM, floppy disk, and so on. You can also transfer the file using the `DBMS_FILE_TRANSFER` package.

The script generated by this function is stored in the `USER_REPCAT_TEMP_OUTPUT` temporary view and is used by several Oracle tools, including the Advanced Replication interface in Oracle Enterprise Manager, during the distribution of deployment templates. The number returned by this function is used to retrieve the appropriate information from the `USER_REPCAT_TEMP_OUTPUT` view.

The user who executes this public function becomes the "registered" user of the instantiated template at the specified site.

---

**Note:** This function is used in performing an offline instantiation of a deployment template.

This function should not be confused with the procedures in the DBMS_OFFLINE_OG package (used for performing an offline instantiation of a master table). See the documentation for this package for more information about their usage.

---

**See Also:**

- "Packaging a Deployment Template for Instantiation" on page 4-9
- *Oracle Database Advanced Replication*
- The Advanced Replication interface's online Help in Oracle Enterprise Manager

## Syntax

```
DBMS_REPCAT_INSTANTIATE.INSTANTIATE_OFFLINE(
     refresh_template_name   IN   VARCHAR2,
     site_name               IN   VARCHAR2,
     runtime_parm_id         IN   NUMBER    := -1e-130,
     next_date               IN   DATE      := SYSDATE,
     interval                IN   VARCHAR2  := 'SYSDATE + 1',
     use_default_gowner      IN   BOOLEAN   := TRUE)
     return NUMBER;
```

## Parameters

*Table 19–3 INSTANTIATE_OFFLINE Function Parameters*

| Parameter | Description |
|---|---|
| refresh_template_name | The name of the deployment template to be instantiated. |
| site_name | The name of the remote site that is instantiating the deployment template. |
| runtime_parm_id | If you have defined run-time parameter values using the INSERT_RUNTIME_PARMS procedure, specify the identification used when creating the run-time parameters (the identification was retrieved by using the GET_RUNTIME_PARM_ID function). |
| next_date | The next refresh date value to be used when creating the refresh group. |
| interval | The refresh interval to be used when creating the refresh group. |
| use_default_gowner | If TRUE, then any materialized view groups created are owned by the default user PUBLIC. If FALSE, then any materialized view groups created are owned by the user performing the instantiation. |

## Exceptions

*Table 19–4 INSTANTIATE_OFFLINE Function Exceptions*

| Exception | Description |
|---|---|
| miss_refresh_template | The deployment template name specified is invalid or does not exist. |
| dupl_template_site | The deployment template has already been instantiated at the materialized view site. A deployment template can be instantiated only once at a particular materialized view site. |
| not_authorized | The user attempting to instantiate the deployment template is not authorized to do so. |

## Returns

*Table 19–5 INSTANTIATE_OFFLINE Function Returns*

| Return Value | Description |
|---|---|
| <*system-generated number*> | Specifies the generated system number for the output_id when you select from the USER_REPCAT_TEMP_OUTPUT view to retrieve the generated instantiation script. |

# INSTANTIATE_ONLINE Function

This function generates a script at the master site that is used to create the materialized view environment at the remote materialized view site while online. This generated script should be used at remote materialized view sites that are able to remain connected to the master site for an extended amount of time, as the instantiation process at the remote materialized view site might be lengthy (depending on the amount of data that is populated to the new materialized views).

The script generated by this function is stored in the `USER_REPCAT_TEMP_OUTPUT` temporary view and is used by several Oracle tools, including the Advanced Replication interface in Oracle Enterprise Manager, during the distribution of deployment templates. The number returned by this function is used to retrieve the appropriate information from the `USER_REPCAT_TEMP_OUTPUT` view.

The user who executes this public function becomes the "registered" user of the instantiated template at the specified site.

**See Also:**

- "Packaging a Deployment Template for Instantiation" on page 4-9
- *Oracle Database Advanced Replication*
- The Advanced Replication interface's online Help in Oracle Enterprise Manager

## Syntax

```
DBMS_REPCAT_INSTANTIATE.INSTANTIATE_ONLINE(
     refresh_template_name   IN   VARCHAR2,
     site_name               IN   VARCHAR2,
     runtime_parm_id         IN   NUMBER    := -1e-130,
     next_date               IN   DATE      := SYSDATE,
     interval                IN   VARCHAR2  := 'SYSDATE + 1',
     use_default_gowner      IN   BOOLEAN   := TRUE)
     return NUMBER;
```

## Parameters

*Table 19–6 INSTANTIATE_ONLINE Function Parameters*

| Parameter | Description |
|---|---|
| refresh_template_name | The name of the deployment template to be instantiated. |
| site_name | The name of the remote site that is instantiating the deployment template. |
| runtime_parm_id | If you have defined run-time parameter values using the INSERT_RUNTIME_PARMS procedure, specify the identification used when creating the run-time parameters (the identification was retrieved by using the GET_RUNTIME_PARM_ID function). |
| next_date | Specifies the next refresh date value to be used when creating the refresh group. |
| interval | Specifies the refresh interval to be used when creating the refresh group. |
| use_default_gowner | If TRUE, then any materialized view groups created are owned by the default user PUBLIC. If FALSE, then any materialized view groups created are owned by the user performing the instantiation. |

## Exceptions

*Table 19–7 INSTANTIATE_ONLINE Function Exceptions*

| Exception | Description |
|---|---|
| miss_refresh_template | The deployment template name specified is invalid or does not exist. |
| dupl_template_site | The deployment template has already been instantiated at the materialized view site. A deployment template can be instantiated only once at a particular materialized view site. |
| not_authorized | The user attempting to instantiate the deployment template is not authorized to do so. |

## Returns

*Table 19–8 INSTANTIATE_ONLINE Function Returns*

| Return Value | Description |
|---|---|
| <*system-generated number*> | Specifies the generated system number for the output_id when you select from the USER_REPCAT_TEMP_OUTPUT view to retrieve the generated instantiation script. |

# 20

# DBMS_REPCAT_ADMIN

DBMS_REPCAT_ADMIN enables you to create users with the privileges needed by the symmetric replication facility.

This chapter contains this topic:


- Summary of DBMS_REPCAT_ADMIN Subprograms

# Summary of DBMS_REPCAT_ADMIN Subprograms

*Table 20–1 DBMS_REPCAT_ADMIN Package Subprograms*

| Subprogram | Description |
|---|---|
| "GRANT_ADMIN_ANY_SCHEMA Procedure" on page 20-3 | Grants the necessary privileges to the replication administrator to administer any replication group at the current site. |
| "GRANT_ADMIN_SCHEMA Procedure" on page 20-4 | Grants the necessary privileges to the replication administrator to administer a schema at the current site. |
| "REGISTER_USER_REPGROUP Procedure" on page 5 | Assigns proxy materialized view administrator or receiver privileges at the master site or master materialized view site for use with remote sites. |
| "REVOKE_ADMIN_ANY_SCHEMA Procedure" on page 20-7 | Revokes the privileges and roles from the replication administrator that were granted by `GRANT_ADMIN_ANY_SCHEMA`. |
| "REVOKE_ADMIN_SCHEMA Procedure" on page 20-8 | Revokes the privileges and roles from the replication administrator that were granted by `GRANT_ADMIN_SCHEMA`. |
| "UNREGISTER_USER_REPGROUP Procedure" on page 20-9 | Revokes the privileges and roles from the proxy materialized view administrator or receiver that were granted by the `REGISTER_USER_REPGROUP` procedure. |

## GRANT_ADMIN_ANY_SCHEMA Procedure

This procedure grants the necessary privileges to the replication administrator to administer any replication groups at the current site.

### Syntax

```
DBMS_REPCAT_ADMIN.GRANT_ADMIN_ANY_SCHEMA (
   username  IN   VARCHAR2);
```

### Parameters

*Table 20–2 GRANT_ADMIN_ANY_SCHEMA Procedure Parameters*

| Parameter | Description |
|---|---|
| username | Name of the replication administrator to whom you want to grant the necessary privileges and roles to administer any replication groups at the current site. |

### Exceptions

*Table 20–3 GRANT_ADMIN_ANY_SCHEMA Procedure Exceptions*

| Exception | Description |
|---|---|
| ORA-01917 | User does not exist. |

# GRANT_ADMIN_SCHEMA Procedure

This procedure grants the necessary privileges to the replication administrator to administer a schema at the current site. This procedure is most useful if your replication group does not span schemas.

The privileges granted by this procedure are more limited than the privileges granted by GRANT_ADMIN_ANY_SCHEMA. However, a replication administrator who is granted privileges with GRANT_ADMIN_SCHEMA still can perform certain administrative activities on replication groups owned by other replication administrators. For example, a replication administrator who is granted privileges with GRANT_ADMIN_SCHEMA can drop replication groups and replication objects owned by other replication administrators.

---

**Note:** To restrict different users to different replicated groups, you can write a wrapper package on top of the DBMS_REPCAT package and grant EXECUTE privilege on the new package, but not on the DBMS_REPCAT package, to each user. The new package performs security checks. For example, the new package can dictate that hr can administer the hr_rg replication group, but no other replication group, and that hr only can administer objects in the hr schema. If the security checks are passed, then the new package calls a subprogram in the DBMS_REPCAT package. If the security checks are not passed, then the new package could log the failure, commit, and raise an exception.

---

**See Also:** "GRANT_ADMIN_ANY_SCHEMA Procedure" on page 20-3

## Syntax

```
DBMS_REPCAT_ADMIN.GRANT_ADMIN_SCHEMA (
   username IN VARCHAR2);
```

## Parameters

***Table 20–4 GRANT_ADMIN_SCHEMA Procedure Parameters***

| Parameter | Description |
|---|---|
| username | Name of the replication administrator. This user is then granted the necessary privileges and roles to administer the schema of the same name within a replication group at the current site. |

## Exceptions

***Table 20–5 GRANT_ADMIN_SCHEMA Procedure Exceptions***

| Exception | Description |
|---|---|
| ORA-01917 | User does not exist. |

## REGISTER_USER_REPGROUP Procedure

This procedure assigns proxy materialized view administrator or receiver privileges at the master site or master materialized view site for use with remote sites. This procedure grants only the necessary privileges to the proxy materialized view administrator or receiver. It does not grant the powerful privileges granted by the GRANT_ADMIN_SCHEMA or GRANT_ADMIN_ANY_SCHEMA procedures.

> **See Also:** Appendix A, "Security Options" for more information about trusted versus untrusted security models

### Syntax

```
DBMS_REPCAT_ADMIN.REGISTER_USER_REPGROUP (
   username         IN   VARCHAR2,
   privilege_type   IN   VARCHAR2,
   {list_of_gnames  IN   VARCHAR2 |
   table_of_gnames  IN   DBMS_UTILITY.NAME_ARRAY)};
```

> **Note:** This procedure is overloaded. The list_of_gnames and table_of_gnames parameters are mutually exclusive.

### Parameters

***Table 20–6 REGISTER_USER_REPGROUP Procedure Parameters***

| Parameter | Description |
|---|---|
| username | Name of the user to whom you are giving either proxy materialized view administrator or receiver privileges. |
| privilege_type | Specifies the privilege type you are assigning. Use the following values for to define your privilege_type:<br>■ receiver for receiver privileges<br>■ proxy_snapadmin for proxy materialized view administration privileges |
| list_of_gnames | Comma-delimited list of replication groups you want a user registered for receiver privileges. There must be no spaces between entries in the list. If you set list_of_gnames to NULL, then the user is registered for all replication groups, even replication groups that are not yet known when this procedure is called. You must use named notation to set list_of_gnames to NULL. An invalid replication group in the list causes registration to fail for the entire list. |
| table_of_gnames | PL/SQL associative array of replication groups you want a user registered for receiver privileges. The PL/SQL associative array must be of type DBMS_UTILITY.NAME_ARRAY. This table is 1-based (the positions start at 1 and increment by 1). Use the single value NULL to register the user for all replication groups. An invalid replication group in the table causes registration to fail for the entire table. |

## Exceptions

*Table 20–7 REGISTER_USER_REPGROUP Procedure Exceptions*

| Exception | Description |
|---|---|
| `nonmaster` | Specified replication group does not exist or the invocation database is not a master site or master materialized view site. |
| `ORA-01917` | User does not exist. |
| `typefailure` | Incorrect privilege type was specified. |

## REVOKE_ADMIN_ANY_SCHEMA Procedure

This procedure revokes the privileges and roles from the replication administrator that were granted by GRANT_ADMIN_ANY_SCHEMA.

> **Note:** Identical privileges and roles that were granted independently of GRANT_ADMIN_ANY_SCHEMA are also revoked.

### Syntax

```
DBMS_REPCAT_ADMIN.REVOKE_ADMIN_ANY_SCHEMA (
   username IN VARCHAR2);
```

### Parameters

***Table 20–8 REVOKE_ADMIN_ANY_SCHEMA Procedure Parameters***

| Parameter | Description |
|---|---|
| username | Name of the replication administrator whose privileges you want to revoke. |

### Exceptions

***Table 20–9 REVOKE_ADMIN_ANY_SCHEMA Procedure Exceptions***

| Exception | Description |
|---|---|
| ORA-01917 | User does not exist. |

# REVOKE_ADMIN_SCHEMA Procedure

This procedure revokes the privileges and roles from the replication administrator that were granted by GRANT_ADMIN_SCHEMA.

> **Note:** Identical privileges and roles that were granted independently of GRANT_ADMIN_SCHEMA are also revoked.

## Syntax

```
DBMS_REPCAT_ADMIN.REVOKE_ADMIN_SCHEMA (
   username IN VARCHAR2);
```

## Parameters

***Table 20–10 REVOKE_ADMIN_SCHEMA Procedure Parameters***

| Parameter | Description |
|---|---|
| username | Name of the replication administrator whose privileges you want to revoke. |

## Exceptions

***Table 20–11 REVOKE_ADMIN_SCHEMA Procedure Exceptions***

| Exception | Description |
|---|---|
| ORA-01917 | User does not exist. |

## UNREGISTER_USER_REPGROUP Procedure

This procedure revokes the privileges and roles from the proxy materialized view administrator or receiver that were granted by the REGISTER_USER_REPGROUP procedure.

### Syntax

```
DBMS_REPCAT_ADMIN.UNREGISTER_USER_REPGROUP (
   username         IN   VARCHAR2,
   privilege_type   IN   VARCHAR2,
   {list_of_gnames  IN   VARCHAR2 |
   table_of_gnames  IN   DBMS_UTILITY.NAME_ARRAY)};
```

**Note:** This procedure is overloaded. The list_of_gnames and table_of_gnames parameters are mutually exclusive.

### Parameters

***Table 20–12 UNREGISTER_USER_REPGROUP Procedure Parameters***

| Parameter | Description |
|---|---|
| username | Name of the user you are unregistering. |
| privilege_type | Specifies the privilege type you are revoking. Use the following values for to define your privilege_type:<br>■ receiver for receiver privileges<br>■ proxy_snapadmin for proxy materialized view administration privileges |
| list_of_gnames | Comma-delimited list of replication groups you want a user unregistered for receiver privileges. There must be no spaces between entries in the list. If you set list_of_gnames to NULL, then the user is unregistered for all replication groups registered. You must use named notation to set list_of_gnames to NULL. An invalid replication group in the list causes unregistration to fail for the entire list. |
| table_of_gnames | PL/SQL associative array of replication groups you want a user unregistered for receiver privileges. The PL/SQL associative array must be of type DBMS_UTILITY.NAME_ARRAY. This table is 1-based (the positions start at 1 and increment by 1). Use the single value NULL to unregister the user for all replication groups registered. An invalid replication group in the table causes unregistration to fail for the entire table. |

### Exceptions

***Table 20–13 UNREGISTER_USER_REPGROUP Procedure Exceptions***

| Exception | Description |
|---|---|
| nonmaster | Specified replication group does not exist or the invocation database is not a master site or master materialized view site. |
| ORA-01917 | User does not exist. |
| typefailure | Incorrect privilege type was specified. |

# 21

# DBMS_REPCAT_RGT

DBMS_REPCAT_RGT controls the maintenance and definition of refresh group templates.

This chapter contains this topic:


- Summary of DBMS_REPCAT_RGT Subprograms

# Summary of DBMS_REPCAT_RGT Subprograms

*Table 21–1 DBMS_REPCAT_RGT Package Subprograms*

| Subprogram | Description |
|---|---|
| "ALTER_REFRESH_TEMPLATE Procedure" on page 21-4 | Allows the DBA to alter existing deployment templates. |
| "ALTER_TEMPLATE_OBJECT Procedure" on page 21-6 | Alters objects that have been added to a specified deployment template. |
| "ALTER_TEMPLATE_PARM Procedure" on page 21-8 | Allows the DBA to alter the parameters for a specific deployment template. |
| "ALTER_USER_ AUTHORIZATION Procedure" on page 21-10 | Alters the contents of the DBA_REPCAT_USER_ AUTHORIZATIONS view. |
| "ALTER_USER_PARM_VALUE Procedure" on page 21-11 | Changes existing parameter values that have been defined for a specific user. |
| "COMPARE_TEMPLATES Function" on page 21-13 | Allows the DBA to compare the contents of two deployment templates. |
| "COPY_TEMPLATE Function" on page 21-14 | Allows the DBA to copy a deployment template. |
| "CREATE_OBJECT_FROM_ EXISTING Function" on page 21-16 | Creates a template object definition from existing database objects and adds it to a target deployment template. |
| "CREATE_REFRESH_ TEMPLATE Function" on page 21-18 | Creates the deployment template, which allows the DBA to define the template name, private/public status, and target refresh group. |
| "CREATE_TEMPLATE_OBJECT Function" on page 21-20 | Adds object definitions to a target deployment template container. |
| "CREATE_TEMPLATE_PARM Function" on page 21-22 | Creates parameters for a specific deployment template to allow custom data sets to be created at the remote materialized view site. |
| "CREATE_USER_ AUTHORIZATION Function" on page 21-24 | Authorizes specific users to instantiate private deployment templates. |
| "CREATE_USER_PARM_ VALUE Function" on page 21-25 | Predefines deployment template parameter values for specific users. |
| "DELETE_RUNTIME_PARMS Procedure" on page 21-27 | Deletes a run-time parameter value that you defined using the INSERT_RUNTIME_PARMS procedure. |
| "DROP_ALL_OBJECTS Procedure" on page 21-28 | Allows the DBA to drop all objects or specific object types from a deployment template. |
| "DROP_ALL_TEMPLATE_ PARMS Procedure" on page 21-29 | Allows the DBA to drop template parameters for a specified deployment template. |
| "DROP_ALL_TEMPLATE_ SITES Procedure" on page 21-30 | Removes all entries from the DBA_REPCAT_TEMPLATE_ SITES view. |
| "DROP_ALL_TEMPLATES Procedure" on page 21-31 | Removes all deployment templates at the site where the procedure is called. |
| "DROP_ALL_USER_ AUTHORIZATIONS Procedure" on page 21-32 | Allows the DBA to drop all user authorizations for a specified deployment template. |

*Table 21–1 (Cont.) DBMS_REPCAT_RGT Package Subprograms*

| Subprogram | Description |
|---|---|
| "DROP_ALL_USER_PARM_ VALUES Procedure" on page 21-33 | Drops user parameter values for a specific deployment template. |
| "DROP_REFRESH_TEMPLATE Procedure" on page 21-34 | Drops a deployment template. |
| "DROP_SITE_INSTANTIATION Procedure" on page 21-35 | Removes the target site from the `DBA_REPCAT_TEMPLATE_ SITES` view. |
| "DROP_TEMPLATE_OBJECT Procedure" on page 21-36 | Removes a template object from a specific deployment template. |
| "DROP_TEMPLATE_PARM Procedure" on page 21-37 | Removes an existing template parameter from the `DBA_ REPCAT_TEMPLATE_PARMS` view. |
| DROP_USER_ AUTHORIZATION Procedure on page 21-38 | Removes a user authorization entry from the `DBA_ REPCAT_USER_AUTHORIZATIONS` view. |
| "DROP_USER_PARM_VALUE Procedure" on page 21-39 | Removes a predefined user parameter value for a specific deployment template. |
| "GET_RUNTIME_PARM_ID Function" on page 21-40 | Retrieves an identification to be used when defining a run-time parameter value. |
| "INSERT_RUNTIME_PARMS Procedure" on page 21-41 | Defines run-time parameter values before instantiating a template. |
| "INSTANTIATE_OFFLINE Function" on page 21-43 | Generates a script at the master site that is used to create the materialized view environment at the remote materialized view site while offline. |
| "INSTANTIATE_ONLINE Function" on page 21-45 | Generates a script at the master site that is used to create the materialized view environment at the remote materialized view site while online. |
| "LOCK_TEMPLATE_ EXCLUSIVE Procedure" on page 47 | Prevents users from reading or instantiating the template when a deployment template is being updated or modified. |
| "LOCK_TEMPLATE_SHARED Procedure" on page 21-48 | Makes a specified deployment template read-only. |

# ALTER_REFRESH_TEMPLATE Procedure

This procedure allows the DBA to alter existing deployment templates. Alterations can include defining a new deployment template name, a new refresh group, or a new owner and changing the public/private status.

## Syntax

```
DBMS_REPCAT_RGT.ALTER_REFRESH_TEMPLATE (
   refresh_template_name      IN   VARCHAR2,
   new_owner                  IN   VARCHAR2 := '-',
   new_refresh_group_name     IN   VARCHAR2 := '-',
   new_refresh_template_name  IN   VARCHAR2 := '-',
   new_template_comment       IN   VARCHAR2 := '-',
   new_public_template        IN   VARCHAR2 := '-',
   new_last_modified          IN   DATE := to_date('1', 'J'),
   new_modified_by            IN   NUMBER := -1e-130);
```

## Parameters

***Table 21–2 ALTER_REFRESH_TEMPLATE Procedure Parameters***

| Parameter | Description |
|---|---|
| refresh_template_name | The name of the deployment template that you want to alter. |
| new_owner | The name of the new deployment template owner. Do not specify a value to keep the current owner. |
| new_refresh_group_name | If necessary, use this parameter to specify a new refresh group name to which the template objects will be added. Do not specify a value to keep the current refresh group. |
| new_refresh_template_name | Use this parameter to specify a new deployment template name. Do not specify a value to keep the current deployment template name. |
| new_template_comment | New deployment template comments. Do not specify a value to keep the current template comment. |
| new_public_template | Determines whether the deployment template is public or private. Only acceptable values are 'Y' and 'N' ('Y' = public and 'N' = private). Do not specify a value to keep the current value. |
| new_last_modified | Contains the date of the last modification made to this deployment template. If a value is not specified, then the current date is automatically used. |
| new_modified_by | Contains the name of the user who last modified this deployment template. If a value is not specified, then the current user is automatically used. |

## Exceptions

***Table 21–3 ALTER_REFRESH_TEMPLATE Procedure Exceptions***

| Exception | Description |
|---|---|
| miss_refresh_template | Deployment template name specified is invalid or does not exist. |

*Table 21–3 (Cont.) ALTER_REFRESH_TEMPLATE Procedure Exceptions*

| Exception | Description |
|---|---|
| `bad_public_template` | The `public_template` parameter is specified incorrectly. The `public_template` parameter must be specified as a `'Y'` for a public template or an `'N'` for a private template. |
| `dupl_refresh_template` | A template with the specified name exists. See the ALL_REPCAT_REFRESH_TEMPLATES view. |

# ALTER_TEMPLATE_OBJECT Procedure

This procedure alters objects that have been added to a specified deployment template. The most common changes are altering the object DDL and assigning the object to a different deployment template.

Changes made to the template are reflected only at new sites instantiating the deployment template. Remote sites that have already instantiated the template must reinstantiate the deployment template to apply the changes.

## Syntax

```
DBMS_REPCAT_RGT.ALTER_TEMPLATE_OBJECT (
   refresh_template_name       IN   VARCHAR2,
   object_name                 IN   VARCHAR2,
   object_type                 IN   VARCHAR2,
   new_refresh_template_name   IN   VARCHAR2 := '-',
   new_object_name             IN   VARCHAR2 := '-',
   new_object_type             IN   VARCHAR2 := '-',
   new_ddl_text                IN   CLOB     := '-',
   new_master_rollback_seg     IN   VARCHAR2 := '-',
   new_flavor_id               IN   NUMBER   := -1e-130);
```

## Parameters

***Table 21–4 ALTER_TEMPLATE_OBJECT Procedure Parameters***

| Parameter | Description |
|---|---|
| refresh_template_name | Deployment template name that contains the object that you want to alter. |
| object_name | Name of the template object that you want to alter. |
| object_type | Type of object that you want to alter. |
| new_refresh_template_name | Name of the new deployment template to which you want to reassign this object. Do not specify a value to keep the object assigned to the current deployment template. |
| new_object_name | New name of the template object. Do not specify a value to keep the current object name. |
| new_object_type | If specified, then the new object type. Objects of the following type can be specified:<br>MATERIALIZED VIEW<br>INDEX<br>TABLE<br>VIEW<br>SYNONYM<br>SEQUENCE<br>PROCEDURE<br>FUNCTION<br>PACKAGE<br>PACKAGE BODY<br>TRIGGER<br>DATABASE LINK |
| new_ddl_text | New object DDL for specified object. Do not specify any new DDL text to keep the current object DDL. |
| new_master_rollback_seg | New master rollback segment for specified object. Do not specify a value to keep the current rollback segment. |
| new_flavor_id | This parameter is for internal use only.<br>**Note:** Do not set this parameter unless directed to do so by Oracle Support Services. |

## Exceptions

*Table 21–5 ALTER_TEMPLATE_OBJECT Procedure Exceptions*

| Exception | Description |
|---|---|
| miss_refresh_template | Deployment template name specified is invalid or does not exist. |
| miss_flavor_id | If you receive this exception, contact Oracle Support Services. |
| bad_object_type | Object type is specified incorrectly. See Table 21–4 for a list of valid object types. |
| miss_template_object | Template object name specified is invalid or does not exist. |
| dupl_template_object | New template name specified in the new_refresh_template_name parameter exists. |

## Usage Notes

Because the ALTER_TEMPLATE_OBJECT procedure utilizes a CLOB, you must use the DBMS_LOB package when using the ALTER_TEMPLATE_OBJECT procedure. The following example illustrates how to use the DBMS_LOB package with the ALTER_TEMPLATE_OBJECT procedure:

```
DECLARE
   tempstring VARCHAR2(100);
   templob CLOB;
BEGIN
   DBMS_LOB.CREATETEMPORARY(templob, TRUE, DBMS_LOB.SESSION);
   tempstring := 'CREATE MATERIALIZED VIEW mview_sales AS SELECT *
      FROM sales WHERE salesperson = :salesid and region_id = :region';
   DBMS_LOB.WRITE(templob, length(tempstring), 1, tempstring);
   DBMS_REPCAT_RGT.ALTER_TEMPLATE_OBJECT(
      refresh_template_name => 'rgt_personnel',
      object_name => 'MVIEW_SALES',
      object_type => 'MATERIALIZED VIEW',
      new_ddl_text => templob);
   DBMS_LOB.FREETEMPORARY(templob);
END;
/
```

# ALTER_TEMPLATE_PARM Procedure

This procedure allows the DBA to alter the parameters for a specific deployment template. Alterations include renaming the parameter and redefining the default value and prompt string.

## Syntax

```
DBMS_REPCAT_RGT.ALTER_TEMPLATE_PARM (
   refresh_template_name       IN   VARCHAR2,
   parameter_name              IN   VARCHAR2,
   new_refresh_template_name   IN   VARCHAR2 := '-',
   new_parameter_name          IN   VARCHAR2 := '-',
   new_default_parm_value      IN   CLOB := NULL,
   new_prompt_string           IN   VARCHAR2 := '-',
   new_user_override           IN   VARCHAR2 := '-');
```

## Parameters

***Table 21–6 ALTER_TEMPLATE_PARM Procedure Parameters***

| Parameter | Description |
|---|---|
| refresh_template_name | Name of the deployment template that contains the parameter that you want to alter. |
| parameter_name | Name of the parameter that you want to alter. |
| new_refresh_template_name | Name of the deployment template that the specified parameter should be reassigned to (useful when you want to move a parameter from one template to another). Do not specify a value to keep the parameter assigned to the current template. |
| new_parameter_name | New name of the template parameter. Do not specify a value to keep the current parameter name. |
| new_default_parm_value | New default value for the specified parameter. Do not specify a value to keep the current default value. |
| new_prompt_string | New prompt text for the specified parameter. Do not specify a value to keep the current prompt string. |
| new_user_override | Determines whether the user can override the default value if prompted during the instantiation process. The user is prompted if no user parameter value has been defined for this parameter. Set this parameter to 'Y' to allow a user to override the default value or set this parameter to 'N' to prevent an override. |

## Exceptions

***Table 21–7 ALTER_TEMPLATE_PARM Procedure Exceptions***

| Exception | Description |
|---|---|
| miss_refresh_template | Deployment template name specified is invalid or does not exist. |
| miss_template_parm | Template parameter specified is invalid or does not exist. |
| dupl_template_parm | Combination of new_refresh_template_name and new_parameter_name exists. |

## Usage Notes

Because the ALTER_TEMPLATE_PARM procedure utilizes a CLOB, you must use the DBMS_LOB package when using the ALTER_TEMPLATE_PARM procedure. The following example illustrates how to use the DBMS_LOB package with the ALTER_TEMPLATE_PARM procedure:

```
DECLARE
   tempstring VARCHAR2(100);
   templob CLOB;
BEGIN
   DBMS_LOB.CREATETEMPORARY(templob, TRUE, DBMS_LOB.SESSION);
   tempstring := 'REGION 20';
   DBMS_LOB.WRITE(templob, length(tempstring), 1, tempstring);
   DBMS_REPCAT_RGT.ALTER_TEMPLATE_PARM(
      refresh_template_name => 'rgt_personnel',
      parameter_name => 'region',
      new_default_parm_value => templob);
   DBMS_LOB.FREETEMPORARY(templob);
END;
/
```

# ALTER_USER_AUTHORIZATION Procedure

This procedure alters the contents of the DBA_REPCAT_USER_AUTHORIZATIONS view. Specifically, you can change user/deployment template authorization assignments. This procedure is helpful, for example, if an employee is reassigned and requires the materialized view environment of another deployment template. The DBA simply assigns the employee the new deployment template and the user is authorized to instantiate the target template.

## Syntax

```
DBMS_REPCAT_RGT.ALTER_USER_AUTHORIZATION (
   user_name                    IN   VARCHAR2,
   refresh_template_name        IN   VARCHAR2,
   new_user_name                IN   VARCHAR2 := '-',
   new_refresh_template_name    IN   VARCHAR2 := '-');
```

## Parameters

***Table 21–8 ALTER_USER_AUTHORIZATION Procedure Parameters***

| Parameter | Description |
|---|---|
| user_name | Name of the user whose authorization you want to alter. |
| refresh_template_name | Name of the deployment template that is currently assigned to the specified user that you want to alter. |
| new_user_name | Use this parameter to define a new user for this template authorization. Do not specify a value to keep the current user. |
| new_refresh_template_name | The deployment template that the specified user (either the existing or, if specified, the new user) is authorized to instantiate. Do not specify a value to keep the current deployment template. |

## Exceptions

***Table 21–9 ALTER_USER_AUTHORIZATION Procedure Exceptions***

| Exception | Description |
|---|---|
| miss_user_authorization | The combination of user_name and refresh_template_name values specified does not exist in the DBA_REPCAT_USER_AUTHORIZATIONS view. |
| miss_user | The user name specified for the new_user_name or user_name parameter is invalid or does not exist. |
| miss_refresh_template | The deployment template specified for the new_refresh_template parameter is invalid or does not exist. |
| dupl_user_authorization | A row exists for the specified user name and deployment template name. See the ALL_REPCAT_USER_AUTHORIZATIONS view. |

## ALTER_USER_PARM_VALUE Procedure

This procedure changes existing parameter values that have been defined for a specific user. This procedure is especially helpful if your materialized view environment uses assignment tables. Change a user parameter value to quickly and securely change the data set of a remote materialized view site.

> **See Also:** *Oracle Database Advanced Replication* for more information about using assignment tables

### Syntax

```
DBMS_REPCAT_RGT.ALTER_USER_PARM_VALUE(
   refresh_template_name         IN   VARCHAR2,
   parameter_name                IN   VARCHAR2,
   user_name                     IN   VARCHAR2,
   new_refresh_template_name     IN   VARCHAR2 := '-',
   new_parameter_name            IN   VARCHAR2 := '-',
   new_user_name                 IN   VARCHAR2 := '-',
   new_parm_value                IN   CLOB := NULL);
```

### Parameters

***Table 21–10 ALTER_USER_PARM_VALUE Procedure Parameters***

| Parameter | Description |
|---|---|
| refresh_template_name | Name of the deployment template that contains the user parameter value that you want to alter. |
| parameter_name | Name of the parameter that you want to alter. |
| user_name | Name of the user whose parameter value you want to alter. |
| new_refresh_template_name | Name of the deployment template that the specified user parameter value should be reassigned to (useful when you are authorizing a user for a different template). Do not specify a value to keep the parameter assigned to the current template. |
| new_parameter_name | The new template parameter name. Do not specify a value to keep the user value defined for the existing parameter. |
| new_user_name | The new user name that this parameter value is for. Do not specify a value to keep the parameter value assigned to the current user. |
| new_parm_value | The new parameter value for the specified user parameter. Do not specify a value to keep the current parameter value. |

## Exceptions

*Table 21–11 ALTER_USER_PARM_VALUE Procedure Exceptions*

| Exception | Description |
|---|---|
| miss_refresh_template | Deployment template name specified is invalid or does not exist. |
| miss_template_parm | Template parameter specified is invalid or does not exist. |
| miss_user | User name specified for the user_name or new_user_name parameters is invalid or does not exist. |
| miss_user_parm_values | User parameter value specified does not exist. |
| dupl_user_parm_values | New user parameter specified exists. |

## Usage Notes

Because the ALTER_USER_PARM_VALUE procedure utilizes a CLOB, you must use the DBMS_LOB package when using the ALTER_USER_PARM_VALUE procedure. The following example illustrates how to use the DBMS_LOB package with the ALTER_USER_PARM_VALUE procedure:

```
DECLARE
   tempstring VARCHAR2(100);
   templob CLOB;
BEGIN
   DBMS_LOB.CREATETEMPORARY(templob, TRUE, DBMS_LOB.SESSION);
   tempstring := 'REGION 20';
   DBMS_LOB.WRITE(templob, length(tempstring), 1, tempstring);
   DBMS_REPCAT_RGT.ALTER_USER_PARM_VALUE(
      refresh_template_name => 'rgt_personnel',
      parameter_name => 'region',
      user_name => 'BOB',
      new_parm_value => templob);
   DBMS_LOB.FREETEMPORARY(templob);
END;
/
```

## COMPARE_TEMPLATES Function

This function allows a DBA to compare the contents of two deployment templates. Any discrepancies between the two deployment templates is stored in the USER_REPCAT_TEMP_OUTPUT temporary view.

The COMPARE_TEMPLATES function returns a number that you specify in the WHERE clause when querying the USER_REPCAT_TEMP_OUTPUT temporary view. For example, if the COMPARE_TEMPLATES procedure returns the number 10, you would execute the following SELECT statement to view all discrepancies between two specified templates (your SELECT statement returns no rows if the templates are identical):

```
SELECT TEXT FROM USER_REPCAT_TEMP_OUTPUT
   WHERE OUTPUT_ID =  10 ORDER BY LINE;
```

The contents of the USER_REPCAT_TEMP_OUTPUT temporary view are lost after you disconnect or a rollback has been performed.

### Syntax

```
DBMS_REPCAT_RGT.COMPARE_TEMPLATES (
   source_template_name    IN   VARCHAR2,
   compare_template_name   IN   VARCHAR2)
  RETURN NUMBER;
```

### Parameters

***Table 21–12 COMPARE_TEMPLATES Function Parameters***

| Parameter | Description |
|---|---|
| source_template_name | Name of the first deployment template to be compared. |
| compare_template_name | Name of the second deployment template to be compared. |

### Exceptions

***Table 21–13 COMPARE_TEMPLATES Function Exceptions***

| Exception | Description |
|---|---|
| miss_refresh_template | The deployment template name to be compared is invalid or does not exist. |

### Returns

***Table 21–14 COMPARE_TEMPLATES Function Returns***

| Return Value | Description |
|---|---|
| <*system-generated number*> | Specifies the number returned for the output_id value when you select from the USER_REPCAT_TEMP_OUTPUT temporary view to view the discrepancies between the compared templates. |

# COPY_TEMPLATE Function

This function enables you to copy a deployment template and is helpful when a new deployment template uses many of the objects contained in an existing deployment template. This function copies the deployment template, template objects, template parameters, and user parameter values. The DBA can optionally have the function copy the user authorizations for this template. The number returned by this function is used internally by Oracle to manage deployment templates.

> **Note:** The values in the DBA_REPCAT_TEMPLATE_SITES view are not copied.

This function also allows the DBA to copy a deployment template to another master site, which is helpful for deployment template distribution and to split network loads between multiple sites.

## Syntax

```
DBMS_REPCAT_RGT.COPY_TEMPLATE (
   old_refresh_template_name     IN   VARCHAR2,
   new_refresh_template_name     IN   VARCHAR2,
   copy_user_authorizations      IN   VARCHAR2,
   dblink                        IN   VARCHAR2 := NULL)
  RETURN NUMBER;
```

## Parameters

***Table 21–15 COPY_TEMPLATE Function Parameters***

| Parameter | Description |
|---|---|
| old_refresh_template_name | Name of the deployment template to be copied. |
| new_refresh_template_name | Name of the new deployment template. |
| copy_user_authorizations | Specifies whether the template authorizations for the original template should be copied for the new deployment template. Valid values for this parameter are Y, N, and NULL.<br>**Note**: All users must exist at the target database. |
| dblink | Optionally defines where the deployment template should be copied from (this is helpful to distribute deployment templates to other master sites). If none is specified, then the deployment template is copied from the local master site. |

## Exceptions

***Table 21–16 COPY_TEMPLATE Function Exceptions***

| Exception | Description |
|---|---|
| miss_refresh_template | Deployment template name to be copied is invalid or does not exist. |
| dupl_refresh_template | Name of the new refresh template specified exists. |
| bad_copy_auth | Value specified for the copy_user_authorization parameter is invalid. Valid values are Y, N, and NULL. |

## Returns

***Table 21–17 COPY_TEMPLATE Function Returns***

| Return Value | Description |
|---|---|
| *<system-generated number>* | System-generated number used internally by Oracle. |

# CREATE_OBJECT_FROM_EXISTING Function

This function creates a template object definition from existing database objects and adds it to a target deployment template. The object DDL that created the original database object is executed when the target deployment template is instantiated at the remote materialized view site. This is ideal for adding existing triggers and procedures to your template. The number returned by this function is used internally by Oracle to manage deployment templates.

## Syntax

```
DBMS_REPCAT_RGT.CREATE_OBJECT_FROM_EXISTING(
     refresh_template_name  IN   VARCHAR2,
     object_name            IN   VARCHAR2,
     sname                  IN   VARCHAR2,
     oname                  IN   VARCHAR2,
     otype                  IN   VARCHAR2)
   RETURN NUMBER;
```

## Parameters

*Table 21–18 CREATE_OBJECT_FROM_EXISTING Function Parameters*

| Parameter | Description |
|---|---|
| refresh_template_name | Name of the deployment template to which you want to add this object. |
| object_name | Optionally, the new name of the existing object that you are adding to your deployment template (enables you to define a new name for an existing object). |
| sname | The schema that contains the object that you are creating your template object from. |
| oname | Name of the object that you are creating your template object from. |
| otype | The type of database object that you are adding to the template (that is, PROCEDURE, TRIGGER, and so on). Objects of the following types can be specified (DATABASE LINK, MATERIALIZED VIEW, and MATERIALIZED VIEW are not valid object types for this function):<br>SEQUENCE PROCEDURE<br>INDEX FUNCTION<br>TABLE PACKAGE<br>VIEW PACKAGE BODY<br>SYNONYM TRIGGER |

## Exceptions

*Table 21–19 CREATE_OBJECT_FROM_EXISTING Function Exceptions*

| Exception | Description |
|---|---|
| miss_refresh_template | The specified refresh template name is invalid or missing. Query the DBA_REPCAT_REFRESH_TEMPLATES view for a list of existing deployment templates. |
| bad_object_type | The object type is specified incorrectly. |
| dupl_template_object | An object of the same name and type has already been added to the specified deployment template. |

*Table 21–19 (Cont.) CREATE_OBJECT_FROM_EXISTING Function Exceptions*

| Exception | Description |
|---|---|
| `objectmissing` | The object specified does not exist. |

## Returns

*Table 21–20 CREATE_OBJECT_FROM_EXISTING Function Returns*

| Return Value | Description |
|---|---|
| *<system-generated number>* | System-generated number used internally by Oracle. |

# CREATE_REFRESH_TEMPLATE Function

This function creates the deployment template, which enables you to define the template name, private/public status, and target refresh group. Each time that you create a template object, user authorization, or template parameter, you reference the deployment template created with this function. This function adds a row to the DBA_REPCAT_REFRESH_TEMPLATES view. The number returned by this function is used internally by Oracle to manage deployment templates.

## Syntax

```
DBMS_REPCAT_RGT.CREATE_REFRESH_TEMPLATE (
   owner                  IN   VARCHAR2,
   refresh_group_name     IN   VARCHAR2,
   refresh_template_name  IN   VARCHAR2,
   template_comment       IN   VARCHAR2 := NULL,
   public_template        IN   VARCHAR2 := NULL,
   last_modified          IN   DATE := SYSDATE,
   modified_by            IN   VARCHAR2 := USER,
   creation_date          IN   DATE := SYSDATE,
   created_by             IN   VARCHAR2 := USER)
  RETURN NUMBER;
```

## Parameters

***Table 21–21 CREATE_REFRESH_TEMPLATE Function Parameters***

| Parameter | Description |
|---|---|
| owner | User name of the deployment template owner is specified with this parameter. If an owner is not specified, then the name of the user creating the template is automatically used. |
| refresh_group_name | Name of the refresh group that is created when this template is instantiated. All objects created by this template are assigned to the specified refresh group. |
| refresh_template_name | Name of the deployment template that you are creating. This name is referenced in all activities that involve this deployment template. |
| template_comment | User comments defined with this parameter are listed in the DBA_REPCAT_REFRESH_TEMPLATES view. |
| public_template | Specifies whether the deployment template is public or private. Only acceptable values are 'Y' and 'N' ('Y' = public and 'N' = private). |
| last_modified | The date of the last modification made to this deployment template. If a value is not specified, then the current date is automatically used. |
| modified_by | Name of the user who last modified this deployment template. If a value is not specified, then the current user is automatically used. |
| creation_date | The date that this deployment template was created. If a value is not specified, then the current date is automatically used. |
| created_by | Name of the user who created this deployment template. If a value is not specified, then the current user is automatically used. |

## Exceptions

*Table 21–22 CREATE_REFRESH_TEMPLATE Function Exceptions*

| Exception | Description |
|---|---|
| `dupl_refresh_template` | A template with the specified name exists. See the ALL_REPCAT_REFRESH_TEMPLATES view to see a list of existing templates. |
| `bad_public_template` | The `public_template` parameter is specified incorrectly. The `public_template` parameter must be specified as a `'Y'` for a public template or an `'N'` for a private template. |

## Returns

*Table 21–23 CREATE_REFRESH_TEMPLATE Function Returns*

| Return Value | Description |
|---|---|
| *<system-generated number>* | System-generated number used internally by Oracle. |

# CREATE_TEMPLATE_OBJECT Function

This function adds object definitions to a target deployment template container. The specified object DDL is executed when the target deployment template is instantiated at the remote materialized view site. In addition to adding materialized views, this function can add tables, procedures, and other objects to your template. The number returned by this function is used internally by Oracle to manage deployment templates.

## Syntax

```
DBMS_REPCAT_RGT.CREATE_TEMPLATE_OBJECT (
   refresh_template_name  IN   VARCHAR2,
   object_name            IN   VARCHAR2,
   object_type            IN   VARCHAR2,
   ddl_text               IN   CLOB,
   master_rollback_seg    IN   VARCHAR2 := NULL,
   flavor_id              IN   NUMBER := -1e-130)
  RETURN NUMBER;
```

## Parameters

***Table 21–24 CREATE_TEMPLATE_OBJECT Function Parameters***

| Parameter | Description |
|---|---|
| refresh_template_name | Name of the deployment template to which you want to add this object. |
| object_name | Name of the template object that you are creating. |
| object_type | The type of database object that you are adding to the template (that is, MATERIALIZED VIEW, TRIGGER, PROCEDURE, and so on). Objects of the following types can be specified:<br>MATERIALIZED VIEW — PROCEDURE<br>INDEX — FUNCTION<br>TABLE — PACKAGE<br>VIEW — PACKAGE BODY<br>SYNONYM — TRIGGER<br>SEQUENCE — DATABASE LINK |
| ddl_text | Contains the DDL that creates the object that you are adding to the template. Be sure to end your DDL with a semicolon. You can use a colon (:) to create a template parameter for your template object. See Chapter 4, "Creating a Deployment Template" for more information.<br>When you add a materialized view with a CREATE MATERIALIZED VIEW statement, ensure that you specify the schema name of the owner of the master table in the materialized view query. |
| master_rollback_seg | Specifies the name of the rollback segment to use when executing the defined object DDL at the remote materialized view site. |
| flavor_id | This parameter is for internal use only.<br>**Note:** Do not set this parameter unless directed to do so by Oracle Support Services. |

## Exceptions

*Table 21–25 CREATE_TEMPLATE_OBJECT Function Exceptions*

| Exception | Description |
|---|---|
| miss_refresh_template | Specified refresh template name is invalid or missing. Query the DBA_REPCAT_REFRESH_TEMPLATES view for a list of existing deployment templates. |
| bad_object_type | Object type is specified incorrectly. See Table 21–24 for a list of valid object types. |
| dupl_template_object | An object of the same name and type has already been added to the specified deployment template. |

## Returns

*Table 21–26 CREATE_TEMPLATE_OBJECT Function Returns*

| Return Value | Description |
|---|---|
| <*system-generated number*> | System-generated number used internally by Oracle. |

## Usage Notes

Because CREATE_TEMPLATE_OBJECT utilizes a CLOB, you must use the DBMS_LOB package when using the CREATE_TEMPLATE_OBJECT function. The following example illustrates how to use the DBMS_LOB package with the CREATE_TEMPLATE_OBJECT function:

```
DECLARE
   tempstring VARCHAR2(100);
   templob CLOB;
   a NUMBER;
BEGIN
   DBMS_LOB.CREATETEMPORARY(templob, TRUE, DBMS_LOB.SESSION);
   tempstring := 'CREATE MATERIALIZED VIEW mview_sales AS SELECT *
        FROM sales WHERE salesperson = :salesid';
   DBMS_LOB.WRITE(templob, length(tempstring), 1, tempstring);
   a := DBMS_REPCAT_RGT.CREATE_TEMPLATE_OBJECT(
        refresh_template_name => 'rgt_personnel',
        object_name => 'mview_sales',
        object_type => 'MATERIALIZED VIEW',
        ddl_text => templob,
        master_rollback_seg => 'RBS');
   DBMS_LOB.FREETEMPORARY(templob);
END;
/
```

# CREATE_TEMPLATE_PARM Function

This function creates parameters for a specific deployment template to allow custom data sets to be created at the remote materialized view site. This function is only required when the DBA wants to define a set of template variables before adding any template objects. When objects are added to the template using the CREATE_TEMPLATE_OBJECT function, any variables in the object DDL are automatically added to the DBA_REPCAT_TEMPLATE_PARMS view.

The DBA typically uses the ALTER_TEMPLATE_PARM procedure to modify the default parameter values or prompt strings or both (see "ALTER_TEMPLATE_PARM Procedure" on page 21-8 for more information). The number returned by this function is used internally by Oracle to manage deployment templates.

## Syntax

```
DBMS_REPCAT_RGT.CREATE_TEMPLATE_PARM (
   refresh_template_name  IN   VARCHAR2,
   parameter_name         IN   VARCHAR2,
   default_parm_value     IN   CLOB := NULL,
   prompt_string          IN   VARCHAR2 := NULL,
   user_override          IN   VARCHAR2 := NULL)
 RETURN NUMBER;
```

## Parameters

*Table 21–27 CREATE_TEMPLATE_PARM Function Parameters*

| Parameter | Description |
|---|---|
| refresh_template_name | Name of the deployment template for which you want to create the parameter. |
| parameter_name | Name of the parameter you are creating. |
| default_parm_value | Default values for this parameter are defined using this parameter. If a user parameter value or run-time parameter value is not present, then this default value is used during the instantiation process. |
| prompt_string | The descriptive prompt text that is displayed for this template parameter during the instantiation process. |
| user_override | Determines whether the user can override the default value if prompted during the instantiation process. The user is prompted if no user parameter value has been defined for this parameter. Set this parameter to 'Y' to allow a user to override the default value or set this parameter to 'N' to not allow an override. |

## Exceptions

*Table 21–28 CREATE_TEMPLATE_PARM Function Exceptions*

| Exception | Description |
|---|---|
| miss_refresh_template | The specified refresh template name is invalid or missing. |
| dupl_template_parm | A parameter of the same name has already been defined for the specified deployment template. |

## Returns

*Table 21–29 CREATE_TEMPLATE_PARM Function Returns*

| Return Value | Description |
|---|---|
| *<system-generated number>* | System-generated number used internally by Oracle. |

## Usage Notes

Because the CREATE_TEMPLATE_PARM function utilizes a CLOB, you must use the DBMS_LOB package when using the CREATE_TEMPLATE_PARM function. The following example illustrates how to use the DBMS_LOB package with the CREATE_TEMPLATE_PARM function:

```
DECLARE
   tempstring VARCHAR2(100);
   templob CLOB;
   a NUMBER;
BEGIN
   DBMS_LOB.CREATETEMPORARY(templob, TRUE, DBMS_LOB.SESSION);
   tempstring := 'REGION 20';
   DBMS_LOB.WRITE(templob, length(tempstring), 1, tempstring);
   a := DBMS_REPCAT_RGT.CREATE_TEMPLATE_PARM(
        refresh_template_name => 'rgt_personnel',
        parameter_name => 'region',
        default_parm_value => templob,
        prompt_string => 'Enter your region ID:',
        user_override => 'Y');
   DBMS_LOB.FREETEMPORARY(templob);
END;
/
```

# CREATE_USER_AUTHORIZATION Function

This function authorizes specific users to instantiate private deployment templates. Users not authorized for a private deployment template are not able to instantiate the private template. This function adds a row to the DBA_REPCAT_USER_AUTHORIZATIONS view.

Before you authorize a user, verify that the user exists at the master site where the user will instantiate the deployment template. The number returned by this function is used internally by Oracle to manage deployment templates.

## Syntax

```
DBMS_REPCAT_RGT.CREATE_USER_AUTHORIZATION (
   user_name               IN   VARCHAR2,
   refresh_template_name   IN   VARCHAR2)
 RETURN NUMBER;
```

## Parameters

***Table 21–30 CREATE_USER_AUTHORIZATION Function Parameters***

| Parameter | Description |
|---|---|
| user_name | Name of the user that you want to authorize to instantiate the specified template. Specify multiple users by separating user names with a comma (for example, 'john, mike, bob') |
| refresh_template_name | Name of the template that you want to authorize the specified user to instantiate. |

## Exceptions

***Table 21–31 CREATE_USER_AUTHORIZATION Function Exceptions***

| Exception | Description |
|---|---|
| miss_user | User name supplied is invalid or does not exist. |
| miss_refresh_template | Refresh template name supplied is invalid or does not exist. |
| dupl_user_authorization | An authorization has already been created for the specified user and deployment template. See the ALL_REPCAT_USER_AUTHORIZATIONS view for a listing of template authorizations. |

## Returns

***Table 21–32 CREATE_USER_AUTHORIZATION Function Returns***

| Return Value | Description |
|---|---|
| *<system-generated number>* | System-generated number used internally by Oracle. |

## CREATE_USER_PARM_VALUE Function

This function predefines deployment template parameter values for specific users. For example, to predefine the region parameter as `west` for user `33456`, use the this function.

Any values specified with this function take precedence over default values specified for the template parameter. The number returned by this function is used internally by Oracle to manage deployment templates.

### Syntax

```
DBMS_REPCAT_RGT.CREATE_USER_PARM_VALUE (
   refresh_template_name    IN   VARCHAR2,
   parameter_name           IN   VARCHAR2,
   user_name                IN   VARCHAR2,
   parm_value               IN   CLOB := NULL)
 RETURN NUMBER;
```

### Parameters

***Table 21–33 CREATE_USER_PARM_VALUE Function Parameters***

| Parameter | Description |
|---|---|
| `refresh_template_name` | Specifies the name of the deployment template that contains the parameter you are creating a user parameter value for. |
| `parameter_name` | Name of the template parameter that you are defining a user parameter value for. |
| `user_name` | Specifies the name of the user that you are predefining a user parameter value for. |
| `parm_value` | The predefined parameter value that will be used during the instantiation process initiated by the specified user. |

### Exceptions

***Table 21–34 CREATE_USER_PARM_VALUE Function Exceptions***

| Exception | Description |
|---|---|
| `miss_refresh_template` | Specified deployment template name is invalid or missing. |
| `dupl_user_parm_values` | A parameter value for the specified user, parameter, and deployment template has already been defined. Query the `DBA_REPCAT_USER_PARM_VALUES` view for a listing of existing user parameter values. |
| `miss_template_parm` | Specified deployment template parameter name is invalid or missing. |
| `miss_user` | Specified user name is invalid or missing. |

## Returns

*Table 21–35 CREATE_USER_PARM_VALUE Function Returns*

| Return Value | Description |
|---|---|
| *<system-generated number>* | System-generated number used internally by Oracle. |

## Usage Notes

Because the CREATE_USER_PARM_VALUE function utilizes a CLOB, you must use the DBMS_LOB package when using the this function. The following example illustrates how to use the DBMS_LOB package with the CREATE_USER_PARM_VALUE function:

```
DECLARE
   tempstring VARCHAR2(100);
   templob CLOB;
   a NUMBER;
BEGIN
   DBMS_LOB.CREATETEMPORARY(templob, TRUE, DBMS_LOB.SESSION);
   tempstring := 'REGION 20';
   DBMS_LOB.WRITE(templob, length(tempstring), 1, tempstring);
   a := DBMS_REPCAT_RGT.CREATE_USER_PARM_VALUE(
        refresh_template_name => 'rgt_personnel',
        parameter_name => 'region',
        user_name => 'BOB',
        user_parm_value => templob);
   DBMS_LOB.FREETEMPORARY(templob);
END;
/
```

## DELETE_RUNTIME_PARMS Procedure

Use this procedure before instantiating a deployment template to delete a run-time parameter value that you defined using the INSERT_RUNTIME_PARMS procedure.

### Syntax

```
DBMS_REPCAT_RGT.DELETE_RUNTIME_PARMS(
   runtime_parm_id     IN   NUMBER,
   parameter_name      IN   VARCHAR2);
```

### Parameters

*Table 21–36 DELETE_RUNTIME_PARMS Procedure Parameters*

| Parameter | Description |
|---|---|
| runtime_parm_id | Specifies the identification that you previously assigned the run-time parameter value to (this value was retrieved using the GET_RUNTIME_PARM_ID function). |
| parameter_name | Specifies the name of the parameter value that you want to drop (query the DBA_REPCAT_TEMPLATE_PARMS view for a list of deployment template parameters). |

### Exceptions

*Table 21–37 DELETE_RUNTIME_PARMS Procedure Exceptions*

| Exception | Description |
|---|---|
| miss_template_parm | The specified deployment template parameter name is invalid or missing. |

# DROP_ALL_OBJECTS Procedure

This procedure allows the DBA to drop all objects or specific object types from a deployment template.

**Caution:** This is a dangerous procedure that cannot be undone.

## Syntax

```
DBMS_REPCAT_RGT.DROP_ALL_OBJECTS (
   refresh_template_name   IN   VARCHAR2,
   object_type             IN   VARCHAR2 := NULL);
```

## Parameters

***Table 21–38 DROP_ALL_OBJECTS Procedure Parameters***

| Parameter | Description |
|---|---|
| refresh_template_name | Name of the deployment template that contains the objects that you want to drop. |
| object_type | If NULL, then all objects in the template are dropped. If an object type is specified, then only objects of that type are dropped. Objects of the following types can be specified:<br>MATERIALIZED VIEW<br>INDEX<br>TABLE<br>VIEW<br>SYNONYM<br>SEQUENCE<br>PROCEDURE<br>FUNCTION<br>PACKAGE<br>PACKAGE BODY<br>TRIGGER<br>DATABASE LINK |

## Exceptions

***Table 21–39 DROP_ALL_OBJECTS Procedure Exceptions***

| Exception | Description |
|---|---|
| miss_refresh_template | Specified deployment template name is invalid or does not exist. |
| bad_object_type | Object type is specified incorrectly. See Table 21–38 for a list of valid object types. |

## DROP_ALL_TEMPLATE_PARMS Procedure

This procedure lets you drop template parameters for a specified deployment template. You can use this procedure to drop all parameters that are not referenced by a template object or to drop from the template all objects that reference any parameter, along with all of the parameters themselves.

**Caution:** This is a dangerous procedure that cannot be undone.

### Syntax

```
DBMS_REPCAT_RGT.DROP_ALL_TEMPLATE_PARMS (
   refresh_template_name   IN   VARCHAR2,
   drop_objects            IN   VARCHAR2 := N);
```

### Parameters

*Table 21–40 DROP_ALL_TEMPLATE_PARMS Procedure Parameters*

| Parameter | Description |
|---|---|
| refresh_template_name | Name of the deployment template that contains the parameters and objects that you want to drop. |
| drop_objects | If no value is specified, then this parameter defaults to N, which drops all parameters not referenced by a template object.<br>If Y is specified, then all objects that reference any template parameter and the template parameters themselves are dropped. The objects are dropped from the template, not from the database. |

### Exceptions

*Table 21–41 DROP_ALL_TEMPLATE_PARMS Procedure Exceptions*

| Exception | Description |
|---|---|
| miss_refresh_template | Specified deployment template name is invalid or does not exist. |

# DROP_ALL_TEMPLATE_SITES Procedure

This procedure removes all entries from the DBA_REPCAT_TEMPLATE_SITES view, which keeps a record of sites that have instantiated a particular deployment template.

**Caution:** This is a dangerous procedure that cannot be undone.

## Syntax

```
DBMS_REPCAT_RGT.DROP_ALL_TEMPLATE_SITES (
   refresh_template_name   IN   VARCHAR2);
```

## Parameter

*Table 21–42 DROP_ALL_TEMPLATE_SITES Procedure Parameter*

| Parameter | Description |
|---|---|
| refresh_template_name | Name of the deployment template that contains the sites that you want to drop. |

## Exceptions

*Table 21–43 DROP_ALL_TEMPLATE_SITES Procedure Exceptions*

| Exception | Description |
|---|---|
| miss_refresh_template | Specified deployment template name is invalid or does not exist. |

## DROP_ALL_TEMPLATES Procedure

This procedure removes all deployment templates at the site where the procedure is called.

**Caution: This is a dangerous procedure that cannot be undone.**

### Syntax

```
DBMS_REPCAT_RGT.DROP_ALL_TEMPLATES;
```

### Parameters

None

# DROP_ALL_USER_AUTHORIZATIONS Procedure

This procedure enables the DBA to drop all user authorizations for a specified deployment template. Executing this procedure removes rows from the DBA_REPCAT_USER_AUTHORIZATIONS view.

This procedure might be implemented after converting a private template to a public template and the user authorizations are no longer required.

## Syntax

```
DBMS_REPCAT_RGT.DROP_ALL_USER_AUTHORIZATIONS (
   refresh_template_name   IN   VARCHAR2);
```

## Parameters

*Table 21–44 DROP_ALL_USER_AUTHORIZATIONS Procedure Parameters*

| Parameter | Description |
|---|---|
| refresh_template_name | Name of the deployment template that contains the user authorizations that you want to drop. |

## Exceptions

*Table 21–45 DROP_ALL_USER_AUTHORIZATIONS Procedure Exceptions*

| Exception | Description |
|---|---|
| miss_refresh_template | Specified deployment template name is invalid or does not exist. |

## DROP_ALL_USER_PARM_VALUES Procedure

This procedure drops user parameter values for a specific deployment template. This procedure is very flexible and enables you to define a set of user parameter values to be deleted.

For example, defining the parameters shown in the following table has the described results.

| Parameter | Result of Defining the Parameter |
|---|---|
| refresh_template_name | Drops all user parameters for the specified deployment template |
| refresh_template_name and user_name | Drops all of the specified user parameters for the specified deployment template |
| refresh_template_name and parameter_name | Drops all user parameter values for the specified deployment template parameter |
| refresh_template_name, parameter_name, and user_name | Drops the specified user's value for the specified deployment template parameter (equivalent to drop_user_parm) |

### Syntax

```
DBMS_REPCAT_RGT.DROP_ALL_USER_PARM_VALUES (
   refresh_template_name   IN   VARCHAR2,
   user_name               IN   VARCHAR2,
   parameter_name          IN   VARCHAR2);
```

### Parameters

*Table 21–46 DROP_ALL_USER_PARM_VALUES Procedure Parameters*

| Parameter | Description |
|---|---|
| refresh_template_name | Name of the deployment template that contains the parameter values that you want to drop. |
| user_name | Name of the user whose parameter values you want to drop. |
| parameter_name | Template parameter that contains the values that you want to drop. |

### Exceptions

*Table 21–47 DROP_ALL_USER_PARM_VALUES Procedure Exceptions*

| Exception | Description |
|---|---|
| miss_refresh_template | Deployment template name specified is invalid or does not exist. |
| miss_user | User name specified is invalid or does not exist. |
| miss_user_parm_values | Deployment template, user, and parameter combination does not exist in the DBA_REPCAT_USER_PARM_VALUES view. |

# DROP_REFRESH_TEMPLATE Procedure

This procedure drops a deployment template. Dropping a deployment template has a cascading effect, removing all related template parameters, user authorizations, template objects, and user parameters (this procedure does not drop template sites).

## Syntax

```
DBMS_REPCAT_RGT.DROP_REFRESH_TEMPLATE (
   refresh_template_name  IN   VARCHAR2);
```

## Parameters

***Table 21–48 DROP_REFRESH_TEMPLATE Procedure Parameters***

| Parameter | Description |
|---|---|
| refresh_template_name | Name of the deployment template to be dropped. |

## Exceptions

***Table 21–49 DROP_REFRESH_TEMPLATE Procedure Exceptions***

| Exception | Description |
|---|---|
| miss_refresh_template | The deployment template name specified is invalid or does not exist. Query the DBA_REPCAT_REFRESH_TEMPLATES view for a list of deployment templates. |

## DROP_SITE_INSTANTIATION Procedure

This procedure drops a template instantiation at any target site. This procedure removes all related metadata at the master site and disables the specified site from refreshing its materialized views.

### Syntax

```
DBMS_REPCAT_RGT.DROP_SITE_INSTANTIATION (
     refresh_template_name  IN   VARCHAR2,
     user_name              IN   VARCHAR2,
     site_name              IN   VARCHAR2);
```

### Parameters

***Table 21–50 DROP_SITE_INSTANTIATION Procedure Parameters***

| Parameter | Description |
|---|---|
| refresh_template_name | The name of the template that contains the site to be dropped. |
| user_name | The name of the user who originally instantiated the template at the remote materialized view site. Query the ALL_REPCAT_TEMPLATE_SITES view to see the users that instantiated templates. See the ALL_REPCAT_TEMPLATE_SITES view on page 23-11 for more information. |
| site_name | Identifies the template site to be dropped. |

### Exceptions

***Table 21–51 DROP_SITE_INSTANTIATION Procedure Exceptions***

| Exception | Description |
|---|---|
| miss_refresh_template | The deployment template name specified is invalid or does not exist. |
| miss_user | The user name specified does not exist. |
| miss_template_site | The deployment template has not been instantiated for user and site. |

# DROP_TEMPLATE_OBJECT Procedure

This procedure removes a template object from a specific deployment template. For example, a DBA would use this procedure to remove an outdated materialized view from a deployment template. Changes made to the template are reflected at new sites instantiating the deployment template. Remote sites that have already instantiated the template must reinstantiate the deployment template to apply the changes.

## Syntax

```
DBMS_REPCAT_RGT.DROP_TEMPLATE_OBJECT (
   refresh_template_name  IN   VARCHAR2,
   object_name            IN   VARCHAR2,
   object_type            IN   VARCHAR2);
```

## Parameters

***Table 21–52 DROP_TEMPLATE_OBJECT Procedure Parameters***

| Parameter | Description |
|---|---|
| refresh_template_name | Name of the deployment template from which you are dropping the object. |
| object_name | Name of the template object to be dropped. |
| object_type | The type of object that is to be dropped. Objects of the following types can be specified:<br>MATERIALIZED VIEW<br>INDEX<br>TABLE<br>VIEW<br>SYNONYM<br>SEQUENCE<br>PROCEDURE<br>FUNCTION<br>PACKAGE<br>PACKAGE BODY<br>TRIGGER<br>DATABASE LINK |

## Exceptions

***Table 21–53 DROP_TEMPLATE_OBJECT Procedure Exceptions***

| Exception | Description |
|---|---|
| miss_refresh_template | The deployment template name specified is invalid or does not exist. |
| miss_template_object | The template object specified is invalid or does not exist. Query the DBA_REPCAT_TEMPLATE_OBJECTS view to see a list of deployment template objects. |

## DROP_TEMPLATE_PARM Procedure

This procedure removes an existing template parameter from the DBA_REPCAT_TEMPLATE_PARMS view. This procedure is useful when you have dropped a template object and a particular parameter is no longer needed.

### Syntax

```
DBMS_REPCAT_RGT.DROP_TEMPLATE_PARM (
   refresh_template_name  IN   VARCHAR2,
   parameter_name         IN   VARCHAR2);
```

### Parameters

*Table 21–54 DROP_TEMPLATE_PARM Procedure Parameters*

| Parameter | Description |
|---|---|
| refresh_template_name | The deployment template name that has the parameter that you want to drop |
| parameter_name | Name of the parameter that you want to drop. |

### Exceptions

*Table 21–55 DROP_TEMPLATE_PARM Procedure Exceptions*

| Exception | Description |
|---|---|
| miss_refresh_template | The deployment template name specified is invalid or does not exist. |
| miss_template_parm | The parameter name specified is invalid or does not exist. Query the DBA_REPCAT_TEMPLATE_PARMS view to see a list of template parameters. |

# DROP_USER_AUTHORIZATION Procedure

This procedure removes a user authorization entry from the DBA_REPCAT_USER_AUTHORIZATIONS view. This procedure is used when removing a user's template authorization. If a user's authorization is removed, then the user is no longer able to instantiate the target deployment template.

**See Also:** "DROP_ALL_USER_AUTHORIZATIONS Procedure" on page 21-32

## Syntax

```
DBMS_REPCAT_RGT.DROP_USER_AUTHORIZATION (
   refresh_template_name   IN   VARCHAR2,
   user_name               IN   VARCHAR2);
```

## Parameters

***Table 21–56 DROP_USER_AUTHORIZATION Procedure Parameters***

| Parameter | Description |
|---|---|
| refresh_template_name | Name of the deployment template from which the user's authorization is being removed. |
| user_name | Name of the user whose authorization is being removed. |

## Exceptions

***Table 21–57 DROP_USER_AUTHORIZATION Procedure Exceptions***

| Exception | Description |
|---|---|
| miss_user | Specified user name is invalid or does not exist. |
| miss_user_authorization | Specified user and deployment template combination does not exist. Query the DBA_REPCAT_USER_AUTHORIZATIONS view to see a list of user/deployment template authorizations. |
| miss_refresh_template | Specified deployment template name is invalid or does not exist. |

## DROP_USER_PARM_VALUE Procedure

This procedure removes a predefined user parameter value for a specific deployment template. This procedure is often executed after a user's template authorization has been removed.

### Syntax

```
DBMS_REPCAT_RGT.DROP_USER_PARM_VALUE (
   refresh_template_name    IN   VARCHAR2,
   parameter_name           IN   VARCHAR2,
   user_name                IN   VARCHAR2);
```

### Parameters

*Table 21–58 DROP_USER_PARM_VALUE Procedure Parameters*

| Parameter | Description |
|---|---|
| refresh_template_name | Deployment template name that contains the parameter value that you want to drop. |
| parameter_name | Parameter name that contains the predefined value that you want to drop. |
| user_name | Name of the user whose parameter value you want to drop. |

### Exceptions

*Table 21–59 DROP_USER_PARM_VALUE Procedure Exceptions*

| Exception | Description |
|---|---|
| miss_refresh_template | Deployment template name specified is invalid or does not exist. |
| miss_user | User name specified is invalid or does not exist. |
| miss_user_parm_values | Deployment template, user, and parameter combination does not exist in the DBA_REPCAT_USER_PARM_VALUES view. |

# GET_RUNTIME_PARM_ID Function

This function retrieves an identification to be used when defining a run-time parameter value. All run-time parameter values are assigned to this identification and are also used during the instantiation process.

## Syntax

```
DBMS_REPCAT_RGT.GET_RUNTIME_PARM_ID
  RETURN NUMBER;
```

## Parameters

None

## Returns

***Table 21–60 GET_RUNTIME_PARM_ID Function Returns***

| Return Value | Corresponding Data Type |
|---|---|
| *<system-generated number>* | Runtime parameter values are assigned to the system-generated number and are also used during the instantiation process. |

## INSERT_RUNTIME_PARMS Procedure

This procedure defines run-time parameter values before instantiating a template. Use this procedure to define parameter values when no user parameter values have been defined and you do not want to accept the default parameter values.

Before using the this procedure, be sure to execute the GET_RUNTIME_PARM_ID function to retrieve a parameter identification to use when inserting a run-time parameter. This identification is used for defining run-time parameter values and instantiating deployment templates.

### Syntax

```
DBMS_REPCAT_RGT.INSERT_RUNTIME_PARMS (
   runtime_parm_id    IN   NUMBER,
   parameter_name     IN   VARCHAR2,
   parameter_value    IN   CLOB);
```

### Parameters

*Table 21–61 INSERT_RUNTIME_PARMS Procedure Parameters*

| Parameter | Description |
|---|---|
| runtime_parm_id | The identification retrieved by the GET_RUNTIME_PARM_ID function. This identification is also used when instantiating the deployment template. Be sure to use the same identification for all parameter values for a deployment template. |
| parameter_name | Name of the template parameter for which you are defining a run-time parameter value. Query the DBA_REPCAT_TEMPLATE_PARMS view for a list of template parameters. |
| parameter_value | The run-time parameter value that you want to use during the deployment template instantiation process. |

### Exceptions

*Table 21–62 INSERT_RUNTIME_PARMS Procedure Exceptions*

| Exception | Description |
|---|---|
| miss_refresh_template | The deployment template name specified is invalid or does not exist. |
| miss_user | The user name specified is invalid or does not exist. |
| miss_user_parm_values | The deployment template, user, and parameter combination does not exist in the DBA_REPCAT_USER_PARM_VALUES view. |

### Usage Notes

Because the this procedure utilizes a CLOB, you must use the DBMS_LOB package when using the INSERT_RUNTIME_PARMS procedure. The following example illustrates how to use the DBMS_LOB package with the INSERT_RUNTIME_PARMS procedure:

```
DECLARE
   tempstring VARCHAR2(100);
   templob CLOB;
BEGIN
   DBMS_LOB.CREATETEMPORARY(templob, TRUE, DBMS_LOB.SESSION);
   tempstring := 'REGION 20';
```

```
   DBMS_LOB.WRITE(templob, length(tempstring), 1, tempstring);
   DBMS_REPCAT_RGT.INSERT_RUNTIME_PARMS(
      runtime_parm_id => 20,
      parameter_name => 'region',
      parameter_value => templob);
   DBMS_LOB.FREETEMPORARY(templob);
END;
/
```

## INSTANTIATE_OFFLINE Function

This function generates a script at the master site that is used to create the materialized view environment at the remote materialized view site while the materialized view site disconnected from the master (that is, while the materialized view site is offline). This generated script should be used at remote materialized view sites that are not able to remain connected to the master site for an extended amount of time, as the instantiation process at the remote materialized view site might be lengthy (depending on the amount of data that is populated to the new materialized views). This function must be executed separately for each user instantiation.

The script generated by this function is stored in the `USER_REPCAT_TEMP_OUTPUT` temporary view and is used by several Oracle tools, including the Advanced Replication interface in Oracle Enterprise Manager, during the distribution of deployment templates. The number returned by this function is used to retrieve the appropriate information from the `USER_REPCAT_TEMP_OUTPUT` temporary view.

> **Note:** This function is used to perform an offline instantiation of a deployment template. Additionally, this function is for replication administrators who are instantiating for another user. Users wanting to perform their own instantiation should use the public version of the `INSTANTIATE_OFFLINE` function. See the "INSTANTIATE_OFFLINE Function" on page 21-43 for more information.
>
> This function should not be confused with the procedures in the DBMS_OFFLINE_OG package (used for performing an offline instantiation of a master table). See the documentation for this package for more information about their usage.

### Syntax

```
DBMS_REPCAT_RGT.INSTANTIATE_OFFLINE(
     refresh_template_name   IN   VARCHAR2,
     site_name               IN   VARCHAR2,
     user_name               IN   VARCHAR2  := NULL,
     runtime_parm_id         IN   NUMBER    := -1e-130,
     next_date               IN   DATE      := SYSDATE,
     interval                IN   VARCHAR2  := 'SYSDATE + 1',
     use_default_gowner      IN   BOOLEAN   := TRUE)
   RETURN NUMBER;
```

## Parameters

*Table 21–63 INSTANTIATE_OFFLINE Function Parameters*

| Parameter | Description |
|---|---|
| refresh_template_name | Name of the deployment template to be instantiated. |
| site_name | Name of the remote site that is instantiating the deployment template. |
| user_name | Name of the authorized user who is instantiating the deployment template. |
| runtime_parm_id | If you have defined run-time parameter values using the INSERT_RUNTIME_PARMS procedure, then specify the identification used when creating the run-time parameters (the identification was retrieved by using the GET_RUNTIME_PARM_ID function). |
| next_date | Specifies the next refresh date value to be used when creating the refresh group. |
| interval | Specifies the refresh interval to be used when creating the refresh group. |
| use_default_gowner | If TRUE, then any materialized view groups created are owned by the default user PUBLIC. If FALSE, then any materialized view groups created are owned by the user performing the instantiation. |

## Exceptions

*Table 21–64 INSTANTIATE_OFFLINE Function Exceptions*

| Exception | Description |
|---|---|
| miss_refresh_template | Deployment template name specified is invalid or does not exist. |
| miss_user | Name of the authorized user is invalid or does not exist. Verify that the specified user is listed in the DBA_REPCAT_USER_AUTHORIZATIONS view. If user is not listed, then the specified user is not authorized to instantiate the target deployment template. |

## Returns

*Table 21–65 INSTANTIATE_OFFLINE Function Returns*

| Return Value | Description |
|---|---|
| <*system-generated number*> | Specifies the generated system number for the output_id when you select from the USER_REPCAT_TEMP_OUTPUT temporary view to retrieve the generated instantiation script. |

## INSTANTIATE_ONLINE Function

This function generates a script at the master site that is used to create the materialized view environment at the remote materialized view site while the materialized view site is connected to the master (that is, while the materialized view site is online). This generated script should be used at remote materialized view sites that are able to remain connected to the master site for an extended amount of time, as the instantiation process at the remote materialized view site might be lengthy (depending on the amount of data that is populated to the new materialized views). This function must be executed separately for each user instantiation.

The script generated by this function is stored in the `USER_REPCAT_TEMP_OUTPUT` temporary view and is used by several Oracle tools, including the Advanced Replication interface in Oracle Enterprise Manager, during the distribution of deployment templates. The number returned by this function is used to retrieve the appropriate information from the `USER_REPCAT_TEMP_OUTPUT` temporary view.

> **Note:** This function is for replication administrators who are instantiating for another user. Users wanting to perform their own instantiation should use the public version of the `INSTANTIATE_OFFLINE` function, described in "INSTANTIATE_OFFLINE Function" on page 21-43 section.

### Syntax

```
DBMS_REPCAT_RGT.INSTANTIATE_ONLINE(
   refresh_template_name   IN   VARCHAR2,
   site_name               IN   VARCHAR2  := NULL,
   user_name               IN   VARCHAR2  := NULL,
   runtime_parm_id         IN   NUMBER    := -1e-130,
   next_date               IN   DATE      := SYSDATE,
   interval                IN   VARCHAR2  := 'SYSDATE + 1',
   use_default_gowner      IN   BOOLEAN   := TRUE)
 RETURN NUMBER;
```

## Parameters

*Table 21–66 INSTANTIATE_ONLINE Function Parameters*

| Parameter | Description |
|---|---|
| refresh_template_name | Name of the deployment template to be instantiated. |
| site_name | Name of the remote site that is instantiating the deployment template. |
| user_name | Name of the authorized user who is instantiating the deployment template. |
| runtime_parm_id | If you have defined run-time parameter values using the INSERT_RUNTIME_PARMS procedure, then specify the identification used when creating the run-time parameters (the identification was retrieved by using the GET_RUNTIME_PARM_ID function). |
| next_date | Specifies the next refresh date value to be used when creating the refresh group. |
| interval | Specifies the refresh interval to be used when creating the refresh group. |
| use_default_gowner | If TRUE, then any materialized view groups created are owned by the default user PUBLIC. If FALSE, then any materialized view groups created are owned by the user performing the instantiation. |

## Exceptions

*Table 21–67 INSTANTIATE_ONLINE Function Exceptions*

| Exception | Description |
|---|---|
| miss_refresh_template | Specified deployment template name is invalid or does not exist. |
| miss_user | Name of the authorized user is invalid or does not exist. Verify that the specified user is listed in the DBA_REPCAT_USER_AUTHORIZATIONS view. If user is not listed, then the specified user is not authorized to instantiate the target deployment template. |
| bad_parms | Not all of the template parameters were populated by the defined user parameter values or template default values or both. The number of predefined values might not have matched the number of template parameters or a predefined value was invalid for the target parameter (that is, type mismatch). |

## Returns

*Table 21–68 INSTANTIATE_ONLINE Function Returns*

| Return Value | Description |
|---|---|
| *<system-generated number>* | Specifies the system-generated number for the output_id when you select from the USER_REPCAT_TEMP_OUTPUT temporary view to retrieve the generated instantiation script. |

## LOCK_TEMPLATE_EXCLUSIVE Procedure

When a deployment template is being updated or modified, use the `LOCK_TEMPLATE_EXCLUSIVE` procedure to prevent users from reading or instantiating the template.

The lock is released when a `ROLLBACK` or `COMMIT` is performed.

> **Note:** This procedure should be executed before you make any modifications to your deployment template.

### Syntax

```
DBMS_REPCAT_RGT.LOCK_TEMPLATE_EXCLUSIVE();
```

### Parameters

None

# LOCK_TEMPLATE_SHARED Procedure

The LOCK_TEMPLATE_SHARED procedure is used to make a specified deployment template "read-only." This procedure should be called before instantiating a template, as this ensures that nobody can change the deployment template while it is being instantiated.

The lock is released when a ROLLBACK or COMMIT is performed.

## Syntax

```
DBMS_REPCAT_RGT.LOCK_TEMPLATE_SHARED();
```

## Parameters

None

# 22

# DBMS_REPUTIL

DBMS_REPUTIL contains subprograms to generate shadow tables, triggers, and packages for table replication, as well as subprograms to generate wrappers for replication of standalone procedure invocations and packaged procedure invocations. This package is referenced only by the generated code.

This chapter contains this topic:


- Summary of DBMS_REPUTIL Subprograms

# Summary of DBMS_REPUTIL Subprograms

***Table 22–1 DBMS_REPUTIL Package Subprograms***

| Subprogram | Description |
|---|---|
| "REPLICATION_OFF Procedure" on page 22-3 | Modifies tables without replicating the modifications to any other sites in the replication environment, or disables row-level replication when using procedural replication. |
| "REPLICATION_ON Procedure" on page 22-4 | Reenables replication of changes after replication has been temporarily suspended. |
| "REPLICATION_IS_ON Function" on page 22-5 | Determines whether replication is running. |
| FROM_REMOTE Function on page 22-6 | Returns `TRUE` at the beginning of procedures in the internal replication packages, and returns `FALSE` at the end of these procedures. |
| "GLOBAL_NAME Function" on page 22-7 | Determines the global database name of the local database (the global name is the returned value). |
| "MAKE_INTERNAL_PKG Procedure" on page 22-8 | Synchronizes internal packages and tables in the replication catalog.<br>**Note:** Do not execute this procedure unless directed to do so by Oracle Support Services. |
| "SYNC_UP_REP Procedure" on page 22-9 | Synchronizes internal triggers and tables/materialized views in the replication catalog.<br>**Note:** Do not execute this procedure unless directed to do so by Oracle Support Services. |

## REPLICATION_OFF Procedure

This procedure enables you to modify tables without replicating the modifications to any other sites in the replication environment. It also disables row-level replication when using procedural replication. In general, you should suspend replication activity for all master groups in your replication environment before setting this flag.

### Syntax

```
DBMS_REPUTIL.REPLICATION_OFF();
```

### Parameters

None

# REPLICATION_ON Procedure

This procedure reenables replication of changes after replication has been temporarily suspended.

## Syntax

```
DBMS_REPUTIL.REPLICATION_ON();
```

## Parameters

None

## REPLICATION_IS_ON Function

This function determines whether replication is running. A returned value of TRUE indicates that the generated replication triggers are enabled. A return value of FALSE indicates that replication is disabled at the current site for the replication group.

The returning value of this function is set by calling the REPLICATION_ON or REPLICATION_OFF procedures in the DBMS_REPUTIL package.

### Syntax

```
DBMS_REPUTIL.REPLICATION_IS_ON()
  return BOOLEAN;
```

### Parameters

None

## FROM_REMOTE Function

This function returns `TRUE` at the beginning of procedures in the internal replication packages, and returns `FALSE` at the end of these procedures. You might need to check this function if you have any triggers that could be fired as the result of an update by an internal package.

### Syntax

```
DBMS_REPUTIL.FROM_REMOTE()
   return BOOLEAN;
```

### Parameters

None

## GLOBAL_NAME Function

This function determines the global database name of the local database (the global name is the returned value).

### Syntax

```
DBMS_REPUTIL.GLOBAL_NAME()
   return VARCHAR2;
```

### Parameters

None

# MAKE_INTERNAL_PKG Procedure

This procedure synchronizes the existence of an internal package with a table or materialized view in the replication catalog. If the table has replication support, then execute this procedure to create the internal package. If replication support does not exist, then this procedure destroys any related internal package. This procedure does not accept the storage table of a nested table.

**Caution:** Do not execute this procedure unless directed to do so by Oracle Support Services.

## Syntax

```
DBMS_REPUTIL.MAKE_INTERNAL_PKG (
   canon_sname    IN   VARCHAR2,
   canon_oname    IN   VARCHAR2);
```

## Parameters

***Table 22–2 MAKE_INTERNAL_PKG Procedure Parameters***

| Parameter | Description |
|---|---|
| canon_sname | Schema containing the table to be synchronized.<br>This parameter value must be canonically defined (capitalization must match object and must not be enclosed in double quotes). |
| canon_oname | Name of the table to be synchronized.<br>This parameter value must be canonically defined (capitalization must match object and must not be enclosed in double quotes). |

## SYNC_UP_REP Procedure

This procedure synchronizes the existence of an internal trigger with a table or materialized view in the replication catalog. If the table or materialized view has replication support, then execute this procedure to create the internal replication trigger. If replication support does not exist, then this procedure destroys any related internal trigger. This procedure does not accept the storage table of a nested table.

**Caution:** Do not execute this procedure unless directed to do so by Oracle Support Services.

### Syntax

```
DBMS_REPUTIL.SYNC_UP_REP (
   canon_sname    IN   VARCHAR2,
   canon_oname    IN   VARCHAR2);
```

### Parameters

***Table 22–3 SYNC_UP_REP Procedure Parameters***

| Parameter | Description |
|---|---|
| canon_sname | Schema containing the table or materialized view to be synchronized.<br>This parameter value must be canonically defined (capitalization must match object and must not be enclosed in double quotes). |
| canon_oname | Name of the table or materialized view to be synchronized.<br>This parameter value must be canonically defined (capitalization must match object and must not be enclosed in double quotes). |

# Part IV

## Replication Data Dictionary Reference

Part IV describes data dictionary views that provide information about your replication environment.

Part IV contains the following chapters:


- Chapter 23, "Replication Catalog Views"
- Chapter 24, "Replication Dynamic Performance Views"
- Chapter 25, "Materialized View and Refresh Group Views"
- Chapter 26, "Deferred Transaction Views"

# 23

# Replication Catalog Views

When you install replication capabilities at a site, Oracle installs the replication catalog, which consists of tables and views, at that site.

This chapter contains this topic:

- Summary of Replication Catalog Views

---

**Caution:** Do not modify the replication catalog tables directly. Instead, use the procedures provided in the `DBMS_REPCAT` package.

---

**See Also:** Chapter 10, "Monitoring a Replication Environment"

## Summary of Replication Catalog Views

Many data dictionary tables have three corresponding views:

- An ALL_ view displays all the information accessible to the current user, including information from the current user's schema as well as information from objects in other schemas, if the current user has access to those objects by way of grants of privileges or roles.
- A DBA_ view displays all relevant information in the entire database. DBA_ views are intended only for administrators. They can be accessed only by users with the SELECT ANY TABLE privilege. This privilege is assigned to the DBA role when Oracle is initially installed.
- A USER_ view displays all the information from the schema of the current user. No special privileges are required to query these views.

The columns of the ALL_, DBA_, and USER_ views corresponding to a single data dictionary table are usually nearly identical. Therefore, these views are described in full only once in this chapter (for the ALL_ view). The views are listed without the full description for DBA_ and USER_ views, but differences are noted.

As shown in Figure 23–1 on page 23-3, the replication catalog views are used by master sites and materialized view sites to determine such information as what objects are being replicated, where they are being replicated, and if any errors have occurred during replication. Table 23–1 on page 23-3 lists all of the replication catalog views.

***Figure 23–1 Replication Catalog Views and Replicated Objects***

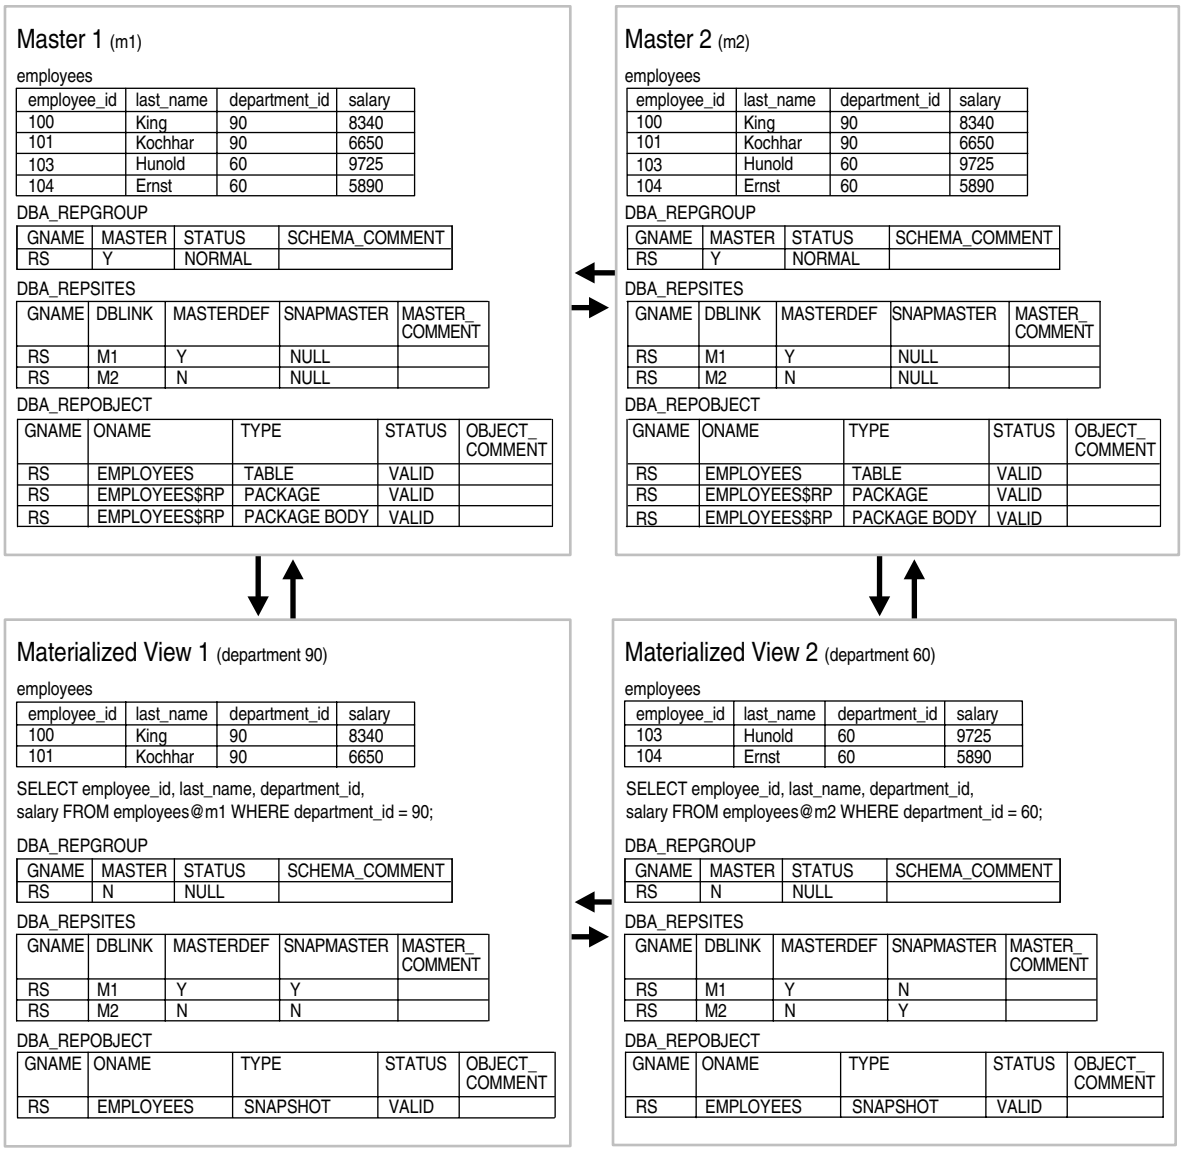


***Table 23–1 Replication Catalog Views***

| ALL_ Views | DBA_ Views | USER_ Views |
|---|---|---|
| N/A | DBA_REGISTERED_MVIEW_GROUPS | N/A |
| ALL_REPCAT_REFRESH_TEMPLATES | DBA_REPCAT_REFRESH_TEMPLATES | USER_REPCAT_REFRESH_TEMPLATES |
| ALL_REPCAT_TEMPLATE_OBJECTS | DBA_REPCAT_TEMPLATE_OBJECTS | USER_REPCAT_TEMPLATE_OBJECTS |
| ALL_REPCAT_TEMPLATE_PARMS | DBA_REPCAT_TEMPLATE_PARMS | USER_REPCAT_TEMPLATE_PARMS |
| ALL_REPCAT_TEMPLATE_SITES | DBA_REPCAT_TEMPLATE_SITES | USER_REPCAT_TEMPLATE_SITES |
| ALL_REPCAT_USER_AUTHORIZATIONS | DBA_REPCAT_USER_AUTHORIZATIONS | USER_REPCAT_USER_AUTHORIZATION |

*Table 23–1 (Cont.) Replication Catalog Views*

| ALL_ Views | DBA_ Views | USER_ Views |
|---|---|---|
| ALL_REPCAT_USER_PARM_VALUES | DBA_REPCAT_USER_PARM_VALUES | USER_REPCAT_USER_PARM_VALUES |
| ALL_REPCATLOG | DBA_REPCATLOG | USER_REPCATLOG |
| ALL_REPCOLUMN | DBA_REPCOLUMN | USER_REPCOLUMN |
| ALL_REPCOLUMN_GROUP | DBA_REPCOLUMN_GROUP | USER_REPCOLUMN_GROUP |
| ALL_REPCONFLICT | DBA_REPCONFLICT | USER_REPCONFLICT |
| ALL_REPDDL | DBA_REPDDL | USER_REPDDL |
| N/A | DBA_REPEXTENSIONS | N/A |
| ALL_REPGENOBJECTS | DBA_REPGENOBJECTS | USER_REPGENOBJECTS |
| ALL_REPGROUP | DBA_REPGROUP | USER_REPGROUP |
| ALL_REPGROUP_PRIVILEGES | DBA_REPGROUP_PRIVILEGES | USER_REPGROUP_PRIVILEGES |
| ALL_REPGROUPED_COLUMN | DBA_REPGROUPED_COLUMN | USER_REPGROUPED_COLUMN |
| ALL_REPKEY_COLUMNS | DBA_REPKEY_COLUMNS | USER_REPKEY_COLUMNS |
| ALL_REPOBJECT | DBA_REPOBJECT | USER_REPOBJECT |
| ALL_REPPARAMETER_COLUMN | DBA_REPPARAMETER_COLUMN | USER_REPPARAMETER_COLUMN |
| ALL_REPPRIORITY | DBA_REPPRIORITY | USER_REPPRIORITY |
| ALL_REPPRIORITY_GROUP | DBA_REPPRIORITY_GROUP | USER_REPPRIORITY_GROUP |
| ALL_REPPROP | DBA_REPPROP | USER_REPPROP |
| ALL_REPRESOL_STATS_CONTROL | DBA_REPRESOL_STATS_CONTROL | USER_REPRESOL_STATS_CONTROL |
| ALL_REPRESOLUTION | DBA_REPRESOLUTION | USER_REPRESOLUTION |
| ALL_REPRESOLUTION_METHOD | DBA_REPRESOLUTION_METHOD | USER_REPRESOLUTION_METHOD |
| ALL_REPRESOLUTION_STATISTICS | DBA_REPRESOLUTION_STATISTICS | USER_REPRESOLUTION_STATISTICS |
| ALL_REPSITES | DBA_REPSITES | USER_REPSITES |
| N/A | DBA_REPSITES_NEW | N/A |

## DBA_REGISTERED_MVIEW_GROUPS

DBA_REGISTERED_MVIEW_GROUPS lists all the registered materialized view groups at the master site or master materialized view site.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| NAME | VARCHAR2(30) | - | Name of the materialized view replication group |
| MVIEW_SITE | VARCHAR2(128) | - | Site of the materialized view replication group |
| GROUP_COMMENT | VARCHAR2(80) | - | Description of the materialized view replication group |
| VERSION | VARCHAR2(8) | - | Oracle version of the materialized view replication group<br>**Note:** Oracle9*i* Database or later materialized view groups show Oracle8. |
| FNAME | VARCHAR2(30) | - | Name of the flavor of the materialized view group |
| OWNER | VARCHAR2(30) | - | Owner of the materialized view replication group |

# ALL_REPCAT_REFRESH_TEMPLATES

Contains global information about each deployment template accessible to the current user, such as the template name, template owner, what refresh group the template objects belong to, and the type of template (private or public).

When the DBA adds materialized view definitions to the template container, the DBA references the appropriate REFRESH_TEMPLATE_NAME. Any materialized views added to a specific template are added to the refresh group specified in REFRESH_GROUP_NAME.

Furthermore, deployment templates created as public are available to all users who can connect to the master site. Deployment templates created as private are limited to those users listed in the ALL_REPCAT_USER_AUTHORIZATIONS view.

## Related Views:

- DBA_REPCAT_REFRESH_TEMPLATES describes all deployment templates in the database.
- USER_REPCAT_REFRESH_TEMPLATES describes all deployment templates owned by the current user.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| REFRESH_TEMPLATE_NAME | VARCHAR2(30) | - | Name of the deployment template. |
| OWNER | VARCHAR2(30) | - | Owner of the deployment template. |
| REFRESH_GROUP_NAME | VARCHAR2(30) | - | Name of the refresh group to which the template objects are added during the instantiation process. |
| TEMPLATE_COMMENT | VARCHAR2(2000) | - | User supplied comment. |
| PUBLIC_TEMPLATE | VARCHAR2(1) | - | If Y then the deployment template is public.<br>If N then the deployment template is private. |

## ALL_REPCAT_TEMPLATE_OBJECTS

Contains the individual object definitions that are contained in each deployment template accessible to the current user. Individual objects are added to a template by specifying the target template in REFRESH_TEMPLATE_NAME.

DDL_TEXT can contain variables to create parameterized templates. Variables are created by placing a colon (:) at the beginning of the variable name (for example, :region). Templates that use parameters allow for greater flexibility during the template instantiation process (that is, in defining data sets specific for a materialized view site).

When the object is added to the template, the specified DDL is examined and if any parameters have been defined, Oracle automatically adds the parameter to the ALL_REPCAT_TEMPLATE_PARMS view.

### Related Views:

- DBA_REPCAT_TEMPLATE_OBJECTS describes the object definitions for all deployment templates in the database.
- USER_REPCAT_TEMPLATE_OBJECTS describes the object definitions for each deployment template owned by the current user.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| REFRESH_TEMPLATE_NAME | VARCHAR2(30) | NOT NULL | The name of the deployment template. |
| OBJECT_NAME | VARCHAR2(30) | NOT NULL | The name of the deployment template object. |
| OBJECT_TYPE | VARCHAR2(17) | - | The object type of the deployment template object:<br>FUNCTION MATERIALIZED VIEW<br>INDEX SYNONYM<br>INDEXTYPE TABLE<br>OPERATOR TRIGGER<br>PACKAGE TYPE<br>PACKAGE BODY TYPE BODY<br>PROCEDURE VIEW |
| DDL_NUM | NUMBER | NOT NULL | Indicates the order in which to execute the DDL statements stored in the DDL_TEXT column when multiple DDL statements are used to create the object. |
| DDL_TEXT | CLOB(4000) | - | The DDL that is executed to create the deployment template object. |
| MASTER_ROLLBACK_SEGMENT | VARCHAR2(30) | - | The name of the rollback segment that is used during the instantiation of the deployment template object. |
| DERIVED_FROM_SNAME | VARCHAR2(30) | - | If applicable, displays the schema that contains the object from which the template object was created. |
| DERIVED_FROM_ONAME | VARCHAR2(30) | - | If applicable, displays the name of the object from which the template object was created. |
| FLAVOR_ID | NUMBER | - | The flavor ID of the deployment template object. |

Because the DDL_TEXT column is defined as a CLOB, you receive an error if you simply try to perform a SELECT on the ALL_REPCAT_TEMPLATE_OBJECTS view. If you do not need to see the object DDL, then use the following select statement (be sure to exclude the DDL_TEXT parameter):

```
SELECT REFRESH_TEMPLATE_NAME, OBJECT_NAME, OBJECT_TYPE, MASTER_ROLLBACK_SEG,
FLAVOR_ID FROM DBA_REPCAT_TEMPLATE_OBJECTS;
```

The following script uses cursors and the DBMS_LOB package to view the entire contents of the ALL_REPCAT_TEMPLATE_OBJECTS view. Use this script to view the entire contents of the ALL_REPCAT_TEMPLATE_OBJECTS view, including the DDL_TEXT column:

```
SET SERVEROUTPUT ON

DECLARE
  CURSOR mycursor IS
        SELECT REFRESH_TEMPLATE_NAME, OBJECT_NAME, OBJECT_TYPE, DDL_TEXT,
          MASTER_ROLLBACK_SEG, FLAVOR_ID
          FROM DBA_REPCAT_TEMPLATE_OBJECTS;
  tempstring VARCHAR2(1000);
  len NUMBER;
BEGIN
  FOR myrec IN mycursor LOOP
        len := DBMS_LOB.GETLENGTH(myrec.ddl_text);
        DBMS_LOB.READ(myrec.ddl_text, len, 1, tempstring);
        DBMS_OUTPUT.PUT_LINE(myrec.refresh_template_name||' '||
          myrec.object_name||' '||myrec.object_type||' '||tempstring||' '||
          myrec.master_rollback_seg||' '||myrec.flavor_id);
  END LOOP;
END;
/
```

**See Also:** *Oracle Database Advanced Application Developer's Guide* for more information about using cursors. Also, see *Oracle Database SecureFiles and Large Objects Developer's Guide* for more information about using the DBMS_LOB package and LOBs in general.

## ALL_REPCAT_TEMPLATE_PARMS

Contains parameters defined in the object DDL for all templates accessible to the current user. When an object is added to a template, the DDL is examined for variables. Any found parameters are automatically added to this view.

You can also define default parameter values and a prompt string in this view. These can make the templates easier to use during the instantiation process.

**See Also:** ALL_REPCAT_TEMPLATE_OBJECTS on page 23-7

### Related Views:

- DBA_REPCAT_TEMPLATE_PARMS describes the template parameters for all deployment templates in the database.
- USER_REPCAT_TEMPLATE_PARMS describes the template parameters for all deployment templates owned by the current user.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| REFRESH_TEMPLATE_NAME | VARCHAR2(30) | NOT NULL | The name of the deployment template. |
| OWNER | VARCHAR2(30) | NOT NULL | The owner of the deployment template. |
| REFRESH_GROUP_NAME | VARCHAR2(30) | NOT NULL | Name of the refresh group to which the template objects are added to during the instantiation process. |
| TEMPLATE_COMMENTS | VARCHAR2(2000) | - | User specified comments. |
| PUBLIC_TEMPLATE | VARCHAR2(1) | - | If Y then the deployment template is public.<br>If N then the deployment template is private. |
| PARAMETER_NAME | VARCHAR2(30) | NOT NULL | The name of the parameter. |
| DEFAULT_PARM_VALUE | CLOB(4000) | - | The default parameter value. |
| PROMPT_STRING | VARCHAR2(2000) | - | The prompt string for the parameter. |
| USER_OVERRIDE | VARCHAR2(1) | - | If Y then the user can override the default parameter value.<br>If N then the user cannot override the default parameter value. |

Because the DEFAULT_PARM_VALUE column is defined as a CLOB, you receive an error if you simply try to perform a SELECT on the ALL_REPCAT_TEMPLATE_PARMS view. If you do not need to see the default parameter value, then use the following select statement (be sure to exclude DEFAULT_PARM_VALUE):

```
SELECT REFRESH_TEMPLATE_NAME, OWNER, REFRESH_GROUP_NAME, TEMPLATE_COMMENT,
  PUBLIC_TEMPLATE, PARAMETER_NAME, PROMPT_STRING, USER_OVERRIDE
  FROM DBA_REPCAT_TEMPLATE_PARMS;
```

The following script uses cursors and the DBMS_LOB package to view the entire contents of the ALL_REPCAT_TEMPLATE_PARMS view. Use this script to view the entire contents of the ALL_REPCAT_TEMPLATE_PARMS view, including the DEFAULT_PARM_VALUE column:

```
SET SERVEROUTPUT ON

DECLARE
  CURSOR mycursor IS
```

```
        SELECT REFRESH_TEMPLATE_NAME, OWNER, REFRESH_GROUP_NAME,
          TEMPLATE_COMMENT, PUBLIC_TEMPLATE, PARAMETER_NAME, DEFAULT_PARM_VALUE,
          PROMPT_STRING, USER_OVERRIDE
          FROM DBA_REPCAT_TEMPLATE_PARMS;
  tempstring VARCHAR2(1000);
  len NUMBER;
BEGIN
  FOR myrec IN mycursor LOOP
        len := DBMS_LOB.GETLENGTH(myrec.default_parm_value);
        DBMS_LOB.READ(myrec.default_parm_value, len, 1, tempstring);
        DBMS_OUTPUT.PUT_LINE(myrec.refresh_template_name||' '||
          myrec.owner||' '||myrec.refresh_group_name||' '||
          myrec.template_comment||' '||myrec.public_template||' '||
          myrec.parameter_name||' '||tempstring||' '||myrec.prompt_string||' '||
          myrec.user_override);
  END LOOP;
END;
/
```

**See Also:** *Oracle Database Advanced Application Developer's Guide* for more information about using cursors. Also, see *Oracle Database SecureFiles and Large Objects Developer's Guide* for more information about using the `DBMS_LOB` package and LOBs in general.

## ALL_REPCAT_TEMPLATE_SITES

Contains information about the current status of template instantiation among the sites of an enterprise network. This view contains information about instantiation sites for deployment templates that are accessible to the current user. Specifically, the DBA can monitor the installation and deletion of templates at specific sites.

### Related Views:

- DBA_REPCAT_TEMPLATE_SITES describes all remote instantiation sites for all templates in the database.
- USER_REPCAT_TEMPLATE_SITES describes remote instantiation sites for all templates owned by the current user.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| REFRESH_TEMPLATE_NAME | VARCHAR2(30) | NOT NULL | Name of the deployment template. |
| REFRESH_GROUP_NAME | VARCHAR2(30) | - | Name of the refresh group to which template objects are added during the instantiation process. |
| TEMPLATE_OWNER | VARCHAR2(30) | - | Name of the user who is considered the owner of the deployment template. |
| USER_NAME | VARCHAR2(30) | NOT NULL | The name of the user who instantiated the deployment template. |
| SITE_NAME | VARCHAR2(128) | - | Target materialized view site of the deployment template. |
| REPAPI_SITE_NAME | VARCHAR2(128) | - | This column is intended for use in a future release of Oracle. |
| STATUS | VARCHAR2(10) | - | Displays the status of the deployment template at the target materialized view site:<br>0 = Not Installed<br>1 = Installed<br>-1 = Installed with errors |
| INSTANTIATION_DATE | DATE | - | Displays when the template was instantiated. Is NULL if the template has not yet been instantiated. |

# ALL_REPCAT_USER_AUTHORIZATIONS

Lists the authorized users for private deployment templates accessible to the current user. Users listed in this view have the ability to instantiate the specified template. Users not listed in this view cannot instantiate the deployment template.

## Related Views:

- DBA_REPCAT_USER_AUTHORIZATIONS lists the authorized users for all the private deployment templates in the database.
- USER_REPCAT_USER_AUTHORIZATION lists the authorized users for private deployment templates owned by the current user.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| REFRESH_TEMPLATE_NAME | VARCHAR2(30) | NOT NULL | Name of the deployment template that a user has been authorized to instantiate. |
| OWNER | VARCHAR2(30) | NOT NULL | Name of the owner of the deployment template. |
| REFRESH_GROUP_NAME | VARCHAR2(30) | NOT NULL | Name of the refresh group to which template objects are added during the instantiation process. |
| TEMPLATE_COMMENT | VARCHAR2(2000) | - | User specified comment. |
| PUBLIC_TEMPLATE | VARCHAR2(1) | - | If Y then the deployment template is public.<br>If N then the deployment template is private. |
| USER_NAME | VARCHAR2(30) | NOT NULL | Name of the user who has been authorized to instantiate the deployment template. |

## ALL_REPCAT_USER_PARM_VALUES

This view describes the template parameters for all deployment templates accessible to the current user. The DBA has the option of building a table of user parameters before distributing the template for instantiation. When a template is instantiated by a specified user, the values stored in the ALL_REPCAT_USER_PARM_VALUES view for the specified user are used automatically.

### Related Views:

- DBA_REPCAT_USER_PARM_VALUES describes the template parameters for all deployment templates in the database.
- USER_REPCAT_USER_PARM_VALUES describes the template parameters for all deployment templates owned by the current user.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| REFRESH_TEMPLATE_NAME | VARCHAR2(30) | NOT NULL | The name of the deployment template for which a user parameter value has been defined. |
| OWNER | VARCHAR2(30) | NOT NULL | The name of the owner of the deployment template. |
| REFRESH_GROUP_NAME | VARCHAR2(30) | NOT NULL | Name of the refresh group to which the template objects are added to during the instantiation process. |
| TEMPATE_COMMENT | VARCHAR2(2000) | - | User specified comment. |
| PUBLIC_TEMPLATE | VARCHAR2(1) | - | If Y then the deployment template is public.<br>If N then the deployment template is private. |
| PARAMETER_NAME | VARCHAR2(30) | NOT NULL | The name of the parameter for which a user parameter value has been defined. |
| DEFAULT_PARM_VALUE | CLOB(4000) | - | The default value for the parameter. |
| PROMPT_STRING | VARCHAR2(2000) | - | The prompt string for the parameter. |
| PARM_VALUE | CLOB(4000) | - | The parameter value that has been defined for the specified user. |
| USER_NAME | VARCHAR2(30) | NOT NULL | The user name of the user for whom the specified parameter value has been defined. |

Because DEFAULT_PARM_VALUE and PARM_VALUE columns are defined as CLOB data types, you receive an error if you simply try to perform a SELECT on the ALL_REPCAT_USER_PARM_VALUES view. If you do not need to see the default or user parameter values, then use the following select statement (be sure to exclude DEFAULT_PARM_VALUE and PARM_VALUE):

```
SELECT REFRESH_TEMPLATE_NAME, OWNER, REFRESH_GROUP_NAME, TEMPLATE_COMMENT,
  PUBLIC_TEMPLATE, PARAMETER_NAME, PROMPT_STRING, USER_NAME
  FROM DBA_REPCAT_USER_PARM_VALUES;
```

The following script uses cursors and the DBMS_LOB package to view the entire contents of the ALL_REPCAT_USER_PARM_VALUES view. Use this script to view the entire contents of the ALL_REPCAT_TEMPLATE_PARMS view, including the DEFAULT_PARM_VALUE and PARM_VALUE columns:

```
SET SERVEROUTPUT ON
```

```
DECLARE
  CURSOR mycursor IS
        SELECT REFRESH_TEMPLATE_NAME, OWNER, REFRESH_GROUP_NAME,
          TEMPLATE_COMMENT, PUBLIC_TEMPLATE, PARAMETER_NAME, DEFAULT_PARM_VALUE,
          PROMPT_STRING, PARM_VALUE, USER_NAME
          FROM DBA_REPCAT_USER_PARM_VALUES;
  tempstring VARCHAR2(1000);
  tempstring2 varchar2(1000);
  len NUMBER;
BEGIN
  FOR myrec IN mycursor LOOP
        len := DBMS_LOB.GETLENGTH(myrec.default_parm_value);
        DBMS_LOB.READ(myrec.default_parm_value, len, 1, tempstring);
        DBMS_OUTPUT.PUT_LINE(myrec.refresh_template_name||' '||
          myrec.owner||' '||myrec.refresh_group_name||' '||
          myrec.template_comment||' '||myrec.public_template||' '||
          myrec.parameter_name||' '||tempstring||' '||myrec.prompt_string||' '||
          tempstring2||' '||myrec.user_name);
  END LOOP;
END;
/
```

**See Also:** *Oracle Database Advanced Application Developer's Guide* for more information about using cursors. Also, see *Oracle Database SecureFiles and Large Objects Developer's Guide* for more information about using the `DBMS_LOB` package and LOBs in general.

## ALL_REPCATLOG

Contains the interim status of any asynchronous administrative requests and any error messages generated at each master site. All messages encountered while executing a request are eventually transferred to the ALL_REPCATLOG view at the master site that originated the request. If an administrative request completes without error, then ultimately all traces of this request are removed from the ALL_REPCATLOG view. This view contains administrative requests and error messages that are accessible to the current user.

### Related Views:

- DBA_REPCATLOG describes the status for all asynchronous administrative requests and all error messages in the database.
- USER_REPCATLOG describes the status for all asynchronous administrative requests and all error messages owned by the current user.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| ID | NUMBER | - | A sequence number. Together, the ID and SOURCE columns identify all log records at all master sites that pertain to a single administrative request. |
| SOURCE | VARCHAR2(128) | - | Location where the request originated. |
| USERID | VARCHAR2(30) | - | Name of the user making the request. |
| TIMESTAMP | DATE | - | When the request was made. |
| ROLE | VARCHAR2(9) | - | Indicates if site is the master definition site (masterdef) or a master site (master). |
| MASTER | VARCHAR2(128) | - | If the role is 'masterdef' and the task is remote, then indicates which master site is performing the task. |
| SNAME | VARCHAR2(30) | - | The name of the schema for the replicated object, if applicable. |
| REQUEST | VARCHAR2(29) | - | The name of the DBMS_REPCAT administrative procedure that was run. |
| ONAME | VARCHAR2(30) | - | The name of the replicated object, if applicable. |
| TYPE | VARCHAR2(12) | - | The type of replicated object:<br>FUNCTION MATERIALIZED VIEW<br>INDEX SYNONYM<br>INDEXTYPE TABLE<br>OPERATOR TRIGGER<br>PACKAGE TYPE<br>PACKAGE BODY TYPE BODY<br>PROCEDURE VIEW |
| STATUS | VARCHAR2(14) | - | The status of the administrative request: READY, DO_CALLBACK, AWAIT_CALLBACK, or ERROR. |
| MESSAGE | VARCHAR2(200) | - | Any error message that has been returned. |
| ERRNUM | NUMBER | - | The Oracle error number for the message. |
| GNAME | VARCHAR2(30) | - | The name of the replication group. |

# ALL_REPCOLUMN

Lists the replicated columns for the tables accessible to the current user.

If the table contains a column object, then this view displays a placeholder for the type and one row for each type attribute. If the table contains a nested table, then this view displays the storage table for the nested table as an independent table. If a table is an object table, then this view displays the hidden object identifier column.

## Related Views:

- DBA_REPCOLUMN describes the replicated columns for all the tables in the database.
- USER_REPCOLUMN describes the replicated columns for all the tables owned by the current user.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| SNAME | VARCHAR2(30) | NOT NULL | The name of the object owner. |
| ONAME | VARCHAR2(30) | NOT NULL | The name of the object. |
| TYPE | VARCHAR2(8) | - | The type of the object, either MATERIALIZED VIEW or TABLE. |
| CNAME | VARCHAR2(4000) | - | The name of the replicated column. |
| ID | NUMBER | - | The ID number of the replicated column. |
| POS | NUMBER | - | The ordering of the replicated column. |
| COMPARE_OLD_ON_DELETE | VARCHAR2(1) | - | Indicates whether Oracle compares the old value of the column in replicated deletes. |
| COMPARE_OLD_ON_UPDATE | VARCHAR2(1) | - | Indicates whether Oracle compares the old value of the column in replicated updates. |
| SEND_OLD_ON_DELETE | VARCHAR2(1) | - | Indicates whether Oracle sends the old value of the column in replicated deletes. |
| SEND_OLD_ON_UPDATE | VARCHAR2(1) | - | Indicates whether Oracle sends the old value of the column in replicated updates. |
| CTYPE | VARCHAR2(30) | - | Displays the column type. For user-defined types, displays the user-defined type name. |
| CTYPE_TOID | RAW(16) | - | If user-defined type, displays the object identifier (OID) of the type. Otherwise, this field is NULL. |
| CTYPE_OWNER | VARCHAR2(30) | - | If user-defined type, displays the owner of a user-defined type. Otherwise, this field is NULL. |
| CTYPE_HASHCODE | VARCHAR2(34) | - | If user-defined type, displays the type's hashcode. Otherwise, this field is NULL. |
| CTYPE_MOD | VARCHAR2(3) | - | Displays REF for REF columns. Otherwise, this field is NULL. |
| DATA_LENGTH | VARCHAR2(40) | - | Displays the length of the column in bytes. |
| DATA_PRECISION | VARCHAR2(40) | - | Displays the column precision in terms of decimal digits for NUMBER columns or binary digits for FLOAT columns. |
| DATA_SCALE | VARCHAR2(40) | - | Displays the digits to right of decimal point in a number. |

| Column | Data Type | NULL | Description |
|---|---|---|---|
| NULLABLE | VARCHAR2(1) | - | Indicates if the column allow NULL values. |
| CHARACTER_SET_NAME | VARCHAR2(44) | - | If applicable, displays the name of character set for the column. |
| TOP | VARCHAR2(30) | - | Displays the top column for an attribute in a column object. For example, in the oe.customers table, cust_address is a column object and street_address is one of its attributes. For the street_address attribute, cust_address is the TOP column.<br>For built-in data types, this field is NULL |
| CHAR_LENGTH | NUMBER | - | Displays the length of the column in characters. This value only applies to the following data types:<br>■ CHAR<br>■ VARCHAR2<br>■ NCHAR<br>■ NVARCHAR2 |
| CHAR_USED | VARCHAR2(1) | - | B indicates that the column uses BYTE length semantics. C indicates that the column uses CHAR length semantics. NULL indicates that the data type is not any of the following:<br>■ CHAR<br>■ VARCHAR2<br>■ NCHAR<br>■ NVARCHAR2 |

# ALL_REPCOLUMN_GROUP

Describes the column groups for each replicated table accessible to the current user.

## Related Views:

- DBA_REPCOLUMN_GROUP describes the column groups for all the tables in the database.
- USER_REPCOLUMN_GROUP describes the column groups for all the tables owned by the current user.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| SNAME | VARCHAR2(30) | NOT NULL | The name of the schema containing the replicated table. |
| ONAME | VARCHAR2(30) | NOT NULL | The name of the replicated table. |
| GROUP_NAME | VARCHAR2(30) | NOT NULL | The column group name. |
| GROUP_COMMENT | VARCHAR2(80) | - | Any user-supplied comments. |

**Note:** The SNAME column is not present in the USER_REPCOLUMN_GROUP view.

## ALL_REPCONFLICT

Contains the name of each table accessible to the current user for which a conflict resolution method has been defined and the type of conflict that the method is used to resolve.

### Related Views:

- DBA_REPCONFLICT describes the conflict resolution method for all the tables in the database on which a conflict resolution method has been defined.
- USER_REPCONFLICT describes the conflict resolution method for all the tables owned by the current user on which a conflict resolution method has been defined.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| SNAME | VARCHAR2(30) | NOT NULL | The name of the schema containing the replicated table. |
| ONAME | VARCHAR2(30) | NOT NULL | The name of the table for which a conflict resolution method has been defined. |
| CONFLICT_TYPE | VARCHAR2(10) | - | The type of conflict that the conflict resolution method is used to resolve: delete, uniqueness, or update. |
| REFERENCE_NAME | VARCHAR2(30) | NOT NULL | The object to which the method applies. For delete conflicts, this is the table name. For uniqueness conflicts, this is the constraint name. For update conflicts, this is the column group name. |

**Note:** The SNAME column is not present in the USER_REPCONFLICT view.

# ALL_REPDDL

Contains the DDL for each replication object accessible to the current user.

## Related Views:

- DBA_REPDDL contains the DDL for each replicated object in the database.
- USER_REPDDL contains the DDL for each replicated object owned by the current user.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| LOG_ID | NUMBER | - | Identifying number of the ALL_REPCATLOG record. |
| SOURCE | VARCHAR2(128) | - | Name of the database at which the request originated. |
| ROLE | VARCHAR2(1) | - | If Y then this database is the master definition site (masterdef) for the request.<br>If N then this database is a master site. |
| MASTER | VARCHAR2(128) | - | Name of the database that processes this request. |
| LINE | NUMBER(38) | - | Ordering of records within a single request. |
| TEXT | VARCHAR2(2000) | - | Portion of an argument or DDL text. |
| DDL_NUM | NUMBER(38) | - | Indicates the order in which to execute the DDL statements stored in the TEXT column when multiple DDL statements are used. |

## ALL_REPGENOBJECTS

Describes each object accessible to the current user that was generated to support replication.

### Related Views:

- DBA_REPGENOBJECTS describes each object in the database that was generated to support replication.
- USER_REPGENOBJECTS describes each object owned by the current user that was generated to support replication.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| SNAME | VARCHAR2(30) | - | The name of the replicated schema. |
| ONAME | VARCHAR2(30) | - | The name of the generated object. |
| TYPE | VARCHAR2(12) | - | The type of the generated object, either PACKAGE, PACKAGE BODY, TRIGGER, or INTERNAL PACKAGE. |
| BASE_SNAME | VARCHAR2(30) | - | The base object's owner. |
| BASE_ONAME | VARCHAR2(30) | - | The object for which this object was generated. |
| BASE_TYPE | VARCHAR2(12) | - | The type of the base object. |
| PACKAGE_PREFIX | VARCHAR2(30) | - | The prefix for the package wrapper. |
| PROCEDURE_PREFIX | VARCHAR2(30) | - | The procedure prefix for the package wrapper. |
| DISTRIBUTED | VARCHAR2(1) | - | This column is obsolete. |
| REASON | VARCHAR2(30) | - | The reason the object was generated. |

# ALL_REPGROUP

Describes all of the replication groups that are accessible to the current user. The members of each replication group are listed in a different view: ALL_REPOBJECT.

## Related Views:

- DBA_REPGROUP describes all of the replication groups in the database that are being replicated.
- USER_REPGROUP describes all of the replication groups owned by the current user that are being replicated.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| SNAME | VARCHAR2(30) | NOT NULL | The name of the replicated schema. This column is obsolete. |
| MASTER | VARCHAR2(1) | - | Y indicates that the current site is a master site.<br>N indicates the current site is a materialized view site. |
| STATUS | VARCHAR2(9) | - | Used at master sites only. Status can be: normal, quiescing, or quiesced. |
| SCHEMA_COMMENT | VARCHAR2(80) | - | Any user-supplied comments. |
| GNAME | VARCHAR2(30) | NOT NULL | The name of the replication group. |
| FNAME | VARCHAR2(30) | - | Flavor name. |
| RPC_PROCESSING_DISABLED | VARCHAR2(1) | - | N indicates that this site can receive and apply deferred remote procedure calls (RPCs).<br>Y indicates that this site cannot receive and apply deferred remote procedure calls (RPCs). |
| OWNER | VARCHAR2(30) | NOT NULL | Owner of the replication group. |

## ALL_REPGROUP_PRIVILEGES

Contains information about the users who are registered for privileges in replication groups. Shows only those replication groups accessible to the current user.

### Related Views:

- DBA_REPGROUP_PRIVILEGES contains information about the users who are registered for privileges in all the replication groups in the database.
- USER_REPGROUP_PRIVILEGES contains information about the users who are registered for privileges in the replication groups owned by the current user.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| USERNAME | VARCHAR2(30) | NOT NULL | Displays the name of the user. |
| GNAME | VARCHAR2(30) | - | Displays the name of the replication group. |
| CREATED | DATE | NOT NULL | Displays the date that the replication group was registered. |
| RECEIVER | VARCHAR2(1) | - | Indicates whether the user has receiver privileges. |
| PROXY_SNAPADMIN | VARCHAR2(1) | - | Indicates whether the user has proxy_ snapadmin privileges. |
| OWNER | VARCHAR2(30) | - | Owner of the replication group. |

# ALL_REPGROUPED_COLUMN

Describes all of the columns that comprise the column groups for each table accessible to the current user.

## Related Views:

- DBA_REPGROUPED_COLUMN describes all of the columns that comprise the column groups for each table in the database.
- USER_REPGROUPED_COLUMN describes all of the columns that comprise the column groups for each table owned by the current user.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| SNAME | VARCHAR2(30) | NOT NULL | The name of the schema containing the replicated table. |
| ONAME | VARCHAR2(30) | NOT NULL | The name of the replicated table. |
| GROUP_NAME | VARCHAR2(30) | NOT NULL | The name of the column group. |
| COLUMN_NAME | VARCHAR2(30) | NOT NULL | The name of the column in the column group. |

**Note:** The SNAME column is not present in the USER_REPGROUPED_COLUMN version of the view.

## ALL_REPKEY_COLUMNS

Describes the replication key column(s) accessible to the current user in each table.

The replication key column(s) is an alternate column or group of columns, instead of the primary key, used to determine which columns of a table to compare when using row-level replication. You can set the replication key columns using the SET_COLUMNS procedure in the DBMS_REPCAT package.

The following types of columns cannot be replication key columns:

- LOB or LOB attribute of a column object
- Collection or collection attribute of a column object
- REF
- An entire column object

**See Also:** "SET_COLUMNS Procedure" on page 18-91

### Related Views:

- DBA_REPKEY_COLUMNS describes the replication key column(s) in each table in the database.
- USER_REPKEY_COLUMNS describes the replication key column(s) in each table owned by the current user.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| SNAME | VARCHAR2(30) | NOT NULL | Owner of the replicated table. |
| ONAME | VARCHAR2(30) | NOT NULL | Name of the replicated table. |
| COL | VARCHAR2(4000) | - | Replication key column(s) in the table. |

# ALL_REPOBJECT

Contains information about the objects in each replication group accessible to the current user. An object can belong to only one replication group. A replication group can span multiple schemas.

## Related Views:

- DBA_REPOBJECT contains information about the objects in each replication group in the database.
- USER_REPOBJECT contains information about the objects owned by the current user in each replication group.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| SNAME | VARCHAR2(30) | - | The name of the schema containing the replicated object. |
| ONAME | VARCHAR2(30) | - | The name of the replicated object. |
| TYPE | VARCHAR2(16) | - | The type of replicated object:<br>FUNCTION MATERIALIZED VIEW<br>INDEX SYNONYM<br>INDEXTYPE TABLE<br>OPERATOR TRIGGER<br>PACKAGE TYPE<br>PACKAGE BODY TYPE BODY<br>PROCEDURE VIEW |
| STATUS | VARCHAR2(10) | - | CREATE indicates that Oracle is applying user supplied or Oracle-generated DDL to the local database in an attempt to create the object locally. When a local replica exists, Oracle COMPAREs the replica with the master definition to ensure that they are consistent. When creation or comparison complete successfully, Oracle updates the status to VALID. Otherwise, it updates the status to ERROR. If you drop an object, then Oracle updates its status to DROPPED before deleting the row from the ALL_REPOBJECT view. |
| GENERATION_STATUS | VARCHAR2(9) | - | Specifies whether the object must generate replication packages. |
| ID | NUMBER | - | The identifier of the local database object, if one exists. |
| OBJECT_COMMENT | VARCHAR2(80) | - | Any user supplied comments. |
| GNAME | VARCHAR2(30) | - | The name of the replication group to which the object belongs. |
| MIN_COMMUNICATION | VARCHAR2(1) | - | If Y then use minimum communication for an update.<br>If N then send all old and all new values for an update. |
| REPLICATION_TRIGGER_EXISTS | VARCHAR2(1) | - | If Y then internal replication trigger exists.<br>If N then internal replication trigger does not exist. |

| Column | Data Type | NULL | Description |
|---|---|---|---|
| INTERNAL_PACKAGE_EXISTS | VARCHAR2(1) | - | If Y then internal package exists.<br>If N then internal package does not exist. |
| GROUP_OWNER | VARCHAR2(30) | - | Owner of the replication group. |
| NESTED_TABLE | VARCHAR2(1) | - | If Y then the replicated object is the storage table of a nested table.<br>If N then the replicated object is not a storage table. |

# ALL_REPPARAMETER_COLUMN

In addition to the information contained in the ALL_REPRESOLUTION view, the ALL_REPPARAMETER_COLUMN view contains information about the columns that are used to resolve conflicts for each replicated table accessible to the current user. These are the column values that are passed as the list_of_column_names argument to the ADD_*conflicttype*_RESOLUTION procedures in the DBMS_REPCAT package.

## Related Views:

- DBA_REPPARAMETER_COLUMN contains information about the columns that are used to resolve conflicts for each replicated table in the database.
- USER_REPPARAMETER_COLUMN contains information about the columns that are used to resolve conflicts for each replicated table owned by the current user.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| SNAME | VARCHAR2(30) | NOT NULL | The name of the schema containing the replicated table. |
| ONAME | VARCHAR2(30) | NOT NULL | The name of the replicated table. |
| CONFLICT_TYPE | VARCHAR2(10) | - | The type of conflict that the method is used to resolve: delete, uniqueness, or update. |
| REFERENCE_NAME | VARCHAR2(30) | NOT NULL | The object to which the method applies. For delete conflicts, this is the table name. For uniqueness conflicts, this is the constraint name. For update conflicts, this is the column group name. |
| SEQUENCE_NO | NUMBER | NOT NULL | The order in which resolution methods are applied, with 1 applied first. |
| METHOD_NAME | VARCHAR2(80) | NOT NULL | The name of an Oracle-supplied conflict resolution method. For user-supplied methods, this value is 'user function'. |
| FUNCTION_NAME | VARCHAR2(92) | NOT NULL | For methods of type 'user function', the name of the user-supplied conflict resolution method. |
| PRIORITY_GROUP | VARCHAR2(30) | - | For methods of name 'priority group', the name of the priority group. |
| PARAMETER_TABLE_NAME | VARCHAR2(30) | NOT NULL | Displays the name of the table to which the parameter column belongs. |
| PARAMETER_COLUMN_NAME | VARCHAR2(4000) | - | The name of the column used as the IN parameter for the conflict resolution method. |
| PARAMETER_SEQUENCE_NO | NUMBER | NOT NULL | Ordering of column used as IN parameter. |

**Note:** The SNAME column is not present in the USER_REPPARAMETER_COLUMN view.

## ALL_REPPRIORITY

Contains the value and priority level of each priority group member in each priority group accessible to the current user. Priority group names must be unique within a replication group. Priority levels and values must each be unique within a given priority group.

### Related Views:

- DBA_REPPRIORITY contains the value and priority level of each priority group member in each priority group in the database.
- USER_REPPRIORITY contains the value and priority level of each priority group member in each priority group owned by the current user.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| SNAME | VARCHAR2(30) | NOT NULL | The name of the replicated schema. This column is obsolete. |
| PRIORITY_GROUP | VARCHAR2(30) | NOT NULL | The name of the priority group or site priority group. |
| PRIORITY | NUMBER | NOT NULL | The priority level of the member. The highest number has the highest priority. |
| DATA_TYPE | VARCHAR2(9) | - | The data type of the values in the priority group. |
| FIXED_DATA_LENGTH | NUMBER(38) | - | The maximum length of values of data type CHAR. |
| CHAR_VALUE | CHAR(255) | - | The value of the priority group member, if DATA_TYPE is CHAR. |
| VARCHAR2_VALUE | VARCHAR2(4000) | - | The value of the priority group member, if DATA_TYPE is VARCHAR2. |
| NUMBER_VALUE | NUMBER | - | The value of the priority group member, if DATA_TYPE is NUMBER. |
| DATE_VALUE | DATE | - | The value of the priority group member, if DATA_TYPE is DATE. |
| RAW_VALUE | RAW(2000) | - | The value of the priority group member, if DATA_TYPE is RAW. |
| GNAME | VARCHAR2(30) | NOT NULL | The name of the replication group. |
| NCHAR_VALUE | NCHAR(500) | - | The value of the priority group member, if DATA_TYPE is NCHAR. |
| NVARCHAR2_VALUE | VARCHAR2(1000) | - | The value of the priority group member, if DATA_TYPE is NVARCHAR2. |
| LARGE_CHAR_VALUE | CHAR(2000) | - | The value of the priority group member, for blank-padded character strings over 255 characters. |

**Note:** The SNAME and GNAME columns are not present in the USER_REPPRIORITY view.

# ALL_REPPRIORITY_GROUP

Describes the priority group or site priority group defined for each replication group accessible to the current user.

## Related Views:

- DBA_REPPRIORITY_GROUP describes the priority group or site priority group defined for each replication group in the database.
- USER_REPPRIORITY_GROUP describes the priority group or site priority group defined for each replication group owned by the current user.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| SNAME | VARCHAR2(30) | NOT NULL | The name of the replicated schema. This column is obsolete. |
| PRIORITY_GROUP | VARCHAR2(30) | NOT NULL | The name of the priority group or site priority group. |
| DATA_TYPE | VARCHAR2(9) | - | The data type of each value in the priority group. |
| FIXED_DATA_LENGTH | NUMBER(38) | - | The maximum length for values of data type CHAR. |
| PRIORITY_COMMENT | VARCHAR2(80) | - | Any user-supplied comments. |
| GNAME | VARCHAR2(30) | NOT NULL | The name of the replication group. |

**Note:** The SNAME and GNAME columns are not present in the USER_REPPRIORITY view.

## ALL_REPPROP

Indicates the technique used to propagate operations on each replicated object to the same object at another master site. These operations might have resulted from a call to a stored procedure or procedure wrapper, or might have been issued against a table directly. This view shows objects accessible to the current user.

### Related Views:

- DBA_REPPROP indicates the technique used to propagate operations on each replicated object to the same object at another master site. This view shows all objects in the database.
- USER_REPPROP indicates the technique used to propagate operations on each replicated object to the same object at another master site. This view shows objects owned by the current user

| Column | Data Type | NULL | Description |
|---|---|---|---|
| SNAME | VARCHAR2(30) | NOT NULL | The name of the schema containing the replicated object. |
| ONAME | VARCHAR2(30) | NOT NULL | The name of the replicated object. |
| TYPE | VARCHAR2(16) | - | The type of object being replicated:<br>FUNCTION PROCEDURE<br>INDEXTYPE MATERIALIZED VIEW<br>OPERATOR TABLE<br>PACKAGE TYPE<br>PACKAGE BODY TYPE BODY |
| DBLINK | VARCHAR2(128) | NOT NULL | The fully qualified database name of the master site to which changes are being propagated. |
| HOW | VARCHAR2(13) | - | How propagation is performed. Values recognized are 'none' for the local master site, and 'synchronous' or 'asynchronous' for all others. |
| PROPAGATE_COMMENT | VARCHAR2(80) | - | Any user-supplied comments. |

# ALL_REPRESOL_STATS_CONTROL

Describes statistics collection for conflict resolutions for all replicated tables accessible to the current user.

## Related Views:

- DBA_REPRESOL_STATS_CONTROL describes statistics collection for conflict resolutions for all replicated tables in the database.
- USER_REPRESOL_STATS_CONTROL describes statistics collection for conflict resolutions for all replicated tables owned by the current user.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| SNAME | VARCHAR2(30) | NOT NULL | Owner of the table. |
| ONAME | VARCHAR2(30) | NOT NULL | Table name. |
| CREATED | DATE | NOT NULL | Timestamp for when statistics collection was first started. |
| STATUS | VARCHAR2(9) | - | Status of statistics collection: ACTIVE or CANCELLED. |
| STATUS_UPDATE_DATE | DATE | NOT NULL | Timestamp for when the status was last updated. |
| PURGED_DATE | DATE | - | Timestamp for the last purge of statistics data. |
| LAST_PURGE_START_DATE | DATE | - | The last start date of the statistics purging date range. |
| LAST_PURGE_END_DATE | DATE | - | The last end date of the statistics purging date range. |

**Note:** The SNAME column is not present in the USER_REPRESOL_STATS_CONTROL view.

## ALL_REPRESOLUTION

Indicates the methods used to resolve update, uniqueness, or delete conflicts for each table accessible to the current user that is replicated using row-level replication for a given schema.

### Related Views:

- DBA_REPRESOLUTION indicates the methods used to resolve update, uniqueness, or delete conflicts for each table in the database that is replicated using row-level replication for a given schema.
- USER_REPRESOLUTION indicates the methods used to resolve update, uniqueness, or delete conflicts for each table owned by the current user that is replicated using row-level replication.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| SNAME | VARCHAR2(30) | NOT NULL | The name of the replicated schema. |
| ONAME | VARCHAR2(30) | NOT NULL | The name of the replicated table. |
| CONFLICT_TYPE | VARCHAR2(10) | - | The type of conflict that the method is used to resolve: delete, uniqueness, or update. |
| REFERENCE_NAME | VARCHAR2(30) | NOT NULL | The object to which the method applies. For delete conflicts, this is the table name. For uniqueness conflicts, this is the constraint name. For update conflicts, this is the column group name. |
| SEQUENCE_NO | NUMBER | NOT NULL | The order that resolution methods are applied, with 1 applied first. |
| METHOD_NAME | VARCHAR2(80) | NOT NULL | The name of an Oracle-supplied conflict resolution method. For user-supplied methods, this value is 'user function'. |
| FUNCTION_NAME | VARCHAR2(92) | NOT NULL | For methods of type 'user function', the name of the user-supplied conflict resolution method. |
| PRIORITY_GROUP | VARCHAR2(30) | - | For methods of type 'priority group', the name of the priority group. |
| RESOLUTION_COMMENT | VARCHAR2(80) | - | Any user-supplied comments. |

**Note:** The SNAME column is not present in the USER_REPRESOLUTION view.

# ALL_REPRESOLUTION_METHOD

Lists all of the conflict resolution methods available in the database. Initially, this view lists the standard methods provided with Advanced Replication. As you create new user functions and add them as conflict resolution methods for an object in the database, these functions are added to this view.

## Related Views:

- DBA_REPRESOLUTION_METHOD lists all of the conflict resolution methods available in the database.
- USER_REPRESOLUTION_METHOD lists all of the conflict resolution methods available in the database.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| CONFLICT_TYPE | VARCHAR2(10) | - | The type of conflict that the resolution method is designed to resolve: update, uniqueness, or delete. |
| METHOD_NAME | VARCHAR2(80) | NOT NULL | The name of the Oracle-supplied method, or the name of the user-supplied method. |

## ALL_REPRESOLUTION_STATISTICS

Lists information about successfully resolved update, uniqueness, and delete conflicts for all replicated tables accessible to the current user. These statistics are gathered for a table only if you have called the DBMS_REPCAT.REGISTER_STATISTICS procedure.

### Related Views:

- DBA_REPRESOLUTION_STATISTICS lists information about successfully resolved update, uniqueness, and delete conflicts for all replicated tables in the database.
- USER_REPRESOLUTION_STATISTICS lists information about successfully resolved update, uniqueness, and delete conflicts for all replicated tables owned by the current user.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| SNAME | VARCHAR2(30) | NOT NULL | The name of the replicated schema. |
| ONAME | VARCHAR2(30) | NOT NULL | The name of the replicated table. |
| CONFLICT_TYPE | VARCHAR2(10) | - | The type of conflict that was successfully resolved: delete, uniqueness, or update. |
| REFERENCE_NAME | VARCHAR2(30) | NOT NULL | The object to which the conflict resolution method applies. For delete conflicts, this is the table name. For uniqueness conflicts, this is the constraint name. For update conflicts, this is the column group name. |
| METHOD_NAME | VARCHAR2(80) | NOT NULL | The name of an Oracle-supplied conflict resolution method. For user-supplied methods, this value is 'user function'. |
| FUNCTION_NAME | VARCHAR2(92) | - | For methods of type 'user function', the name of the user supplied conflict resolution method. |
| PRIORITY_GROUP | VARCHAR2(30) | - | For methods of type 'priority group', the name of the priority group. |
| RESOLVED_DATE | DATE | NOT NULL | Date on which the conflict for this row was resolved. |
| PRIMARY_KEY_VALUE | VARCHAR2(2000) | NOT NULL | A concatenated representation of the row's primary key. |

**Note:** The SNAME column is not present in the USER_REPRESOLUTION_STATISTICS view.

# ALL_REPSITES

Lists the members of each replication group accessible to the current user.

## Related Views:

- DBA_REPSITES lists the members of each replication group in the database.
- USER_REPSITES lists the members of each replication group owned by the current user.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| GNAME | VARCHAR2(30) | NOT NULL | The name of the replication group. |
| DBLINK | VARCHAR2(128) | NOT NULL | The database link to a master site for this replication group. |
| MASTERDEF | VARCHAR2(1) | - | Indicates which of the DBLINKs is the master definition site. |
| SNAPMASTER | VARCHAR2(1) | - | Used by materialized view sites to indicate which of the DBLINKs to use when refreshing. |
| MASTER_COMMENT | VARCHAR2(80) | - | User-supplied comments. |
| MASTER | VARCHAR2(1) | - | If Y then the site is a master site for the replicated group.<br>If N then the site is not a master site for the replicated group. |
| GROUP_OWNER | VARCHAR2(30) | NOT NULL | Owner of the replication group. |

The DBA_REPSITES view has the following additional columns:

| Column | Data Type | NULL | Description |
|---|---|---|---|
| PROP_UPDATES | NUMBER | - | Encoding of propagating technique for master site. |
| MY_DBLINK | VARCHAR2(1) | - | Used to detect problems after import. If Y then the DBLINK is the global name. |

## DBA_REPCAT_REFRESH_TEMPLATES

This view contains global information about each deployment template in the database, such as the template name, template owner, what refresh group the template objects belong to, and the type of template (private or public).

Its columns are the same as those in `ALL_REPCAT_REFRESH_TEMPLATES`. For detailed information about this view and its columns, see ALL_REPCAT_REFRESH_TEMPLATES on page 23-6.

# DBA_REPCAT_TEMPLATE_OBJECTS

The DBA_REPCAT_TEMPLATE_OBJECTS view contains the individual object definitions that are contained in all deployment templates in the database. Individual objects are added to a template by specifying the target template in REFRESH_TEMPLATE_NAME.

Its columns are the same as those in ALL_REPCAT_TEMPLATE_OBJECTS. For detailed information about this view and its columns, see ALL_REPCAT_TEMPLATE_OBJECTS on page 23-7.

## DBA_REPCAT_TEMPLATE_PARMS

Parameters defined in the object DDL for all templates in the database are stored in the DBA_REPCAT_TEMPLATE_PARMS table. When an object is added to a template, the DDL is examined for variables. Any found parameters are automatically added to this view.

Its columns are the same as those in ALL_REPCAT_TEMPLATE_PARMS. For detailed information about this view and its columns, see ALL_REPCAT_TEMPLATE_PARMS on page 23-9.

# DBA_REPCAT_TEMPLATE_SITES

The DBA_REPCAT_TEMPLATE_SITES view provides the DBA with information about the current status of template instantiation for all the sites of a enterprise network. This view contains information about instantiation sites for all deployment templates in the database. Specifically, the DBA can monitor the installation and deletion of templates at specific sites. Its columns are the same as those in ALL_REPCAT_TEMPLATE_SITES on page 23-11.

## DBA_REPCAT_USER_AUTHORIZATIONS

The DBA_REPCAT_USER_AUTHORIZATIONS view lists the authorized users for all templates in the database specified for private use. Users listed in this view have the ability to instantiate the specified template. Users not contained in this view cannot instantiate the template. Its columns are the same as those in ALL_REPCAT_USER_AUTHORIZATIONS on page 23-12.

# DBA_REPCAT_USER_PARM_VALUES

The DBA_REPCAT_USER_PARM_VALUES view describes the template parameters for all deployment templates in the database. The DBA has the option of building a table of user parameters before distributing the template for instantiation. When a template is instantiated by a specified user, the values stored in the DBA_REPCAT_USER_PARM_VALUES table for the specified user are used automatically.

Its columns are the same as those in ALL_REPCAT_USER_PARM_VALUES. For detailed information about this view and its columns, see ALL_REPCAT_USER_PARM_VALUES on page 23-13.

## DBA_REPCATLOG

The DBA_REPCATLOG view at each master site contains the interim status of any asynchronous administrative requests and any error messages generated. All messages encountered while executing a request are eventually transferred to the DBA_REPCATLOG view at the master site that originated the request. If an administrative request completes without error, then ultimately all traces of this request are removed from the DBA_REPCATLOG view. Its columns are the same as those in ALL_REPCATLOG on page 23-15.

# DBA_REPCOLUMN

The DBA_REPCOLUMN view lists the replicated columns for all the tables in the database. Its columns are the same as those in ALL_REPCOLUMN on page 23-16.

## DBA_REPCOLUMN_GROUP

The DBA_REPCOLUMN_GROUP view lists all the column groups each replicated table in the database. Its columns are the same as those in ALL_REPCOLUMN_GROUP on page 23-18.

# DBA_REPCONFLICT

The DBA_REPCONFLICT view displays the name of each table in the database on which a conflict resolution method has been defined and the type of conflict that the method is used to resolve. Its columns are the same as those in ALL_REPCONFLICT on page 23-19.

## DBA_REPDDL

The DBA_REPDDL contains the DDL for each replication object in the database. Its columns are the same as those in ALL_REPDDL on page 23-20.

# DBA_REPEXTENSIONS

The DBA_REPEXTENSIONS view contains information about current operations that are adding new master sites to a master group without quiescing the master group.

**See Also:** "Adding New Master Sites without Quiescing the Master Group" on page 7-3 for information about the procedure that adds new master sites to a replication environment

| Column | Data Type | NULL | Description |
|---|---|---|---|
| EXTENSION_ID | RAW(16) | NOT NULL | The identifier for a current pending request to add master databases to a master group without quiesce. |
| REQUEST | VARCHAR2(15) | - | Extension request type. Currently, the only possible value is ADD_NEW_MASTERS, which indicates a request to add new master sites to a master group without quiescing. |
| MASTERDEF | VARCHAR2(128) | - | The global name of the master definition site of the master groups to which new master sites are being added. |
| EXPORT_REQUIRED | VARCHAR2(3) | - | YES indicates that one or more new master sites will be added using export/import of either the entire database or at the table level.<br>NO indicates that all new master sites will be added using change-based recovery. |
| REPCATLOG_ID | NUMBER | - | Identifier of replication catalog records related to a replication extension, on which the master definition site is waiting. This value is only meaningful at the master definition site. |
| EXTENSION_STATUS | VARCHAR2(13) | - | Status of each replication extension. This value is only meaningful at the master definition site.<br>The possible values are:<br>READY: The extension request has been created and is ready.<br>STOPPING: The new master sites have been added to the master group and the master definition site is attempting to stop propagation from existing masters to new master sites and to the master definition site.<br>EXPORTING: The propagation of deferred transactions has been stopped from existing master sites to new master sites and to the master definition site. The master definition site is waiting for the export to finish.<br>INSTANTIATING: The DBMS_REPCAT.RESUME_PROPAGATION_TO_MDEF procedure has been invoked (if export was used), and the master definition site is waiting for the new masters to instantiate.<br>ERROR: An error occurred during the execution of this extension request. |

| Column | Data Type | NULL | Description |
|---|---|---|---|
| FLASHBACK_SCN | NUMBER | - | The system change number (SCN) that must be used during export or change-based recovery when the new master sites are added. The new master sites must be consistent with the SCN listed. |
| BREAK_TRANS_TO_MASTERDEF | VARCHAR2(3) | - | This value is meaningful only if EXPORT_REQUIRED is TRUE.<br>If BREAK_TRANS_TO_MASTERDEF is TRUE, then existing masters can continue to propagate their deferred transactions to the master definition site for replication groups that are not adding master sites. Deferred transactions for replication groups that are adding master sites cannot be propagated until the export completes.<br>If BREAK_TRANS_TO_MASTERDEF is FALSE, then existing masters cannot propagate any deferred transactions to the master definition site. |
| BREAK_TRANS_TO_NEW_MASTERS | VARCHAR2(3) | - | If BREAK_TRANS_TO_NEW_MASTERS is TRUE, then existing master sites can continue to propagate deferred transactions to the new master sites for replication groups that are not adding master sites.<br>If BREAK_TRANS_TO_NEW_MASTERS is FALSE, then propagation of deferred transaction queues to the new masters is disabled. |
| PERCENTAGE_FOR_CATCHUP_MDEF | NUMBER | - | This value is meaningful only if BREAK_TRANS_TO_MASTERDEF is TRUE.<br>The percentage of propagation resources to use for catching up propagation to the master definition site. |
| CYCLE_SECONDS_MDEF | NUMBER | - | This value is meaningful when PERCENTAGE_FOR_CATCHUP_MDEF is both meaningful and is a value between 10 and 90, inclusive. In this case, propagation to the master definition site alternates between replication groups that are not being extended and replication groups that are being extended, with one push to each during each cycle. This value indicates the length of the cycle in seconds. |
| PERCENTAGE_FOR_CATCHUP_NEW | NUMBER | - | This value is meaningful only if BREAK_TRANS_TO_NEW_MASTERS is TRUE.<br>The percentage of propagation resources to use for catching up propagation to new master sites. |
| CYCLE_SECONDS_NEW | NUMBER | - | This value is meaningful when PERCENTAGE_FOR_CATCHUP_NEW is both meaningful and is a value between 10 and 90, inclusive. In this case, propagation to a new master alternates between replication groups that are not being extended and replication groups that are being extended, with one push to each during each cycle. This value indicates the length of the cycle in seconds. |

# DBA_REPGENOBJECTS

The DBA_REPGENOBJECTS view describes each object in the database that was generated to support replication. Its columns are the same as those in ALL_REPGENOBJECTS on page 23-21.

## DBA_REPGROUP

The DBA_REPGROUP view describes all of the replication groups in the database. The members of each replication group are listed in a different view, DBA_REPOBJECT. The DBA_REPGROUP view's columns are the same as those in ALL_REPGROUP on page 23-22.

# DBA_REPGROUP_PRIVILEGES

The DBA_REPGROUP_PRIVILEGES view contains information about the users who are registered for privileges in replication groups. Shows all replication groups in the database. Its columns are the same as those in ALL_REPGROUP_PRIVILEGES on page 23-23.

## DBA_REPGROUPED_COLUMN

The DBA_REPGROUPED_COLUMN view lists all of the columns that comprise the column groups for each table in the database. Its columns are the same as those in ALL_REPGROUPED_COLUMN on page 23-24.

# DBA_REPKEY_COLUMNS

The DBA_REPKEY_COLUMNS view describes the replication key column(s) in each table in the database. Its columns are the same as those in ALL_REPKEY_COLUMNS on page 23-25.

## DBA_REPOBJECT

The DBA_REPOBJECT view contains information about the objects in each replication group in the database. An object can belong to only one replication group. A replication group can span multiple schemas. Its columns are the same as those in ALL_REPOBJECT on page 23-26.

# DBA_REPPARAMETER_COLUMN

In addition to the information contained in the DBA_REPRESOLUTION view, the DBA_REPPARAMETER_COLUMN view contains information about the columns that are used to resolve conflicts for each replicated table in the database. These are the column values that are passed as the list_of_column_names argument to the ADD_*conflicttype*_RESOLUTION procedures in the DBMS_REPCAT package. Its columns are the same as those in ALL_REPPARAMETER_COLUMN on page 23-28.

## DBA_REPPRIORITY

The DBA_REPPRIORITY view contains the value and priority level of each priority group member in each priority group in the database. Priority group names must be unique within a replication group. Priority levels and values must each be unique within a given priority group. Its columns are the same as those in ALL_REPPRIORITY on page 23-29.

# DBA_REPPRIORITY_GROUP

The DBA_REPPRIORITY_GROUP view describes the priority group or site priority group defined for each replication group in the database. Its columns are the same as those in ALL_REPPRIORITY_GROUP on page 23-30.

## DBA_REPPROP

The DBA_REPPROP view indicates the technique used to propagate operations on each replicated object to the same object at another master site. These operations might have resulted from a call to a stored procedure or procedure wrapper, or might have been issued against a table directly. This view shows all objects in the database. Its columns are the same as those in ALL_REPPROP on page 23-31.

# DBA_REPRESOL_STATS_CONTROL

The DBA_REPRESOL_STATS_CONTROL view describes statistics collection for conflict resolutions for all replicated tables in the database. Its columns are the same as those in ALL_REPRESOL_STATS_CONTROL on page 23-32.

## DBA_REPRESOLUTION

The DBA_REPRESOLUTION view indicates the methods used to resolve update, uniqueness, or delete conflicts for each table in the database that is replicated using row-level replication for a given schema. Its columns are the same as those in ALL_REPRESOLUTION on page 23-33.

# DBA_REPRESOLUTION_METHOD

The DBA_REPRESOLUTION_METHOD view lists all of the conflict resolution methods available in the database. Initially, this view lists the standard methods provided with the advanced replication facility. As you create new user functions and add them as conflict resolution methods for an object in the database, these functions are added to this view. Its columns are the same as those in ALL_REPRESOLUTION_METHOD on page 23-34.

## DBA_REPRESOLUTION_STATISTICS

The DBA_REPRESOLUTION_STATISTICS view lists information about successfully resolved update, uniqueness, and delete conflicts for all replicated tables in the database. These statistics are only gathered for a table if you have called the DBMS_REPCAT.REGISTER_STATISTICS procedure. The DBA_REPRESOLUTION_STATISTICS view's columns are the same as those in ALL_REPRESOLUTION_STATISTICS on page 23-35.

# DBA_REPSITES

The DBA_REPSITES view lists the members of each replication group in the database.

This view has the following additional columns that are not included in the ALL_REPSITES and USER_REPSITES views:

| Column | Data Type | NULL | Description |
|---|---|---|---|
| PROP_UPDATES | NUMBER | - | Encoding of propagating technique for master site. |
| MY_DBLINK | VARCHAR2(1) | - | Used to detect problem after import. If Y then the dblink is the global name. |

Except for these additional columns, its columns are the same as those in ALL_REPSITES on page 23-36.

## DBA_REPSITES_NEW

The DBA_REPSITES_NEW view lists the new replication sites that you plan to add to your replication environment.

**See Also:** "Adding New Master Sites without Quiescing the Master Group" on page 7-3 for information about the procedure that adds new master sites to a replication environment

| Column | Data Type | NULL | Description |
|---|---|---|---|
| EXTENSION_ID | RAW(16) | NOT NULL | The identifier for a current pending request to add master databases to a master group without quiesce. |
| GOWNER | VARCHAR2(30) | NOT NULL | The name of the user who owns the master group. |
| GNAME | VARCHAR2(30) | NOT NULL | The name of the master group. |
| DBLINK | VARCHAR2(128) | NOT NULL | The database link for a new master site. |
| FULL_INSTANTIATION | VARCHAR2(1) | - | Y indicates that the new database in DBLINK is to be added using full database export/import or change-based recovery.<br>N indicates that the new database in DBLINK is to be added using object-level export/import. |
| MASTER_STATUS | VARCHAR2(13) | - | The instantiation status of a new master site. This value is only meaningful at the master definition site.<br>The possible values are:<br>READY: The new master site is ready.<br>INSTANTIATING: The new master site is in the process of being instantiated.<br>INSTANTIATED: The new master has been instantiated and is being prepared for replication activity. That is, the DBMS_REPCAT.PREPARE_INSTANTIATED_MASTER procedure has been run.<br>PREPARED: The propagation of deferred transactions is enabled from the new master site to other prepared masters, to existing masters, and to the master definition site. The new master is now prepared to participate in the replication environment. |

# USER_REPCAT_REFRESH_TEMPLATES

This view contains global information about each deployment template owned by the current user, such as the template name, template owner, what refresh group the template objects belong to, and the type of template (private or public).

Its columns are the same as those in `ALL_REPCAT_REFRESH_TEMPLATES`. For detailed information about this view and its columns, see ALL_REPCAT_REFRESH_TEMPLATES on page 23-6.

## USER_REPCAT_TEMPLATE_OBJECTS

The USER_REPCAT_TEMPLATE_OBJECTS view contains the individual object definitions that are contained in each deployment template owned by the current user. Individual objects are added to a template by specifying the target template in REFRESH_TEMPLATE_NAME.

Its columns are the same as those in ALL_REPCAT_TEMPLATE_OBJECTS. For detailed information about this view and its columns, see ALL_REPCAT_TEMPLATE_OBJECTS on page 23-7.

# USER_REPCAT_TEMPLATE_PARMS

Parameters defined in the object DDL for all templates owned by the current user are stored in the USER_REPCAT_TEMPLATE_PARMS table. When an object is added to a template, the DDL is examined for variables; any found parameters are automatically added to this view.

Its columns are the same as those in ALL_REPCAT_TEMPLATE_PARMS. For detailed information about this view and its columns, see ALL_REPCAT_TEMPLATE_PARMS on page 23-9.

## USER_REPCAT_TEMPLATE_SITES

The USER_REPCAT_TEMPLATE_SITES view provides the user with information about the current status of template instantiation among the sites of a enterprise network. This view contains information about instantiation sites for deployment templates that are owned by the current user. Specifically, the user can monitor the installation and deletion of templates at specific sites. Its columns are the same as those in ALL_REPCAT_TEMPLATE_SITES on page 23-11.

# USER_REPCAT_USER_AUTHORIZATION

The `USER_REPCAT_USER_AUTHORIZATION` view lists the authorized users for all of the templates that are owned by the current user and specified for private use. Users listed in this view have the ability to instantiate the specified template. Users not contained in this view cannot instantiate the template. Its columns are the same as those in ALL_REPCAT_USER_AUTHORIZATIONS on page 23-12.

## USER_REPCAT_USER_PARM_VALUES

The USER_REPCAT_USER_PARM_VALUES view describes the template parameters for all deployment templates owned by the current user. The DBA has the option of building a table of user parameters before distributing the template for instantiation. When a template is instantiated by a specified user, the values stored in the USER_REPCAT_USER_PARM_VALUES view for the specified user are used automatically.

Its columns are the same as those in ALL_REPCAT_USER_PARM_VALUES. For detailed information about this view and its columns, see ALL_REPCAT_USER_PARM_VALUES on page 23-13.

# USER_REPCATLOG

The USER_REPCATLOG view at each master site contains the interim status of any asynchronous administrative requests and any error messages generated. All messages encountered while executing a request are eventually transferred to the USER_REPCATLOG view at the master site that originated the request. If an administrative request completes without error, then ultimately all traces of this request are removed from the USER_REPCATLOG view.

This view contains asynchronous administrative requests and error messages that are owned by the current user. Its columns are the same as those in ALL_REPCATLOG on page 23-15.

## USER_REPCOLUMN

The USER_REPCOLUMN view lists the replicated columns for all the tables owned by the current user. Its columns are the same as those in ALL_REPCOLUMN on page 23-16.

# USER_REPCOLUMN_GROUP

The USER_REPCOLUMN_GROUP view lists the column groups for each replicated table owned by the current user. Its columns are the same as those in ALL_REPCOLUMN_GROUP on page 23-18.

**Note:** The SNAME column is not present in the USER_REPCOLUMN_GROUP view. This column is available in the ALL_REPCOLUMN_GROUP and DBA_REPCOLUMN_GROUP views.

## USER_REPCONFLICT

The USER_REPCONFLICT view displays the name of each table owned by the current user on which a conflict resolution method has been defined and the type of conflict that the method is used to resolve. Its columns are the same as those in ALL_REPCONFLICT on page 23-19.

---

**Note:** The SNAME column is not present in the USER_REPCONFLICT view. This column is available in the ALL_REPCONFLICT and DBA_REPCONFLICT views.

---

# USER_REPDDL

The USER_REPDDL contains the DDL for each replication object owned by the current user. Its columns are the same as those in ALL_REPDDL on page 23-20.

## USER_REPGENOBJECTS

The USER_REPGENOBJECTS view describes each object owned by the current user that was generated to support replication. Its columns are the same as those in ALL_REPGENOBJECTS on page 23-21.

# USER_REPGROUP

The USER_REPGROUP view describes all of the replication groups owned by the current user. The members of each replication group are listed in a different view, USER_REPOBJECT. The USER_REPGROUP view's columns are the same as those in ALL_REPGROUP on page 23-22.

## USER_REPGROUP_PRIVILEGES

The `USER_REPGROUP_PRIVILEGES` view contains information about the users who are registered for privileges in replication groups. Shows only those replication groups owned by the current user. Its columns are the same as those in ALL_REPGROUP_PRIVILEGES on page 23-23.

# USER_REPGROUPED_COLUMN

The USER_REPGROUPED_COLUMN view lists all of the columns that comprise the column groups for each table. Its columns are the same as those in ALL_REPGROUPED_COLUMN on page 23-24.

**Note:** The SNAME column is not present in the USER_REPGROUPED_COLUMN view. This column is available in the ALL_REPGROUPED_COLUMN and DBA_REPGROUPED_COLUMN views.

## USER_REPKEY_COLUMNS

The USER_REPKEY_COLUMNS view describes the replication key column(s) in each table owned by the current user. Its columns are the same as those in ALL_REPKEY_COLUMNS on page 23-25.

# USER_REPOBJECT

The USER_REPOBJECT view contains information about the objects owned by the current user in each replication group. An object can belong to only one replication group. A replication group can span multiple schemas. Its columns are the same as those in ALL_REPOBJECT on page 23-26.

## USER_REPPARAMETER_COLUMN

In addition to the information contained in the USER_REPRESOLUTION view, the USER_REPPARAMETER_COLUMN view contains information about the columns that are used to resolve conflicts for each replicated table owned by the current user. These are the column values that are passed as the list_of_column_names argument to the ADD_*conflicttype*_RESOLUTION procedures in the DBMS_REPCAT package. Its columns are the same as those in ALL_REPPARAMETER_COLUMN on page 23-28.

---

**Note:** The SNAME column is not present in the USER_REPPARAMETER_COLUMN view. This column is available in the ALL_REPPARAMETER_COLUMN and DBA_REPPARAMETER_COLUMN views.

---

# USER_REPPRIORITY

The USER_REPPRIORITY view contains the value and priority level of each priority group member in each priority group owned by the current user. Priority group names must be unique within a replication group. Priority levels and values must each be unique within a given priority group. Its columns are the same as those in ALL_REPPRIORITY on page 23-29.

---

**Note:** The SNAME column is not present in the USER_REPPRIORITY view. This column is available in the ALL_REPPRIORITY and DBA_REPPRIORITY views.

---

## USER_REPPRIORITY_GROUP

The `USER_REPPRIORITY_GROUP` view describes the priority group or site priority group defined for each replication group owned by the current user. Its columns are the same as those in ALL_REPPRIORITY_GROUP on page 23-30.

# USER_REPPROP

The USER_REPPROP view indicates the technique used to propagate operations on each replicated object to the same object at another master site. These operations might have resulted from a call to a stored procedure or procedure wrapper, or might have been issued against a table directly. This view shows objects owned by the current user. Its columns are the same as those in ALL_REPPROP on page 23-31.

## USER_REPRESOL_STATS_CONTROL

The USER_REPRESOL_STATS_CONTROL view describes statistics collection for conflict resolutions for all replicated tables owned by the current user. Its columns are the same as those in ALL_REPRESOL_STATS_CONTROL on page 23-32.

---

**Note:** The SNAME column is not present in the USER_REPRESOL_STATS_CONTROL view. This column is available in the ALL_REPRESOL_STATS_CONTROL and DBA_REPRESOL_STATS_CONTROL views.

---

# USER_REPRESOLUTION

The USER_REPRESOLUTION view indicates the methods used to resolve update, uniqueness, or delete conflicts for each table owned by the current user that is replicated using row-level replication for a given schema. Its columns are the same as those in ALL_REPRESOLUTION on page 23-33.

---

**Note:** The SNAME column is not present in the USER_REPREPRESOLUTION view. This column is available in the ALL_REPREPRESOLUTION and DBA_REPREPRESOLUTION views.

---

## USER_REPRESOLUTION_METHOD

The USER_REPRESOLUTION_METHOD view lists all of the conflict resolution methods available in the database. Initially, this view lists the standard methods provided with the advanced replication facility. As you create new user functions and add them as conflict resolution methods for an object in the database, these functions are added to this view. Its columns are the same as those in ALL_REPRESOLUTION_METHOD on page 23-34.

# USER_REPRESOLUTION_STATISTICS

The USER_REPRESOLUTION_STATISTICS view lists information about successfully resolved update, uniqueness, and delete conflicts for all replicated tables owned by the current user. These statistics are only gathered for a table if you have called the DBMS_REPCAT.REGISTER_STATISTICS procedure. The USER_REPRESOLUTION_STATISTICS view's columns are the same as those in ALL_REPRESOLUTION_STATISTICS on page 23-35.

---

**Note:** The SNAME column is not present in the USER_REPRESOLUTION_STATISTICS view. This column is available in the ALL_REPRESOLUTION_STATISTICS and DBA_REPRESOLUTION_STATISTICS views.

---

## USER_REPSITES

The USER_REPSITES view lists the members of each replication group owned by the current user. Its columns are the same as those in ALL_REPSITES on page 23-36.

# 24

# Replication Dynamic Performance Views

All Oracle installations include the dynamic performance views, often referred to as V$ views, described in this chapter. These views are used by master sites and materialized view sites to determine such information as which materialized views are being refreshed currently and statistics about the deferred transaction queue.

This chapter contains these topics:

- V$MVREFRESH
- V$REPLPROP
- V$REPLQUEUE

**See Also:** Chapter 10, "Monitoring a Replication Environment"

# V$MVREFRESH

Contains information about the materialized views currently being refreshed.

| Column | Data Type | Description |
|---|---|---|
| SID | NUMBER | Session identifier. |
| SERIAL# | NUMBER | Session serial number, which is used to identify uniquely a session's objects. Guarantees that session-level commands are applied to the correct session objects if the session ends and another session begins with the same session ID. |
| CURRMVOWNER | VARCHAR2(31) | Owner of the materialized view currently being refreshed. The materialized view resides in this user's schema. |
| CURRMVNAME | VARCHAR2(31) | Name of the materialized view currently being refreshed. |

# V$REPLPROP

Contains information about the parallel propagation currently in progress at the replication site. Use this view to determine which transactions are currently being propagated, the number of calls propagated in each transaction, and the current activity of the parallel propagation processes or parallel propagation coordinator process.

**Note:** This view only contains data when deferred transactions are being pushed using parallel propagation at the current site. The `parallelism` parameter must be set to `1` or higher in the `DBMS_DEFER_SYS.PUSH` function for a push to use parallel propagation. Otherwise, the push uses serial propagation, and no data appears in this view during the push.

| Column | Data Type | Description |
|---|---|---|
| SID | NUMBER | Session identifier. |
| SERIAL# | NUMBER | Session serial number. Used to identify uniquely a session's objects. Guarantees that session-level commands are applied to the correct session objects if the session ends and another session begins with the same session ID. |
| NAME | VARCHAR2(71) | `Replication Parallel Prop Slave` *n* indicates that the process is active, either waiting, pushing deferred transactions, purging metadata, or creating an error transaction.<br>`Replication Parallel Prop Coordinator` indicates that the coordinator process is active, either waiting, sleeping, or scheduling processes to perform operations.<br>The `Replication Parallel Prop Coordinator` reads transactions from the deferred transaction queue and assigns them to the `Replication Parallel Prop Slaves`. Then, the processes propagate the transactions to the destination site. When the processes push transactions in a push session, the processes remain active until the push session completes, even if there are no more transactions to push. |
| DBLINK | VARCHAR2(128) | Database link on which this replication session is propagating. |

| Column | Data Type | Description |
|---|---|---|
| STATE | VARCHAR2(12) | WAIT indicates that either the slave or coordinator process is waiting for an event (that is, a message).<br><br>SLEEP indicates that the coordinator process is sleeping for the duration of the delay_seconds setting. You set delay_seconds with the SCHEDULE_PUSH procedure in the DBMS_DEFER_SYS package.<br><br>PUSH indicates that the slave process is pushing transactions from the deferred transaction queue to the remote site.<br><br>PURGE indicates that the slave process is purging metadata related to successfully applied transactions from the remote site.<br><br>CREATE ERROR indicates that the slave process is creating an error transaction. In this case, an error or a conflict occurred while the slave was pushing deferred transactions to the remote site.<br><br>SCHEDULE TXN indicates that the coordinator process is determining the order that transactions are applied and assigning slave processes to execute the transactions. |
| XID | VARCHAR2(22) | If the session is a slave session, then indicates the transaction id of the transaction that the slave is currently propagating. |
| SEQUENCE | NUMBER | If the process is a slave process, then the sequence number of the calls propagated in the current operation, if relevant. Each transaction must process one or more calls, and the value of SEQUENCE starts at zero and increases as each call is processed. So, the SEQUENCE value shows the call that is currently being processed in each transaction. This value increases until the slave has processed all of the calls in a transaction. |

# V$REPLQUEUE

Contains statistics about the replication deferred transactions queue. All values are stored since the start of the current database instance.

| Column | Data Type | Description |
|---|---|---|
| TXNS_ENQUEUED | NUMBER | Number of transactions enqueued in the deferred transactions queue. |
| CALLS_ENQUEUED | NUMBER | Number of calls enqueued into the deferred transactions queue. |
| TXNS_PURGED | NUMBER | Number of transactions purged from the deferred transactions queue. |
| LAST_ENQUEUE_TIME | DATE | Date when the last transaction was enqueued into the deferred transaction queue. NULL if no transactions have been enqueued into the deferred transaction queue since the instance started. |
| LAST_PURGE_TIME | DATE | Date when the last transaction was purged from the deferred transaction queue. NULL if no transactions have been purged from the deferred transaction queue since the instance started. |

# 25

# Materialized View and Refresh Group Views

This chapter lists the following data dictionary views, which provide information about materialized views and materialized view refresh groups.

| ALL_ Views | DBA_ Views | USER_ Views |
|---|---|---|
| ALL_BASE_TABLE_MVIEWS | DBA_BASE_TABLE_MVIEWS | USER_BASE_TABLE_MVIEWS |
| N/A | DBA_MVIEW_LOG_FILTER_COLS | N/A |
| ALL_MVIEW_LOGS | DBA_MVIEW_LOGS | USER_MVIEW_LOGS |
| ALL_MVIEW_REFRESH_TIMES | DBA_MVIEW_REFRESH_TIMES | USER_MVIEW_REFRESH_TIMES |
| ALL_MVIEWS | DBA_MVIEWS | USER_MVIEWS |
| N/A | DBA_RCHILD | N/A |
| ALL_REFRESH | DBA_REFRESH | USER_REFRESH |
| ALL_REFRESH_CHILDREN | DBA_REFRESH_CHILDREN | USER_REFRESH_CHILDREN |
| ALL_REGISTERED_MVIEWS | DBA_REGISTERED_MVIEWS | USER_REGISTERED_MVIEWS |
| N/A | DBA_RGROUP | N/A |

**See Also:**

- *Oracle Database Reference* for more information about these views
- Chapter 10, "Monitoring a Replication Environment" for example queries that use some of these views

# 26

# Deferred Transaction Views

Oracle provides several views for you to use when administering deferred transactions. These views provide information about each deferred transaction, such as the transaction destinations, the deferred calls that comprise the transactions, and any errors encountered during attempted execution of the transaction.

This chapter contains these topics:

- DEFCALL
- DEFCALLDEST
- DEFDEFAULTDEST
- DEFERRCOUNT
- DEFERROR
- DEFLOB
- DEFPROPAGATOR
- DEFSCHEDULE
- DEFTRAN
- DEFTRANDEST

---

**Caution:** You should not modify the tables directly. Instead, use the procedures provided in the `DBMS_DEFER` and `DBMS_DEFER_SYS` packages.

---

**See Also:** Chapter 10, "Monitoring a Replication Environment"

# DEFCALL

Records all deferred remote procedure calls.

For calls placed in the queue using asynchronous replication, Oracle uses null compression for column objects and object tables that contain three or more consecutive nulls. Therefore, this view might show fewer attributes than the total number of attributes in a column object and fewer columns than the total number for an object table. For example, null compression can cause a column object with eight attributes to show only five attributes.

Null compression does not apply to error transactions.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| CALLNO | NUMBER | - | The unique ID of a call within a transaction. |
| DEFERRED_TRAN_ID | VARCHAR2(30) | - | The unique ID of the associated transaction. |
| SCHEMANAME | VARCHAR2(30) | - | The schema name of the deferred call. |
| PACKAGENAME | VARCHAR2(30) | - | The package name of the deferred call. For a replicated table, this can refer to the table name. |
| PROCNAME | VARCHAR2(30) | - | The procedure name of the deferred call. For a replicated table, this can refer to an operation name. |
| ARGCOUNT | NUMBER | - | The number of arguments in the deferred call. |

# DEFCALLDEST

Lists the destinations for each deferred remote procedure call.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| CALLNO | NUMBER | NOT NULL | Unique ID of a call within a transaction. |
| DEFERRED_TRAN_ID | VARCHAR2(30) | NOT NULL | Corresponds to the DEFERRED_TRAN_ID column in the DEFTRAN view. Each deferred transaction is made up of one or more deferred calls. |
| DBLINK | VARCHAR2(128) | NOT NULL | The fully qualified database name of the destination database. |

# DEFDEFAULTDEST

If you are not using Advanced Replication and do not supply a destination for a deferred transaction or the calls within that transaction, then Oracle uses the DEFDEFAULTDEST view to determine the destination databases to which you want to defer a remote procedure call.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| DBLINK | VARCHAR2(128) | NOT NULL | The fully qualified database name to which a transaction is replicated. |

# DEFERRCOUNT

Contains information about the error transactions for a destination.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| ERRCOUNT | NUMBER | - | Number of existing transactions that caused an error for the destination. |
| DESTINATION | VARCHAR2(128) | - | Database link used to address destination. |

# DEFERROR

Contains the ID of each transaction that could not be applied. You can use this ID to locate the queued calls associated with this transaction. These calls are stored in the DEFCALL view. You can use the procedures in the DBMS_DEFER_QUERY package to determine the arguments to the procedures listed in the DEFCALL view.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| DEFERRED_TRAN_ID | VARCHAR2(22) | NOT NULL | The ID of the transaction causing the error. |
| ORIGIN_TRAN_DB | VARCHAR2(128) | - | The database originating the deferred transaction. |
| ORIGIN_TRAN_ID | VARCHAR2(22) | - | The original ID of the transaction. |
| CALLNO | NUMBER | - | Unique ID of the call at DEFERRED_TRAN_ID. |
| DESTINATION | VARCHAR2(128) | - | Database link used to address destination. |
| START_TIME | DATE | - | Time when the original transaction was enqueued. |
| ERROR_NUMBER | NUMBER | - | Oracle error number. |
| ERROR_MSG | VARCHAR2(2000) | - | Error message text. |
| RECEIVER | VARCHAR2(30) | - | Original receiver of the deferred transaction. |

# DEFLOB

Contains the LOB parameters to deferred remote procedure calls (RPCs).

| Column | Data Type | NULL | Description |
|---|---|---|---|
| ID | RAW(16) | NOT NULL | Identifier of the LOB parameter. |
| DEFERRED_TRAN_ID | VARCHAR2(22) | - | Transaction ID for deferred remote procedure calls (RPCs) with this LOB parameter. |
| BLOB_COL | BLOB(4000) | - | The binary LOB parameter. |
| CLOB_COL | CLOB(4000) | - | The character LOB parameter. |
| NCLOB_COL | NCLOB(4000) | - | The national character LOB parameter. |

# DEFPROPAGATOR

Contains information about the local propagator.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| USERNAME | VARCHAR2(30) | NOT NULL | User name of the propagator. |
| USERID | NUMBER | NOT NULL | User ID of the propagator. |
| STATUS | VARCHAR2(7) | - | Status of the propagator. |
| CREATED | DATE | NOT NULL | Time when the propagator was registered. |

# DEFSCHEDULE

Contains information about when a job is next scheduled to be executed and also includes propagation statistics. The propagation statistics are for propagation of deferred transactions from the current site to the site specified in the DBLINK column.

To clear the propagation statistics for a remote site and start fresh, use the CLEAR_PROP_STATISTICS procedure in the DBMS_DEFER_SYS package.

**Note:** The statistics in this view are populated only if parallel propagation is used with a database link. To use parallel propagation, set the parallelism parameter to 1 or greater when you run the SCHEDULE_PUSH procedure in the DBMS_DEFER_SYS package.

**See Also:**

- *Oracle Database Advanced Replication* for information about parallel propagation
- "SCHEDULE_PUSH Procedure" on page 14-22
- "CLEAR_PROP_STATISTICS Procedure" on page 14-5

| Column | Data Type | NULL | Description |
|---|---|---|---|
| DBLINK | VARCHAR2(128) | NOT NULL | Fully qualified path name to the master site for which you have scheduled periodic execution of deferred remote procedure calls. |
| JOB | NUMBER | - | Number assigned to job when you created it by calling DBMS_DEFER_SYS.SCHEDULE_PUSH. Query the WHAT column of the USER_JOBS view to determine what is executed when the job is run. |
| INTERVAL | VARCHAR2(200) | - | Function used to calculate the next time to push the deferred transaction queue to destination. |
| NEXT_DATE | DATE | - | Next date that job is scheduled to be executed. |
| LAST_DATE | DATE | - | Last time the queue was pushed (or attempted to push) remote procedure calls to this destination. |
| DISABLED | CHAR(1) | - | If Y then propagation to destination is disabled.<br>If N then propagation to the destination is enabled. |
| LAST_TXN_COUNT | NUMBER | - | Number of transactions pushed during last attempt. |
| LAST_ERROR_NUMBER | NUMBER | - | Oracle error number from last push. |
| LAST_ERROR_MESSAGE | VARCHAR2(2000) | - | Error message from last push. |

| Column | Data Type | NULL | Description |
|---|---|---|---|
| CATCHUP | RAW(16) | NOT NULL | The extension identifier associated with a new master site that is being added to a master group without quiescing the master group. If there is no extension identifier for a master site, then the value is 00. |
| TOTAL_TXN_COUNT | NUMBER | - | Total combined number of successful transactions and error transactions. |
| AVG_THROUGHPUT | NUMBER | - | The average number of transactions for each second that are propagated using parallel propagation. The transactions include both successfully applied transactions and error transactions created on the remote site. Time that has elapsed when the propagation coordinator is inactive (sleeping) is included in the calculation. |
| AVG_LATENCY | NUMBER | - | If the transaction is successfully applied at the remote site, then the average number of seconds between the first call of a transaction on the current site and the confirmation that the transaction was applied at the remote site. The first call begins when the user makes the first data manipulation language (DML) change, not when the transaction is committed.<br><br>If the transaction is an error transaction, then the average number of seconds between the first call of a transaction on the current site and the confirmation that the error transaction is committed on the remote site. |
| TOTAL_BYTES_SENT | NUMBER | - | Total number of bytes sent, including replicated data and metadata. |
| TOTAL_BYTES_RECEIVED | NUMBER | - | Total number of bytes received in propagation confirmation messages. |
| TOTAL_ROUND_TRIPS | NUMBER | - | Total number of network round trips completed to replicate data. A round trip is one or more consecutively sent messages followed by one or more consecutively received messages. So, if site A sends 20 messages to site B and then site B sends one message to site A, then that is that one round trip. |
| TOTAL_ADMIN_COUNT | NUMBER | - | Total number of administrative requests sent to maintain information about transactions applied at the receiving site. The receiving site is the site specified in the DBLINK column. This special administration is only required for parallel propagation. |
| TOTAL_ERROR_COUNT | NUMBER | - | Total number of unresolved conflicts for which a remote error was created. |
| TOTAL_SLEEP_TIME | NUMBER | - | Total number of seconds the propagation coordinator was inactive (sleeping). You control the amount of time that the propagation coordinator sleeps using the delay_seconds parameter in the DBMS_DEFER_SYS.PUSH function. |

| Column | Data Type | NULL | Description |
|---|---|---|---|
| DISABLED_INTERNALLY_SET | VARCHAR2(1) | - | This value is relevant only if DISABLED is Y.<br><br>If DISABLED_INTERNALLY_SET is Y then propagation to destination was set to disabled internally by Oracle for propagation synchronization when adding a new master site to a master group without quiescing the master group. Oracle will enable propagation automatically at a later time.<br><br>If DISABLED_INTERNALLY_SET is N then propagation was not disabled internally. |

# DEFTRAN

Records all deferred transactions in the deferred transactions queue at the current site.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| DEFERRED_TRAN_ID | VARCHAR2(30) | - | The transaction ID that enqueued the calls. |
| DELIVERY_ORDER | NUMBER | - | An identifier that determines the order of deferred transactions in the queue. The identifier is derived from the system change number (SCN) of the originating transaction. |
| DESTINATION_LIST | VARCHAR2(1) | - | R indicates that the destinations are determined by the ALL_REPSITES view.<br>D indicates that the destinations were determined by the DEFDEFAULTDEST view or the NODE_LIST argument to the TRANSACTION or CALL procedures. |
| START_TIME | DATE | - | The time that the original transaction was enqueued. |

# DEFTRANDEST

Lists the destinations for each deferred transaction in the deferred transactions queue at the current site.

| Column | Data Type | NULL | Description |
|---|---|---|---|
| DEFERRED_TRAN_ID | VARCHAR2(30) | NOT NULL | The transaction ID of the transaction to replicate to the given database link. |
| DELIVERY_ORDER | NUMBER | - | An identifier that determines the order of deferred transactions in the queue. The identifier is derived from the system change number (SCN) of the originating transaction. |
| DBLINK | VARCHAR2(128) | NOT NULL | The fully qualified database name of the destination database. |

# Part V

## Appendixes

Part V contains the following appendixes:

- Appendix A, "Security Options"
- Appendix B, "User-Defined Conflict Resolution Methods"

# A

# Security Options

This appendix describes security options for multimaster and materialized view replication environments.

This appendix contains these topics:

- Security Setup for Multimaster Replication
- Security Setup for Materialized View Replication

## Security Setup for Multimaster Replication

Nearly all users should find it easiest to use the configuration wizards in the Advanced Replication interface in Oracle Enterprise Manager when configuring multimaster replication security. However, in certain cases you might need to use the replication management API to perform these setup operations.

To configure a replication environment, the database administrator must connect with DBA privileges to grant the necessary privileges to the replication administrator.

First set up user accounts at each master site with the appropriate privileges to configure and maintain the replication environment and to propagate and apply replicated changes. You must also define links for users at each master site.

In addition to the end users who access replicated objects, there are three special categories of "users" in a replication environment:

- Replication administrators, who are responsible for configuring and maintaining a replication environment.
- Propagators, who are responsible for propagating deferred transactions.
- Receivers at remote sites, who are responsible for applying these transactions.

Typically, a single user acts as administrator, propagator, and receiver. However, you can have separate users perform each of these functions. You can choose to have a single, global replication administrator or, if your replication groups do not span schema boundaries, you might prefer to have separate replication administrators for different schemas. Note, however, that you can have only one registered propagator for each database.

Table A–1 on page A-3 describes the necessary privileges that must be assigned to these specialized accounts. Most privileges needed by these users are granted to them through calls to the replication management API. You also must grant certain privileges directly, such as the privileges required to connect to the database and manage database objects.

## Trusted Compared with Untrusted Security

In addition to the different types of users, you also must determine which type of security model you will implement: trusted or untrusted. With a trusted security model, the receiver has access to all local master groups. Because the receiver performs database activities at the local master site on behalf of the propagator at the remote site, the propagator also has access to all master groups at the receiver's site. Remember that a single receiver is used for all incoming transactions.

For example, consider the scenario in Figure A–1. Even though only Master Groups A and C exist at Master Site B, the propagator has access to Master Groups A, B, C, and D at Master Site A because the trusted security model has been used. While this greatly increases the flexibility of database administration, due to the mobility of remote database administration, it also increases the chances of a malicious user at a remote site viewing or corrupting data at the master site.

Regardless of the security model used, Oracle automatically grants the appropriate privileges for objects as they are added to or removed from a replication environment.

***Figure A–1 Trusted Security: Multimaster Replication***

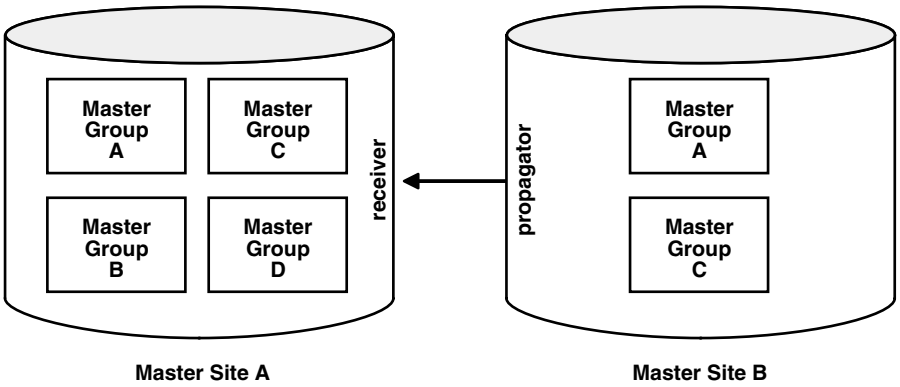


Untrusted security assigns only the privileges to the receiver that are required to work with specified master groups. The propagator, therefore, can only access the specified master groups that are local to the receiver. Figure A–2 illustrates an untrusted security model. Because Master Site B contains only Master Groups A and C, the receiver at Master Site A has been granted privileges for Master Groups A and C only, thereby limiting the propagator's access at Master Site A.

***Figure A–2 Untrusted Security: Multimaster Replication***

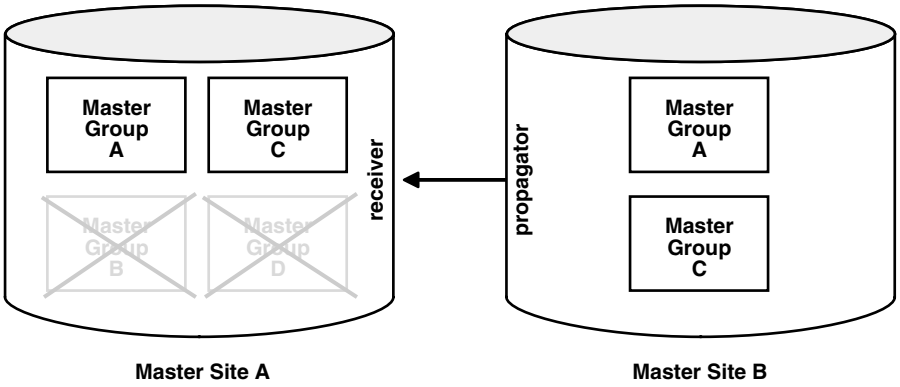

Typically, master sites are considered trusted and therefore the trusted security model is used. If, however, your remote master sites are untrusted, then you might want to use the untrusted model and assign your receiver limited privileges. A site might be considered untrusted, for example, if a consulting shop performs work for multiple customers. Use the appropriate API call listed for the receiver in Table A–1 to assign the different users the appropriate privileges.

***Table A–1 Required User Accounts***

| User | Privileges |
|---|---|
| global replication administrator | `DBMS_REPCAT_ADMIN.GRANT_ADMIN_ANY_SCHEMA` |
| schema-level replication administrator | `DBMS_REPCAT_ADMIN.GRANT_ADMIN_SCHEMA` |
| propagator | `DBMS_DEFER_SYS.REGISTER_PROPAGATOR` |
| receiver | See "REGISTER_USER_REPGROUP Procedure" on page 20-5 for details.<br>**Trusted**:<br>`DBMS_REPCAT_ADMIN.REGISTER_USER_REPGROUP`<br>`privilege => 'receiver'`<br>`list_of_gnames => NULL,`<br>`...`<br>**Untrusted**:<br>`DBMS_REPCAT_ADMIN.REGISTER_USER_REPGROUP`<br>`privilege => 'receiver'`<br>`list_of_gnames => '`*`mastergroupname`*`',`<br>`...` |

After you have created these accounts and assigned the appropriate privileges, create the following private database links, including user name and password between each site:

- From the local replication administrator to the remote replication administrator.
- From the local propagator to the remote receiver.

Assuming you have designated a single user account to act as replication administrator, propagator, and receiver, you must create N(N-1) links, where N is the number of master sites in your replication environment.

After creating these links, you must call `DBMS_DEFER_SYS.SCHEDULE_PUSH` and `DBMS_DEFER_SYS.SCHEDULE_PURGE`, at each location, to define how frequently you want to propagate your deferred transaction queue to each remote location, and how frequently you want to purge this queue. You must call `DBMS_DEFER_SYS.SCHEDULE_PUSH` multiple times at each site, once for each remote location.

A sample script for setting up multimaster replication between `hq.example.com` and `sales.example.com` is shown as follows:

```
/*--- Create global replication administrator at HQ ---*/
CONNECT system@hq.example.com
ACCEPT password PROMPT 'Enter password for user: ' HIDE
CREATE USER repadmin IDENTIFIED BY &password;
EXECUTE DBMS_REPCAT_ADMIN.GRANT_ADMIN_ANY_SCHEMA(username => 'repadmin');

/*--- Create global replication administrator at Sales ---*/
CONNECT system@sales.example.com
CREATE USER repadmin IDENTIFIED BY &password;
```

```
EXECUTE DBMS_REPCAT_ADMIN.GRANT_ADMIN_ANY_SCHEMA(username => 'repadmin');

/*--- Create single user to act as both propagator and receiver at HQ ---*/
CONNECT system@hq.example.com
CREATE USER prop_rec IDENTIFIED BY &password;
/*--- Grant privileges necessary to act as propagator ---*/
EXECUTE DBMS_DEFER_SYS.REGISTER_PROPAGATOR(username => 'prop_rec');
/*--- Grant privileges necessary to act as receiver ---*/
BEGIN
  DBMS_REPCAT_ADMIN.REGISTER_USER_REPGROUP(
        username => 'prop_rec',
        privilege_type => 'receiver',
        list_of_gnames => NULL);
END;
/

/*--- Create single user to act as both propagator and receiver at Sales ---*/
CONNECT system@sales.example.com
CREATE USER prop_rec IDENTIFIED BY &password;
/*--- Grant privileges necessary to act as propagator ---*/execute
EXECUTE DBMS_DEFER_SYS.REGISTER_PROPAGATOR(username => 'prop_rec');
/*--- Grant privileges necessary to act as receiver ---*/
BEGIN
  DBMS_REPCAT_ADMIN.REGISTER_USER_REPGROUP(
        username => 'prop_rec',
        privilege_type => 'receiver',
        list_of_gnames => NULL);
END;
/

/*--- Create public link from HQ to Sales with necessary USING clause ---*/
CONNECT system@hq.example.com
CREATE PUBLIC DATABASE LINK sales.example.com USING 'sales.example.com';

/*--- Create private repadmin to repadmin link ---*/
CONNECT repadmin@hq.example.com
CREATE DATABASE LINK sales.example.com CONNECT TO repadmin
 IDENTIFIED BY &password;

/*--- Schedule replication from HQ to Sales ---*/
BEGIN
  DBMS_DEFER_SYS.SCHEDULE_PUSH(
     destination => 'sales.example.com',
     interval => 'sysdate + 1/24',
     next_date => sysdate,
     stop_on_error => FALSE,
     parallelism => 1);
END;
/

/*--- Schedule purge of def tran queue at HQ ---*/
BEGIN
  DBMS_DEFER_SYS.SCHEDULE_PURGE(
     next_date => sysdate,
     interval => 'sysdate + 1',
     delay_seconds => 0,
     rollback_segment => '');
END;
/
```

```
/*--- Create link from propagator to receiver for scheduled push ---*/
CONNECT prop_rec/prop_rec@hq.example.com
CREATE DATABASE LINK sales.example.com CONNECT TO prop_rec
 IDENTIFIED BY &password;

/*--- Create public link from Sales to HQ with necessary USING clause ---*/
CONNECT system@sales.example.com
CREATE PUBLIC DATABASE LINK hq.example.com USING 'hq.example.com';

/*--- Create private repadmin to repadmin link ---*/
CONNECT repadmin@sales.example.com
CREATE DATABASE LINK hq.example.com CONNECT TO repadmin IDENTIFIED BY &password;

/*--- Schedule replication from Sales to HQ ---*/
BEGIN
  DBMS_DEFER_SYS.SCHEDULE_PUSH(
     destination => 'hq.example.com',
     interval => 'sysdate + 1/24',
     next_date => sysdate,
     stop_on_error => FALSE,
     parallelism => 1);
END;
/

/*--- Schedule purge of def tran queue at Sales ---*/
BEGIN
  DBMS_DEFER_SYS.SCHEDULE_PURGE(
     next_date => sysdate,
     interval => 'sysdate + 1',
     delay_seconds => 0,
     rollback_segment =>'');
END;
/

/*--- Create link from propagator to receiver for scheduled push ---*/
CONNECT prop_rec/prop_rec@sales.example.com
CREATE DATABASE LINK hq.example.com connect TO prop_rec IDENTIFIED BY &password;
```

## Security Setup for Materialized View Replication

Nearly all users should find it easiest to use the configuration wizards in the Advanced Replication interface in Oracle Enterprise Manager when configuring materialized view replication security. However, for certain specialized cases, you might need to use the replication management API to perform these setup operations. To configure a replication environment, the database administrator must connect with DBA privileges to grant the necessary privileges to the replication administrator.

First set up user accounts at each materialized view site with the appropriate privileges to configure and maintain the replication environment and to propagate replicated changes. You must also define links for these users to the associated master site or master materialized view site. You might need to create additional users, or assign additional privileges to users at the associated master site or master materialized view site.

In addition to end users who will be accessing replicated objects, there are three special categories of "users" at a materialized view site:

- Replication administrators, who are responsible for configuring and maintaining a replication environment.

- Propagators, who are responsible for propagating deferred transactions.
- Refreshers, who are responsible for pulling down changes to the materialized views from the associated master tables or master materialized views.

Typically, a single user performs each of these functions. However, there might be situations where you need different users performing these functions. For example, materialized views can be created by a materialized view site administrator and refreshed by another end user.

Table A–2 describes the privileges needed to create and maintain a materialized view site.

***Table A–2 Required Materialized View Site User Accounts***

| User | Privileges |
|---|---|
| Materialized view site replication administrator | `DBMS_REPCAT_ADMIN.GRANT_ADMIN_ANY_SCHEMA` |
| Propagator | `DBMS_DEFER_SYS.REGISTER_PROPAGATOR` |
| Refresher | `CREATE ANY MATERIALIZED VIEW`<br>`ALTER ANY MATERIALIZED VIEW` |

In addition to creating the appropriate users at the materialized view site, you might need to create additional users at the associated master site or master materialized view site, as well. Table A–3 on  on page A-8 describes the privileges need by master site or master materialized view site users to support a new materialized view site.

## Trusted Compared with Untrusted Security

In addition to the different users at the master site or master materialized view site, you also must determine which type of security model you will implement: trusted or untrusted. With a trusted security model, the receiver and proxy materialized view administrator have access to all local replication groups. The receiver and proxy materialized view administrator perform database activities at the local master site or master materialized view site on behalf of the propagator and materialized view administrator, respectively, at the remote materialized view site. Therefore, the propagator and materialized view administrator at the remote materialized view site also have access to all replication groups at the master site or master materialized view site. Remember that a single receiver is used for all incoming transactions.

For example, consider the scenario in Figure A–3. Even though Materialized View Groups A and C exist at the materialized view site (based on Master Groups A and C at the Master Site), the propagator and materialized view administrator have access to Master Groups A, B, C, and D at the Master Site because the trusted security model has been used. While this greatly increases the flexibility of database administration, because the DBA can perform administrative functions at any of these remote sites and have these changes propagated to the master sites, it also increases the chances of a malicious user at a remote site viewing or corrupting data at the master site.

Regardless of the security model used, Oracle automatically grants the appropriate privileges for objects as they are added to or removed from a replication environment.

*Figure A–3 Trusted Security: Materialized View Replication*

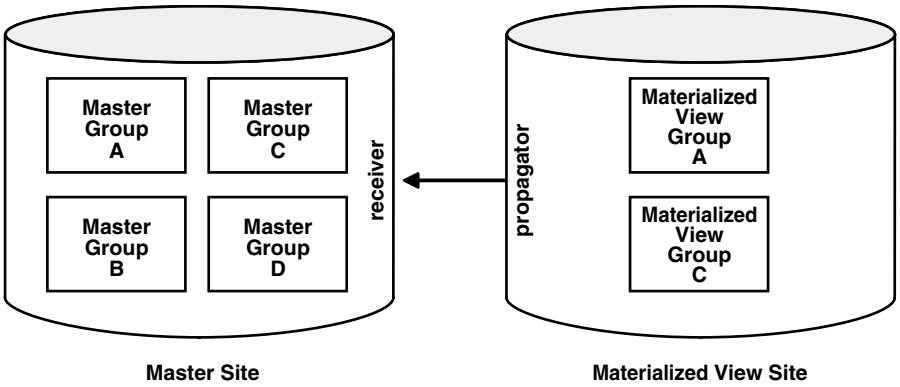


Untrusted security assigns only the privileges to the proxy materialized view administrator and receiver that are required to work with specified replication groups. The propagator and materialized view administrator, therefore, can only access these specified replication groups at the Master Site. Figure A–4 illustrates an untrusted security model with materialized view replication. Because the Materialized View Site contains Materialized View Groups A and C, access to only Master Groups A and C are required. Using untrusted security does not allow the propagator or the materialized view administrator at the Materialized View Site to access Master Groups B and D at the Master Site.

*Figure A–4 Untrusted Security: Materialized View Replication*

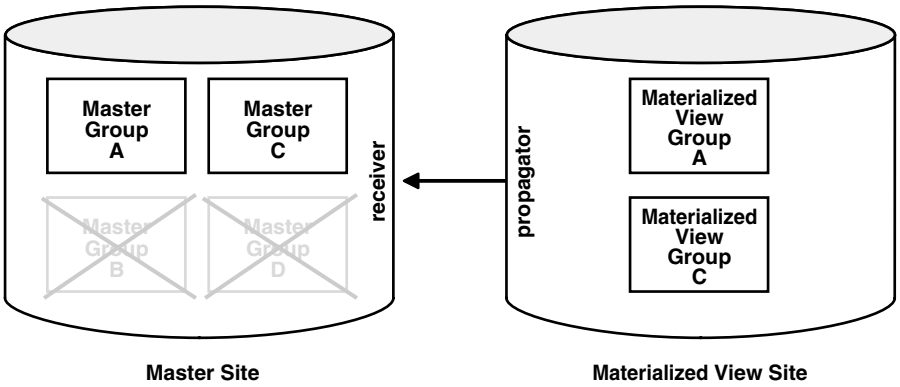


Typically, materialized view sites are more vulnerable to security breaches and therefore the untrusted security model is used. There are very few reasons why you would want to use a trusted security model with your materialized view site and it is recommended that you use the untrusted security model with materialized view sites.

One reason you might choose to use a trusted security model is if your materialized view site is considered a master site in every way (security, constant network connectivity, resources) but is a materialized view only because of data subsetting requirements. Remember that row and column subsetting are not supported in a multimaster configuration.

Use the appropriate API calls listed for the proxy materialized view administrator and receiver in Table A–3 to assign the different users the appropriate privileges.

***Table A–3 Required Master Site or Master Materialized View Site User Accounts***

| User | Privileges |
|---|---|
| proxy materialized view site administrator | See "REGISTER_USER_REPGROUP Procedure" on page 20-5 for details.<br>**Trusted**:<br>`DBMS_REPCAT_ADMIN.REGISTER_USER_REPGROUP`<br>`privilege => 'proxy_snapadmin'`<br>`list_of_gnames => NULL,`<br>`...`<br>**Untrusted**:<br>`DBMS_REPCAT_ADMIN.REGISTER_USER_REPGROUP`<br>`privilege => 'proxy_snapadmin'`<br>`list_of_gnames => '`*`mastergroupname`*`',`<br>`...` |
| receiver | See "REGISTER_USER_REPGROUP Procedure" on page 20-5 for details.<br>**Trusted**:<br>`DBMS_REPCAT_ADMIN.REGISTER_USER_REPGROUP`<br>`privilege => 'receiver'`<br>`list_of_gnames => NULL,`<br>`...`<br>**Untrusted**:<br>`DBMS_REPCAT_ADMIN.REGISTER_USER_REPGROUP`<br>`privilege => 'receiver'`<br>`list_of_gnames => '`*`mastergroupname`*`',`<br>`...` |
| proxy refresher | **Trusted**:<br>Grant `CREATE SESSION`<br>Grant `SELECT ANY TABLE`<br>**Untrusted**:<br>Grant `CREATE SESSION`<br>Grant `SELECT` on necessary master tables or master materialized views and materialized view logs |

After creating the accounts at both the materialized view and associated master sites or master materialized view sites, you must create the following private database links, including user name and password, from the materialized view site to the master site or master materialized view site:

- From the materialized view replication administrator to the proxy materialized view replication administrator.
- From the propagator to the receiver.
- From the refresher to the proxy refresher.
- From the materialized view owner to the master site or master materialized view site for refreshes.

Assuming you have designated a single user account to act as materialized view administrator, propagator, and refresher, you must create one link for each materialized view site for those functions. You do not need a link from the master site or master materialized view site to the materialized view site.

After creating these links, you must call DBMS_DEFER_SYS.SCHEDULE_PUSH and DBMS_DEFER_SYS.SCHEDULE_PURGE at the materialized view site to define how frequently you want to propagate your deferred transaction queue to the associated master site or master materialized view site, and how frequently you want to purge this queue. You must also call DBMS_REFRESH.REFRESH at the materialized view site to schedule how frequently to pull changes from the associated master site or master materialized view site.

# B

# User-Defined Conflict Resolution Methods

This appendix describes how to build user-defined conflict resolution methods and user-defined conflict notification methods.

This appendix contains these topics:


- User-Defined Conflict Resolution Methods
- User-Defined Conflict Notification Methods
- Viewing Conflict Resolution Information


## User-Defined Conflict Resolution Methods

Oracle enables you to write your own conflict resolution or notification methods. A user-defined conflict resolution method is a PL/SQL function that returns either `TRUE` or `FALSE`. `TRUE` indicates that the method has successfully resolved all conflicting modifications for a column group. If the method cannot successfully resolve a conflict, then it should return `FALSE`. Oracle continues to evaluate available conflict resolution methods, in sequence order, until either a method returns `TRUE` or there are no more methods available.

If the conflict resolution method raises an exception, then Oracle stops evaluation of the method, and, if any other methods were provided to resolve the conflict with a later sequence number, then Oracle does not evaluate them.

### Conflict Resolution Method Parameters

The parameters needed by a user-defined conflict resolution method are determined by the type of conflict being resolved (uniqueness, update, or delete) and the columns of the table being replicated. All conflict resolution methods take some combination of old, new, and current column values for the table.

- The old value represents the value of the row at the initiating site before you made the change.
- The new value represents the value of the row at the initiating site after you made the change.
- The current value represents the value of the equivalent row at the receiving site.

---

**Note:** Recall that Oracle uses the primary key, or the key specified by `SET_COLUMNS`, to determine which rows to compare.

---

The conflict resolution function should accept as parameters the values for the columns specified in the PARAMETER_COLUMN_NAME argument to the DBMS_REPCAT.ADD_*conflicttype*_RESOLUTION procedures. The column parameters are passed to the conflict resolution method in the order listed in the PARAMETER_COLUMN_NAME argument, or in ascending alphabetical order if you specified '*' for this argument. When both old and new column values are passed as parameters (for update conflicts), the old value of the column immediately precedes the new value.

**Note:**

- Type checking of parameter columns in user-defined conflict resolution methods is not performed until you regenerate replication support for the associated replicated table.
- Attributes of column objects cannot be defined as column parameters for user-defined conflict resolution methods.

## Resolving Update Conflicts

For update conflicts, a user-defined function should accept the following values for each column in the column group:

- Old column value from the initiating site. The mode for this parameter is IN. This value should not be changed.
- New column value from the initiating site. The mode for this parameter is IN OUT. If the function can resolve the conflict successfully, then it should modify the new column value as needed.
- Current column value from the receiving site. The mode for this parameter is IN.

The old, new, and current values for a column are received consecutively. The final argument to the conflict resolution method should be a Boolean flag. If this flag is FALSE, then it indicates that you have updated the value of the IN OUT parameter (new) and that you should update the current column value with this new value. If this flag is TRUE, then it indicates that the current column value should not be changed.

## Resolving Uniqueness Conflicts

Uniqueness conflicts can occur as the result of an INSERT or UPDATE. Your uniqueness conflict resolution method should accept the new column value from the initiating site in IN OUT mode for each column in the column group. The final parameter to the conflict resolution method should be a Boolean flag.

If the method can resolve the conflict, then it should modify the new column values so that Oracle can insert or update the current row with the new column values. Your function should set the Boolean flag to TRUE if it wants to discard the new column values, and FALSE otherwise.

Because a conflict resolution method cannot guarantee convergence for uniqueness conflicts, a user-defined uniqueness resolution method should include a notification mechanism.

## Resolving Delete Conflicts

Delete conflicts occur when you successfully delete from the local site, but the associated row cannot be found at the remote site (for example, because it had been updated). For delete conflicts, the function should accept old column values in IN OUT

mode for the entire row. The final parameter to the conflict resolution method should be a Boolean flag.

If the conflict resolution method can resolve the conflict, then it modifies the old column values so that Oracle can delete the current row that matches all old column values. Your function should set the Boolean flag to `TRUE` if it wants to discard these column values, and `FALSE` otherwise.

If you perform a delete at the local site and an update at the remote site, then the remote site detects the delete conflict, but the local site detects an unresolvable update conflict. This type of conflict cannot be handled automatically. The conflict raises a `NO_DATA_FOUND` exception and Oracle logs the transaction as an error transaction.

Designing a mechanism to properly handle these types of update/delete conflicts is difficult. It is far easier to avoid these types of conflicts entirely, by simply "marking" deleted rows, and then purging them using procedural replication.

> **See Also:** "Creating Conflict Avoidance Methods for Delete Conflicts" on page 6-23

## Multitier Materialized Views and User-Defined Conflict Resolution Methods

When you use user-defined conflict resolution methods with multitier materialized views, the information about these methods is pulled down to the master materialized view sites automatically. This information is stored in the data dictionary at the master materialized view site. However, the user-defined conflict resolution methods themselves cannot be pulled down from the master site. Therefore, you must re-create these methods at the master materialized view site.

> **See Also:**
>
> - "Viewing Conflict Resolution Information" on page B-8 for information about the data dictionary views that store information about user-defined conflict resolution methods
> - *Oracle Database Advanced Replication* for more information about conflict resolution and multitier materialized views

## Restrictions for User-Defined Conflict Resolution Methods

The following sections describe restrictions for user-defined conflict resolution methods.

### SQL Statement Restrictions for User-Defined Conflict Resolution Methods

Avoid the following types of SQL statements in user-defined conflict resolution methods. Use of such statements can result in unpredictable results.

- Data definition language (DDL) statements (such as `CREATE`, `ALTER`, `DROP`)
- Transaction control statements (such as `COMMIT`, `ROLLBACK`)
- Session control (such as `ALTER SESSION`)
- System control (such as `ALTER SYSTEM`)

### Column Subsetting Restrictions for User-Defined Conflict Resolution Methods

Avoid subsetting the columns in a column group when you create updatable multitier materialized views. Column subsetting excludes columns that are in master tables or master materialized views from a materialized view based on these masters. You do

this by specifying certain select columns in the SELECT statement during materialized view creation.

When you use conflict resolution with multitier materialized views, you cannot define the conflict resolution methods at the materialized view site. Conflict resolution methods are always pulled down from the master site. Therefore, if you subset the columns in a column group that has a user-defined conflict resolution applied to it, the conflict resolution method will not be able to find all of the columns in the column group at a master materialized view site. When this happens, the conflict resolution method returns the following error:

```
ORA-23460 missing value for column in resolution method
```

For example, consider a case where the job_id, salary, and commission_pct columns in the hr.employees table are part of a column group name employees_cg1 that has a user-defined conflict resolution method applied to it at the master site hq.example.com. To protect the privacy of your sales staff, you create a level 1 updatable materialized view that uses column subsetting to exclude the salary and commission_pct columns at the ca.us office. When you create this materialized view at the ca.us office, the conflict resolution method is pulled down from hq.example.com. You then create an updatable multitier materialized view at the sf.ca office based on the level 1 materialized view at the ca.us office.

Given this replication environment, if a conflict arises for a job_id value at the level 1 materialized view at the ca.us office, then the conflict resolution method fails to find the salary and commission_pct columns and returns the ORA-23460 error mentioned previously.

> **See Also:** *Oracle Database Advanced Replication* for more information about column subsetting

## Examples of User-Defined Conflict Resolution Method

The following examples show user-defined methods that are variations on the standard maximum and additive prebuilt conflict resolution methods. Unlike the standard methods, these custom functions can handle nulls in the columns used to resolve the conflict.

### Maximum User Function

```
-- User function similar to MAXIMUM method.
-- If curr is null or curr < new, then use new values.
-- If new is null or new < curr, then use current values.
-- If both are null, then no resolution.
-- Does not converge with > 2 masters, unless
-- always increasing.

CREATE OR REPLACE FUNCTION max_null_loses(old  IN       NUMBER,
                       new                     IN  OUT  NUMBER,
                       cur                     IN       NUMBER,
                       ignore_discard_flag     OUT      BOOLEAN)
  RETURN BOOLEAN IS
BEGIN
   IF (new IS NULL AND cur IS NULL) OR new = cur THEN
       RETURN FALSE;
   END IF;
   IF new IS NULL THEN
       ignore_discard_flag := TRUE;
   ELSIF cur IS NULL THEN
```

```
       ignore_discard_flag := FALSE;
   ELSIF new < cur THEN
       ignore_discard_flag := TRUE;
   ELSE
       ignore_discard_flag := FALSE;
   END IF;
   RETURN TRUE;
END max_null_loses;
/
```

### Additive User Function

```
-- User function similar to ADDITIVE method.
-- If old is null, then old = 0.
-- If new is null, then new = 0.
-- If curr is null, then curr = 0.
-- new = curr + (new - old) -> just like ADDITIVE method.

CREATE OR REPLACE FUNCTION additive_nulls(old  IN       NUMBER,
                        new                    IN  OUT  NUMBER,
                        cur                    IN       NUMBER,
                        ignore_discard_flag    OUT      BOOLEAN)
   RETURN BOOLEAN IS
   old_val NUMBER := 0.0;
   new_val NUMBER := 0.0;
   cur_val NUMBER := 0.0;
BEGIN
   IF old IS NOT NULL THEN
      old_val := old;
   END IF;
   IF new IS NOT NULL THEN
      new_val := new;
   END IF;
   IF cur IS NOT NULL THEN
      cur_val := cur;
   END IF;
   new := cur_val + (new_val - old_val);
   ignore_discard_flag := FALSE;
   RETURN TRUE;
END additive_nulls;
/
```

## User-Defined Conflict Notification Methods

A conflict notification method is a user-defined function that provides conflict notification rather than or in addition to conflict resolution. For example, you can write your own conflict notification methods to log conflict information in a database table, send an email message, or page an administrator. After you write a conflict notification method, you can assign it to a column group (or constraint) in a specific order so that Oracle notifies you when a conflict happens, before attempting subsequent conflict resolution methods, or after Oracle attempts to resolve a conflict but cannot do so.

To configure a replicated table with a user-defined conflict notification mechanism, you must complete the following steps:

1. Create a conflict notification log.
2. Create the user-defined conflict notification method in a package.

The following sections explain each step.

## Creating a Conflict Notification Log

When configuring a replicated table to use a user-defined conflict notification method, the first step is to create a database table that can record conflict notifications. You can create a table to log conflict notifications for one or many tables in a master group.

To create a conflict notification log table at all master sites, use the replication execute DDL facility. For more information, see "EXECUTE_DDL Procedure" on page 18-68. Do *not* generate replication support for the conflict notification tables because their entries are specific to the site that detects a conflict.

### Sample Conflict Notification Log Table

The following `CREATE TABLE` statement creates a table that you can use to log conflict notifications from several tables in a master group.

```
CREATE TABLE sales.conf_report (
 line          NUMBER(2),    --- used to order message text
 txt           VARCHAR2(80), --- conflict notification message
 timestamp     DATE,         --- time of conflict
 table_name    VARCHAR2(30), --- table in which the
                             --- conflict occurred
 table_owner   VARCHAR2(30), --- owner of the table
 conflict_type VARCHAR2(6)   --- INSERT, DELETE or UNIQUE
);
```

## Creating a Conflict Notification Package

To create a conflict notification method, you must define the method in a PL/SQL package and then replicate the package as part of a master group along with the associated replicated table.

A conflict notification method can perform conflict notification only, or both conflict notification and resolution. If possible, always use one of Oracle's prebuilt conflict resolution methods to resolve conflicts. When a user-defined conflict notification method performs only conflict notification, assign the user-defined method to a column group (or constraint) along with conflict resolution methods that can resolve conflicts.

---

**Note:** If Oracle cannot ultimately resolve a replication conflict, then Oracle rolls back the entire transaction, including any updates to a notification table. If notification is necessary independent of transactions, then you can design a notification mechanism to use the Oracle `DBMS_PIPES` package.

---

### Sample Conflict Notification Package

The following package and package body perform a simple form of *conflict notification* by logging uniqueness conflicts for a `CUSTOMERS` table into the previously defined `CONF_REPORT` table.

**Note:** This example of *conflict notification* does not resolve any conflicts. You should either provide a method to resolve conflicts (such as *discard* or *overwrite*), or provide a notification mechanism that will succeed (for example, using e-mail) even if the error is not resolved and the transaction is rolled back. With simple modifications, the following user-defined conflict notification method can take more active steps. For example, instead of just recording the notification message, the package can use the `DBMS_OFFICE` utility package to send an Oracle Office email message to an administrator.

```
CREATE OR REPLACE PACKAGE notify AS
 -- Report uniqueness constraint violations on CUSTOMERS table
 FUNCTION customers_unique_violation (
   first_name         IN OUT VARCHAR2,
   last_name          IN OUT VARCHAR2,
   discard_new_values IN OUT BOOLEAN)
 RETURN BOOLEAN;
END notify;
/

CREATE OR REPLACE PACKAGE BODY notify AS
 -- Define a PL/SQL associative array to hold the notification message
 TYPE message_table IS TABLE OF VARCHAR2(80) INDEX BY BINARY_INTEGER;
 PROCEDURE report_conflict (
   conflict_report IN MESSAGE_TABLE,
   report_length   IN NUMBER,
   conflict_time   IN DATE,
   conflict_table  IN VARCHAR2,
   table_owner     IN VARCHAR2,
   conflict_type   IN VARCHAR2) IS
 BEGIN
   FOR idx IN 1..report_length LOOP
     BEGIN
       INSERT INTO sales.conf_report
         (line, txt, timestamp, table_name, table_owner, conflict_type)
         VALUES (idx, SUBSTR(conflict_report(idx),1,80), conflict_time,
                 conflict_table, table_owner, conflict_type);
     EXCEPTION WHEN others THEN NULL;
     END;
   END LOOP;
 END report_conflict;
 -- This is the conflict resolution method that is called first when
 -- a uniqueness constraint violated is detected in the CUSTOMERS table.
 FUNCTION customers_unique_violation (
   first_name  IN OUT VARCHAR2,
   last_name  IN OUT VARCHAR2,
   discard_new_values  IN OUT BOOLEAN)
  RETURN BOOLEAN IS
   local_node  VARCHAR2(128);
   conf_report MESSAGE_TABLE;
   conf_time   DATE := SYSDATE;
  BEGIN
  -- Get the global name of the local site
    BEGIN
      SELECT global_name INTO local_node FROM global_name;
    EXCEPTION WHEN others THEN local_node := '?';
    END;
```

```
  -- Generate a message for the DBA
  conf_report(1) := 'UNIQUENESS CONFLICT DETECTED IN TABLE CUSTOMERS ON ' ||
                    TO_CHAR(conf_time, 'MM-DD-YYYY HH24:MI:SS');
  conf_report(2) := '  AT NODE ' || local_node;
  conf_report(3) := 'ATTEMPTING TO RESOLVE CONFLICT USING ' ||
                    'APPEND SEQUENCE METHOD';
  conf_report(4) := 'FIRST NAME: ' || first_name;
  conf_report(5) := 'LAST NAME:  ' || last_name;
  conf_report(6) := NULL;
  --- Report the conflict
  report_conflict(conf_report, 5, conf_time, 'CUSTOMERS',
                'OFF_SHORE_ACCOUNTS', 'UNIQUE');
  --- Do not discard the new column values. They are still needed by
  --- other conflict resolution methods.
  discard_new_values := FALSE;
  --- Indicate that the conflict was not resolved.
    RETURN FALSE;
  END customers_unique_violation;
END notify;
/
```

## Viewing Conflict Resolution Information

Oracle provides replication catalog (REPCAT) views that you can use to determine what conflict resolution methods are being used by each of the tables and column groups in your replication environment. Each view has three versions: `USER_*`, `ALL_*`, `SYS.DBA_*`. The available views are shown in the following table.

| View | Description |
|---|---|
| ALL_REPRESOLUTION_METHOD | Lists all of the available conflict resolution methods. |
| ALL_REPCOLUMN_GROUP | Lists all of the column groups defined for the database. |
| ALL_REPGROUPED_COLUMN | Lists all of the columns in each column group in the database. |
| ALL_REPPRIORITY_GROUP | Lists all of the priority groups and site priority groups defined for the database. |
| ALL_REPPRIORITY | Lists the values and corresponding priority levels for each priority or site priority group. |
| ALL_REPCONFLICT | Lists the types of conflicts (delete, update, or uniqueness) for which you have specified a resolution method, for the tables, column groups, and unique constraints in the database. |
| ALL_REPRESOLUTION | Shows more specific information about the conflict resolution method used to resolve conflicts on each object. |
| ALL_REPPARAMETER_COLUMN | Shows which columns are used by the conflict resolution methods to resolve a conflict. |

**See Also:** Chapter 23, "Replication Catalog Views"

# Index

## A


ADD procedure, 5-12, 8-25, 17-3
ADD_DEFAULT_DEST procedure, 14-4
ADD_DELETE_RESOLUTION procedure, 18-16
ADD_GROUPED_COLUMN procedure, 18-6
ADD_MASTER_DATABASE procedure, 3-8, 7-23, 7-25, 18-7
ADD_NEW_MASTERS procedure, 7-8, 7-19, 18-8
ADD_PRIORITY_CHAR procedure, 18-13
ADD_PRIORITY_*datatype* procedure, 18-13
ADD_PRIORITY_DATE procedure, 18-13
ADD_PRIORITY_NUMBER procedure, 18-13
ADD_PRIORITY_VARCHAR2 procedure, 18-13
ADD_SITE_PRIORITY_SITE procedure, 6-17, 18-15
ADD_UNIQUENESS_RESOLUTION procedure, 18-16
ADD_UPDATE_RESOLUTION procedure, 6-3, 6-5, 6-9, 6-11, 6-14, 6-18, 18-16
administrative requests
  ALL_REPCATLOG view, 23-15
  executing, 7-27, 18-54
  monitoring, 10-12
    errors, 10-13
    jobs, 10-14
  purging, 18-76
administrators
  for materialized view sites
    creating, 2-17
Advanced Replication
  migrating to Streams, 18-95
Advanced Replication interface
  monitoring replication, 10-1
ALL_REPCAT_REFRESH_TEMPLATES view, 23-6
ALL_REPCAT_TEMPLATE_OBJECTS view, 23-7
ALL_REPCAT_TEMPLATE_PARMS view, 23-9
ALL_REPCAT_TEMPLATE_SITES view, 23-11
ALL_REPCAT_USER_AUTHORIZATIONS view, 23-12
ALL_REPCAT_USER_PARM_VALUES view, 23-13
ALL_REPCATLOG view
  administrative requests, 23-15
ALL_REPCOLUMN view, 23-16
ALL_REPCOLUMN_GROUP view, 23-18
ALL_REPCONFLICT view, 23-19
ALL_REPDDL view, 23-20
ALL_REPGENOBJECTS view, 23-21
ALL_REPGROUP view, 23-22
ALL_REPGROUP_PRIVILEGES view, 23-23
ALL_REPGROUPED_COLUMN view, 23-24
ALL_REPKEY_COLUMNS view, 23-25
ALL_REPOBJECT view, 23-26
ALL_REPPARAMETER_COLUMN view, 23-28
ALL_REPPRIORITY view, 23-29
ALL_REPPRIORITY_GROUP view, 23-30
ALL_REPPROP view, 23-31
ALL_REPRESOL_STATS_CONTROL view, 23-32
ALL_REPRESOLUTION view, 23-33
ALL_REPRESOLUTION_METHOD view, 23-34
ALL_REPRESOLUTION_STATISTICS view, 23-35
  gathering statistics, 6-30
ALL_REPSITES view, 23-36
ALTER MATERIALIZED VIEW LOG statement, 8-12
ALTER_CATCHUP_PARAMETERS procedure, 18-20
ALTER_MASTER_PROPAGATION procedure, 18-22
ALTER_MASTER_REPOBJECT procedure, 6-7, 6-16, 6-24, 9-1, 18-23
ALTER_MVIEW_PROPAGATION procedure, 18-25
ALTER_PRIORITY procedure, 18-26
ALTER_PRIORITY_CHAR procedure, 18-27
ALTER_PRIORITY_*datatype* procedure, 18-27
ALTER_PRIORITY_DATE procedure, 18-27
ALTER_PRIORITY_NUMBER procedure, 18-27
ALTER_PRIORITY_RAW procedure, 18-27
ALTER_REFRESH_TEMPLATE procedure, 21-4
ALTER_SITE_PRIORITY procedure, 18-28
ALTER_SITE_PRIORITY_SITE procedure, 18-29
ALTER_TEMPLATE_OBJECT procedure, 21-6
ALTER_TEMPLATE_PARM procedure, 21-8
ALTER_USER_AUTHORIZATION procedure, 21-10
ALTER_USER_PARM_VALUE procedure, 21-11
ANY_CHAR_ARG procedure, 12-5
ANY_CLOB_ARG procedure, 12-5
ANY_VARCHAR2_ARG procedure, 12-5
ANYDATA
  GET_ANYDATA_ARG function, 13-8
ANYDATA data type
  replication, 9-13

ANYDATA_ARG procedure, 12-5
authorization
  template users, 4-9
availability
  extended, 7-3, 18-8, 18-24, 18-74, 18-88, 18-93, 18-98

## B

BEGIN_INSTANTIATION procedure, 7-27, 15-3
BEGIN_LOAD procedure, 7-29, 15-5
BEGIN_TABLE_REORGANIZATION, 8-15
BLOB_ARG procedure, 12-5

## C

CALL procedure, 12-3
CANCEL_STATISTICS procedure, 6-30, 18-30
CHANGE procedure, 17-4
CHAR_ARG procedure, 12-5
CLEAR_PROP_STATISTICS procedure, 10-24, 14-5
CLOB_ARG procedure, 12-5
column objects
  user-defined conflict resolution, B-2
column subsetting
  user-defined conflict resolution methods, B-3
columns
  adding to master tables, 18-85
  column groups, 6-3, 6-5, 6-8, 6-10, 6-12, 6-17
    adding members to, 18-6
    creating, 18-51, 18-73
    dropping, 18-55
    removing members from, 18-56
COMMENT_ON_COLUMN_GROUP procedure, 18-31
COMMENT_ON_DELETE_RESOLUTION procedure, 18-38
COMMENT_ON_MVIEW_REPSITES procedure, 18-32
COMMENT_ON_PRIORITY procedure, 18-33
COMMENT_ON_REPGROUP procedure, 18-34
COMMENT_ON_REPOBJECT procedure, 18-35
COMMENT_ON_REPSITES procedure, 18-36
COMMENT_ON_SITE_PRIORITY procedure, 18-37
COMMENT_ON_UNIQUE_RESOLUTION procedure, 18-38
COMMENT_ON_UPDATE_RESOLUTION procedure, 18-38
comments
  comments field
    updating in views, 7-33
  updating, 7-33
COMMIT_WORK procedure, 12-4
COMPARE_OLD_VALUES procedure, 18-40
COMPARE_TEMPLATES function, 21-13
comparing
  tables, 16-3
conflict resolution, 6-1
  additive method, 6-9, 18-16
  auditing, 6-30
  average method, 6-9
  column groups, 6-3, 6-5, 6-8, 6-10, 6-12, 6-17
  DBA_REPRESOLUTION_STATISTICS view, 6-31
  discard method, 6-2
  information
    viewing, B-8
  maximum method, 6-4
  minimum method, 6-4
  overwrite method, 6-2
  preparing for, 6-1
  priority groups method, 6-11
  procedural replication and, 7-37
  site priority method, 6-15
    sample trigger, 6-16
  statistics, 18-30, 18-81
    canceling, 6-30
    collecting, 6-30
    viewing, 6-30
  time stamp method
    sample trigger, 6-8
  timestamp method, 6-6
  uniqueness, 6-19
  user-defined methods, B-1
    column objects, B-2
    column subsetting, B-3
    example, B-4
    for delete conflicts, B-2
    for uniqueness conflicts, B-2
    for update conflicts, B-2
    multitier materialized views, B-3
    parameters, B-1
    restrictions, B-3
  viewing information, B-8
conflicts
  avoiding
    delete, 6-23
    dynamic ownership, 6-26
  notification log table
    creating, B-6
    sample, B-6
  notification methods
    user-defined, B-5
  notification package
    creating, B-6
    sample, B-6
  token passing, 6-27
  workflow, 6-27
COPY_TEMPLATE function, 21-14
CREATE_MASTER_REPGROUP procedure, 3-4, 18-42
CREATE_MASTER_REPOBJECT procedure, 6-8, 6-16, 18-43
CREATE_MVIEW_REPGROUP procedure, 5-5, 8-22, 8-27, 8-29, 18-46
CREATE_MVIEW_REPOBJECT procedure, 5-6, 5-7, 5-11, 5-12, 8-27, 8-30, 18-48
CREATE_OBJECT_FROM_EXISTING function, 21-16
CREATE_REFRESH_TEMPLATE function, 21-18
CREATE_REFRESH_TEMPLATE procedure, 4-4

CREATE_TEMPLATE_OBJECT function, 21-20
CREATE_TEMPLATE_OBJECT procedure, 4-5
CREATE_TEMPLATE_PARM function, 21-22
CREATE_USER_AUTHORIZATION function, 21-24
CREATE_USER_AUTHORIZATION procedure, 4-9
CREATE_USER_PARM_VALUE function, 21-25

## D

data definition language
  altering replicated objects, 18-23
  asynchronous, 18-68
data dictionary views
  comments
    updating, 7-33
  materialized views, 25-1
  refresh groups, 25-1
  replication, 10-1, 23-1
database links
  creating, 2-13, 2-22, 4-15, 5-4, 5-9, 8-28
*datatype*_ARG procedure, 12-5
date expressions, 2-5
DATE_ARG procedure, 12-5
DBA_REGISTERED_MVIEW_GROUPS view, 23-5
DBA_REPCAT_REFRESH_TEMPLATES view, 23-37
DBA_REPCAT_TEMPLATE_OBJECTS view, 23-38
DBA_REPCAT_TEMPLATE_PARMS view, 23-39
DBA_REPCAT_TEMPLATE_SITES view, 23-40
DBA_REPCAT_USER_AUTHORIZATIONS
    view, 23-41
DBA_REPCAT_USER_PARM_VALUES view, 23-42
DBA_REPCATLOG view, 23-43
  purging requests from, 18-76
DBA_REPCOLUMN view, 23-44
DBA_REPCOLUMN_GROUP view, 23-45
  updating, 18-31
DBA_REPCONFLICT view, 23-46
DBA_REPDDL view, 23-47
DBA_REPEXTENSIONS view, 23-48
DBA_REPGENOBJECTS view, 23-50
DBA_REPGROUP view, 23-51
  updating, 18-34
DBA_REPGROUP_PRIVILEGES view, 23-52
DBA_REPGROUPED_COLUMN view, 23-53
DBA_REPKEY_COLUMNS view, 23-54
DBA_REPOBJECT view, 23-55
  updating, 18-35
DBA_REPPARAMETER_COLUMN view, 23-56
DBA_REPPRIORITY view, 23-57
DBA_REPPRIORITY_GROUP view, 23-58
DBA_REPPRIORITYGROUP view
  updating, 18-33, 18-37
DBA_REPPROP view, 23-59
DBA_REPRESOL_STATS_CONTROL view, 23-60
DBA_REPRESOLUTION view, 23-61
  updating, 18-38
DBA_REPRESOLUTION_METHOD view, 23-62
DBA_REPRESOLUTION_STATISTICS view, 23-63
  purging, 6-31, 18-77
DBA_REPSITES view, 23-64
  updating, 18-36
DBA_REPSITES_NEW view, 23-65
DBMS_DEFER package, 12-1
DBMS_DEFER_QUERY package, 13-1
  GET_ANYDATA_ARG function, 9-13
DBMS_DEFER_SYS package, 14-1
  CLEAR_PROP_STATISTICS procedure, 10-24
  EXECUTE_ERROR procedure, 7-27, 9-15
  EXECUTE_ERROR_AS_USER procedure, 9-15
  PURGE function, 9-12
  PUSH function, 9-11
  REGISTER_PROPAGATOR procedure, 2-5, 2-18,
    2-22
  SCHEDULE_PURGE procedure, 2-6, 2-19, 2-23
  SCHEDULE_PUSH procedure, 2-14, 2-20, 2-24
DBMS_MVIEW package
  BEGIN_TABLE_REORGANIZATION
    procedure, 8-15
  END_TABLE_REORGANIZATION
    procedure, 8-15
  PURGE_LOG procedure, 8-13
  PURGE_MVIEW_FROM_LOG procedure, 8-8,
    8-10, 8-11, 8-13
  REFRESH procedure, 8-1, 8-26
  UNREGISTER_MVIEW procedure, 8-10
DBMS_OFFLINE_OG package, 15-1
  BEGIN_INSTANTIATION procedure, 7-27
  BEGIN_LOAD procedure, 7-29
  END_INSTANTIATION procedure, 7-30
  END_LOAD procedure, 7-30
  RESUME_SUBSET_OF_MASTERS
    procedure, 7-29
DBMS_OFFLINE_SNAPSHOT package
  END_LOAD procedure, 8-24, 8-25
DBMS_RECTIFIER_DIFF package, 9-7, 16-1
  DIFFERENCES procedure, 9-8
  RECTIFY procedure, 9-8
DBMS_REFRESH package, 17-1
  ADD procedure, 5-12, 8-25
  MAKE procedure, 5-5, 5-10, 8-23
  REFRESH procedure, 8-1
DBMS_REPCAT package, 7-33, 18-1
  ADD_MASTER_DATABASE procedure, 3-8,
    7-23, 7-25
  ADD_NEW_MASTERS procedure, 7-8, 7-19
  ADD_SITE_PRIORITY_SITE procedure, 6-17
  ADD_UPDATE_RESOLUTION procedure, 6-3,
    6-5, 6-9, 6-11, 6-14, 6-18
  ALTER_MASTER_REPOBJECT procedure, 6-7,
    6-16, 6-24, 9-1
  CANCEL_STATISTICS procedure, 6-30
  CREATE_MASTER_REPGROUP procedure, 3-4
  CREATE_MASTER_REPOBJECT procedure, 6-8,
    6-16
  CREATE_MVIEW_REPGROUP procedure, 5-5,
    8-22, 8-27, 8-29
  CREATE_MVIEW_REPOBJECT procedure, 5-6,
    5-7, 5-11, 5-12, 8-27, 8-30
  DEFINE_SITE_PRIORITY procedure, 6-17
  DO_DEFERRED_REPCAT_ADMIN

procedure, 6-26, 7-27
DROP_MVIEW_REPGROUP procedure, 8-3, 8-4, 8-6
DROP_MVIEW_REPOBJECT procedure, 8-7
GENERATE_REPLICATION_SUPPORT procedure, 3-9, 3-10, 9-3
MAKE_COLUMN_GROUP procedure, 6-3, 6-5, 6-8, 6-10, 6-12, 6-17
PREPARE_INSTANTIATED_MASTER procedure, 7-13, 7-22
PURGE_STATISTICS procedure, 6-31
REGISTER_STATISTICS procedure, 6-30
RELOCATE_MASTERDEF procedure, 7-2
REMOVE_MASTER_DATABASE procedure, 7-31
RESUME_MASTER_ACTIVITY procedure, 3-10
RESUME_PROPAGATION_TO_MDEF procedure, 7-11, 7-21
SPECIFY_NEW_MASTERS procedure, 7-8, 7-19
SWITCH_MVIEW_MASTER procedure, 8-2
UNREGISTER_MVIEW_REPGROUP procedure, 8-8
DBMS_REPCAT_ADMIN package, 20-1
GRANT_ADMIN_ANY_SCHEMA procedure, 2-4, 2-17, 2-22
REGISTER_USER_REPGROUP procedure, 2-5, 2-6, 2-9, 2-12, 2-18, 2-20
DBMS_REPCAT_INSTANTIATE package, 19-1
DROP_SITE_INSTANTIATION procedure, 8-3, 8-4
DBMS_REPCAT_RGT package, 21-1
CREATE_REFRESH_TEMPLATE procedure, 4-4
CREATE_TEMPLATE_OBJECT procedure, 4-5
CREATE_USER_AUTHORIZATION procedure, 4-9
INSTANTIATE_OFFLINE procedure, 4-11
INSTANTIATE_ONLINE procedure, 4-12
DBMS_REPUTIL package, 22-1
REPLICATION_OFF procedure, 7-37, 9-5
REPLICATION_ON procedure, 7-37, 9-5
DDL. *See* data definition language
DEFCALL view, 26-2
DEFCALLDEST view, 26-3
DEFDEFAULTDEST view, 26-4
adding destinations to, 14-4
removing destinations from, 14-6, 14-7
DEFERRCOUNT view, 26-5
deferred transaction queues
deferred calls
determining value of, 9-13
managing, 9-11
purging propagated transactions, 9-12
pushing, 9-11
deferred transactions
data dictionary views, 26-1
DEFDEFAULTDEST view
adding destination to, 14-4
removing destinations from, 14-6
deferred remote procedure calls (RPCs)
argument types, 13-4
argument values, 13-8
arguments to, 12-5
building, 12-3
executing immediately, 14-16
DEFSCHEDULE view
clearing statistics, 14-5
removing destinations from, 14-7
deleting from queue, 14-9
monitoring, 10-15
purge job, 10-17, 10-18
push jobs, 10-15, 10-16
reexecuting, 14-12
scheduling execution, 14-22
starting, 12-7
DEFERROR view, 9-15, 26-6
deleting transactions from, 14-8
DEFINE_COLUMN_GROUP procedure, 18-51
DEFINE_PRIORITY_GROUP procedure, 18-52
DEFINE_SITE_PRIORITY procedure, 6-17, 18-53
DEFLOB view, 26-7
DEFPROPAGATOR view, 26-8
DEFSCHEDULE view, 26-9
clearing statistics, 10-24, 14-5
DEFTRAN view, 26-12
DEFTRANDEST view, 26-13
DELETE_DEF_DESTINATION procedure, 14-7
DELETE_DEFAULT_DEST procedure, 14-6
DELETE_ERROR procedure, 14-8
DELETE_RUNTIME_PARMS procedure, 21-27
DELETE_TRAN procedure, 14-9
deployment templates
alter object, 21-6
alter parameters, 21-8
alter template, 21-4
alter user authorization, 21-10
alter user parameter values, 21-11
authorize users, 4-9
compare templates, 21-13
concepts, 4-1
copy template, 21-14
create object from existing, 21-16
create template, 21-18
creating, 4-2, 4-4
data dictionary views for, 23-6
distributing files, 4-14
drop site instantiation, 19-3
dropping, 21-34
dropping all, 21-31
dropping materialized view group, 8-3
flowchart for creating, 4-3
instantiating, 4-14
instantiation script, 4-12
lock template, 21-47, 21-48
monitoring, 10-7
objects
creating, 21-20
dropping, 21-36
dropping all, 21-28
offline instantiation, 4-9, 19-4, 21-43
online instantiation, 19-6, 21-45

packaging, 4-9, 4-10
for offline instantiation, 4-11
for online instantiation, 4-11
parameters
creating, 4-7, 21-22
dropping, 21-37
dropping all, 21-29
user values, 4-8
run-time parameters
creating, 21-41
deleting, 21-27
get ID, 21-40
inserting, 21-41
sites
dropping, 21-35
dropping all, 21-30
user authorizations
creating, 21-24
dropping, 21-38
dropping all, 21-32
user parameter values
creating, 21-25
dropping, 21-39
dropping all, 21-33
user-defined types, 4-2
DESTROY procedure, 17-6
differences
between tables, 16-3
rectifying, 16-6
DIFFERENCES procedure, 9-8, 16-3
DISABLED function, 14-10
disabling
propagation, 14-24
DO_DEFERRED_REPCAT_ADMIN procedure, 6-26, 7-27, 18-54
DROP MATERIALIZED VIEW LOG statement, 8-17
DROP_ALL_OBJECTS procedure, 21-28
DROP_ALL_TEMPLATE_PARMS procedure, 21-29
DROP_ALL_TEMPLATE_SITES procedure, 21-30
DROP_ALL_TEMPLATES procedure, 21-31
DROP_ALL_USER_AUTHORIZATIONS procedure, 21-32
DROP_ALL_USER_PARM_VALUES procedure, 21-33
DROP_COLUMN_GROUP procedure, 18-55
DROP_DELETE_RESOLUTION procedure, 18-66
DROP_GROUPED_COLUMN procedure, 18-56
DROP_MASTER_REPGROUP procedure, 18-57
DROP_MASTER_REPOBJECT procedure, 18-58
DROP_MVIEW_REPGROUP procedure, 8-3, 8-4, 18-59
DROP_MVIEW_REPOBJECT procedure, 8-7, 18-60
DROP_PRIORITY procedure, 18-61
DROP_PRIORITY_CHAR procedure, 18-63
DROP_PRIORITY_*datatype* procedure, 18-63
DROP_PRIORITY_DATE procedure, 18-63
DROP_PRIORITY_GROUP procedure, 18-62
DROP_PRIORITY_NUMBER procedure, 18-63
DROP_PRIORITY_VARCHAR2 procedure, 18-63
DROP_REFRESH_TEMPLATE procedure, 21-34
DROP_SITE_INSTANTIATION procedure, 8-3, 8-4, 19-3, 21-35
DROP_SITE_PRIORITY procedure, 18-64
DROP_SITE_PRIORITY_SITE procedure, 18-65
DROP_TEMPLATE_OBJECT procedure, 21-36
DROP_TEMPLATE_PARM procedure, 21-37
DROP_UNIQUE_RESOLUTION procedure, 18-66
DROP_UPDATE_RESOLUTION procedure, 18-66
DROP_USER_AUTHORIZATION procedure, 21-38
DROP_USER_PARM_VALUE procedure, 21-39
dynamic ownership
conflict avoidance and, 6-26
locating owner of a row, 6-28
obtaining ownership, 6-29
workflow partitioning, 6-27
dynamic performance views
replication, 24-1


## E


END_INSTANTIATION procedure, 7-30, 15-6
END_LOAD procedure, 7-30, 8-24, 8-25, 15-7
END_TABLE_REORGANIZATION procedure, 8-15
Enterprise Manager
Advanced Replication interface, 10-1
errors
error queues
DEFERROR view, 9-15
managing, 9-15
error transactions
monitoring, 10-18
reexecuting as alternate user, 9-15
reexecuting as receiver, 9-15
EXCLUDE_PUSH function, 14-11
EXECUTE_DDL procedure, 18-68
EXECUTE_ERROR procedure, 7-27, 9-15, 14-12
EXECUTE_ERROR_AS_USER procedure, 9-15, 14-13
extended availability, 7-3, 18-8, 18-24, 18-74, 18-88, 18-93, 18-98


## F


foreign key constraints
adding master sites, 7-2
FROM_REMOTE function, 22-6


## G


GENERATE_MVIEW_SUPPORT procedure, 18-69
GENERATE_REPLICATION_SUPPORT procedure, 3-9, 3-10, 9-3, 18-71
generating
replication support, 3-9
procedural replication, 7-37
GET_ANYDATA_ARG function, 9-13, 13-8
GET_ARG_FORM function, 13-3
GET_ARG_TYPE function, 13-4
GET_BLOB_ARG function, 13-8
GET_CALL_ARGS procedure, 13-6
GET_CHAR_ARG function, 13-8

GET_CLOB_ARG function, 13-8
GET_*datatype*_ARG function, 13-8
GET_DATE_ARG function, 13-8
GET_IDS_ARG function, 13-8
GET_IYM_ARG function, 13-8
GET_NCHAR_ARG function, 13-8
GET_NCLOB_ARG function, 13-8
GET_NUMBER_ARG function, 13-8
GET_NVARCHAR2_ARG function, 13-8
GET_OBJECT_NULL_VECTOR_ARG function, 13-10
GET_RAW_ARG function, 13-8
GET_ROWID_ARG function, 13-8
GET_RUNTIME_PARM_ID function, 21-40
GET_TIMESTAMP_ARG function, 13-8
GET_TSLTZ_ARG function, 13-8
GET_TSTZ_ARG function, 13-8
GET_VARCHAR2_ARG function, 13-8
GLOBAL_NAME function, 22-7
GRANT_ADMIN_ANY_SCHEMA procedure, 2-4, 2-17, 2-22, 20-3
GRANT_ADMIN_SCHEMA procedure, 20-4

## I

IDS_ARG procedure, 12-5
Import
  replication groups
    offline instantiation and, 15-5, 15-7
  status check, 18-86
INSERT_RUNTIME_PARMS procedure, 21-41
INSTANTIATE_OFFLINE function, 19-4, 21-43
INSTANTIATE_OFFLINE procedure, 4-11
INSTANTIATE_ONLINE function, 19-6, 21-45
INSTANTIATE_ONLINE procedure, 4-12
instantiation, 4-14
  DROP_SITE_INSTANTIATION procedure, 19-3, 21-35
  offline, 4-9
    INSTANTIATE_OFFLINE function, 19-4, 21-43
  online
    INSTANTIATE_ONLINE function, 19-6, 21-45
  refreshing after, 4-16
  script, 4-12
IYM_ARG procedure, 12-5

## J

jobs
  queues for
    removing jobs from, 14-27, 14-28

## L

LOCK_TEMPLATE_EXCLUSIVE procedure, 21-47
LOCK_TEMPLATE_SHARED procedure, 21-48
LONG columns
  replication, 9-6

## M

MAKE procedure, 5-5, 5-10, 8-23, 17-7
MAKE_COLUMN_GROUP procedure, 6-3, 6-5, 6-8, 6-10, 6-17, 18-73
MAKE_INTERNAL_PKG procedure, 22-8
master definition site
  relocating, 18-82
master groups
  adding master sites to
    with quiesce, 7-23
    without quiesce, 7-3
  adding objects to, 3-4
  creating, 3-1, 3-4, 18-42
  dropping, 18-57
  flowchart for creating, 3-3
  monitoring, 10-2
  quiescing, 18-96
  removing master sites from, 7-31
  resuming replication activity, 18-87
master materialized views
  monitoring, 10-4
  reorganizing, 8-14
master sites
  adding, 3-7, 7-2
    circular dependencies, 3-8, 7-2
    flowchart for, 7-15
    flowchart for determining method, 7-4
    foreign key constraints, 7-2
    restrictions, 7-5
    restrictions for change-based recovery, 7-3
    restrictions for full database export/import, 7-3
    self-referential constraints, 3-8, 7-2
    using change-based recovery, 7-5
    using full database export/import, 7-5
    using object-level export/import, 7-14
    using offline instantiation, 7-25
    with quiesce, 7-23
    without quiesce, 7-3
  changing master definition site, 7-1
  cleaning up, 8-8
  creating, 18-7
  creating users for, 2-6, 2-9, 2-12, 2-20
  database links, 2-13
  determining differences, 9-7
  dropping, 18-84
  flowchart for setting up, 2-3
  monitoring, 10-1, 10-4
  propagating changes between, 14-22
  removing, 7-31
  scheduled links for, 2-13
  scheduled purges for, 2-5
  setup, 2-3
master tables
  adding columns to, 18-85
  redefining online, 8-14
  reorganizing, 8-14
    methods, 8-15
  truncating, 8-15
materialized view groups

adding objects to, 5-6, 5-11, 8-30
changing masters, 8-2
creating, 5-2, 5-4, 5-10, 18-46
dropping, 8-3, 8-6
group owner, 8-27
monitoring, 10-8
refreshing, 18-78
materialized view logs
adding columns, 8-12
altering, 8-12
privileges required, 8-12
dropping, 8-17
managing, 8-11
space, 8-12
monitoring, 10-5
purging
materialized views from, 8-8, 8-10
privileges required, 8-13
purging rows from
manually, 8-13
reducing space allocated to, 8-13
reorganizing masters with, 8-14
truncating, 8-13
truncating master table with, 8-15
materialized view sites
adding
using offline instantiation, 8-17
administrators
creating, 2-17
changing masters, 18-97
database links
creating, 2-22, 4-15, 5-4, 5-9, 8-28
dropping, 8-2, 18-59
dropping objects from, 8-6
flowchart for setting up, 2-16
group owner
using, 8-27
monitoring, 10-8
multitier
setting up, 2-16
propagating changes to master, 14-22
refresher
creating, 2-17, 2-21
schedule purge, 2-19, 2-23
users
creating, 2-17
materialized views
data dictionary views, 25-1
deployment templates
user-defined types, 4-2
dropping, 8-7
generating support for, 18-69
monitoring, 10-8, 10-9
multitier
setting up, 5-2
user-defined conflict resolution, B-3
purging from materialized view logs, 8-8, 8-10
refresh groups
creating, 5-5, 5-10
refreshing, 4-16, 8-1, 8-26
security, A-5
trusted compared with untrusted, A-6
unregistering from master, 8-10
monitoring replication, 10-1
Advanced Replication interface, 10-1
multimaster replication
monitoring, 10-1
security
trusted compared with untrusted, A-2
multitier materialized views
setting up, 2-16

## N

NCHAR_ARG procedure, 12-5
NCLOB_ARG procedure, 12-5
notification log table
conflicts
creating, B-6
sample, B-6
notification methods
user-defined, B-5
notification package
conflicts
creating, B-6
NUMBER_ARG procedure, 12-5
NVARCHAR2_ARG procedure, 12-5

## O

objects
adding to materialized view sites, 18-48
altering, 18-23
creating
for master group, 18-42, 18-43
for materialized view sites, 18-48
dropping
from materialized view site, 8-6, 18-60
generating replication support for, 18-71
offline instantiation
adding a master site, 7-25
adding a materialized view site, 8-17
INSTANTIATE_OFFLINE function, 19-4, 21-43
replication groups, 15-3, 15-5, 15-6, 15-7, 15-9
online instantiation
INSTANTIATE_ONLINE function, 19-6, 21-45
online redefinition of tables, 8-14
Oracle Streams
migrating to, 18-95

## P

packaging
deployment templates, 4-9
parallel propagation
monitoring, 10-23, 10-24
parameters
deployment templates, 4-7
user values, 4-8
performance
replication

monitoring, 10-21
planning
for replication, 1-2
PREPARE_INSTANTIATED_MASTER procedure, 7-13, 7-22
PREPARE_INSTANTIATED_MASTERS procedure, 18-74
PRESERVE MATERIALIZED VIEW LOG option
TRUNCATE TABLE statement, 8-15
priority groups
adding members to, 18-13
altering members
priorities, 18-26
values, 18-27
creating, 18-52
dropping, 18-62
removing members from, 18-61, 18-63
site priority groups
adding members to, 18-15
procedural replication
conflicts and, 7-37
generating replication support for, 7-37
restrictions, 7-34
serialization of transactions, 7-37
user-defined types, 7-36
using, 7-34
propagation
altering method, 18-22, 18-25
disabling, 14-24
of changes, 18-22
parallel
monitoring, 10-23, 10-24
status of, 14-10
propagator
registering, 2-5, 14-19
proxy materialized view administrator
creating, 2-6, 2-9, 2-12, 2-20
PURGE function, 9-12, 14-14
PURGE_LOG procedure, 8-13
PURGE_MASTER_LOG procedure, 18-76
PURGE_MVIEW_FROM_LOG procedure, 8-8, 8-10, 8-11, 8-13
PURGE_STATISTICS procedure, 6-31, 18-77
purges
DBA_REPCATLOG table, 18-76
deferred transaction queue, 9-12
master sites, 2-5
materialized view sites, 2-19, 2-23
monitoring, 10-17
PUSH function, 9-11, 14-16
pushes
deferred transaction queue, 9-11

## Q

quiescing
adding master sites with, 7-23
adding master sites without, 7-3
master groups, 18-96

## R

RAW_ARG procedure, 12-5
receiver
registering, 2-5
RECTIFY procedure, 9-8, 16-6
rectifying
tables, 9-8, 16-6
redefining tables
online
replication, 8-14
refresh
materialized view sites, 18-78
materialized views, 8-1, 8-26
monitoring, 10-11, 10-12
refresh groups
adding members to, 17-3
adding objects to, 5-7, 5-12, 8-25
creating, 5-5, 5-10, 17-7
data dictionary views, 25-1
deleting, 17-6
monitoring, 10-10
refresh, 8-1
refresh interval
changing, 17-4
refreshing
manually, 17-9
removing members from, 17-10
REFRESH procedure, 8-1, 8-26, 17-9
REFRESH_ALL_MVIEWS procedure, 8-2
REFRESH_DEPENDENT procedure, 8-2
REFRESH_MVIEW_REPGROUP procedure, 18-78
refresher
creating, 2-17, 2-21
REGISTER_MVIEW_REPGROUP procedure, 18-80
REGISTER_PROPAGATOR procedure, 2-5, 2-18, 2-22, 14-19
REGISTER_STATISTICS procedure, 6-30, 18-81
REGISTER_USER_REPGROUP procedure, 2-5, 2-6, 2-9, 2-12, 2-18, 2-20, 20-5
RELOCATE_MASTERDEF procedure, 7-2, 18-82
REMOVE_MASTER_DATABASE procedure, 7-31
REMOVE_MASTER_DATABASES procedure, 18-84
RENAME_SHADOW_COLUMN_GROUP procedure, 18-85
REPCAT_IMPORT_CHECK procedure, 18-86
replication
catalog views, 10-1, 23-1
column groups, 6-3, 6-5, 6-8, 6-10, 6-12, 6-17
conflict resolution, 6-1
uniqueness, 6-19
creating an environment, 1-1
data dictionary views, 10-1, 23-1
database links
creating, 2-13
datetime data types
abbreviations, 11-2
deferred transaction queues
managing, 9-11
deferred transactions
data dictionary views, 26-1

deployment templates
user-defined types, 4-2
determining differences between tables, 9-7
disabling, 7-37, 9-4, 9-5, 22-3
dynamic performance views, 24-1
enabling, 7-37, 9-4, 9-5, 22-4
error queues
managing, 9-15
flowchart for creating environment, 1-1
generating support for, 3-9
interval data types
abbreviations, 11-2
LONG column
converting to LOB, 9-6
managing an environment, 6-1
master groups
creating, 3-1
master sites
adding, 3-7
materialized view groups
creating, 5-2, 5-4, 5-10
materialized view logs
managing, 8-11
monitoring, 10-1
deferred transactions, 10-15
error transactions, 10-18
master environments, 10-1
materialized view environments, 10-8
performance, 10-21
objects
adding to master group, 3-4
dropping from master sites, 18-58
parallel propagation
monitoring, 10-23, 10-24
planning for, 1-2
procedural replication, 7-34
restrictions, 7-34
user-defined types, 7-36
propagator
registering, 2-5
receiver
registering, 2-5
replicated objects, 9-1
replication queues, 9-1
resuming, 3-10
scheduled links
creating, 2-13
security, A-1
setting up sites, 2-1
sites
setup, 2-1
statistics
clearing, 10-24
triggers, 9-6
replication catalog views, 23-1
comments
updating, 7-33
monitoring replication, 10-1
replication management API, 11-1
conflict resolution, 6-1
deployment templates
creating, 4-2
instantiating, 4-14
packaging, 4-9
examples, 11-1
managing a replication environment, 6-1
managing replicated objects, 9-1
managing replication queues, 9-1
master groups
creating, 3-1
materialized view groups
creating, 5-2
overview, 1-1
packages, 10-1
setting up replication sites, 2-1
replication objects
altering, 9-1
tables
altering, 9-4
REPLICATION_IS_ON function, 22-5
REPLICATION_OFF procedure, 22-3
REPLICATION_ON procedure, 7-37, 22-4
RESUME_MASTER_ACTIVITY procedure, 3-10, 18-87
RESUME_PROPAGATION_TO_MDEF procedure, 7-11, 7-21, 18-88
RESUME_SUBSET_OF_MASTERS procedure, 7-29, 15-9
resuming replication activity, 18-87
REVOKE_ADMIN_ANY_SCHEMA procedure, 20-7
REVOKE_ADMIN_SCHEMA procedure, 20-8
ROWID_ARG procedure, 12-5


## S


SCHEDULE_PURGE procedure, 2-6, 2-19, 2-23, 14-20
SCHEDULE_PUSH procedure, 2-14, 2-20, 2-24, 14-22
scheduled links
creating, 2-13
security
for materialized view replication, A-5
trusted compared with untrusted, A-6
for multimaster replication, A-1
trusted compared with untrusted, A-2
replication, A-1
trusted compared with untrusted, A-2, A-6
SEND_OLD_VALUES procedure, 18-89
serialization
of transactions, 7-37
SET_COLUMNS procedure, 18-41, 18-91
SET_DISABLED procedure, 14-24
site priority
altering, 18-28
site priority groups
adding members to, 18-15
creating
syntax, 18-53
dropping, 18-64

removing members from, 18-65
snapshots. *See* materialized views
SPECIFY_NEW_MASTERS procedure, 7-8, 18-93
statistics
for conflict resolution
auditing, 6-30
cancelling, 6-30
clearing, 6-31, 18-77
collecting, 6-30, 18-81
viewing, 6-30
for propagation
clearing, 10-24, 14-5
status
propagation, 14-10
storage parameters
materialized view log
altering, 8-12
STREAMS_MIGRATION procedure, 18-95
SUBTRACT procedure, 17-10
SUSPEND_MASTER_ACTIVITY procedure, 18-96
SWITCH_MVIEW_MASTER procedure, 8-2, 18-97
SYNC_UP_REP procedure, 22-9

## T

tables
altering
without replicating changes, 9-4
altering replicated, 9-1
comparing, 16-3
differences between, 9-7
rectifying, 9-8, 16-6
redefining online
replication, 8-14
updating comments, 7-33
templates. *See* deployment templates
TIMESTAMP_ARG procedure, 12-5
token passing, 6-27
sample implementation, 6-26
TRANSACTION procedure, 12-7
transactions
serialization of, 7-37
triggers
for site priority conflict resolution, 6-16
for time stamp conflict resolution, 6-8
replicating, 9-6
TRUNCATE statement, 8-14
TRUNCATE TABLE statement
PRESERVE MATERIALIZED VIEW LOG
option, 8-15
trusted security, A-2, A-6
TSLTZ_ARG procedure, 12-5

## U

UNDO_ADD_NEW_MASTERS_REQUEST
procedure, 18-98
UNREGISTER_MVIEW procedure, 8-10
UNREGISTER_MVIEW_REPGROUP
procedure, 18-100
UNREGISTER_PROPAGATOR procedure, 14-26
UNREGISTER_USER_REPGROUP procedure, 20-9
UNSCHEDULE_PURGE procedure, 14-27
UNSCHEDULE_PUSH procedure, 14-28
USER_REPCAT_REFRESH_TEMPLATES
view, 23-66
USER_REPCAT_TEMP_OUTPUT view, 4-10
USER_REPCAT_TEMPLATE_OBJECTS view, 23-67
USER_REPCAT_TEMPLATE_PARMS view, 23-68
USER_REPCAT_TEMPLATE_SITES view, 23-69
USER_REPCAT_USER_AUTHORIZATIONS
view, 23-70
USER_REPCAT_USER_PARM_VALUES
view, 23-71
USER_REPCATLOG view, 23-72
USER_REPCOLUMN view, 23-73
USER_REPCOLUMN_GROUP view, 23-74
USER_REPCONFLICT view, 23-75
USER_REPDDL view, 23-76
USER_REPGENOBJECTS view, 23-77
USER_REPGROUP view, 23-78
USER_REPGROUP_PRIVILEGES view, 23-79
USER_REPGROUPED_COLUMN view, 23-80
USER_REPKEY_COLUMNS view, 23-81
USER_REPOBJECT view, 23-82
USER_REPPARAMETER_COLUMN view, 23-83
USER_REPPRIORITY view, 23-84
USER_REPPRIORITY_GROUP view, 23-85
USER_REPPROP view, 23-86
USER_REPRESOL_STATS_CONTROL view, 23-87
USER_REPRESOLUTION view, 23-88
USER_REPRESOLUTION_METHOD view, 23-89
USER_REPRESOLUTION_STATISTICS view, 23-90
USER_REPSITES view, 23-91
users
authorize for deployment template, 4-9
master materialized view sites, 2-20
master sites, 2-6, 2-9, 2-12
materialized view sites, 2-17

## V

V$MVREFRESH view, 24-2
V$REPLPROP view, 10-24, 24-3
V$REPLQUEUE view, 24-5
VALIDATE procedure, 18-101
VARCHAR2_ARG procedure, 12-5

## W

WAIT_MASTER_LOG procedure, 18-103
workflow, 6-27